

XV6

0.1

Generated by Doxygen 1.9.6

1 ENVIRONMENT PREPARETION	1
1.1 LINUX	1
1.2 BUILDING AND RUNNING XV6	1
2 Class Index	3
2.1 Class List	3
3 File Index	5
3.1 File List	5
4 Class Documentation	9
4.1 backcmd	9
4.1.1 Detailed Description	9
4.1.2 Member Data Documentation	9
4.1.2.1 cmd	9
4.1.2.2 type	9
4.2 buf	10
4.2.1 Detailed Description	10
4.2.2 Member Data Documentation	10
4.2.2.1 blockno	10
4.2.2.2 data	10
4.2.2.3 dev	11
4.2.2.4 flags	11
4.2.2.5 lock	11
4.2.2.6 next	11
4.2.2.7 prev	11
4.2.2.8 qnext	12
4.2.2.9 refcnt	12
4.3 cmd	12
4.3.1 Detailed Description	12
4.3.2 Member Data Documentation	12
4.3.2.1 type	12
4.4 context	13
4.4.1 Detailed Description	13
4.4.2 Member Data Documentation	13
4.4.2.1 ebp	13
4.4.2.2 ebx	13
4.4.2.3 edi	13
4.4.2.4 eip	14
4.4.2.5 esi	14
4.5 cpu	14
4.5.1 Detailed Description	14
4.5.2 Member Data Documentation	14

4.5.2.1 apicid	15
4.5.2.2 gdt	15
4.5.2.3 intena	15
4.5.2.4 ncli	15
4.5.2.5 proc	15
4.5.2.6 scheduler	16
4.5.2.7 started	16
4.5.2.8 ts	16
4.6 devsw	16
4.6.1 Detailed Description	16
4.6.2 Member Data Documentation	17
4.6.2.1 read	17
4.6.2.2 write	17
4.7 dinode	17
4.7.1 Detailed Description	17
4.7.2 Member Data Documentation	17
4.7.2.1 addrs	18
4.7.2.2 major	18
4.7.2.3 minor	18
4.7.2.4 nlink	18
4.7.2.5 size	18
4.7.2.6 type	19
4.8 dirent	19
4.8.1 Detailed Description	19
4.8.2 Member Data Documentation	19
4.8.2.1 inum	19
4.8.2.2 name	19
4.9 elfhdr	20
4.9.1 Detailed Description	20
4.9.2 Member Data Documentation	20
4.9.2.1 ehsize	20
4.9.2.2 elf	20
4.9.2.3 entry	21
4.9.2.4 flags	21
4.9.2.5 machine	21
4.9.2.6 magic	21
4.9.2.7 phentsize	21
4.9.2.8 phnum	21
4.9.2.9 phoff	22
4.9.2.10 shentsize	22
4.9.2.11 shnum	22
4.9.2.12 shoff	22

4.9.2.13 shstrndx	22
4.9.2.14 type	22
4.9.2.15 version	23
4.10 execcmd	23
4.10.1 Detailed Description	23
4.10.2 Member Data Documentation	23
4.10.2.1 argv	23
4.10.2.2 eargv	23
4.10.2.3 type	24
4.11 file	24
4.11.1 Detailed Description	24
4.11.2 Member Enumeration Documentation	24
4.11.2.1 anonymous enum	24
4.11.3 Member Data Documentation	25
4.11.3.1 ip	25
4.11.3.2 off	25
4.11.3.3 pipe	25
4.11.3.4 readable	25
4.11.3.5 ref	26
4.11.3.6	26
4.11.3.7 writable	26
4.12 gatedesc	26
4.12.1 Detailed Description	27
4.12.2 Member Data Documentation	27
4.12.2.1 args	27
4.12.2.2 cs	27
4.12.2.3 dpl	27
4.12.2.4 off_15_0	27
4.12.2.5 off_31_16	27
4.12.2.6 p	28
4.12.2.7 rsv1	28
4.12.2.8 s	28
4.12.2.9 type	28
4.13 header	28
4.13.1 Detailed Description	29
4.13.2 Member Data Documentation	29
4.13.2.1 ptr	29
4.13.2.2	29
4.13.2.3 size	29
4.13.2.4 x	29
4.14 inode	30
4.14.1 Detailed Description	30

4.14.2 Member Data Documentation	30
4.14.2.1 addrs	30
4.14.2.2 dev	30
4.14.2.3 inum	31
4.14.2.4 lock	31
4.14.2.5 major	31
4.14.2.6 minor	31
4.14.2.7 nlink	31
4.14.2.8 ref	32
4.14.2.9 size	32
4.14.2.10 type	32
4.14.2.11 valid	32
4.15 ioapic	32
4.15.1 Detailed Description	33
4.15.2 Member Data Documentation	33
4.15.2.1 data	33
4.15.2.2 pad	33
4.15.2.3 reg	33
4.16 kmap	33
4.16.1 Detailed Description	34
4.16.2 Member Data Documentation	34
4.16.2.1 perm	34
4.16.2.2 phys_end	34
4.16.2.3 phys_start	34
4.16.2.4 virt	34
4.17 listcmd	35
4.17.1 Detailed Description	35
4.17.2 Member Data Documentation	35
4.17.2.1 left	35
4.17.2.2 right	35
4.17.2.3 type	35
4.18 log	36
4.18.1 Detailed Description	36
4.18.2 Member Data Documentation	36
4.18.2.1 committing	36
4.18.2.2 dev	36
4.18.2.3 lh	36
4.18.2.4 lock	37
4.18.2.5 outstanding	37
4.18.2.6 size	37
4.18.2.7 start	37
4.19 logheader	37

4.19.1 Detailed Description	38
4.19.2 Member Data Documentation	38
4.19.2.1 block	38
4.19.2.2 n	38
4.20 mp	38
4.20.1 Detailed Description	39
4.20.2 Member Data Documentation	39
4.20.2.1 checksum	39
4.20.2.2 imcrp	39
4.20.2.3 length	39
4.20.2.4 physaddr	39
4.20.2.5 reserved	40
4.20.2.6 signature	40
4.20.2.7 specrev	40
4.20.2.8 type	40
4.21 mpconf	40
4.21.1 Detailed Description	41
4.21.2 Member Data Documentation	41
4.21.2.1 checksum	41
4.21.2.2 entry	41
4.21.2.3 lapicaddr	41
4.21.2.4 length	41
4.21.2.5 oemlength	42
4.21.2.6 oemtable	42
4.21.2.7 product	42
4.21.2.8 reserved	42
4.21.2.9 signature	42
4.21.2.10 version	42
4.21.2.11 xchecksum	43
4.21.2.12 xlength	43
4.22 mpioapic	43
4.22.1 Detailed Description	43
4.22.2 Member Data Documentation	43
4.22.2.1 addr	43
4.22.2.2 apicno	44
4.22.2.3 flags	44
4.22.2.4 type	44
4.22.2.5 version	44
4.23 mpproc	44
4.23.1 Detailed Description	45
4.23.2 Member Data Documentation	45
4.23.2.1 apicid	45

4.23.2.2 feature	45
4.23.2.3 flags	45
4.23.2.4 reserved	45
4.23.2.5 signature	45
4.23.2.6 type	46
4.23.2.7 version	46
4.24 pipe	46
4.24.1 Detailed Description	46
4.24.2 Member Data Documentation	46
4.24.2.1 data	46
4.24.2.2 lock	47
4.24.2.3 nread	47
4.24.2.4 nwrite	47
4.24.2.5 readopen	47
4.24.2.6 writeopen	47
4.25 pipecmd	48
4.25.1 Detailed Description	48
4.25.2 Member Data Documentation	48
4.25.2.1 left	48
4.25.2.2 right	48
4.25.2.3 type	48
4.26 proc	49
4.26.1 Detailed Description	49
4.26.2 Member Data Documentation	49
4.26.2.1 chan	49
4.26.2.2 context	49
4.26.2.3 cwd	50
4.26.2.4 killed	50
4.26.2.5 kstack	50
4.26.2.6 name	50
4.26.2.7 ofile	50
4.26.2.8 parent	51
4.26.2.9 pgdir	51
4.26.2.10 pid	51
4.26.2.11 priority	51
4.26.2.12 state	51
4.26.2.13 sz	52
4.26.2.14 tf	52
4.27 proghdr	52
4.27.1 Detailed Description	52
4.27.2 Member Data Documentation	52
4.27.2.1 align	53

4.27.2.2 filesz	53
4.27.2.3 flags	53
4.27.2.4 memsz	53
4.27.2.5 off	53
4.27.2.6 paddr	54
4.27.2.7 type	54
4.27.2.8 vaddr	54
4.28 redircmd	54
4.28.1 Detailed Description	54
4.28.2 Member Data Documentation	54
4.28.2.1 cmd	55
4.28.2.2 efile	55
4.28.2.3 fd	55
4.28.2.4 file	55
4.28.2.5 mode	55
4.28.2.6 type	56
4.29 rtcdate	56
4.29.1 Detailed Description	56
4.29.2 Member Data Documentation	56
4.29.2.1 day	56
4.29.2.2 hour	57
4.29.2.3 minute	57
4.29.2.4 month	57
4.29.2.5 second	57
4.29.2.6 year	57
4.30 run	58
4.30.1 Detailed Description	58
4.30.2 Member Data Documentation	58
4.30.2.1 next	58
4.31 segdesc	58
4.31.1 Detailed Description	59
4.31.2 Member Data Documentation	59
4.31.2.1 avl	59
4.31.2.2 base_15_0	59
4.31.2.3 base_23_16	59
4.31.2.4 base_31_24	59
4.31.2.5 db	59
4.31.2.6 dpl	60
4.31.2.7 g	60
4.31.2.8 lim_15_0	60
4.31.2.9 lim_19_16	60
4.31.2.10 p	60

4.31.2.11 rsv1	60
4.31.2.12 s	61
4.31.2.13 type	61
4.32 sleeplock	61
4.32.1 Detailed Description	61
4.32.2 Member Data Documentation	61
4.32.2.1 lk	61
4.32.2.2 locked	62
4.32.2.3 name	62
4.32.2.4 pid	62
4.33 spinlock	62
4.33.1 Detailed Description	62
4.33.2 Member Data Documentation	63
4.33.2.1 cpu	63
4.33.2.2 locked	63
4.33.2.3 name	63
4.33.2.4 pcs	63
4.34 stat	63
4.34.1 Detailed Description	64
4.34.2 Member Data Documentation	64
4.34.2.1 dev	64
4.34.2.2 ino	64
4.34.2.3 nlink	64
4.34.2.4 size	65
4.34.2.5 type	65
4.35 superblock	65
4.35.1 Detailed Description	65
4.35.2 Member Data Documentation	65
4.35.2.1 bmapstart	66
4.35.2.2 inodestart	66
4.35.2.3 logstart	66
4.35.2.4 nblocks	66
4.35.2.5 ninodes	66
4.35.2.6 nlog	67
4.35.2.7 size	67
4.36 taskstate	67
4.36.1 Detailed Description	68
4.36.2 Member Data Documentation	68
4.36.2.1 cr3	68
4.36.2.2 cs	68
4.36.2.3 ds	68
4.36.2.4 eax	69

4.36.2.5 ebp	69
4.36.2.6 ebx	69
4.36.2.7 ecx	69
4.36.2.8 edi	69
4.36.2.9 edx	69
4.36.2.10 eflags	70
4.36.2.11 eip	70
4.36.2.12 es	70
4.36.2.13 esi	70
4.36.2.14 esp	70
4.36.2.15 esp0	70
4.36.2.16 esp1	71
4.36.2.17 esp2	71
4.36.2.18 fs	71
4.36.2.19 gs	71
4.36.2.20 iomb	71
4.36.2.21 ldt	71
4.36.2.22 link	72
4.36.2.23 padding1	72
4.36.2.24 padding10	72
4.36.2.25 padding2	72
4.36.2.26 padding3	72
4.36.2.27 padding4	72
4.36.2.28 padding5	73
4.36.2.29 padding6	73
4.36.2.30 padding7	73
4.36.2.31 padding8	73
4.36.2.32 padding9	73
4.36.2.33 ss	73
4.36.2.34 ss0	74
4.36.2.35 ss1	74
4.36.2.36 ss2	74
4.36.2.37 t	74
4.37 trapframe	74
4.37.1 Detailed Description	75
4.37.2 Member Data Documentation	75
4.37.2.1 cs	75
4.37.2.2 ds	76
4.37.2.3 eax	76
4.37.2.4 ebp	76
4.37.2.5 ebx	76
4.37.2.6 ecx	76

4.37.2.7 edi	76
4.37.2.8 edx	77
4.37.2.9 eflags	77
4.37.2.10 eip	77
4.37.2.11 err	77
4.37.2.12 es	77
4.37.2.13 esi	78
4.37.2.14 esp	78
4.37.2.15 fs	78
4.37.2.16 gs	78
4.37.2.17 oesp	78
4.37.2.18 padding1	78
4.37.2.19 padding2	79
4.37.2.20 padding3	79
4.37.2.21 padding4	79
4.37.2.22 padding5	79
4.37.2.23 padding6	79
4.37.2.24 ss	79
4.37.2.25 trapno	80
4.38 uproc	80
4.38.1 Detailed Description	80
4.38.2 Member Data Documentation	80
4.38.2.1 name	80
4.38.2.2 pid	80
4.38.2.3 ppid	80
5 File Documentation	81
5.1 asm.h File Reference	81
5.1.1 Macro Definition Documentation	81
5.1.1.1 SEG_ASM	81
5.1.1.2 SEG_NULLASM	81
5.1.1.3 STA_R	82
5.1.1.4 STA_W	82
5.1.1.5 STA_X	82
5.2 asm.h	82
5.3 bio.c File Reference	82
5.3.1 Function Documentation	83
5.3.1.1 bget()	83
5.3.1.2 binit()	84
5.3.1.3 bread()	84
5.3.1.4 brelse()	84
5.3.1.5 bwrite()	85

5.3.2 Variable Documentation	85
5.3.2.1	85
5.3.2.2 buf	85
5.3.2.3 head	85
5.3.2.4 lock	86
5.4 bio.c	86
5.5 bio.d File Reference	88
5.6 bio.d	88
5.7 bootasm.d File Reference	88
5.8 bootasm.d	88
5.9 bootmain.c File Reference	88
5.9.1 Macro Definition Documentation	88
5.9.1.1 SECTSIZE	88
5.9.2 Function Documentation	89
5.9.2.1 bootmain()	89
5.9.2.2 readsect()	89
5.9.2.3 readseg()	90
5.9.2.4 waitdisk()	90
5.10 bootmain.c	90
5.11 bootmain.d File Reference	91
5.12 bootmain.d	91
5.13 buf.h File Reference	92
5.13.1 Macro Definition Documentation	92
5.13.1.1 B_DIRTY	92
5.13.1.2 B_VALID	92
5.14 buf.h	92
5.15 cat.c File Reference	93
5.15.1 Function Documentation	93
5.15.1.1 cat()	93
5.15.1.2 main()	94
5.15.2 Variable Documentation	94
5.15.2.1 buf	94
5.16 cat.c	94
5.17 cat.d File Reference	95
5.18 cat.d	95
5.19 console.c File Reference	95
5.19.1 Macro Definition Documentation	96
5.19.1.1 BACKSPACE	96
5.19.1.2 C	96
5.19.1.3 CRTPORT	96
5.19.1.4 INPUT_BUF	96
5.19.2 Function Documentation	97

5.19.2.1 cgaputc()	97
5.19.2.2 consoleinit()	97
5.19.2.3 consoleintr()	98
5.19.2.4 consoleread()	98
5.19.2.5 consolewrite()	99
5.19.2.6 consputc()	99
5.19.2.7 cprintf()	100
5.19.2.8 panic()	100
5.19.2.9 printint()	101
5.19.3 Variable Documentation	101
5.19.3.1 buf	101
5.19.3.2	102
5.19.3.3 crt	102
5.19.3.4 e	102
5.19.3.5	102
5.19.3.6 lock	102
5.19.3.7 locking	102
5.19.3.8 panicked	103
5.19.3.9 r	103
5.19.3.10 w	103
5.20 console.c	103
5.21 console.d File Reference	107
5.22 console.d	107
5.23 date.h File Reference	107
5.24 date.h	107
5.25 defs.h File Reference	107
5.25.1 Macro Definition Documentation	110
5.25.1.1 NELEM	110
5.25.2 Function Documentation	110
5.25.2.1 acquire()	111
5.25.2.2 acquiresleep()	111
5.25.2.3 allocvm()	111
5.25.2.4 argint()	112
5.25.2.5 argptr()	112
5.25.2.6 argstr()	113
5.25.2.7 begin_op()	113
5.25.2.8 binit()	113
5.25.2.9 bread()	114
5.25.2.10 brelse()	114
5.25.2.11 bwrite()	114
5.25.2.12 chpr()	115
5.25.2.13 clearpteu()	115

5.25.2.14 cmostime()	115
5.25.2.15 consoleinit()	116
5.25.2.16 consoleintr()	116
5.25.2.17 copyout()	117
5.25.2.18 copyuvm()	117
5.25.2.19 cprintf()	118
5.25.2.20 cps()	119
5.25.2.21 cpuid()	119
5.25.2.22 deallocuvm()	119
5.25.2.23 dirlink()	120
5.25.2.24 dirlookup()	120
5.25.2.25 end_op()	121
5.25.2.26 exec()	121
5.25.2.27 exit()	123
5.25.2.28 fetchint()	124
5.25.2.29 fetchstr()	124
5.25.2.30 filealloc()	124
5.25.2.31 fileclose()	125
5.25.2.32 filedup()	125
5.25.2.33 fileinit()	125
5.25.2.34 fileread()	126
5.25.2.35 filestat()	126
5.25.2.36 filewrite()	126
5.25.2.37 fork()	127
5.25.2.38 freevm()	128
5.25.2.39 getcallerpcs()	128
5.25.2.40 getprocs()	128
5.25.2.41 growproc()	129
5.25.2.42 holding()	129
5.25.2.43 holdingsleep()	129
5.25.2.44 ialloc()	130
5.25.2.45 ideinit()	130
5.25.2.46 ideintr()	130
5.25.2.47 iderw()	131
5.25.2.48 idtinit()	132
5.25.2.49 idup()	132
5.25.2.50 iinit()	132
5.25.2.51 ilock()	133
5.25.2.52 initlock()	133
5.25.2.53 initlog()	133
5.25.2.54 initsleeplock()	134
5.25.2.55 inituvm()	134

5.25.2.56 ioapicenable()	134
5.25.2.57 ioapicinit()	135
5.25.2.58 iput()	135
5.25.2.59 iunlock()	135
5.25.2.60 iunlockput()	136
5.25.2.61 iupdate()	136
5.25.2.62 kalloc()	136
5.25.2.63 kbdtintr()	137
5.25.2.64 kfree()	137
5.25.2.65 kill()	137
5.25.2.66 kinit1()	138
5.25.2.67 kinit2()	138
5.25.2.68 kvmalloc()	138
5.25.2.69 lapiceoi()	138
5.25.2.70 lapicid()	139
5.25.2.71 lapicinit()	139
5.25.2.72 lapicstartap()	140
5.25.2.73 loadvm()	140
5.25.2.74 log_write()	141
5.25.2.75 memcmp()	141
5.25.2.76 memmove()	141
5.25.2.77 memset()	142
5.25.2.78 microdelay()	142
5.25.2.79 mpinit()	143
5.25.2.80 mycpu()	143
5.25.2.81 myproc()	144
5.25.2.82 namecmp()	144
5.25.2.83 namei()	144
5.25.2.84 nameiparent()	144
5.25.2.85 panic()	145
5.25.2.86 picenable()	145
5.25.2.87 picinit()	145
5.25.2.88 pinit()	145
5.25.2.89 pipealloc()	146
5.25.2.90 pipeclose()	146
5.25.2.91 piperead()	147
5.25.2.92 pipewrite()	147
5.25.2.93 popcli()	148
5.25.2.94 procdump()	148
5.25.2.95 pushcli()	148
5.25.2.96 readi()	149
5.25.2.97 readsb()	149

5.25.2.98 release()	149
5.25.2.99 releasesleep()	150
5.25.2.100 safestrcpy()	150
5.25.2.101 sched()	151
5.25.2.102 scheduler()	151
5.25.2.103 seginit()	152
5.25.2.104 setproc()	152
5.25.2.105 setupkvm()	152
5.25.2.106 sleep()	153
5.25.2.107 stati()	153
5.25.2.108 strlen()	154
5.25.2.109 strncmp()	154
5.25.2.110 strncpy()	154
5.25.2.111 switchkvm()	155
5.25.2.112 switchvm()	155
5.25.2.113 swtch()	155
5.25.2.114 syscall()	156
5.25.2.115 timerinit()	156
5.25.2.116 tvinit()	156
5.25.2.117 uartinit()	156
5.25.2.118 uartintr()	157
5.25.2.119 uartputc()	157
5.25.2.120 userinit()	158
5.25.2.121 uva2ka()	158
5.25.2.122 wait()	159
5.25.2.123 wakeup()	159
5.25.2.124 writei()	160
5.25.2.125 yield()	160
5.25.3 Variable Documentation	160
5.25.3.1 ioapicid	161
5.25.3.2 ismp	161
5.25.3.3 lapic	161
5.25.3.4 ticks	161
5.25.3.5 tickslock	161
5.26 defs.h	162
5.27 echo.c File Reference	164
5.27.1 Function Documentation	164
5.27.1.1 main()	164
5.28 echo.c	165
5.29 echo.d File Reference	165
5.30 echo.d	165
5.31 elf.h File Reference	165

5.31.1 Macro Definition Documentation	165
5.31.1.1 ELF_MAGIC	165
5.31.1.2 ELF_PROG_FLAG_EXEC	166
5.31.1.3 ELF_PROG_FLAG_READ	166
5.31.1.4 ELF_PROG_FLAG_WRITE	166
5.31.1.5 ELF_PROG_LOAD	166
5.32 elf.h	166
5.33 entryother.d File Reference	167
5.34 entryother.d	167
5.35 exec.c File Reference	167
5.35.1 Function Documentation	167
5.35.1.1 exec()	168
5.36 exec.c	169
5.37 exec.d File Reference	170
5.38 exec.d	170
5.39 exit.c File Reference	171
5.39.1 Function Documentation	171
5.39.1.1 main()	171
5.40 exit.c	171
5.41 exit.d File Reference	171
5.42 exit.d	171
5.43 fcntl.h File Reference	172
5.43.1 Macro Definition Documentation	172
5.43.1.1 O_CREATE	172
5.43.1.2 O_RDONLY	172
5.43.1.3 O_RDWR	172
5.43.1.4 O_WRONLY	172
5.44 fcntl.h	173
5.45 file.c File Reference	173
5.45.1 Function Documentation	173
5.45.1.1 filealloc()	174
5.45.1.2 fileclose()	174
5.45.1.3 filedup()	174
5.45.1.4 fileinit()	175
5.45.1.5 fileread()	175
5.45.1.6 filestat()	175
5.45.1.7 filewrite()	176
5.45.2 Variable Documentation	176
5.45.2.1 devsw	176
5.45.2.2 file	176
5.45.2.3	177
5.45.2.4 lock	177

5.46 file.c	177
5.47 file.d File Reference	179
5.48 file.d	179
5.49 file.h File Reference	179
5.49.1 Macro Definition Documentation	179
5.49.1.1 CONSOLE	179
5.49.2 Variable Documentation	180
5.49.2.1 devsw	180
5.50 file.h	180
5.51 foo.c File Reference	180
5.51.1 Function Documentation	181
5.51.1.1 main()	181
5.52 foo.c	181
5.53 foo.d File Reference	182
5.54 foo.d	182
5.55 forktest.c File Reference	182
5.55.1 Macro Definition Documentation	182
5.55.1.1 N	182
5.55.2 Function Documentation	183
5.55.2.1 forktest()	183
5.55.2.2 main()	183
5.55.2.3 printf()	184
5.56 forktest.c	184
5.57 forktest.d File Reference	185
5.58 forktest.d	185
5.59 fs.c File Reference	185
5.59.1 Macro Definition Documentation	186
5.59.1.1 min	186
5.59.2 Function Documentation	186
5.59.2.1 balloc()	186
5.59.2.2 bfree()	187
5.59.2.3 bmap()	187
5.59.2.4 bzero()	188
5.59.2.5 dirlink()	188
5.59.2.6 dirlookup()	189
5.59.2.7 ialloc()	189
5.59.2.8 idup()	190
5.59.2.9 iget()	190
5.59.2.10 iinit()	191
5.59.2.11 ilock()	191
5.59.2.12 iput()	192
5.59.2.13 itrunc()	192

5.59.2.14 iunlock()	193
5.59.2.15 iunlockput()	193
5.59.2.16 iupdate()	193
5.59.2.17 namecmp()	194
5.59.2.18 namei()	194
5.59.2.19 nameiparent()	194
5.59.2.20 namex()	195
5.59.2.21 readi()	195
5.59.2.22 readsb()	196
5.59.2.23 skipelem()	196
5.59.2.24 stati()	197
5.59.2.25 writei()	197
5.59.3 Variable Documentation	197
5.59.3.1	198
5.59.3.2 inode	198
5.59.3.3 lock	198
5.59.3.4 sb	198
5.60 fs.c	198
5.61 fs.d File Reference	206
5.62 fs.d	206
5.63 fs.h File Reference	206
5.63.1 Macro Definition Documentation	207
5.63.1.1 BBLOCK	207
5.63.1.2 BPB	207
5.63.1.3 BSIZE	207
5.63.1.4 DIRSIZ	207
5.63.1.5 IBLOCK	207
5.63.1.6 IPB	208
5.63.1.7 MAXFILE	208
5.63.1.8 NDIRECT	208
5.63.1.9 NINDIRECT	208
5.63.1.10 ROOTINO	208
5.64 fs.h	209
5.65 grep.c File Reference	209
5.65.1 Function Documentation	210
5.65.1.1 grep()	210
5.65.1.2 main()	210
5.65.1.3 match()	211
5.65.1.4 matchhere()	211
5.65.1.5 matchstar()	212
5.65.2 Variable Documentation	212
5.65.2.1 buf	212

5.66 grep.c	212
5.67 grep.d File Reference	213
5.68 grep.d	213
5.69 ide.c File Reference	214
5.69.1 Macro Definition Documentation	214
5.69.1.1 IDE_BSY	215
5.69.1.2 IDE_CMD_RDMUL	215
5.69.1.3 IDE_CMD_READ	215
5.69.1.4 IDE_CMD_WRITE	215
5.69.1.5 IDE_CMD_WRMUL	215
5.69.1.6 IDE_DF	215
5.69.1.7 IDE_DRDY	216
5.69.1.8 IDE_ERR	216
5.69.1.9 SECTOR_SIZE	216
5.69.2 Function Documentation	216
5.69.2.1 ideinit()	216
5.69.2.2 ideintr()	217
5.69.2.3 iderw()	217
5.69.2.4 idestart()	218
5.69.2.5 idewait()	218
5.69.3 Variable Documentation	218
5.69.3.1 havdisk1	218
5.69.3.2 idelock	219
5.69.3.3 idequeue	219
5.70 ide.c	219
5.71 ide.d File Reference	221
5.72 ide.d	221
5.73 init.c File Reference	221
5.73.1 Function Documentation	221
5.73.1.1 main()	222
5.73.2 Variable Documentation	222
5.73.2.1 argv	222
5.74 init.c	222
5.75 init.d File Reference	223
5.76 init.d	223
5.77 initcode.d File Reference	223
5.78 initcode.d	223
5.79 ioapic.c File Reference	223
5.79.1 Macro Definition Documentation	224
5.79.1.1 INT_ACTIVELOW	224
5.79.1.2 INT_DISABLED	224
5.79.1.3 INT_LEVEL	224

5.79.1.4 INT_LOGICAL	224
5.79.1.5 IOAPIC	225
5.79.1.6 REG_ID	225
5.79.1.7 REG_TABLE	225
5.79.1.8 REG_VER	225
5.79.2 Function Documentation	225
5.79.2.1 ioapicenable()	225
5.79.2.2 ioapicinit()	226
5.79.2.3 ioapicread()	226
5.79.2.4 ioapicwrite()	226
5.79.3 Variable Documentation	226
5.79.3.1 ioapic	227
5.80 ioapic.c	227
5.81 ioapic.d File Reference	228
5.82 ioapic.d	228
5.83 kalloc.c File Reference	228
5.83.1 Function Documentation	228
5.83.1.1 freerange()	229
5.83.1.2 kalloc()	229
5.83.1.3 kfree()	229
5.83.1.4 kinit1()	230
5.83.1.5 kinit2()	230
5.83.2 Variable Documentation	230
5.83.2.1 end	230
5.83.2.2 freelist	230
5.83.2.3	230
5.83.2.4 lock	231
5.83.2.5 use_lock	231
5.84 kalloc.c	231
5.85 kalloc.d File Reference	232
5.86 kalloc.d	232
5.87 kbd.c File Reference	232
5.87.1 Function Documentation	233
5.87.1.1 kbdgetc()	233
5.87.1.2 kbdtintr()	233
5.88 kbd.c	234
5.89 kbd.d File Reference	234
5.90 kbd.d	234
5.91 kbd.h File Reference	234
5.91.1 Macro Definition Documentation	235
5.91.1.1 ALT	235
5.91.1.2 C	235

5.91.1.3 CAPSLOCK	236
5.91.1.4 CTL	236
5.91.1.5 E0ESC	236
5.91.1.6 KBDATAP	236
5.91.1.7 KBS_DIB	236
5.91.1.8 KBSTATP	236
5.91.1.9 KEY_DEL	237
5.91.1.10 KEY_DN	237
5.91.1.11 KEY_END	237
5.91.1.12 KEY_HOME	237
5.91.1.13 KEY_INS	237
5.91.1.14 KEY_LF	237
5.91.1.15 KEY_PGDN	238
5.91.1.16 KEY_PGUP	238
5.91.1.17 KEY_RT	238
5.91.1.18 KEY_UP	238
5.91.1.19 NO	238
5.91.1.20 NUMLOCK	238
5.91.1.21 SCROLLLOCK	239
5.91.1.22 SHIFT	239
5.91.2 Variable Documentation	239
5.91.2.1 ctlmap	239
5.91.2.2 normalmap	240
5.91.2.3 shiftcode	240
5.91.2.4 shiftmap	240
5.91.2.5 togglecode	241
5.92 kbd.h	241
5.93 kill.c File Reference	243
5.93.1 Function Documentation	243
5.93.1.1 main()	243
5.94 kill.c	243
5.95 kill.d File Reference	244
5.96 kill.d	244
5.97 lapic.c File Reference	244
5.97.1 Macro Definition Documentation	245
5.97.1.1 ASSERT	245
5.97.1.2 BCAST	245
5.97.1.3 BUSY	246
5.97.1.4 CMOS_PORT	246
5.97.1.5 CMOS_RETURN	246
5.97.1.6 CMOS_STATA	246
5.97.1.7 CMOS_STATB	246

5.97.1.8 CMOS_UIP	246
5.97.1.9 CONV	247
5.97.1.10 DAY	247
5.97.1.11 DEASSERT	247
5.97.1.12 DELIVS	247
5.97.1.13 ENABLE	247
5.97.1.14 EOI	247
5.97.1.15 ERROR	248
5.97.1.16 ESR	248
5.97.1.17 FIXED	248
5.97.1.18 HOURS	248
5.97.1.19 ICRHI	248
5.97.1.20 ICRLO	248
5.97.1.21 ID	249
5.97.1.22 INIT	249
5.97.1.23 LEVEL	249
5.97.1.24 LINT0	249
5.97.1.25 LINT1	249
5.97.1.26 MASKED	249
5.97.1.27 MINS	250
5.97.1.28 MONTH	250
5.97.1.29 PCINT	250
5.97.1.30 PERIODIC	250
5.97.1.31 SECS	250
5.97.1.32 STARTUP	250
5.97.1.33 SVR	251
5.97.1.34 TCCR	251
5.97.1.35 TDCR	251
5.97.1.36 TICR	251
5.97.1.37 TIMER	251
5.97.1.38 TPR	251
5.97.1.39 VER	252
5.97.1.40 X1	252
5.97.1.41 YEAR	252
5.97.2 Function Documentation	252
5.97.2.1 cmos_read()	252
5.97.2.2 cmostime()	253
5.97.2.3 fill_rtcddate()	253
5.97.2.4 lapiceoi()	253
5.97.2.5 lapicid()	254
5.97.2.6 lapicinit()	254
5.97.2.7 lapicstartap()	255

5.97.2.8 lapicw()	255
5.97.2.9 microdelay()	255
5.97.3 Variable Documentation	256
5.97.3.1 lapic	256
5.98 lapic.c	256
5.99 lapic.d File Reference	259
5.100 lapic.d	259
5.101 In.c File Reference	259
5.101.1 Function Documentation	259
5.101.1.1 main()	259
5.102 In.c	259
5.103 In.d File Reference	260
5.104 In.d	260
5.105 log.c File Reference	260
5.105.1 Function Documentation	260
5.105.1.1 begin_op()	261
5.105.1.2 commit()	261
5.105.1.3 end_op()	261
5.105.1.4 initlog()	262
5.105.1.5 install_trans()	262
5.105.1.6 log_write()	263
5.105.1.7 read_head()	263
5.105.1.8 recover_from_log()	263
5.105.1.9 write_head()	264
5.105.1.10 write_log()	264
5.105.2 Variable Documentation	264
5.105.2.1 log	264
5.106 log.c	265
5.107 log.d File Reference	267
5.108 log.d	267
5.109 ls.c File Reference	268
5.109.1 Function Documentation	268
5.109.1.1 fmtname()	268
5.109.1.2 ls()	268
5.109.1.3 main()	269
5.110 ls.c	269
5.111 ls.d File Reference	271
5.112 ls.d	271
5.113 main.c File Reference	271
5.113.1 Function Documentation	271
5.113.1.1 __attribute__()	271
5.113.1.2 mpenter()	272

5.113.1.3 mpmain()	272
5.113.1.4 startothers()	272
5.113.2 Variable Documentation	273
5.113.2.1 entrypmdir	273
5.114 main.c	273
5.115 main.d File Reference	275
5.116 main.d	275
5.117 memide.c File Reference	275
5.117.1 Function Documentation	275
5.117.1.1 ideinit()	275
5.117.1.2 ideintr()	276
5.117.1.3 iderw()	276
5.117.2 Variable Documentation	276
5.117.2.1 _binary_fs_img_size	276
5.117.2.2 _binary_fs_img_start	277
5.117.2.3 disksize	277
5.117.2.4 memdisk	277
5.118 memide.c	277
5.119 memlayout.h File Reference	278
5.119.1 Macro Definition Documentation	278
5.119.1.1 DEVSPACE	278
5.119.1.2 EXTMEM	278
5.119.1.3 KERNBASE	279
5.119.1.4 KERNLINK	279
5.119.1.5 P2V	279
5.119.1.6 P2V_WO	279
5.119.1.7 PHYSTOP	279
5.119.1.8 V2P	279
5.119.1.9 V2P_WO	280
5.120 memlayout.h	280
5.121 mkdir.c File Reference	280
5.121.1 Function Documentation	280
5.121.1.1 main()	280
5.122 mkdir.c	281
5.123 mkdir.d File Reference	281
5.124 mkdir.d	281
5.125 mkfs.c File Reference	281
5.125.1 Macro Definition Documentation	282
5.125.1.1 min	282
5.125.1.2 NINODES	282
5.125.1.3 stat	283
5.125.1.4 static_assert	283

5.125.2 Function Documentation	283
5.125.2.1 <code>balloc()</code>	283
5.125.2.2 <code>ialloc()</code>	283
5.125.2.3 <code>iappend()</code>	284
5.125.2.4 <code>main()</code>	284
5.125.2.5 <code>rinode()</code>	286
5.125.2.6 <code>rsect()</code>	286
5.125.2.7 <code>winode()</code>	286
5.125.2.8 <code>wsect()</code>	287
5.125.2.9 <code>xint()</code>	287
5.125.2.10 <code>xshort()</code>	287
5.125.3 Variable Documentation	287
5.125.3.1 <code>freeblock</code>	288
5.125.3.2 <code>freeinode</code>	288
5.125.3.3 <code>fsfd</code>	288
5.125.3.4 <code>nbitmap</code>	288
5.125.3.5 <code>nblocks</code>	288
5.125.3.6 <code>ninodeblocks</code>	289
5.125.3.7 <code>nlog</code>	289
5.125.3.8 <code>nmeta</code>	289
5.125.3.9 <code>sb</code>	289
5.125.3.10 <code>zeroes</code>	289
5.126 <code>mkfs.c</code>	290
5.127 <code>mmu.h</code> File Reference	293
5.127.1 Macro Definition Documentation	294
5.127.1.1 <code>CR0_PE</code>	294
5.127.1.2 <code>CR0_PG</code>	294
5.127.1.3 <code>CR0_WP</code>	295
5.127.1.4 <code>CR4_PSE</code>	295
5.127.1.5 <code>DPL_USER</code>	295
5.127.1.6 <code>FL_IF</code>	295
5.127.1.7 <code>NPENTRIES</code>	295
5.127.1.8 <code>NPTENTRIES</code>	295
5.127.1.9 <code>NSEGS</code>	296
5.127.1.10 <code>PDX</code>	296
5.127.1.11 <code>PDXSHIFT</code>	296
5.127.1.12 <code>PGADDR</code>	296
5.127.1.13 <code>PGROUNDDOWN</code>	296
5.127.1.14 <code>PGROUNDUP</code>	296
5.127.1.15 <code>PGSIZE</code>	297
5.127.1.16 <code>PTE_ADDR</code>	297
5.127.1.17 <code>PTE_FLAGS</code>	297

5.127.1.18 PTE_P	297
5.127.1.19 PTE_PS	297
5.127.1.20 PTE_U	297
5.127.1.21 PTE_W	298
5.127.1.22 PTX	298
5.127.1.23 PTXSHIFT	298
5.127.1.24 SEG	298
5.127.1.25 SEG16	298
5.127.1.26 SEG_KCODE	299
5.127.1.27 SEG_KDATA	299
5.127.1.28 SEG_TSS	299
5.127.1.29 SEG_UCODE	299
5.127.1.30 SEG_UDATA	299
5.127.1.31 SETGATE	300
5.127.1.32 STA_R	300
5.127.1.33 STA_W	300
5.127.1.34 STA_X	300
5.127.1.35 STS_IG32	300
5.127.1.36 STS_T32A	301
5.127.1.37 STS_TG32	301
5.127.2 Typedef Documentation	301
5.127.2.1 pte_t	301
5.128 mmu.h	301
5.129 mp.c File Reference	303
5.129.1 Function Documentation	304
5.129.1.1 mpconfig()	304
5.129.1.2 mpinit()	304
5.129.1.3 mpsearch()	305
5.129.1.4 mpsearch1()	305
5.129.1.5 sum()	306
5.129.2 Variable Documentation	306
5.129.2.1 cpus	306
5.129.2.2 ioapicid	306
5.129.2.3 ncpu	307
5.130 mp.c	307
5.131 mp.d File Reference	308
5.132 mp.d	308
5.133 mp.h File Reference	309
5.133.1 Macro Definition Documentation	309
5.133.1.1 MPBOOT	309
5.133.1.2 MPBUS	309
5.133.1.3 MPIOAPIC	309

5.133.1.4 MPIOINTR	310
5.133.1.5 MPLINTR	310
5.133.1.6 MPPROC	310
5.134 mp.h	310
5.135 nice.c File Reference	311
5.135.1 Function Documentation	311
5.135.1.1 main()	311
5.136 nice.c	312
5.137 nice.d File Reference	312
5.138 nice.d	312
5.139 param.h File Reference	312
5.139.1 Macro Definition Documentation	312
5.139.1.1 FSSIZE	313
5.139.1.2 KSTACKSIZE	313
5.139.1.3 LOGSIZE	313
5.139.1.4 MAXARG	313
5.139.1.5 MAXOPBLOCKS	313
5.139.1.6 NBUF	313
5.139.1.7 NCPU	314
5.139.1.8 NDEV	314
5.139.1.9 NFILE	314
5.139.1.10 NINODE	314
5.139.1.11 NOFILE	314
5.139.1.12 NPROC	314
5.139.1.13 ROOTDEV	315
5.140 param.h	315
5.141 picirq.c File Reference	315
5.141.1 Macro Definition Documentation	315
5.141.1.1 IO_PIC1	315
5.141.1.2 IO_PIC2	316
5.141.2 Function Documentation	316
5.141.2.1 picinit()	316
5.142 picirq.c	316
5.143 picirq.d File Reference	316
5.144 picirq.d	316
5.145 pipe.c File Reference	317
5.145.1 Macro Definition Documentation	317
5.145.1.1 PIPESIZE	317
5.145.2 Function Documentation	317
5.145.2.1 pipealloc()	318
5.145.2.2 pipeclose()	318
5.145.2.3 piperead()	319

5.145.2.4 <code>pipewrite()</code>	319
5.146 <code>pipe.c</code>	320
5.147 <code>pipe.d</code> File Reference	321
5.148 <code>pipe.d</code>	321
5.149 <code>printf.c</code> File Reference	321
5.149.1 Function Documentation	322
5.149.1.1 <code>printf()</code>	322
5.149.1.2 <code>printint()</code>	322
5.149.1.3 <code>putc()</code>	323
5.150 <code>printf.c</code>	323
5.151 <code>printf.d</code> File Reference	324
5.152 <code>printf.d</code>	324
5.153 <code>proc.c</code> File Reference	324
5.153.1 Function Documentation	325
5.153.1.1 <code>allocproc()</code>	326
5.153.1.2 <code>chpr()</code>	326
5.153.1.3 <code>cps()</code>	327
5.153.1.4 <code>cpuid()</code>	327
5.153.1.5 <code>exit()</code>	327
5.153.1.6 <code>fork()</code>	328
5.153.1.7 <code>forkret()</code>	329
5.153.1.8 <code>growproc()</code>	329
5.153.1.9 <code>kill()</code>	330
5.153.1.10 <code>mycpu()</code>	330
5.153.1.11 <code>myproc()</code>	330
5.153.1.12 <code>pinit()</code>	331
5.153.1.13 <code>procdump()</code>	331
5.153.1.14 <code>sched()</code>	331
5.153.1.15 <code>scheduler()</code>	332
5.153.1.16 <code>sleep()</code>	332
5.153.1.17 <code>trapret()</code>	333
5.153.1.18 <code>userinit()</code>	333
5.153.1.19 <code>wait()</code>	334
5.153.1.20 <code>wakeup()</code>	334
5.153.1.21 <code>wakeup1()</code>	335
5.153.1.22 <code>yield()</code>	335
5.153.2 Variable Documentation	335
5.153.2.1 <code>initproc</code>	335
5.153.2.2 <code>lock</code>	335
5.153.2.3 <code>nextpid</code>	336
5.153.2.4 <code>proc</code>	336
5.153.2.5	336

5.154 proc.c	336
5.155 proc.d File Reference	343
5.156 proc.d	343
5.157 proc.h File Reference	343
5.157.1 Enumeration Type Documentation	343
5.157.1.1 procstate	343
5.157.2 Variable Documentation	344
5.157.2.1 cpus	344
5.157.2.2 ncpu	344
5.158 proc.h	344
5.159 ps.c File Reference	345
5.159.1 Function Documentation	345
5.159.1.1 main()	345
5.160 ps.c	346
5.161 ps.d File Reference	346
5.162 ps.d	346
5.163 pstree.c File Reference	346
5.163.1 Macro Definition Documentation	346
5.163.1.1 MAXPROC	346
5.163.2 Function Documentation	347
5.163.2.1 main()	347
5.164 pstree.c	347
5.165 pstree.d File Reference	348
5.166 pstree.d	348
5.167 README.md File Reference	348
5.168 rm.c File Reference	348
5.168.1 Function Documentation	348
5.168.1.1 main()	349
5.169 rm.c	349
5.170 rm.d File Reference	349
5.171 rm.d	349
5.172 sh.c File Reference	349
5.172.1 Macro Definition Documentation	350
5.172.1.1 BACK	351
5.172.1.2 EXEC	351
5.172.1.3 LIST	351
5.172.1.4 MAXARGS	351
5.172.1.5 PIPE	351
5.172.1.6 REDIR	351
5.172.2 Function Documentation	352
5.172.2.1 backcmd()	352
5.172.2.2 execcmd()	352

5.172.2.3 fork1()	352
5.172.2.4 getcmd()	353
5.172.2.5 gettoken()	353
5.172.2.6 listcmd()	354
5.172.2.7 main()	354
5.172.2.8 nulterminate()	354
5.172.2.9 panic()	355
5.172.2.10 parseblock()	355
5.172.2.11 parsecmd()	356
5.172.2.12 parseexec()	356
5.172.2.13 parseline()	357
5.172.2.14 parsepipe()	357
5.172.2.15 parseredirs()	357
5.172.2.16 peek()	358
5.172.2.17 pipecmd()	358
5.172.2.18 redircmd()	359
5.172.2.19 runcmd()	359
5.172.3 Variable Documentation	360
5.172.3.1 symbols	360
5.172.3.2 whitespace	360
5.173 sh.c	361
5.174 sh.d File Reference	366
5.175 sh.d	366
5.176 shutdown.c File Reference	366
5.176.1 Function Documentation	367
5.176.1.1 main()	367
5.177 shutdown.c	367
5.178 shutdown.d File Reference	367
5.179 shutdown.d	367
5.180 sleeplock.c File Reference	368
5.180.1 Function Documentation	368
5.180.1.1 acquiresleep()	368
5.180.1.2 holdingsleep()	368
5.180.1.3 initsleeplock()	369
5.180.1.4 releasesleep()	369
5.181 sleeplock.c	369
5.182 sleeplock.d File Reference	370
5.183 sleeplock.d	370
5.184 sleeplock.h File Reference	370
5.185 sleeplock.h	370
5.186 spinlock.c File Reference	370
5.186.1 Function Documentation	371

5.186.1.1 acquire()	371
5.186.1.2 getcallerpcs()	371
5.186.1.3 holding()	372
5.186.1.4 initlock()	372
5.186.1.5 popcli()	372
5.186.1.6 pushcli()	373
5.186.1.7 release()	373
5.187 spinlock.c	373
5.188 spinlock.d File Reference	375
5.189 spinlock.d	375
5.190 spinlock.h File Reference	375
5.191 spinlock.h	375
5.192 stat.h File Reference	375
5.192.1 Macro Definition Documentation	376
5.192.1.1 T_DEV	376
5.192.1.2 T_DIR	376
5.192.1.3 T_FILE	376
5.193 stat.h	376
5.194 stressfs.c File Reference	377
5.194.1 Function Documentation	377
5.194.1.1 main()	377
5.195 stressfs.c	378
5.196 stressfs.d File Reference	378
5.197 stressfs.d	378
5.198 string.c File Reference	378
5.198.1 Function Documentation	379
5.198.1.1 memcmp()	379
5.198.1.2 memcpy()	379
5.198.1.3 memmove()	380
5.198.1.4 memset()	380
5.198.1.5 safestrcpy()	380
5.198.1.6 strlen()	381
5.198.1.7 strncmp()	381
5.198.1.8 strncpy()	381
5.199 string.c	382
5.200 string.d File Reference	383
5.201 string.d	383
5.202 syscall.c File Reference	383
5.202.1 Function Documentation	384
5.202.1.1 argint()	384
5.202.1.2 argptr()	385
5.202.1.3 argstr()	385

5.202.1.4	fetchint()	385
5.202.1.5	fetchstr()	386
5.202.1.6	sys_chdir()	386
5.202.1.7	sys_chpr()	386
5.202.1.8	sys_close()	387
5.202.1.9	sys_cps()	387
5.202.1.10	sys_dup()	387
5.202.1.11	sys_exec()	388
5.202.1.12	sys_exit()	388
5.202.1.13	sys_fork()	388
5.202.1.14	sys_fstat()	388
5.202.1.15	sys_getpid()	389
5.202.1.16	sys_getprocs()	389
5.202.1.17	sys_halt()	389
5.202.1.18	sys_kill()	390
5.202.1.19	sys_link()	390
5.202.1.20	sys_mkdir()	391
5.202.1.21	sys_mknod()	391
5.202.1.22	sys_open()	391
5.202.1.23	sys_pipe()	392
5.202.1.24	sys_read()	392
5.202.1.25	sys_sbrk()	393
5.202.1.26	sys_sleep()	393
5.202.1.27	sys_unlink()	394
5.202.1.28	sys_uptime()	394
5.202.1.29	sys_wait()	395
5.202.1.30	sys_write()	395
5.202.1.31	syscall()	395
5.202.2	Variable Documentation	395
5.202.2.1	syscalls	396
5.203	syscall.c	396
5.204	syscall.d File Reference	398
5.205	syscall.d	398
5.206	syscall.h File Reference	398
5.206.1	Macro Definition Documentation	399
5.206.1.1	SYS_chdir	399
5.206.1.2	SYS_chpr	399
5.206.1.3	SYS_close	399
5.206.1.4	SYS_cps	399
5.206.1.5	SYS_dup	400
5.206.1.6	SYS_exec	400
5.206.1.7	SYS_exit	400

5.206.1.8 SYS_fork	400
5.206.1.9 SYS_fstat	400
5.206.1.10 SYS_getpid	400
5.206.1.11 SYS_getprocs	401
5.206.1.12 SYS_halt	401
5.206.1.13 SYS_kill	401
5.206.1.14 SYS_link	401
5.206.1.15 SYS_mkdir	401
5.206.1.16 SYS_mknod	401
5.206.1.17 SYS_open	402
5.206.1.18 SYS_pipe	402
5.206.1.19 SYS_read	402
5.206.1.20 SYS_sbrk	402
5.206.1.21 SYS_sleep	402
5.206.1.22 SYS_unlink	402
5.206.1.23 SYS_uptime	403
5.206.1.24 SYS_wait	403
5.206.1.25 SYS_write	403
5.207 syscall.h	403
5.208 sysfile.c File Reference	404
5.208.1 Function Documentation	404
5.208.1.1 argfd()	404
5.208.1.2 create()	405
5.208.1.3 fdalloc()	406
5.208.1.4 isdirempty()	406
5.208.1.5 sys_chdir()	406
5.208.1.6 sys_close()	407
5.208.1.7 sys_dup()	407
5.208.1.8 sys_exec()	407
5.208.1.9 sys_fstat()	408
5.208.1.10 sys_link()	408
5.208.1.11 sys_mkdir()	409
5.208.1.12 sys_mknod()	409
5.208.1.13 sys_open()	409
5.208.1.14 sys_pipe()	410
5.208.1.15 sys_read()	410
5.208.1.16 sys_unlink()	411
5.208.1.17 sys_write()	412
5.209 sysfile.c	412
5.210 sysfile.d File Reference	417
5.211 sysfile.d	417
5.212 sysproc.c File Reference	417

5.212.1 Function Documentation	418
5.212.1.1 sys_chpr()	418
5.212.1.2 sys_cps()	418
5.212.1.3 sys_exit()	419
5.212.1.4 sys_fork()	419
5.212.1.5 sys_getpid()	419
5.212.1.6 sys_getprocs()	419
5.212.1.7 sys_halt()	420
5.212.1.8 sys_kill()	420
5.212.1.9 sys_sbrk()	420
5.212.1.10 sys_sleep()	421
5.212.1.11 sys_uptime()	421
5.212.1.12 sys_wait()	421
5.212.2 Variable Documentation	421
5.212.2.1 lock	422
5.212.2.2 proc	422
5.212.2.3	422
5.213 sysproc.c	422
5.214 sysproc.d File Reference	424
5.215 sysproc.d	424
5.216 trap.c File Reference	424
5.216.1 Function Documentation	425
5.216.1.1 idtinit()	425
5.216.1.2 trap()	425
5.216.1.3 tvinit()	426
5.216.2 Variable Documentation	426
5.216.2.1 idt	427
5.216.2.2 ticks	427
5.216.2.3 tickslock	427
5.216.2.4 vectors	427
5.217 trap.c	428
5.218 trap.d File Reference	429
5.219 trap.d	429
5.220 traps.h File Reference	429
5.220.1 Macro Definition Documentation	430
5.220.1.1 IRQ_COM1	430
5.220.1.2 IRQ_ERROR	430
5.220.1.3 IRQ_IDE	430
5.220.1.4 IRQ_KBD	430
5.220.1.5 IRQ_SPURIOUS	431
5.220.1.6 IRQ_TIMER	431
5.220.1.7 T_ALIGN	431

5.220.1.8 T_BOUND	431
5.220.1.9 T_BRKPT	431
5.220.1.10 T_DBLFLT	431
5.220.1.11 T_DEBUG	432
5.220.1.12 T_DEFAULT	432
5.220.1.13 T_DEVICE	432
5.220.1.14 T_DIVIDE	432
5.220.1.15 T_FPERR	432
5.220.1.16 T_GPFLT	432
5.220.1.17 T_ILLOP	433
5.220.1.18 T_IRQ0	433
5.220.1.19 T_MCHK	433
5.220.1.20 T_NMI	433
5.220.1.21 T_OFLOW	433
5.220.1.22 T_PGFLT	433
5.220.1.23 T_SEGNP	434
5.220.1.24 T_SIMDERR	434
5.220.1.25 T_STACK	434
5.220.1.26 T_SYSCALL	434
5.220.1.27 T_TSS	434
5.221 traps.h	435
5.222 types.h File Reference	435
5.222.1 Typedef Documentation	435
5.222.1.1 pde_t	435
5.222.1.2 uchar	436
5.222.1.3 uint	436
5.222.1.4 ushort	436
5.223 types.h	436
5.224 uart.c File Reference	436
5.224.1 Macro Definition Documentation	437
5.224.1.1 COM1	437
5.224.2 Function Documentation	437
5.224.2.1 uartgetc()	437
5.224.2.2 uartinit()	438
5.224.2.3 uartintr()	438
5.224.2.4 uartputc()	438
5.224.3 Variable Documentation	439
5.224.3.1 uart	439
5.225 uart.c	439
5.226 uart.d File Reference	440
5.227 uart.d	440
5.228 ulib.c File Reference	440

5.228.1 Function Documentation	440
5.228.1.1 atoi()	441
5.228.1.2 gets()	441
5.228.1.3 memmove()	441
5.228.1.4 memset()	442
5.228.1.5 stat()	442
5.228.1.6 strchr()	442
5.228.1.7 strcmp()	443
5.228.1.8 strcpy()	443
5.228.1.9 strlen()	443
5.229 ulib.c	444
5.230 ulib.d File Reference	445
5.231 ulib.d	445
5.232 umalloc.c File Reference	445
5.232.1 Typedef Documentation	446
5.232.1.1 Align	446
5.232.1.2 Header	446
5.232.2 Function Documentation	446
5.232.2.1 free()	446
5.232.2.2 malloc()	447
5.232.2.3 morecore()	447
5.232.3 Variable Documentation	447
5.232.3.1 base	448
5.232.3.2 freep	448
5.233 umalloc.c	448
5.234 umalloc.d File Reference	449
5.235 umalloc.d	449
5.236 uproc.h File Reference	449
5.237 uproc.h	449
5.238 user.h File Reference	450
5.238.1 Function Documentation	450
5.238.1.1 atoi()	451
5.238.1.2 chdir()	451
5.238.1.3 chpr()	451
5.238.1.4 close()	451
5.238.1.5 cps()	452
5.238.1.6 dup()	452
5.238.1.7 exec()	452
5.238.1.8 exit()	453
5.238.1.9 fork()	454
5.238.1.10 free()	455
5.238.1.11 fstat()	455

5.238.1.12 getpid()	455
5.238.1.13 getprocs()	456
5.238.1.14 gets()	456
5.238.1.15 halt()	456
5.238.1.16 kill()	456
5.238.1.17 link()	457
5.238.1.18 malloc()	457
5.238.1.19 memmove()	457
5.238.1.20 memset()	458
5.238.1.21 mkdir()	458
5.238.1.22 mknod()	458
5.238.1.23 open()	458
5.238.1.24 pipe()	459
5.238.1.25 printf()	459
5.238.1.26 read()	459
5.238.1.27 sbrk()	459
5.238.1.28 sleep()	459
5.238.1.29 stat()	460
5.238.1.30 strchr()	460
5.238.1.31 strcmp()	460
5.238.1.32 strcpy()	461
5.238.1.33 strlen()	461
5.238.1.34 unlink()	461
5.238.1.35 uptime()	461
5.238.1.36 wait()	462
5.238.1.37 write()	462
5.239 user.h	462
5.240 usertests.c File Reference	463
5.240.1 Macro Definition Documentation	464
5.240.1.1 BIG	464
5.240.1.2 RTC_ADDR	465
5.240.1.3 RTC_DATA	465
5.240.2 Function Documentation	465
5.240.2.1 argptest()	465
5.240.2.2 bigargtest()	465
5.240.2.3 bigdir()	466
5.240.2.4 bigfile()	467
5.240.2.5 bigwrite()	467
5.240.2.6 bsstest()	468
5.240.2.7 concreate()	468
5.240.2.8 createdelete()	469
5.240.2.9 createtest()	470

5.240.2.10	dirfile()	471
5.240.2.11	dirtest()	472
5.240.2.12	exectest()	472
5.240.2.13	exitputtest()	473
5.240.2.14	exitwait()	473
5.240.2.15	forktest()	474
5.240.2.16	fourfiles()	474
5.240.2.17	fourteen()	475
5.240.2.18	fsfull()	476
5.240.2.19	iputtest()	476
5.240.2.20	iref()	477
5.240.2.21	linktest()	478
5.240.2.22	linkunlink()	478
5.240.2.23	main()	479
5.240.2.24	mem()	480
5.240.2.25	openiputtest()	480
5.240.2.26	opentest()	481
5.240.2.27	pipe1()	481
5.240.2.28	preempt()	482
5.240.2.29	rand()	483
5.240.2.30	rmdot()	483
5.240.2.31	sbrktest()	484
5.240.2.32	sharedfd()	485
5.240.2.33	subdir()	486
5.240.2.34	uio()	488
5.240.2.35	unlinkread()	489
5.240.2.36	validateint()	490
5.240.2.37	validatetest()	490
5.240.2.38	writetest()	491
5.240.2.39	writetest1()	491
5.240.3	Variable Documentation	492
5.240.3.1	buf	492
5.240.3.2	echoargv	492
5.240.3.3	name	493
5.240.3.4	randstate	493
5.240.3.5	stdout	493
5.240.3.6	unit	493
5.241	usertests.c	494
5.242	usertests.d File Reference	514
5.243	usertests.d	514
5.244	vm.c File Reference	515
5.244.1	Function Documentation	515

5.244.1.1 allocvm()	516
5.244.1.2 clearpteu()	516
5.244.1.3 copyout()	517
5.244.1.4 copyvm()	517
5.244.1.5 deallocvm()	518
5.244.1.6 freevm()	518
5.244.1.7 initvm()	519
5.244.1.8 kvmalloc()	519
5.244.1.9 loadvm()	519
5.244.1.10 mappages()	520
5.244.1.11 seginit()	520
5.244.1.12 setupkvm()	521
5.244.1.13 switchkvm()	521
5.244.1.14 switchvm()	521
5.244.1.15 uva2ka()	522
5.244.1.16 walkpgdir()	522
5.244.2 Variable Documentation	522
5.244.2.1 data	522
5.244.2.2 kmap	523
5.244.2.3 kpgdir	523
5.245 vm.c	523
5.246 vm.d File Reference	528
5.247 vm.d	528
5.248 wc.c File Reference	528
5.248.1 Function Documentation	528
5.248.1.1 main()	528
5.248.1.2 wc()	529
5.248.2 Variable Documentation	529
5.248.2.1 buf	529
5.249 wc.c	529
5.250 wc.d File Reference	530
5.251 wc.d	530
5.252 x86.h File Reference	530
5.252.1 Function Documentation	531
5.252.1.1 cli()	531
5.252.1.2 inb()	531
5.252.1.3 insl()	531
5.252.1.4 lcr3()	532
5.252.1.5 lgdt()	532
5.252.1.6 lidt()	532
5.252.1.7 loadgs()	533
5.252.1.8 ltr()	533

5.252.1.9 outb()	533
5.252.1.10 outsl()	533
5.252.1.11 outw()	534
5.252.1.12 rcr2()	534
5.252.1.13 readeflags()	534
5.252.1.14 sti()	534
5.252.1.15 stosb()	535
5.252.1.16 stosl()	535
5.252.1.17 xchg()	535
5.253 x86.h	536
5.254 xv6.dox File Reference	538
5.255 zombie.c File Reference	538
5.255.1 Function Documentation	538
5.255.1.1 main()	538
5.256 zombie.c	538
5.257 zombie.d File Reference	539
5.258 zombie.d	539

Index	541
--------------	------------

Chapter 1

ENVIRONMENT PREPARETION

1.1 LINUX

1. `sudo apt-get install git wget qemu qemu`
2. `sudo apt-get install libc6-dev:i386 gcc`
3. `sudo apt-get install gdb`
4. `git clone https://github.com/subhrendu1987/xv6-public/`
5. `cd xv6-public && make qemu`
6. If you have 64 bit OS there is a chance Makefile will not be able to find qemu. Check using the command `make qemu-nox` and see the output. In that case you should edit the Makefile at line 54 and add the following code: `QEMU = qemu-system-x86_64`

1. `make qemu-gdb`

1. Open new terminal `cd xv6-public && gdb ./kernel`

1.2 BUILDING AND RUNNING XV6

- To build xv6 on an x86 ELF machine (like Linux or FreeBSD), run "make".
- On non-x86 or non-ELF machines (like OS X, even on x86), you may need to install a cross-compiler gcc suite capable of producing x86 ELF binaries (see <https://pdos.csail.mit.edu/6.828/>).
- Then run "make TOOLPREFIX=i386-jos-elf-". Now install the QEMU PC simulator and run "make qemu".

Chapter 2

Class Index

2.1 Class List

Here are the classes, structs, unions and interfaces with brief descriptions:

backcmd	9
buf	10
cmd	12
context	13
cpu	14
devsw	16
dinode	17
dirent	19
elfhdr	20
execcmd	23
file	24
gatedesc	26
header	28
inode	30
ioapic	32
kmap	33
listcmd	35
log	36
logheader	37
mp	38
mpconf	40
mpioapic	43
mproc	44
pipe	46
pipecmd	48
proc	49
proghdr	52
redircmd	54
rtcdat	56
run	58
segdesc	58
sleeplock	61
spinlock	62
stat	63
superblock	65
taskstate	67
trapframe	74
uproc	80

Chapter 3

File Index

3.1 File List

Here is a list of all files with brief descriptions:

asm.h	81
bio.c	82
bio.d	88
bootasm.d	88
bootmain.c	88
bootmain.d	91
buf.h	92
cat.c	93
cat.d	95
console.c	95
console.d	107
date.h	107
defs.h	107
echo.c	164
echo.d	165
elf.h	165
entryother.d	167
exec.c	167
exec.d	170
exit.c	171
exit.d	171
fcntl.h	172
file.c	173
file.d	179
file.h	179
foo.c	180
foo.d	182
forktest.c	182
forktest.d	185
fs.c	185
fs.d	206
fs.h	206
grep.c	209
grep.d	213
ide.c	214

ide.d	221
init.c	221
init.d	223
initcode.d	223
ioapic.c	223
ioapic.d	228
kalloc.c	228
kalloc.d	232
kbd.c	232
kbd.d	234
kbd.h	234
kill.c	243
kill.d	244
lapic.c	244
lapic.d	259
ln.c	259
ln.d	260
log.c	260
log.d	267
ls.c	268
ls.d	271
main.c	271
main.d	275
memide.c	275
memlayout.h	278
mkdir.c	280
mkdir.d	281
mkfs.c	281
mmu.h	293
mp.c	303
mp.d	308
mp.h	309
nice.c	311
nice.d	312
param.h	312
picirq.c	315
picirq.d	316
pipe.c	317
pipe.d	321
printf.c	321
printf.d	324
proc.c	324
proc.d	343
proc.h	343
ps.c	345
ps.d	346
pstree.c	346
pstree.d	348
rm.c	348
rm.d	349
sh.c	349
sh.d	366
shutdown.c	366
shutdown.d	367
sleeplock.c	368
sleeplock.d	370
sleeplock.h	370
spinlock.c	370

spinlock.d	375
spinlock.h	375
stat.h	375
stressfs.c	377
stressfs.d	378
string.c	378
string.d	383
syscall.c	383
syscall.d	398
syscall.h	398
sysfile.c	404
sysfile.d	417
sysproc.c	417
sysproc.d	424
trap.c	424
trap.d	429
traps.h	429
types.h	435
uart.c	436
uart.d	440
ulib.c	440
ulib.d	445
umalloc.c	445
umalloc.d	449
uproc.h	449
user.h	450
usertests.c	463
usertests.d	514
vm.c	515
vm.d	528
wc.c	528
wc.d	530
x86.h	530
zombie.c	538
zombie.d	539

Chapter 4

Class Documentation

4.1 backcmd

Public Attributes

- struct [cmd](#) * [cmd](#)
- int [type](#)

4.1.1 Detailed Description

Definition at line [47](#) of file [sh.c](#).

4.1.2 Member Data Documentation

4.1.2.1 cmd

```
struct cmd* backcmd::cmd
```

Definition at line [49](#) of file [sh.c](#).

Referenced by [backcmd\(\)](#), [nulterminate\(\)](#), and [runcmd\(\)](#).

4.1.2.2 type

```
int backcmd::type
```

Definition at line [48](#) of file [sh.c](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.2 buf

```
#include <buf.h>
```

Public Attributes

- [uint blockno](#)
- [uchar data \[BSIZE\]](#)
- [uint dev](#)
- [int flags](#)
- [struct sleeplock lock](#)
- [struct buf * next](#)
- [struct buf * prev](#)
- [struct buf * qnext](#)
- [uint refcnt](#)

4.2.1 Detailed Description

Definition at line 1 of file [buf.h](#).

4.2.2 Member Data Documentation

4.2.2.1 blockno

```
uint buf::blockno
```

Definition at line 4 of file [buf.h](#).

Referenced by [bget\(\)](#), [bread\(\)](#), [iderw\(\)](#), [idestart\(\)](#), and [log_write\(\)](#).

4.2.2.2 data

```
uchar buf::data[BSIZE]
```

Definition at line 10 of file [buf.h](#).

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idestart\(\)](#), [ilock\(\)](#), [install_trans\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [read_head\(\)](#), [readi\(\)](#), [readsb\(\)](#), [write_head\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

4.2.2.3 dev

```
uint buf::dev
```

Definition at line 3 of file [buf.h](#).

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bget\(\)](#), [bread\(\)](#), [bzero\(\)](#), [iderw\(\)](#), [idestart\(\)](#), and [readsb\(\)](#).

4.2.2.4 flags

```
int buf::flags
```

Definition at line 2 of file [buf.h](#).

Referenced by [bget\(\)](#), [bread\(\)](#), [bwrite\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idestart\(\)](#), and [log_write\(\)](#).

4.2.2.5 lock

```
struct sleeplock buf::lock
```

Definition at line 5 of file [buf.h](#).

Referenced by [bget\(\)](#), [binit\(\)](#), [brelse\(\)](#), [bwrite\(\)](#), and [iderw\(\)](#).

4.2.2.6 next

```
struct buf\* buf::next
```

Definition at line 8 of file [buf.h](#).

Referenced by [bget\(\)](#), [binit\(\)](#), and [brelse\(\)](#).

4.2.2.7 prev

```
struct buf\* buf::prev
```

Definition at line 7 of file [buf.h](#).

Referenced by [bget\(\)](#), [binit\(\)](#), and [brelse\(\)](#).

4.2.2.8 qnext

```
struct buf* buf::qnext
```

Definition at line 9 of file [buf.h](#).

Referenced by [ideintr\(\)](#), and [iderw\(\)](#).

4.2.2.9 refcnt

```
uint buf::refcnt
```

Definition at line 6 of file [buf.h](#).

Referenced by [bget\(\)](#), and [brelse\(\)](#).

The documentation for this struct was generated from the following file:

- [buf.h](#)

4.3 cmd

Public Attributes

- int [type](#)

4.3.1 Detailed Description

Definition at line 16 of file [sh.c](#).

4.3.2 Member Data Documentation

4.3.2.1 type

```
int cmd::type
```

Definition at line 17 of file [sh.c](#).

Referenced by [backcmd\(\)](#), [execcmd\(\)](#), [listcmd\(\)](#), [nulterminate\(\)](#), [pipecmd\(\)](#), [redircmd\(\)](#), and [runcmd\(\)](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.4 context

```
#include <proc.h>
```

Public Attributes

- [uint ebp](#)
- [uint ebx](#)
- [uint edi](#)
- [uint eip](#)
- [uint esi](#)

4.4.1 Detailed Description

Definition at line [27](#) of file [proc.h](#).

4.4.2 Member Data Documentation

4.4.2.1 ebp

```
uint context::ebp
```

Definition at line [31](#) of file [proc.h](#).

Referenced by [procdump\(\)](#).

4.4.2.2 ebx

```
uint context::ebx
```

Definition at line [30](#) of file [proc.h](#).

4.4.2.3 edi

```
uint context::edi
```

Definition at line [28](#) of file [proc.h](#).

4.4.2.4 eip

```
uint context::eip
```

Definition at line 32 of file [proc.h](#).

Referenced by [allocproc\(\)](#).

4.4.2.5 esi

```
uint context::esi
```

Definition at line 29 of file [proc.h](#).

The documentation for this struct was generated from the following file:

- [proc.h](#)

4.5 cpu

```
#include <proc.h>
```

Public Attributes

- [uchar](#) [apicid](#)
- [struct](#) [segdesc](#) [gdt](#) [[NSEGS](#)]
- [int](#) [intena](#)
- [int](#) [ncli](#)
- [struct](#) [proc](#) * [proc](#)
- [struct](#) [context](#) * [scheduler](#)
- [volatile](#) [uint](#) [started](#)
- [struct](#) [taskstate](#) [ts](#)

4.5.1 Detailed Description

Definition at line 2 of file [proc.h](#).

4.5.2 Member Data Documentation

4.5.2.1 apicid

```
uchar cpu::apicid
```

Definition at line 3 of file [proc.h](#).

Referenced by [mpinit\(\)](#), [mycpu\(\)](#), and [startothers\(\)](#).

4.5.2.2 gdt

```
struct segdesc cpu::gdt[NSEGS]
```

Definition at line 6 of file [proc.h](#).

Referenced by [seginit\(\)](#), and [switchvm\(\)](#).

4.5.2.3 intena

```
int cpu::intena
```

Definition at line 9 of file [proc.h](#).

Referenced by [pushcli\(\)](#), and [sched\(\)](#).

4.5.2.4 ncli

```
int cpu::ncli
```

Definition at line 8 of file [proc.h](#).

Referenced by [pushcli\(\)](#).

4.5.2.5 proc

```
struct proc* cpu::proc
```

Definition at line 10 of file [proc.h](#).

Referenced by [myproc\(\)](#), and [scheduler\(\)](#).

4.5.2.6 scheduler

```
struct context* cpu::scheduler
```

Definition at line 4 of file [proc.h](#).

Referenced by [scheduler\(\)](#).

4.5.2.7 started

```
volatile uint cpu::started
```

Definition at line 7 of file [proc.h](#).

Referenced by [startothers\(\)](#).

4.5.2.8 ts

```
struct taskstate cpu::ts
```

Definition at line 5 of file [proc.h](#).

Referenced by [switchvm\(\)](#).

The documentation for this struct was generated from the following file:

- [proc.h](#)

4.6 devsw

```
#include <file.h>
```

Public Attributes

- [int](#)(* [read](#))(struct [inode](#) *, char *, int)
- [int](#)(* [write](#))(struct [inode](#) *, char *, int)

4.6.1 Detailed Description

Definition at line 30 of file [file.h](#).

4.6.2 Member Data Documentation

4.6.2.1 read

```
int (* devsw::read) (struct inode *, char *, int)
```

Definition at line 31 of file [file.h](#).

Referenced by [consoleinit\(\)](#), and [readi\(\)](#).

4.6.2.2 write

```
int (* devsw::write) (struct inode *, char *, int)
```

Definition at line 32 of file [file.h](#).

Referenced by [consoleinit\(\)](#), and [writei\(\)](#).

The documentation for this struct was generated from the following file:

- [file.h](#)

4.7 dinode

```
#include <fs.h>
```

Public Attributes

- [uint](#) [addr](#)s [NDIRECT+1]
- short [major](#)
- short [minor](#)
- short [nlink](#)
- [uint](#) [size](#)
- short [type](#)

4.7.1 Detailed Description

Definition at line 29 of file [fs.h](#).

4.7.2 Member Data Documentation

4.7.2.1 `addr`

```
uint dinode::addr[NDIRECT+1]
```

Definition at line 35 of file [fs.h](#).

Referenced by [iappend\(\)](#), [ilock\(\)](#), and [iupdate\(\)](#).

4.7.2.2 `major`

```
short dinode::major
```

Definition at line 31 of file [fs.h](#).

Referenced by [ilock\(\)](#), and [iupdate\(\)](#).

4.7.2.3 `minor`

```
short dinode::minor
```

Definition at line 32 of file [fs.h](#).

Referenced by [ilock\(\)](#), and [iupdate\(\)](#).

4.7.2.4 `nlink`

```
short dinode::nlink
```

Definition at line 33 of file [fs.h](#).

Referenced by [ialloc\(\)](#), [ilock\(\)](#), and [iupdate\(\)](#).

4.7.2.5 `size`

```
uint dinode::size
```

Definition at line 34 of file [fs.h](#).

Referenced by [ialloc\(\)](#), [iappend\(\)](#), [ilock\(\)](#), [iupdate\(\)](#), and [main\(\)](#).

4.7.2.6 type

```
short dinode::type
```

Definition at line 30 of file [fs.h](#).

Referenced by [ialloc\(\)](#), [ilock\(\)](#), and [iupdate\(\)](#).

The documentation for this struct was generated from the following file:

- [fs.h](#)

4.8 dirent

```
#include <fs.h>
```

Public Attributes

- [ushort](#) [inum](#)
- [char](#) [name](#) [[DIRSIZ](#)]

4.8.1 Detailed Description

Definition at line 53 of file [fs.h](#).

4.8.2 Member Data Documentation

4.8.2.1 inum

```
ushort dirent::inum
```

Definition at line 54 of file [fs.h](#).

Referenced by [dirlink\(\)](#), [dirlookup\(\)](#), [ls\(\)](#), and [main\(\)](#).

4.8.2.2 name

```
char dirent::name[DIRSIZ]
```

Definition at line 55 of file [fs.h](#).

Referenced by [dirlink\(\)](#), [ls\(\)](#), and [main\(\)](#).

The documentation for this struct was generated from the following file:

- [fs.h](#)

4.9 elfhdr

```
#include <elf.h>
```

Public Attributes

- [ushort ehsize](#)
- [uchar elf](#) [12]
- [uint entry](#)
- [uint flags](#)
- [ushort machine](#)
- [uint magic](#)
- [ushort phentsize](#)
- [ushort phnum](#)
- [uint phoff](#)
- [ushort shentsize](#)
- [ushort shnum](#)
- [uint shoff](#)
- [ushort shstrndx](#)
- [ushort type](#)
- [uint version](#)

4.9.1 Detailed Description

Definition at line 6 of file [elf.h](#).

4.9.2 Member Data Documentation

4.9.2.1 ehsize

```
ushort elfhdr::ehsize
```

Definition at line 16 of file [elf.h](#).

4.9.2.2 elf

```
uchar elfhdr::elf[12]
```

Definition at line 8 of file [elf.h](#).

Referenced by [bootmain\(\)](#), and [exec\(\)](#).

4.9.2.3 entry

```
uint elfhdr::entry
```

Definition at line 12 of file [elf.h](#).

Referenced by [bootmain\(\)](#), and [exec\(\)](#).

4.9.2.4 flags

```
uint elfhdr::flags
```

Definition at line 15 of file [elf.h](#).

4.9.2.5 machine

```
ushort elfhdr::machine
```

Definition at line 10 of file [elf.h](#).

4.9.2.6 magic

```
uint elfhdr::magic
```

Definition at line 7 of file [elf.h](#).

Referenced by [exec\(\)](#).

4.9.2.7 phentsize

```
ushort elfhdr::phentsize
```

Definition at line 17 of file [elf.h](#).

4.9.2.8 phnum

```
ushort elfhdr::phnum
```

Definition at line 18 of file [elf.h](#).

Referenced by [bootmain\(\)](#), and [exec\(\)](#).

4.9.2.9 phoff

`uint elfhdr::phoff`

Definition at line 13 of file [elf.h](#).

Referenced by [bootmain\(\)](#), and [exec\(\)](#).

4.9.2.10 shentsize

`ushort elfhdr::shentsize`

Definition at line 19 of file [elf.h](#).

4.9.2.11 shnum

`ushort elfhdr::shnum`

Definition at line 20 of file [elf.h](#).

4.9.2.12 shoff

`uint elfhdr::shoff`

Definition at line 14 of file [elf.h](#).

4.9.2.13 shstrndx

`ushort elfhdr::shstrndx`

Definition at line 21 of file [elf.h](#).

4.9.2.14 type

`ushort elfhdr::type`

Definition at line 9 of file [elf.h](#).

4.9.2.15 version

```
uint elfhdr::version
```

Definition at line 11 of file [elf.h](#).

The documentation for this struct was generated from the following file:

- [elf.h](#)

4.10 execcmd

Public Attributes

- char * [argv](#) [[MAXARGS](#)]
- char * [eargv](#) [[MAXARGS](#)]
- int [type](#)

4.10.1 Detailed Description

Definition at line 20 of file [sh.c](#).

4.10.2 Member Data Documentation

4.10.2.1 argv

```
char* execcmd::argv[MAXARGS]
```

Definition at line 22 of file [sh.c](#).

Referenced by [nulterminate\(\)](#), and [runcmd\(\)](#).

4.10.2.2 eargv

```
char* execcmd::eargv[MAXARGS]
```

Definition at line 23 of file [sh.c](#).

Referenced by [nulterminate\(\)](#).

4.10.2.3 type

```
int execcmd::type
```

Definition at line 21 of file [sh.c](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.11 file

```
#include <file.h>
```

Public Types

- enum { [FD_NONE](#) , [FD_PIPE](#) , [FD_INODE](#) }

Public Attributes

- struct [inode](#) * [ip](#)
- [uint](#) [off](#)
- struct [pipe](#) * [pipe](#)
- char [readable](#)
- int [ref](#)
- enum file:: { ... } [type](#)
- char [writable](#)

4.11.1 Detailed Description

Definition at line 1 of file [file.h](#).

4.11.2 Member Enumeration Documentation

4.11.2.1 anonymous enum

```
anonymous enum
```

Enumerator

FD_NONE	
FD_PIPE	
FD_INODE	

Definition at line 2 of file [file.h](#).

```
00002 { FD_NONE, FD_PIPE, FD_INODE } type;
```

4.11.3 Member Data Documentation

4.11.3.1 ip

```
struct inode* file::ip
```

Definition at line 7 of file [file.h](#).

Referenced by [fileclose\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), and [sys_open\(\)](#).

4.11.3.2 off

```
uint file::off
```

Definition at line 8 of file [file.h](#).

Referenced by [fileread\(\)](#), [filewrite\(\)](#), and [sys_open\(\)](#).

4.11.3.3 pipe

```
struct pipe* file::pipe
```

Definition at line 6 of file [file.h](#).

Referenced by [fileclose\(\)](#), [fileread\(\)](#), and [filewrite\(\)](#).

4.11.3.4 readable

```
char file::readable
```

Definition at line 4 of file [file.h](#).

Referenced by [fileread\(\)](#), and [sys_open\(\)](#).

4.11.3.5 ref

```
int file::ref
```

Definition at line 3 of file [file.h](#).

Referenced by [filealloc\(\)](#), [fileclose\(\)](#), and [filedup\(\)](#).

4.11.3.6

```
enum { ... } file::type
```

Referenced by [fileclose\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), and [sys_open\(\)](#).

4.11.3.7 writable

```
char file::writable
```

Definition at line 5 of file [file.h](#).

Referenced by [fileclose\(\)](#), [filewrite\(\)](#), and [sys_open\(\)](#).

The documentation for this struct was generated from the following file:

- [file.h](#)

4.12 gatedesc

```
#include <mmu.h>
```

Public Attributes

- [uint args](#): 5
- [uint cs](#): 16
- [uint dpl](#): 2
- [uint off_15_0](#): 16
- [uint off_31_16](#): 16
- [uint p](#): 1
- [uint rsv1](#): 3
- [uint s](#): 1
- [uint type](#): 4

4.12.1 Detailed Description

Definition at line 148 of file [mmu.h](#).

4.12.2 Member Data Documentation

4.12.2.1 args

`uint gatedesc::args`

Definition at line 151 of file [mmu.h](#).

4.12.2.2 cs

`uint gatedesc::cs`

Definition at line 150 of file [mmu.h](#).

4.12.2.3 dpl

`uint gatedesc::dpl`

Definition at line 155 of file [mmu.h](#).

4.12.2.4 off_15_0

`uint gatedesc::off_15_0`

Definition at line 149 of file [mmu.h](#).

4.12.2.5 off_31_16

`uint gatedesc::off_31_16`

Definition at line 157 of file [mmu.h](#).

4.12.2.6 p

`uint gatedesc::p`

Definition at line 156 of file [mmu.h](#).

Referenced by [lidt\(\)](#).

4.12.2.7 rsv1

`uint gatedesc::rsv1`

Definition at line 152 of file [mmu.h](#).

4.12.2.8 s

`uint gatedesc::s`

Definition at line 154 of file [mmu.h](#).

4.12.2.9 type

`uint gatedesc::type`

Definition at line 153 of file [mmu.h](#).

The documentation for this struct was generated from the following file:

- [mmu.h](#)

4.13 header

Public Attributes

- struct {
 union [header](#) * [ptr](#)
 [uint](#) [size](#)
} [s](#)
- [Align](#) [x](#)

4.13.1 Detailed Description

Definition at line 11 of file [umalloc.c](#).

4.13.2 Member Data Documentation

4.13.2.1 ptr

```
union header* header::ptr
```

Definition at line 13 of file [umalloc.c](#).

Referenced by [free\(\)](#), and [malloc\(\)](#).

4.13.2.2

```
struct { ... } header::s
```

Referenced by [free\(\)](#), [malloc\(\)](#), and [morecore\(\)](#).

4.13.2.3 size

```
uint header::size
```

Definition at line 14 of file [umalloc.c](#).

Referenced by [free\(\)](#), [malloc\(\)](#), and [morecore\(\)](#).

4.13.2.4 x

```
Align header::x
```

Definition at line 16 of file [umalloc.c](#).

The documentation for this union was generated from the following file:

- [umalloc.c](#)

4.14 inode

```
#include <file.h>
```

Public Attributes

- [uint](#) [addr](#)s [[NDIRECT](#)+1]
- [uint](#) [dev](#)
- [uint](#) [inum](#)
- [struct](#) [sleeplock](#) [lock](#)
- [short](#) [major](#)
- [short](#) [minor](#)
- [short](#) [nlink](#)
- [int](#) [ref](#)
- [uint](#) [size](#)
- [short](#) [type](#)
- [int](#) [valid](#)

4.14.1 Detailed Description

Definition at line 13 of file [file.h](#).

4.14.2 Member Data Documentation

4.14.2.1 [addr](#)s

```
uint inode::addr[s]\[NDIRECT+1\]
```

Definition at line 25 of file [file.h](#).

Referenced by [bmap\(\)](#), [ilock\(\)](#), [itrunc\(\)](#), and [iupdate\(\)](#).

4.14.2.2 [dev](#)

```
uint inode::dev
```

Definition at line 14 of file [file.h](#).

Referenced by [bmap\(\)](#), [create\(\)](#), [dirlookup\(\)](#), [iget\(\)](#), [iinit\(\)](#), [ilock\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [readi\(\)](#), [stati\(\)](#), [sys_link\(\)](#), and [writei\(\)](#).

4.14.2.3 inum

```
uint inode::inum
```

Definition at line 15 of file [file.h](#).

Referenced by [create\(\)](#), [dirlink\(\)](#), [dirlookup\(\)](#), [ialloc\(\)](#), [iget\(\)](#), [ilock\(\)](#), [iupdate\(\)](#), [stati\(\)](#), and [sys_link\(\)](#).

4.14.2.4 lock

```
struct sleeplock inode::lock
```

Definition at line 17 of file [file.h](#).

Referenced by [ilock\(\)](#), [iput\(\)](#), and [iunlock\(\)](#).

4.14.2.5 major

```
short inode::major
```

Definition at line 21 of file [file.h](#).

Referenced by [create\(\)](#), [ilock\(\)](#), [iupdate\(\)](#), [readi\(\)](#), [sys_mknod\(\)](#), and [writei\(\)](#).

4.14.2.6 minor

```
short inode::minor
```

Definition at line 22 of file [file.h](#).

Referenced by [create\(\)](#), [ilock\(\)](#), [iupdate\(\)](#), and [sys_mknod\(\)](#).

4.14.2.7 nlink

```
short inode::nlink
```

Definition at line 23 of file [file.h](#).

Referenced by [create\(\)](#), [ilock\(\)](#), [iput\(\)](#), [iupdate\(\)](#), [stati\(\)](#), [sys_link\(\)](#), and [sys_unlink\(\)](#).

4.14.2.8 ref

```
int inode::ref
```

Definition at line 16 of file [file.h](#).

Referenced by [idup\(\)](#), [iget\(\)](#), [ilock\(\)](#), [iput\(\)](#), and [iunlock\(\)](#).

4.14.2.9 size

```
uint inode::size
```

Definition at line 24 of file [file.h](#).

Referenced by [dirlink\(\)](#), [dirlookup\(\)](#), [ilock\(\)](#), [isdirempty\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [readi\(\)](#), [stati\(\)](#), and [writei\(\)](#).

4.14.2.10 type

```
short inode::type
```

Definition at line 20 of file [file.h](#).

Referenced by [create\(\)](#), [dirlookup\(\)](#), [ilock\(\)](#), [iput\(\)](#), [iupdate\(\)](#), [namex\(\)](#), [readi\(\)](#), [stati\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_open\(\)](#), [sys_unlink\(\)](#), and [writei\(\)](#).

4.14.2.11 valid

```
int inode::valid
```

Definition at line 18 of file [file.h](#).

Referenced by [iget\(\)](#), [ilock\(\)](#), and [iput\(\)](#).

The documentation for this struct was generated from the following file:

- [file.h](#)

4.15 ioapic

Public Attributes

- [uint data](#)
- [uint pad](#) [3]
- [uint reg](#)

4.15.1 Detailed Description

Definition at line 28 of file [ioapic.c](#).

4.15.2 Member Data Documentation

4.15.2.1 data

```
uint ioapic::data
```

Definition at line 31 of file [ioapic.c](#).

Referenced by [ioapicread\(\)](#), and [ioapicwrite\(\)](#).

4.15.2.2 pad

```
uint ioapic::pad[3]
```

Definition at line 30 of file [ioapic.c](#).

4.15.2.3 reg

```
uint ioapic::reg
```

Definition at line 29 of file [ioapic.c](#).

Referenced by [ioapicread\(\)](#), and [ioapicwrite\(\)](#).

The documentation for this struct was generated from the following file:

- [ioapic.c](#)

4.16 kmap

Public Attributes

- int [perm](#)
- uint [phys_end](#)
- uint [phys_start](#)
- void * [virt](#)

4.16.1 Detailed Description

Definition at line 105 of file [vm.c](#).

4.16.2 Member Data Documentation

4.16.2.1 perm

```
int kmap::perm
```

Definition at line 109 of file [vm.c](#).

Referenced by [setupkvm\(\)](#).

4.16.2.2 phys_end

```
uint kmap::phys_end
```

Definition at line 108 of file [vm.c](#).

Referenced by [setupkvm\(\)](#).

4.16.2.3 phys_start

```
uint kmap::phys_start
```

Definition at line 107 of file [vm.c](#).

Referenced by [setupkvm\(\)](#).

4.16.2.4 virt

```
void* kmap::virt
```

Definition at line 106 of file [vm.c](#).

Referenced by [setupkvm\(\)](#).

The documentation for this struct was generated from the following file:

- [vm.c](#)

4.17 listcmd

Public Attributes

- struct `cmd` * `left`
- struct `cmd` * `right`
- int `type`

4.17.1 Detailed Description

Definition at line 41 of file [sh.c](#).

4.17.2 Member Data Documentation

4.17.2.1 left

```
struct cmd* listcmd::left
```

Definition at line 43 of file [sh.c](#).

Referenced by [listcmd\(\)](#), [nulterminate\(\)](#), and [runcmd\(\)](#).

4.17.2.2 right

```
struct cmd* listcmd::right
```

Definition at line 44 of file [sh.c](#).

Referenced by [listcmd\(\)](#), [nulterminate\(\)](#), and [runcmd\(\)](#).

4.17.2.3 type

```
int listcmd::type
```

Definition at line 42 of file [sh.c](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.18 log

Public Attributes

- int [committing](#)
- int [dev](#)
- struct [logheader](#) lh
- struct [spinlock](#) lock
- int [outstanding](#)
- int [size](#)
- int [start](#)

4.18.1 Detailed Description

Definition at line [39](#) of file [log.c](#).

4.18.2 Member Data Documentation

4.18.2.1 committing

```
int log::committing
```

Definition at line [44](#) of file [log.c](#).

Referenced by [begin_op\(\)](#), and [end_op\(\)](#).

4.18.2.2 dev

```
int log::dev
```

Definition at line [45](#) of file [log.c](#).

Referenced by [initlog\(\)](#), [install_trans\(\)](#), [read_head\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

4.18.2.3 lh

```
struct logheader log::lh
```

Definition at line [46](#) of file [log.c](#).

Referenced by [begin_op\(\)](#), [commit\(\)](#), [install_trans\(\)](#), [log_write\(\)](#), [read_head\(\)](#), [recover_from_log\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

4.18.2.4 lock

```
struct spinlock log::lock
```

Definition at line 40 of file [log.c](#).

Referenced by [begin_op\(\)](#), [end_op\(\)](#), [initlog\(\)](#), and [log_write\(\)](#).

4.18.2.5 outstanding

```
int log::outstanding
```

Definition at line 43 of file [log.c](#).

Referenced by [begin_op\(\)](#), [end_op\(\)](#), and [log_write\(\)](#).

4.18.2.6 size

```
int log::size
```

Definition at line 42 of file [log.c](#).

Referenced by [initlog\(\)](#), and [log_write\(\)](#).

4.18.2.7 start

```
int log::start
```

Definition at line 41 of file [log.c](#).

Referenced by [initlog\(\)](#), [install_trans\(\)](#), [read_head\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

The documentation for this struct was generated from the following file:

- [log.c](#)

4.19 logheader

Public Attributes

- int [block](#) [[LOGSIZE](#)]
- int [n](#)

4.19.1 Detailed Description

Definition at line 34 of file [log.c](#).

4.19.2 Member Data Documentation

4.19.2.1 block

```
int logheader::block[LOGSIZE]
```

Definition at line 36 of file [log.c](#).

Referenced by [install_trans\(\)](#), [log_write\(\)](#), [read_head\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

4.19.2.2 n

```
int logheader::n
```

Definition at line 35 of file [log.c](#).

Referenced by [begin_op\(\)](#), [commit\(\)](#), [install_trans\(\)](#), [log_write\(\)](#), [read_head\(\)](#), [recover_from_log\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

The documentation for this struct was generated from the following file:

- [log.c](#)

4.20 mp

```
#include <mp.h>
```

Public Attributes

- [uchar checksum](#)
- [uchar imcrp](#)
- [uchar length](#)
- [void * physaddr](#)
- [uchar reserved](#) [3]
- [uchar signature](#) [4]
- [uchar specrev](#)
- [uchar type](#)

4.20.1 Detailed Description

Definition at line 3 of file [mp.h](#).

4.20.2 Member Data Documentation

4.20.2.1 checksum

```
uchar mp::checksum
```

Definition at line 8 of file [mp.h](#).

4.20.2.2 imcrp

```
uchar mp::imcrp
```

Definition at line 10 of file [mp.h](#).

Referenced by [mpinit\(\)](#).

4.20.2.3 length

```
uchar mp::length
```

Definition at line 6 of file [mp.h](#).

4.20.2.4 physaddr

```
void* mp::physaddr
```

Definition at line 5 of file [mp.h](#).

Referenced by [mpconfig\(\)](#).

4.20.2.5 reserved

```
uchar mp::reserved[3]
```

Definition at line 11 of file [mp.h](#).

4.20.2.6 signature

```
uchar mp::signature[4]
```

Definition at line 4 of file [mp.h](#).

4.20.2.7 specrev

```
uchar mp::specrev
```

Definition at line 7 of file [mp.h](#).

4.20.2.8 type

```
uchar mp::type
```

Definition at line 9 of file [mp.h](#).

The documentation for this struct was generated from the following file:

- [mp.h](#)

4.21 mpconf

```
#include <mp.h>
```

Public Attributes

- [uchar checksum](#)
- [ushort entry](#)
- [uint * lapicaddr](#)
- [ushort length](#)
- [ushort oemlength](#)
- [uint * oemtable](#)
- [uchar product](#) [20]
- [uchar reserved](#)
- [uchar signature](#) [4]
- [uchar version](#)
- [uchar xchecksum](#)
- [ushort xlength](#)

4.21.1 Detailed Description

Definition at line 14 of file [mp.h](#).

4.21.2 Member Data Documentation

4.21.2.1 checksum

```
uchar mpconf::checksum
```

Definition at line 18 of file [mp.h](#).

4.21.2.2 entry

```
ushort mpconf::entry
```

Definition at line 22 of file [mp.h](#).

4.21.2.3 lapicaddr

```
uint* mpconf::lapicaddr
```

Definition at line 23 of file [mp.h](#).

Referenced by [mpinit\(\)](#).

4.21.2.4 length

```
ushort mpconf::length
```

Definition at line 16 of file [mp.h](#).

Referenced by [mpconfig\(\)](#), and [mpinit\(\)](#).

4.21.2.5 oemlength

```
ushort mpconf::oemlength
```

Definition at line 21 of file [mp.h](#).

4.21.2.6 oemtable

```
uint* mpconf::oemtable
```

Definition at line 20 of file [mp.h](#).

4.21.2.7 product

```
uchar mpconf::product[20]
```

Definition at line 19 of file [mp.h](#).

4.21.2.8 reserved

```
uchar mpconf::reserved
```

Definition at line 26 of file [mp.h](#).

4.21.2.9 signature

```
uchar mpconf::signature[4]
```

Definition at line 15 of file [mp.h](#).

4.21.2.10 version

```
uchar mpconf::version
```

Definition at line 17 of file [mp.h](#).

Referenced by [mpconfig\(\)](#).

4.21.2.11 xchecksum

`uchar` `mpconf::xchecksum`

Definition at line 25 of file [mp.h](#).

4.21.2.12 xlength

`ushort` `mpconf::xlength`

Definition at line 24 of file [mp.h](#).

The documentation for this struct was generated from the following file:

- [mp.h](#)

4.22 mpioapic

```
#include <mp.h>
```

Public Attributes

- `uint * addr`
- `uchar apicno`
- `uchar flags`
- `uchar type`
- `uchar version`

4.22.1 Detailed Description

Definition at line 40 of file [mp.h](#).

4.22.2 Member Data Documentation

4.22.2.1 addr

`uint*` `mpioapic::addr`

Definition at line 45 of file [mp.h](#).

4.22.2.2 apicno

```
uchar mpioapic::apicno
```

Definition at line 42 of file [mp.h](#).

4.22.2.3 flags

```
uchar mpioapic::flags
```

Definition at line 44 of file [mp.h](#).

4.22.2.4 type

```
uchar mpioapic::type
```

Definition at line 41 of file [mp.h](#).

4.22.2.5 version

```
uchar mpioapic::version
```

Definition at line 43 of file [mp.h](#).

The documentation for this struct was generated from the following file:

- [mp.h](#)

4.23 mpproc

```
#include <mp.h>
```

Public Attributes

- [uchar apicid](#)
- [uint feature](#)
- [uchar flags](#)
- [uchar reserved](#) [8]
- [uchar signature](#) [4]
- [uchar type](#)
- [uchar version](#)

4.23.1 Detailed Description

Definition at line 29 of file [mp.h](#).

4.23.2 Member Data Documentation

4.23.2.1 apicid

```
uchar mpproc::apicid
```

Definition at line 31 of file [mp.h](#).

4.23.2.2 feature

```
uint mpproc::feature
```

Definition at line 36 of file [mp.h](#).

4.23.2.3 flags

```
uchar mpproc::flags
```

Definition at line 33 of file [mp.h](#).

4.23.2.4 reserved

```
uchar mpproc::reserved[8]
```

Definition at line 37 of file [mp.h](#).

4.23.2.5 signature

```
uchar mpproc::signature[4]
```

Definition at line 35 of file [mp.h](#).

4.23.2.6 type

`uchar` `mpproc::type`

Definition at line 30 of file [mp.h](#).

4.23.2.7 version

`uchar` `mpproc::version`

Definition at line 32 of file [mp.h](#).

The documentation for this struct was generated from the following file:

- [mp.h](#)

4.24 pipe

Public Attributes

- `char` [data](#) [[PIPESIZE](#)]
- `struct` [spinlock](#) [lock](#)
- `uint` [nread](#)
- `uint` [nwrite](#)
- `int` [readopen](#)
- `int` [writeopen](#)

4.24.1 Detailed Description

Definition at line 13 of file [pipe.c](#).

4.24.2 Member Data Documentation

4.24.2.1 data

`char` `pipe::data` [[PIPESIZE](#)]

Definition at line 15 of file [pipe.c](#).

Referenced by [piperead\(\)](#), and [pipewrite\(\)](#).

4.24.2.2 lock

```
struct spinlock pipe::lock
```

Definition at line 14 of file [pipe.c](#).

Referenced by [pipealloc\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), and [pipewrite\(\)](#).

4.24.2.3 nread

```
uint pipe::nread
```

Definition at line 16 of file [pipe.c](#).

Referenced by [pipealloc\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), and [pipewrite\(\)](#).

4.24.2.4 nwrite

```
uint pipe::nwrite
```

Definition at line 17 of file [pipe.c](#).

Referenced by [pipealloc\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), and [pipewrite\(\)](#).

4.24.2.5 readopen

```
int pipe::readopen
```

Definition at line 18 of file [pipe.c](#).

Referenced by [pipealloc\(\)](#), [pipeclose\(\)](#), and [pipewrite\(\)](#).

4.24.2.6 writeopen

```
int pipe::writeopen
```

Definition at line 19 of file [pipe.c](#).

Referenced by [pipealloc\(\)](#), [pipeclose\(\)](#), and [piperead\(\)](#).

The documentation for this struct was generated from the following file:

- [pipe.c](#)

4.25 pipecmd

Public Attributes

- struct `cmd` * `left`
- struct `cmd` * `right`
- int `type`

4.25.1 Detailed Description

Definition at line 35 of file [sh.c](#).

4.25.2 Member Data Documentation

4.25.2.1 left

```
struct cmd* pipecmd::left
```

Definition at line 37 of file [sh.c](#).

Referenced by [nulterminate\(\)](#), [pipecmd\(\)](#), and [runcmd\(\)](#).

4.25.2.2 right

```
struct cmd* pipecmd::right
```

Definition at line 38 of file [sh.c](#).

Referenced by [nulterminate\(\)](#), [pipecmd\(\)](#), and [runcmd\(\)](#).

4.25.2.3 type

```
int pipecmd::type
```

Definition at line 36 of file [sh.c](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.26 proc

```
#include <proc.h>
```

Public Attributes

- void * [chan](#)
- struct [context](#) * [context](#)
- struct [inode](#) * [cwd](#)
- int [killed](#)
- char * [kstack](#)
- char [name](#) [16]
- struct [file](#) * [ofile](#) [NOFILE]
- struct [proc](#) * [parent](#)
- [pde_t](#) * [pgdir](#)
- int [pid](#)
- int [priority](#)
- enum [procstate](#) [state](#)
- [uint](#) [sz](#)
- struct [trapframe](#) * [tf](#)

4.26.1 Detailed Description

Definition at line 38 of file [proc.h](#).

4.26.2 Member Data Documentation

4.26.2.1 chan

```
void* proc::chan
```

Definition at line 47 of file [proc.h](#).

Referenced by [sleep\(\)](#), [wakeup\(\)](#), and [wakeup1\(\)](#).

4.26.2.2 context

```
struct context* proc::context
```

Definition at line 46 of file [proc.h](#).

Referenced by [allocproc\(\)](#), [procdump\(\)](#), [sched\(\)](#), and [scheduler\(\)](#).

4.26.2.3 `cwd`

```
struct inode* proc::cwd
```

Definition at line 50 of file `proc.h`.

Referenced by `exit()`, `fork()`, `sys_chdir()`, and `userinit()`.

4.26.2.4 `killed`

```
int proc::killed
```

Definition at line 48 of file `proc.h`.

Referenced by `kill()`, `sys_sleep()`, `trap()`, and `wait()`.

4.26.2.5 `kstack`

```
char* proc::kstack
```

Definition at line 41 of file `proc.h`.

Referenced by `allocproc()`, `fork()`, `switchvm()`, and `wait()`.

4.26.2.6 `name`

```
char proc::name[16]
```

Definition at line 51 of file `proc.h`.

Referenced by `cps()`, `exec()`, `fork()`, `procdump()`, `sys_getprocs()`, `syscall()`, `userinit()`, and `wait()`.

4.26.2.7 `ofile`

```
struct file* proc::ofile[NOFILE]
```

Definition at line 49 of file `proc.h`.

Referenced by `exit()`, `fdalloc()`, `fork()`, `sys_close()`, and `sys_pipe()`.

4.26.2.8 parent

```
struct proc* proc::parent
```

Definition at line 44 of file [proc.h](#).

Referenced by [exit\(\)](#), [fork\(\)](#), [sys_getprocs\(\)](#), and [wait\(\)](#).

4.26.2.9 pgdir

```
pde_t* proc::pgdir
```

Definition at line 40 of file [proc.h](#).

Referenced by [exec\(\)](#), [fork\(\)](#), [growproc\(\)](#), [switchvm\(\)](#), [userinit\(\)](#), and [wait\(\)](#).

4.26.2.10 pid

```
int proc::pid
```

Definition at line 43 of file [proc.h](#).

Referenced by [acquiresleep\(\)](#), [allocproc\(\)](#), [chpr\(\)](#), [cps\(\)](#), [fork\(\)](#), [holdingsleep\(\)](#), [kill\(\)](#), [procdump\(\)](#), [sys_chpr\(\)](#), [sys_getpid\(\)](#), [sys_getprocs\(\)](#), [sys_kill\(\)](#), [syscall\(\)](#), and [wait\(\)](#).

4.26.2.11 priority

```
int proc::priority
```

Definition at line 52 of file [proc.h](#).

Referenced by [chpr\(\)](#), and [cps\(\)](#).

4.26.2.12 state

```
enum procstate proc::state
```

Definition at line 42 of file [proc.h](#).

Referenced by [allocproc\(\)](#), [cps\(\)](#), [exit\(\)](#), [fork\(\)](#), [kill\(\)](#), [procdump\(\)](#), [sched\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [sys_getprocs\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup1\(\)](#), and [yield\(\)](#).

4.26.2.13 sz

```
uint proc::sz
```

Definition at line 39 of file [proc.h](#).

Referenced by [argptr\(\)](#), [exec\(\)](#), [fetchint\(\)](#), [fetchstr\(\)](#), [fork\(\)](#), [growproc\(\)](#), [sys_sbrk\(\)](#), and [userinit\(\)](#).

4.26.2.14 tf

```
struct trapframe* proc::tf
```

Definition at line 45 of file [proc.h](#).

Referenced by [allocproc\(\)](#), [argint\(\)](#), [exec\(\)](#), [fork\(\)](#), [syscall\(\)](#), [trap\(\)](#), and [userinit\(\)](#).

The documentation for this struct was generated from the following file:

- [proc.h](#)

4.27 proghdr

```
#include <elf.h>
```

Public Attributes

- [uint align](#)
- [uint filesz](#)
- [uint flags](#)
- [uint memsz](#)
- [uint off](#)
- [uint paddr](#)
- [uint type](#)
- [uint vaddr](#)

4.27.1 Detailed Description

Definition at line 25 of file [elf.h](#).

4.27.2 Member Data Documentation

4.27.2.1 align

`uint proghdr::align`

Definition at line 33 of file [elf.h](#).

4.27.2.2 filesz

`uint proghdr::filesz`

Definition at line 30 of file [elf.h](#).

Referenced by [bootmain\(\)](#).

4.27.2.3 flags

`uint proghdr::flags`

Definition at line 32 of file [elf.h](#).

4.27.2.4 memsz

`uint proghdr::memsz`

Definition at line 31 of file [elf.h](#).

Referenced by [bootmain\(\)](#).

4.27.2.5 off

`uint proghdr::off`

Definition at line 27 of file [elf.h](#).

Referenced by [bootmain\(\)](#).

4.27.2.6 paddr

`uint proghdr::paddr`

Definition at line 29 of file [elf.h](#).

Referenced by [bootmain\(\)](#).

4.27.2.7 type

`uint proghdr::type`

Definition at line 26 of file [elf.h](#).

4.27.2.8 vaddr

`uint proghdr::vaddr`

Definition at line 28 of file [elf.h](#).

The documentation for this struct was generated from the following file:

- [elf.h](#)

4.28 redircmd

Public Attributes

- struct [cmd](#) * [cmd](#)
- char * [efile](#)
- int [fd](#)
- char * [file](#)
- int [mode](#)
- int [type](#)

4.28.1 Detailed Description

Definition at line 26 of file [sh.c](#).

4.28.2 Member Data Documentation

4.28.2.1 cmd

```
struct cmd* redircmd::cmd
```

Definition at line 28 of file [sh.c](#).

Referenced by [nulterminate\(\)](#), [redircmd\(\)](#), and [runcmd\(\)](#).

4.28.2.2 efile

```
char* redircmd::efile
```

Definition at line 30 of file [sh.c](#).

Referenced by [nulterminate\(\)](#), and [redircmd\(\)](#).

4.28.2.3 fd

```
int redircmd::fd
```

Definition at line 32 of file [sh.c](#).

Referenced by [redircmd\(\)](#), and [runcmd\(\)](#).

4.28.2.4 file

```
char* redircmd::file
```

Definition at line 29 of file [sh.c](#).

Referenced by [redircmd\(\)](#), and [runcmd\(\)](#).

4.28.2.5 mode

```
int redircmd::mode
```

Definition at line 31 of file [sh.c](#).

Referenced by [redircmd\(\)](#), and [runcmd\(\)](#).

4.28.2.6 type

```
int redircmd::type
```

Definition at line 27 of file [sh.c](#).

The documentation for this struct was generated from the following file:

- [sh.c](#)

4.29 rtcddate

```
#include <date.h>
```

Public Attributes

- [uint day](#)
- [uint hour](#)
- [uint minute](#)
- [uint month](#)
- [uint second](#)
- [uint year](#)

4.29.1 Detailed Description

Definition at line 1 of file [date.h](#).

4.29.2 Member Data Documentation

4.29.2.1 day

```
uint rtcddate::day
```

Definition at line 5 of file [date.h](#).

Referenced by [cmostime\(\)](#).

4.29.2.2 hour

```
uint rtdatetime::hour
```

Definition at line 4 of file [date.h](#).

Referenced by [cmstime\(\)](#).

4.29.2.3 minute

```
uint rtdatetime::minute
```

Definition at line 3 of file [date.h](#).

Referenced by [cmstime\(\)](#).

4.29.2.4 month

```
uint rtdatetime::month
```

Definition at line 6 of file [date.h](#).

Referenced by [cmstime\(\)](#).

4.29.2.5 second

```
uint rtdatetime::second
```

Definition at line 2 of file [date.h](#).

Referenced by [cmstime\(\)](#).

4.29.2.6 year

```
uint rtdatetime::year
```

Definition at line 7 of file [date.h](#).

Referenced by [cmstime\(\)](#).

The documentation for this struct was generated from the following file:

- [date.h](#)

4.30 run

Public Attributes

- struct [run](#) * [next](#)

4.30.1 Detailed Description

Definition at line 16 of file [kalloc.c](#).

4.30.2 Member Data Documentation

4.30.2.1 next

```
struct run* run::next
```

Definition at line 17 of file [kalloc.c](#).

The documentation for this struct was generated from the following file:

- [kalloc.c](#)

4.31 segdesc

```
#include <mmu.h>
```

Public Attributes

- [uint](#) [avl](#): 1
- [uint](#) [base_15_0](#): 16
- [uint](#) [base_23_16](#): 8
- [uint](#) [base_31_24](#): 8
- [uint](#) [db](#): 1
- [uint](#) [dpl](#): 2
- [uint](#) [g](#): 1
- [uint](#) [lim_15_0](#): 16
- [uint](#) [lim_19_16](#): 4
- [uint](#) [p](#): 1
- [uint](#) [rsv1](#): 1
- [uint](#) [s](#): 1
- [uint](#) [type](#): 4

4.31.1 Detailed Description

Definition at line 26 of file [mmu.h](#).

4.31.2 Member Data Documentation

4.31.2.1 avl

```
uint segdesc::avl
```

Definition at line 35 of file [mmu.h](#).

4.31.2.2 base_15_0

```
uint segdesc::base_15_0
```

Definition at line 28 of file [mmu.h](#).

4.31.2.3 base_23_16

```
uint segdesc::base_23_16
```

Definition at line 29 of file [mmu.h](#).

4.31.2.4 base_31_24

```
uint segdesc::base_31_24
```

Definition at line 39 of file [mmu.h](#).

4.31.2.5 db

```
uint segdesc::db
```

Definition at line 37 of file [mmu.h](#).

4.31.2.6 dpl

```
uint segdesc::dpl
```

Definition at line 32 of file [mmu.h](#).

4.31.2.7 g

```
uint segdesc::g
```

Definition at line 38 of file [mmu.h](#).

4.31.2.8 lim_15_0

```
uint segdesc::lim_15_0
```

Definition at line 27 of file [mmu.h](#).

4.31.2.9 lim_19_16

```
uint segdesc::lim_19_16
```

Definition at line 34 of file [mmu.h](#).

4.31.2.10 p

```
uint segdesc::p
```

Definition at line 33 of file [mmu.h](#).

Referenced by [lgdt\(\)](#).

4.31.2.11 rsv1

```
uint segdesc::rsv1
```

Definition at line 36 of file [mmu.h](#).

4.31.2.12 s

```
uint segdesc::s
```

Definition at line 31 of file [mmu.h](#).

Referenced by [switchvm\(\)](#).

4.31.2.13 type

```
uint segdesc::type
```

Definition at line 30 of file [mmu.h](#).

The documentation for this struct was generated from the following file:

- [mmu.h](#)

4.32 sleeplock

```
#include <sleeplock.h>
```

Public Attributes

- struct [spinlock lk](#)
- [uint locked](#)
- char * [name](#)
- int [pid](#)

4.32.1 Detailed Description

Definition at line 2 of file [sleeplock.h](#).

4.32.2 Member Data Documentation

4.32.2.1 lk

```
struct spinlock sleeplock::lk
```

Definition at line 4 of file [sleeplock.h](#).

Referenced by [acquiresleep\(\)](#), [holdingsleep\(\)](#), [initsleeplock\(\)](#), and [releasesleep\(\)](#).

4.32.2.2 locked

```
uint sleeplock::locked
```

Definition at line 3 of file [sleeplock.h](#).

Referenced by [acquiresleep\(\)](#), [holdingsleep\(\)](#), [initsleeplock\(\)](#), and [releasesleep\(\)](#).

4.32.2.3 name

```
char* sleeplock::name
```

Definition at line 7 of file [sleeplock.h](#).

Referenced by [initsleeplock\(\)](#).

4.32.2.4 pid

```
int sleeplock::pid
```

Definition at line 8 of file [sleeplock.h](#).

Referenced by [acquiresleep\(\)](#), [holdingsleep\(\)](#), [initsleeplock\(\)](#), and [releasesleep\(\)](#).

The documentation for this struct was generated from the following file:

- [sleeplock.h](#)

4.33 spinlock

```
#include <spinlock.h>
```

Public Attributes

- struct [cpu](#) * [cpu](#)
- [uint](#) [locked](#)
- char * [name](#)
- [uint](#) [pcs](#) [10]

4.33.1 Detailed Description

Definition at line 2 of file [spinlock.h](#).

4.33.2 Member Data Documentation

4.33.2.1 cpu

```
struct cpu* spinlock::cpu
```

Definition at line 7 of file [spinlock.h](#).

Referenced by [acquire\(\)](#), [holding\(\)](#), [initlock\(\)](#), and [release\(\)](#).

4.33.2.2 locked

```
uint spinlock::locked
```

Definition at line 3 of file [spinlock.h](#).

Referenced by [acquire\(\)](#), [holding\(\)](#), [initlock\(\)](#), and [release\(\)](#).

4.33.2.3 name

```
char* spinlock::name
```

Definition at line 6 of file [spinlock.h](#).

Referenced by [initlock\(\)](#).

4.33.2.4 pcs

```
uint spinlock::pcs[10]
```

Definition at line 8 of file [spinlock.h](#).

Referenced by [acquire\(\)](#), [panic\(\)](#), and [release\(\)](#).

The documentation for this struct was generated from the following file:

- [spinlock.h](#)

4.34 stat

```
#include <stat.h>
```

Public Attributes

- int [dev](#)
- uint [ino](#)
- short [nlink](#)
- uint [size](#)
- short [type](#)

4.34.1 Detailed Description

Definition at line 5 of file [stat.h](#).

4.34.2 Member Data Documentation

4.34.2.1 dev

```
int stat::dev
```

Definition at line 7 of file [stat.h](#).

Referenced by [stati\(\)](#).

4.34.2.2 ino

```
uint stat::ino
```

Definition at line 8 of file [stat.h](#).

Referenced by [ls\(\)](#), and [stati\(\)](#).

4.34.2.3 nlink

```
short stat::nlink
```

Definition at line 9 of file [stat.h](#).

Referenced by [stati\(\)](#).

4.34.2.4 size

```
uint stat::size
```

Definition at line 10 of file [stat.h](#).

Referenced by [ls\(\)](#), and [stati\(\)](#).

4.34.2.5 type

```
short stat::type
```

Definition at line 6 of file [stat.h](#).

Referenced by [ls\(\)](#), and [stati\(\)](#).

The documentation for this struct was generated from the following file:

- [stat.h](#)

4.35 superblock

```
#include <fs.h>
```

Public Attributes

- [uint bmapstart](#)
- [uint inodestart](#)
- [uint logstart](#)
- [uint nblocks](#)
- [uint ninodes](#)
- [uint nlog](#)
- [uint size](#)

4.35.1 Detailed Description

Definition at line 14 of file [fs.h](#).

4.35.2 Member Data Documentation

4.35.2.1 bmapstart

`uint superblock::bmapstart`

Definition at line 21 of file [fs.h](#).

Referenced by [balloc\(\)](#), [iinit\(\)](#), and [main\(\)](#).

4.35.2.2 inodestart

`uint superblock::inodestart`

Definition at line 20 of file [fs.h](#).

Referenced by [iinit\(\)](#), and [main\(\)](#).

4.35.2.3 logstart

`uint superblock::logstart`

Definition at line 19 of file [fs.h](#).

Referenced by [iinit\(\)](#), [initlog\(\)](#), and [main\(\)](#).

4.35.2.4 nblocks

`uint superblock::nblocks`

Definition at line 16 of file [fs.h](#).

Referenced by [iinit\(\)](#), and [main\(\)](#).

4.35.2.5 ninodes

`uint superblock::ninodes`

Definition at line 17 of file [fs.h](#).

Referenced by [ialloc\(\)](#), [iinit\(\)](#), and [main\(\)](#).

4.35.2.6 nlog

```
uint superblock::nlog
```

Definition at line 18 of file [fs.h](#).

Referenced by [iinit\(\)](#), [initlog\(\)](#), and [main\(\)](#).

4.35.2.7 size

```
uint superblock::size
```

Definition at line 15 of file [fs.h](#).

Referenced by [balloc\(\)](#), [iinit\(\)](#), and [main\(\)](#).

The documentation for this struct was generated from the following file:

- [fs.h](#)

4.36 taskstate

```
#include <mmu.h>
```

Public Attributes

- void * [cr3](#)
- ushort [cs](#)
- ushort [ds](#)
- uint [eax](#)
- uint * [ebp](#)
- uint [ebx](#)
- uint [ecx](#)
- uint [edi](#)
- uint [edx](#)
- uint [eflags](#)
- uint * [eip](#)
- ushort [es](#)
- uint [esi](#)
- uint * [esp](#)
- uint [esp0](#)
- uint * [esp1](#)
- uint * [esp2](#)
- ushort [fs](#)
- ushort [gs](#)
- ushort [iomb](#)
- ushort [ldt](#)
- uint [link](#)

- [ushort padding1](#)
- [ushort padding10](#)
- [ushort padding2](#)
- [ushort padding3](#)
- [ushort padding4](#)
- [ushort padding5](#)
- [ushort padding6](#)
- [ushort padding7](#)
- [ushort padding8](#)
- [ushort padding9](#)
- [ushort ss](#)
- [ushort ss0](#)
- [ushort ss1](#)
- [ushort ss2](#)
- [ushort t](#)

4.36.1 Detailed Description

Definition at line 107 of file [mmu.h](#).

4.36.2 Member Data Documentation

4.36.2.1 cr3

```
void* taskstate::cr3
```

Definition at line 118 of file [mmu.h](#).

4.36.2.2 cs

```
ushort taskstate::cs
```

Definition at line 131 of file [mmu.h](#).

4.36.2.3 ds

```
ushort taskstate::ds
```

Definition at line 135 of file [mmu.h](#).

4.36.2.4 eax

```
uint taskstate::eax
```

Definition at line 121 of file [mmu.h](#).

4.36.2.5 ebp

```
uint* taskstate::ebp
```

Definition at line 126 of file [mmu.h](#).

4.36.2.6 ebx

```
uint taskstate::ebx
```

Definition at line 124 of file [mmu.h](#).

4.36.2.7 ecx

```
uint taskstate::ecx
```

Definition at line 122 of file [mmu.h](#).

4.36.2.8 edi

```
uint taskstate::edi
```

Definition at line 128 of file [mmu.h](#).

4.36.2.9 edx

```
uint taskstate::edx
```

Definition at line 123 of file [mmu.h](#).

4.36.2.10 eflags

```
uint taskstate::eflags
```

Definition at line 120 of file [mmu.h](#).

4.36.2.11 eip

```
uint* taskstate::eip
```

Definition at line 119 of file [mmu.h](#).

4.36.2.12 es

```
ushort taskstate::es
```

Definition at line 129 of file [mmu.h](#).

4.36.2.13 esi

```
uint taskstate::esi
```

Definition at line 127 of file [mmu.h](#).

4.36.2.14 esp

```
uint* taskstate::esp
```

Definition at line 125 of file [mmu.h](#).

4.36.2.15 esp0

```
uint taskstate::esp0
```

Definition at line 109 of file [mmu.h](#).

Referenced by [switchvm\(\)](#).

4.36.2.16 esp1

```
uint* taskstate::esp1
```

Definition at line 112 of file [mmu.h](#).

4.36.2.17 esp2

```
uint* taskstate::esp2
```

Definition at line 115 of file [mmu.h](#).

4.36.2.18 fs

```
ushort taskstate::fs
```

Definition at line 137 of file [mmu.h](#).

4.36.2.19 gs

```
ushort taskstate::gs
```

Definition at line 139 of file [mmu.h](#).

4.36.2.20 iomb

```
ushort taskstate::iomb
```

Definition at line 144 of file [mmu.h](#).

Referenced by [switchvm\(\)](#).

4.36.2.21 ldt

```
ushort taskstate::ldt
```

Definition at line 141 of file [mmu.h](#).

4.36.2.22 link

```
uint taskstate::link
```

Definition at line 108 of file [mmu.h](#).

4.36.2.23 padding1

```
ushort taskstate::padding1
```

Definition at line 111 of file [mmu.h](#).

4.36.2.24 padding10

```
ushort taskstate::padding10
```

Definition at line 142 of file [mmu.h](#).

4.36.2.25 padding2

```
ushort taskstate::padding2
```

Definition at line 114 of file [mmu.h](#).

4.36.2.26 padding3

```
ushort taskstate::padding3
```

Definition at line 117 of file [mmu.h](#).

4.36.2.27 padding4

```
ushort taskstate::padding4
```

Definition at line 130 of file [mmu.h](#).

4.36.2.28 padding5

```
ushort taskstate::padding5
```

Definition at line 132 of file [mmu.h](#).

4.36.2.29 padding6

```
ushort taskstate::padding6
```

Definition at line 134 of file [mmu.h](#).

4.36.2.30 padding7

```
ushort taskstate::padding7
```

Definition at line 136 of file [mmu.h](#).

4.36.2.31 padding8

```
ushort taskstate::padding8
```

Definition at line 138 of file [mmu.h](#).

4.36.2.32 padding9

```
ushort taskstate::padding9
```

Definition at line 140 of file [mmu.h](#).

4.36.2.33 ss

```
ushort taskstate::ss
```

Definition at line 133 of file [mmu.h](#).

4.36.2.34 ss0

```
ushort taskstate::ss0
```

Definition at line 110 of file [mmu.h](#).

Referenced by [switchvm\(\)](#).

4.36.2.35 ss1

```
ushort taskstate::ss1
```

Definition at line 113 of file [mmu.h](#).

4.36.2.36 ss2

```
ushort taskstate::ss2
```

Definition at line 116 of file [mmu.h](#).

4.36.2.37 t

```
ushort taskstate::t
```

Definition at line 143 of file [mmu.h](#).

The documentation for this struct was generated from the following file:

- [mmu.h](#)

4.37 trapframe

```
#include <x86.h>
```

Public Attributes

- [ushort cs](#)
- [ushort ds](#)
- [uint eax](#)
- [uint ebp](#)
- [uint ebx](#)
- [uint ecx](#)
- [uint edi](#)
- [uint edx](#)
- [uint eflags](#)
- [uint eip](#)
- [uint err](#)
- [ushort es](#)
- [uint esi](#)
- [uint esp](#)
- [ushort fs](#)
- [ushort gs](#)
- [uint oesp](#)
- [ushort padding1](#)
- [ushort padding2](#)
- [ushort padding3](#)
- [ushort padding4](#)
- [ushort padding5](#)
- [ushort padding6](#)
- [ushort ss](#)
- [uint trapno](#)

4.37.1 Detailed Description

Definition at line 150 of file [x86.h](#).

4.37.2 Member Data Documentation

4.37.2.1 cs

```
ushort trapframe::cs
```

Definition at line 175 of file [x86.h](#).

Referenced by [trap\(\)](#), and [userinit\(\)](#).

4.37.2.2 ds

```
ushort trapframe::ds
```

Definition at line 168 of file [x86.h](#).

Referenced by [userinit\(\)](#).

4.37.2.3 eax

```
uint trapframe::eax
```

Definition at line 159 of file [x86.h](#).

Referenced by [fork\(\)](#), and [syscall\(\)](#).

4.37.2.4 ebp

```
uint trapframe::ebp
```

Definition at line 154 of file [x86.h](#).

4.37.2.5 ebx

```
uint trapframe::ebx
```

Definition at line 156 of file [x86.h](#).

4.37.2.6 ecx

```
uint trapframe::ecx
```

Definition at line 158 of file [x86.h](#).

4.37.2.7 edi

```
uint trapframe::edi
```

Definition at line 152 of file [x86.h](#).

4.37.2.8 edx

```
uint trapframe::edx
```

Definition at line 157 of file [x86.h](#).

4.37.2.9 eflags

```
uint trapframe::eflags
```

Definition at line 177 of file [x86.h](#).

Referenced by [userinit\(\)](#).

4.37.2.10 eip

```
uint trapframe::eip
```

Definition at line 174 of file [x86.h](#).

Referenced by [exec\(\)](#), [trap\(\)](#), and [userinit\(\)](#).

4.37.2.11 err

```
uint trapframe::err
```

Definition at line 173 of file [x86.h](#).

Referenced by [trap\(\)](#).

4.37.2.12 es

```
ushort trapframe::es
```

Definition at line 166 of file [x86.h](#).

Referenced by [userinit\(\)](#).

4.37.2.13 esi

```
uint trapframe::esi
```

Definition at line 153 of file [x86.h](#).

4.37.2.14 esp

```
uint trapframe::esp
```

Definition at line 180 of file [x86.h](#).

Referenced by [argint\(\)](#), [exec\(\)](#), and [userinit\(\)](#).

4.37.2.15 fs

```
ushort trapframe::fs
```

Definition at line 164 of file [x86.h](#).

4.37.2.16 gs

```
ushort trapframe::gs
```

Definition at line 162 of file [x86.h](#).

4.37.2.17 oesp

```
uint trapframe::oesp
```

Definition at line 155 of file [x86.h](#).

4.37.2.18 padding1

```
ushort trapframe::padding1
```

Definition at line 163 of file [x86.h](#).

4.37.2.19 padding2

```
ushort trapframe::padding2
```

Definition at line 165 of file [x86.h](#).

4.37.2.20 padding3

```
ushort trapframe::padding3
```

Definition at line 167 of file [x86.h](#).

4.37.2.21 padding4

```
ushort trapframe::padding4
```

Definition at line 169 of file [x86.h](#).

4.37.2.22 padding5

```
ushort trapframe::padding5
```

Definition at line 176 of file [x86.h](#).

4.37.2.23 padding6

```
ushort trapframe::padding6
```

Definition at line 182 of file [x86.h](#).

4.37.2.24 ss

```
ushort trapframe::ss
```

Definition at line 181 of file [x86.h](#).

Referenced by [userinit\(\)](#).

4.37.2.25 trapno

```
uint trapframe::trapno
```

Definition at line 170 of file [x86.h](#).

Referenced by [trap\(\)](#).

The documentation for this struct was generated from the following file:

- [x86.h](#)

4.38 uproc

```
#include <uproc.h>
```

Public Attributes

- char [name](#) [16]
- int [pid](#)
- int [ppid](#)

4.38.1 Detailed Description

Definition at line 1 of file [uproc.h](#).

4.38.2 Member Data Documentation

4.38.2.1 name

```
char uproc::name[16]
```

Definition at line 4 of file [uproc.h](#).

4.38.2.2 pid

```
int uproc::pid
```

Definition at line 2 of file [uproc.h](#).

Referenced by [main\(\)](#), and [sys_getprocs\(\)](#).

4.38.2.3 ppid

```
int uproc::ppid
```

Definition at line 3 of file [uproc.h](#).

Referenced by [main\(\)](#), and [sys_getprocs\(\)](#).

The documentation for this struct was generated from the following file:

- [uproc.h](#)

Chapter 5

File Documentation

5.1 asm.h File Reference

Macros

- #define [SEG_ASM](#)(type, [base](#), lim)
- #define [SEG_NULLASM](#)
- #define [STA_R](#) 0x2
- #define [STA_W](#) 0x2
- #define [STA_X](#) 0x8

5.1.1 Macro Definition Documentation

5.1.1.1 SEG_ASM

```
#define SEG_ASM(  
    type,  
    base,  
    lim )
```

Value:

```
.word (((lim) >> 12) & 0xffff), ((base) & 0xffff); \  
.byte (((base) >> 16) & 0xff), (0x90 | (type)), \  
      (0xc0 | (((lim) >> 28) & 0xf)), (((base) >> 24) & 0xff)
```

Definition at line 11 of file [asm.h](#).

5.1.1.2 SEG_NULLASM

```
#define SEG_NULLASM
```

Value:

```
.word 0, 0; \  
.byte 0, 0, 0, 0
```

Definition at line 5 of file [asm.h](#).

5.1.1.3 STA_R

```
#define STA_R 0x2
```

Definition at line 18 of file [asm.h](#).

5.1.1.4 STA_W

```
#define STA_W 0x2
```

Definition at line 17 of file [asm.h](#).

5.1.1.5 STA_X

```
#define STA_X 0x8
```

Definition at line 16 of file [asm.h](#).

5.2 asm.h

[Go to the documentation of this file.](#)

```
00001 //
00002 // assembler macros to create x86 segments
00003 //
00004
00005 #define SEG_NULLASM
00006         .word 0, 0;
00007         .byte 0, 0, 0, 0
00008
00009 // The 0xC0 means the limit is in 4096-byte units
00010 // and (for executable segments) 32-bit mode.
00011 #define SEG_ASM(type,base,lim)
00012         .word (((lim) >> 12) & 0xffff), ((base) & 0xffff);
00013         .byte (((base) >> 16) & 0xff), (0x90 | (type)),
00014         (0xC0 | (((lim) >> 28) & 0xf)), (((base) >> 24) & 0xff)
00015
00016 #define STA_X      0x8      // Executable segment
00017 #define STA_W      0x2      // Writeable (non-executable segments)
00018 #define STA_R      0x2      // Readable (executable segments)
```

5.3 bio.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "buf.h"
```

Functions

- static struct `buf` * `bget` (`uint` dev, `uint` blockno)
- void `binit` (void)
- struct `buf` * `bread` (`uint` dev, `uint` blockno)
- void `brelse` (struct `buf` *b)
- void `bwrite` (struct `buf` *b)

Variables

- struct {
 struct `buf` `buf` [`NBUF`]
 struct `buf` `head`
 struct `spinlock` `lock`
} `bcache`

5.3.1 Function Documentation

5.3.1.1 bget()

```
static struct buf * bget (
    uint dev,
    uint blockno ) [static]
```

Definition at line 62 of file `bio.c`.

```
00063 {
00064     struct buf *b;
00065
00066     acquire(&bcache.lock);
00067
00068     // Is the block already cached?
00069     for(b = bcache.head.next; b != &bcache.head; b = b->next){
00070         if(b->dev == dev && b->blockno == blockno){
00071             b->refcnt++;
00072             release(&bcache.lock);
00073             acquiresleep(&b->lock);
00074             return b;
00075         }
00076     }
00077
00078     // Not cached; recycle an unused buffer.
00079     // Even if refcnt==0, B_DIRTY indicates a buffer is in use
00080     // because log.c has modified it but not yet committed it.
00081     for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
00082         if(b->refcnt == 0 && (b->flags & B_DIRTY) == 0) {
00083             b->dev = dev;
00084             b->blockno = blockno;
00085             b->flags = 0;
00086             b->refcnt = 1;
00087             release(&bcache.lock);
00088             acquiresleep(&b->lock);
00089             return b;
00090         }
00091     }
00092     panic("bget: no buffers");
00093 }
```

Referenced by `bread`).

5.3.1.2 binit()

```
void binit (
    void )
```

Definition at line 39 of file [bio.c](#).

```
00040 {
00041     struct buf *b;
00042
00043     initlock(&bcache.lock, "bcache");
00044
00045     //PAGEBREAK!
00046     // Create linked list of buffers
00047     bcache.head.prev = &bcache.head;
00048     bcache.head.next = &bcache.head;
00049     for(b = bcache.buf; b < bcache.buf+NBUF; b++){
00050         b->next = bcache.head.next;
00051         b->prev = &bcache.head;
00052         initsleeplock(&b->lock, "buffer");
00053         bcache.head.next->prev = b;
00054         bcache.head.next = b;
00055     }
00056 }
```

5.3.1.3 bread()

```
struct buf * bread (
    uint dev,
    uint blockno )
```

Definition at line 97 of file [bio.c](#).

```
00098 {
00099     struct buf *b;
00100
00101     b = bget(dev, blockno);
00102     if((b->flags & B_INVALID) == 0) {
00103         iderw(b);
00104     }
00105     return b;
00106 }
```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [ilock\(\)](#), [install_trans\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [read_head\(\)](#), [readi\(\)](#), [readsb\(\)](#), [write_head\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.3.1.4 brelse()

```
void brelse (
    struct buf * b )
```

Definition at line 121 of file [bio.c](#).

```
00122 {
00123     if(!holdingsleep(&b->lock))
00124         panic("brelse");
00125
00126     releasesleep(&b->lock);
00127
00128     acquire(&bcache.lock);
00129     b->refcnt--;
00130     if (b->refcnt == 0) {
00131         // no one is waiting for it.
00132         b->next->prev = b->prev;
00133         b->prev->next = b->next;
00134         b->next = bcache.head.next;
00135         b->prev = &bcache.head;
```

```

00136     bcache.head.next->prev = b;
00137     bcache.head.next = b;
00138 }
00139
00140 release(&bcache.lock);
00141 }

```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [ilock\(\)](#), [install_trans\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [read_head\(\)](#), [readi\(\)](#), [readsb\(\)](#), [write_head\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.3.1.5 bwrite()

```

void bwrite (
    struct buf * b )

```

Definition at line 110 of file [bio.c](#).

```

00111 {
00112     if(!holdingsleep(&b->lock))
00113         panic("bwrite");
00114     b->flags |= B_DIRTY;
00115     iderw(b);
00116 }

```

Referenced by [install_trans\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

5.3.2 Variable Documentation

5.3.2.1

```

struct { ... } bcache

```

Referenced by [bget\(\)](#), [binit\(\)](#), and [brelse\(\)](#).

5.3.2.2 buf

```

struct buf buf[NBUF]

```

Definition at line 31 of file [bio.c](#).

5.3.2.3 head

```

struct buf head

```

Definition at line 35 of file [bio.c](#).

5.3.2.4 lock

struct `spinlock` lock

Definition at line 30 of file `bio.c`.

Referenced by `holding()`.

5.4 bio.c

[Go to the documentation of this file.](#)

```

00001 // Buffer cache.
00002 //
00003 // The buffer cache is a linked list of buf structures holding
00004 // cached copies of disk block contents. Caching disk blocks
00005 // in memory reduces the number of disk reads and also provides
00006 // a synchronization point for disk blocks used by multiple processes.
00007 //
00008 // Interface:
00009 // * To get a buffer for a particular disk block, call bread.
00010 // * After changing buffer data, call bwrite to write it to disk.
00011 // * When done with the buffer, call brelse.
00012 // * Do not use the buffer after calling brelse.
00013 // * Only one process at a time can use a buffer,
00014 //   so do not keep them longer than necessary.
00015 //
00016 // The implementation uses two state flags internally:
00017 // * B_VALID: the buffer data has been read from the disk.
00018 // * B_DIRTY: the buffer data has been modified
00019 //   and needs to be written to disk.
00020
00021 #include "types.h"
00022 #include "defs.h"
00023 #include "param.h"
00024 #include "spinlock.h"
00025 #include "sleeplock.h"
00026 #include "fs.h"
00027 #include "buf.h"
00028
00029 struct {
00030     struct spinlock lock;
00031     struct buf buf[NBUF];
00032
00033     // Linked list of all buffers, through prev/next.
00034     // head.next is most recently used.
00035     struct buf head;
00036 } bcache;
00037
00038 void
00039 binit(void)
00040 {
00041     struct buf *b;
00042
00043     initlock(&bcache.lock, "bcache");
00044
00045     //PAGEBREAK!
00046     // Create linked list of buffers
00047     bcache.head.prev = &bcache.head;
00048     bcache.head.next = &bcache.head;
00049     for(b = bcache.buf; b < bcache.buf+NBUF; b++){
00050         b->next = bcache.head.next;
00051         b->prev = &bcache.head;
00052         initsleeplock(&b->lock, "buffer");
00053         bcache.head.next->prev = b;
00054         bcache.head.next = b;
00055     }
00056 }
00057
00058 // Look through buffer cache for block on device dev.
00059 // If not found, allocate a buffer.
00060 // In either case, return locked buffer.
00061 static struct buf*
00062 bget(uint dev, uint blockno)
00063 {
00064     struct buf *b;
00065
00066     acquire(&bcache.lock);

```



```

00067
00068 // Is the block already cached?
00069 for(b = bcache.head.next; b != &bcache.head; b = b->next){
00070     if(b->dev == dev && b->blockno == blockno){
00071         b->refcnt++;
00072         release(&bcache.lock);
00073         acquiresleep(&b->lock);
00074         return b;
00075     }
00076 }
00077
00078 // Not cached; recycle an unused buffer.
00079 // Even if refcnt==0, B_DIRTY indicates a buffer is in use
00080 // because log.c has modified it but not yet committed it.
00081 for(b = bcache.head.prev; b != &bcache.head; b = b->prev){
00082     if(b->refcnt == 0 && (b->flags & B_DIRTY) == 0) {
00083         b->dev = dev;
00084         b->blockno = blockno;
00085         b->flags = 0;
00086         b->refcnt = 1;
00087         release(&bcache.lock);
00088         acquiresleep(&b->lock);
00089         return b;
00090     }
00091 }
00092 panic("bget: no buffers");
00093 }
00094
00095 // Return a locked buf with the contents of the indicated block.
00096 struct buf*
00097 bread(uint dev, uint blockno)
00098 {
00099     struct buf *b;
00100
00101     b = bget(dev, blockno);
00102     if((b->flags & B_VALID) == 0) {
00103         iderw(b);
00104     }
00105     return b;
00106 }
00107
00108 // Write b's contents to disk. Must be locked.
00109 void
00110 bwrite(struct buf *b)
00111 {
00112     if(!holdingsleep(&b->lock))
00113         panic("bwrite");
00114     b->flags |= B_DIRTY;
00115     iderw(b);
00116 }
00117
00118 // Release a locked buffer.
00119 // Move to the head of the MRU list.
00120 void
00121 brelse(struct buf *b)
00122 {
00123     if(!holdingsleep(&b->lock))
00124         panic("brelse");
00125
00126     releasesleep(&b->lock);
00127
00128     acquire(&bcache.lock);
00129     b->refcnt--;
00130     if (b->refcnt == 0) {
00131         // no one is waiting for it.
00132         b->next->prev = b->prev;
00133         b->prev->next = b->next;
00134         b->next = bcache.head.next;
00135         b->prev = &bcache.head;
00136         bcache.head.next->prev = b;
00137         bcache.head.next = b;
00138     }
00139
00140     release(&bcache.lock);
00141 }
00142 //PAGEBREAK!
00143 // Blank page.
00144

```

5.5 bio.d File Reference

5.6 bio.d

[Go to the documentation of this file.](#)

```
00001 bio.o: bio.c /usr/include/stdc-predef.h types.h defs.h param.h spinlock.h \  
00002  sleeplock.h fs.h buf.h
```

5.7 bootasm.d File Reference

5.8 bootasm.d

[Go to the documentation of this file.](#)

```
00001 bootasm.o: bootasm.S asm.h memlayout.h mmu.h
```

5.9 bootmain.c File Reference

```
#include "types.h"  
#include "elf.h"  
#include "x86.h"  
#include "memlayout.h"
```

Macros

- `#define` [SECTSIZE](#) 512

Functions

- void [bootmain](#) (void)
- void [readsect](#) (void *dst, [uint](#) offset)
- void [readseg](#) ([uchar](#) *, [uint](#), [uint](#))
- void [waitdisk](#) (void)

5.9.1 Macro Definition Documentation

5.9.1.1 SECTSIZE

```
#define SECTSIZE 512
```

Definition at line 13 of file [bootmain.c](#).

5.9.2 Function Documentation

5.9.2.1 bootmain()

```
void bootmain (
    void )
```

Definition at line 18 of file [bootmain.c](#).

```
00019 {
00020     struct elfhdr *elf;
00021     struct proghdr *ph, *eph;
00022     void (*entry) (void);
00023     uchar* pa;
00024
00025     elf = (struct elfhdr*)0x10000; // scratch space
00026
00027     // Read 1st page off disk
00028     readseg((uchar*)elf, 4096, 0);
00029
00030     // Is this an ELF executable?
00031     if(elf->magic != ELF_MAGIC)
00032         return; // let bootasm.S handle error
00033
00034     // Load each program segment (ignores ph flags).
00035     ph = (struct proghdr*)((uchar*)elf + elf->phoff);
00036     eph = ph + elf->phnum;
00037     for(; ph < eph; ph++){
00038         pa = (uchar*)ph->paddr;
00039         readseg(pa, ph->filesz, ph->off);
00040         if(ph->memsz > ph->filesz)
00041             stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
00042     }
00043
00044     // Call the entry point from the ELF header.
00045     // Does not return!
00046     entry = (void(*) (void)) (elf->entry);
00047     entry();
00048 }
```

5.9.2.2 readsect()

```
void readsect (
    void * dst,
    uint offset )
```

Definition at line 60 of file [bootmain.c](#).

```
00061 {
00062     // Issue command.
00063     waitdisk();
00064     outb(0x1F2, 1); // count = 1
00065     outb(0x1F3, offset);
00066     outb(0x1F4, offset » 8);
00067     outb(0x1F5, offset » 16);
00068     outb(0x1F6, (offset » 24) | 0xE0);
00069     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
00070
00071     // Read data.
00072     waitdisk();
00073     insl(0x1F0, dst, SECTSIZE/4);
00074 }
```

Referenced by [readseg\(\)](#).

5.9.2.3 readseg()

```
void readseg (
    uchar * pa,
    uint count,
    uint offset )
```

Definition at line 79 of file [bootmain.c](#).

```
00080 {
00081     uchar* epa;
00082
00083     epa = pa + count;
00084
00085     // Round down to sector boundary.
00086     pa -= offset % SECTSIZE;
00087
00088     // Translate from bytes to sectors; kernel starts at sector 1.
00089     offset = (offset / SECTSIZE) + 1;
00090
00091     // If this is too slow, we could read lots of sectors at a time.
00092     // We'd write more to memory than asked, but it doesn't matter --
00093     // we load in increasing order.
00094     for(; pa < epa; pa += SECTSIZE, offset++)
00095         readsect(pa, offset);
00096 }
```

Referenced by [bootmain\(\)](#).

5.9.2.4 waitdisk()

```
void waitdisk (
    void )
```

Definition at line 51 of file [bootmain.c](#).

```
00052 {
00053     // Wait for disk ready.
00054     while((inb(0x1F7) & 0xC0) != 0x40)
00055         ;
00056 }
```

Referenced by [readsect\(\)](#).

5.10 bootmain.c

[Go to the documentation of this file.](#)

```
00001 // Boot loader.
00002 //
00003 // Part of the boot block, along with bootasm.S, which calls bootmain().
00004 // bootasm.S has put the processor into protected 32-bit mode.
00005 // bootmain() loads an ELF kernel image from the disk starting at
00006 // sector 1 and then jumps to the kernel entry routine.
00007
00008 #include "types.h"
00009 #include "elf.h"
00010 #include "x86.h"
00011 #include "memlayout.h"
00012
00013 #define SECTSIZE 512
00014
00015 void readseg(uchar*, uint, uint);
00016
00017 void
00018 bootmain(void)
00019 {
00020     struct elfhdr *elf;
00021     struct proghdr *ph, *eph;
```

```

00022 void (*entry)(void);
00023 uchar* pa;
00024
00025 elf = (struct elfhdr*)0x10000; // scratch space
00026
00027 // Read 1st page off disk
00028 readseg((uchar*)elf, 4096, 0);
00029
00030 // Is this an ELF executable?
00031 if(elf->magic != ELF_MAGIC)
00032     return; // let bootasm.S handle error
00033
00034 // Load each program segment (ignores ph flags).
00035 ph = (struct proghdr*)((uchar*)elf + elf->phoff);
00036 eph = ph + elf->phnum;
00037 for(; ph < eph; ph++){
00038     pa = (uchar*)ph->paddr;
00039     readseg(pa, ph->filesz, ph->off);
00040     if(ph->memsz > ph->filesz)
00041         stosb(pa + ph->filesz, 0, ph->memsz - ph->filesz);
00042 }
00043
00044 // Call the entry point from the ELF header.
00045 // Does not return!
00046 entry = (void(*) (void)) (elf->entry);
00047 entry();
00048 }
00049
00050 void
00051 waitdisk(void)
00052 {
00053     // Wait for disk ready.
00054     while((inb(0x1F7) & 0xC0) != 0x40)
00055         ;
00056 }
00057
00058 // Read a single sector at offset into dst.
00059 void
00060 readsect(void *dst, uint offset)
00061 {
00062     // Issue command.
00063     waitdisk();
00064     outb(0x1F2, 1); // count = 1
00065     outb(0x1F3, offset);
00066     outb(0x1F4, offset » 8);
00067     outb(0x1F5, offset » 16);
00068     outb(0x1F6, (offset » 24) | 0xE0);
00069     outb(0x1F7, 0x20); // cmd 0x20 - read sectors
00070
00071     // Read data.
00072     waitdisk();
00073     insl(0x1F0, dst, SECTSIZE/4);
00074 }
00075
00076 // Read 'count' bytes at 'offset' from kernel into physical address 'pa'.
00077 // Might copy more than asked.
00078 void
00079 readseg(uchar* pa, uint count, uint offset)
00080 {
00081     uchar* epa;
00082
00083     epa = pa + count;
00084
00085     // Round down to sector boundary.
00086     pa -= offset % SECTSIZE;
00087
00088     // Translate from bytes to sectors; kernel starts at sector 1.
00089     offset = (offset / SECTSIZE) + 1;
00090
00091     // If this is too slow, we could read lots of sectors at a time.
00092     // We'd write more to memory than asked, but it doesn't matter --
00093     // we load in increasing order.
00094     for(; pa < epa; pa += SECTSIZE, offset++){
00095         readsect(pa, offset);
00096     }

```

5.11 bootmain.d File Reference

5.12 bootmain.d

[Go to the documentation of this file.](#)

```
00001 bootmain.o: bootmain.c types.h elf.h x86.h memlayout.h
```

5.13 buf.h File Reference

Classes

- struct [buf](#)

Macros

- #define [B_DIRTY](#) 0x4
- #define [B_VALID](#) 0x2

5.13.1 Macro Definition Documentation

5.13.1.1 B_DIRTY

```
#define B_DIRTY 0x4
```

Definition at line 13 of file [buf.h](#).

5.13.1.2 B_VALID

```
#define B_VALID 0x2
```

Definition at line 12 of file [buf.h](#).

5.14 buf.h

[Go to the documentation of this file.](#)

```
00001 struct buf {
00002     int flags;
00003     uint dev;
00004     uint blockno;
00005     struct sleeplock lock;
00006     uint refcnt;
00007     struct buf *prev; // LRU cache list
00008     struct buf *next;
00009     struct buf *qnext; // disk queue
00010     uchar data[BSIZE];
00011 };
00012 #define B_VALID 0x2 // buffer has been read from disk
00013 #define B_DIRTY 0x4 // buffer needs to be written to disk
00014
```

5.15 cat.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- void [cat](#) (int fd)
- int [main](#) (int argc, char *[argv](#)[])

Variables

- char [buf](#) [512]

5.15.1 Function Documentation

5.15.1.1 cat()

```
void cat (
    int fd )
```

Definition at line 8 of file [cat.c](#).

```
00009 {
00010     int n;
00011
00012     while ((n = read(fd, buf, sizeof(buf))) > 0) {
00013         if (write(1, buf, n) != n) {
00014             printf(1, "cat: write error\n");
00015             exit();
00016         }
00017     }
00018     if (n < 0) {
00019         printf(1, "cat: read error\n");
00020         exit();
00021     }
00022 }
```

Referenced by [main\(\)](#).

5.15.1.2 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 25 of file [cat.c](#).

```
00026 {
00027     int fd, i;
00028
00029     if(argc <= 1){
00030         cat(0);
00031         exit();
00032     }
00033
00034     for(i = 1; i < argc; i++){
00035         if((fd = open(argv[i], 0)) < 0){
00036             printf(1, "cat: cannot open %s\n", argv[i]);
00037             exit();
00038         }
00039         cat(fd);
00040         close(fd);
00041     }
00042     exit();
00043 }
```

5.15.2 Variable Documentation

5.15.2.1 buf

```
char buf[512]
```

Definition at line 5 of file [cat.c](#).

5.16 cat.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 char buf[512];
00006
00007 void
00008 cat(int fd)
00009 {
00010     int n;
00011
00012     while((n = read(fd, buf, sizeof(buf))) > 0) {
00013         if(write(1, buf, n) != n) {
00014             printf(1, "cat: write error\n");
00015             exit();
00016         }
00017     }
00018     if(n < 0){
00019         printf(1, "cat: read error\n");
00020         exit();
00021     }
00022 }
00023
00024 int
00025 main(int argc, char *argv[])
00026 {
00027     int fd, i;
00028 }
```



```

00029  if(argc <= 1){
00030      cat(0);
00031      exit();
00032  }
00033
00034  for(i = 1; i < argc; i++){
00035      if((fd = open(argv[i], 0)) < 0){
00036          printf(1, "cat: cannot open %s\n", argv[i]);
00037          exit();
00038      }
00039      cat(fd);
00040      close(fd);
00041  }
00042  exit();
00043 }

```

5.17 cat.d File Reference

5.18 cat.d

[Go to the documentation of this file.](#)

```
00001 cat.o: cat.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.19 console.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "traps.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "file.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"

```

Macros

- #define [BACKSPACE](#) 0x100
- #define [C\(x\)](#) ((x)-'@')
- #define [CRTPORT](#) 0x3d4
- #define [INPUT_BUF](#) 128

Functions

- static void [cgaputc](#) (int c)
- void [consoleinit](#) (void)
- void [consoleintr](#) (int(*getc)(void))
- int [consoleread](#) (struct [inode](#) *ip, char *dst, int n)
- int [consolewrite](#) (struct [inode](#) *ip, char *buf, int n)
- static void [consputc](#) (int)
- void [cprintf](#) (char *fmt,...)
- void [panic](#) (char *s)
- static void [printint](#) (int xx, int [base](#), int sign)

Variables

- struct {
 struct [spinlock](#) lock
 int [locking](#)
} [cons](#)
- static [ushort](#) * [crt](#) = ([ushort](#)*)[P2V](#)(0xb8000)
- struct {
 char [buf](#) [[INPUT_BUF](#)]
 [uint](#) [e](#)
 [uint](#) [r](#)
 [uint](#) [w](#)
} [input](#)
- static int [panicked](#) = 0

5.19.1 Macro Definition Documentation

5.19.1.1 BACKSPACE

```
#define BACKSPACE 0x100
```

Definition at line [127](#) of file [console.c](#).

5.19.1.2 C

```
#define C(  
    x ) ((x) - '@')
```

Definition at line [189](#) of file [console.c](#).

5.19.1.3 CRTPORT

```
#define CRTPORT 0x3d4
```

Definition at line [128](#) of file [console.c](#).

5.19.1.4 INPUT_BUF

```
#define INPUT_BUF 128
```

Definition at line [181](#) of file [console.c](#).

5.19.2 Function Documentation

5.19.2.1 cgaputc()

```
static void cgaputc (
    int c ) [static]
```

Definition at line 132 of file [console.c](#).

```
00133 {
00134     int pos;
00135
00136     // Cursor position: col + 80*row.
00137     outb(CRTPORT, 14);
00138     pos = inb(CRTPORT+1) << 8;
00139     outb(CRTPORT, 15);
00140     pos |= inb(CRTPORT+1);
00141
00142     if(c == '\n')
00143         pos += 80 - pos%80;
00144     else if(c == BACKSPACE){
00145         if(pos > 0) --pos;
00146     } else
00147         crt[pos++] = (c&0xff) | 0x0700; // black on white
00148
00149     if(pos < 0 || pos > 25*80)
00150         panic("pos under/overflow");
00151
00152     if((pos/80) >= 24){ // Scroll up.
00153         memmove(crt, crt+80, sizeof(crt[0])*23*80);
00154         pos -= 80;
00155         memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
00156     }
00157
00158     outb(CRTPORT, 14);
00159     outb(CRTPORT+1, pos>>8);
00160     outb(CRTPORT, 15);
00161     outb(CRTPORT+1, pos);
00162     crt[pos] = ' ' | 0x0700;
00163 }
```

Referenced by [consputc\(\)](#).

5.19.2.2 consoleinit()

```
void consoleinit (
    void )
```

Definition at line 289 of file [console.c](#).

```
00290 {
00291     initlock(&cons.lock, "console");
00292
00293     devsw[CONSOLE].write = consolewrite;
00294     devsw[CONSOLE].read = consoleread;
00295     cons.locking = 1;
00296
00297     ioapicenable(Irq_KBD, 0);
00298 }
```

5.19.2.3 consoleintr()

```
void consoleintr (
    int(*) (void) getc )
```

Definition at line 192 of file [console.c](#).

```
00193 {
00194     int c, doprocdump = 0;
00195
00196     acquire(&cons.lock);
00197     while((c = getc()) >= 0){
00198         switch(c){
00199             case C('P'): // Process listing.
00200                 // procdump() locks cons.lock indirectly; invoke later
00201                 doprocdump = 1;
00202                 break;
00203             case C('U'): // Kill line.
00204                 while(input.e != input.w &&
00205                     input.buf[(input.e-1) % INPUT_BUF] != '\n'){
00206                     input.e--;
00207                     consputc(BACKSPACE);
00208                 }
00209                 break;
00210             case C('H'): case '\x7f': // Backspace
00211                 if(input.e != input.w){
00212                     input.e--;
00213                     consputc(BACKSPACE);
00214                 }
00215                 break;
00216             default:
00217                 if(c != 0 && input.e-input.r < INPUT_BUF){
00218                     c = (c == '\r') ? '\n' : c;
00219                     input.buf[input.e++ % INPUT_BUF] = c;
00220                     consputc(c);
00221                     if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
00222                         input.w = input.e;
00223                         wakeup(&input.r);
00224                     }
00225                 }
00226                 break;
00227         }
00228     }
00229     release(&cons.lock);
00230     if(doprocdump) {
00231         procdump(); // now call procdump() wo. cons.lock held
00232     }
00233 }
```

Referenced by [kbdintr\(\)](#), and [uartintr\(\)](#).

5.19.2.4 consoleread()

```
int consoleread (
    struct inode * ip,
    char * dst,
    int n )
```

Definition at line 236 of file [console.c](#).

```
00237 {
00238     uint target;
00239     int c;
00240
00241     iunlock(ip);
00242     target = n;
00243     acquire(&cons.lock);
00244     while(n > 0){
00245         while(input.r == input.w){
00246             if(myproc()->killed){
00247                 release(&cons.lock);
00248                 ilock(ip);
00249                 return -1;
00250             }
00251             sleep(&input.r, &cons.lock);
```

```

00252     }
00253     c = input.buf[input.r++ % INPUT_BUF];
00254     if(c == C('D')){ // EOF
00255         if(n < target){
00256             // Save ^D for next time, to make sure
00257             // caller gets a 0-byte result.
00258             input.r--;
00259         }
00260         break;
00261     }
00262     *dst++ = c;
00263     --n;
00264     if(c == '\n')
00265         break;
00266 }
00267 release(&cons.lock);
00268 ilock(ip);
00269
00270 return target - n;
00271 }

```

Referenced by [consoleinit\(\)](#).

5.19.2.5 consolewrite()

```

int consolewrite (
    struct inode * ip,
    char * buf,
    int n )

```

Definition at line 274 of file [console.c](#).

```

00275 {
00276     int i;
00277
00278     iunlock(ip);
00279     acquire(&cons.lock);
00280     for(i = 0; i < n; i++)
00281         consputc(buf[i] & 0xff);
00282     release(&cons.lock);
00283     ilock(ip);
00284
00285     return n;
00286 }

```

Referenced by [consoleinit\(\)](#).

5.19.2.6 consputc()

```

void consputc (
    int c ) [static]

```

Definition at line 166 of file [console.c](#).

```

00167 {
00168     if(panicked){
00169         cli();
00170         for(;;)
00171             ;
00172     }
00173
00174     if(c == BACKSPACE){
00175         uartputc('\b'); uartputc(' '); uartputc('\b');
00176     } else
00177         uartputc(c);
00178     cgaputc(c);
00179 }

```

Referenced by [consoleintr\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), and [printint\(\)](#).

5.19.2.7 cprintf()

```
void cprintf (
    char * fmt,
    ... )
```

Definition at line 55 of file [console.c](#).

```
00056 {
00057     int i, c, locking;
00058     uint *argp;
00059     char *s;
00060
00061     locking = cons.locking;
00062     if(locking)
00063         acquire(&cons.lock);
00064
00065     if (fmt == 0)
00066         panic("null fmt");
00067
00068     argp = (uint*)(void*)&fmt + 1;
00069     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
00070         if(c != '%'){
00071             consputc(c);
00072             continue;
00073         }
00074         c = fmt[++i] & 0xff;
00075         if(c == 0)
00076             break;
00077         switch(c){
00078             case 'd':
00079             printint(*argp++, 10, 1);
00080             break;
00081             case 'x':
00082             case 'p':
00083             printint(*argp++, 16, 0);
00084             break;
00085             case 's':
00086                 if((s = (char*)*argp++) == 0)
00087                     s = "(null)";
00088                 for(; *s; s++)
00089                     consputc(*s);
00090                 break;
00091             case '%':
00092                 consputc('%');
00093                 break;
00094             default:
00095                 // Print unknown % sequence to draw attention.
00096                 consputc('%');
00097                 consputc(c);
00098                 break;
00099         }
00100     }
00101
00102     if(locking)
00103         release(&cons.lock);
00104 }
```

Referenced by [allocuvn\(\)](#), [cps\(\)](#), [exec\(\)](#), [iinit\(\)](#), [ioapicinit\(\)](#), [panic\(\)](#), [procdump\(\)](#), [syscall\(\)](#), and [trap\(\)](#).

5.19.2.8 panic()

```
void panic (
    char * s )
```

Definition at line 107 of file [console.c](#).

```
00108 {
00109     int i;
00110     uint pcs[10];
00111
00112     cli();
00113     cons.locking = 0;
00114     // use lapiccpunum so that we can call panic from mycpu()
00115     cprintf("lapicid %d: panic: ", lapicid());
```

```

00116     cprintf(s);
00117     cprintf("\n");
00118     getcallerpcs(&s, pcs);
00119     for(i=0; i<10; i++)
00120         cprintf(" %p", pcs[i]);
00121     panicked = 1; // freeze other CPU
00122     for(;;)
00123         ;
00124 }

```

Referenced by [acquire\(\)](#), [balloc\(\)](#), [bfree\(\)](#), [bget\(\)](#), [bmap\(\)](#), [brelse\(\)](#), [bwrite\(\)](#), [cgaputc\(\)](#), [clearpteu\(\)](#), [copyuvvm\(\)](#), [cprintf\(\)](#), [create\(\)](#), [deallocuvvm\(\)](#), [dirlink\(\)](#), [dirlookup\(\)](#), [end_op\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fileread\(\)](#), [filewrite\(\)](#), [freevm\(\)](#), [ialloc\(\)](#), [iderw\(\)](#), [idestart\(\)](#), [iget\(\)](#), [ilock\(\)](#), [initlog\(\)](#), [inituvvm\(\)](#), [isdirempty\(\)](#), [iunlock\(\)](#), [kfree\(\)](#), [loaduvvm\(\)](#), [log_write\(\)](#), [mappages\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), [popcli\(\)](#), [release\(\)](#), [sched\(\)](#), [setupkvm\(\)](#), [sleep\(\)](#), [switchuvvm\(\)](#), [sys_unlink\(\)](#), [trap\(\)](#), and [userinit\(\)](#).

5.19.2.9 printint()

```

static void printint (
    int xx,
    int base,
    int sign ) [static]

```

Definition at line 28 of file [console.c](#).

```

00029 {
00030     static char digits[] = "0123456789abcdef";
00031     char buf[16];
00032     int i;
00033     uint x;
00034
00035     if(sign && (sign = xx < 0))
00036         x = -xx;
00037     else
00038         x = xx;
00039
00040     i = 0;
00041     do{
00042         buf[i++] = digits[x % base];
00043     }while((x /= base) != 0);
00044
00045     if(sign)
00046         buf[i++] = '-';
00047
00048     while(--i >= 0)
00049         consputc(buf[i]);
00050 }

```

Referenced by [cprintf\(\)](#).

5.19.3 Variable Documentation

5.19.3.1 buf

```
char buf[INPUT_BUF]
```

Definition at line 183 of file [console.c](#).

5.19.3.2

```
struct { ... } cons [static]
```

Referenced by [consoleinit\(\)](#), [consoleintr\(\)](#), [consoleread\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), and [panic\(\)](#).

5.19.3.3 crt

```
ushort* crt = (ushort*)P2V(0xb8000) [static]
```

Definition at line 129 of file [console.c](#).

Referenced by [cgaputc\(\)](#).

5.19.3.4 e

```
uint e
```

Definition at line 186 of file [console.c](#).

Referenced by [mpinit\(\)](#), and [mpsearch1\(\)](#).

5.19.3.5

```
struct { ... } input
```

Referenced by [consoleintr\(\)](#), and [consoleread\(\)](#).

5.19.3.6 lock

```
struct spinlock lock
```

Definition at line 23 of file [console.c](#).

5.19.3.7 locking

```
int locking
```

Definition at line 24 of file [console.c](#).

Referenced by [cprintf\(\)](#).

5.19.3.8 panicked

```
int panicked = 0 [static]
```

Definition at line 20 of file [console.c](#).

Referenced by [consputc\(\)](#), and [panic\(\)](#).

5.19.3.9 r

```
uint r
```

Definition at line 184 of file [console.c](#).

Referenced by [cmostime\(\)](#), [fileread\(\)](#), [filewrite\(\)](#), [fill_rtcddate\(\)](#), [holding\(\)](#), [holdingsleep\(\)](#), [idewait\(\)](#), [iput\(\)](#), [kalloc\(\)](#), [kfree\(\)](#), and [stat\(\)](#).

5.19.3.10 w

```
uint w
```

Definition at line 185 of file [console.c](#).

Referenced by [wc\(\)](#).

5.20 console.c

[Go to the documentation of this file.](#)

```
00001 // Console input and output.
00002 // Input is from the keyboard or serial port.
00003 // Output is written to the screen and serial port.
00004
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "param.h"
00008 #include "traps.h"
00009 #include "spinlock.h"
00010 #include "sleeplock.h"
00011 #include "fs.h"
00012 #include "file.h"
00013 #include "memlayout.h"
00014 #include "mmu.h"
00015 #include "proc.h"
00016 #include "x86.h"
00017
00018 static void consputc(int);
00019
00020 static int panicked = 0;
00021
00022 static struct {
00023     struct spinlock lock;
00024     int locking;
00025 } cons;
00026
00027 static void
00028 printint(int xx, int base, int sign)
00029 {
00030     static char digits[] = "0123456789abcdef";
```

```

00031 char buf[16];
00032 int i;
00033 uint x;
00034
00035 if(sign && (sign = xx < 0))
00036     x = -xx;
00037 else
00038     x = xx;
00039
00040 i = 0;
00041 do{
00042     buf[i++] = digits[x % base];
00043 }while((x /= base) != 0);
00044
00045 if(sign)
00046     buf[i++] = '-';
00047
00048 while(--i >= 0)
00049     consputc(buf[i]);
00050 }
00051 //PAGEBREAK: 50
00052
00053 // Print to the console. only understands %d, %x, %p, %s.
00054 void
00055 cprintf(char *fmt, ...)
00056 {
00057     int i, c, locking;
00058     uint *argp;
00059     char *s;
00060
00061     locking = cons.locking;
00062     if(locking)
00063         acquire(&cons.lock);
00064
00065     if (fmt == 0)
00066         panic("null fmt");
00067
00068     argp = (uint*)(void*)&fmt + 1;
00069     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
00070         if(c != '%'){
00071             consputc(c);
00072             continue;
00073         }
00074         c = fmt[++i] & 0xff;
00075         if(c == 0)
00076             break;
00077         switch(c){
00078             case 'd':
00079                 printint(*argp++, 10, 1);
00080                 break;
00081             case 'x':
00082             case 'p':
00083                 printint(*argp++, 16, 0);
00084                 break;
00085             case 's':
00086                 if((s = (char*)*argp++) == 0)
00087                     s = "(null)";
00088                 for(; *s; s++)
00089                     consputc(*s);
00090                 break;
00091             case '%':
00092                 consputc('%');
00093                 break;
00094             default:
00095                 // Print unknown % sequence to draw attention.
00096                 consputc('%');
00097                 consputc(c);
00098                 break;
00099         }
00100     }
00101
00102     if(locking)
00103         release(&cons.lock);
00104 }
00105
00106 void
00107 panic(char *s)
00108 {
00109     int i;
00110     uint pcs[10];
00111
00112     cli();
00113     cons.locking = 0;
00114     // use lapiccpunum so that we can call panic from mycpu()
00115     cprintf("lapicid %d: panic: ", lapicid());
00116     cprintf(s);
00117     cprintf("\n");

```

```

00118     getcallerpcs(&s, pcs);
00119     for(i=0; i<10; i++)
00120         cprintf(" %p", pcs[i]);
00121     panicked = 1; // freeze other CPU
00122     for(;;)
00123         ;
00124 }
00125
00126 //PAGEBREAK: 50
00127 #define BACKSPACE 0x100
00128 #define CRTPORT 0x3d4
00129 static ushort *crt = (ushort*)P2V(0xb8000); // CGA memory
00130
00131 static void
00132 cgaputc(int c)
00133 {
00134     int pos;
00135
00136     // Cursor position: col + 80*row.
00137     outb(CRTPORT, 14);
00138     pos = inb(CRTPORT+1) << 8;
00139     outb(CRTPORT, 15);
00140     pos |= inb(CRTPORT+1);
00141
00142     if(c == '\n')
00143         pos += 80 - pos%80;
00144     else if(c == BACKSPACE){
00145         if(pos > 0) --pos;
00146     } else
00147         crt[pos++] = (c&0xff) | 0x0700; // black on white
00148
00149     if(pos < 0 || pos > 25*80)
00150         panic("pos under/overflow");
00151
00152     if((pos/80) >= 24){ // Scroll up.
00153         memmove(crt, crt+80, sizeof(crt[0])*23*80);
00154         pos -= 80;
00155         memset(crt+pos, 0, sizeof(crt[0])*(24*80 - pos));
00156     }
00157
00158     outb(CRTPORT, 14);
00159     outb(CRTPORT+1, pos>>8);
00160     outb(CRTPORT, 15);
00161     outb(CRTPORT+1, pos);
00162     crt[pos] = ' ' | 0x0700;
00163 }
00164
00165 void
00166 consputc(int c)
00167 {
00168     if(panicked){
00169         cli();
00170         for(;;)
00171             ;
00172     }
00173
00174     if(c == BACKSPACE){
00175         uartputc('\b'); uartputc(' '); uartputc('\b');
00176     } else
00177         uartputc(c);
00178     cgaputc(c);
00179 }
00180
00181 #define INPUT_BUF 128
00182 struct {
00183     char buf[INPUT_BUF];
00184     uint r; // Read index
00185     uint w; // Write index
00186     uint e; // Edit index
00187 } input;
00188
00189 #define C(x) ((x)-'@') // Control-x
00190
00191 void
00192 consoleintr(int (*getc)(void))
00193 {
00194     int c, doprocdump = 0;
00195
00196     acquire(&cons.lock);
00197     while((c = getc()) >= 0){
00198         switch(c){
00199             case C('P'): // Process listing.
00200                 // procdump() locks cons.lock indirectly; invoke later
00201                 doprocdump = 1;
00202                 break;
00203             case C('U'): // Kill line.
00204                 while(input.e != input.w &&

```

```

00205         input.buf[(input.e-1) % INPUT_BUF] != '\n'){
00206             input.e--;
00207             consputc(BACKSPACE);
00208         }
00209         break;
00210     case C('H'): case '\x7f': // Backspace
00211         if(input.e != input.w){
00212             input.e--;
00213             consputc(BACKSPACE);
00214         }
00215         break;
00216     default:
00217         if(c != 0 && input.e-input.r < INPUT_BUF){
00218             c = (c == '\r') ? '\n' : c;
00219             input.buf[input.e++] = c;
00220             consputc(c);
00221             if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
00222                 input.w = input.e;
00223                 wakeup(&input.r);
00224             }
00225         }
00226         break;
00227     }
00228 }
00229 release(&cons.lock);
00230 if(doprocDump) {
00231     procDump(); // now call procDump() wo. cons.lock held
00232 }
00233 }
00234
00235 int
00236 consleread(struct inode *ip, char *dst, int n)
00237 {
00238     uint target;
00239     int c;
00240
00241     iunlock(ip);
00242     target = n;
00243     acquire(&cons.lock);
00244     while(n > 0){
00245         while(input.r == input.w){
00246             if(myproc()->killed){
00247                 release(&cons.lock);
00248                 ilock(ip);
00249                 return -1;
00250             }
00251             sleep(&input.r, &cons.lock);
00252         }
00253         c = input.buf[input.r++ % INPUT_BUF];
00254         if(c == C('D')){ // EOF
00255             if(n < target){
00256                 // Save ^D for next time, to make sure
00257                 // caller gets a 0-byte result.
00258                 input.r--;
00259             }
00260             break;
00261         }
00262         *dst++ = c;
00263         --n;
00264         if(c == '\n')
00265             break;
00266     }
00267     release(&cons.lock);
00268     ilock(ip);
00269
00270     return target - n;
00271 }
00272
00273 int
00274 consolewrite(struct inode *ip, char *buf, int n)
00275 {
00276     int i;
00277
00278     iunlock(ip);
00279     acquire(&cons.lock);
00280     for(i = 0; i < n; i++)
00281         consputc(buf[i] & 0xff);
00282     release(&cons.lock);
00283     ilock(ip);
00284
00285     return n;
00286 }
00287
00288 void
00289 consoleinit(void)
00290 {
00291     initlock(&cons.lock, "console");

```

```
00292
00293     devsw[CONSOLE].write = consolewrite;
00294     devsw[CONSOLE].read = consoleread;
00295     cons.locking = 1;
00296
00297     ioapicenable(IRQ_KBD, 0);
00298 }
00299
```

5.21 console.d File Reference

5.22 console.d

[Go to the documentation of this file.](#)

```
00001 console.o: console.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 traps.h spinlock.h sleeplock.h fs.h file.h memlayout.h mmu.h proc.h \
00003 x86.h
```

5.23 date.h File Reference

Classes

- struct [rtcdat](#)e

5.24 date.h

[Go to the documentation of this file.](#)

```
00001 struct rtcdat {
00002     uint second;
00003     uint minute;
00004     uint hour;
00005     uint day;
00006     uint month;
00007     uint year;
00008 };
```

5.25 defs.h File Reference

Macros

- #define [NELEM](#)(x) (sizeof(x)/sizeof((x)[0]))

Functions

- void [acquire](#) (struct [spinlock](#) *)
- void [acquiresleep](#) (struct [sleeplock](#) *)
- int [allocuvm](#) ([pde_t](#) *, [uint](#), [uint](#))
- int [argint](#) (int, int *)
- int [argptr](#) (int, char **, int)
- int [argstr](#) (int, char **)
- void [begin_op](#) ()
- void [binit](#) (void)
- struct [buf](#) * [bread](#) ([uint](#), [uint](#))
- void [brelse](#) (struct [buf](#) *)
- void [bwrite](#) (struct [buf](#) *)
- int [chpr](#) (int pid, int priority)
- void [clearpteu](#) ([pde_t](#) *pgdir, char *uva)
- void [cmostime](#) (struct [rtcddate](#) *)
- void [consoleinit](#) (void)
- void [consoleintr](#) (int(*) (void))
- int [copyout](#) ([pde_t](#) *, [uint](#), void *, [uint](#))
- [pde_t](#) * [copyuvm](#) ([pde_t](#) *, [uint](#))
- void [cprintf](#) (char *,...)
- int [cps](#) (void)
- int [cpuid](#) (void)
- int [deallocuvm](#) ([pde_t](#) *, [uint](#), [uint](#))
- int [dirlink](#) (struct [inode](#) *, char *, [uint](#))
- struct [inode](#) * [dirlookup](#) (struct [inode](#) *, char *, [uint](#) *)
- void [end_op](#) ()
- int [exec](#) (char *, char **)
- void [exit](#) (void)
- int [fetchint](#) ([uint](#), int *)
- int [fetchstr](#) ([uint](#), char **)
- struct [file](#) * [filealloc](#) (void)
- void [fileclose](#) (struct [file](#) *)
- struct [file](#) * [filedup](#) (struct [file](#) *)
- void [fileinit](#) (void)
- int [fileread](#) (struct [file](#) *, char *, int n)
- int [filestat](#) (struct [file](#) *, struct [stat](#) *)
- int [filewrite](#) (struct [file](#) *, char *, int n)
- int [fork](#) (void)
- void [freevm](#) ([pde_t](#) *)
- void [getcallerpcs](#) (void *, [uint](#) *)
- int [getprocs](#) (int max, struct [uproc](#) *)
- int [growproc](#) (int)
- int [holding](#) (struct [spinlock](#) *)
- int [holdingsleep](#) (struct [sleeplock](#) *)
- struct [inode](#) * [ialloc](#) ([uint](#), short)
- void [ideinit](#) (void)
- void [ideintr](#) (void)
- void [iderw](#) (struct [buf](#) *)
- void [idtinit](#) (void)
- struct [inode](#) * [idup](#) (struct [inode](#) *)
- void [iinit](#) (int dev)
- void [ilock](#) (struct [inode](#) *)
- void [initlock](#) (struct [spinlock](#) *, char *)
- void [initlog](#) (int dev)

- void [initsleeplock](#) (struct [sleeplock](#) *, char *)
- void [inituvm](#) ([pde_t](#) *, char *, [uint](#))
- void [ioapicenable](#) (int irq, int [cpu](#))
- void [ioapicinit](#) (void)
- void [iput](#) (struct [inode](#) *)
- void [iunlock](#) (struct [inode](#) *)
- void [iunlockput](#) (struct [inode](#) *)
- void [iupdate](#) (struct [inode](#) *)
- char * [kalloc](#) (void)
- void [kbdintr](#) (void)
- void [kfree](#) (char *)
- int [kill](#) (int)
- void [kinit1](#) (void *, void *)
- void [kinit2](#) (void *, void *)
- void [kvmalloc](#) (void)
- void [lapiceoi](#) (void)
- int [lapicid](#) (void)
- void [lapicinit](#) (void)
- void [lapicstartap](#) ([uchar](#), [uint](#))
- int [loaduvm](#) ([pde_t](#) *, char *, struct [inode](#) *, [uint](#), [uint](#))
- void [log_write](#) (struct [buf](#) *)
- int [memcmp](#) (const void *, const void *, [uint](#))
- void * [memmove](#) (void *, const void *, [uint](#))
- void * [memset](#) (void *, int, [uint](#))
- void [microdelay](#) (int)
- void [mpinit](#) (void)
- struct [cpu](#) * [mycpu](#) (void)
- struct [proc](#) * [myproc](#) ()
- int [namecmp](#) (const char *, const char *)
- struct [inode](#) * [namei](#) (char *)
- struct [inode](#) * [nameiparent](#) (char *, char *)
- void [panic](#) (char *) [__attribute__\(\(noreturn\)\)](#)
- void [picenable](#) (int)
- void [picinit](#) (void)
- void [pinit](#) (void)
- int [pipealloc](#) (struct [file](#) **, struct [file](#) **)
- void [pipeclose](#) (struct [pipe](#) *, int)
- int [piperead](#) (struct [pipe](#) *, char *, int)
- int [pipewrite](#) (struct [pipe](#) *, char *, int)
- void [popcli](#) (void)
- void [procdump](#) (void)
- void [pushcli](#) (void)
- int [readi](#) (struct [inode](#) *, char *, [uint](#), [uint](#))
- void [readsb](#) (int dev, struct [superblock](#) *sb)
- void [release](#) (struct [spinlock](#) *)
- void [releasesleep](#) (struct [sleeplock](#) *)
- char * [safestrcpy](#) (char *, const char *, int)
- void [sched](#) (void)
- void [scheduler](#) (void) [__attribute__\(\(noreturn\)\)](#)
- void [segininit](#) (void)
- void [setproc](#) (struct [proc](#) *)
- [pde_t](#) * [setupkvm](#) (void)
- void [sleep](#) (void *, struct [spinlock](#) *)
- void [stati](#) (struct [inode](#) *, struct [stat](#) *)
- int [strlen](#) (const char *)

- int [strncmp](#) (const char *, const char *, [uint](#))
- char * [strncpy](#) (char *, const char *, int)
- void [switchkvm](#) (void)
- void [switchuvm](#) (struct [proc](#) *)
- void [switch](#) (struct [context](#) **, struct [context](#) *)
- void [syscall](#) (void)
- void [timerinit](#) (void)
- void [tvinit](#) (void)
- void [uartinit](#) (void)
- void [uartintr](#) (void)
- void [uartputc](#) (int)
- void [userinit](#) (void)
- char * [uva2ka](#) ([pde_t](#) *, char *)
- int [wait](#) (void)
- void [wakeup](#) (void *)
- int [writei](#) (struct [inode](#) *, char *, [uint](#), [uint](#))
- void [yield](#) (void)

Variables

- [uchar](#) [ioapicid](#)
- int [ismp](#)
- volatile [uint](#) * [lapic](#)
- [uint](#) [ticks](#)
- struct [spinlock](#) [tickslock](#)

5.25.1 Macro Definition Documentation

5.25.1.1 NELEM

```
#define NELEM(  
    x ) (sizeof(x)/sizeof((x)[0]))
```

Definition at line [195](#) of file [defs.h](#).

5.25.2 Function Documentation

5.25.2.1 acquire()

```
void acquire (
    struct spinlock * lk )
```

Definition at line 25 of file [spinlock.c](#).

```
00026 {
00027     pushcli(); // disable interrupts to avoid deadlock.
00028     if(holding(lk))
00029         panic("acquire");
00030
00031     // The xchg is atomic.
00032     while(xchg(&lk->locked, 1) != 0)
00033         ;
00034
00035     // Tell the C compiler and the processor to not move loads or stores
00036     // past this point, to ensure that the critical section's memory
00037     // references happen after the lock is acquired.
00038     __sync_synchronize();
00039
00040     // Record info about lock acquisition for debugging.
00041     lk->cpu = mycpu();
00042     getcallerpcs(&lk, lk->pcs);
00043 }
```

Referenced by [acquiresleep\(\)](#), [allocproc\(\)](#), [begin_op\(\)](#), [bget\(\)](#), [brelse\(\)](#), [chpr\(\)](#), [consoleintr\(\)](#), [consoleread\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), [cps\(\)](#), [end_op\(\)](#), [exit\(\)](#), [filealloc\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fork\(\)](#), [holdingsleep\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idup\(\)](#), [iget\(\)](#), [iput\(\)](#), [kalloc\(\)](#), [kfree\(\)](#), [kill\(\)](#), [log_write\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup\(\)](#), and [yield\(\)](#).

5.25.2.2 acquiresleep()

```
void acquiresleep (
    struct sleeplock * lk )
```

Definition at line 23 of file [sleeplock.c](#).

```
00024 {
00025     acquire(&lk->lk);
00026     while (lk->locked) {
00027         sleep(lk, &lk->lk);
00028     }
00029     lk->locked = 1;
00030     lk->pid = myproc()->pid;
00031     release(&lk->lk);
00032 }
```

Referenced by [bget\(\)](#), [ilock\(\)](#), and [iput\(\)](#).

5.25.2.3 allocvm()

```
int allocvm (
    pde_t * pgdir,
    uint oldsz,
    uint newsz )
```

Definition at line 222 of file [vm.c](#).

```
00223 {
00224     char *mem;
00225     uint a;
00226
00227     if(newsz >= KERNBASE)
00228         return 0;
```

```

00229     if(newsz < oldsz)
00230         return oldsz;
00231
00232     a = PGROUNDUP(oldsz);
00233     for(; a < newsz; a += PGSIZE){
00234         mem = kalloc();
00235         if(mem == 0){
00236             cprintf("allocuvn out of memory\n");
00237             deallocvm(pgdir, newsz, oldsz);
00238             return 0;
00239         }
00240         memset(mem, 0, PGSIZE);
00241         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
00242             cprintf("allocuvn out of memory (2)\n");
00243             deallocvm(pgdir, newsz, oldsz);
00244             kfree(mem);
00245             return 0;
00246         }
00247     }
00248     return newsz;
00249 }

```

Referenced by [exec\(\)](#), and [growproc\(\)](#).

5.25.2.4 argint()

```

int argint (
    int n,
    int * ip )

```

Definition at line 50 of file [syscall.c](#).

```

00051 {
00052     return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);
00053 }

```

Referenced by [argfd\(\)](#), [argptr\(\)](#), [argstr\(\)](#), [sys_chpr\(\)](#), [sys_exec\(\)](#), [sys_getprocs\(\)](#), [sys_kill\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), [sys_read\(\)](#), [sys_sbrk\(\)](#), [sys_sleep\(\)](#), and [sys_write\(\)](#).

5.25.2.5 argptr()

```

int argptr (
    int n,
    char ** pp,
    int size )

```

Definition at line 59 of file [syscall.c](#).

```

00060 {
00061     int i;
00062     struct proc *curproc = myproc();
00063
00064     if(argint(n, &i) < 0)
00065         return -1;
00066     if(size < 0 || (uint)i >= curproc->sz || (uint)i+size > curproc->sz)
00067         return -1;
00068     *pp = (char*)i;
00069     return 0;
00070 }

```

Referenced by [sys_fstat\(\)](#), [sys_getprocs\(\)](#), [sys_pipe\(\)](#), [sys_read\(\)](#), and [sys_write\(\)](#).

5.25.2.6 argstr()

```
int argstr (
    int n,
    char ** pp )
```

Definition at line 77 of file [syscall.c](#).

```
00078 {
00079     int addr;
00080     if(argint(n, &addr) < 0)
00081         return -1;
00082     return fetchstr(addr, pp);
00083 }
```

Referenced by [sys_chdir\(\)](#), [sys_exec\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.25.2.7 begin_op()

```
void begin_op ( )
```

Definition at line 126 of file [log.c](#).

```
00127 {
00128     acquire(&log.lock);
00129     while(1){
00130         if(log.committing){
00131             sleep(&log, &log.lock);
00132         } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){
00133             // this op might exhaust log space; wait for commit.
00134             sleep(&log, &log.lock);
00135         } else {
00136             log.outstanding += 1;
00137             release(&log.lock);
00138             break;
00139         }
00140     }
00141 }
```

Referenced by [exec\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [filewrite\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.25.2.8 binit()

```
void binit (
    void )
```

Definition at line 39 of file [bio.c](#).

```
00040 {
00041     struct buf *b;
00042
00043     initlock(&bcache.lock, "bcache");
00044
00045     //PAGEBREAK!
00046     // Create linked list of buffers
00047     bcache.head.prev = &bcache.head;
00048     bcache.head.next = &bcache.head;
00049     for(b = bcache.buf; b < bcache.buf+NBUF; b++){
00050         b->next = bcache.head.next;
00051         b->prev = &bcache.head;
00052         initsleeplock(&b->lock, "buffer");
00053         bcache.head.next->prev = b;
00054         bcache.head.next = b;
00055     }
00056 }
```

5.25.2.9 bread()

```
struct buf * bread (
    uint dev,
    uint blockno )
```

Definition at line 97 of file [bio.c](#).

```
00098 {
00099     struct buf *b;
00100
00101     b = bget(dev, blockno);
00102     if((b->flags & B_VALID) == 0) {
00103         iderw(b);
00104     }
00105     return b;
00106 }
```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [ilock\(\)](#), [install_trans\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [read_head\(\)](#), [readi\(\)](#), [readsb\(\)](#), [write_head\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.25.2.10 brelse()

```
void brelse (
    struct buf * b )
```

Definition at line 121 of file [bio.c](#).

```
00122 {
00123     if(!holdingsleep(&b->lock))
00124         panic("brelse");
00125
00126     releasesleep(&b->lock);
00127
00128     acquire(&bcache.lock);
00129     b->refcnt--;
00130     if (b->refcnt == 0) {
00131         // no one is waiting for it.
00132         b->next->prev = b->prev;
00133         b->prev->next = b->next;
00134         b->next = bcache.head.next;
00135         b->prev = &bcache.head;
00136         bcache.head.next->prev = b;
00137         bcache.head.next = b;
00138     }
00139
00140     release(&bcache.lock);
00141 }
```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [ilock\(\)](#), [install_trans\(\)](#), [itrunc\(\)](#), [iupdate\(\)](#), [read_head\(\)](#), [readi\(\)](#), [readsb\(\)](#), [write_head\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.25.2.11 bwrite()

```
void bwrite (
    struct buf * b )
```

Definition at line 110 of file [bio.c](#).

```
00111 {
00112     if(!holdingsleep(&b->lock))
00113         panic("bwrite");
00114     b->flags |= B_DIRTY;
00115     iderw(b);
00116 }
```

Referenced by [install_trans\(\)](#), [write_head\(\)](#), and [write_log\(\)](#).

5.25.2.12 chpr()

```
int chpr (
    int pid,
    int priority )
```

Definition at line 559 of file [proc.c](#).

```
00560 {
00561     struct proc *p;
00562     acquire(&ptable.lock);
00563     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00564         if(p->pid == pid){
00565             p->priority = priority;
00566             break;
00567         }
00568     }
00569     release(&ptable.lock);
00570     return pid;
00571 }
```

Referenced by [main\(\)](#), and [sys_chpr\(\)](#).

5.25.2.13 clearpteu()

```
void clearpteu (
    pde_t * pgdir,
    char * uva )
```

Definition at line 303 of file [vm.c](#).

```
00304 {
00305     pte_t *pte;
00306
00307     pte = walkpgdir(pgdir, uva, 0);
00308     if(pte == 0)
00309         panic("clearpteu");
00310     *pte &= ~PTE_U;
00311 }
```

Referenced by [exec\(\)](#).

5.25.2.14 cmostime()

```
void cmostime (
    struct rtcdate * r )
```

Definition at line 196 of file [lapic.c](#).

```
00197 {
00198     struct rtcdate t1, t2;
00199     int sb, bcd;
00200
00201     sb = cmos_read(CMOS_STATB);
00202
00203     bcd = (sb & (1 << 2)) == 0;
00204
00205     // make sure CMOS doesn't modify time while we read it
00206     for(;;) {
00207         fill_rtcdate(&t1);
00208         if(cmos_read(CMOS_STATB) & CMOS_UIP)
00209             continue;
00210         fill_rtcdate(&t2);
00211         if(memcmp(&t1, &t2, sizeof(t1)) == 0)
00212             break;
00213     }
```

```

00214
00215 // convert
00216 if(bcd) {
00217 #define CONV(x)      ((t1.x >> 4) * 10) + (t1.x & 0xf)
00218     CONV(second);
00219     CONV(minute);
00220     CONV(hour );
00221     CONV(day );
00222     CONV(month );
00223     CONV(year );
00224 #undef CONV
00225 }
00226
00227 *r = t1;
00228 r->year += 2000;
00229 }

```

5.25.2.15 consoleinit()

```

void consoleinit (
    void )

```

Definition at line 289 of file [console.c](#).

```

00290 {
00291     initlock(&cons.lock, "console");
00292
00293     devsw[CONSOLE].write = consolewrite;
00294     devsw[CONSOLE].read = consoleread;
00295     cons.locking = 1;
00296
00297     ioapicenable(Irq_KBD, 0);
00298 }

```

5.25.2.16 consoleintr()

```

void consoleintr (
    int(*) (void) getc )

```

Definition at line 192 of file [console.c](#).

```

00193 {
00194     int c, doprocdump = 0;
00195
00196     acquire(&cons.lock);
00197     while((c = getc()) >= 0){
00198         switch(c){
00199             case C('P'): // Process listing.
00200                 // procdump() locks cons.lock indirectly; invoke later
00201                 doprocdump = 1;
00202                 break;
00203             case C('U'): // Kill line.
00204                 while(input.e != input.w &&
00205                     input.buf[(input.e-1) % INPUT_BUF] != '\n'){
00206                     input.e--;
00207                     consputc(BACKSPACE);
00208                 }
00209                 break;
00210             case C('H'): case '\x7f': // Backspace
00211                 if(input.e != input.w){
00212                     input.e--;
00213                     consputc(BACKSPACE);
00214                 }
00215                 break;
00216             default:
00217                 if(c != 0 && input.e-input.r < INPUT_BUF){
00218                     c = (c == '\r') ? '\n' : c;
00219                     input.buf[input.e++ % INPUT_BUF] = c;
00220                     consputc(c);
00221                     if(c == '\n' || c == C('D') || input.e == input.r+INPUT_BUF){
00222                         input.w = input.e;

```

```

00223         wakeup(&input.r);
00224     }
00225 }
00226     break;
00227 }
00228 }
00229 release(&cons.lock);
00230 if(doprocDump) {
00231     procdump(); // now call procdump() wo. cons.lock held
00232 }
00233 }

```

Referenced by [kbdintr\(\)](#), and [uartintr\(\)](#).

5.25.2.17 copyout()

```

int copyout (
    pde_t * pgdir,
    uint va,
    void * p,
    uint len )

```

Definition at line 366 of file [vm.c](#).

```

00367 {
00368     char *buf, *pa0;
00369     uint n, va0;
00370
00371     buf = (char*)p;
00372     while(len > 0){
00373         va0 = (uint)PGROUNDDOWN(va);
00374         pa0 = uva2ka(pgdir, (char*)va0);
00375         if(pa0 == 0)
00376             return -1;
00377         n = PGSIZE - (va - va0);
00378         if(n > len)
00379             n = len;
00380         memmove(pa0 + (va - va0), buf, n);
00381         len -= n;
00382         buf += n;
00383         va = va0 + PGSIZE;
00384     }
00385     return 0;
00386 }

```

Referenced by [exec\(\)](#).

5.25.2.18 copyvm()

```

pde_t * copyvm (
    pde_t * pgdir,
    uint sz )

```

Definition at line 316 of file [vm.c](#).

```

00317 {
00318     pde_t *d;
00319     pte_t *pte;
00320     uint pa, i, flags;
00321     char *mem;
00322
00323     if((d = setupkvm()) == 0)
00324         return 0;
00325     for(i = 0; i < sz; i += PGSIZE){
00326         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
00327             panic("copyvm: pte should exist");
00328         if(!(*pte & PTE_P))

```

```

00329     panic("copyvm: page not present");
00330     pa = PTE_ADDR(*pte);
00331     flags = PTE_FLAGS(*pte);
00332     if((mem = kalloc()) == 0)
00333         goto bad;
00334     memmove(mem, (char*)P2V(pa), PGSIZE);
00335     if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
00336         kfree(mem);
00337         goto bad;
00338     }
00339 }
00340 return d;
00341
00342 bad:
00343     freevm(d);
00344     return 0;
00345 }

```

Referenced by [fork\(\)](#).

5.25.2.19 cprintf()

```

void cprintf (
    char * fmt,
    ... )

```

Definition at line 55 of file [console.c](#).

```

00056 {
00057     int i, c, locking;
00058     uint *argp;
00059     char *s;
00060
00061     locking = cons.locking;
00062     if(locking)
00063         acquire(&cons.lock);
00064
00065     if (fmt == 0)
00066         panic("null fmt");
00067
00068     argp = (uint*)(void*)&fmt + 1;
00069     for(i = 0; (c = fmt[i] & 0xff) != 0; i++){
00070         if(c != '%'){
00071             consputc(c);
00072             continue;
00073         }
00074         c = fmt[++i] & 0xff;
00075         if(c == 0)
00076             break;
00077         switch(c){
00078             case 'd':
00079             printint(*argp++, 10, 1);
00080             break;
00081             case 'x':
00082             case 'p':
00083             printint(*argp++, 16, 0);
00084             break;
00085             case 's':
00086                 if((s = (char*)argp++) == 0)
00087                     s = "(null)";
00088                 for(; *s; s++)
00089                     consputc(*s);
00090             break;
00091             case '%':
00092                 consputc('%');
00093             break;
00094             default:
00095                 // Print unknown % sequence to draw attention.
00096                 consputc('%');
00097                 consputc(c);
00098             break;
00099         }
00100     }
00101
00102     if(locking)
00103         release(&cons.lock);
00104 }

```

Referenced by [allocvm\(\)](#), [cps\(\)](#), [exec\(\)](#), [iinit\(\)](#), [ioapicinit\(\)](#), [panic\(\)](#), [procdump\(\)](#), [syscall\(\)](#), and [trap\(\)](#).

5.25.2.20 cps()

```
int cps (
    void )
```

Definition at line 537 of file [proc.c](#).

```
00538 {
00539     struct proc *p;
00540     //Enables interrupts on this processor.
00541     sti();
00542
00543     //Loop over process table looking for process with pid.
00544     acquire(&ptable.lock);
00545     cprintf("name \t pid \t state \t priority \n");
00546     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00547         if(p->state == SLEEPING)
00548             cprintf("%s \t %d \t SLEEPING \t %d \n ", p->name, p->pid, p->priority);
00549         else if(p->state == RUNNING)
00550             cprintf("%s \t %d \t RUNNING \t %d \n ", p->name, p->pid, p->priority);
00551         else if(p->state == RUNNABLE)
00552             cprintf("%s \t %d \t RUNNABLE \t %d \n ", p->name, p->pid, p->priority);
00553     }
00554     release(&ptable.lock);
00555     return 23;
00556 }
```

Referenced by [main\(\)](#), and [sys_cps\(\)](#).

5.25.2.21 cpuid()

```
int cpuid (
    void )
```

Definition at line 31 of file [proc.c](#).

```
00031 {
00032     return mycpu() - cpus;
00033 }
```

Referenced by [seginit\(\)](#), and [trap\(\)](#).

5.25.2.22 deallocvm()

```
int deallocvm (
    pde_t * pgdir,
    uint oldsz,
    uint newsz )
```

Definition at line 256 of file [vm.c](#).

```
00257 {
00258     pte_t *pte;
00259     uint a, pa;
00260
00261     if(newsz >= oldsz)
00262         return oldsz;
00263
00264     a = PGROUNDUP(newsz);
00265     for(; a < oldsz; a += PGSIZE){
00266         pte = walkpgdir(pgdir, (char*)a, 0);
00267         if(!pte)
00268             a = PGADDR(PDX(a) + 1, 0, 0) - PGSIZE;
00269         else if((*pte & PTE_P) != 0){
00270             pa = PTE_ADDR(*pte);
00271             if(pa == 0)
```

```

00272         panic("kfree");
00273         char *v = P2V(pa);
00274         kfree(v);
00275         *pte = 0;
00276     }
00277 }
00278 return newsz;
00279 }

```

Referenced by [allocuvm\(\)](#), [freevm\(\)](#), and [growproc\(\)](#).

5.25.2.23 dirlink()

```

int dirlink (
    struct inode * dp,
    char * name,
    uint inum )

```

Definition at line 552 of file [fs.c](#).

```

00553 {
00554     int off;
00555     struct dirent de;
00556     struct inode *ip;
00557
00558     // Check that name is not present.
00559     if((ip = dirlookup(dp, name, 0)) != 0){
00560         iput(ip);
00561         return -1;
00562     }
00563
00564     // Look for an empty dirent.
00565     for(off = 0; off < dp->size; off += sizeof(de)){
00566         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00567             panic("dirlink read");
00568         if(de.inum == 0)
00569             break;
00570     }
00571
00572     strncpy(de.name, name, DIRSIZ);
00573     de.inum = inum;
00574     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00575         panic("dirlink");
00576
00577     return 0;
00578 }

```

Referenced by [create\(\)](#), and [sys_link\(\)](#).

5.25.2.24 dirlookup()

```

struct inode * dirlookup (
    struct inode * dp,
    char * name,
    uint * poff )

```

Definition at line 525 of file [fs.c](#).

```

00526 {
00527     uint off, inum;
00528     struct dirent de;
00529
00530     if(dp->type != T_DIR)
00531         panic("dirlookup not DIR");
00532
00533     for(off = 0; off < dp->size; off += sizeof(de)){
00534         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))

```

```

00535     panic("dirlookup read");
00536     if(de.inum == 0)
00537         continue;
00538     if(namecmp(name, de.name) == 0){
00539         // entry matches path element
00540         if(poff)
00541             *poff = off;
00542         inum = de.inum;
00543         return iget(dp->dev, inum);
00544     }
00545 }
00546
00547 return 0;
00548 }

```

Referenced by [create\(\)](#), [dirlink\(\)](#), [namex\(\)](#), and [sys_unlink\(\)](#).

5.25.2.25 end_op()

```
void end_op ( )
```

Definition at line 146 of file [log.c](#).

```

00147 {
00148     int do_commit = 0;
00149
00150     acquire(&log.lock);
00151     log.outstanding -= 1;
00152     if(log.committing)
00153         panic("log.committing");
00154     if(log.outstanding == 0){
00155         do_commit = 1;
00156         log.committing = 1;
00157     } else {
00158         // begin_op() may be waiting for log space,
00159         // and decrementing log.outstanding has decreased
00160         // the amount of reserved space.
00161         wakeup(&log);
00162     }
00163     release(&log.lock);
00164
00165     if(do_commit){
00166         // call commit w/o holding locks, since not allowed
00167         // to sleep with locks.
00168         commit();
00169         acquire(&log.lock);
00170         log.committing = 0;
00171         wakeup(&log);
00172         release(&log.lock);
00173     }
00174 }

```

Referenced by [exec\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [filewrite\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.25.2.26 exec()

```
int exec (
    char * path,
    char ** argv )
```

Definition at line 11 of file [exec.c](#).

```

00012 {
00013     char *s, *last;
00014     int i, off;
00015     uint argc, sz, sp, ustack[3+MAXARG+1];
00016     struct elfhdr elf;
00017     struct inode *ip;

```

```

00018 struct proghdr ph;
00019 pde_t *pgdir, *oldpgdir;
00020 struct proc *curproc = myproc();
00021
00022 begin_op();
00023
00024 if((ip = namei(path)) == 0){
00025     end_op();
00026     cprintf("exec: fail\n");
00027     return -1;
00028 }
00029 ilock(ip);
00030 pgdir = 0;
00031
00032 // Check ELF header
00033 if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
00034     goto bad;
00035 if(elf.magic != ELF_MAGIC)
00036     goto bad;
00037
00038 if((pgdir = setupkvm()) == 0)
00039     goto bad;
00040
00041 // Load program into memory.
00042 sz = 0;
00043 for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
00044     if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
00045         goto bad;
00046     if(ph.type != ELF_PROG_LOAD)
00047         continue;
00048     if(ph.memsz < ph.filesz)
00049         goto bad;
00050     if(ph.vaddr + ph.memsz < ph.vaddr)
00051         goto bad;
00052     if((sz = allocuvvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
00053         goto bad;
00054     if(ph.vaddr % PGSIZE != 0)
00055         goto bad;
00056     if(loaduvvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
00057         goto bad;
00058 }
00059 iunlockput(ip);
00060 end_op();
00061 ip = 0;
00062
00063 // Allocate two pages at the next page boundary.
00064 // Make the first inaccessible. Use the second as the user stack.
00065 sz = PGROUNDUP(sz);
00066 if((sz = allocuvvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
00067     goto bad;
00068 clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
00069 sp = sz;
00070
00071 // Push argument strings, prepare rest of stack in ustack.
00072 for(argc = 0; argv[argc]; argc++) {
00073     if(argc >= MAXARG)
00074         goto bad;
00075     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
00076     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
00077         goto bad;
00078     ustack[3+argc] = sp;
00079 }
00080 ustack[3+argc] = 0;
00081
00082 ustack[0] = 0xffffffff; // fake return PC
00083 ustack[1] = argc;
00084 ustack[2] = sp - (argc+1)*4; // argv pointer
00085
00086 sp -= (3+argc+1) * 4;
00087 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
00088     goto bad;
00089
00090 // Save program name for debugging.
00091 for(last=s=path; *s; s++)
00092     if(*s == '/')
00093         last = s+1;
00094 safestrcpy(curproc->name, last, sizeof(curproc->name));
00095
00096 // Commit to the user image.
00097 oldpgdir = curproc->pgdir;
00098 curproc->pgdir = pgdir;
00099 curproc->sz = sz;
00100 curproc->tf->eip = elf.entry; // main
00101 curproc->tf->esp = sp;
00102 switchvm(curproc);
00103 freevm(oldpgdir);
00104 return 0;

```

```

00105
00106 bad:
00107     if(pgdir)
00108         freevm(pgdir);
00109     if(ip) {
00110         iunlockput(ip);
00111         end_op();
00112     }
00113     return -1;
00114 }

```

Referenced by [bigargtest\(\)](#), [exectest\(\)](#), [main\(\)](#), [runcmd\(\)](#), and [sys_exec\(\)](#).

5.25.2.27 exit()

```

void exit (
    void )

```

Definition at line 228 of file [proc.c](#).

```

00229 {
00230     struct proc *curproc = myproc();
00231     struct proc *p;
00232     int fd;
00233
00234     if(curproc == initproc)
00235         panic("init exiting");
00236
00237     // Close all open files.
00238     for(fd = 0; fd < NOFILE; fd++){
00239         if(curproc->ofile[fd]){
00240             fileclose(curproc->ofile[fd]);
00241             curproc->ofile[fd] = 0;
00242         }
00243     }
00244
00245     begin_op();
00246     iput(curproc->cwd);
00247     end_op();
00248     curproc->cwd = 0;
00249
00250     acquire(&ptable.lock);
00251
00252     // Parent might be sleeping in wait().
00253     wakeup1(curproc->parent);
00254
00255     // Pass abandoned children to init.
00256     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00257         if(p->parent == curproc){
00258             p->parent = initproc;
00259             if(p->state == ZOMBIE)
00260                 wakeup1(initproc);
00261         }
00262     }
00263
00264     // Jump into the scheduler, never to return.
00265     curproc->state = ZOMBIE;
00266     sched();
00267     panic("zombie exit");
00268 }

```

Referenced by [argptest\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [bsstest\(\)](#), [cat\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [dirfile\(\)](#), [dirtest\(\)](#), [exectest\(\)](#), [exitputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [panic\(\)](#), [pipe1\(\)](#), [rmdot\(\)](#), [rsect\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [sys_exit\(\)](#), [trap\(\)](#), [uio\(\)](#), [unlinkread\(\)](#), [validatetest\(\)](#), [wc\(\)](#), [writetest\(\)](#), [writetest1\(\)](#), and [wsect\(\)](#).

5.25.2.28 fetchint()

```
int fetchint (
    uint addr,
    int * ip )
```

Definition at line 18 of file [syscall.c](#).

```
00019 {
00020     struct proc *curproc = myproc();
00021
00022     if(addr >= curproc->sz || addr+4 > curproc->sz)
00023         return -1;
00024     *ip = *(int*)(addr);
00025     return 0;
00026 }
```

Referenced by [argint\(\)](#), and [sys_exec\(\)](#).

5.25.2.29 fetchstr()

```
int fetchstr (
    uint addr,
    char ** pp )
```

Definition at line 32 of file [syscall.c](#).

```
00033 {
00034     char *s, *ep;
00035     struct proc *curproc = myproc();
00036
00037     if(addr >= curproc->sz)
00038         return -1;
00039     *pp = (char*)addr;
00040     ep = (char*)curproc->sz;
00041     for(s = *pp; s < ep; s++){
00042         if(*s == 0)
00043             return s - *pp;
00044     }
00045     return -1;
00046 }
```

Referenced by [argstr\(\)](#), and [sys_exec\(\)](#).

5.25.2.30 filealloc()

```
struct file * filealloc (
    void )
```

Definition at line 27 of file [file.c](#).

```
00028 {
00029     struct file *f;
00030
00031     acquire(&ftable.lock);
00032     for(f = ftable.file; f < ftable.file + NFILE; f++){
00033         if(f->ref == 0){
00034             f->ref = 1;
00035             release(&ftable.lock);
00036             return f;
00037         }
00038     }
00039     release(&ftable.lock);
00040     return 0;
00041 }
```

Referenced by [pipealloc\(\)](#), and [sys_open\(\)](#).

5.25.2.31 fclose()

```
void fclose (
    struct file * f )
```

Definition at line 57 of file [file.c](#).

```
00058 {
00059     struct file ff;
00060
00061     acquire(&ftable.lock);
00062     if(f->ref < 1)
00063         panic("fclose");
00064     if(--f->ref > 0){
00065         release(&ftable.lock);
00066         return;
00067     }
00068     ff = *f;
00069     f->ref = 0;
00070     f->type = FD_NONE;
00071     release(&ftable.lock);
00072
00073     if(ff.type == FD_PIPE)
00074         pipeclose(ff.pipe, ff.writable);
00075     else if(ff.type == FD_INODE){
00076         begin_op();
00077         iput(ff.ip);
00078         end_op();
00079     }
00080 }
```

Referenced by [exit\(\)](#), [pipealloc\(\)](#), [sys_close\(\)](#), [sys_open\(\)](#), and [sys_pipe\(\)](#).

5.25.2.32 fildup()

```
struct file * fildup (
    struct file * f )
```

Definition at line 45 of file [file.c](#).

```
00046 {
00047     acquire(&ftable.lock);
00048     if(f->ref < 1)
00049         panic("fildup");
00050     f->ref++;
00051     release(&ftable.lock);
00052     return f;
00053 }
```

Referenced by [fork\(\)](#), and [sys_dup\(\)](#).

5.25.2.33 fileinit()

```
void fileinit (
    void )
```

Definition at line 20 of file [file.c](#).

```
00021 {
00022     initlock(&ftable.lock, "ftable");
00023 }
```

5.25.2.34 fileread()

```
int fileread (
    struct file * f,
    char * addr,
    int n )
```

Definition at line 97 of file [file.c](#).

```
00098 {
00099     int r;
00100
00101     if(f->readable == 0)
00102         return -1;
00103     if(f->type == FD_PIPE)
00104         return piperead(f->pipe, addr, n);
00105     if(f->type == FD_INODE) {
00106         ilock(f->ip);
00107         if((r = readi(f->ip, addr, f->off, n)) > 0)
00108             f->off += r;
00109         iunlock(f->ip);
00110         return r;
00111     }
00112     panic("fileread");
00113 }
```

Referenced by [sys_read\(\)](#).

5.25.2.35 filestat()

```
int filestat (
    struct file * f,
    struct stat * st )
```

Definition at line 84 of file [file.c](#).

```
00085 {
00086     if(f->type == FD_INODE) {
00087         ilock(f->ip);
00088         stati(f->ip, st);
00089         iunlock(f->ip);
00090         return 0;
00091     }
00092     return -1;
00093 }
```

Referenced by [sys_fstat\(\)](#).

5.25.2.36 filewrite()

```
int filewrite (
    struct file * f,
    char * addr,
    int n )
```

Definition at line 118 of file [file.c](#).

```
00119 {
00120     int r;
00121
00122     if(f->writable == 0)
00123         return -1;
00124     if(f->type == FD_PIPE)
00125         return pipewrite(f->pipe, addr, n);
00126     if(f->type == FD_INODE) {
```



```

00127     // write a few blocks at a time to avoid exceeding
00128     // the maximum log transaction size, including
00129     // i-node, indirect block, allocation blocks,
00130     // and 2 blocks of slop for non-aligned writes.
00131     // this really belongs lower down, since writei()
00132     // might be writing a device like the console.
00133     int max = ((MAXOPBLOCKS-1-1-2) / 2) * 512;
00134     int i = 0;
00135     while(i < n){
00136         int n1 = n - i;
00137         if(n1 > max)
00138             n1 = max;
00139
00140         begin_op();
00141         ilock(f->ip);
00142         if ((r = writei(f->ip, addr + i, f->off, n1)) > 0)
00143             f->off += r;
00144         iunlock(f->ip);
00145         end_op();
00146
00147         if(r < 0)
00148             break;
00149         if(r != n1)
00150             panic("short filewrite");
00151         i += r;
00152     }
00153     return i == n ? n : -1;
00154 }
00155 panic("filewrite");
00156 }

```

Referenced by [sys_write\(\)](#).

5.25.2.37 fork()

```

int fork (
    void )

```

Definition at line 181 of file [proc.c](#).

```

00182 {
00183     int i, pid;
00184     struct proc *np;
00185     struct proc *curproc = myproc();
00186
00187     // Allocate process.
00188     if((np = allocproc()) == 0){
00189         return -1;
00190     }
00191
00192     // Copy process state from proc.
00193     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
00194         kfree(np->kstack);
00195         np->kstack = 0;
00196         np->state = UNUSED;
00197         return -1;
00198     }
00199     np->sz = curproc->sz;
00200     np->parent = curproc;
00201     *np->tf = *curproc->tf;
00202
00203     // Clear %eax so that fork returns 0 in the child.
00204     np->tf->eax = 0;
00205
00206     for(i = 0; i < NOFILE; i++)
00207         if(curproc->ofile[i])
00208             np->ofile[i] = filedup(curproc->ofile[i]);
00209     np->cwd = idup(curproc->cwd);
00210
00211     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
00212
00213     pid = np->pid;
00214
00215     acquire(&ptable.lock);
00216
00217     np->state = RUNNABLE;
00218
00219     release(&ptable.lock);

```

```

00220
00221     return pid;
00222 }

```

Referenced by [bigargtest\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [exitiputtest\(\)](#), [exitwait\(\)](#), [fork1\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [sys_fork\(\)](#), [uio\(\)](#), and [validateetest\(\)](#).

5.25.2.38 freevm()

```

void freevm (
    pde_t * pgdir )

```

Definition at line 284 of file [vm.c](#).

```

00285 {
00286     uint i;
00287
00288     if(pgdir == 0)
00289         panic("freevm: no pgdir");
00290     deallocuvvm(pgdir, KERNBASE, 0);
00291     for(i = 0; i < NPENTRIES; i++){
00292         if(pgdir[i] & PTE_P){
00293             char * v = P2V(PTE_ADDR(pgdir[i]));
00294             kfree(v);
00295         }
00296     }
00297     kfree((char*)pgdir);
00298 }

```

Referenced by [copyuvm\(\)](#), [exec\(\)](#), [setupkvm\(\)](#), and [wait\(\)](#).

5.25.2.39 getcallerpcs()

```

void getcallerpcs (
    void * ,
    uint * )

```

Referenced by [panic\(\)](#), and [procdump\(\)](#).

5.25.2.40 getprocs()

```

int getprocs (
    int max,
    struct uproc * )

```

Referenced by [main\(\)](#).

5.25.2.41 growproc()

```
int growproc (
    int n )
```

Definition at line 159 of file [proc.c](#).

```
00160 {
00161     uint sz;
00162     struct proc *curproc = myproc();
00163
00164     sz = curproc->sz;
00165     if(n > 0){
00166         if((sz = allocuvn(curproc->pgdir, sz, sz + n)) == 0)
00167             return -1;
00168     } else if(n < 0){
00169         if((sz = deallocuvn(curproc->pgdir, sz, sz + n)) == 0)
00170             return -1;
00171     }
00172     curproc->sz = sz;
00173     switchuvn(curproc);
00174     return 0;
00175 }
```

Referenced by [sys_sbrk\(\)](#).

5.25.2.42 holding()

```
int holding (
    struct spinlock * lock )
```

Definition at line 90 of file [spinlock.c](#).

```
00091 {
00092     int r;
00093     pushcli();
00094     r = lock->locked && lock->cpu == mycpu();
00095     popcli();
00096     return r;
00097 }
```

Referenced by [acquire\(\)](#), [release\(\)](#), and [sched\(\)](#).

5.25.2.43 holdingsleep()

```
int holdingsleep (
    struct sleeplock * lk )
```

Definition at line 45 of file [sleeplock.c](#).

```
00046 {
00047     int r;
00048
00049     acquire(&lk->lk);
00050     r = lk->locked && (lk->pid == myproc()->pid);
00051     release(&lk->lk);
00052     return r;
00053 }
```

Referenced by [brelse\(\)](#), [bwrite\(\)](#), [iderw\(\)](#), and [iunlock\(\)](#).

5.25.2.44 ialloc()

```
struct inode * ialloc (
    uint dev,
    short type )
```

Definition at line 195 of file [fs.c](#).

```
00196 {
00197     int inum;
00198     struct buf *bp;
00199     struct dinode *dip;
00200
00201     for(inum = 1; inum < sb.ninodes; inum++){
00202         bp = bread(dev, IBLOCK(inum, sb));
00203         dip = (struct dinode*)bp->data + inum%IPB;
00204         if(dip->type == 0){ // a free inode
00205             memset(dip, 0, sizeof(*dip));
00206             dip->type = type;
00207             log_write(bp); // mark it allocated on the disk
00208             brelse(bp);
00209             return iget(dev, inum);
00210         }
00211         brelse(bp);
00212     }
00213     panic("ialloc: no inodes");
00214 }
```

Referenced by [create\(\)](#).

5.25.2.45 ideinit()

```
void ideinit (
    void )
```

Definition at line 51 of file [ide.c](#).

```
00052 {
00053     int i;
00054
00055     initlock(&idelock, "ide");
00056     ioapicenable(IRQ_IDE, ncpu - 1);
00057     idewait(0);
00058
00059     // Check if disk 1 is present
00060     outb(0x1f6, 0xe0 | (1<<4));
00061     for(i=0; i<1000; i++){
00062         if(inb(0x1f7) != 0){
00063             havedisk1 = 1;
00064             break;
00065         }
00066     }
00067
00068     // Switch back to disk 0.
00069     outb(0x1f6, 0xe0 | (0<<4));
00070 }
```

5.25.2.46 ideintr()

```
void ideintr (
    void )
```

Definition at line 104 of file [ide.c](#).

```
00105 {
00106     struct buf *b;
00107 }
```

```

00108 // First queued buffer is the active request.
00109 acquire(&idelock);
00110
00111 if((b = idequeue) == 0){
00112     release(&idelock);
00113     return;
00114 }
00115 idequeue = b->qnext;
00116
00117 // Read data if needed.
00118 if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
00119     insl(0x1f0, b->data, BSIZE/4);
00120
00121 // Wake process waiting for this buf.
00122 b->flags |= B_VALID;
00123 b->flags &= ~B_DIRTY;
00124 wakeup(b);
00125
00126 // Start disk on next buf in queue.
00127 if(idequeue != 0)
00128     idestart(idequeue);
00129
00130 release(&idelock);
00131 }

```

Referenced by [trap\(\)](#).

5.25.2.47 iderw()

```

void iderw (
    struct buf * b )

```

Definition at line 138 of file [ide.c](#).

```

00139 {
00140     struct buf **pp;
00141
00142     if(!holdingsleep(&b->lock))
00143         panic("iderw: buf not locked");
00144     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
00145         panic("iderw: nothing to do");
00146     if(b->dev != 0 && !havedisk1)
00147         panic("iderw: ide disk 1 not present");
00148
00149     acquire(&idelock); //DOC:acquire-lock
00150
00151     // Append b to idequeue.
00152     b->qnext = 0;
00153     for(pp=&idequeue; *pp; pp=&(*pp)->qnext) //DOC:insert-queue
00154         ;
00155     *pp = b;
00156
00157     // Start disk if necessary.
00158     if(idequeue == b)
00159         idestart(b);
00160
00161     // Wait for request to finish.
00162     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
00163         sleep(b, &idelock);
00164     }
00165
00166     release(&idelock);
00167 }
00168 }

```

Referenced by [bread\(\)](#), and [bwrite\(\)](#).

5.25.2.48 idtinit()

```
void idtinit (
    void )
```

Definition at line 30 of file [trap.c](#).

```
00031 {
00032     lidt(idt, sizeof(idt));
00033 }
```

5.25.2.49 idup()

```
struct inode * idup (
    struct inode * ip )
```

Definition at line 277 of file [fs.c](#).

```
00278 {
00279     acquire(&icache.lock);
00280     ip->ref++;
00281     release(&icache.lock);
00282     return ip;
00283 }
```

Referenced by [fork\(\)](#), and [namex\(\)](#).

5.25.2.50 iinit()

```
void iinit (
    int dev )
```

Definition at line 172 of file [fs.c](#).

```
00173 {
00174     int i = 0;
00175
00176     initlock(&icache.lock, "icache");
00177     for(i = 0; i < NINODE; i++) {
00178         initsleeplock(&icache.inode[i].lock, "inode");
00179     }
00180
00181     readsb(dev, &sb);
00182     cprintf("sb: size %d nblocks %d ninodes %d nlog %d logstart %d\
00183 inodestart %d bmap start %d\n", sb.size, sb.nblocks,
00184         sb.ninodes, sb.nlog, sb.logstart, sb.inodestart,
00185         sb.bmapstart);
00186 }
```

Referenced by [forkret\(\)](#).

5.25.2.51 ilock()

```
void ilock (
    struct inode * ip )
```

Definition at line 288 of file [fs.c](#).

```
00289 {
00290     struct buf *bp;
00291     struct dinode *dip;
00292
00293     if(ip == 0 || ip->ref < 1)
00294         panic("ilock");
00295
00296     acquiresleep(&ip->lock);
00297
00298     if(ip->valid == 0){
00299         bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00300         dip = (struct dinode*)bp->data + ip->inum%IPB;
00301         ip->type = dip->type;
00302         ip->major = dip->major;
00303         ip->minor = dip->minor;
00304         ip->nlink = dip->nlink;
00305         ip->size = dip->size;
00306         memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
00307         brelse(bp);
00308         ip->valid = 1;
00309         if(ip->type == 0)
00310             panic("ilock: no type");
00311     }
00312 }
```

Referenced by [consoleread\(\)](#), [consolewrite\(\)](#), [create\(\)](#), [exec\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.25.2.52 initlock()

```
void initlock (
    struct spinlock * lk,
    char * name )
```

Definition at line 13 of file [spinlock.c](#).

```
00014 {
00015     lk->name = name;
00016     lk->locked = 0;
00017     lk->cpu = 0;
00018 }
```

Referenced by [binit\(\)](#), [consoleinit\(\)](#), [fileinit\(\)](#), [ideinit\(\)](#), [iinit\(\)](#), [initlog\(\)](#), [initsleeplock\(\)](#), [kinit1\(\)](#), [pinit\(\)](#), [pipealloc\(\)](#), and [tvinit\(\)](#).

5.25.2.53 initlog()

```
void initlog (
    int dev )
```

Definition at line 54 of file [log.c](#).

```
00055 {
00056     if (sizeof(struct logheader) >= BSIZE)
00057         panic("initlog: too big logheader");
00058
00059     struct superblock sb;
00060     initlock(&log.lock, "log");
00061     readsb(dev, &sb);
00062     log.start = sb.logstart;
00063     log.size = sb.nlog;
00064     log.dev = dev;
00065     recover_from_log();
00066 }
```

Referenced by [forkret\(\)](#).

5.25.2.54 initsleeplock()

```
void initsleeplock (
    struct sleeplock * lk,
    char * name )
```

Definition at line 14 of file [sleeplock.c](#).

```
00015 {
00016     initlock(&lk->lk, "sleep lock");
00017     lk->name = name;
00018     lk->locked = 0;
00019     lk->pid = 0;
00020 }
```

Referenced by [binit\(\)](#), and [iinit\(\)](#).

5.25.2.55 initvm()

```
void initvm (
    pde_t * pgdir,
    char * init,
    uint sz )
```

Definition at line 183 of file [vm.c](#).

```
00184 {
00185     char *mem;
00186
00187     if(sz >= PGSIZE)
00188         panic("initvm: more than a page");
00189     mem = kalloc();
00190     memset(mem, 0, PGSIZE);
00191     mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_W|PTE_U);
00192     memmove(mem, init, sz);
00193 }
```

Referenced by [userinit\(\)](#).

5.25.2.56 ioapicenable()

```
void ioapicenable (
    int irq,
    int cpu )
```

Definition at line 68 of file [ioapic.c](#).

```
00069 {
00070     // Mark interrupt edge-triggered, active high,
00071     // enabled, and routed to the given cpunum,
00072     // which happens to be that cpu's APIC ID.
00073     ioapicwrite(REG_TABLE+2*irq, T_IRQ0 + irq);
00074     ioapicwrite(REG_TABLE+2*irq+1, cpunum << 24);
00075 }
```

Referenced by [consoleinit\(\)](#), [ideinit\(\)](#), and [uartinit\(\)](#).

5.25.2.57 ioapicinit()

```
void ioapicinit (
    void )
```

Definition at line 49 of file [ioapic.c](#).

```
00050 {
00051     int i, id, maxintr;
00052
00053     ioapic = (volatile struct ioapic*)IOAPIC;
00054     maxintr = (ioapicread(REG_VER) » 16) & 0xFF;
00055     id = ioapicread(REG_ID) » 24;
00056     if(id != ioapicid)
00057         cprintf("ioapicinit: id isn't equal to ioapicid; not a MP\n");
00058
00059     // Mark all interrupts edge-triggered, active high, disabled,
00060     // and not routed to any CPUs.
00061     for(i = 0; i <= maxintr; i++){
00062         ioapicwrite(REG_TABLE+2*i, INT_DISABLED | (T_IRQ0 + i));
00063         ioapicwrite(REG_TABLE+2*i+1, 0);
00064     }
00065 }
```

5.25.2.58 iput()

```
void iput (
    struct inode * ip )
```

Definition at line 332 of file [fs.c](#).

```
00333 {
00334     acquiresleep(&ip->lock);
00335     if(ip->valid && ip->nlink == 0){
00336         acquire(&icache.lock);
00337         int r = ip->ref;
00338         release(&icache.lock);
00339         if(r == 1){
00340             // inode has no links and no other references: truncate and free.
00341             itrunc(ip);
00342             ip->type = 0;
00343             iupdate(ip);
00344             ip->valid = 0;
00345         }
00346     }
00347     releasesleep(&ip->lock);
00348
00349     acquire(&icache.lock);
00350     ip->ref--;
00351     release(&icache.lock);
00352 }
```

Referenced by [dirlink\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [iunlockput\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), and [sys_link\(\)](#).

5.25.2.59 iunlock()

```
void iunlock (
    struct inode * ip )
```

Definition at line 316 of file [fs.c](#).

```
00317 {
00318     if(ip == 0 || !holdingsleep(&ip->lock) || ip->ref < 1)
00319         panic("iunlock");
00320
00321     releasesleep(&ip->lock);
00322 }
```

Referenced by [consoleread\(\)](#), [consolewrite\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), [iunlockput\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), and [sys_open\(\)](#).

5.25.2.60 iunlockput()

```
void iunlockput (
    struct inode * ip )
```

Definition at line 356 of file [fs.c](#).

```
00357 {
00358     iunlock(ip);
00359     iput(ip);
00360 }
```

Referenced by [create\(\)](#), [exec\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.25.2.61 iupdate()

```
void iupdate (
    struct inode * ip )
```

Definition at line 221 of file [fs.c](#).

```
00222 {
00223     struct buf *bp;
00224     struct dinode *dip;
00225
00226     bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00227     dip = (struct dinode*)bp->data + ip->inum%IPB;
00228     dip->type = ip->type;
00229     dip->major = ip->major;
00230     dip->minor = ip->minor;
00231     dip->nlink = ip->nlink;
00232     dip->size = ip->size;
00233     memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
00234     log_write(bp);
00235     brelse(bp);
00236 }
```

Referenced by [create\(\)](#), [iput\(\)](#), [itrunc\(\)](#), [sys_link\(\)](#), [sys_unlink\(\)](#), and [writei\(\)](#).

5.25.2.62 kalloc()

```
char * kalloc (
    void )
```

Definition at line 83 of file [kalloc.c](#).

```
00084 {
00085     struct run *r;
00086
00087     if(kmem.use_lock)
00088         acquire(&kmem.lock);
00089     r = kmem.freelist;
00090     if(r)
00091         kmem.freelist = r->next;
00092     if(kmem.use_lock)
00093         release(&kmem.lock);
00094     return (char*)r;
00095 }
```

Referenced by [allocproc\(\)](#), [allocuvm\(\)](#), [copyuvm\(\)](#), [inituvm\(\)](#), [pipealloc\(\)](#), [setupkvm\(\)](#), [startothers\(\)](#), and [walkpgdir\(\)](#).

5.25.2.63 kbdintr()

```
void kbdintr (
    void )
```

Definition at line 47 of file [kbd.c](#).

```
00048 {
00049     consoleintr(kbdgetc);
00050 }
```

Referenced by [trap\(\)](#).

5.25.2.64 kfree()

```
void kfree (
    char * v )
```

Definition at line 60 of file [kalloc.c](#).

```
00061 {
00062     struct run *r;
00063
00064     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
00065         panic("kfree");
00066
00067     // Fill with junk to catch dangling refs.
00068     memset(v, 1, PGSIZE);
00069
00070     if(kmem.use_lock)
00071         acquire(&kmem.lock);
00072     r = (struct run*)v;
00073     r->next = kmem.freelist;
00074     kmem.freelist = r;
00075     if(kmem.use_lock)
00076         release(&kmem.lock);
00077 }
```

Referenced by [allocuvn\(\)](#), [copyuvn\(\)](#), [deallocuvn\(\)](#), [fork\(\)](#), [freerange\(\)](#), [freevm\(\)](#), [pipealloc\(\)](#), [pipeclose\(\)](#), and [wait\(\)](#).

5.25.2.65 kill()

```
int kill (
    int pid )
```

Definition at line 480 of file [proc.c](#).

```
00481 {
00482     struct proc *p;
00483
00484     acquire(&ptable.lock);
00485     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
00486         if(p->pid == pid) {
00487             p->killed = 1;
00488             // Wake process from sleep if necessary.
00489             if(p->state == SLEEPING)
00490                 p->state = RUNNABLE;
00491             release(&ptable.lock);
00492             return 0;
00493         }
00494     }
00495     release(&ptable.lock);
00496     return -1;
00497 }
```

Referenced by [main\(\)](#), [mem\(\)](#), [preempt\(\)](#), [sbrktest\(\)](#), [sys_kill\(\)](#), and [validatetest\(\)](#).

5.25.2.66 kinit1()

```
void kinit1 (
    void * vstart,
    void * vend )
```

Definition at line 32 of file [kalloc.c](#).

```
00033 {
00034     initlock(&kmem.lock, "kmem");
00035     kmem.use_lock = 0;
00036     freerange(vstart, vend);
00037 }
```

5.25.2.67 kinit2()

```
void kinit2 (
    void * vstart,
    void * vend )
```

Definition at line 40 of file [kalloc.c](#).

```
00041 {
00042     freerange(vstart, vend);
00043     kmem.use_lock = 1;
00044 }
```

5.25.2.68 kvmalloc()

```
void kvmalloc (
    void )
```

Definition at line 141 of file [vm.c](#).

```
00142 {
00143     kpgdir = setupkvm();
00144     switchkvm();
00145 }
```

5.25.2.69 lapiceoi()

```
void lapiceoi (
    void )
```

Definition at line 110 of file [lapic.c](#).

```
00111 {
00112     if(lapic)
00113         lapicw(EOI, 0);
00114 }
```

Referenced by [trap\(\)](#).

5.25.2.70 lapicid()

```
int lapicid (
    void )
```

Definition at line 101 of file [lapic.c](#).

```
00102 {
00103     if (!lapic)
00104         return 0;
00105     return lapic[ID] » 24;
00106 }
```

Referenced by [mycpu\(\)](#), and [panic\(\)](#).

5.25.2.71 lapicinit()

```
void lapicinit (
    void )
```

Definition at line 55 of file [lapic.c](#).

```
00056 {
00057     if (!lapic)
00058         return;
00059
00060     // Enable local APIC; set spurious interrupt vector.
00061     lapicw(SVR, ENABLE | (T_IRQ0 + IRQ_SPURIOUS));
00062
00063     // The timer repeatedly counts down at bus frequency
00064     // from lapic[TICR] and then issues an interrupt.
00065     // If xv6 cared more about precise timekeeping,
00066     // TICR would be calibrated using an external time source.
00067     lapicw(TDCR, X1);
00068     lapicw(TIMER, PERIODIC | (T_IRQ0 + IRQ_TIMER));
00069     lapicw(TICR, 10000000);
00070
00071     // Disable logical interrupt lines.
00072     lapicw(LINT0, MASKED);
00073     lapicw(LINT1, MASKED);
00074
00075     // Disable performance counter overflow interrupts
00076     // on machines that provide that interrupt entry.
00077     if (((lapic[VER] » 16) & 0xFF) >= 4)
00078         lapicw(PCINT, MASKED);
00079
00080     // Map error interrupt to IRQ_ERROR.
00081     lapicw(ERROR, T_IRQ0 + IRQ_ERROR);
00082
00083     // Clear error status register (requires back-to-back writes).
00084     lapicw(ESR, 0);
00085     lapicw(ESR, 0);
00086
00087     // Ack any outstanding interrupts.
00088     lapicw(EOI, 0);
00089
00090     // Send an Init Level De-Assert to synchronise arbitration ID's.
00091     lapicw(ICRHI, 0);
00092     lapicw(ICRLO, BCAST | INIT | LEVEL);
00093     while(lapic[ICRLO] & DELIVS)
00094         ;
00095
00096     // Enable interrupts on the APIC (but not on the processor).
00097     lapicw(TPR, 0);
00098 }
```

Referenced by [mpenter\(\)](#).

5.25.2.72 lapicstartap()

```
void lapicstartap (
    uchar apicid,
    uint addr )
```

Definition at line 129 of file [lapic.c](#).

```
00130 {
00131     int i;
00132     ushort *wrv;
00133
00134     // "The BSP must initialize CMOS shutdown code to 0AH
00135     // and the warm reset vector (DWORD based at 40:67) to point at
00136     // the AP startup code prior to the [universal startup algorithm]."
00137     outb(CMOS_PORT, 0xF); // offset 0xF is shutdown code
00138     outb(CMOS_PORT+1, 0x0A);
00139     wrv = (ushort*)P2V((0x40«4 | 0x67)); // Warm reset vector
00140     wrv[0] = 0;
00141     wrv[1] = addr » 4;
00142
00143     // "Universal startup algorithm."
00144     // Send INIT (level-triggered) interrupt to reset other CPU.
00145     lapicw(ICRHI, apicid«24);
00146     lapicw(ICRLO, INIT | LEVEL | ASSERT);
00147     microdelay(200);
00148     lapicw(ICRLO, INIT | LEVEL);
00149     microdelay(100); // should be 10ms, but too slow in Bochs!
00150
00151     // Send startup IPI (twice!) to enter code.
00152     // Regular hardware is supposed to only accept a STARTUP
00153     // when it is in the halted state due to an INIT. So the second
00154     // should be ignored, but it is part of the official Intel algorithm.
00155     // Bochs complains about the second one. Too bad for Bochs.
00156     for(i = 0; i < 2; i++){
00157         lapicw(ICRHI, apicid«24);
00158         lapicw(ICRLO, STARTUP | (addr»12));
00159         microdelay(200);
00160     }
00161 }
```

Referenced by [startothers\(\)](#).

5.25.2.73 loaduvm()

```
int loaduvm (
    pde_t * pgdir,
    char * addr,
    struct inode * ip,
    uint offset,
    uint sz )
```

Definition at line 198 of file [vm.c](#).

```
00199 {
00200     uint i, pa, n;
00201     pte_t *pte;
00202
00203     if((uint) addr % PGSIZE != 0)
00204         panic("loaduvm: addr must be page aligned");
00205     for(i = 0; i < sz; i += PGSIZE){
00206         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
00207             panic("loaduvm: address should exist");
00208         pa = PTE_ADDR(*pte);
00209         if(sz - i < PGSIZE)
00210             n = sz - i;
00211         else
00212             n = PGSIZE;
00213         if(readi(ip, P2V(pa), offset+i, n) != n)
00214             return -1;
00215     }
00216     return 0;
00217 }
```

Referenced by [exec\(\)](#).

5.25.2.74 log_write()

```
void log_write (
    struct buf * b )
```

Definition at line 214 of file [log.c](#).

```
00215 {
00216     int i;
00217
00218     if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
00219         panic("too big a transaction");
00220     if (log.outstanding < 1)
00221         panic("log_write outside of trans");
00222
00223     acquire(&log.lock);
00224     for (i = 0; i < log.lh.n; i++) {
00225         if (log.lh.block[i] == b->blockno)    // log absorbtion
00226             break;
00227     }
00228     log.lh.block[i] = b->blockno;
00229     if (i == log.lh.n)
00230         log.lh.n++;
00231     b->flags |= B_DIRTY; // prevent eviction
00232     release(&log.lock);
00233 }
```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [iupdate\(\)](#), and [writei\(\)](#).

5.25.2.75 memcmp()

```
int memcmp (
    const void * v1,
    const void * v2,
    uint n )
```

Definition at line 16 of file [string.c](#).

```
00017 {
00018     const uchar *s1, *s2;
00019
00020     s1 = v1;
00021     s2 = v2;
00022     while (n-- > 0) {
00023         if (*s1 != *s2)
00024             return *s1 - *s2;
00025         s1++, s2++;
00026     }
00027
00028     return 0;
00029 }
```

Referenced by [cmostime\(\)](#), [mpconfig\(\)](#), and [mpsearch1\(\)](#).

5.25.2.76 memmove()

```
void * memmove (
    void * dst,
    const void * src,
    uint n )
```

Definition at line 32 of file [string.c](#).

```
00033 {
00034     const char *s;
```

```

00035  char *d;
00036
00037  s = src;
00038  d = dst;
00039  if (s < d && s + n > d) {
00040      s += n;
00041      d += n;
00042      while (n-- > 0)
00043          *--d = *--s;
00044  } else
00045      while (n-- > 0)
00046          *d++ = *s++;
00047
00048  return dst;
00049 }

```

Referenced by [cgaputc\(\)](#), [copyout\(\)](#), [copyuvm\(\)](#), [fmtname\(\)](#), [grep\(\)](#), [iderw\(\)](#), [ilock\(\)](#), [inituvm\(\)](#), [install_trans\(\)](#), [iupdate\(\)](#), [ls\(\)](#), [main\(\)](#), [memcpy\(\)](#), [readi\(\)](#), [readsb\(\)](#), [skipelem\(\)](#), [startothers\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.25.2.77 `memset()`

```

void * memset (
    void * dst,
    int c,
    uint n )

```

Definition at line 5 of file [string.c](#).

```

00006 {
00007     if ((int)dst%4 == 0 && n%4 == 0) {
00008         c &= 0xFF;
00009         stosl(dst, (c<<24) | (c<<16) | (c<<8) | c, n/4);
00010     } else
00011         stosb(dst, c, n);
00012     return dst;
00013 }

```

Referenced by [allocproc\(\)](#), [allocuvm\(\)](#), [backcmd\(\)](#), [bigfile\(\)](#), [bzero\(\)](#), [cgaputc\(\)](#), [concreate\(\)](#), [execcmd\(\)](#), [fmtname\(\)](#), [fourfiles\(\)](#), [getcmd\(\)](#), [ialloc\(\)](#), [inituvm\(\)](#), [kfree\(\)](#), [listcmd\(\)](#), [main\(\)](#), [pipecmd\(\)](#), [redircmd\(\)](#), [setupkvm\(\)](#), [sharedfd\(\)](#), [sys_exec\(\)](#), [sys_unlink\(\)](#), [userinit\(\)](#), and [walkpgdir\(\)](#).

5.25.2.78 `microdelay()`

```

void microdelay (
    int us )

```

Definition at line 119 of file [lapic.c](#).

```

00120 {
00121 }

```

Referenced by [cmos_read\(\)](#), [lapicstartap\(\)](#), and [uartputc\(\)](#).

5.25.2.79 mpinit()

```
void mpinit (
    void )
```

Definition at line 92 of file [mp.c](#).

```
00093 {
00094     uchar *p, *e;
00095     int ismp;
00096     struct mp *mp;
00097     struct mpconf *conf;
00098     struct mpproc *proc;
00099     struct mpioapic *ioapic;
00100
00101     if((conf = mpconfig(&mp)) == 0)
00102         panic("Expect to run on an SMP");
00103     ismp = 1;
00104     lapic = (uint*)conf->lapicaddr;
00105     for(p=(uchar*)(conf+1), e=(uchar*)conf+conf->length; p<e; ){
00106         switch(*p){
00107             case MPPROC:
00108                 proc = (struct mpproc*)p;
00109                 if(ncpu < NCPU) {
00110                     cpus[ncpu].apicid = proc->apicid; // apicid may differ from ncpu
00111                     ncpu++;
00112                 }
00113                 p += sizeof(struct mpproc);
00114                 continue;
00115             case MPIOAPIC:
00116                 ioapic = (struct mpioapic*)p;
00117                 ioapicid = ioapic->apicno;
00118                 p += sizeof(struct mpioapic);
00119                 continue;
00120             case MPBUS:
00121             case MPIINTR:
00122             case MPLINTR:
00123                 p += 8;
00124                 continue;
00125             default:
00126                 ismp = 0;
00127                 break;
00128         }
00129     }
00130     if(!ismp)
00131         panic("Didn't find a suitable machine");
00132
00133     if(mp->imcrp){
00134         // Bochs doesn't support IMCR, so this doesn't run on Bochs.
00135         // But it would on real hardware.
00136         outb(0x22, 0x70); // Select IMCR
00137         outb(0x23, inb(0x23) | 1); // Mask external interrupts.
00138     }
00139 }
```

5.25.2.80 mycpu()

```
struct cpu * mycpu (
    void )
```

Definition at line 38 of file [proc.c](#).

```
00039 {
00040     int apicid, i;
00041
00042     if(readeflags() & FL_IF)
00043         panic("mycpu called with interrupts enabled\n");
00044
00045     apicid = lapicid();
00046     // APIC IDs are not guaranteed to be contiguous. Maybe we should have
00047     // a reverse map, or reserve a register to store &cpus[i].
00048     for (i = 0; i < ncpu; ++i) {
00049         if (cpus[i].apicid == apicid)
00050             return &cpus[i];
00051     }
00052     panic("unknown apicid\n");
00053 }
```

Referenced by [acquire\(\)](#), [cpuid\(\)](#), [holding\(\)](#), [myproc\(\)](#), [popcli\(\)](#), [pushcli\(\)](#), [sched\(\)](#), [scheduler\(\)](#), [startothers\(\)](#), and [switchvm\(\)](#).

5.25.2.81 myproc()

```
struct proc * myproc ( )
```

Definition at line 58 of file [proc.c](#).

```
00058 {
00059     struct cpu *c;
00060     struct proc *p;
00061     pushcli();
00062     c = mycpu();
00063     p = c->proc;
00064     popcli();
00065     return p;
00066 }
```

Referenced by [acquiresleep\(\)](#), [argfd\(\)](#), [argint\(\)](#), [argptr\(\)](#), [consoleread\(\)](#), [exec\(\)](#), [exit\(\)](#), [fdalloc\(\)](#), [fetchint\(\)](#), [fetchstr\(\)](#), [fork\(\)](#), [growproc\(\)](#), [holdingsleep\(\)](#), [namex\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [sched\(\)](#), [sleep\(\)](#), [sys_chdir\(\)](#), [sys_close\(\)](#), [sys_getpid\(\)](#), [sys_pipe\(\)](#), [sys_sbrk\(\)](#), [sys_sleep\(\)](#), [syscall\(\)](#), [trap\(\)](#), [wait\(\)](#), and [yield\(\)](#).

5.25.2.82 namecmp()

```
int namecmp (
    const char * s,
    const char * t )
```

Definition at line 517 of file [fs.c](#).

```
00518 {
00519     return strcmp(s, t, DIRSIZ);
00520 }
```

Referenced by [dirlookup\(\)](#), and [sys_unlink\(\)](#).

5.25.2.83 namei()

```
struct inode * namei (
    char * path )
```

Definition at line 660 of file [fs.c](#).

```
00661 {
00662     char name[DIRSIZ];
00663     return namex(path, 0, name);
00664 }
```

Referenced by [exec\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_open\(\)](#), and [userinit\(\)](#).

5.25.2.84 nameiparent()

```
struct inode * nameiparent (
    char * path,
    char * name )
```

Definition at line 667 of file [fs.c](#).

```
00668 {
00669     return namex(path, 1, name);
00670 }
```

Referenced by [create\(\)](#), [namex\(\)](#), [sys_link\(\)](#), and [sys_unlink\(\)](#).

5.25.2.85 panic()

```
void panic (
    char * s )
```

Definition at line 107 of file [console.c](#).

```
00108 {
00109     int i;
00110     uint pcs[10];
00111
00112     cli();
00113     cons.locking = 0;
00114     // use lapiccpunum so that we can call panic from mycpu()
00115     cprintf("lapicid %d: panic: ", lapicid());
00116     cprintf(s);
00117     cprintf("\n");
00118     getcallerpcs(&s, pcs);
00119     for(i=0; i<10; i++)
00120         cprintf(" %p", pcs[i]);
00121     panicked = 1; // freeze other CPU
00122     for(;;)
00123         ;
00124 }
```

Referenced by [acquire\(\)](#), [balloc\(\)](#), [bfree\(\)](#), [bget\(\)](#), [bmap\(\)](#), [brelse\(\)](#), [bwrite\(\)](#), [cgaputc\(\)](#), [clearpteu\(\)](#), [copyvm\(\)](#), [cprintf\(\)](#), [create\(\)](#), [deallocvm\(\)](#), [dirlink\(\)](#), [dirlookup\(\)](#), [end_op\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fileread\(\)](#), [filewrite\(\)](#), [fork1\(\)](#), [freevm\(\)](#), [ialloc\(\)](#), [iderw\(\)](#), [idestart\(\)](#), [iget\(\)](#), [ilock\(\)](#), [initlog\(\)](#), [initvm\(\)](#), [isdirempty\(\)](#), [iunlock\(\)](#), [kfree\(\)](#), [loadvm\(\)](#), [log_write\(\)](#), [mappages\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), [parseblock\(\)](#), [parsecmd\(\)](#), [parseexec\(\)](#), [parseredirs\(\)](#), [popcli\(\)](#), [release\(\)](#), [runcmd\(\)](#), [sched\(\)](#), [setupkvm\(\)](#), [sleep\(\)](#), [switchvm\(\)](#), [sys_unlink\(\)](#), [trap\(\)](#), and [userinit\(\)](#).

5.25.2.86 picenable()

```
void picenable (
    int )
```

5.25.2.87 picinit()

```
void picinit (
    void )
```

Definition at line 11 of file [picirq.c](#).

```
00012 {
00013     // mask all interrupts
00014     outb(IO_PIC1+1, 0xFF);
00015     outb(IO_PIC2+1, 0xFF);
00016 }
```

5.25.2.88 pinit()

```
void pinit (
    void )
```

Definition at line 24 of file [proc.c](#).

```
00025 {
00026     initlock(&ptable.lock, "ptable");
00027 }
```

5.25.2.89 pipealloc()

```
int pipealloc (
    struct file ** f0,
    struct file ** f1 )
```

Definition at line 23 of file [pipe.c](#).

```
00024 {
00025     struct pipe *p;
00026
00027     p = 0;
00028     *f0 = *f1 = 0;
00029     if ((*f0 = filealloc()) == 0 || (*f1 = filealloc()) == 0)
00030         goto bad;
00031     if ((p = (struct pipe*)kalloc()) == 0)
00032         goto bad;
00033     p->readopen = 1;
00034     p->writeopen = 1;
00035     p->nwrite = 0;
00036     p->nread = 0;
00037     initlock(&p->lock, "pipe");
00038     (*f0)->type = FD_PIPE;
00039     (*f0)->readable = 1;
00040     (*f0)->writable = 0;
00041     (*f0)->pipe = p;
00042     (*f1)->type = FD_PIPE;
00043     (*f1)->readable = 0;
00044     (*f1)->writable = 1;
00045     (*f1)->pipe = p;
00046     return 0;
00047
00048 //PAGEBREAK: 20
00049 bad:
00050     if (p)
00051         kfree((char*)p);
00052     if (*f0)
00053         fileclose(*f0);
00054     if (*f1)
00055         fileclose(*f1);
00056     return -1;
00057 }
```

Referenced by [sys_pipe\(\)](#).

5.25.2.90 pipeclose()

```
void pipeclose (
    struct pipe * p,
    int writable )
```

Definition at line 60 of file [pipe.c](#).

```
00061 {
00062     acquire(&p->lock);
00063     if (writable) {
00064         p->writeopen = 0;
00065         wakeup(&p->nread);
00066     } else {
00067         p->readopen = 0;
00068         wakeup(&p->nwrite);
00069     }
00070     if (p->readopen == 0 && p->writeopen == 0) {
00071         release(&p->lock);
00072         kfree((char*)p);
00073     } else
00074         release(&p->lock);
00075 }
```

Referenced by [fileclose\(\)](#).

5.25.2.91 `piperead()`

```
int piperead (
    struct pipe * p,
    char * addr,
    int n )
```

Definition at line 101 of file `pipe.c`.

```
00102 {
00103     int i;
00104
00105     acquire(&p->lock);
00106     while(p->nread == p->nwrite && p->writeopen){ //DOC: pipe-empty
00107         if(myproc()->killed){
00108             release(&p->lock);
00109             return -1;
00110         }
00111         sleep(&p->nread, &p->lock); //DOC: piperead-sleep
00112     }
00113     for(i = 0; i < n; i++){ //DOC: piperead-copy
00114         if(p->nread == p->nwrite)
00115             break;
00116         addr[i] = p->data[p->nread++ % PIPESIZE];
00117     }
00118     wakeup(&p->nwrite); //DOC: piperead-wakeup
00119     release(&p->lock);
00120     return i;
00121 }
```

Referenced by `fileread()`.

5.25.2.92 `pipewrite()`

```
int pipewrite (
    struct pipe * p,
    char * addr,
    int n )
```

Definition at line 79 of file `pipe.c`.

```
00080 {
00081     int i;
00082
00083     acquire(&p->lock);
00084     for(i = 0; i < n; i++){
00085         while(p->nwrite == p->nread + PIPESIZE){ //DOC: pipewrite-full
00086             if(p->readopen == 0 || myproc()->killed){
00087                 release(&p->lock);
00088                 return -1;
00089             }
00090             wakeup(&p->nread);
00091             sleep(&p->nwrite, &p->lock); //DOC: pipewrite-sleep
00092         }
00093         p->data[p->nwrite++ % PIPESIZE] = addr[i];
00094     }
00095     wakeup(&p->nread); //DOC: pipewrite-wakeup1
00096     release(&p->lock);
00097     return n;
00098 }
```

Referenced by `filewrite()`.

5.25.2.93 popcli()

```
void popcli (
    void )
```

Definition at line 117 of file [spinlock.c](#).

```
00118 {
00119     if(readeflags() & FL_IF)
00120         panic("popcli - interruptible");
00121     if(--mycpu()->ncli < 0)
00122         panic("popcli");
00123     if(mycpu()->ncli == 0 && mycpu()->intena)
00124         sti();
00125 }
```

Referenced by [holding\(\)](#), [myproc\(\)](#), [release\(\)](#), and [switchvm\(\)](#).

5.25.2.94 procdump()

```
void procdump (
    void )
```

Definition at line 504 of file [proc.c](#).

```
00505 {
00506     static char *states[] = {
00507         [UNUSED]    "unused",
00508         [EMBRYO]     "embryo",
00509         [SLEEPING]   "sleep ",
00510         [RUNNABLE]   "runble",
00511         [RUNNING]    "run   ",
00512         [ZOMBIE]     "zombie"
00513     };
00514     int i;
00515     struct proc *p;
00516     char *state;
00517     uint pc[10];
00518
00519     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
00520         if(p->state == UNUSED)
00521             continue;
00522         if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
00523             state = states[p->state];
00524         else
00525             state = "???";
00526         cprintf("%d %s %s", p->pid, state, p->name);
00527         if(p->state == SLEEPING) {
00528             getcallerpcs((uint*)p->context->ebp+2, pc);
00529             for(i=0; i<10 && pc[i] != 0; i++)
00530                 cprintf(" %p", pc[i]);
00531         }
00532         cprintf("\n");
00533     }
00534 }
```

Referenced by [consoleintr\(\)](#).

5.25.2.95 pushcli()

```
void pushcli (
    void )
```

Definition at line 105 of file [spinlock.c](#).

```
00106 {
00107     int eflags;
00108
00109     eflags = readeflags();
00110     cli();
00111     if(mycpu()->ncli == 0)
00112         mycpu()->intena = eflags & FL_IF;
00113     mycpu()->ncli += 1;
00114 }
```

Referenced by [acquire\(\)](#), [holding\(\)](#), [myproc\(\)](#), and [switchvm\(\)](#).

5.25.2.96 readi()

```
int readi (
    struct inode * ip,
    char * dst,
    uint off,
    uint n )
```

Definition at line 453 of file [fs.c](#).

```
00454 {
00455     uint tot, m;
00456     struct buf *bp;
00457
00458     if(ip->type == T_DEV){
00459         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
00460             return -1;
00461         return devsw[ip->major].read(ip, dst, n);
00462     }
00463
00464     if(off > ip->size || off + n < off)
00465         return -1;
00466     if(off + n > ip->size)
00467         n = ip->size - off;
00468
00469     for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
00470         bp = bread(ip->dev, bmap(ip, off/BSIZE));
00471         m = min(n - tot, BSIZE - off%BSIZE);
00472         memmove(dst, bp->data + off%BSIZE, m);
00473         brelse(bp);
00474     }
00475     return n;
00476 }
```

Referenced by [dirlink\(\)](#), [dirlookup\(\)](#), [exec\(\)](#), [fileread\(\)](#), [isdirempty\(\)](#), and [loaduvm\(\)](#).

5.25.2.97 readsb()

```
void readsb (
    int dev,
    struct superblock * sb )
```

Definition at line 32 of file [fs.c](#).

```
00033 {
00034     struct buf *bp;
00035
00036     bp = bread(dev, 1);
00037     memmove(sb, bp->data, sizeof(*sb));
00038     brelse(bp);
00039 }
```

Referenced by [iinit\(\)](#), and [initlog\(\)](#).

5.25.2.98 release()

```
void release (
    struct spinlock * lk )
```

Definition at line 47 of file [spinlock.c](#).

```
00048 {
00049     if(!holding(lk))
00050         panic("release");
00051 }
```

```

00052     lk->pcs[0] = 0;
00053     lk->cpu = 0;
00054
00055     // Tell the C compiler and the processor to not move loads or stores
00056     // past this point, to ensure that all the stores in the critical
00057     // section are visible to other cores before the lock is released.
00058     // Both the C compiler and the hardware may re-order loads and
00059     // stores; __sync_synchronize() tells them both not to.
00060     __sync_synchronize();
00061
00062     // Release the lock, equivalent to lk->locked = 0.
00063     // This code can't use a C assignment, since it might
00064     // not be atomic. A real OS would use C atomics here.
00065     asm volatile("movl $0, %0" : "+m" (lk->locked) : );
00066
00067     popcli();
00068 }

```

Referenced by [acquiresleep\(\)](#), [allocproc\(\)](#), [begin_op\(\)](#), [bget\(\)](#), [brelse\(\)](#), [chpr\(\)](#), [consoleintr\(\)](#), [consoleread\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), [cps\(\)](#), [end_op\(\)](#), [filealloc\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fork\(\)](#), [forkret\(\)](#), [holdingsleep\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idup\(\)](#), [iget\(\)](#), [iput\(\)](#), [kalloc\(\)](#), [kfree\(\)](#), [kill\(\)](#), [log_write\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup\(\)](#), and [yield\(\)](#).

5.25.2.99 releasesleep()

```

void releasesleep (
    struct sleeplock * lk )

```

Definition at line 35 of file [sleeplock.c](#).

```

00036 {
00037     acquire(&lk->lk);
00038     lk->locked = 0;
00039     lk->pid = 0;
00040     wakeup(lk);
00041     release(&lk->lk);
00042 }

```

Referenced by [brelse\(\)](#), [iput\(\)](#), and [iunlock\(\)](#).

5.25.2.100 safestrcpy()

```

char * safestrcpy (
    char * s,
    const char * t,
    int n )

```

Definition at line 83 of file [string.c](#).

```

00084 {
00085     char *os;
00086
00087     os = s;
00088     if(n <= 0)
00089         return os;
00090     while(--n > 0 && (*s++ = *t++) != 0)
00091         ;
00092     *s = 0;
00093     return os;
00094 }

```

Referenced by [exec\(\)](#), [fork\(\)](#), and [userinit\(\)](#).

5.25.2.101 sched()

```
void sched (
    void )
```

Definition at line 366 of file [proc.c](#).

```
00367 {
00368     int intena;
00369     struct proc *p = myproc();
00370
00371     if(!holding(&ptable.lock))
00372         panic("sched ptable.lock");
00373     if(mycpu()->ncli != 1)
00374         panic("sched locks");
00375     if(p->state == RUNNING)
00376         panic("sched running");
00377     if(readeflags() & FL_IF)
00378         panic("sched interruptible");
00379     intena = mycpu()->intena;
00380     swtch(&p->context, mycpu()->scheduler);
00381     mycpu()->intena = intena;
00382 }
```

Referenced by [exit\(\)](#), [sleep\(\)](#), and [yield\(\)](#).

5.25.2.102 scheduler()

```
void scheduler (
    void )
```

Definition at line 323 of file [proc.c](#).

```
00324 {
00325     struct proc *p;
00326     struct cpu *c = mycpu();
00327     c->proc = 0;
00328
00329     for(;;){
00330         // Enable interrupts on this processor.
00331         sti();
00332
00333         // Loop over process table looking for process to run.
00334         acquire(&ptable.lock);
00335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00336             if(p->state != RUNNABLE)
00337                 continue;
00338
00339             // Switch to chosen process. It is the process's job
00340             // to release ptable.lock and then reacquire it
00341             // before jumping back to us.
00342             c->proc = p;
00343             switchvm(p);
00344             p->state = RUNNING;
00345
00346             swtch(&(c->scheduler), p->context);
00347             switchkvm();
00348
00349             // Process is done running for now.
00350             // It should have changed its p->state before coming back.
00351             c->proc = 0;
00352         }
00353         release(&ptable.lock);
00354     }
00355 }
00356 }
```

Referenced by [sched\(\)](#).

5.25.2.103 seginit()

```
void seginit (
    void )
```

Definition at line 16 of file [vm.c](#).

```
00017 {
00018     struct cpu *c;
00019
00020     // Map "logical" addresses to virtual addresses using identity map.
00021     // Cannot share a CODE descriptor for both kernel and user
00022     // because it would have to have DPL_USR, but the CPU forbids
00023     // an interrupt from CPL=0 to DPL=3.
00024     c = &cpus[cpuid()];
00025     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
00026     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
00027     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
00028     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
00029     lgdt(c->gdt, sizeof(c->gdt));
00030 }
```

Referenced by [mpenter\(\)](#).

5.25.2.104 setproc()

```
void setproc (
    struct proc * )
```

5.25.2.105 setupkvm()

```
pde_t * setupkvm (
    void )
```

Definition at line 119 of file [vm.c](#).

```
00120 {
00121     pde_t *pgdir;
00122     struct kmap *k;
00123
00124     if((pgdir = (pde_t*)kalloc()) == 0)
00125         return 0;
00126     memset(pgdir, 0, PGSIZE);
00127     if (P2V(PHYSTOP) > (void*)DEVSPACE)
00128         panic("PHYSTOP too high");
00129     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
00130         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
00131             (uint)k->phys_start, k->perm) < 0) {
00132             freevm(pgdir);
00133             return 0;
00134         }
00135     return pgdir;
00136 }
```

Referenced by [copyuvvm\(\)](#), [exec\(\)](#), [kvmalloc\(\)](#), and [userinit\(\)](#).

5.25.2.106 sleep()

```
void sleep (
    void * chan,
    struct spinlock * lk )
```

Definition at line 418 of file [proc.c](#).

```
00419 {
00420     struct proc *p = myproc();
00421
00422     if(p == 0)
00423         panic("sleep");
00424
00425     if(lk == 0)
00426         panic("sleep without lk");
00427
00428     // Must acquire ptable.lock in order to
00429     // change p->state and then call sched.
00430     // Once we hold ptable.lock, we can be
00431     // guaranteed that we won't miss any wakeup
00432     // (wakeup runs with ptable.lock locked),
00433     // so it's okay to release lk.
00434     if(lk != &ptable.lock){ //DOC: sleeplock0
00435         acquire(&ptable.lock); //DOC: sleeplock1
00436         release(lk);
00437     }
00438     // Go to sleep.
00439     p->chan = chan;
00440     p->state = SLEEPING;
00441
00442     sched();
00443
00444     // Tidy up.
00445     p->chan = 0;
00446
00447     // Reacquire original lock.
00448     if(lk != &ptable.lock){ //DOC: sleeplock2
00449         release(&ptable.lock);
00450         acquire(lk);
00451     }
00452 }
```

Referenced by [acquiresleep\(\)](#), [begin_op\(\)](#), [consoleread\(\)](#), [iderw\(\)](#), [main\(\)](#), [openiputtest\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [sbrktest\(\)](#), [sys_sleep\(\)](#), [validatetest\(\)](#), and [wait\(\)](#).

5.25.2.107 stati()

```
void stati (
    struct inode * ip,
    struct stat * st )
```

Definition at line 440 of file [fs.c](#).

```
00441 {
00442     st->dev = ip->dev;
00443     st->ino = ip->inum;
00444     st->type = ip->type;
00445     st->nlink = ip->nlink;
00446     st->size = ip->size;
00447 }
```

Referenced by [filestat\(\)](#).

5.25.2.108 strlen()

```
int strlen (
    const char * s )
```

Definition at line 97 of file [string.c](#).

```
00098 {
00099     int n;
00100
00101     for(n = 0; s[n]; n++)
00102         ;
00103     return n;
00104 }
```

Referenced by [exec\(\)](#), [fmtname\(\)](#), [ls\(\)](#), [main\(\)](#), [parsecmd\(\)](#), and [printf\(\)](#).

5.25.2.109 strncmp()

```
int strncmp (
    const char * p,
    const char * q,
    uint n )
```

Definition at line 59 of file [string.c](#).

```
00060 {
00061     while(n > 0 && *p && *p == *q)
00062         n--, p++, q++;
00063     if(n == 0)
00064         return 0;
00065     return (uchar)*p - (uchar)*q;
00066 }
```

Referenced by [namecmp\(\)](#).

5.25.2.110 strncpy()

```
char * strncpy (
    char * s,
    const char * t,
    int n )
```

Definition at line 69 of file [string.c](#).

```
00070 {
00071     char *os;
00072
00073     os = s;
00074     while(n-- > 0 && (*s++ = *t++) != 0)
00075         ;
00076     while(n-- > 0)
00077         *s++ = 0;
00078     return os;
00079 }
```

Referenced by [dirlink\(\)](#), [main\(\)](#), and [sys_getprocs\(\)](#).

5.25.2.111 switchkvm()

```
void switchkvm (  
    void )
```

Definition at line 150 of file [vm.c](#).

```
00151 {  
00152     lcr3(V2P(kpgdir));    // switch to the kernel page table  
00153 }
```

Referenced by [kvmalloc\(\)](#), [mpenter\(\)](#), and [scheduler\(\)](#).

5.25.2.112 switchvmm()

```
void switchvmm (  
    struct proc * p )
```

Definition at line 157 of file [vm.c](#).

```
00158 {  
00159     if(p == 0)  
00160         panic("switchvmm: no process");  
00161     if(p->kstack == 0)  
00162         panic("switchvmm: no kstack");  
00163     if(p->pgdir == 0)  
00164         panic("switchvmm: no pgdir");  
00165  
00166     pushcli();  
00167     mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A, &mycpu()->ts,  
00168                                   sizeof(mycpu()->ts)-1, 0);  
00169     mycpu()->gdt[SEG_TSS].s = 0;  
00170     mycpu()->ts.ss0 = SEG_KDATA << 3;  
00171     mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;  
00172     // setting IOPL=0 in eflags *and* iomb beyond the tss segment limit  
00173     // forbids I/O instructions (e.g., inb and outb) from user space  
00174     mycpu()->ts.iomb = (ushort) 0xFFFF;  
00175     ltr(SEG_TSS << 3);  
00176     lcr3(V2P(p->pgdir));    // switch to process's address space  
00177     popcli();  
00178 }
```

Referenced by [exec\(\)](#), [growproc\(\)](#), and [scheduler\(\)](#).

5.25.2.113 switch()

```
void swtch (  
    struct context **,  
    struct context * )
```

Referenced by [sched\(\)](#), and [scheduler\(\)](#).

5.25.2.114 syscall()

```
void syscall (
    void )
```

Definition at line 140 of file [syscall.c](#).

```
00141 {
00142     int num;
00143     struct proc *curproc = myproc();
00144
00145     num = curproc->tf->eax;
00146     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
00147         curproc->tf->eax = syscalls[num]();
00148     } else {
00149         cprintf("%d %s: unknown sys call %d\n",
00150             curproc->pid, curproc->name, num);
00151         curproc->tf->eax = -1;
00152     }
00153 }
```

Referenced by [trap\(\)](#).

5.25.2.115 timerinit()

```
void timerinit (
    void )
```

5.25.2.116 tvinit()

```
void tvinit (
    void )
```

Definition at line 18 of file [trap.c](#).

```
00019 {
00020     int i;
00021
00022     for(i = 0; i < 256; i++)
00023         SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);
00024     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3, vectors[T_SYSCALL], DPL_USER);
00025
00026     initlock(&tickslock, "time");
00027 }
```

5.25.2.117 uartinit()

```
void uartinit (
    void )
```

Definition at line 20 of file [uart.c](#).

```
00021 {
00022     char *p;
00023
00024     // Turn off the FIFO
00025     outb(COM1+2, 0);
00026
00027     // 9600 baud, 8 data bits, 1 stop bit, parity off.
00028     outb(COM1+3, 0x80); // Unlock divisor
```

```
00029     outb(COM1+0, 115200/9600);
00030     outb(COM1+1, 0);
00031     outb(COM1+3, 0x03);    // Lock divisor, 8 data bits.
00032     outb(COM1+4, 0);
00033     outb(COM1+1, 0x01);    // Enable receive interrupts.
00034
00035     // If status is 0xFF, no serial port.
00036     if(inb(COM1+5) == 0xFF)
00037         return;
00038     uart = 1;
00039
00040     // Acknowledge pre-existing interrupt conditions;
00041     // enable interrupts.
00042     inb(COM1+2);
00043     inb(COM1+0);
00044     ioapicenable(IRQ_COM1, 0);
00045
00046     // Announce that we're here.
00047     for(p="xv6...\n"; *p; p++)
00048         uartputc(*p);
00049 }
```

5.25.2.118 uartintr()

```
void uartintr (
    void )
```

Definition at line 74 of file [uart.c](#).

```
00075 {
00076     consoleintr(uartgetc);
00077 }
```

Referenced by [trap\(\)](#).

5.25.2.119 uartputc()

```
void uartputc (
    int c )
```

Definition at line 52 of file [uart.c](#).

```
00053 {
00054     int i;
00055
00056     if(!uart)
00057         return;
00058     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)
00059         microdelay(10);
00060     outb(COM1+0, c);
00061 }
```

Referenced by [consputc\(\)](#), and [uartinit\(\)](#).

5.25.2.120 userinit()

```
void userinit (
    void )
```

Definition at line 121 of file [proc.c](#).

```
00122 {
00123     struct proc *p;
00124     extern char _binary_initcode_start[], _binary_initcode_size[];
00125
00126     p = allocproc();
00127
00128     initproc = p;
00129     if((p->pgdir = setupkvm()) == 0)
00130         panic("userinit: out of memory?");
00131     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
00132     p->sz = PGSIZE;
00133     memset(p->tf, 0, sizeof(*p->tf));
00134     p->tf->cs = (SEG_UCODE « 3) | DPL_USER;
00135     p->tf->ds = (SEG_UDATA « 3) | DPL_USER;
00136     p->tf->es = p->tf->ds;
00137     p->tf->ss = p->tf->ds;
00138     p->tf->eflags = FL_IF;
00139     p->tf->esp = PGSIZE;
00140     p->tf->eip = 0; // beginning of initcode.S
00141
00142     safestrcpy(p->name, "initcode", sizeof(p->name));
00143     p->cwd = namei("/");
00144
00145     // this assignment to p->state lets other cores
00146     // run this process. the acquire forces the above
00147     // writes to be visible, and the lock is also needed
00148     // because the assignment might not be atomic.
00149     acquire(&ptable.lock);
00150
00151     p->state = RUNNABLE;
00152
00153     release(&ptable.lock);
00154 }
```

5.25.2.121 uva2ka()

```
char * uva2ka (
    pde_t * pgdir,
    char * uva )
```

Definition at line 350 of file [vm.c](#).

```
00351 {
00352     pte_t *pte;
00353
00354     pte = walkpgdir(pgdir, uva, 0);
00355     if((*pte & PTE_P) == 0)
00356         return 0;
00357     if((*pte & PTE_U) == 0)
00358         return 0;
00359     return (char*)P2V(PTE_ADDR(*pte));
00360 }
```

Referenced by [copyout\(\)](#).

5.25.2.122 wait()

```
int wait (
    void )
```

Definition at line 273 of file [proc.c](#).

```
00274 {
00275     struct proc *p;
00276     int havekids, pid;
00277     struct proc *curproc = myproc();
00278
00279     acquire(&ptable.lock);
00280     for(;;){
00281         // Scan through table looking for exited children.
00282         havekids = 0;
00283         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00284             if(p->parent != curproc)
00285                 continue;
00286             havekids = 1;
00287             if(p->state == ZOMBIE){
00288                 // Found one.
00289                 pid = p->pid;
00290                 kfree(p->kstack);
00291                 p->kstack = 0;
00292                 freevm(p->pgdir);
00293                 p->pid = 0;
00294                 p->parent = 0;
00295                 p->name[0] = 0;
00296                 p->killed = 0;
00297                 p->state = UNUSED;
00298                 release(&ptable.lock);
00299                 return pid;
00300             }
00301         }
00302
00303         // No point waiting if we don't have any children.
00304         if(!havekids || curproc->killed){
00305             release(&ptable.lock);
00306             return -1;
00307         }
00308
00309         // Wait for children to exit. (See wakeupl call in proc_exit.)
00310         sleep(curproc, &ptable.lock); //DOC: wait-sleep
00311     }
00312 }
```

Referenced by [bigargtest\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [exitiputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [sys_wait\(\)](#), [uio\(\)](#), and [validatetest\(\)](#).

5.25.2.123 wakeup()

```
void wakeup (
    void * chan )
```

Definition at line 469 of file [proc.c](#).

```
00470 {
00471     acquire(&ptable.lock);
00472     wakeupl(chan);
00473     release(&ptable.lock);
00474 }
```

Referenced by [consoleintr\(\)](#), [end_op\(\)](#), [ideintr\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), and [trap\(\)](#).

5.25.2.124 writei()

```
int writei (
    struct inode * ip,
    char * src,
    uint off,
    uint n )
```

Definition at line 482 of file [fs.c](#).

```
00483 {
00484     uint tot, m;
00485     struct buf *bp;
00486
00487     if(ip->type == T_DEV){
00488         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
00489             return -1;
00490         return devsw[ip->major].write(ip, src, n);
00491     }
00492
00493     if(off > ip->size || off + n < off)
00494         return -1;
00495     if(off + n > MAXFILE*BSIZE)
00496         return -1;
00497
00498     for(tot=0; tot<n; tot+=m, off+=m, src+=m){
00499         bp = bread(ip->dev, bmap(ip, off/BSIZE));
00500         m = min(n - tot, BSIZE - off%BSIZE);
00501         memmove(bp->data + off%BSIZE, src, m);
00502         log_write(bp);
00503         brelse(bp);
00504     }
00505
00506     if(n > 0 && off > ip->size){
00507         ip->size = off;
00508         iupdate(ip);
00509     }
00510     return n;
00511 }
```

Referenced by [dirlink\(\)](#), [filewrite\(\)](#), and [sys_unlink\(\)](#).

5.25.2.125 yield()

```
void yield (
    void )
```

Definition at line 386 of file [proc.c](#).

```
00387 {
00388     acquire(&ptable.lock); //DOC: yieldlock
00389     myproc()->state = RUNNABLE;
00390     sched();
00391     release(&ptable.lock);
00392 }
```

Referenced by [trap\(\)](#).

5.25.3 Variable Documentation

5.25.3.1 ioapicid

```
uchar ioapicid [extern]
```

Definition at line 16 of file [mp.c](#).

Referenced by [ioapicinit\(\)](#), and [mpinit\(\)](#).

5.25.3.2 ismp

```
int ismp [extern]
```

Referenced by [mpinit\(\)](#).

5.25.3.3 lapic

```
volatile uint* lapic [extern]
```

Definition at line 44 of file [lapic.c](#).

Referenced by [lapiceoi\(\)](#), [lapicid\(\)](#), [lapicinit\(\)](#), [lapicw\(\)](#), and [mpinit\(\)](#).

5.25.3.4 ticks

```
uint ticks [extern]
```

Definition at line 15 of file [trap.c](#).

Referenced by [sys_sleep\(\)](#), [sys_uptime\(\)](#), and [trap\(\)](#).

5.25.3.5 tickslock

```
struct spinlock tickslock [extern]
```

Definition at line 14 of file [trap.c](#).

Referenced by [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), and [tvinit\(\)](#).

5.26 defs.h

[Go to the documentation of this file.](#)

```

00001 struct buf;
00002 struct context;
00003 struct file;
00004 struct inode;
00005 struct pipe;
00006 struct proc;
00007 struct rtcdate;
00008 struct spinlock;
00009 struct sleeplock;
00010 struct stat;
00011 struct superblock;
00012 struct uproc;
00013
00014 // bio.c
00015 void          binit(void);
00016 struct buf*   bread(uint, uint);
00017 void          brelse(struct buf*);
00018 void          bwrite(struct buf*);
00019
00020 // console.c
00021 void          consoleinit(void);
00022 void          cprintf(char*, ...);
00023 void          consoleintr(int (*)(void));
00024 void          panic(char*) __attribute__((noreturn));
00025
00026 // exec.c
00027 int           exec(char*, char**);
00028
00029 // file.c
00030 struct file*  filealloc(void);
00031 void          fileclose(struct file*);
00032 struct file*  filedup(struct file*);
00033 void          fileinit(void);
00034 int           fileread(struct file*, char*, int n);
00035 int           filestat(struct file*, struct stat*);
00036 int           filewrite(struct file*, char*, int n);
00037
00038 // fs.c
00039 void          readsb(int dev, struct superblock *sb);
00040 int           dirlink(struct inode*, char*, uint);
00041 struct inode* dirlookup(struct inode*, char*, uint*);
00042 struct inode* ialloc(uint, short);
00043 struct inode* idup(struct inode*);
00044 void          iinit(int dev);
00045 void          ilock(struct inode*);
00046 void          iput(struct inode*);
00047 void          iunlock(struct inode*);
00048 void          iunlockput(struct inode*);
00049 void          iupdate(struct inode*);
00050 int           namecmp(const char*, const char*);
00051 struct inode* namei(char*);
00052 struct inode* nameiparent(char*, char*);
00053 int           readi(struct inode*, char*, uint, uint);
00054 void          stati(struct inode*, struct stat*);
00055 int           writei(struct inode*, char*, uint, uint);
00056
00057 // ide.c
00058 void          ideinit(void);
00059 void          ideintr(void);
00060 void          iderw(struct buf*);
00061
00062 // ioapic.c
00063 void          ioapicenable(int irq, int cpu);
00064 extern uchar  ioapicid;
00065 void          ioapicinit(void);
00066
00067 // kalloc.c
00068 char*         kalloc(void);
00069 void          kfree(char*);
00070 void          kinit1(void*, void*);
00071 void          kinit2(void*, void*);
00072
00073 // kbd.c
00074 void          kbdintr(void);
00075
00076 // lapic.c
00077 void          cmostime(struct rtcdate *r);
00078 int           lapicid(void);
00079 extern volatile uint* lapic;
00080 void          lapiceoi(void);
00081 void          lapicinit(void);
00082 void          lapicstartap(uchar, uint);

```

```

00083 void                microdelay(int);
00084
00085 // log.c
00086 void                initlog(int dev);
00087 void                log_write(struct buf*);
00088 void                begin_op();
00089 void                end_op();
00090
00091 // mp.c
00092 extern int          ismp;
00093 void                mpinit(void);
00094
00095 // picirq.c
00096 void                picenable(int);
00097 void                picinit(void);
00098
00099 // pipe.c
00100 int                pipealloc(struct file**, struct file**);
00101 void                pipeclose(struct pipe*, int);
00102 int                piperead(struct pipe*, char*, int);
00103 int                pipewrite(struct pipe*, char*, int);
00104
00105 //PAGEBREAK: 16
00106 // proc.c
00107 int                cpuid(void);
00108 void                exit(void);
00109 int                fork(void);
00110 int                growproc(int);
00111 int                kill(int);
00112 struct cpu*        mycpu(void);
00113 struct proc*        myproc();
00114 void                pinit(void);
00115 void                procdump(void);
00116 void                scheduler(void) __attribute__((noreturn));
00117 void                sched(void);
00118 void                setproc(struct proc*);
00119 void                sleep(void*, struct spinlock*);
00120 void                userinit(void);
00121 int                wait(void);
00122 void                wakeup(void*);
00123 void                yield(void);
00124 int                chpr(int pid, int priority);
00125 int                cps(void);
00126 int                getprocs(int max, struct uproc*);
00127 //int                halt(void);
00128
00129 // switch.S
00130 void                swtch(struct context**, struct context*);
00131
00132 // spinlock.c
00133 void                acquire(struct spinlock*);
00134 void                getcallerpcs(void*, uint*);
00135 int                holding(struct spinlock*);
00136 void                initlock(struct spinlock*, char*);
00137 void                release(struct spinlock*);
00138 void                pushcli(void);
00139 void                popcli(void);
00140
00141 // sleeplock.c
00142 void                acquiresleep(struct sleeplock*);
00143 void                releasesleep(struct sleeplock*);
00144 int                holdingsleep(struct sleeplock*);
00145 void                initsleeplock(struct sleeplock*, char*);
00146
00147 // string.c
00148 int                memcmp(const void*, const void*, uint);
00149 void*                memmove(void*, const void*, uint);
00150 void*                memset(void*, int, uint);
00151 char*                safestrcpy(char*, const char*, int);
00152 int                strlen(const char*);
00153 int                strncmp(const char*, const char*, uint);
00154 char*                strncpy(char*, const char*, int);
00155
00156 // syscall.c
00157 int                argint(int, int*);
00158 int                argptr(int, char**, int);
00159 int                argstr(int, char**);
00160 int                fetchint(uint, int*);
00161 int                fetchstr(uint, char**);
00162 void                syscall(void);
00163
00164 // timer.c
00165 void                timerinit(void);
00166
00167 // trap.c
00168 void                idtinit(void);
00169 extern uint        ticks;

```

```

00170 void                tvinit(void);
00171 extern struct spinlock tickslock;
00172
00173 // uart.c
00174 void                uartinit(void);
00175 void                uartintr(void);
00176 void                uartputc(int);
00177
00178 // vm.c
00179 void                seginit(void);
00180 void                kvmalloc(void);
00181 pde_t*              setupkvm(void);
00182 char*               uva2ka(pde_t*, char*);
00183 int                 allocvm(pde_t*, uint, uint);
00184 int                 deallocvm(pde_t*, uint, uint);
00185 void                freevm(pde_t*);
00186 void                initvm(pde_t*, char*, uint);
00187 int                 loadvm(pde_t*, char*, struct inode*, uint, uint);
00188 pde_t*              copyvm(pde_t*, uint);
00189 void                switchvm(struct proc*);
00190 void                switchkvm(void);
00191 int                 copyout(pde_t*, uint, void*, uint);
00192 void                clearpteu(pde_t *pgdir, char *uva);
00193
00194 // number of elements in fixed-size array
00195 #define NELEM(x) (sizeof(x)/sizeof((x)[0]))

```

5.27 echo.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"

```

Functions

- int [main](#) (int argc, char *argv[])

5.27.1 Function Documentation

5.27.1.1 main()

```

int main (
    int argc,
    char * argv[] )

```

Definition at line 6 of file [echo.c](#).

```

00007 {
00008     int i;
00009
00010     for(i = 1; i < argc; i++)
00011         printf(1, "%s%s", argv[i], i+1 < argc ? " " : "\n");
00012     exit();
00013 }

```

5.28 echo.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 int
00006 main(int argc, char *argv[])
00007 {
00008     int i;
00009
00010     for(i = 1; i < argc; i++)
00011         printf(1, "%s%s", argv[i], i+1 < argc ? " " : "\n");
00012     exit();
00013 }
```

5.29 echo.d File Reference

5.30 echo.d

[Go to the documentation of this file.](#)

```
00001 echo.o: echo.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.31 elf.h File Reference

Classes

- struct [elfhdr](#)
- struct [proghdr](#)

Macros

- #define [ELF_MAGIC](#) 0x464C457FU
- #define [ELF_PROG_FLAG_EXEC](#) 1
- #define [ELF_PROG_FLAG_READ](#) 4
- #define [ELF_PROG_FLAG_WRITE](#) 2
- #define [ELF_PROG_LOAD](#) 1

5.31.1 Macro Definition Documentation

5.31.1.1 ELF_MAGIC

```
#define ELF_MAGIC 0x464C457FU
```

Definition at line 3 of file [elf.h](#).

5.31.1.2 ELF_PROG_FLAG_EXEC

```
#define ELF_PROG_FLAG_EXEC 1
```

Definition at line 40 of file [elf.h](#).

5.31.1.3 ELF_PROG_FLAG_READ

```
#define ELF_PROG_FLAG_READ 4
```

Definition at line 42 of file [elf.h](#).

5.31.1.4 ELF_PROG_FLAG_WRITE

```
#define ELF_PROG_FLAG_WRITE 2
```

Definition at line 41 of file [elf.h](#).

5.31.1.5 ELF_PROG_LOAD

```
#define ELF_PROG_LOAD 1
```

Definition at line 37 of file [elf.h](#).

5.32 elf.h

[Go to the documentation of this file.](#)

```
00001 // Format of an ELF executable file
00002
00003 #define ELF_MAGIC 0x464C457FU // "\x7FELF" in little endian
00004
00005 // File header
00006 struct elfhdr {
00007     uint magic; // must equal ELF_MAGIC
00008     uchar elf[12];
00009     ushort type;
00010     ushort machine;
00011     uint version;
00012     uint entry;
00013     uint phoff;
00014     uint shoff;
00015     uint flags;
00016     ushort ehsize;
00017     ushort phentsize;
00018     ushort phnum;
00019     ushort shentsize;
00020     ushort shnum;
00021     ushort shstrndx;
00022 };
00023
00024 // Program section header
00025 struct proghdr {
00026     uint type;
```



```
00027     uint off;
00028     uint vaddr;
00029     uint paddr;
00030     uint filesz;
00031     uint memsz;
00032     uint flags;
00033     uint align;
00034 };
00035
00036 // Values for Proghdr type
00037 #define ELF_PROG_LOAD      1
00038
00039 // Flag bits for Proghdr flags
00040 #define ELF_PROG_FLAG_EXEC 1
00041 #define ELF_PROG_FLAG_WRITE 2
00042 #define ELF_PROG_FLAG_READ 4
```

5.33 entryother.d File Reference

5.34 entryother.d

[Go to the documentation of this file.](#)

00001 entryother.o: entryother.S asm.h memlayout.h mmu.h

5.35 exec.c File Reference

```
#include "types.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "defs.h"
#include "x86.h"
#include "elf.h"
```

Functions

- int [exec](#) (char *path, char **argv)

5.35.1 Function Documentation

5.35.1.1 exec()

```
int exec (
    char * path,
    char ** argv )
```

Definition at line 11 of file [exec.c](#).

```
00012 {
00013     char *s, *last;
00014     int i, off;
00015     uint argc, sz, sp, ustack[3+MAXARG+1];
00016     struct elfhdr elf;
00017     struct inode *ip;
00018     struct proghdr ph;
00019     pde_t *pgdir, *oldpgdir;
00020     struct proc *curproc = myproc();
00021
00022     begin_op();
00023
00024     if((ip = namei(path)) == 0){
00025         end_op();
00026         cprintf("exec: fail\n");
00027         return -1;
00028     }
00029     ilock(ip);
00030     pgdir = 0;
00031
00032     // Check ELF header
00033     if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
00034         goto bad;
00035     if(elf.magic != ELF_MAGIC)
00036         goto bad;
00037
00038     if((pgdir = setupkvm()) == 0)
00039         goto bad;
00040
00041     // Load program into memory.
00042     sz = 0;
00043     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
00044         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
00045             goto bad;
00046         if(ph.type != ELF_PROG_LOAD)
00047             continue;
00048         if(ph.memsz < ph.filesz)
00049             goto bad;
00050         if(ph.vaddr + ph.memsz < ph.vaddr)
00051             goto bad;
00052         if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
00053             goto bad;
00054         if(ph.vaddr % PGSIZE != 0)
00055             goto bad;
00056         if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
00057             goto bad;
00058     }
00059     iunlockput(ip);
00060     end_op();
00061     ip = 0;
00062
00063     // Allocate two pages at the next page boundary.
00064     // Make the first inaccessible. Use the second as the user stack.
00065     sz = PGROUNDUP(sz);
00066     if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
00067         goto bad;
00068     clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
00069     sp = sz;
00070
00071     // Push argument strings, prepare rest of stack in ustack.
00072     for(argc = 0; argv[argc]; argc++) {
00073         if(argc >= MAXARG)
00074             goto bad;
00075         sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
00076         if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
00077             goto bad;
00078         ustack[3+argc] = sp;
00079     }
00080     ustack[3+argc] = 0;
00081
00082     ustack[0] = 0xffffffff; // fake return PC
00083     ustack[1] = argc;
00084     ustack[2] = sp - (argc+1)*4; // argv pointer
00085
00086     sp -= (3+argc+1) * 4;
00087     if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
00088         goto bad;
```

```

00089
00090 // Save program name for debugging.
00091 for(last=s=path; *s; s++)
00092     if(*s == '/')
00093         last = s+1;
00094 safestrcpy(curproc->name, last, sizeof(curproc->name));
00095
00096 // Commit to the user image.
00097 oldpgdir = curproc->pgdir;
00098 curproc->pgdir = pgdir;
00099 curproc->sz = sz;
00100 curproc->tf->eip = elf.entry; // main
00101 curproc->tf->esp = sp;
00102 switchvm(curproc);
00103 freevm(oldpgdir);
00104 return 0;
00105
00106 bad:
00107 if(pgdir)
00108     freevm(pgdir);
00109 if(ip){
00110     iunlockput(ip);
00111     end_op();
00112 }
00113 return -1;
00114 }

```

Referenced by [bigargtest\(\)](#), [exectest\(\)](#), [main\(\)](#), [runcmd\(\)](#), and [sys_exec\(\)](#).

5.36 exec.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "param.h"
00003 #include "memlayout.h"
00004 #include "mmu.h"
00005 #include "proc.h"
00006 #include "defs.h"
00007 #include "x86.h"
00008 #include "elf.h"
00009
00010 int
00011 exec(char *path, char **argv)
00012 {
00013     char *s, *last;
00014     int i, off;
00015     uint argc, sz, sp, ustack[3+MAXARG+1];
00016     struct elfhdr elf;
00017     struct inode *ip;
00018     struct proghdr ph;
00019     pde_t *pgdir, *oldpgdir;
00020     struct proc *curproc = myproc();
00021
00022     begin_op();
00023
00024     if((ip = namei(path)) == 0){
00025         end_op();
00026         cprintf("exec: fail\n");
00027         return -1;
00028     }
00029     ilock(ip);
00030     pgdir = 0;
00031
00032     // Check ELF header
00033     if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
00034         goto bad;
00035     if(elf.magic != ELF_MAGIC)
00036         goto bad;
00037
00038     if((pgdir = setupkvm()) == 0)
00039         goto bad;
00040
00041     // Load program into memory.
00042     sz = 0;
00043     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
00044         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
00045             goto bad;
00046         if(ph.type != ELF_PROG_LOAD)
00047             continue;
00048         if(ph.memsz < ph.filesz)

```

```

00049     goto bad;
00050     if(ph.vaddr + ph.memsz < ph.vaddr)
00051         goto bad;
00052     if((sz = allocuvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
00053         goto bad;
00054     if(ph.vaddr % PGSIZE != 0)
00055         goto bad;
00056     if(loaduvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
00057         goto bad;
00058 }
00059 iunlockput(ip);
00060 end_op();
00061 ip = 0;
00062
00063 // Allocate two pages at the next page boundary.
00064 // Make the first inaccessible. Use the second as the user stack.
00065 sz = PGROUNDUP(sz);
00066 if((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
00067     goto bad;
00068 clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
00069 sp = sz;
00070
00071 // Push argument strings, prepare rest of stack in ustack.
00072 for(argc = 0; argv[argc]; argc++) {
00073     if(argc >= MAXARG)
00074         goto bad;
00075     sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
00076     if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
00077         goto bad;
00078     ustack[3+argc] = sp;
00079 }
00080 ustack[3+argc] = 0;
00081
00082 ustack[0] = 0xffffffff; // fake return PC
00083 ustack[1] = argc;
00084 ustack[2] = sp - (argc+1)*4; // argv pointer
00085
00086 sp -= (3+argc+1) * 4;
00087 if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
00088     goto bad;
00089
00090 // Save program name for debugging.
00091 for(last=s=path; *s; s++)
00092     if(*s == '/')
00093         last = s+1;
00094 safestrcpy(curproc->name, last, sizeof(curproc->name));
00095
00096 // Commit to the user image.
00097 oldpgdir = curproc->pgdir;
00098 curproc->pgdir = pgdir;
00099 curproc->sz = sz;
00100 curproc->tf->eip = elf.entry; // main
00101 curproc->tf->esp = sp;
00102 switchvm(curproc);
00103 freevm(oldpgdir);
00104 return 0;
00105
00106 bad:
00107 if(pgdir)
00108     freevm(pgdir);
00109 if(ip){
00110     iunlockput(ip);
00111     end_op();
00112 }
00113 return -1;
00114 }

```

5.37 exec.d File Reference

5.38 exec.d

[Go to the documentation of this file.](#)

```

00001 exec.o: exec.c /usr/include/stdc-predef.h types.h param.h memlayout.h \
00002 mmu.h proc.h defs.h x86.h elf.h

```

5.39 exit.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- int [main](#) (int argc, char *argv[])

5.39.1 Function Documentation

5.39.1.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 10 of file [exit.c](#).

```
00011 {
00012     halt();
00013     exit();
00014 }
```

5.40 exit.c

[Go to the documentation of this file.](#)

```
00001 // Shuts down the system by using the halt() system call
00002 // to send a special signal to QEMU.
00003 // Added by Bill Katsak
00004 // Copied from: http://pdos.csail.mit.edu/6.828/2012/homework/xv6-syscall.html
00005
00006 #include "types.h"
00007 #include "stat.h"
00008 #include "user.h"
00009
00010 int main(int argc, char *argv[])
00011 {
00012     halt();
00013     exit();
00014 }
```

5.41 exit.d File Reference

5.42 exit.d

[Go to the documentation of this file.](#)

```
00001 exit.o: exit.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.43 fcntl.h File Reference

Macros

- `#define O_CREATE 0x200`
- `#define O_RDONLY 0x000`
- `#define O_RDWR 0x002`
- `#define O_WRONLY 0x001`

5.43.1 Macro Definition Documentation

5.43.1.1 O_CREATE

```
#define O_CREATE 0x200
```

Definition at line 4 of file [fcntl.h](#).

5.43.1.2 O_RDONLY

```
#define O_RDONLY 0x000
```

Definition at line 1 of file [fcntl.h](#).

5.43.1.3 O_RDWR

```
#define O_RDWR 0x002
```

Definition at line 3 of file [fcntl.h](#).

5.43.1.4 O_WRONLY

```
#define O_WRONLY 0x001
```

Definition at line 2 of file [fcntl.h](#).

5.44 fcntl.h

[Go to the documentation of this file.](#)

```
00001 #define O_RDONLY 0x000
00002 #define O_WRONLY 0x001
00003 #define O_RDWR 0x002
00004 #define O_CREATE 0x200
```

5.45 file.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "fs.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "file.h"
```

Functions

- struct [file](#) * [filealloc](#) (void)
- void [fileclose](#) (struct [file](#) *f)
- struct [file](#) * [filedup](#) (struct [file](#) *f)
- void [fileinit](#) (void)
- int [fileread](#) (struct [file](#) *f, char *addr, int n)
- int [filestat](#) (struct [file](#) *f, struct [stat](#) *st)
- int [filewrite](#) (struct [file](#) *f, char *addr, int n)

Variables

- struct [devsw](#) [devsw](#) [[NDEV](#)]
- struct {
 - struct [file](#) [file](#) [[NFILE](#)]
 - struct [spinlock](#) [lock](#)
- } [ftable](#)

5.45.1 Function Documentation

5.45.1.1 filealloc()

```
struct file * filealloc (
    void )
```

Definition at line 27 of file [file.c](#).

```
00028 {
00029     struct file *f;
00030
00031     acquire(&ftable.lock);
00032     for(f = ftable.file; f < ftable.file + NFILE; f++){
00033         if(f->ref == 0){
00034             f->ref = 1;
00035             release(&ftable.lock);
00036             return f;
00037         }
00038     }
00039     release(&ftable.lock);
00040     return 0;
00041 }
```

Referenced by [pipealloc\(\)](#), and [sys_open\(\)](#).

5.45.1.2 fileclose()

```
void fileclose (
    struct file * f )
```

Definition at line 57 of file [file.c](#).

```
00058 {
00059     struct file ff;
00060
00061     acquire(&ftable.lock);
00062     if(f->ref < 1)
00063         panic("fileclose");
00064     if(--f->ref > 0){
00065         release(&ftable.lock);
00066         return;
00067     }
00068     ff = *f;
00069     f->ref = 0;
00070     f->type = FD_NONE;
00071     release(&ftable.lock);
00072
00073     if(ff.type == FD_PIPE)
00074         pipeclose(ff.pipe, ff.writable);
00075     else if(ff.type == FD_INODE){
00076         begin_op();
00077         iput(ff.ip);
00078         end_op();
00079     }
00080 }
```

Referenced by [exit\(\)](#), [pipealloc\(\)](#), [sys_close\(\)](#), [sys_open\(\)](#), and [sys_pipe\(\)](#).

5.45.1.3 filedup()

```
struct file * filedup (
    struct file * f )
```

Definition at line 45 of file [file.c](#).

```
00046 {
00047     acquire(&ftable.lock);
00048     if(f->ref < 1)
00049         panic("filedup");
00050     f->ref++;
00051     release(&ftable.lock);
00052     return f;
00053 }
```

Referenced by [fork\(\)](#), and [sys_dup\(\)](#).

5.45.1.4 fileinit()

```
void fileinit (
    void )
```

Definition at line 20 of file [file.c](#).

```
00021 {
00022     initlock(&ftable.lock, "ftable");
00023 }
```

5.45.1.5 fileread()

```
int fileread (
    struct file * f,
    char * addr,
    int n )
```

Definition at line 97 of file [file.c](#).

```
00098 {
00099     int r;
00100
00101     if(f->readable == 0)
00102         return -1;
00103     if(f->type == FD_PIPE)
00104         return piperead(f->pipe, addr, n);
00105     if(f->type == FD_INODE) {
00106         ilock(f->ip);
00107         if((r = readi(f->ip, addr, f->off, n)) > 0)
00108             f->off += r;
00109         iunlock(f->ip);
00110         return r;
00111     }
00112     panic("fileread");
00113 }
```

Referenced by [sys_read\(\)](#).

5.45.1.6 filestat()

```
int filestat (
    struct file * f,
    struct stat * st )
```

Definition at line 84 of file [file.c](#).

```
00085 {
00086     if(f->type == FD_INODE) {
00087         ilock(f->ip);
00088         stati(f->ip, st);
00089         iunlock(f->ip);
00090         return 0;
00091     }
00092     return -1;
00093 }
```

Referenced by [sys_fstat\(\)](#).

5.45.1.7 filewrite()

```
int filewrite (
    struct file * f,
    char * addr,
    int n )
```

Definition at line 118 of file [file.c](#).

```
00119 {
00120     int r;
00121
00122     if(f->writable == 0)
00123         return -1;
00124     if(f->type == FD_PIPE)
00125         return pipewrite(f->pipe, addr, n);
00126     if(f->type == FD_INODE){
00127         // write a few blocks at a time to avoid exceeding
00128         // the maximum log transaction size, including
00129         // i-node, indirect block, allocation blocks,
00130         // and 2 blocks of slop for non-aligned writes.
00131         // this really belongs lower down, since writei()
00132         // might be writing a device like the console.
00133         int max = ((MAXOPBLOCKS-1-1-2) / 2) * 512;
00134         int i = 0;
00135         while(i < n){
00136             int n1 = n - i;
00137             if(n1 > max)
00138                 n1 = max;
00139
00140             begin_op();
00141             ilock(f->ip);
00142             if ((r = writei(f->ip, addr + i, f->off, n1)) > 0)
00143                 f->off += r;
00144             iunlock(f->ip);
00145             end_op();
00146
00147             if(r < 0)
00148                 break;
00149             if(r != n1)
00150                 panic("short filewrite");
00151             i += r;
00152         }
00153         return i == n ? n : -1;
00154     }
00155     panic("filewrite");
00156 }
```

Referenced by [sys_write\(\)](#).

5.45.2 Variable Documentation

5.45.2.1 devsw

```
struct devsw devsw[NDEV]
```

Definition at line 13 of file [file.c](#).

5.45.2.2 file

```
struct file file[NFILE]
```

Definition at line 16 of file [file.c](#).

5.45.2.3

```
struct { ... } ftable
```

Referenced by [filealloc\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), and [fileinit\(\)](#).

5.45.2.4 lock

```
struct spinlock lock
```

Definition at line 15 of file [file.c](#).

5.46 file.c

[Go to the documentation of this file.](#)

```
00001 //
00002 // File descriptors
00003 //
00004
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "param.h"
00008 #include "fs.h"
00009 #include "spinlock.h"
00010 #include "sleeplock.h"
00011 #include "file.h"
00012
00013 struct devsw devsw[NDEV];
00014 struct {
00015     struct spinlock lock;
00016     struct file file[NFILE];
00017 } ftable;
00018
00019 void
00020 fileinit(void)
00021 {
00022     initlock(&ftable.lock, "ftable");
00023 }
00024
00025 // Allocate a file structure.
00026 struct file*
00027 filealloc(void)
00028 {
00029     struct file *f;
00030
00031     acquire(&ftable.lock);
00032     for(f = ftable.file; f < ftable.file + NFILE; f++){
00033         if(f->ref == 0){
00034             f->ref = 1;
00035             release(&ftable.lock);
00036             return f;
00037         }
00038     }
00039     release(&ftable.lock);
00040     return 0;
00041 }
00042
00043 // Increment ref count for file f.
00044 struct file*
00045 filedup(struct file *f)
00046 {
00047     acquire(&ftable.lock);
00048     if(f->ref < 1)
00049         panic("filedup");
00050     f->ref++;
00051     release(&ftable.lock);
00052     return f;
00053 }
00054
00055 // Close file f. (Decrement ref count, close when reaches 0.)
```

```

00056 void
00057 fclose(struct file *f)
00058 {
00059     struct file ff;
00060
00061     acquire(&ftable.lock);
00062     if(f->ref < 1)
00063         panic("fclose");
00064     if(--f->ref > 0){
00065         release(&ftable.lock);
00066         return;
00067     }
00068     ff = *f;
00069     f->ref = 0;
00070     f->type = FD_NONE;
00071     release(&ftable.lock);
00072
00073     if(ff.type == FD_PIPE)
00074         pipeclose(ff.pipe, ff.writable);
00075     else if(ff.type == FD_INODE){
00076         begin_op();
00077         iput(ff.ip);
00078         end_op();
00079     }
00080 }
00081
00082 // Get metadata about file f.
00083 int
00084 filestat(struct file *f, struct stat *st)
00085 {
00086     if(f->type == FD_INODE){
00087         ilock(f->ip);
00088         stati(f->ip, st);
00089         iunlock(f->ip);
00090         return 0;
00091     }
00092     return -1;
00093 }
00094
00095 // Read from file f.
00096 int
00097 fileread(struct file *f, char *addr, int n)
00098 {
00099     int r;
00100
00101     if(f->readable == 0)
00102         return -1;
00103     if(f->type == FD_PIPE)
00104         return piperead(f->pipe, addr, n);
00105     if(f->type == FD_INODE){
00106         ilock(f->ip);
00107         if((r = readi(f->ip, addr, f->off, n)) > 0)
00108             f->off += r;
00109         iunlock(f->ip);
00110         return r;
00111     }
00112     panic("fileread");
00113 }
00114
00115 //PAGEBREAK!
00116 // Write to file f.
00117 int
00118 filewrite(struct file *f, char *addr, int n)
00119 {
00120     int r;
00121
00122     if(f->writable == 0)
00123         return -1;
00124     if(f->type == FD_PIPE)
00125         return pipewrite(f->pipe, addr, n);
00126     if(f->type == FD_INODE){
00127         // write a few blocks at a time to avoid exceeding
00128         // the maximum log transaction size, including
00129         // i-node, indirect block, allocation blocks,
00130         // and 2 blocks of slop for non-aligned writes.
00131         // this really belongs lower down, since writei()
00132         // might be writing a device like the console.
00133         int max = ((MAXOPBLOCKS-1-1-2) / 2) * 512;
00134         int i = 0;
00135         while(i < n){
00136             int n1 = n - i;
00137             if(n1 > max)
00138                 n1 = max;
00139
00140             begin_op();
00141             ilock(f->ip);
00142             if ((r = writei(f->ip, addr + i, f->off, n1)) > 0)

```

```

00143         f->off += r;
00144         iunlock(f->ip);
00145         end_op();
00146
00147         if(r < 0)
00148             break;
00149         if(r != nl)
00150             panic("short filewrite");
00151         i += r;
00152     }
00153     return i == n ? n : -1;
00154 }
00155 panic("filewrite");
00156 }
00157

```

5.47 file.d File Reference

5.48 file.d

[Go to the documentation of this file.](#)

```

00001 file.o: file.c /usr/include/stdc-predef.h types.h defs.h param.h fs.h \
00002 spinlock.h sleeplock.h file.h

```

5.49 file.h File Reference

Classes

- struct [devsw](#)
- struct [file](#)
- struct [inode](#)

Macros

- #define [CONSOLE](#) 1

Variables

- struct [devsw](#) [devsw](#) []

5.49.1 Macro Definition Documentation

5.49.1.1 CONSOLE

```
#define CONSOLE 1
```

Definition at line 37 of file [file.h](#).

5.49.2 Variable Documentation

5.49.2.1 devsw

```
struct devsw devsw[] [extern]
```

Definition at line 13 of file [file.c](#).

5.50 file.h

[Go to the documentation of this file.](#)

```
00001 struct file {
00002     enum { FD_NONE, FD_PIPE, FD_INODE } type;
00003     int ref; // reference count
00004     char readable;
00005     char writable;
00006     struct pipe *pipe;
00007     struct inode *ip;
00008     uint off;
00009 };
00010
00011
00012 // in-memory copy of an inode
00013 struct inode {
00014     uint dev;           // Device number
00015     uint inum;          // Inode number
00016     int ref;            // Reference count
00017     struct sleeplock lock; // protects everything below here
00018     int valid;          // inode has been read from disk?
00019
00020     short type;         // copy of disk inode
00021     short major;
00022     short minor;
00023     short nlink;
00024     uint size;
00025     uint addrs[NDIRECT+1];
00026 };
00027
00028 // table mapping major device number to
00029 // device functions
00030 struct devsw {
00031     int (*read)(struct inode*, char*, int);
00032     int (*write)(struct inode*, char*, int);
00033 };
00034
00035 extern struct devsw devsw[];
00036
00037 #define CONSOLE 1
```

5.51 foo.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"
```

Functions

- [int main](#) (int argc, char *argv[])

5.51.1 Function Documentation

5.51.1.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 6 of file [foo.c](#).

```
00006 {
00007     int pid;
00008     int k, n;
00009     int x, z;
00010
00011     if (argc != 2) {
00012         printf(2, "usage: %s n\n", argv[0]);
00013     }
00014
00015     n = atoi(argv[1]);
00016
00017     for ( k = 0; k < n; k++ ) {
00018         pid = fork ();
00019         if ( pid < 0 ) {
00020             printf(1, "%d failed in fork!\n", getpid());
00021             exit();
00022         } else if (pid == 0) {
00023             // child
00024             printf(1, "Child (with pid=%d) created\n",getpid());
00025             for ( z = 0; z < 100.0; z += 0.1 ){
00026                 x = x + 3.14 * 89.64;    // useless calculations to consume CPU time
00027                 if ((z== (int)z) && ((z%10) ==0)){
00028                     printf(1, "[pid=%d] %d \n",getpid(),z);
00029                 }
00030             }
00031
00032             exit();
00033         }
00034     }
00035
00036     for (k = 0; k < n; k++) {
00037         wait();
00038     }
00039
00040     exit();
00041 }
```

5.52 foo.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "fcntl.h"
00005
00006 int main(int argc, char *argv[]) {
00007     int pid;
00008     int k, n;
00009     int x, z;
00010
00011     if (argc != 2) {
00012         printf(2, "usage: %s n\n", argv[0]);
00013     }
00014
00015     n = atoi(argv[1]);
00016
00017     for ( k = 0; k < n; k++ ) {
00018         pid = fork ();
00019         if ( pid < 0 ) {
00020             printf(1, "%d failed in fork!\n", getpid());
00021             exit();
00022         } else if (pid == 0) {
```

```

00023     // child
00024     printf(1, "Child (with pid=%d) created\n",getpid());
00025     for ( z = 0; z < 100.0; z += 0.1 ){
00026         x = x + 3.14 * 89.64; // useless calculations to consume CPU time
00027         if((z==(int)z) && ((z%10) ==0)){
00028             printf(1, "[pid=%d] %d \n",getpid(),z);
00029         }
00030     }
00031
00032     exit();
00033 }
00034 }
00035
00036 for (k = 0; k < n; k++) {
00037     wait();
00038 }
00039
00040 exit();
00041 }

```

5.53 foo.d File Reference

5.54 foo.d

[Go to the documentation of this file.](#)

00001 foo.o: foo.c /usr/include/stdc-predef.h types.h stat.h user.h fcntl.h

5.55 forktest.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"

```

Macros

- `#define N 1000`

Functions

- void `forktest` (void)
- int `main` (void)
- void `printf` (int fd, const char *s,...)

5.55.1 Macro Definition Documentation

5.55.1.1 N

```
#define N 1000
```

Definition at line 8 of file [forktest.c](#).

5.55.2 Function Documentation

5.55.2.1 forktest()

```
void forktest (
    void )
```

Definition at line 17 of file [forktest.c](#).

```
00018 {
00019     int n, pid;
00020
00021     printf(1, "fork test\n");
00022
00023     for(n=0; n<N; n++){
00024         pid = fork();
00025         if(pid < 0)
00026             break;
00027         if(pid == 0)
00028             exit();
00029     }
00030
00031     if(n == N){
00032         printf(1, "fork claimed to work N times!\n", N);
00033         exit();
00034     }
00035
00036     for(; n > 0; n--){
00037         if(wait() < 0){
00038             printf(1, "wait stopped early\n");
00039             exit();
00040         }
00041     }
00042
00043     if(wait() != -1){
00044         printf(1, "wait got too many\n");
00045         exit();
00046     }
00047
00048     printf(1, "fork test OK\n");
00049 }
```

Referenced by [main\(\)](#).

5.55.2.2 main()

```
int main (
    void )
```

Definition at line 52 of file [forktest.c](#).

```
00053 {
00054     forktest();
00055     exit();
00056 }
```

5.55.2.3 printf()

```
void printf (
    int fd,
    const char * s,
    ... )
```

Definition at line 11 of file [forktest.c](#).

```
00012 {
00013     write(fd, s, strlen(s));
00014 }
```

Referenced by [argptest\(\)](#), [balloc\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [bsstest\(\)](#), [cat\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirfile\(\)](#), [dirtest\(\)](#), [exectest\(\)](#), [exitiputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [fsfull\(\)](#), [getcmd\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [ls\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [panic\(\)](#), [parsecmd\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [rmdot\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [uio\(\)](#), [unlinkread\(\)](#), [validatetest\(\)](#), [wc\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.56 forktest.c

[Go to the documentation of this file.](#)

```
00001 // Test that fork fails gracefully.
00002 // Tiny executable so that the limit can be filling the proc table.
00003
00004 #include "types.h"
00005 #include "stat.h"
00006 #include "user.h"
00007
00008 #define N 1000
00009
00010 void
00011 printf(int fd, const char *s, ...)
00012 {
00013     write(fd, s, strlen(s));
00014 }
00015
00016 void
00017 forktest(void)
00018 {
00019     int n, pid;
00020
00021     printf(1, "fork test\n");
00022
00023     for(n=0; n<N; n++){
00024         pid = fork();
00025         if(pid < 0)
00026             break;
00027         if(pid == 0)
00028             exit();
00029     }
00030
00031     if(n == N){
00032         printf(1, "fork claimed to work N times!\n", N);
00033         exit();
00034     }
00035
00036     for(; n > 0; n--){
00037         if(wait() < 0){
00038             printf(1, "wait stopped early\n");
00039             exit();
00040         }
00041     }
00042
00043     if(wait() != -1){
00044         printf(1, "wait got too many\n");
00045         exit();
00046     }
00047
00048     printf(1, "fork test OK\n");
00049 }
00050
00051 int
00052 main(void)
00053 {
00054     forktest();
00055     exit();
00056 }
```

5.57 forktest.d File Reference

5.58 forktest.d

[Go to the documentation of this file.](#)

```
00001 forktest.o: forktest.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.59 fs.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "stat.h"
#include "mmu.h"
#include "proc.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "buf.h"
#include "file.h"
```

Macros

- #define `min(a, b)` `((a) < (b) ? (a) : (b))`

Functions

- static `uint` `balloc` (`uint` dev)
- static void `bfree` (int dev, `uint` b)
- static `uint` `bmap` (struct `inode` *ip, `uint` bn)
- static void `bzero` (int dev, int bno)
- int `dirlink` (struct `inode` *dp, char *name, `uint` inum)
- struct `inode` * `dirlookup` (struct `inode` *dp, char *name, `uint` *poff)
- struct `inode` * `ialloc` (`uint` dev, short type)
- struct `inode` * `idup` (struct `inode` *ip)
- static struct `inode` * `iget` (`uint` dev, `uint` inum)
- void `iinit` (int dev)
- void `ilock` (struct `inode` *ip)
- void `iput` (struct `inode` *ip)
- static void `itrunc` (struct `inode` *)
- void `iunlock` (struct `inode` *ip)
- void `iunlockput` (struct `inode` *ip)
- void `iupdate` (struct `inode` *ip)
- int `namecmp` (const char *s, const char *t)
- struct `inode` * `namei` (char *path)
- struct `inode` * `nameiparent` (char *path, char *name)
- static struct `inode` * `nameex` (char *path, int nameiparent, char *name)
- int `readi` (struct `inode` *ip, char *dst, `uint` off, `uint` n)
- void `readsb` (int dev, struct `superblock` *sb)
- static char * `skipelem` (char *path, char *name)
- void `stati` (struct `inode` *ip, struct `stat` *st)
- int `writei` (struct `inode` *ip, char *src, `uint` off, `uint` n)

Variables

- struct {
 struct [inode](#) [inode](#) [[NINODE](#)]
 struct [spinlock](#) [lock](#)
} [icache](#)
- struct [superblock](#) [sb](#)

5.59.1 Macro Definition Documentation

5.59.1.1 min

```
#define min(  
    a,  
    b ) ((a) < (b) ? (a) : (b))
```

Definition at line 24 of file [fs.c](#).

5.59.2 Function Documentation

5.59.2.1 balloc()

```
static uint balloc (  
    uint dev ) [static]
```

Definition at line 57 of file [fs.c](#).

```
00058 {  
00059     int b, bi, m;  
00060     struct buf *bp;  
00061  
00062     bp = 0;  
00063     for(b = 0; b < sb.size; b += BPB){  
00064         bp = bread(dev, BBLOCK(b, sb));  
00065         for(bi = 0; bi < BPB && b + bi < sb.size; bi++){  
00066             m = 1 << (bi % 8);  
00067             if((bp->data[bi/8] & m) == 0){ // Is block free?  
00068                 bp->data[bi/8] |= m; // Mark block in use.  
00069                 log\_write(bp);  
00070                 brelse(bp);  
00071                 bzero(dev, b + bi);  
00072                 return b + bi;  
00073             }  
00074         }  
00075         brelse(bp);  
00076     }  
00077     panic("balloc: out of blocks");  
00078 }
```

Referenced by [bmap\(\)](#).

5.59.2.2 bfree()

```
static void bfree (
    int dev,
    uint b ) [static]
```

Definition at line 82 of file [fs.c](#).

```
00083 {
00084     struct buf *bp;
00085     int bi, m;
00086
00087     bp = bread(dev, BBLOCK(b, sb));
00088     bi = b % BPB;
00089     m = 1 « (bi % 8);
00090     if((bp->data[bi/8] & m) == 0)
00091         panic("freeing free block");
00092     bp->data[bi/8] &= ~m;
00093     log_write(bp);
00094     brelse(bp);
00095 }
```

Referenced by [itrunc\(\)](#).

5.59.2.3 bmap()

```
static uint bmap (
    struct inode *ip,
    uint bn ) [static]
```

Definition at line 373 of file [fs.c](#).

```
00374 {
00375     uint addr, *a;
00376     struct buf *bp;
00377
00378     if(bn < NDIRECT){
00379         if((addr = ip->addrs[bn]) == 0)
00380             ip->addrs[bn] = addr = balloc(ip->dev);
00381         return addr;
00382     }
00383     bn -= NDIRECT;
00384
00385     if(bn < NINDIRECT){
00386         // Load indirect block, allocating if necessary.
00387         if((addr = ip->addrs[NINDIRECT]) == 0)
00388             ip->addrs[NINDIRECT] = addr = balloc(ip->dev);
00389         bp = bread(ip->dev, addr);
00390         a = (uint*)bp->data;
00391         if((addr = a[bn]) == 0){
00392             a[bn] = addr = balloc(ip->dev);
00393             log_write(bp);
00394         }
00395         brelse(bp);
00396         return addr;
00397     }
00398     panic("bmap: out of range");
00400 }
```

Referenced by [readi\(\)](#), and [writei\(\)](#).

5.59.2.4 bzero()

```
static void bzero (
    int dev,
    int bno ) [static]
```

Definition at line 43 of file [fs.c](#).

```
00044 {
00045     struct buf *bp;
00046
00047     bp = bread(dev, bno);
00048     memset(bp->data, 0, BSIZE);
00049     log_write(bp);
00050     brelse(bp);
00051 }
```

Referenced by [balloc\(\)](#), [ialloc\(\)](#), and [main\(\)](#).

5.59.2.5 dirlink()

```
int dirlink (
    struct inode * dp,
    char * name,
    uint inum )
```

Definition at line 552 of file [fs.c](#).

```
00553 {
00554     int off;
00555     struct dirent de;
00556     struct inode *ip;
00557
00558     // Check that name is not present.
00559     if((ip = dirlookup(dp, name, 0)) != 0){
00560         iput(ip);
00561         return -1;
00562     }
00563
00564     // Look for an empty dirent.
00565     for(off = 0; off < dp->size; off += sizeof(de)){
00566         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00567             panic("dirlink read");
00568         if(de.inum == 0)
00569             break;
00570     }
00571
00572     strncpy(de.name, name, DIRSIZ);
00573     de.inum = inum;
00574     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00575         panic("dirlink");
00576
00577     return 0;
00578 }
```

Referenced by [create\(\)](#), and [sys_link\(\)](#).

5.59.2.6 dirlookup()

```
struct inode * dirlookup (
    struct inode * dp,
    char * name,
    uint * poff )
```

Definition at line 525 of file [fs.c](#).

```
00526 {
00527     uint off, inum;
00528     struct dirent de;
00529
00530     if(dp->type != T_DIR)
00531         panic("dirlookup not DIR");
00532
00533     for(off = 0; off < dp->size; off += sizeof(de)){
00534         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00535             panic("dirlookup read");
00536         if(de.inum == 0)
00537             continue;
00538         if(namecmp(name, de.name) == 0){
00539             // entry matches path element
00540             if(poff)
00541                 *poff = off;
00542             inum = de.inum;
00543             return iget(dp->dev, inum);
00544         }
00545     }
00546
00547     return 0;
00548 }
```

Referenced by [create\(\)](#), [dirlink\(\)](#), [namex\(\)](#), and [sys_unlink\(\)](#).

5.59.2.7 ialloc()

```
struct inode * ialloc (
    uint dev,
    short type )
```

Definition at line 195 of file [fs.c](#).

```
00196 {
00197     int inum;
00198     struct buf *bp;
00199     struct dinode *dip;
00200
00201     for(inum = 1; inum < sb.ninodes; inum++){
00202         bp = bread(dev, IBLOCK(inum, sb));
00203         dip = (struct dinode*)bp->data + inum%IPB;
00204         if(dip->type == 0){ // a free inode
00205             memset(dip, 0, sizeof(*dip));
00206             dip->type = type;
00207             log_write(bp); // mark it allocated on the disk
00208             brelse(bp);
00209             return iget(dev, inum);
00210         }
00211         brelse(bp);
00212     }
00213     panic("ialloc: no inodes");
00214 }
```

Referenced by [create\(\)](#).

5.59.2.8 idup()

```
struct inode * idup (
    struct inode * ip )
```

Definition at line 277 of file [fs.c](#).

```
00278 {
00279     acquire(&icache.lock);
00280     ip->ref++;
00281     release(&icache.lock);
00282     return ip;
00283 }
```

Referenced by [fork\(\)](#), and [namex\(\)](#).

5.59.2.9 iget()

```
static struct inode * iget (
    uint dev,
    uint inum ) [static]
```

Definition at line 242 of file [fs.c](#).

```
00243 {
00244     struct inode *ip, *empty;
00245
00246     acquire(&icache.lock);
00247
00248     // Is the inode already cached?
00249     empty = 0;
00250     for(ip = &icache.inode[0]; ip < &icache.inode[NINODE]; ip++){
00251         if(ip->ref > 0 && ip->dev == dev && ip->inum == inum){
00252             ip->ref++;
00253             release(&icache.lock);
00254             return ip;
00255         }
00256         if(empty == 0 && ip->ref == 0)    // Remember empty slot.
00257             empty = ip;
00258     }
00259
00260     // Recycle an inode cache entry.
00261     if(empty == 0)
00262         panic("iget: no inodes");
00263
00264     ip = empty;
00265     ip->dev = dev;
00266     ip->inum = inum;
00267     ip->ref = 1;
00268     ip->valid = 0;
00269     release(&icache.lock);
00270
00271     return ip;
00272 }
```

Referenced by [dirlookup\(\)](#), [ialloc\(\)](#), and [namex\(\)](#).

5.59.2.10 iinit()

```
void iinit (
    int dev )
```

Definition at line 172 of file [fs.c](#).

```
00173 {
00174     int i = 0;
00175
00176     initlock(&icache.lock, "icache");
00177     for(i = 0; i < NINODE; i++) {
00178         initsleeplock(&icache.inode[i].lock, "inode");
00179     }
00180
00181     readsb(dev, &sb);
00182     cprintf("sb: size %d nblocks %d ninodes %d nlog %d logstart %d\
00183 inodestart %d bmap start %d\n", sb.size, sb.nblocks,
00184         sb.ninodes, sb.nlog, sb.logstart, sb.inodestart,
00185         sb.bmapstart);
00186 }
```

Referenced by [forkret\(\)](#).

5.59.2.11 ilock()

```
void ilock (
    struct inode * ip )
```

Definition at line 288 of file [fs.c](#).

```
00289 {
00290     struct buf *bp;
00291     struct dinode *dip;
00292
00293     if(ip == 0 || ip->ref < 1)
00294         panic("ilock");
00295
00296     acquiresleep(&ip->lock);
00297
00298     if(ip->valid == 0){
00299         bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00300         dip = (struct dinode*)bp->data + ip->inum%IPB;
00301         ip->type = dip->type;
00302         ip->major = dip->major;
00303         ip->minor = dip->minor;
00304         ip->nlink = dip->nlink;
00305         ip->size = dip->size;
00306         memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
00307         brelse(bp);
00308         ip->valid = 1;
00309         if(ip->type == 0)
00310             panic("ilock: no type");
00311     }
00312 }
```

Referenced by [consoleread\(\)](#), [consolewrite\(\)](#), [create\(\)](#), [exec\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.59.2.12 iput()

```
void iput (
    struct inode * ip )
```

Definition at line 332 of file [fs.c](#).

```
00333 {
00334     acquiresleep(&ip->lock);
00335     if(ip->valid && ip->nlink == 0){
00336         acquire(&icache.lock);
00337         int r = ip->ref;
00338         release(&icache.lock);
00339         if(r == 1){
00340             // inode has no links and no other references: truncate and free.
00341             itrunc(ip);
00342             ip->type = 0;
00343             iupdate(ip);
00344             ip->valid = 0;
00345         }
00346     }
00347     releasesleep(&ip->lock);
00348
00349     acquire(&icache.lock);
00350     ip->ref--;
00351     release(&icache.lock);
00352 }
```

Referenced by [dirlink\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [iunlockput\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), and [sys_link\(\)](#).

5.59.2.13 itrunc()

```
static void itrunc (
    struct inode * ip ) [static]
```

Definition at line 408 of file [fs.c](#).

```
00409 {
00410     int i, j;
00411     struct buf *bp;
00412     uint *a;
00413
00414     for(i = 0; i < NDIRECT; i++){
00415         if(ip->addrs[i]){
00416             bfree(ip->dev, ip->addrs[i]);
00417             ip->addrs[i] = 0;
00418         }
00419     }
00420
00421     if(ip->addrs[NDIRECT]){
00422         bp = bread(ip->dev, ip->addrs[NDIRECT]);
00423         a = (uint*)bp->data;
00424         for(j = 0; j < NINDIRECT; j++){
00425             if(a[j])
00426                 bfree(ip->dev, a[j]);
00427         }
00428         brelse(bp);
00429         bfree(ip->dev, ip->addrs[NDIRECT]);
00430         ip->addrs[NDIRECT] = 0;
00431     }
00432
00433     ip->size = 0;
00434     iupdate(ip);
00435 }
```

Referenced by [iput\(\)](#).

5.59.2.14 iunlock()

```
void iunlock (
    struct inode * ip )
```

Definition at line 316 of file [fs.c](#).

```
00317 {
00318     if(ip == 0 || !holdingsleep(&ip->lock) || ip->ref < 1)
00319         panic("iunlock");
00320
00321     releasesleep(&ip->lock);
00322 }
```

Referenced by [consoleread\(\)](#), [consolewrite\(\)](#), [fileread\(\)](#), [filestat\(\)](#), [filewrite\(\)](#), [iunlockput\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), and [sys_open\(\)](#).

5.59.2.15 iunlockput()

```
void iunlockput (
    struct inode * ip )
```

Definition at line 356 of file [fs.c](#).

```
00357 {
00358     iunlock(ip);
00359     iput(ip);
00360 }
```

Referenced by [create\(\)](#), [exec\(\)](#), [namex\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.59.2.16 iupdate()

```
void iupdate (
    struct inode * ip )
```

Definition at line 221 of file [fs.c](#).

```
00222 {
00223     struct buf *bp;
00224     struct dinode *dip;
00225
00226     bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00227     dip = (struct dinode*)bp->data + ip->inum%IPB;
00228     dip->type = ip->type;
00229     dip->major = ip->major;
00230     dip->minor = ip->minor;
00231     dip->nlink = ip->nlink;
00232     dip->size = ip->size;
00233     memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
00234     log_write(bp);
00235     brelse(bp);
00236 }
```

Referenced by [create\(\)](#), [iput\(\)](#), [itrunc\(\)](#), [sys_link\(\)](#), [sys_unlink\(\)](#), and [writei\(\)](#).

5.59.2.17 namecmp()

```
int namecmp (
    const char * s,
    const char * t )
```

Definition at line 517 of file [fs.c](#).

```
00518 {
00519     return strcmp(s, t, DIRSIZ);
00520 }
```

Referenced by [dirlookup\(\)](#), and [sys_unlink\(\)](#).

5.59.2.18 namei()

```
struct inode * namei (
    char * path )
```

Definition at line 660 of file [fs.c](#).

```
00661 {
00662     char name[DIRSIZ];
00663     return namex(path, 0, name);
00664 }
```

Referenced by [exec\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_open\(\)](#), and [userinit\(\)](#).

5.59.2.19 nameiparent()

```
struct inode * nameiparent (
    char * path,
    char * name )
```

Definition at line 667 of file [fs.c](#).

```
00668 {
00669     return namex(path, 1, name);
00670 }
```

Referenced by [create\(\)](#), [namex\(\)](#), [sys_link\(\)](#), and [sys_unlink\(\)](#).

5.59.2.20 namex()

```
static struct inode * namex (
    char * path,
    int nameiparent,
    char * name ) [static]
```

Definition at line 625 of file [fs.c](#).

```
00626 {
00627     struct inode *ip, *next;
00628
00629     if(*path == '/')
00630         ip = iget(ROOTDEV, ROOTINO);
00631     else
00632         ip = idup(myproc()->cwd);
00633
00634     while((path = skipelem(path, name)) != 0){
00635         ilock(ip);
00636         if(ip->type != T_DIR){
00637             iunlockput(ip);
00638             return 0;
00639         }
00640         if(nameiparent && *path == '\0'){
00641             // Stop one level early.
00642             iunlock(ip);
00643             return ip;
00644         }
00645         if((next = dirlookup(ip, name, 0)) == 0){
00646             iunlockput(ip);
00647             return 0;
00648         }
00649         iunlockput(ip);
00650         ip = next;
00651     }
00652     if(nameiparent){
00653         iput(ip);
00654         return 0;
00655     }
00656     return ip;
00657 }
```

Referenced by [namei\(\)](#), and [nameiparent\(\)](#).

5.59.2.21 readi()

```
int readi (
    struct inode * ip,
    char * dst,
    uint off,
    uint n )
```

Definition at line 453 of file [fs.c](#).

```
00454 {
00455     uint tot, m;
00456     struct buf *bp;
00457
00458     if(ip->type == T_DEV){
00459         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
00460             return -1;
00461         return devsw[ip->major].read(ip, dst, n);
00462     }
00463
00464     if(off > ip->size || off + n < off)
00465         return -1;
00466     if(off + n > ip->size)
00467         n = ip->size - off;
00468
00469     for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
00470         bp = bread(ip->dev, bmap(ip, off/BSIZE));
00471         m = min(n - tot, BSIZE - off%BSIZE);
00472         memmove(dst, bp->data + off%BSIZE, m);
```

```

00473     brelse(bp);
00474 }
00475 return n;
00476 }

```

Referenced by [dirlink\(\)](#), [dirlookup\(\)](#), [exec\(\)](#), [fileread\(\)](#), [isdirempty\(\)](#), and [loaduvm\(\)](#).

5.59.2.22 readsb()

```

void readsb (
    int dev,
    struct superblock * sb )

```

Definition at line 32 of file [fs.c](#).

```

00033 {
00034     struct buf *bp;
00035
00036     bp = bread(dev, 1);
00037     memmove(sb, bp->data, sizeof(*sb));
00038     brelse(bp);
00039 }

```

Referenced by [iinit\(\)](#), and [initlog\(\)](#).

5.59.2.23 skipelem()

```

static char * skipelem (
    char * path,
    char * name ) [static]

```

Definition at line 596 of file [fs.c](#).

```

00597 {
00598     char *s;
00599     int len;
00600
00601     while(*path == '/')
00602         path++;
00603     if(*path == 0)
00604         return 0;
00605     s = path;
00606     while(*path != '/' && *path != 0)
00607         path++;
00608     len = path - s;
00609     if(len >= DIRSIZ)
00610         memmove(name, s, DIRSIZ);
00611     else {
00612         memmove(name, s, len);
00613         name[len] = 0;
00614     }
00615     while(*path == '/')
00616         path++;
00617     return path;
00618 }

```

Referenced by [namex\(\)](#).

5.59.2.24 stati()

```
void stati (
    struct inode * ip,
    struct stat * st )
```

Definition at line 440 of file [fs.c](#).

```
00441 {
00442     st->dev = ip->dev;
00443     st->ino = ip->inum;
00444     st->type = ip->type;
00445     st->nlink = ip->nlink;
00446     st->size = ip->size;
00447 }
```

Referenced by [filestat\(\)](#).

5.59.2.25 writei()

```
int writei (
    struct inode * ip,
    char * src,
    uint off,
    uint n )
```

Definition at line 482 of file [fs.c](#).

```
00483 {
00484     uint tot, m;
00485     struct buf *bp;
00486
00487     if(ip->type == T_DEV){
00488         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
00489             return -1;
00490         return devsw[ip->major].write(ip, src, n);
00491     }
00492
00493     if(off > ip->size || off + n < off)
00494         return -1;
00495     if(off + n > MAXFILE*BSIZE)
00496         return -1;
00497
00498     for(tot=0; tot<n; tot+=m, off+=m, src+=m){
00499         bp = bread(ip->dev, bmap(ip, off/BSIZE));
00500         m = min(n - tot, BSIZE - off%BSIZE);
00501         memmove(bp->data + off%BSIZE, src, m);
00502         log_write(bp);
00503         brelse(bp);
00504     }
00505
00506     if(n > 0 && off > ip->size){
00507         ip->size = off;
00508         iupdate(ip);
00509     }
00510     return n;
00511 }
```

Referenced by [dirlink\(\)](#), [filewrite\(\)](#), and [sys_unlink\(\)](#).

5.59.3 Variable Documentation

5.59.3.1

```
struct { ... } icache
```

Referenced by [idup\(\)](#), [iget\(\)](#), [iinit\(\)](#), and [iput\(\)](#).

5.59.3.2 inode

```
struct inode inode[NINODE]
```

Definition at line 168 of file [fs.c](#).

5.59.3.3 lock

```
struct spinlock lock
```

Definition at line 167 of file [fs.c](#).

5.59.3.4 sb

```
struct superblock sb
```

Definition at line 28 of file [fs.c](#).

Referenced by [balloc\(\)](#), [bfree\(\)](#), [cmostime\(\)](#), [ialloc\(\)](#), [iinit\(\)](#), [ilock\(\)](#), [initlog\(\)](#), [iupdate\(\)](#), and [readsb\(\)](#).

5.60 fs.c

[Go to the documentation of this file.](#)

```
00001 // File system implementation. Five layers:
00002 //   + Blocks: allocator for raw disk blocks.
00003 //   + Log: crash recovery for multi-step updates.
00004 //   + Files: inode allocator, reading, writing, metadata.
00005 //   + Directories: inode with special contents (list of other inodes!)
00006 //   + Names: paths like /usr/rtn/xv6/fs.c for convenient naming.
00007 //
00008 // This file contains the low-level file system manipulation
00009 // routines. The (higher-level) system call implementations
00010 // are in sysfile.c.
00011
00012 #include "types.h"
00013 #include "defs.h"
00014 #include "param.h"
00015 #include "stat.h"
00016 #include "mmu.h"
00017 #include "proc.h"
00018 #include "spinlock.h"
00019 #include "sleeplock.h"
00020 #include "fs.h"
00021 #include "buf.h"
00022 #include "file.h"
00023
```



```

00024 #define min(a, b) ((a) < (b) ? (a) : (b))
00025 static void itrunc(struct inode*);
00026 // there should be one superblock per disk device, but we run with
00027 // only one device
00028 struct superblock sb;
00029
00030 // Read the super block.
00031 void
00032 readsb(int dev, struct superblock *sb)
00033 {
00034     struct buf *bp;
00035
00036     bp = bread(dev, 1);
00037     memmove(sb, bp->data, sizeof(*sb));
00038     brelse(bp);
00039 }
00040
00041 // Zero a block.
00042 static void
00043 bzero(int dev, int bno)
00044 {
00045     struct buf *bp;
00046
00047     bp = bread(dev, bno);
00048     memset(bp->data, 0, BSIZE);
00049     log_write(bp);
00050     brelse(bp);
00051 }
00052
00053 // Blocks.
00054
00055 // Allocate a zeroed disk block.
00056 static uint
00057 balloc(uint dev)
00058 {
00059     int b, bi, m;
00060     struct buf *bp;
00061
00062     bp = 0;
00063     for(b = 0; b < sb.size; b += BPB){
00064         bp = bread(dev, BBLOCK(b, sb));
00065         for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
00066             m = 1 << (bi % 8);
00067             if((bp->data[bi/8] & m) == 0){ // Is block free?
00068                 bp->data[bi/8] |= m; // Mark block in use.
00069                 log_write(bp);
00070                 brelse(bp);
00071                 bzero(dev, b + bi);
00072                 return b + bi;
00073             }
00074         }
00075         brelse(bp);
00076     }
00077     panic("balloc: out of blocks");
00078 }
00079
00080 // Free a disk block.
00081 static void
00082 bfree(int dev, uint b)
00083 {
00084     struct buf *bp;
00085     int bi, m;
00086
00087     bp = bread(dev, BBLOCK(b, sb));
00088     bi = b % BPB;
00089     m = 1 << (bi % 8);
00090     if((bp->data[bi/8] & m) == 0)
00091         panic("freeing free block");
00092     bp->data[bi/8] &= ~m;
00093     log_write(bp);
00094     brelse(bp);
00095 }
00096
00097 // Inodes.
00098 //
00099 // An inode describes a single unnamed file.
00100 // The inode disk structure holds metadata: the file's type,
00101 // its size, the number of links referring to it, and the
00102 // list of blocks holding the file's content.
00103 //
00104 // The inodes are laid out sequentially on disk at
00105 // sb.startinode. Each inode has a number, indicating its
00106 // position on the disk.
00107 //
00108 // The kernel keeps a cache of in-use inodes in memory
00109 // to provide a place for synchronizing access
00110 // to inodes used by multiple processes. The cached

```

```

00111 // inodes include book-keeping information that is
00112 // not stored on disk: ip->ref and ip->valid.
00113 //
00114 // An inode and its in-memory representation go through a
00115 // sequence of states before they can be used by the
00116 // rest of the file system code.
00117 //
00118 // * Allocation: an inode is allocated if its type (on disk)
00119 // is non-zero. ialloc() allocates, and iput() frees if
00120 // the reference and link counts have fallen to zero.
00121 //
00122 // * Referencing in cache: an entry in the inode cache
00123 // is free if ip->ref is zero. Otherwise ip->ref tracks
00124 // the number of in-memory pointers to the entry (open
00125 // files and current directories). iget() finds or
00126 // creates a cache entry and increments its ref; iput()
00127 // decrements ref.
00128 //
00129 // * Valid: the information (type, size, &c) in an inode
00130 // cache entry is only correct when ip->valid is 1.
00131 // ilock() reads the inode from
00132 // the disk and sets ip->valid, while iput() clears
00133 // ip->valid if ip->ref has fallen to zero.
00134 //
00135 // * Locked: file system code may only examine and modify
00136 // the information in an inode and its content if it
00137 // has first locked the inode.
00138 //
00139 // Thus a typical sequence is:
00140 //   ip = iget(dev, inum)
00141 //   ilock(ip)
00142 //   ... examine and modify ip->xxx ...
00143 //   iunlock(ip)
00144 //   iput(ip)
00145 //
00146 // ilock() is separate from iget() so that system calls can
00147 // get a long-term reference to an inode (as for an open file)
00148 // and only lock it for short periods (e.g., in read()).
00149 // The separation also helps avoid deadlock and races during
00150 // pathname lookup. iget() increments ip->ref so that the inode
00151 // stays cached and pointers to it remain valid.
00152 //
00153 // Many internal file system functions expect the caller to
00154 // have locked the inodes involved; this lets callers create
00155 // multi-step atomic operations.
00156 //
00157 // The icache.lock spin-lock protects the allocation of icache
00158 // entries. Since ip->ref indicates whether an entry is free,
00159 // and ip->dev and ip->inum indicate which i-node an entry
00160 // holds, one must hold icache.lock while using any of those fields.
00161 //
00162 // An ip->lock sleep-lock protects all ip-> fields other than ref,
00163 // dev, and inum. One must hold ip->lock in order to
00164 // read or write that inode's ip->valid, ip->size, ip->type, &c.
00165
00166 struct {
00167     struct spinlock lock;
00168     struct inode inode[NINODE];
00169 } icache;
00170
00171 void
00172 iinit(int dev)
00173 {
00174     int i = 0;
00175
00176     initlock(&icache.lock, "icache");
00177     for(i = 0; i < NINODE; i++) {
00178         initsleeplock(&icache.inode[i].lock, "inode");
00179     }
00180
00181     readsb(dev, &sb);
00182     cprintf("sb: size %d nblocks %d ninodes %d nlog %d logstart %d\
00183 inodestart %d bmap start %d\n", sb.size, sb.nblocks,
00184         sb.ninodes, sb.nlog, sb.logstart, sb.inodestart,
00185         sb.bmapstart);
00186 }
00187
00188 static struct inode* iget(uint dev, uint inum);
00189
00190 //PAGEBREAK!
00191 // Allocate an inode on device dev.
00192 // Mark it as allocated by giving it type type.
00193 // Returns an unlocked but allocated and referenced inode.
00194 struct inode*
00195 ialloc(uint dev, short type)
00196 {
00197     int inum;

```

```

00198 struct buf *bp;
00199 struct dinode *dip;
00200
00201 for(inum = 1; inum < sb.ninodes; inum++){
00202     bp = bread(dev, IBLOCK(inum, sb));
00203     dip = (struct dinode*)bp->data + inum%IPB;
00204     if(dip->type == 0){ // a free inode
00205         memset(dip, 0, sizeof(*dip));
00206         dip->type = type;
00207         log_write(bp); // mark it allocated on the disk
00208         brelse(bp);
00209         return iget(dev, inum);
00210     }
00211     brelse(bp);
00212 }
00213 panic("ialloc: no inodes");
00214 }
00215
00216 // Copy a modified in-memory inode to disk.
00217 // Must be called after every change to an ip->xxx field
00218 // that lives on disk, since i-node cache is write-through.
00219 // Caller must hold ip->lock.
00220 void
00221 iupdate(struct inode *ip)
00222 {
00223     struct buf *bp;
00224     struct dinode *dip;
00225
00226     bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00227     dip = (struct dinode*)bp->data + ip->inum%IPB;
00228     dip->type = ip->type;
00229     dip->major = ip->major;
00230     dip->minor = ip->minor;
00231     dip->nlink = ip->nlink;
00232     dip->size = ip->size;
00233     memmove(dip->addrs, ip->addrs, sizeof(ip->addrs));
00234     log_write(bp);
00235     brelse(bp);
00236 }
00237
00238 // Find the inode with number inum on device dev
00239 // and return the in-memory copy. Does not lock
00240 // the inode and does not read it from disk.
00241 static struct inode*
00242 iget(uint dev, uint inum)
00243 {
00244     struct inode *ip, *empty;
00245
00246     acquire(&icache.lock);
00247
00248     // Is the inode already cached?
00249     empty = 0;
00250     for(ip = &icache.inode[0]; ip < &icache.inode[NINODE]; ip++){
00251         if(ip->ref > 0 && ip->dev == dev && ip->inum == inum){
00252             ip->ref++;
00253             release(&icache.lock);
00254             return ip;
00255         }
00256         if(empty == 0 && ip->ref == 0) // Remember empty slot.
00257             empty = ip;
00258     }
00259
00260     // Recycle an inode cache entry.
00261     if(empty == 0)
00262         panic("iget: no inodes");
00263
00264     ip = empty;
00265     ip->dev = dev;
00266     ip->inum = inum;
00267     ip->ref = 1;
00268     ip->valid = 0;
00269     release(&icache.lock);
00270
00271     return ip;
00272 }
00273
00274 // Increment reference count for ip.
00275 // Returns ip to enable ip = idup(ip1) idiom.
00276 struct inode*
00277 idup(struct inode *ip)
00278 {
00279     acquire(&icache.lock);
00280     ip->ref++;
00281     release(&icache.lock);
00282     return ip;
00283 }
00284

```

```

00285 // Lock the given inode.
00286 // Reads the inode from disk if necessary.
00287 void
00288 ilock(struct inode *ip)
00289 {
00290     struct buf *bp;
00291     struct dinode *dip;
00292
00293     if(ip == 0 || ip->ref < 1)
00294         panic("ilock");
00295
00296     acquiresleep(&ip->lock);
00297
00298     if(ip->valid == 0){
00299         bp = bread(ip->dev, IBLOCK(ip->inum, sb));
00300         dip = (struct dinode*)bp->data + ip->inum%IPB;
00301         ip->type = dip->type;
00302         ip->major = dip->major;
00303         ip->minor = dip->minor;
00304         ip->nlink = dip->nlink;
00305         ip->size = dip->size;
00306         memmove(ip->addrs, dip->addrs, sizeof(ip->addrs));
00307         brelse(bp);
00308         ip->valid = 1;
00309         if(ip->type == 0)
00310             panic("ilock: no type");
00311     }
00312 }
00313
00314 // Unlock the given inode.
00315 void
00316 iunlock(struct inode *ip)
00317 {
00318     if(ip == 0 || !holdingsleep(&ip->lock) || ip->ref < 1)
00319         panic("iunlock");
00320
00321     releasesleep(&ip->lock);
00322 }
00323
00324 // Drop a reference to an in-memory inode.
00325 // If that was the last reference, the inode cache entry can
00326 // be recycled.
00327 // If that was the last reference and the inode has no links
00328 // to it, free the inode (and its content) on disk.
00329 // All calls to iput() must be inside a transaction in
00330 // case it has to free the inode.
00331 void
00332 iput(struct inode *ip)
00333 {
00334     acquiresleep(&ip->lock);
00335     if(ip->valid && ip->nlink == 0){
00336         acquire(&icache.lock);
00337         int r = ip->ref;
00338         release(&icache.lock);
00339         if(r == 1){
00340             // inode has no links and no other references: truncate and free.
00341             itrunc(ip);
00342             ip->type = 0;
00343             iupdate(ip);
00344             ip->valid = 0;
00345         }
00346     }
00347     releasesleep(&ip->lock);
00348
00349     acquire(&icache.lock);
00350     ip->ref--;
00351     release(&icache.lock);
00352 }
00353
00354 // Common idiom: unlock, then put.
00355 void
00356 iunlockput(struct inode *ip)
00357 {
00358     iunlock(ip);
00359     iput(ip);
00360 }
00361
00362 //PAGEBREAK!
00363 // Inode content
00364 //
00365 // The content (data) associated with each inode is stored
00366 // in blocks on the disk. The first NDIRECT block numbers
00367 // are listed in ip->addrs[]. The next NINDIRECT blocks are
00368 // listed in block ip->addrs[NDIRECT].
00369
00370 // Return the disk block address of the nth block in inode ip.
00371 // If there is no such block, bmap allocates one.

```

```

00372 static uint
00373 bmap(struct inode *ip, uint bn)
00374 {
00375     uint addr, *a;
00376     struct buf *bp;
00377
00378     if(bn < NDIRECT){
00379         if((addr = ip->addrs[bn]) == 0)
00380             ip->addrs[bn] = addr = balloc(ip->dev);
00381         return addr;
00382     }
00383     bn -= NDIRECT;
00384
00385     if(bn < NINDIRECT){
00386         // Load indirect block, allocating if necessary.
00387         if((addr = ip->addrs[NDIRECT]) == 0)
00388             ip->addrs[NDIRECT] = addr = balloc(ip->dev);
00389         bp = bread(ip->dev, addr);
00390         a = (uint*)bp->data;
00391         if((addr = a[bn]) == 0){
00392             a[bn] = addr = balloc(ip->dev);
00393             log_write(bp);
00394         }
00395         brelse(bp);
00396         return addr;
00397     }
00398
00399     panic("bmap: out of range");
00400 }
00401
00402 // Truncate inode (discard contents).
00403 // Only called when the inode has no links
00404 // to it (no directory entries referring to it)
00405 // and has no in-memory reference to it (is
00406 // not an open file or current directory).
00407 static void
00408 itrunc(struct inode *ip)
00409 {
00410     int i, j;
00411     struct buf *bp;
00412     uint *a;
00413
00414     for(i = 0; i < NDIRECT; i++){
00415         if(ip->addrs[i]){
00416             bfree(ip->dev, ip->addrs[i]);
00417             ip->addrs[i] = 0;
00418         }
00419     }
00420
00421     if(ip->addrs[NDIRECT]){
00422         bp = bread(ip->dev, ip->addrs[NDIRECT]);
00423         a = (uint*)bp->data;
00424         for(j = 0; j < NINDIRECT; j++){
00425             if(a[j])
00426                 bfree(ip->dev, a[j]);
00427         }
00428         brelse(bp);
00429         bfree(ip->dev, ip->addrs[NDIRECT]);
00430         ip->addrs[NDIRECT] = 0;
00431     }
00432
00433     ip->size = 0;
00434     iupdate(ip);
00435 }
00436
00437 // Copy stat information from inode.
00438 // Caller must hold ip->lock.
00439 void
00440 stati(struct inode *ip, struct stat *st)
00441 {
00442     st->dev = ip->dev;
00443     st->ino = ip->inum;
00444     st->type = ip->type;
00445     st->nlink = ip->nlink;
00446     st->size = ip->size;
00447 }
00448
00449 //PAGEBREAK!
00450 // Read data from inode.
00451 // Caller must hold ip->lock.
00452 int
00453 readi(struct inode *ip, char *dst, uint off, uint n)
00454 {
00455     uint tot, m;
00456     struct buf *bp;
00457
00458     if(ip->type == T_DEV){

```

```

00459     if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].read)
00460         return -1;
00461     return devsw[ip->major].read(ip, dst, n);
00462 }
00463
00464 if(off > ip->size || off + n < off)
00465     return -1;
00466 if(off + n > ip->size)
00467     n = ip->size - off;
00468
00469 for(tot=0; tot<n; tot+=m, off+=m, dst+=m){
00470     bp = bread(ip->dev, bmap(ip, off/BSIZE));
00471     m = min(n - tot, BSIZE - off%BSIZE);
00472     memmove(dst, bp->data + off%BSIZE, m);
00473     brelse(bp);
00474 }
00475 return n;
00476 }
00477
00478 // PAGEBREAK!
00479 // Write data to inode.
00480 // Caller must hold ip->lock.
00481 int
00482 writei(struct inode *ip, char *src, uint off, uint n)
00483 {
00484     uint tot, m;
00485     struct buf *bp;
00486
00487     if(ip->type == T_DEV){
00488         if(ip->major < 0 || ip->major >= NDEV || !devsw[ip->major].write)
00489             return -1;
00490         return devsw[ip->major].write(ip, src, n);
00491     }
00492
00493     if(off > ip->size || off + n < off)
00494         return -1;
00495     if(off + n > MAXFILE*BSIZE)
00496         return -1;
00497
00498     for(tot=0; tot<n; tot+=m, off+=m, src+=m){
00499         bp = bread(ip->dev, bmap(ip, off/BSIZE));
00500         m = min(n - tot, BSIZE - off%BSIZE);
00501         memmove(bp->data + off%BSIZE, src, m);
00502         log_write(bp);
00503         brelse(bp);
00504     }
00505
00506     if(n > 0 && off > ip->size){
00507         ip->size = off;
00508         iupdate(ip);
00509     }
00510     return n;
00511 }
00512
00513 //PAGEBREAK!
00514 // Directories
00515
00516 int
00517 namecmp(const char *s, const char *t)
00518 {
00519     return strncmp(s, t, DIRSIZ);
00520 }
00521
00522 // Look for a directory entry in a directory.
00523 // If found, set *poff to byte offset of entry.
00524 struct inode*
00525 dirlookup(struct inode *dp, char *name, uint *poff)
00526 {
00527     uint off, inum;
00528     struct dirent de;
00529
00530     if(dp->type != T_DIR)
00531         panic("dirlookup not DIR");
00532
00533     for(off = 0; off < dp->size; off += sizeof(de)){
00534         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00535             panic("dirlookup read");
00536         if(de.inum == 0)
00537             continue;
00538         if(namecmp(name, de.name) == 0){
00539             // entry matches path element
00540             if(poff)
00541                 *poff = off;
00542             inum = de.inum;
00543             return iget(dp->dev, inum);
00544         }
00545     }

```

```

00546
00547     return 0;
00548 }
00549
00550 // Write a new directory entry (name, inum) into the directory dp.
00551 int
00552 dirlink(struct inode *dp, char *name, uint inum)
00553 {
00554     int off;
00555     struct dirent de;
00556     struct inode *ip;
00557
00558     // Check that name is not present.
00559     if((ip = dirlookup(dp, name, 0)) != 0){
00560         iput(ip);
00561         return -1;
00562     }
00563
00564     // Look for an empty dirent.
00565     for(off = 0; off < dp->size; off += sizeof(de)){
00566         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00567             panic("dirlink read");
00568         if(de.inum == 0)
00569             break;
00570     }
00571
00572     strncpy(de.name, name, DIRSIZ);
00573     de.inum = inum;
00574     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00575         panic("dirlink");
00576
00577     return 0;
00578 }
00579
00580 //PAGEBREAK!
00581 // Paths
00582
00583 // Copy the next path element from path into name.
00584 // Return a pointer to the element following the copied one.
00585 // The returned path has no leading slashes,
00586 // so the caller can check *path=='\0' to see if the name is the last one.
00587 // If no name to remove, return 0.
00588 //
00589 // Examples:
00590 //   skipelem("a/bb/c", name) = "bb/c", setting name = "a"
00591 //   skipelem("///a//bb", name) = "bb", setting name = "a"
00592 //   skipelem("a", name) = "", setting name = "a"
00593 //   skipelem("", name) = skipelem("///", name) = 0
00594 //
00595 static char*
00596 skipelem(char *path, char *name)
00597 {
00598     char *s;
00599     int len;
00600
00601     while(*path == '/')
00602         path++;
00603     if(*path == 0)
00604         return 0;
00605     s = path;
00606     while(*path != '/' && *path != 0)
00607         path++;
00608     len = path - s;
00609     if(len >= DIRSIZ)
00610         memmove(name, s, DIRSIZ);
00611     else {
00612         memmove(name, s, len);
00613         name[len] = 0;
00614     }
00615     while(*path == '/')
00616         path++;
00617     return path;
00618 }
00619
00620 // Look up and return the inode for a path name.
00621 // If parent != 0, return the inode for the parent and copy the final
00622 // path element into name, which must have room for DIRSIZ bytes.
00623 // Must be called inside a transaction since it calls iput().
00624 static struct inode*
00625 nameex(char *path, int nameiparent, char *name)
00626 {
00627     struct inode *ip, *next;
00628
00629     if(*path == '/')
00630         ip = iget(ROOTDEV, ROOTINO);
00631     else
00632         ip = idup(myproc()->cwd);

```

```

00633
00634     while((path = skipelem(path, name)) != 0){
00635         ilock(ip);
00636         if(ip->type != T_DIR){
00637             iunlockput(ip);
00638             return 0;
00639         }
00640         if(nameiparent && *path == '\\0'){
00641             // Stop one level early.
00642             iunlock(ip);
00643             return ip;
00644         }
00645         if((next = dirlookup(ip, name, 0)) == 0){
00646             iunlockput(ip);
00647             return 0;
00648         }
00649         iunlockput(ip);
00650         ip = next;
00651     }
00652     if(nameiparent){
00653         iput(ip);
00654         return 0;
00655     }
00656     return ip;
00657 }
00658
00659 struct inode*
00660 namei(char *path)
00661 {
00662     char name[DIRSIZ];
00663     return namex(path, 0, name);
00664 }
00665
00666 struct inode*
00667 nameiparent(char *path, char *name)
00668 {
00669     return namex(path, 1, name);
00670 }

```

5.61 fs.d File Reference

5.62 fs.d

[Go to the documentation of this file.](#)

```

00001 fs.o: fs.c /usr/include/stdc-predef.h types.h defs.h param.h stat.h mmu.h \
00002 proc.h spinlock.h sleeplock.h fs.h buf.h file.h

```

5.63 fs.h File Reference

Classes

- struct [dinode](#)
- struct [dirent](#)
- struct [superblock](#)

Macros

- #define [BBLOCK](#)(b, sb) (b/BPB + sb.bmapstart)
- #define [BPB](#) (BSIZE*8)
- #define [BSIZE](#) 512
- #define [DIRSIZ](#) 14
- #define [IBLOCK](#)(i, sb) ((i) / IPB + sb.inodestart)
- #define [IPB](#) (BSIZE / sizeof(struct [dinode](#)))
- #define [MAXFILE](#) (NDIRECT + NINDIRECT)
- #define [NDIRECT](#) 12
- #define [NINDIRECT](#) (BSIZE / sizeof(uint))
- #define [ROOTINO](#) 1

5.63.1 Macro Definition Documentation

5.63.1.1 BBLOCK

```
#define BBLOCK(  
    b,  
    sb ) (b/BPB + sb.bmapstart)
```

Definition at line 48 of file [fs.h](#).

5.63.1.2 BPB

```
#define BPB (BSIZE*8)
```

Definition at line 45 of file [fs.h](#).

5.63.1.3 BSIZE

```
#define BSIZE 512
```

Definition at line 6 of file [fs.h](#).

5.63.1.4 DIRSIZ

```
#define DIRSIZ 14
```

Definition at line 51 of file [fs.h](#).

5.63.1.5 IBLOCK

```
#define IBLOCK(  
    i,  
    sb ) ((i) / IPB + sb.inodestart)
```

Definition at line 42 of file [fs.h](#).

5.63.1.6 IPB

```
#define IPB (BSIZE / sizeof(struct dinode))
```

Definition at line 39 of file [fs.h](#).

5.63.1.7 MAXFILE

```
#define MAXFILE (NDIRECT + NINDIRECT)
```

Definition at line 26 of file [fs.h](#).

5.63.1.8 NDIRECT

```
#define NDIRECT 12
```

Definition at line 24 of file [fs.h](#).

5.63.1.9 NINDIRECT

```
#define NINDIRECT (BSIZE / sizeof(uint))
```

Definition at line 25 of file [fs.h](#).

5.63.1.10 ROOTINO

```
#define ROOTINO 1
```

Definition at line 5 of file [fs.h](#).

5.64 fs.h

[Go to the documentation of this file.](#)

```

00001 // On-disk file system format.
00002 // Both the kernel and user programs use this header file.
00003
00004
00005 #define ROOTINO 1 // root i-number
00006 #define BSIZE 512 // block size
00007
00008 // Disk layout:
00009 // [ boot block | super block | log | inode blocks |
00010 //                               free bit map | data blocks]
00011 //
00012 // mkfs computes the super block and builds an initial file system. The
00013 // super block describes the disk layout:
00014 struct superblock {
00015     uint size; // Size of file system image (blocks)
00016     uint nblocks; // Number of data blocks
00017     uint ninodes; // Number of inodes.
00018     uint nlog; // Number of log blocks
00019     uint logstart; // Block number of first log block
00020     uint inodestart; // Block number of first inode block
00021     uint bmapstart; // Block number of first free map block
00022 };
00023
00024 #define NDIRECT 12
00025 #define NINDIRECT (BSIZE / sizeof(uint))
00026 #define MAXFILE (NDIRECT + NINDIRECT)
00027
00028 // On-disk inode structure
00029 struct dinode {
00030     short type; // File type
00031     short major; // Major device number (T_DEV only)
00032     short minor; // Minor device number (T_DEV only)
00033     short nlink; // Number of links to inode in file system
00034     uint size; // Size of file (bytes)
00035     uint addrs[NDIRECT+1]; // Data block addresses
00036 };
00037
00038 // Inodes per block.
00039 #define IPB (BSIZE / sizeof(struct dinode))
00040
00041 // Block containing inode i
00042 #define IBLOCK(i, sb) ((i) / IPB + sb.inodestart)
00043
00044 // Bitmap bits per block
00045 #define BPB (BSIZE*8)
00046
00047 // Block of free map containing bit for block b
00048 #define BBLOCK(b, sb) (b/BPB + sb.bmapstart)
00049
00050 // Directory is a file containing a sequence of dirent structures.
00051 #define DIRSIZ 14
00052
00053 struct dirent {
00054     ushort inum;
00055     char name[DIRSIZ];
00056 };
00057

```

5.65 grep.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"

```

Functions

- void [grep](#) (char *pattern, int fd)
- int [main](#) (int argc, char *argv[])
- int [match](#) (char *, char *)
- int [matchhere](#) (char *, char *)
- int [matchstar](#) (int, char *, char *)

Variables

- char `buf`[1024]

5.65.1 Function Documentation

5.65.1.1 `grep()`

```
void grep (  
    char * pattern,  
    int fd )
```

Definition at line 11 of file `grep.c`.

```
00012 {  
00013     int n, m;  
00014     char *p, *q;  
00015  
00016     m = 0;  
00017     while((n = read(fd, buf+m, sizeof(buf)-m-1)) > 0){  
00018         m += n;  
00019         buf[m] = '\0';  
00020         p = buf;  
00021         while((q = strchr(p, '\n')) != 0){  
00022             *q = 0;  
00023             if(match(pattern, p)){  
00024                 *q = '\n';  
00025                 write(1, p, q+1 - p);  
00026             }  
00027             p = q+1;  
00028         }  
00029         if(p == buf)  
00030             m = 0;  
00031         if(m > 0){  
00032             m -= p - buf;  
00033             memmove(buf, p, m);  
00034         }  
00035     }  
00036 }
```

Referenced by `main()`.

5.65.1.2 `main()`

```
int main (  
    int argc,  
    char * argv[] )
```

Definition at line 39 of file `grep.c`.

```
00040 {  
00041     int fd, i;  
00042     char *pattern;  
00043  
00044     if(argc <= 1){  
00045         printf(2, "usage: grep pattern [file ...]\n");  
00046         exit();  
00047     }  
00048     pattern = argv[1];  
00049  
00050     if(argc <= 2){  
00051         grep(pattern, 0);  
00052         exit();  
00053     }
```

```
00054
00055     for(i = 2; i < argc; i++){
00056         if((fd = open(argv[i], 0)) < 0){
00057             printf(1, "grep: cannot open %s\n", argv[i]);
00058             exit();
00059         }
00060         grep(pattern, fd);
00061         close(fd);
00062     }
00063     exit();
00064 }
```

5.65.1.3 match()

```
int match (
    char * re,
    char * text )
```

Definition at line 73 of file [grep.c](#).

```
00074 {
00075     if(re[0] == '^')
00076         return matchhere(re+1, text);
00077     do{ // must look at empty string
00078         if(matchhere(re, text))
00079             return 1;
00080     }while(*text++ != '\0');
00081     return 0;
00082 }
```

Referenced by [grep\(\)](#).

5.65.1.4 matchhere()

```
int matchhere (
    char * re,
    char * text )
```

Definition at line 85 of file [grep.c](#).

```
00086 {
00087     if(re[0] == '\0')
00088         return 1;
00089     if(re[1] == '*')
00090         return matchstar(re[0], re+2, text);
00091     if(re[0] == '$' && re[1] == '\0')
00092         return *text == '\0';
00093     if(*text != '\0' && (re[0] == '.' || re[0] == *text))
00094         return matchhere(re+1, text+1);
00095     return 0;
00096 }
```

Referenced by [match\(\)](#), [matchhere\(\)](#), and [matchstar\(\)](#).

5.65.1.5 matchstar()

```
int matchstar (
    int c,
    char * re,
    char * text )
```

Definition at line 99 of file [grep.c](#).

```
00100 {
00101     do{ // a * matches zero or more instances
00102         if(matchhere(re, text))
00103             return 1;
00104     }while(*text!='\0' && (*text++==c || c=='.'));
00105     return 0;
00106 }
```

Referenced by [matchhere\(\)](#).

5.65.2 Variable Documentation

5.65.2.1 buf

```
char buf[1024]
```

Definition at line 7 of file [grep.c](#).

Referenced by [grep\(\)](#).

5.66 grep.c

[Go to the documentation of this file.](#)

```
00001 // Simple grep. Only supports ^ . * $ operators.
00002
00003 #include "types.h"
00004 #include "stat.h"
00005 #include "user.h"
00006
00007 char buf[1024];
00008 int match(char*, char*);
00009
00010 void
00011 grep(char *pattern, int fd)
00012 {
00013     int n, m;
00014     char *p, *q;
00015
00016     m = 0;
00017     while((n = read(fd, buf+m, sizeof(buf)-m-1)) > 0){
00018         m += n;
00019         buf[m] = '\0';
00020         p = buf;
00021         while((q = strchr(p, '\n')) != 0){
00022             *q = 0;
00023             if(match(pattern, p)){
00024                 *q = '\n';
00025                 write(1, p, q+1 - p);
00026             }
00027             p = q+1;
00028         }
00029         if(p == buf)
00030             m = 0;
00031         if(m > 0){
00032             m -= p - buf;
```

```

00033     memmove(buf, p, m);
00034 }
00035 }
00036 }
00037
00038 int
00039 main(int argc, char *argv[])
00040 {
00041     int fd, i;
00042     char *pattern;
00043
00044     if(argc <= 1){
00045         printf(2, "usage: grep pattern [file ...]\n");
00046         exit();
00047     }
00048     pattern = argv[1];
00049
00050     if(argc <= 2){
00051         grep(pattern, 0);
00052         exit();
00053     }
00054
00055     for(i = 2; i < argc; i++){
00056         if((fd = open(argv[i], 0)) < 0){
00057             printf(1, "grep: cannot open %s\n", argv[i]);
00058             exit();
00059         }
00060         grep(pattern, fd);
00061         close(fd);
00062     }
00063     exit();
00064 }
00065
00066 // Regexp matcher from Kernighan & Pike,
00067 // The Practice of Programming, Chapter 9.
00068
00069 int matchhere(char*, char*);
00070 int matchstar(int, char*, char*);
00071
00072 int
00073 match(char *re, char *text)
00074 {
00075     if(re[0] == '^')
00076         return matchhere(re+1, text);
00077     do{ // must look at empty string
00078         if(matchhere(re, text))
00079             return 1;
00080     }while(*text++ != '\0');
00081     return 0;
00082 }
00083
00084 // matchhere: search for re at beginning of text
00085 int matchhere(char *re, char *text)
00086 {
00087     if(re[0] == '\0')
00088         return 1;
00089     if(re[1] == '*')
00090         return matchstar(re[0], re+2, text);
00091     if(re[0] == '$' && re[1] == '\0')
00092         return *text == '\0';
00093     if(*text != '\0' && (re[0] == '.' || re[0] == *text))
00094         return matchhere(re+1, text+1);
00095     return 0;
00096 }
00097
00098 // matchstar: search for c*re at beginning of text
00099 int matchstar(int c, char *re, char *text)
00100 {
00101     do{ // a * matches zero or more instances
00102         if(matchhere(re, text))
00103             return 1;
00104     }while(*text++ != '\0' && (*text++ == c || c == '.'));
00105     return 0;
00106 }
00107

```

5.67 grep.d File Reference

5.68 grep.d

[Go to the documentation of this file.](#)

```
00001 grep.o: grep.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.69 ide.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"
#include "traps.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "buf.h"
```

Macros

- #define [IDE_BSY](#) 0x80
- #define [IDE_CMD_RDMUL](#) 0xc4
- #define [IDE_CMD_READ](#) 0x20
- #define [IDE_CMD_WRITE](#) 0x30
- #define [IDE_CMD_WRMUL](#) 0xc5
- #define [IDE_DF](#) 0x20
- #define [IDE_DRDY](#) 0x40
- #define [IDE_ERR](#) 0x01
- #define [SECTOR_SIZE](#) 512

Functions

- void [ideinit](#) (void)
- void [ideintr](#) (void)
- void [iderw](#) (struct [buf](#) *b)
- static void [idestart](#) (struct [buf](#) *)
- static int [idewait](#) (int checkerr)

Variables

- static int [havedisk1](#)
- static struct [spinlock](#) [idelock](#)
- static struct [buf](#) * [idequeue](#)

5.69.1 Macro Definition Documentation

5.69.1.1 IDE_BSY

```
#define IDE_BSY 0x80
```

Definition at line 17 of file [ide.c](#).

5.69.1.2 IDE_CMD_RDMUL

```
#define IDE_CMD_RDMUL 0xc4
```

Definition at line 24 of file [ide.c](#).

5.69.1.3 IDE_CMD_READ

```
#define IDE_CMD_READ 0x20
```

Definition at line 22 of file [ide.c](#).

5.69.1.4 IDE_CMD_WRITE

```
#define IDE_CMD_WRITE 0x30
```

Definition at line 23 of file [ide.c](#).

5.69.1.5 IDE_CMD_WRMUL

```
#define IDE_CMD_WRMUL 0xc5
```

Definition at line 25 of file [ide.c](#).

5.69.1.6 IDE_DF

```
#define IDE_DF 0x20
```

Definition at line 19 of file [ide.c](#).

5.69.1.7 IDE_DRDY

```
#define IDE_DRDY 0x40
```

Definition at line 18 of file [ide.c](#).

5.69.1.8 IDE_ERR

```
#define IDE_ERR 0x01
```

Definition at line 20 of file [ide.c](#).

5.69.1.9 SECTOR_SIZE

```
#define SECTOR_SIZE 512
```

Definition at line 16 of file [ide.c](#).

5.69.2 Function Documentation

5.69.2.1 ideinit()

```
void ideinit (  
    void )
```

Definition at line 51 of file [ide.c](#).

```
00052 {  
00053     int i;  
00054  
00055     initlock(&idelock, "ide");  
00056     ioapicenable(IRQ_IDE, ncpu - 1);  
00057     idewait(0);  
00058  
00059     // Check if disk 1 is present  
00060     outb(0x1f6, 0xe0 | (1<<4));  
00061     for(i=0; i<1000; i++){  
00062         if(inb(0x1f7) != 0){  
00063             havedisk1 = 1;  
00064             break;  
00065         }  
00066     }  
00067  
00068     // Switch back to disk 0.  
00069     outb(0x1f6, 0xe0 | (0<<4));  
00070 }
```

5.69.2.2 ideintr()

```
void ideintr (
    void )
```

Definition at line 104 of file [ide.c](#).

```
00105 {
00106     struct buf *b;
00107
00108     // First queued buffer is the active request.
00109     acquire(&idelock);
00110
00111     if((b = idequeue) == 0){
00112         release(&idelock);
00113         return;
00114     }
00115     idequeue = b->qnext;
00116
00117     // Read data if needed.
00118     if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
00119         insl(0x1f0, b->data, BSIZE/4);
00120
00121     // Wake process waiting for this buf.
00122     b->flags |= B_VALID;
00123     b->flags &= ~B_DIRTY;
00124     wakeup(b);
00125
00126     // Start disk on next buf in queue.
00127     if(idequeue != 0)
00128         idestart(idequeue);
00129
00130     release(&idelock);
00131 }
```

5.69.2.3 iderw()

```
void iderw (
    struct buf * b )
```

Definition at line 138 of file [ide.c](#).

```
00139 {
00140     struct buf **pp;
00141
00142     if(!holdingsleep(&b->lock))
00143         panic("iderw: buf not locked");
00144     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
00145         panic("iderw: nothing to do");
00146     if(b->dev != 0 && !havedisk1)
00147         panic("iderw: ide disk 1 not present");
00148
00149     acquire(&idelock); //DOC:acquire-lock
00150
00151     // Append b to idequeue.
00152     b->qnext = 0;
00153     for(pp=&idequeue; *pp; pp=(*pp)->qnext) //DOC:insert-queue
00154         ;
00155     *pp = b;
00156
00157     // Start disk if necessary.
00158     if(idequeue == b)
00159         idestart(b);
00160
00161     // Wait for request to finish.
00162     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
00163         sleep(b, &idelock);
00164     }
00165
00166     release(&idelock);
00167
00168 }
```

5.69.2.4 idestart()

```
static void idestart (
    struct buf * b ) [static]
```

Definition at line 74 of file [ide.c](#).

```
00075 {
00076     if(b == 0)
00077         panic("idestart");
00078     if(b->blockno >= FSSIZE)
00079         panic("incorrect blockno");
00080     int sector_per_block = BSIZE/SECTOR_SIZE;
00081     int sector = b->blockno * sector_per_block;
00082     int read_cmd = (sector_per_block == 1) ? IDE_CMD_READ : IDE_CMD_RDMD;
00083     int write_cmd = (sector_per_block == 1) ? IDE_CMD_WRITE : IDE_CMD_WRMD;
00084
00085     if (sector_per_block > 7) panic("idestart");
00086
00087     idewait(0);
00088     outb(0x3f6, 0); // generate interrupt
00089     outb(0x1f2, sector_per_block); // number of sectors
00090     outb(0x1f3, sector & 0xff);
00091     outb(0x1f4, (sector >> 8) & 0xff);
00092     outb(0x1f5, (sector >> 16) & 0xff);
00093     outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((sector>>24)&0x0f));
00094     if(b->flags & B_DIRTY){
00095         outb(0x1f7, write_cmd);
00096         outsl(0x1f0, b->data, BSIZE/4);
00097     } else {
00098         outb(0x1f7, read_cmd);
00099     }
00100 }
```

Referenced by [ideintr\(\)](#), and [iderw\(\)](#).

5.69.2.5 idewait()

```
static int idewait (
    int checkerr ) [static]
```

Definition at line 39 of file [ide.c](#).

```
00040 {
00041     int r;
00042
00043     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)
00044         ;
00045     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)
00046         return -1;
00047     return 0;
00048 }
```

Referenced by [ideinit\(\)](#), [ideintr\(\)](#), and [idestart\(\)](#).

5.69.3 Variable Documentation

5.69.3.1 havedisk1

```
int havedisk1 [static]
```

Definition at line 34 of file [ide.c](#).

Referenced by [ideinit\(\)](#), and [iderw\(\)](#).

5.69.3.2 idelock

```
struct spinlock idelock [static]
```

Definition at line 31 of file [ide.c](#).

Referenced by [ideinit\(\)](#), [ideintr\(\)](#), and [iderw\(\)](#).

5.69.3.3 idequeue

```
struct buf* idequeue [static]
```

Definition at line 32 of file [ide.c](#).

Referenced by [ideintr\(\)](#), and [iderw\(\)](#).

5.70 ide.c

[Go to the documentation of this file.](#)

```
00001 // Simple PIO-based (non-DMA) IDE driver code.
00002
00003 #include "types.h"
00004 #include "defs.h"
00005 #include "param.h"
00006 #include "memlayout.h"
00007 #include "mmu.h"
00008 #include "proc.h"
00009 #include "x86.h"
00010 #include "traps.h"
00011 #include "spinlock.h"
00012 #include "sleeplock.h"
00013 #include "fs.h"
00014 #include "buf.h"
00015
00016 #define SECTOR_SIZE 512
00017 #define IDE_BSY 0x80
00018 #define IDE_DRDY 0x40
00019 #define IDE_DF 0x20
00020 #define IDE_ERR 0x01
00021
00022 #define IDE_CMD_READ 0x20
00023 #define IDE_CMD_WRITE 0x30
00024 #define IDE_CMD_RDMUL 0xc4
00025 #define IDE_CMD_WRMUL 0xc5
00026
00027 // idequeue points to the buf now being read/written to the disk.
00028 // idequeue->qnext points to the next buf to be processed.
00029 // You must hold idelock while manipulating queue.
00030
00031 static struct spinlock idelock;
00032 static struct buf *idequeue;
00033
00034 static int havedisk1;
00035 static void idestart(struct buf*);
00036
00037 // Wait for IDE disk to become ready.
00038 static int
00039 idewait(int checkerr)
00040 {
00041     int r;
00042
00043     while(((r = inb(0x1f7)) & (IDE_BSY|IDE_DRDY)) != IDE_DRDY)
00044         ;
00045     if(checkerr && (r & (IDE_DF|IDE_ERR)) != 0)
00046         return -1;
00047     return 0;
00048 }
00049
```

```

00050 void
00051 ideinit(void)
00052 {
00053     int i;
00054
00055     initlock(&idelock, "ide");
00056     ioapicenable(IRQ_IDE, ncpu - 1);
00057     idewait(0);
00058
00059     // Check if disk 1 is present
00060     outb(0x1f6, 0xe0 | (1<<4));
00061     for(i=0; i<1000; i++){
00062         if(inb(0x1f7) != 0){
00063             havedisk1 = 1;
00064             break;
00065         }
00066     }
00067
00068     // Switch back to disk 0.
00069     outb(0x1f6, 0xe0 | (0<<4));
00070 }
00071
00072 // Start the request for b. Caller must hold idelock.
00073 static void
00074 idestart(struct buf *b)
00075 {
00076     if(b == 0)
00077         panic("idestart");
00078     if(b->blockno >= FSSIZE)
00079         panic("incorrect blockno");
00080     int sector_per_block = BSIZE/SECTOR_SIZE;
00081     int sector = b->blockno * sector_per_block;
00082     int read_cmd = (sector_per_block == 1) ? IDE_CMD_READ : IDE_CMD_RDMUL;
00083     int write_cmd = (sector_per_block == 1) ? IDE_CMD_WRITE : IDE_CMD_WRMUL;
00084
00085     if (sector_per_block > 7) panic("idestart");
00086
00087     idewait(0);
00088     outb(0x3f6, 0); // generate interrupt
00089     outb(0x1f2, sector_per_block); // number of sectors
00090     outb(0x1f3, sector & 0xff);
00091     outb(0x1f4, (sector >> 8) & 0xff);
00092     outb(0x1f5, (sector >> 16) & 0xff);
00093     outb(0x1f6, 0xe0 | ((b->dev<1)<<4) | ((sector<24)&0x0f));
00094     if(b->flags & B_DIRTY){
00095         outb(0x1f7, write_cmd);
00096         outsl(0x1f0, b->data, BSIZE/4);
00097     } else {
00098         outb(0x1f7, read_cmd);
00099     }
00100 }
00101
00102 // Interrupt handler.
00103 void
00104 ideintr(void)
00105 {
00106     struct buf *b;
00107
00108     // First queued buffer is the active request.
00109     acquire(&idelock);
00110
00111     if((b = idequeue) == 0){
00112         release(&idelock);
00113         return;
00114     }
00115     idequeue = b->qnext;
00116
00117     // Read data if needed.
00118     if(!(b->flags & B_DIRTY) && idewait(1) >= 0)
00119         insl(0x1f0, b->data, BSIZE/4);
00120
00121     // Wake process waiting for this buf.
00122     b->flags |= B_VALID;
00123     b->flags &= ~B_DIRTY;
00124     wakeup(b);
00125
00126     // Start disk on next buf in queue.
00127     if(idequeue != 0)
00128         idestart(idequeue);
00129
00130     release(&idelock);
00131 }
00132
00133 //PAGEBREAK!
00134 // Sync buf with disk.
00135 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
00136 // Else if B_VALID is not set, read buf from disk, set B_VALID.

```

```

00137 void
00138 iderw(struct buf *b)
00139 {
00140     struct buf **pp;
00141
00142     if(!holdingsleep(&b->lock))
00143         panic("iderw: buf not locked");
00144     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
00145         panic("iderw: nothing to do");
00146     if(b->dev != 0 && !havedisk1)
00147         panic("iderw: ide disk 1 not present");
00148
00149     acquire(&idelock); //DOC:acquire-lock
00150
00151     // Append b to idequeue.
00152     b->qnext = 0;
00153     for(pp=&idequeue; *pp; pp=(*pp)->qnext) //DOC:insert-queue
00154         ;
00155     *pp = b;
00156
00157     // Start disk if necessary.
00158     if(idequeue == b)
00159         idestart(b);
00160
00161     // Wait for request to finish.
00162     while((b->flags & (B_VALID|B_DIRTY)) != B_VALID){
00163         sleep(b, &idelock);
00164     }
00165
00166
00167     release(&idelock);
00168 }

```

5.71 ide.d File Reference

5.72 ide.d

[Go to the documentation of this file.](#)

```

00001 ide.o: ide.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 memlayout.h mmu.h proc.h x86.h traps.h spinlock.h sleeplock.h fs.h buf.h

```

5.73 init.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

```

Functions

- int [main](#) (void)

Variables

- char * [argv](#) [] = { "sh", 0 }

5.73.1 Function Documentation

5.73.1.1 main()

```
int main (
    void )
```

Definition at line 11 of file [init.c](#).

```
00012 {
00013     int pid, wpid;
00014
00015     if(open("console", O_RDWR) < 0){
00016         mknod("console", 1, 1);
00017         open("console", O_RDWR);
00018     }
00019     dup(0); // stdout
00020     dup(0); // stderr
00021
00022     for(;;){
00023         printf(1, "init: starting sh\n");
00024         pid = fork();
00025         if(pid < 0){
00026             printf(1, "init: fork failed\n");
00027             exit();
00028         }
00029         if(pid == 0){
00030             exec("sh", argv);
00031             printf(1, "init: exec sh failed\n");
00032             exit();
00033         }
00034         while((wpid=wait()) >= 0 && wpid != pid)
00035             printf(1, "zombie!\n");
00036     }
00037 }
```

5.73.2 Variable Documentation

5.73.2.1 argv

```
char* argv[] = { "sh", 0 }
```

Definition at line 8 of file [init.c](#).

Referenced by [exec\(\)](#), [main\(\)](#), and [sys_exec\(\)](#).

5.74 init.c

[Go to the documentation of this file.](#)

```
00001 // init: The initial user-level program
00002
00003 #include "types.h"
00004 #include "stat.h"
00005 #include "user.h"
00006 #include "fcntl.h"
00007
00008 char *argv[] = { "sh", 0 };
00009
00010 int
00011 main(void)
00012 {
00013     int pid, wpid;
00014
00015     if(open("console", O_RDWR) < 0){
00016         mknod("console", 1, 1);
00017         open("console", O_RDWR);
00018     }
```



```

00019  dup(0); // stdout
00020  dup(0); // stderr
00021
00022  for(;;){
00023      printf(1, "init: starting sh\n");
00024      pid = fork();
00025      if(pid < 0){
00026          printf(1, "init: fork failed\n");
00027          exit();
00028      }
00029      if(pid == 0){
00030          exec("sh", argv);
00031          printf(1, "init: exec sh failed\n");
00032          exit();
00033      }
00034      while((wpid=wait()) >= 0 && wpid != pid)
00035          printf(1, "zombie!\n");
00036  }
00037 }

```

5.75 init.d File Reference

5.76 init.d

[Go to the documentation of this file.](#)

```
00001 init.o: init.c /usr/include/stdc-predef.h types.h stat.h user.h fcntl.h
```

5.77 initcode.d File Reference

5.78 initcode.d

[Go to the documentation of this file.](#)

```
00001 initcode.o: initcode.S syscall.h traps.h
```

5.79 ioapic.c File Reference

```

#include "types.h"
#include "defs.h"
#include "traps.h"

```

Classes

- struct [ioapic](#)

Macros

- #define [INT_ACTIVELOW](#) 0x00002000
- #define [INT_DISABLED](#) 0x00010000
- #define [INT_LEVEL](#) 0x00008000
- #define [INT_LOGICAL](#) 0x00000800
- #define [IOAPIC](#) 0xFEC00000
- #define [REG_ID](#) 0x00
- #define [REG_TABLE](#) 0x10
- #define [REG_VER](#) 0x01

Functions

- void [ioapicenable](#) (int irq, int cpunum)
- void [ioapicinit](#) (void)
- static uint [ioapicread](#) (int reg)
- static void [ioapicwrite](#) (int reg, uint data)

Variables

- volatile struct [ioapic](#) * [ioapic](#)

5.79.1 Macro Definition Documentation

5.79.1.1 INT_ACTIVELOW

```
#define INT_ACTIVELOW 0x00002000
```

Definition at line 22 of file [ioapic.c](#).

5.79.1.2 INT_DISABLED

```
#define INT_DISABLED 0x00010000
```

Definition at line 20 of file [ioapic.c](#).

5.79.1.3 INT_LEVEL

```
#define INT_LEVEL 0x00008000
```

Definition at line 21 of file [ioapic.c](#).

5.79.1.4 INT_LOGICAL

```
#define INT_LOGICAL 0x00000800
```

Definition at line 23 of file [ioapic.c](#).

5.79.1.5 IOAPIC

```
#define IOAPIC 0xFEC00000
```

Definition at line 9 of file [ioapic.c](#).

5.79.1.6 REG_ID

```
#define REG_ID 0x00
```

Definition at line 11 of file [ioapic.c](#).

5.79.1.7 REG_TABLE

```
#define REG_TABLE 0x10
```

Definition at line 13 of file [ioapic.c](#).

5.79.1.8 REG_VER

```
#define REG_VER 0x01
```

Definition at line 12 of file [ioapic.c](#).

5.79.2 Function Documentation

5.79.2.1 ioapicenable()

```
void ioapicenable (  
    int irq,  
    int cpunum )
```

Definition at line 68 of file [ioapic.c](#).

```
00069 {  
00070     // Mark interrupt edge-triggered, active high,  
00071     // enabled, and routed to the given cpunum,  
00072     // which happens to be that cpu's APIC ID.  
00073     ioapicwrite(REG_TABLE+2*irq, T_IRQ0 + irq);  
00074     ioapicwrite(REG_TABLE+2*irq+1, cpunum « 24);  
00075 }
```

Referenced by [consoleinit\(\)](#), [ideinit\(\)](#), and [uartinit\(\)](#).

5.79.2.2 ioapicinit()

```
void ioapicinit (
    void )
```

Definition at line 49 of file [ioapic.c](#).

```
00050 {
00051     int i, id, maxintr;
00052
00053     ioapic = (volatile struct ioapic*)IOAPIC;
00054     maxintr = (ioapicread(REG_VER) » 16) & 0xFF;
00055     id = ioapicread(REG_ID) » 24;
00056     if(id != ioapicid)
00057         cprintf("ioapicinit: id isn't equal to ioapicid; not a MP\n");
00058
00059     // Mark all interrupts edge-triggered, active high, disabled,
00060     // and not routed to any CPUs.
00061     for(i = 0; i <= maxintr; i++){
00062         ioapicwrite(REG_TABLE+2*i, INT_DISABLED | (T_IRQ0 + i));
00063         ioapicwrite(REG_TABLE+2*i+1, 0);
00064     }
00065 }
```

5.79.2.3 ioapicread()

```
static uint ioapicread (
    int reg ) [static]
```

Definition at line 35 of file [ioapic.c](#).

```
00036 {
00037     ioapic->reg = reg;
00038     return ioapic->data;
00039 }
```

Referenced by [ioapicinit\(\)](#).

5.79.2.4 ioapicwrite()

```
static void ioapicwrite (
    int reg,
    uint data ) [static]
```

Definition at line 42 of file [ioapic.c](#).

```
00043 {
00044     ioapic->reg = reg;
00045     ioapic->data = data;
00046 }
```

Referenced by [ioapicenable\(\)](#), and [ioapicinit\(\)](#).

5.79.3 Variable Documentation

5.79.3.1 ioapic

volatile struct ioapic* ioapic

Definition at line 25 of file ioapic.c.

5.80 ioapic.c

[Go to the documentation of this file.](#)

```

00001 // The I/O APIC manages hardware interrupts for an SMP system.
00002 // http://www.intel.com/design/chipsets/datashts/29056601.pdf
00003 // See also picirq.c.
00004
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "traps.h"
00008
00009 #define IOAPIC 0xFEC00000 // Default physical address of IO APIC
00010
00011 #define REG_ID 0x00 // Register index: ID
00012 #define REG_VER 0x01 // Register index: version
00013 #define REG_TABLE 0x10 // Redirection table base
00014
00015 // The redirection table starts at REG_TABLE and uses
00016 // two registers to configure each interrupt.
00017 // The first (low) register in a pair contains configuration bits.
00018 // The second (high) register contains a bitmask telling which
00019 // CPUs can serve that interrupt.
00020 #define INT_DISABLED 0x00010000 // Interrupt disabled
00021 #define INT_LEVEL 0x00008000 // Level-triggered (vs edge-)
00022 #define INT_ACTIVELOW 0x00002000 // Active low (vs high)
00023 #define INT_LOGICAL 0x00000800 // Destination is CPU id (vs APIC ID)
00024
00025 volatile struct ioapic *ioapic;
00026
00027 // IO APIC MMIO structure: write reg, then read or write data.
00028 struct ioapic {
00029     uint reg;
00030     uint pad[3];
00031     uint data;
00032 };
00033
00034 static uint
00035 ioapicread(int reg)
00036 {
00037     ioapic->reg = reg;
00038     return ioapic->data;
00039 }
00040
00041 static void
00042 ioapicwrite(int reg, uint data)
00043 {
00044     ioapic->reg = reg;
00045     ioapic->data = data;
00046 }
00047
00048 void
00049 ioapicinit(void)
00050 {
00051     int i, id, maxintr;
00052
00053     ioapic = (volatile struct ioapic*)IOAPIC;
00054     maxintr = (ioapicread(REG_VER) >> 16) & 0xFF;
00055     id = ioapicread(REG_ID) >> 24;
00056     if(id != ioapicid)
00057         cprintf("ioapicinit: id isn't equal to ioapicid; not a MP\n");
00058
00059     // Mark all interrupts edge-triggered, active high, disabled,
00060     // and not routed to any CPUs.
00061     for(i = 0; i <= maxintr; i++){
00062         ioapicwrite(REG_TABLE+2*i, INT_DISABLED | (T_IRQ0 + i));
00063         ioapicwrite(REG_TABLE+2*i+1, 0);
00064     }
00065 }
00066
00067 void
00068 ioapicenable(int irq, int cpunum)
00069 {

```

```

00070 // Mark interrupt edge-triggered, active high,
00071 // enabled, and routed to the given cpunum,
00072 // which happens to be that cpu's APIC ID.
00073 ioapicwrite(REG_TABLE+2*irq, T_IRQ0 + irq);
00074 ioapicwrite(REG_TABLE+2*irq+1, cpunum « 24);
00075 }

```

5.81 ioapic.d File Reference

5.82 ioapic.d

[Go to the documentation of this file.](#)

```

00001 ioapic.o: ioapic.c /usr/include/stdc-predef.h types.h defs.h traps.h

```

5.83 kalloc.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "spinlock.h"

```

Classes

- struct [run](#)

Functions

- void [freerange](#) (void *vstart, void *vend)
- char * [kalloc](#) (void)
- void [kfree](#) (char *v)
- void [kinit1](#) (void *vstart, void *vend)
- void [kinit2](#) (void *vstart, void *vend)

Variables

- char [end](#) []
- struct {
 struct [run](#) * [freelist](#)
 struct [spinlock](#) [lock](#)
 int [use_lock](#)
 } [kmem](#)

5.83.1 Function Documentation

5.83.1.1 freerange()

```
void freerange (
    void * vstart,
    void * vend )
```

Definition at line 47 of file [kalloc.c](#).

```
00048 {
00049     char *p;
00050     p = (char*)PGROUNDUP((uint)vstart);
00051     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
00052         kfree(p);
00053 }
```

Referenced by [kinit1\(\)](#), and [kinit2\(\)](#).

5.83.1.2 kalloc()

```
char * kalloc (
    void )
```

Definition at line 83 of file [kalloc.c](#).

```
00084 {
00085     struct run *r;
00086
00087     if(kmem.use_lock)
00088         acquire(&kmem.lock);
00089     r = kmem.freelist;
00090     if(r)
00091         kmem.freelist = r->next;
00092     if(kmem.use_lock)
00093         release(&kmem.lock);
00094     return (char*)r;
00095 }
```

Referenced by [allocproc\(\)](#), [allocuvm\(\)](#), [copyuvm\(\)](#), [inituvm\(\)](#), [pipealloc\(\)](#), [setupkvm\(\)](#), [startothers\(\)](#), and [walkpgdir\(\)](#).

5.83.1.3 kfree()

```
void kfree (
    char * v )
```

Definition at line 60 of file [kalloc.c](#).

```
00061 {
00062     struct run *r;
00063
00064     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
00065         panic("kfree");
00066
00067     // Fill with junk to catch dangling refs.
00068     memset(v, 1, PGSIZE);
00069
00070     if(kmem.use_lock)
00071         acquire(&kmem.lock);
00072     r = (struct run*)v;
00073     r->next = kmem.freelist;
00074     kmem.freelist = r;
00075     if(kmem.use_lock)
00076         release(&kmem.lock);
00077 }
```

Referenced by [allocuvm\(\)](#), [copyuvm\(\)](#), [deallocuvm\(\)](#), [fork\(\)](#), [freerange\(\)](#), [freevm\(\)](#), [pipealloc\(\)](#), [pipeclose\(\)](#), and [wait\(\)](#).

5.83.1.4 kinit1()

```
void kinit1 (
    void * vstart,
    void * vend )
```

Definition at line 32 of file [kalloc.c](#).

```
00033 {
00034     initlock(&kmem.lock, "kmem");
00035     kmem.use_lock = 0;
00036     freerange(vstart, vend);
00037 }
```

5.83.1.5 kinit2()

```
void kinit2 (
    void * vstart,
    void * vend )
```

Definition at line 40 of file [kalloc.c](#).

```
00041 {
00042     freerange(vstart, vend);
00043     kmem.use_lock = 1;
00044 }
```

5.83.2 Variable Documentation

5.83.2.1 end

```
char end[] [extern]
```

Referenced by [kfree\(\)](#).

5.83.2.2 freelist

```
struct run* freelist
```

Definition at line 23 of file [kalloc.c](#).

5.83.2.3

```
struct { ... } kmem
```

Referenced by [kalloc\(\)](#), [kfree\(\)](#), [kinit1\(\)](#), and [kinit2\(\)](#).

5.83.2.4 lock

```
struct spinlock lock
```

Definition at line 21 of file [kalloc.c](#).

5.83.2.5 use_lock

```
int use_lock
```

Definition at line 22 of file [kalloc.c](#).

5.84 kalloc.c

[Go to the documentation of this file.](#)

```
00001 // Physical memory allocator, intended to allocate
00002 // memory for user processes, kernel stacks, page table pages,
00003 // and pipe buffers. Allocates 4096-byte pages.
00004
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "param.h"
00008 #include "memlayout.h"
00009 #include "mmu.h"
00010 #include "spinlock.h"
00011
00012 void freerange(void *vstart, void *vend);
00013 extern char end[]; // first address after kernel loaded from ELF file
00014 // defined by the kernel linker script in kernel.ld
00015
00016 struct run {
00017     struct run *next;
00018 };
00019
00020 struct {
00021     struct spinlock lock;
00022     int use_lock;
00023     struct run *freelist;
00024 } kmem;
00025
00026 // Initialization happens in two phases.
00027 // 1. main() calls kinit1() while still using entrypgdir to place just
00028 // the pages mapped by entrypgdir on free list.
00029 // 2. main() calls kinit2() with the rest of the physical pages
00030 // after installing a full page table that maps them on all cores.
00031 void
00032 kinit1(void *vstart, void *vend)
00033 {
00034     initlock(&kmem.lock, "kmem");
00035     kmem.use_lock = 0;
00036     freerange(vstart, vend);
00037 }
00038
00039 void
00040 kinit2(void *vstart, void *vend)
00041 {
00042     freerange(vstart, vend);
00043     kmem.use_lock = 1;
00044 }
00045
00046 void
00047 freerange(void *vstart, void *vend)
00048 {
00049     char *p;
00050     p = (char*)PGROUNDUP((uint)vstart);
00051     for(; p + PGSIZE <= (char*)vend; p += PGSIZE)
00052         kfree(p);
00053 }
00054 //PAGEBREAK: 21
00055 // Free the page of physical memory pointed at by v,
```

```

00056 // which normally should have been returned by a
00057 // call to kalloc(). (The exception is when
00058 // initializing the allocator; see kinit above.)
00059 void
00060 kfree(char *v)
00061 {
00062     struct run *r;
00063
00064     if((uint)v % PGSIZE || v < end || V2P(v) >= PHYSTOP)
00065         panic("kfree");
00066
00067     // Fill with junk to catch dangling refs.
00068     memset(v, 1, PGSIZE);
00069
00070     if(kmem.use_lock)
00071         acquire(&kmem.lock);
00072     r = (struct run*)v;
00073     r->next = kmem.freelist;
00074     kmem.freelist = r;
00075     if(kmem.use_lock)
00076         release(&kmem.lock);
00077 }
00078
00079 // Allocate one 4096-byte page of physical memory.
00080 // Returns a pointer that the kernel can use.
00081 // Returns 0 if the memory cannot be allocated.
00082 char*
00083 kalloc(void)
00084 {
00085     struct run *r;
00086
00087     if(kmem.use_lock)
00088         acquire(&kmem.lock);
00089     r = kmem.freelist;
00090     if(r)
00091         kmem.freelist = r->next;
00092     if(kmem.use_lock)
00093         release(&kmem.lock);
00094     return (char*)r;
00095 }
00096

```

5.85 kalloc.d File Reference

5.86 kalloc.d

[Go to the documentation of this file.](#)

```

00001 kalloc.o: kalloc.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 memlayout.h mmu.h spinlock.h

```

5.87 kbd.c File Reference

```

#include "types.h"
#include "x86.h"
#include "defs.h"
#include "kbd.h"

```

Functions

- int [kbdgetc](#) (void)
- void [kbdintr](#) (void)

5.87.1 Function Documentation

5.87.1.1 kbdgetc()

```
int kbdgetc (
    void )
```

Definition at line 7 of file [kbd.c](#).

```
00008 {
00009     static uint shift;
00010     static uchar *charcode[4] = {
00011         normalmap, shiftmap, ctlmap, ctlmap
00012     };
00013     uint st, data, c;
00014
00015     st = inb(KBSTATP);
00016     if((st & KBS_DIB) == 0)
00017         return -1;
00018     data = inb(KBDATAP);
00019
00020     if(data == 0xE0){
00021         shift |= EOESC;
00022         return 0;
00023     } else if(data & 0x80){
00024         // Key released
00025         data = (shift & EOESC ? data : data & 0x7F);
00026         shift &= ~(shiftcode[data] | EOESC);
00027         return 0;
00028     } else if(shift & EOESC){
00029         // Last character was an E0 escape; or with 0x80
00030         data |= 0x80;
00031         shift &= ~EOESC;
00032     }
00033
00034     shift |= shiftcode[data];
00035     shift ^= togglecode[data];
00036     c = charcode[shift & (CTL | SHIFT)][data];
00037     if(shift & CAPSLOCK){
00038         if('a' <= c && c <= 'z')
00039             c += 'A' - 'a';
00040         else if('A' <= c && c <= 'Z')
00041             c += 'a' - 'A';
00042     }
00043     return c;
00044 }
```

Referenced by [kbdintr\(\)](#).

5.87.1.2 kbdintr()

```
void kbdintr (
    void )
```

Definition at line 47 of file [kbd.c](#).

```
00048 {
00049     consoleintr(kbdgetc);
00050 }
```

Referenced by [trap\(\)](#).

5.88 kbd.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "x86.h"
00003 #include "defs.h"
00004 #include "kbd.h"
00005
00006 int
00007 kbdgetc(void)
00008 {
00009     static uint shift;
00010     static uchar *charcode[4] = {
00011         normalmap, shiftmap, ctlmap, ctlmap
00012     };
00013     uint st, data, c;
00014
00015     st = inb(KBSTATP);
00016     if((st & KBS_DIB) == 0)
00017         return -1;
00018     data = inb(KBDATAP);
00019
00020     if(data == 0xE0){
00021         shift |= E0ESC;
00022         return 0;
00023     } else if(data & 0x80){
00024         // Key released
00025         data = (shift & E0ESC ? data : data & 0x7F);
00026         shift &= ~(shiftcode[data] | E0ESC);
00027         return 0;
00028     } else if(shift & E0ESC){
00029         // Last character was an E0 escape; or with 0x80
00030         data |= 0x80;
00031         shift &= ~E0ESC;
00032     }
00033
00034     shift |= shiftcode[data];
00035     shift ^= togglecode[data];
00036     c = charcode[shift & (CTL | SHIFT)][data];
00037     if(shift & CAPSLOCK){
00038         if('a' <= c && c <= 'z')
00039             c += 'A' - 'a';
00040         else if('A' <= c && c <= 'Z')
00041             c += 'a' - 'A';
00042     }
00043     return c;
00044 }
00045
00046 void
00047 kbdintr(void)
00048 {
00049     consoleintr(kbdgetc);
00050 }

```

5.89 kbd.d File Reference

5.90 kbd.d

[Go to the documentation of this file.](#)

```
00001 kbd.o: kbd.c /usr/include/stdc-predef.h types.h x86.h defs.h kbd.h
```

5.91 kbd.h File Reference

Macros

- `#define ALT` (1<<2)
- `#define C(x)` (x - '@')

- `#define CAPSLOCK (1<<3)`
- `#define CTL (1<<1)`
- `#define E0ESC (1<<6)`
- `#define KBDATAP 0x60`
- `#define KBS_DIB 0x01`
- `#define KBSTAP 0x64`
- `#define KEY_DEL 0xE9`
- `#define KEY_DN 0xE3`
- `#define KEY_END 0xE1`
- `#define KEY_HOME 0xE0`
- `#define KEY_INS 0xE8`
- `#define KEY_LF 0xE4`
- `#define KEY_PGDN 0xE7`
- `#define KEY_PGUP 0xE6`
- `#define KEY_RT 0xE5`
- `#define KEY_UP 0xE2`
- `#define NO 0`
- `#define NUMLOCK (1<<4)`
- `#define SCROLLLOCK (1<<5)`
- `#define SHIFT (1<<0)`

Variables

- static `uchar ctlmap` [256]
- static `uchar normalmap` [256]
- static `uchar shiftcode` [256]
- static `uchar shiftmap` [256]
- static `uchar togglecode` [256]

5.91.1 Macro Definition Documentation

5.91.1.1 ALT

```
#define ALT (1<<2)
```

Definition at line 11 of file [kbd.h](#).

5.91.1.2 C

```
#define C(  
    x ) (x - '@')
```

Definition at line 32 of file [kbd.h](#).

5.91.1.3 CAPSLOCK

```
#define CAPSLOCK (1<<3)
```

Definition at line 13 of file [kbd.h](#).

5.91.1.4 CTL

```
#define CTL (1<<1)
```

Definition at line 10 of file [kbd.h](#).

5.91.1.5 E0ESC

```
#define E0ESC (1<<6)
```

Definition at line 17 of file [kbd.h](#).

5.91.1.6 KBDATAP

```
#define KBDATAP 0x60
```

Definition at line 5 of file [kbd.h](#).

5.91.1.7 KBS_DIB

```
#define KBS_DIB 0x01
```

Definition at line 4 of file [kbd.h](#).

5.91.1.8 KBSTATP

```
#define KBSTATP 0x64
```

Definition at line 3 of file [kbd.h](#).

5.91.1.9 KEY_DEL

```
#define KEY_DEL 0xE9
```

Definition at line 29 of file [kbd.h](#).

5.91.1.10 KEY_DN

```
#define KEY_DN 0xE3
```

Definition at line 23 of file [kbd.h](#).

5.91.1.11 KEY_END

```
#define KEY_END 0xE1
```

Definition at line 21 of file [kbd.h](#).

5.91.1.12 KEY_HOME

```
#define KEY_HOME 0xE0
```

Definition at line 20 of file [kbd.h](#).

5.91.1.13 KEY_INS

```
#define KEY_INS 0xE8
```

Definition at line 28 of file [kbd.h](#).

5.91.1.14 KEY_LF

```
#define KEY_LF 0xE4
```

Definition at line 24 of file [kbd.h](#).

5.91.1.15 KEY_PGDN

```
#define KEY_PGDN 0xE7
```

Definition at line 27 of file [kbd.h](#).

5.91.1.16 KEY_PGUP

```
#define KEY_PGUP 0xE6
```

Definition at line 26 of file [kbd.h](#).

5.91.1.17 KEY_RT

```
#define KEY_RT 0xE5
```

Definition at line 25 of file [kbd.h](#).

5.91.1.18 KEY_UP

```
#define KEY_UP 0xE2
```

Definition at line 22 of file [kbd.h](#).

5.91.1.19 NO

```
#define NO 0
```

Definition at line 7 of file [kbd.h](#).

5.91.1.20 NUMLOCK

```
#define NUMLOCK (1<<4)
```

Definition at line 14 of file [kbd.h](#).

5.91.1.21 SCROLLLOCK

```
#define SCROLLLOCK (1<<5)
```

Definition at line 15 of file [kbd.h](#).

5.91.1.22 SHIFT

```
#define SHIFT (1<<0)
```

Definition at line 9 of file [kbd.h](#).

5.91.2 Variable Documentation

5.91.2.1 ctlmap

```
uchar ctlmap[256] [static]
```

Initial value:

```
=
{
    NO,      NO,      NO,      NO,      NO,      NO,      NO,      NO,
    NO,      NO,      NO,      NO,      NO,      NO,      NO,      NO,
    C('Q'),  C('W'),  C('E'),  C('R'),  C('T'),  C('Y'),  C('U'),  C('I'),
    C('O'),  C('P'),  NO,      NO,      '\r',  NO,      C('A'),  C('S'),
    C('D'),  C('F'),  C('G'),  C('H'),  C('J'),  C('K'),  C('L'),  NO,
    NO,      NO,      NO,      C('\'),  C('Z'),  C('X'),  C('C'),  C('V'),
    C('B'),  C('N'),  C('M'),  NO,      NO,      C('/'),  NO,      NO,
    [0x9C] '\r',
    [0xB5] C('/'),
    [0xC8] KEY_UP,    [0xD0] KEY_DN,
    [0xC9] KEY_PGUP,  [0xD1] KEY_PGDN,
    [0xCB] KEY_LF,    [0xCD] KEY_RT,
    [0x97] KEY_HOME,  [0xCF] KEY_END,
    [0xD2] KEY_INS,   [0xD3] KEY_DEL
}
```

Definition at line 95 of file [kbd.h](#).

Referenced by [kbdgetc\(\)](#).

5.91.2.2 normalmap

```
uchar normalmap[256] [static]
```

Initial value:

```
=
{
    NO,    0x1B, '1', '2', '3', '4', '5', '6',
    '7', '8', '9', '0', '-', '=', '\b', '\t',
    'q', 'w', 'e', 'r', 't', 'y', 'u', 'i',
    'o', 'p', '[', ']', '\n', NO, 'a', 's',
    'd', 'f', 'g', 'h', 'j', 'k', 'l', ';',
    '\", ' ', NO, '\\', 'z', 'x', 'c', 'v',
    'b', 'n', 'm', ',', '.', '/', NO, '*',
    NO, ' ', NO, NO, NO, NO, NO, NO,
    NO, NO, NO, NO, NO, NO, NO, '7',
    '8', '9', '-', '4', '5', '6', '+', '1',
    '2', '3', '0', '.', NO, NO, NO, NO,
    [0x9C] '\n',
    [0xB5] '/',
    [0xC8] KEY_UP,    [0xD0] KEY_DN,
    [0xC9] KEY_PGUP,  [0xD1] KEY_PGDN,
    [0xCB] KEY_LF,    [0xCD] KEY_RT,
    [0x97] KEY_HOME,  [0xCF] KEY_END,
    [0xD2] KEY_INS,   [0xD3] KEY_DEL
}
```

Definition at line 51 of file [kbd.h](#).

Referenced by [kbdgetc\(\)](#).

5.91.2.3 shiftcode

```
uchar shiftcode[256] [static]
```

Initial value:

```
=
{
    [0x1D] CTL,
    [0x2A] SHIFT,
    [0x36] SHIFT,
    [0x38] ALT,
    [0x9D] CTL,
    [0xB8] ALT
}
```

Definition at line 34 of file [kbd.h](#).

Referenced by [kbdgetc\(\)](#).

5.91.2.4 shiftmap

```
uchar shiftmap[256] [static]
```

Initial value:

```
=
{
    NO,    033, '!', '@', '#', '$', '%', '^',
    '&', '*', '(', ')', '-', '+', '\b', '\t',
    'Q', 'W', 'E', 'R', 'T', 'Y', 'U', 'I',
    'O', 'P', '[', ']', '\n', NO, 'A', 'S',
    'D', 'F', 'G', 'H', 'J', 'K', 'L', ':',
    '"', '~', NO, '|', 'Z', 'X', 'C', 'V',
    'B', 'N', 'M', '<', '>', '?', NO, '*',

```

```

NO,      ' ', NO, NO, NO, NO, NO, NO,
NO,      NO, NO, NO, NO, NO, NO, '7',
'8',     '9', '-', '4', '5', '6', '+', '1',
'2',     '3', '0', '.', NO, NO, NO, NO,
[0x9C]   '\n',
[0xB5]   '/',
[0xC8]   KEY_UP,    [0xD0] KEY_DN,
[0xC9]   KEY_PGUP,  [0xD1] KEY_PGDN,
[0xCB]   KEY_LF,    [0xCD] KEY_RT,
[0x97]   KEY_HOME,  [0xCF] KEY_END,
[0xD2]   KEY_INS,   [0xD3] KEY_DEL
}

```

Definition at line 73 of file [kbd.h](#).

Referenced by [kbdgetc\(\)](#).

5.91.2.5 togglecode

```
uchar togglecode[256] [static]
```

Initial value:

```

=
{
    [0x3A] CAPSLOCK,
    [0x45] NUMLOCK,
    [0x46] SCROLLLOCK
}

```

Definition at line 44 of file [kbd.h](#).

Referenced by [kbdgetc\(\)](#).

5.92 kbd.h

[Go to the documentation of this file.](#)

```

00001 // PC keyboard interface constants
00002
00003 #define KBSTATP      0x64    // kbd controller status port(I)
00004 #define KBS_DIB      0x01    // kbd data in buffer
00005 #define KBDATAP      0x60    // kbd data port(I)
00006
00007 #define NO            0
00008
00009 #define SHIFT         (1<<0)
00010 #define CTL           (1<<1)
00011 #define ALT           (1<<2)
00012
00013 #define CAPSLOCK      (1<<3)
00014 #define NUMLOCK       (1<<4)
00015 #define SCROLLLOCK    (1<<5)
00016
00017 #define E0ESC         (1<<6)
00018
00019 // Special keycodes
00020 #define KEY_HOME      0xE0
00021 #define KEY_END       0xE1
00022 #define KEY_UP        0xE2
00023 #define KEY_DN        0xE3
00024 #define KEY_LF        0xE4
00025 #define KEY_RT        0xE5
00026 #define KEY_PGUP      0xE6
00027 #define KEY_PGDN      0xE7
00028 #define KEY_INS       0xE8
00029 #define KEY_DEL       0xE9
00030
00031 // C('A') == Control-A
00032 #define C(x) (x - '@')
00033

```

```

00034 static uchar shiftcode[256] =
00035 {
00036     [0x1D] CTL,
00037     [0x2A] SHIFT,
00038     [0x36] SHIFT,
00039     [0x38] ALT,
00040     [0x9D] CTL,
00041     [0xB8] ALT
00042 };
00043
00044 static uchar togglecode[256] =
00045 {
00046     [0x3A] CAPSLOCK,
00047     [0x45] NUMLOCK,
00048     [0x46] SCROLLLOCK
00049 };
00050
00051 static uchar normalmap[256] =
00052 {
00053     NO,    0x1B, '1', '2', '3', '4', '5', '6', // 0x00
00054     '7', '8', '9', '0', '-', '=', '\b', '\t',
00055     'q', 'w', 'e', 'r', 't', 'y', 'u', 'i', // 0x10
00056     'o', 'p', '[', ']', '\n', NO, 'a', 's',
00057     'd', 'f', 'g', 'h', 'j', 'k', 'l', ';', // 0x20
00058     '\'', '`', NO, '\\', 'z', 'x', 'c', 'v',
00059     'b', 'n', 'm', ' ', '.', '/', NO, '*', // 0x30
00060     NO, ' ', NO, NO, NO, NO, NO, NO,
00061     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
00062     '8', '9', '-', '4', '5', '6', '+', '1',
00063     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
00064     [0x9C] '\n', // KP_Enter
00065     [0xB5] '/', // KP_Div
00066     [0xC8] KEY_UP, [0xD0] KEY_DN,
00067     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
00068     [0xCB] KEY_LF, [0xCD] KEY_RT,
00069     [0x97] KEY_HOME, [0xCF] KEY_END,
00070     [0xD2] KEY_INS, [0xD3] KEY_DEL
00071 };
00072
00073 static uchar shiftmap[256] =
00074 {
00075     NO,    033, '!', '@', '#', '$', '%', '^', // 0x00
00076     '&', '*', '(', ')', '-', '=', '\b', '\t',
00077     'Q', 'W', 'E', 'R', 'T', 'Y', 'U', 'I', // 0x10
00078     'O', 'P', '[', ']', '\n', NO, 'A', 'S',
00079     'D', 'F', 'G', 'H', 'J', 'K', 'L', ';', // 0x20
00080     '"', '~', NO, '|', 'Z', 'X', 'C', 'V',
00081     'B', 'N', 'M', '<', '>', '?', NO, '*', // 0x30
00082     NO, ' ', NO, NO, NO, NO, NO, NO,
00083     NO, NO, NO, NO, NO, NO, NO, '7', // 0x40
00084     '8', '9', '-', '4', '5', '6', '+', '1',
00085     '2', '3', '0', '.', NO, NO, NO, NO, // 0x50
00086     [0x9C] '\n', // KP_Enter
00087     [0xB5] '/', // KP_Div
00088     [0xC8] KEY_UP, [0xD0] KEY_DN,
00089     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
00090     [0xCB] KEY_LF, [0xCD] KEY_RT,
00091     [0x97] KEY_HOME, [0xCF] KEY_END,
00092     [0xD2] KEY_INS, [0xD3] KEY_DEL
00093 };
00094
00095 static uchar ctlmap[256] =
00096 {
00097     NO,    NO,    NO,    NO,    NO,    NO,    NO,    NO,
00098     NO,    NO,    NO,    NO,    NO,    NO,    NO,    NO,
00099     C('Q'), C('W'), C('E'), C('R'), C('T'), C('Y'), C('U'), C('I'),
00100     C('O'), C('P'), NO,    NO,    '\r', NO,    C('A'), C('S'),
00101     C('D'), C('F'), C('G'), C('H'), C('J'), C('K'), C('L'), NO,
00102     NO,    NO,    NO,    C('\n'), C('Z'), C('X'), C('C'), C('V'),
00103     C('B'), C('N'), C('M'), NO,    NO,    C('/'), NO,    NO,
00104     [0x9C] '\r', // KP_Enter
00105     [0xB5] C('/'), // KP_Div
00106     [0xC8] KEY_UP, [0xD0] KEY_DN,
00107     [0xC9] KEY_PGUP, [0xD1] KEY_PGDN,
00108     [0xCB] KEY_LF, [0xCD] KEY_RT,
00109     [0x97] KEY_HOME, [0xCF] KEY_END,
00110     [0xD2] KEY_INS, [0xD3] KEY_DEL
00111 };
00112

```

5.93 kill.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- `int main (int argc, char **argv)`

5.93.1 Function Documentation

5.93.1.1 main()

```
int main (
    int argc,
    char ** argv )
```

Definition at line 6 of file [kill.c](#).

```
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "usage: kill pid...\n");
00012         exit();
00013     }
00014     for(i=1; i<argc; i++)
00015         kill(atoi(argv[i]));
00016     exit();
00017 }
```

5.94 kill.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 int
00006 main(int argc, char **argv)
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "usage: kill pid...\n");
00012         exit();
00013     }
00014     for(i=1; i<argc; i++)
00015         kill(atoi(argv[i]));
00016     exit();
00017 }
```

5.95 kill.d File Reference

5.96 kill.d

[Go to the documentation of this file.](#)

```
00001 kill.o: kill.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.97 lapic.c File Reference

```
#include "param.h"
#include "types.h"
#include "defs.h"
#include "date.h"
#include "memlayout.h"
#include "traps.h"
#include "mmu.h"
#include "x86.h"
```

Macros

- #define [ASSERT](#) 0x00004000
- #define [BCAST](#) 0x00080000
- #define [BUSY](#) 0x00001000
- #define [CMOS_PORT](#) 0x70
- #define [CMOS_RETURN](#) 0x71
- #define [CMOS_STATA](#) 0x0a
- #define [CMOS_STATB](#) 0x0b
- #define [CMOS_UIP](#) (1 << 7)
- #define [CONV](#)(x) (t1.x = ((t1.x >> 4) * 10) + (t1.x & 0xf))
- #define [DAY](#) 0x07
- #define [DEASSERT](#) 0x00000000
- #define [DELIVS](#) 0x00001000
- #define [ENABLE](#) 0x00000100
- #define [EOI](#) (0x00B0/4)
- #define [ERROR](#) (0x0370/4)
- #define [ESR](#) (0x0280/4)
- #define [FIXED](#) 0x00000000
- #define [HOURS](#) 0x04
- #define [ICRHI](#) (0x0310/4)
- #define [ICRLO](#) (0x0300/4)
- #define [ID](#) (0x0020/4)
- #define [INIT](#) 0x00000500
- #define [LEVEL](#) 0x00008000
- #define [LINT0](#) (0x0350/4)
- #define [LINT1](#) (0x0360/4)
- #define [MASKED](#) 0x00010000
- #define [MINS](#) 0x02
- #define [MONTH](#) 0x08
- #define [PCINT](#) (0x0340/4)

- `#define PERIODIC 0x00020000`
- `#define SECS 0x00`
- `#define STARTUP 0x00000600`
- `#define SVR (0x00F0/4)`
- `#define TCCR (0x0390/4)`
- `#define TDCR (0x03E0/4)`
- `#define TICR (0x0380/4)`
- `#define TIMER (0x0320/4)`
- `#define TPR (0x0080/4)`
- `#define VER (0x0030/4)`
- `#define X1 0x0000000B`
- `#define YEAR 0x09`

Functions

- static `uint cmos_read (uint reg)`
- void `cmostime (struct rtcddate *r)`
- static void `fill_rtcddate (struct rtcddate *r)`
- void `lapiceoi (void)`
- int `lapicid (void)`
- void `lapicinit (void)`
- void `lapicstartap (uchar apicid, uint addr)`
- static void `lapicw (int index, int value)`
- void `microdelay (int us)`

Variables

- volatile `uint * lapic`

5.97.1 Macro Definition Documentation

5.97.1.1 ASSERT

```
#define ASSERT 0x00004000
```

Definition at line 25 of file [lapic.c](#).

5.97.1.2 BCAST

```
#define BCAST 0x00080000
```

Definition at line 28 of file [lapic.c](#).

5.97.1.3 BUSY

```
#define BUSY 0x00001000
```

Definition at line 29 of file [lapic.c](#).

5.97.1.4 CMOS_PORT

```
#define CMOS_PORT 0x70
```

Definition at line 123 of file [lapic.c](#).

5.97.1.5 CMOS_RETURN

```
#define CMOS_RETURN 0x71
```

Definition at line 124 of file [lapic.c](#).

5.97.1.6 CMOS_STATA

```
#define CMOS_STATA 0x0a
```

Definition at line 163 of file [lapic.c](#).

5.97.1.7 CMOS_STATB

```
#define CMOS_STATB 0x0b
```

Definition at line 164 of file [lapic.c](#).

5.97.1.8 CMOS_UIP

```
#define CMOS_UIP (1 << 7)
```

Definition at line 165 of file [lapic.c](#).

5.97.1.9 CONV

```
#define CONV(  
    x ) (t1.x = ((t1.x >> 4) * 10) + (t1.x & 0xf))
```

5.97.1.10 DAY

```
#define DAY 0x07
```

Definition at line 170 of file [lapic.c](#).

5.97.1.11 DEASSERT

```
#define DEASSERT 0x00000000
```

Definition at line 26 of file [lapic.c](#).

5.97.1.12 DELIVS

```
#define DELIVS 0x00001000
```

Definition at line 24 of file [lapic.c](#).

5.97.1.13 ENABLE

```
#define ENABLE 0x00000100
```

Definition at line 19 of file [lapic.c](#).

5.97.1.14 EOI

```
#define EOI (0x00B0/4)
```

Definition at line 17 of file [lapic.c](#).

5.97.1.15 ERROR

```
#define ERROR (0x0370/4)
```

Definition at line 38 of file [lapic.c](#).

5.97.1.16 ESR

```
#define ESR (0x0280/4)
```

Definition at line 20 of file [lapic.c](#).

5.97.1.17 FIXED

```
#define FIXED 0x00000000
```

Definition at line 30 of file [lapic.c](#).

5.97.1.18 HOURS

```
#define HOURS 0x04
```

Definition at line 169 of file [lapic.c](#).

5.97.1.19 ICRHI

```
#define ICRHI (0x0310/4)
```

Definition at line 31 of file [lapic.c](#).

5.97.1.20 ICRLO

```
#define ICRLO (0x0300/4)
```

Definition at line 21 of file [lapic.c](#).

5.97.1.21 ID

```
#define ID (0x0020/4)
```

Definition at line 14 of file [lapic.c](#).

5.97.1.22 INIT

```
#define INIT 0x00000500
```

Definition at line 22 of file [lapic.c](#).

5.97.1.23 LEVEL

```
#define LEVEL 0x00008000
```

Definition at line 27 of file [lapic.c](#).

5.97.1.24 LINT0

```
#define LINT0 (0x0350/4)
```

Definition at line 36 of file [lapic.c](#).

5.97.1.25 LINT1

```
#define LINT1 (0x0360/4)
```

Definition at line 37 of file [lapic.c](#).

5.97.1.26 MASKED

```
#define MASKED 0x00010000
```

Definition at line 39 of file [lapic.c](#).

5.97.1.27 MINS

```
#define MINS 0x02
```

Definition at line 168 of file [lapic.c](#).

5.97.1.28 MONTH

```
#define MONTH 0x08
```

Definition at line 171 of file [lapic.c](#).

5.97.1.29 PCINT

```
#define PCINT (0x0340/4)
```

Definition at line 35 of file [lapic.c](#).

5.97.1.30 PERIODIC

```
#define PERIODIC 0x00020000
```

Definition at line 34 of file [lapic.c](#).

5.97.1.31 SECS

```
#define SECS 0x00
```

Definition at line 167 of file [lapic.c](#).

5.97.1.32 STARTUP

```
#define STARTUP 0x00000600
```

Definition at line 23 of file [lapic.c](#).

5.97.1.33 SVR

```
#define SVR (0x00F0/4)
```

Definition at line 18 of file [lapic.c](#).

5.97.1.34 TCCR

```
#define TCCR (0x0390/4)
```

Definition at line 41 of file [lapic.c](#).

5.97.1.35 TDCR

```
#define TDCR (0x03E0/4)
```

Definition at line 42 of file [lapic.c](#).

5.97.1.36 TICR

```
#define TICR (0x0380/4)
```

Definition at line 40 of file [lapic.c](#).

5.97.1.37 TIMER

```
#define TIMER (0x0320/4)
```

Definition at line 32 of file [lapic.c](#).

5.97.1.38 TPR

```
#define TPR (0x0080/4)
```

Definition at line 16 of file [lapic.c](#).

5.97.1.39 VER

```
#define VER (0x0030/4)
```

Definition at line 15 of file [lapic.c](#).

5.97.1.40 X1

```
#define X1 0x0000000B
```

Definition at line 33 of file [lapic.c](#).

5.97.1.41 YEAR

```
#define YEAR 0x09
```

Definition at line 172 of file [lapic.c](#).

5.97.2 Function Documentation

5.97.2.1 cmos_read()

```
static uint cmos_read (  
    uint reg ) [static]
```

Definition at line 175 of file [lapic.c](#).

```
00176 {  
00177     outb(CMOS_PORT, reg);  
00178     microdelay(200);  
00179  
00180     return inb(CMOS_RETURN);  
00181 }
```

Referenced by [cmostime\(\)](#), and [fill_rtcddate\(\)](#).

5.97.2.2 cmostime()

```
void cmostime (
    struct rtcdate * r )
```

Definition at line 196 of file [lapic.c](#).

```
00197 {
00198     struct rtcdate t1, t2;
00199     int sb, bcd;
00200
00201     sb = cmos_read(CMOS_STATB);
00202
00203     bcd = (sb & (1 << 2)) == 0;
00204
00205     // make sure CMOS doesn't modify time while we read it
00206     for(;;) {
00207         fill_rtcdate(&t1);
00208         if(cmos_read(CMOS_STATB) & CMOS_UIP)
00209             continue;
00210         fill_rtcdate(&t2);
00211         if(memcmp(&t1, &t2, sizeof(t1)) == 0)
00212             break;
00213     }
00214
00215     // convert
00216     if(bcd) {
00217 #define CONV(x)      (t1.x = ((t1.x >> 4) * 10) + (t1.x & 0xf))
00218         CONV(second);
00219         CONV(minute);
00220         CONV(hour);
00221         CONV(day);
00222         CONV(month);
00223         CONV(year);
00224 #undef CONV
00225     }
00226
00227     *r = t1;
00228     r->year += 2000;
00229 }
```

5.97.2.3 fill_rtcdate()

```
static void fill_rtcdate (
    struct rtcdate * r ) [static]
```

Definition at line 184 of file [lapic.c](#).

```
00185 {
00186     r->second = cmos_read(SECS);
00187     r->minute = cmos_read(MINS);
00188     r->hour = cmos_read(HOURS);
00189     r->day = cmos_read(DAY);
00190     r->month = cmos_read(MONTH);
00191     r->year = cmos_read(YEAR);
00192 }
```

Referenced by [cmostime\(\)](#).

5.97.2.4 lapiceoi()

```
void lapiceoi (
    void )
```

Definition at line 110 of file [lapic.c](#).

```
00111 {
00112     if(lapic)
00113         lapicw(EOI, 0);
00114 }
```

Referenced by [trap\(\)](#).

5.97.2.5 lapicid()

```
int lapicid (
    void )
```

Definition at line 101 of file [lapic.c](#).

```
00102 {
00103     if (!lapic)
00104         return 0;
00105     return lapic[ID] » 24;
00106 }
```

Referenced by [mycpu\(\)](#), and [panic\(\)](#).

5.97.2.6 lapicinit()

```
void lapicinit (
    void )
```

Definition at line 55 of file [lapic.c](#).

```
00056 {
00057     if (!lapic)
00058         return;
00059
00060     // Enable local APIC; set spurious interrupt vector.
00061     lapicw(SVR, ENABLE | (T_IRQ0 + IRQ_SPURIOUS));
00062
00063     // The timer repeatedly counts down at bus frequency
00064     // from lapic[TICR] and then issues an interrupt.
00065     // If xv6 cared more about precise timekeeping,
00066     // TICR would be calibrated using an external time source.
00067     lapicw(TDCR, X1);
00068     lapicw(TIMER, PERIODIC | (T_IRQ0 + IRQ_TIMER));
00069     lapicw(TICR, 10000000);
00070
00071     // Disable logical interrupt lines.
00072     lapicw(LINT0, MASKED);
00073     lapicw(LINT1, MASKED);
00074
00075     // Disable performance counter overflow interrupts
00076     // on machines that provide that interrupt entry.
00077     if (((lapic[VER] » 16) & 0xFF) >= 4)
00078         lapicw(PCINT, MASKED);
00079
00080     // Map error interrupt to IRQ_ERROR.
00081     lapicw(ERROR, T_IRQ0 + IRQ_ERROR);
00082
00083     // Clear error status register (requires back-to-back writes).
00084     lapicw(ESR, 0);
00085     lapicw(ESR, 0);
00086
00087     // Ack any outstanding interrupts.
00088     lapicw(EOI, 0);
00089
00090     // Send an Init Level De-Assert to synchronise arbitration ID's.
00091     lapicw(ICRHI, 0);
00092     lapicw(ICRLO, BCAST | INIT | LEVEL);
00093     while(lapic[ICRLO] & DELIVS)
00094         ;
00095
00096     // Enable interrupts on the APIC (but not on the processor).
00097     lapicw(TPR, 0);
00098 }
```

Referenced by [mpenter\(\)](#).

5.97.2.7 lapicstartap()

```
void lapicstartap (
    uchar apicid,
    uint addr )
```

Definition at line 129 of file [lapic.c](#).

```
00130 {
00131     int i;
00132     ushort *wrv;
00133
00134     // "The BSP must initialize CMOS shutdown code to 0AH
00135     // and the warm reset vector (DWORD based at 40:67) to point at
00136     // the AP startup code prior to the [universal startup algorithm]."
00137     outb(CMOS_PORT, 0xF); // offset 0xF is shutdown code
00138     outb(CMOS_PORT+1, 0x0A);
00139     wrv = (ushort*)P2V((0x40«4 | 0x67)); // Warm reset vector
00140     wrv[0] = 0;
00141     wrv[1] = addr » 4;
00142
00143     // "Universal startup algorithm."
00144     // Send INIT (level-triggered) interrupt to reset other CPU.
00145     lapicw(ICRHI, apicid«24);
00146     lapicw(ICRLO, INIT | LEVEL | ASSERT);
00147     microdelay(200);
00148     lapicw(ICRLO, INIT | LEVEL);
00149     microdelay(100); // should be 10ms, but too slow in Bochs!
00150
00151     // Send startup IPI (twice!) to enter code.
00152     // Regular hardware is supposed to only accept a STARTUP
00153     // when it is in the halted state due to an INIT. So the second
00154     // should be ignored, but it is part of the official Intel algorithm.
00155     // Bochs complains about the second one. Too bad for Bochs.
00156     for(i = 0; i < 2; i++){
00157         lapicw(ICRHI, apicid«24);
00158         lapicw(ICRLO, STARTUP | (addr»12));
00159         microdelay(200);
00160     }
00161 }
```

Referenced by [startothers\(\)](#).

5.97.2.8 lapicw()

```
static void lapicw (
    int index,
    int value ) [static]
```

Definition at line 48 of file [lapic.c](#).

```
00049 {
00050     lapic[index] = value;
00051     lapic[ID]; // wait for write to finish, by reading
00052 }
```

Referenced by [lapiceoi\(\)](#), [lapicinit\(\)](#), and [lapicstartap\(\)](#).

5.97.2.9 microdelay()

```
void microdelay (
    int us )
```

Definition at line 119 of file [lapic.c](#).

```
00120 {
00121 }
```

Referenced by [cmos_read\(\)](#), [lapicstartap\(\)](#), and [uartputc\(\)](#).

5.97.3 Variable Documentation

5.97.3.1 lapic

volatile `uint*` lapic

Definition at line 44 of file `lapic.c`.

Referenced by `lapiceoi()`, `lapicid()`, `lapicinit()`, `lapicw()`, and `mpinit()`.

5.98 lapic.c

[Go to the documentation of this file.](#)

```
00001 // The local APIC manages internal (non-I/O) interrupts.
00002 // See Chapter 8 & Appendix C of Intel processor manual volume 3.
00003
00004 #include "param.h"
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "date.h"
00008 #include "memlayout.h"
00009 #include "traps.h"
00010 #include "mmu.h"
00011 #include "x86.h"
00012
00013 // Local APIC registers, divided by 4 for use as uint[] indices.
00014 #define ID (0x0020/4) // ID
00015 #define VER (0x0030/4) // Version
00016 #define TPR (0x0080/4) // Task Priority
00017 #define EOI (0x00B0/4) // EOI
00018 #define SVR (0x00F0/4) // Spurious Interrupt Vector
00019 #define ENABLE 0x00000100 // Unit Enable
00020 #define ESR (0x0280/4) // Error Status
00021 #define ICRLO (0x0300/4) // Interrupt Command
00022 #define INIT 0x00000500 // INIT/RESET
00023 #define STARTUP 0x00000600 // Startup IPI
00024 #define DELIVS 0x00001000 // Delivery status
00025 #define ASSERT 0x00004000 // Assert interrupt (vs deassert)
00026 #define DEASSERT 0x00000000
00027 #define LEVEL 0x00008000 // Level triggered
00028 #define BCAST 0x00008000 // Send to all APICs, including self.
00029 #define BUSY 0x00001000
00030 #define FIXED 0x00000000
00031 #define ICRHI (0x0310/4) // Interrupt Command [63:32]
00032 #define TIMER (0x0320/4) // Local Vector Table 0 (TIMER)
00033 #define X1 0x0000000B // divide counts by 1
00034 #define PERIODIC 0x00020000 // Periodic
00035 #define PCINT (0x0340/4) // Performance Counter LVT
00036 #define LINT0 (0x0350/4) // Local Vector Table 1 (LINT0)
00037 #define LINT1 (0x0360/4) // Local Vector Table 2 (LINT1)
00038 #define ERROR (0x0370/4) // Local Vector Table 3 (ERROR)
00039 #define MASKED 0x00010000 // Interrupt masked
00040 #define TICR (0x0380/4) // Timer Initial Count
00041 #define TCCR (0x0390/4) // Timer Current Count
00042 #define TDCR (0x03E0/4) // Timer Divide Configuration
00043
00044 volatile uint *lapic; // Initialized in mp.c
00045
00046 //PAGEBREAK!
00047 static void
00048 lapicw(int index, int value)
00049 {
00050     lapic[index] = value;
00051     lapic[ID]; // wait for write to finish, by reading
00052 }
00053
00054 void
00055 lapicinit(void)
00056 {
00057     if(!lapic)
00058         return;
```

```

00059
00060 // Enable local APIC; set spurious interrupt vector.
00061 lapicw(SVR, ENABLE | (T_IRQ0 + IRQ_SPURIOUS));
00062
00063 // The timer repeatedly counts down at bus frequency
00064 // from lapic[TICR] and then issues an interrupt.
00065 // If xv6 cared more about precise timekeeping,
00066 // TICR would be calibrated using an external time source.
00067 lapicw(TDCR, X1);
00068 lapicw(TIMER, PERIODIC | (T_IRQ0 + IRQ_TIMER));
00069 lapicw(TICR, 10000000);
00070
00071 // Disable logical interrupt lines.
00072 lapicw(LINT0, MASKED);
00073 lapicw(LINT1, MASKED);
00074
00075 // Disable performance counter overflow interrupts
00076 // on machines that provide that interrupt entry.
00077 if(((lapic[VER]»16) & 0xFF) >= 4)
00078     lapicw(PCINT, MASKED);
00079
00080 // Map error interrupt to IRQ_ERROR.
00081 lapicw(ERROR, T_IRQ0 + IRQ_ERROR);
00082
00083 // Clear error status register (requires back-to-back writes).
00084 lapicw(ESR, 0);
00085 lapicw(ESR, 0);
00086
00087 // Ack any outstanding interrupts.
00088 lapicw(EOI, 0);
00089
00090 // Send an Init Level De-Assert to synchronise arbitration ID's.
00091 lapicw(ICRHI, 0);
00092 lapicw(ICRLO, BCAST | INIT | LEVEL);
00093 while(lapic[ICRLO] & DELIVS)
00094     ;
00095
00096 // Enable interrupts on the APIC (but not on the processor).
00097 lapicw(TPR, 0);
00098 }
00099
00100 int
00101 lapicid(void)
00102 {
00103     if (!lapic)
00104         return 0;
00105     return lapic[ID] » 24;
00106 }
00107
00108 // Acknowledge interrupt.
00109 void
00110 lapiceoi(void)
00111 {
00112     if(lapic)
00113         lapicw(EOI, 0);
00114 }
00115
00116 // Spin for a given number of microseconds.
00117 // On real hardware would want to tune this dynamically.
00118 void
00119 microdelay(int us)
00120 {
00121 }
00122
00123 #define CMOS_PORT    0x70
00124 #define CMOS_RETURN  0x71
00125
00126 // Start additional processor running entry code at addr.
00127 // See Appendix B of MultiProcessor Specification.
00128 void
00129 lapicstartap(uchar apicid, uint addr)
00130 {
00131     int i;
00132     ushort *wrv;
00133
00134     // "The BSP must initialize CMOS shutdown code to 0AH
00135     // and the warm reset vector (DWORD based at 40:67) to point at
00136     // the AP startup code prior to the [universal startup algorithm]."
00137     outb(CMOS_PORT, 0xF); // offset 0xF is shutdown code
00138     outb(CMOS_PORT+1, 0x0A);
00139     wrv = (ushort*)P2V((0x40«4 | 0x67)); // Warm reset vector
00140     wrv[0] = 0;
00141     wrv[1] = addr » 4;
00142
00143     // "Universal startup algorithm."
00144     // Send INIT (level-triggered) interrupt to reset other CPU.
00145     lapicw(ICRHI, apicid«24);

```

```

00146     lapicw(ICRLO, INIT | LEVEL | ASSERT);
00147     microdelay(200);
00148     lapicw(ICRLO, INIT | LEVEL);
00149     microdelay(100);    // should be 10ms, but too slow in Bochs!
00150
00151     // Send startup IPI (twice!) to enter code.
00152     // Regular hardware is supposed to only accept a STARTUP
00153     // when it is in the halted state due to an INIT.  So the second
00154     // should be ignored, but it is part of the official Intel algorithm.
00155     // Bochs complains about the second one.  Too bad for Bochs.
00156     for(i = 0; i < 2; i++){
00157         lapicw(ICRHI, apicid«24);
00158         lapicw(ICRLO, STARTUP | (addr»12));
00159         microdelay(200);
00160     }
00161 }
00162
00163 #define CMOS_STATA    0x0a
00164 #define CMOS_STATB    0x0b
00165 #define CMOS_UIP      (1 « 7)    // RTC update in progress
00166
00167 #define SECS          0x00
00168 #define MINS          0x02
00169 #define HOURS         0x04
00170 #define DAY           0x07
00171 #define MONTH         0x08
00172 #define YEAR          0x09
00173
00174 static uint
00175 cmos_read(uint reg)
00176 {
00177     outb(CMOS_PORT, reg);
00178     microdelay(200);
00179
00180     return inb(CMOS_RETURN);
00181 }
00182
00183 static void
00184 fill_rtcddate(struct rtcdate *r)
00185 {
00186     r->second = cmos_read(SECS);
00187     r->minute = cmos_read(MINS);
00188     r->hour   = cmos_read(HOURS);
00189     r->day     = cmos_read(DAY);
00190     r->month   = cmos_read(MONTH);
00191     r->year    = cmos_read(YEAR);
00192 }
00193
00194 // qemu seems to use 24-hour GWT and the values are BCD encoded
00195 void
00196 cmostime(struct rtcdate *r)
00197 {
00198     struct rtcdate t1, t2;
00199     int sb, bcd;
00200
00201     sb = cmos_read(CMOS_STATB);
00202
00203     bcd = (sb & (1 « 2)) == 0;
00204
00205     // make sure CMOS doesn't modify time while we read it
00206     for(;;) {
00207         fill_rtcddate(&t1);
00208         if(cmos_read(CMOS_STATA) & CMOS_UIP)
00209             continue;
00210         fill_rtcddate(&t2);
00211         if(memcmp(&t1, &t2, sizeof(t1)) == 0)
00212             break;
00213     }
00214
00215     // convert
00216     if(bcd) {
00217 #define CONV(x)        (t1.x = ((t1.x » 4) * 10) + (t1.x & 0xf))
00218         CONV(second);
00219         CONV(minute);
00220         CONV(hour);
00221         CONV(day);
00222         CONV(month);
00223         CONV(year);
00224 #undef CONV
00225     }
00226
00227     *r = t1;
00228     r->year += 2000;
00229 }

```

5.99 lapic.d File Reference

5.100 lapic.d

[Go to the documentation of this file.](#)

```
00001 lapic.o: lapic.c /usr/include/stdc-predef.h param.h types.h defs.h date.h \  
00002 memlayout.h traps.h mmu.h x86.h
```

5.101 In.c File Reference

```
#include "types.h"  
#include "stat.h"  
#include "user.h"
```

Functions

- `int main (int argc, char *argv[])`

5.101.1 Function Documentation

5.101.1.1 main()

```
int main (  
    int argc,  
    char * argv[] )
```

Definition at line 6 of file [ln.c](#).

```
00007 {  
00008     if(argc != 3){  
00009         printf(2, "Usage: ln old new\n");  
00010         exit();  
00011     }  
00012     if(link(argv[1], argv[2]) < 0)  
00013         printf(2, "link %s %s: failed\n", argv[1], argv[2]);  
00014     exit();  
00015 }
```

5.102 In.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"  
00002 #include "stat.h"  
00003 #include "user.h"  
00004  
00005 int  
00006 main(int argc, char *argv[])  
00007 {  
00008     if(argc != 3){  
00009         printf(2, "Usage: ln old new\n");  
00010         exit();  
00011     }  
00012     if(link(argv[1], argv[2]) < 0)  
00013         printf(2, "link %s %s: failed\n", argv[1], argv[2]);  
00014     exit();  
00015 }
```

5.103 In.d File Reference

5.104 In.d

[Go to the documentation of this file.](#)

```
00001 ln.o: ln.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.105 log.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "buf.h"
```

Classes

- struct [log](#)
- struct [logheader](#)

Functions

- void [begin_op](#) (void)
- static void [commit](#) ()
- void [end_op](#) (void)
- void [initlog](#) (int dev)
- static void [install_trans](#) (void)
- void [log_write](#) (struct [buf](#) *b)
- static void [read_head](#) (void)
- static void [recover_from_log](#) (void)
- static void [write_head](#) (void)
- static void [write_log](#) (void)

Variables

- struct [log](#) [log](#)

5.105.1 Function Documentation

5.105.1.1 begin_op()

```
void begin_op (
    void )
```

Definition at line 126 of file [log.c](#).

```
00127 {
00128     acquire(&log.lock);
00129     while(1){
00130         if(log.committing){
00131             sleep(&log, &log.lock);
00132         } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){
00133             // this op might exhaust log space; wait for commit.
00134             sleep(&log, &log.lock);
00135         } else {
00136             log.outstanding += 1;
00137             release(&log.lock);
00138             break;
00139         }
00140     }
00141 }
```

Referenced by [exec\(\)](#), [exit\(\)](#), [fclose\(\)](#), [fwrite\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.105.1.2 commit()

```
static void commit ( ) [static]
```

Definition at line 193 of file [log.c](#).

```
00194 {
00195     if (log.lh.n > 0) {
00196         write_log(); // Write modified blocks from cache to log
00197         write_head(); // Write header to disk -- the real commit
00198         install_trans(); // Now install writes to home locations
00199         log.lh.n = 0;
00200         write_head(); // Erase the transaction from the log
00201     }
00202 }
```

Referenced by [end_op\(\)](#).

5.105.1.3 end_op()

```
void end_op (
    void )
```

Definition at line 146 of file [log.c](#).

```
00147 {
00148     int do_commit = 0;
00149
00150     acquire(&log.lock);
00151     log.outstanding -= 1;
00152     if(log.committing)
00153         panic("log.committing");
00154     if(log.outstanding == 0){
00155         do_commit = 1;
00156         log.committing = 1;
00157     } else {
00158         // begin_op() may be waiting for log space,
00159         // and decrementing log.outstanding has decreased
00160         // the amount of reserved space.
00161         wakeup(&log);
00162     }
```

```

00163     release(&log.lock);
00164
00165     if(do_commit){
00166         // call commit w/o holding locks, since not allowed
00167         // to sleep with locks.
00168         commit();
00169         acquire(&log.lock);
00170         log.committing = 0;
00171         wakeup(&log);
00172         release(&log.lock);
00173     }
00174 }

```

Referenced by [exec\(\)](#), [exit\(\)](#), [fileclose\(\)](#), [filewrite\(\)](#), [sys_chdir\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.105.1.4 initlog()

```

void initlog (
    int dev )

```

Definition at line 54 of file [log.c](#).

```

00055 {
00056     if (sizeof(struct logheader) >= BSIZE)
00057         panic("initlog: too big logheader");
00058
00059     struct superblock sb;
00060     initlock(&log.lock, "log");
00061     readsb(dev, &sb);
00062     log.start = sb.logstart;
00063     log.size = sb.nlog;
00064     log.dev = dev;
00065     recover_from_log();
00066 }

```

Referenced by [forkret\(\)](#).

5.105.1.5 install_trans()

```

static void install_trans (
    void ) [static]

```

Definition at line 70 of file [log.c](#).

```

00071 {
00072     int tail;
00073
00074     for (tail = 0; tail < log.lh.n; tail++) {
00075         struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block
00076         struct buf *dbuf = bread(log.dev, log.lh.block[tail]); // read dst
00077         memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst
00078         bwrite(dbuf); // write dst to disk
00079         brelse(lbuf);
00080         brelse(dbuf);
00081     }
00082 }

```

Referenced by [commit\(\)](#), and [recover_from_log\(\)](#).

5.105.1.6 log_write()

```
void log_write (
    struct buf * b )
```

Definition at line 214 of file [log.c](#).

```
00215 {
00216     int i;
00217
00218     if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
00219         panic("too big a transaction");
00220     if (log.outstanding < 1)
00221         panic("log_write outside of trans");
00222
00223     acquire(&log.lock);
00224     for (i = 0; i < log.lh.n; i++) {
00225         if (log.lh.block[i] == b->blockno)    // log absorbtion
00226             break;
00227     }
00228     log.lh.block[i] = b->blockno;
00229     if (i == log.lh.n)
00230         log.lh.n++;
00231     b->flags |= B_DIRTY; // prevent eviction
00232     release(&log.lock);
00233 }
```

Referenced by [balloc\(\)](#), [bfree\(\)](#), [bmap\(\)](#), [bzero\(\)](#), [ialloc\(\)](#), [iupdate\(\)](#), and [writei\(\)](#).

5.105.1.7 read_head()

```
static void read_head (
    void ) [static]
```

Definition at line 86 of file [log.c](#).

```
00087 {
00088     struct buf *buf = bread(log.dev, log.start);
00089     struct logheader *lh = (struct logheader *) (buf->data);
00090     int i;
00091     log.lh.n = lh->n;
00092     for (i = 0; i < log.lh.n; i++) {
00093         log.lh.block[i] = lh->block[i];
00094     }
00095     brelse(buf);
00096 }
```

Referenced by [recover_from_log\(\)](#).

5.105.1.8 recover_from_log()

```
static void recover_from_log (
    void ) [static]
```

Definition at line 116 of file [log.c](#).

```
00117 {
00118     read_head();
00119     install_trans(); // if committed, copy from log to disk
00120     log.lh.n = 0;
00121     write_head(); // clear the log
00122 }
```

Referenced by [initlog\(\)](#).

5.105.1.9 write_head()

```
static void write_head (
    void ) [static]
```

Definition at line 102 of file [log.c](#).

```
00103 {
00104     struct buf *buf = bread(log.dev, log.start);
00105     struct logheader *hb = (struct logheader *) (buf->data);
00106     int i;
00107     hb->n = log.lh.n;
00108     for (i = 0; i < log.lh.n; i++) {
00109         hb->block[i] = log.lh.block[i];
00110     }
00111     bwrite(buf);
00112     brelse(buf);
00113 }
```

Referenced by [commit\(\)](#), and [recover_from_log\(\)](#).

5.105.1.10 write_log()

```
static void write_log (
    void ) [static]
```

Definition at line 178 of file [log.c](#).

```
00179 {
00180     int tail;
00181
00182     for (tail = 0; tail < log.lh.n; tail++) {
00183         struct buf *to = bread(log.dev, log.start+tail+1); // log block
00184         struct buf *from = bread(log.dev, log.lh.block[tail]); // cache block
00185         memmove(to->data, from->data, BSIZE);
00186         bwrite(to); // write the log
00187         brelse(from);
00188         brelse(to);
00189     }
00190 }
```

Referenced by [commit\(\)](#).

5.105.2 Variable Documentation

5.105.2.1 log

```
struct log log
```

Definition at line 48 of file [log.c](#).

5.106 log.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "spinlock.h"
00005 #include "sleeplock.h"
00006 #include "fs.h"
00007 #include "buf.h"
00008
00009 // Simple logging that allows concurrent FS system calls.
00010 //
00011 // A log transaction contains the updates of multiple FS system
00012 // calls. The logging system only commits when there are
00013 // no FS system calls active. Thus there is never
00014 // any reasoning required about whether a commit might
00015 // write an uncommitted system call's updates to disk.
00016 //
00017 // A system call should call begin_op()/end_op() to mark
00018 // its start and end. Usually begin_op() just increments
00019 // the count of in-progress FS system calls and returns.
00020 // But if it thinks the log is close to running out, it
00021 // sleeps until the last outstanding end_op() commits.
00022 //
00023 // The log is a physical re-do log containing disk blocks.
00024 // The on-disk log format:
00025 //   header block, containing block #s for block A, B, C, ...
00026 //   block A
00027 //   block B
00028 //   block C
00029 //   ...
00030 // Log appends are synchronous.
00031
00032 // Contents of the header block, used for both the on-disk header block
00033 // and to keep track in memory of logged block# before commit.
00034 struct logheader {
00035     int n;
00036     int block[LOGSIZE];
00037 };
00038
00039 struct log {
00040     struct spinlock lock;
00041     int start;
00042     int size;
00043     int outstanding; // how many FS sys calls are executing.
00044     int committing;  // in commit(), please wait.
00045     int dev;
00046     struct logheader lh;
00047 };
00048 struct log log;
00049
00050 static void recover_from_log(void);
00051 static void commit();
00052
00053 void
00054 initlog(int dev)
00055 {
00056     if (sizeof(struct logheader) >= BSIZE)
00057         panic("initlog: too big logheader");
00058
00059     struct superblock sb;
00060     initlock(&log.lock, "log");
00061     readsb(dev, &sb);
00062     log.start = sb.logstart;
00063     log.size = sb.nlog;
00064     log.dev = dev;
00065     recover_from_log();
00066 }
00067
00068 // Copy committed blocks from log to their home location
00069 static void
00070 install_trans(void)
00071 {
00072     int tail;
00073
00074     for (tail = 0; tail < log.lh.n; tail++) {
00075         struct buf *lbuf = bread(log.dev, log.start+tail+1); // read log block
00076         struct buf *dbuf = bread(log.dev, log.lh.block[tail]); // read dst
00077         memmove(dbuf->data, lbuf->data, BSIZE); // copy block to dst
00078         bwrite(dbuf); // write dst to disk
00079         brelse(lbuf);
00080         brelse(dbuf);
00081     }
00082 }

```

```

00083
00084 // Read the log header from disk into the in-memory log header
00085 static void
00086 read_head(void)
00087 {
00088     struct buf *buf = bread(log.dev, log.start);
00089     struct logheader *lh = (struct logheader *) (buf->data);
00090     int i;
00091     log.lh.n = lh->n;
00092     for (i = 0; i < log.lh.n; i++) {
00093         log.lh.block[i] = lh->block[i];
00094     }
00095     brelse(buf);
00096 }
00097
00098 // Write in-memory log header to disk.
00099 // This is the true point at which the
00100 // current transaction commits.
00101 static void
00102 write_head(void)
00103 {
00104     struct buf *buf = bread(log.dev, log.start);
00105     struct logheader *hb = (struct logheader *) (buf->data);
00106     int i;
00107     hb->n = log.lh.n;
00108     for (i = 0; i < log.lh.n; i++) {
00109         hb->block[i] = log.lh.block[i];
00110     }
00111     bwrite(buf);
00112     brelse(buf);
00113 }
00114
00115 static void
00116 recover_from_log(void)
00117 {
00118     read_head();
00119     install_trans(); // if committed, copy from log to disk
00120     log.lh.n = 0;
00121     write_head(); // clear the log
00122 }
00123
00124 // called at the start of each FS system call.
00125 void
00126 begin_op(void)
00127 {
00128     acquire(&log.lock);
00129     while(1){
00130         if(log.committing){
00131             sleep(&log, &log.lock);
00132         } else if(log.lh.n + (log.outstanding+1)*MAXOPBLOCKS > LOGSIZE){
00133             // this op might exhaust log space; wait for commit.
00134             sleep(&log, &log.lock);
00135         } else {
00136             log.outstanding += 1;
00137             release(&log.lock);
00138             break;
00139         }
00140     }
00141 }
00142
00143 // called at the end of each FS system call.
00144 // commits if this was the last outstanding operation.
00145 void
00146 end_op(void)
00147 {
00148     int do_commit = 0;
00149
00150     acquire(&log.lock);
00151     log.outstanding -= 1;
00152     if(log.committing)
00153         panic("log.committing");
00154     if(log.outstanding == 0){
00155         do_commit = 1;
00156         log.committing = 1;
00157     } else {
00158         // begin_op() may be waiting for log space,
00159         // and decrementing log.outstanding has decreased
00160         // the amount of reserved space.
00161         wakeup(&log);
00162     }
00163     release(&log.lock);
00164
00165     if(do_commit){
00166         // call commit w/o holding locks, since not allowed
00167         // to sleep with locks.
00168         commit();
00169         acquire(&log.lock);

```

```

00170     log.committing = 0;
00171     wakeup(&log);
00172     release(&log.lock);
00173 }
00174 }
00175
00176 // Copy modified blocks from cache to log.
00177 static void
00178 write_log(void)
00179 {
00180     int tail;
00181
00182     for (tail = 0; tail < log.lh.n; tail++) {
00183         struct buf *to = bread(log.dev, log.start+tail+1); // log block
00184         struct buf *from = bread(log.dev, log.lh.block[tail]); // cache block
00185         memmove(to->data, from->data, BSIZE);
00186         bwrite(to); // write the log
00187         brelse(from);
00188         brelse(to);
00189     }
00190 }
00191
00192 static void
00193 commit()
00194 {
00195     if (log.lh.n > 0) {
00196         write_log(); // Write modified blocks from cache to log
00197         write_head(); // Write header to disk -- the real commit
00198         install_trans(); // Now install writes to home locations
00199         log.lh.n = 0;
00200         write_head(); // Erase the transaction from the log
00201     }
00202 }
00203
00204 // Caller has modified b->data and is done with the buffer.
00205 // Record the block number and pin in the cache with B_DIRTY.
00206 // commit()/write_log() will do the disk write.
00207 //
00208 // log_write() replaces bwrite(); a typical use is:
00209 //   bp = bread(...)
00210 //   modify bp->data[]
00211 //   log_write(bp)
00212 //   brelse(bp)
00213 void
00214 log_write(struct buf *b)
00215 {
00216     int i;
00217
00218     if (log.lh.n >= LOGSIZE || log.lh.n >= log.size - 1)
00219         panic("too big a transaction");
00220     if (log.outstanding < 1)
00221         panic("log_write outside of trans");
00222
00223     acquire(&log.lock);
00224     for (i = 0; i < log.lh.n; i++) {
00225         if (log.lh.block[i] == b->blockno) // log absorption
00226             break;
00227     }
00228     log.lh.block[i] = b->blockno;
00229     if (i == log.lh.n)
00230         log.lh.n++;
00231     b->flags |= B_DIRTY; // prevent eviction
00232     release(&log.lock);
00233 }
00234

```

5.107 log.d File Reference

5.108 log.d

[Go to the documentation of this file.](#)

```

00001 log.o: log.c /usr/include/stdc-predef.h types.h defs.h param.h spinlock.h \
00002 sleeplock.h fs.h buf.h

```

5.109 ls.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fs.h"
```

Functions

- char * [fmtname](#) (char *path)
- void [ls](#) (char *path)
- int [main](#) (int argc, char *argv[])

5.109.1 Function Documentation

5.109.1.1 [fmtname\(\)](#)

```
char * fmtname (
    char * path )
```

Definition at line 7 of file [ls.c](#).

```
00008 {
00009     static char buf[DIRSIZ+1];
00010     char *p;
00011
00012     // Find first character after last slash.
00013     for(p=path+strlen(path); p >= path && *p != '/'; p--)
00014         ;
00015     p++;
00016
00017     // Return blank-padded name.
00018     if(strlen(p) >= DIRSIZ)
00019         return p;
00020     memmove(buf, p, strlen(p));
00021     memset(buf+strlen(p), ' ', DIRSIZ-strlen(p));
00022     return buf;
00023 }
```

Referenced by [ls\(\)](#).

5.109.1.2 [ls\(\)](#)

```
void ls (
    char * path )
```

Definition at line 26 of file [ls.c](#).

```
00027 {
00028     char buf[512], *p;
00029     int fd;
00030     struct dirent de;
00031     struct stat st;
00032
00033     if((fd = open(path, 0)) < 0){
00034         printf(2, "ls: cannot open %s\n", path);
```

```

00035     return;
00036 }
00037
00038 if(fstat(fd, &st) < 0){
00039     printf(2, "ls: cannot stat %s\n", path);
00040     close(fd);
00041     return;
00042 }
00043
00044 switch(st.type){
00045 case T_FILE:
00046     printf(1, "%s %d %d %d\n", fmtname(path), st.type, st.ino, st.size);
00047     break;
00048
00049 case T_DIR:
00050     if(strlen(path) + 1 + DIRSIZ + 1 > sizeof buf){
00051         printf(1, "ls: path too long\n");
00052         break;
00053     }
00054     strcpy(buf, path);
00055     p = buf+strlen(buf);
00056     *p++ = '/';
00057     while(read(fd, &de, sizeof(de)) == sizeof(de)){
00058         if(de.inum == 0)
00059             continue;
00060         memmove(p, de.name, DIRSIZ);
00061         p[DIRSIZ] = 0;
00062         if(stat(buf, &st) < 0){
00063             printf(1, "ls: cannot stat %s\n", buf);
00064             continue;
00065         }
00066         printf(1, "%s %d %d %d\n", fmtname(buf), st.type, st.ino, st.size);
00067     }
00068     break;
00069 }
00070 close(fd);
00071 }

```

Referenced by [main\(\)](#).

5.109.1.3 main()

```

int main (
    int argc,
    char * argv[] )

```

Definition at line 74 of file [ls.c](#).

```

00075 {
00076     int i;
00077
00078     if(argc < 2){
00079         ls(".");
00080         exit();
00081     }
00082     for(i=1; i<argc; i++){
00083         ls(argv[i]);
00084     }
00085 }

```

5.110 ls.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "fs.h"
00005
00006 char*
00007 fmtname(char *path)
00008 {
00009     static char buf[DIRSIZ+1];

```

```

00010 char *p;
00011
00012 // Find first character after last slash.
00013 for(p=path+strlen(path); p >= path && *p != '/'; p--)
00014 ;
00015 p++;
00016
00017 // Return blank-padded name.
00018 if(strlen(p) >= DIRSIZ)
00019     return p;
00020 memmove(buf, p, strlen(p));
00021 memset(buf+strlen(p), ' ', DIRSIZ-strlen(p));
00022 return buf;
00023 }
00024
00025 void
00026 ls(char *path)
00027 {
00028     char buf[512], *p;
00029     int fd;
00030     struct dirent de;
00031     struct stat st;
00032
00033     if((fd = open(path, 0)) < 0){
00034         printf(2, "ls: cannot open %s\n", path);
00035         return;
00036     }
00037
00038     if(fstat(fd, &st) < 0){
00039         printf(2, "ls: cannot stat %s\n", path);
00040         close(fd);
00041         return;
00042     }
00043
00044     switch(st.type){
00045     case T_FILE:
00046         printf(1, "%s %d %d %d\n", fmtname(path), st.type, st.ino, st.size);
00047         break;
00048
00049     case T_DIR:
00050         if(strlen(path) + 1 + DIRSIZ + 1 > sizeof buf){
00051             printf(1, "ls: path too long\n");
00052             break;
00053         }
00054         strcpy(buf, path);
00055         p = buf+strlen(buf);
00056         *p++ = '/';
00057         while(read(fd, &de, sizeof(de)) == sizeof(de)){
00058             if(de.inum == 0)
00059                 continue;
00060             memmove(p, de.name, DIRSIZ);
00061             p[DIRSIZ] = 0;
00062             if(stat(buf, &st) < 0){
00063                 printf(1, "ls: cannot stat %s\n", buf);
00064                 continue;
00065             }
00066             printf(1, "%s %d %d %d\n", fmtname(buf), st.type, st.ino, st.size);
00067         }
00068         break;
00069     }
00070     close(fd);
00071 }
00072
00073 int
00074 main(int argc, char *argv[])
00075 {
00076     int i;
00077
00078     if(argc < 2){
00079         ls(".");
00080         exit();
00081     }
00082     for(i=1; i<argc; i++)
00083         ls(argv[i]);
00084     exit();
00085 }

```


5.111 ls.d File Reference

5.112 ls.d

[Go to the documentation of this file.](#)

```
00001 ls.o: ls.c /usr/include/stdc-predef.h types.h stat.h user.h fs.h
```

5.113 main.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"
```

Functions

- [__attribute__](#) ((__aligned__(PGSIZE)))
- static void [mpenter](#) (void)
- static void [mpmain](#) (void)
- static void [startothers](#) (void)

Variables

- [pde_t](#) [entrypgdir](#) []

5.113.1 Function Documentation

5.113.1.1 __attribute__()

```
__attribute__ (
    (__aligned__(PGSIZE)) )
```

Definition at line 102 of file [main.c](#).

```
00103 {
00104     // Map VA's [0, 4MB) to PA's [0, 4MB)
00105     [0] = (0) | PTE_P | PTE_W | PTE_PS,
00106     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
00107     [KERNBASE»PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
00108 };
```

5.113.1.2 mpenter()

```
static void mpenter (
    void ) [static]
```

Definition at line 42 of file [main.c](#).

```
00043 {
00044     switchkvm();
00045     seginit();
00046     lapicinit();
00047     mpmain();
00048 }
```

Referenced by [startothers\(\)](#).

5.113.1.3 mpmain()

```
static void mpmain (
    void ) [static]
```

Definition at line 10 of file [main.c](#).

```
00019 {
00020     kinit1(end, P2V(4*1024*1024)); // phys page allocator
00021     kvmalloc(); // kernel page table
00022     mpinit(); // detect other processors
00023     lapicinit(); // interrupt controller
00024     seginit(); // segment descriptors
00025     picinit(); // disable pic
00026     ioapicinit(); // another interrupt controller
00027     consoleinit(); // console hardware
00028     uartinit(); // serial port
00029     pinit(); // process table
00030     tvinit(); // trap vectors
00031     binit(); // buffer cache
00032     fileinit(); // file table
00033     ideinit(); // disk
00034     startothers(); // start other processors
00035     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
00036     userinit(); // first user process
00037     mpmain(); // finish this processor's setup
00038 }
```

Referenced by [mpenter\(\)](#).

5.113.1.4 startothers()

```
static void startothers (
    void ) [static]
```

Definition at line 64 of file [main.c](#).

```
00065 {
00066     extern uchar _binary_entryother_start[], _binary_entryother_size[];
00067     uchar *code;
00068     struct cpu *c;
00069     char *stack;
00070
00071     // Write entry code to unused memory at 0x7000.
00072     // The linker has placed the image of entryother.S in
00073     // _binary_entryother_start.
00074     code = P2V(0x7000);
00075     memmove(code, _binary_entryother_start, (uint)_binary_entryother_size);
00076
00077     for(c = cpus; c < cpus+ncpu; c++){
00078         if(c == mycpu()) // We've started already.
```

```

00079         continue;
00080
00081         // Tell entryother.S what stack to use, where to enter, and what
00082         // pgdir to use. We cannot use kpgdir yet, because the AP processor
00083         // is running in low memory, so we use entrypgdir for the APs too.
00084         stack = kalloc();
00085         *(void**) (code-4) = stack + KSTACKSIZE;
00086         *(void**) (void) (code-8) = mpenter;
00087         *(int**) (code-12) = (void *) V2P(entrypgdir);
00088
00089         lapicstartap(c->apicid, V2P(code));
00090
00091         // wait for cpu to finish mpmain()
00092         while(c->started == 0)
00093             ;
00094     }
00095 }

```

5.113.2 Variable Documentation

5.113.2.1 entrypgdir

`pde_t entrypgdir[]`

Definition at line 60 of file [main.c](#).

Referenced by [startothers\(\)](#).

5.114 main.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "memlayout.h"
00005 #include "mmu.h"
00006 #include "proc.h"
00007 #include "x86.h"
00008
00009 static void startothers(void);
00010 static void mpmain(void) __attribute__((noreturn));
00011 extern pde_t *kpgdir;
00012 extern char end[]; // first address after kernel loaded from ELF file
00013
00014 // Bootstrap processor starts running C code here.
00015 // Allocate a real stack and switch to it, first
00016 // doing some setup required for memory allocator to work.
00017 int
00018 main(void)
00019 {
00020     kinit1(end, P2V(4*1024*1024)); // phys page allocator
00021     kvmalloc(); // kernel page table
00022     mpinit(); // detect other processors
00023     lapicinit(); // interrupt controller
00024     seginit(); // segment descriptors
00025     picinit(); // disable pic
00026     ioapicinit(); // another interrupt controller
00027     consoleinit(); // console hardware
00028     uartinit(); // serial port
00029     pinit(); // process table
00030     tvinit(); // trap vectors
00031     binit(); // buffer cache
00032     fileinit(); // file table
00033     ideinit(); // disk
00034     startothers(); // start other processors
00035     kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must come after startothers()
00036     userinit(); // first user process
00037     mpmain(); // finish this processor's setup

```

```

00038 }
00039
00040 // Other CPUs jump here from entryother.S.
00041 static void
00042 mpenter(void)
00043 {
00044     switchkvm();
00045     seginit();
00046     lapicinit();
00047     mpmain();
00048 }
00049
00050 // Common CPU setup code.
00051 static void
00052 mpmain(void)
00053 {
00054     cprintf("cpu%d: starting %d\n", cpuid(), cpuid());
00055     idtinit(); // load idt register
00056     xchg(&(mycpu()->started), 1); // tell startothers() we're up
00057     scheduler(); // start running processes
00058 }
00059
00060 pde_t entrypgdir[]; // For entry.S
00061
00062 // Start the non-boot (AP) processors.
00063 static void
00064 startothers(void)
00065 {
00066     extern uchar _binary_entryother_start[], _binary_entryother_size[];
00067     uchar *code;
00068     struct cpu *c;
00069     char *stack;
00070
00071     // Write entry code to unused memory at 0x7000.
00072     // The linker has placed the image of entryother.S in
00073     // _binary_entryother_start.
00074     code = P2V(0x7000);
00075     memmove(code, _binary_entryother_start, (uint)_binary_entryother_size);
00076
00077     for(c = cpus; c < cpus+ncpu; c++){
00078         if(c == mycpu()) // We've started already.
00079             continue;
00080
00081         // Tell entryother.S what stack to use, where to enter, and what
00082         // pgdir to use. We cannot use kpgdir yet, because the AP processor
00083         // is running in low memory, so we use entrypgdir for the APs too.
00084         stack = kalloc();
00085         *(void**) (code-4) = stack + KSTACKSIZE;
00086         *(void**) (code-8) = mpenter;
00087         *(int**) (code-12) = (void *) V2P(entrypgdir);
00088
00089         lapicstartap(c->apicid, V2P(code));
00090
00091         // wait for cpu to finish mpmain()
00092         while(c->started == 0)
00093             ;
00094     }
00095 }
00096
00097 // The boot page table used in entry.S and entryother.S.
00098 // Page directories (and page tables) must start on page boundaries,
00099 // hence the __aligned__ attribute.
00100 // PTE_PS in a page directory entry enables 4Mbyte pages.
00101
00102 __attribute__((__aligned__(PGSIZE)))
00103 pde_t entrypgdir[NPDENTRIES] = {
00104     // Map VA's [0, 4MB) to PA's [0, 4MB)
00105     [0] = (0) | PTE_P | PTE_W | PTE_PS,
00106     // Map VA's [KERNBASE, KERNBASE+4MB) to PA's [0, 4MB)
00107     [KERNBASE>>PDXSHIFT] = (0) | PTE_P | PTE_W | PTE_PS,
00108 };
00109
00110 //PAGEBREAK!
00111 // Blank page.
00112 //PAGEBREAK!
00113 // Blank page.
00114 //PAGEBREAK!
00115 // Blank page.
00116

```

5.115 main.d File Reference

5.116 main.d

[Go to the documentation of this file.](#)

```
00001 main.o: main.c /usr/include/stdc-predef.h types.h defs.h param.h \  
00002 memlayout.h mmu.h proc.h x86.h
```

5.117 memide.c File Reference

```
#include "types.h"  
#include "defs.h"  
#include "param.h"  
#include "mmu.h"  
#include "proc.h"  
#include "x86.h"  
#include "traps.h"  
#include "spinlock.h"  
#include "sleeplock.h"  
#include "fs.h"  
#include "buf.h"
```

Functions

- void [ideinit](#) (void)
- void [ideintr](#) (void)
- void [iderw](#) (struct [buf](#) *b)

Variables

- [uchar](#) [_binary_fs_img_size](#) []
- [uchar](#) [_binary_fs_img_start](#) []
- static int [disksize](#)
- static [uchar](#) * [memdisk](#)

5.117.1 Function Documentation

5.117.1.1 ideinit()

```
void ideinit (  
    void )
```

Definition at line 22 of file [memide.c](#).

```
00023 {  
00024     memdisk = \_binary\_fs\_img\_start;  
00025     disksize = (uint)\_binary\_fs\_img\_size/BSIZE;  
00026 }
```

5.117.1.2 ideintr()

```
void ideintr (
    void )
```

Definition at line 30 of file [memide.c](#).

```
00031 {
00032     // no-op
00033 }
```

Referenced by [trap\(\)](#).

5.117.1.3 iderw()

```
void iderw (
    struct buf * b )
```

Definition at line 39 of file [memide.c](#).

```
00040 {
00041     uchar *p;
00042
00043     if(!holdingsleep(&b->lock))
00044         panic("iderw: buf not locked");
00045     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
00046         panic("iderw: nothing to do");
00047     if(b->dev != 1)
00048         panic("iderw: request not for disk 1");
00049     if(b->blockno >= disksize)
00050         panic("iderw: block out of range");
00051
00052     p = memdisk + b->blockno*BSIZE;
00053
00054     if(b->flags & B_DIRTY){
00055         b->flags &= ~B_DIRTY;
00056         memmove(p, b->data, BSIZE);
00057     } else
00058         memmove(b->data, p, BSIZE);
00059     b->flags |= B_VALID;
00060 }
```

Referenced by [bread\(\)](#), and [bwrite\(\)](#).

5.117.2 Variable Documentation

5.117.2.1 _binary_fs_img_size

```
uchar _binary_fs_img_size[ ]
```

Definition at line 16 of file [memide.c](#).

Referenced by [ideinit\(\)](#).

5.117.2.2 _binary_fs_img_start

```
uchar _binary_fs_img_start[] [extern]
```

Referenced by [ideinit\(\)](#).

5.117.2.3 disksize

```
int disksize [static]
```

Definition at line 18 of file [memide.c](#).

Referenced by [ideinit\(\)](#), and [iderw\(\)](#).

5.117.2.4 memdisk

```
uchar* memdisk [static]
```

Definition at line 19 of file [memide.c](#).

Referenced by [ideinit\(\)](#), and [iderw\(\)](#).

5.118 memide.c

[Go to the documentation of this file.](#)

```
00001 // Fake IDE disk; stores blocks in memory.
00002 // Useful for running kernel without scratch disk.
00003
00004 #include "types.h"
00005 #include "defs.h"
00006 #include "param.h"
00007 #include "mmu.h"
00008 #include "proc.h"
00009 #include "x86.h"
00010 #include "traps.h"
00011 #include "spinlock.h"
00012 #include "sleeplock.h"
00013 #include "fs.h"
00014 #include "buf.h"
00015
00016 extern uchar _binary_fs_img_start[], _binary_fs_img_size[];
00017
00018 static int disksize;
00019 static uchar *memdisk;
00020
00021 void
00022 ideinit(void)
00023 {
00024     memdisk = _binary_fs_img_start;
00025     disksize = (uint)_binary_fs_img_size/BSIZE;
00026 }
00027
00028 // Interrupt handler.
00029 void
00030 ideintr(void)
00031 {
00032     // no-op
00033 }
00034
00035 // Sync buf with disk.
```

```

00036 // If B_DIRTY is set, write buf to disk, clear B_DIRTY, set B_VALID.
00037 // Else if B_VALID is not set, read buf from disk, set B_VALID.
00038 void
00039 iderw(struct buf *b)
00040 {
00041     uchar *p;
00042
00043     if(!holdingsleep(&b->lock))
00044         panic("iderw: buf not locked");
00045     if((b->flags & (B_VALID|B_DIRTY)) == B_VALID)
00046         panic("iderw: nothing to do");
00047     if(b->dev != 1)
00048         panic("iderw: request not for disk 1");
00049     if(b->blockno >= disksize)
00050         panic("iderw: block out of range");
00051
00052     p = memdisk + b->blockno*BSIZE;
00053
00054     if(b->flags & B_DIRTY){
00055         b->flags &= ~B_DIRTY;
00056         memmove(p, b->data, BSIZE);
00057     } else
00058         memmove(b->data, p, BSIZE);
00059     b->flags |= B_VALID;
00060 }

```

5.119 memlayout.h File Reference

Macros

- #define [DEVSPACE](#) 0xFE000000
- #define [EXTMEM](#) 0x100000
- #define [KERNBASE](#) 0x80000000
- #define [KERNLINK](#) ([KERNBASE](#)+[EXTMEM](#))
- #define [P2V](#)(a) ((void *)(((char *) (a)) + [KERNBASE](#)))
- #define [P2V_WO](#)(x) ((x) + [KERNBASE](#))
- #define [PHYSTOP](#) 0xE000000
- #define [V2P](#)(a) (((uint) (a)) - [KERNBASE](#))
- #define [V2P_WO](#)(x) ((x) - [KERNBASE](#))

5.119.1 Macro Definition Documentation

5.119.1.1 DEVSPACE

```
#define DEVSPACE 0xFE000000
```

Definition at line 5 of file [memlayout.h](#).

5.119.1.2 EXTMEM

```
#define EXTMEM 0x100000
```

Definition at line 3 of file [memlayout.h](#).

5.119.1.3 KERNBASE

```
#define KERNBASE 0x80000000
```

Definition at line 8 of file [memlayout.h](#).

5.119.1.4 KERMLINK

```
#define KERMLINK (KERNBASE+EXTMEM)
```

Definition at line 9 of file [memlayout.h](#).

5.119.1.5 P2V

```
#define P2V(  
    a ) ((void *)(((char *) (a)) + KERNBASE))
```

Definition at line 12 of file [memlayout.h](#).

5.119.1.6 P2V_WO

```
#define P2V_WO(  
    x ) ((x) + KERNBASE)
```

Definition at line 15 of file [memlayout.h](#).

5.119.1.7 PHYSTOP

```
#define PHYSTOP 0xE000000
```

Definition at line 4 of file [memlayout.h](#).

5.119.1.8 V2P

```
#define V2P(  
    a ) (((uint) (a)) - KERNBASE)
```

Definition at line 11 of file [memlayout.h](#).

5.119.1.9 V2P_WO

```
#define V2P_WO(
    x ) ((x) - KERNBASE)
```

Definition at line 14 of file [memlayout.h](#).

5.120 memlayout.h

[Go to the documentation of this file.](#)

```
00001 // Memory layout
00002
00003 #define EXTMEM 0x100000 // Start of extended memory
00004 #define PHYSTOP 0xE000000 // Top physical memory
00005 #define DEVSPACE 0xFE000000 // Other devices are at high addresses
00006
00007 // Key addresses for address space layout (see kmap in vm.c for layout)
00008 #define KERNBASE 0x80000000 // First kernel virtual address
00009 #define KERNLINK (KERNBASE+EXTMEM) // Address where kernel is linked
00010
00011 #define V2P(a) (((uint) (a)) - KERNBASE)
00012 #define P2V(a) ((void *) (((char *) (a)) + KERNBASE))
00013
00014 #define V2P_WO(x) ((x) - KERNBASE) // same as V2P, but without casts
00015 #define P2V_WO(x) ((x) + KERNBASE) // same as P2V, but without casts
```

5.121 mkdir.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- [int main](#) (int argc, char *argv[])

5.121.1 Function Documentation

5.121.1.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 6 of file [mkdir.c](#).

```
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "Usage: mkdir files...\n");
00012         exit();
00013     }
00014
00015     for(i = 1; i < argc; i++){
00016         if(mkdir(argv[i]) < 0){
00017             printf(2, "mkdir: %s failed to create\n", argv[i]);
00018             break;
00019         }
00020     }
00021
00022     exit();
00023 }
```

5.122 mkdir.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 int
00006 main(int argc, char *argv[])
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "Usage: mkdir files...\n");
00012         exit();
00013     }
00014
00015     for(i = 1; i < argc; i++){
00016         if(mkdir(argv[i]) < 0){
00017             printf(2, "mkdir: %s failed to create\n", argv[i]);
00018             break;
00019         }
00020     }
00021
00022     exit();
00023 }
```

5.123 mkdir.d File Reference

5.124 mkdir.d

[Go to the documentation of this file.](#)

```
00001 mkdir.o: mkdir.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.125 mkfs.c File Reference

```
#include <stdio.h>
#include <unistd.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <assert.h>
#include "types.h"
#include "fs.h"
#include "stat.h"
#include "param.h"
```

Macros

- #define [min](#)(a, b) ((a) < (b) ? (a) : (b))
- #define [NINODES](#) 200
- #define [stat](#) xv6_stat
- #define [static_assert](#)(a, b) do { switch (0) case 0: case (a): ; } while (0)

Functions

- void `balloc` (int)
- uint `ialloc` (ushort type)
- void `iappend` (uint inum, void *p, int n)
- int `main` (int argc, char *argv[])
- void `rinode` (uint inum, struct `dinode` *ip)
- void `rsect` (uint sec, void *buf)
- void `winode` (uint, struct `dinode` *)
- void `wsect` (uint, void *)
- uint `xint` (uint x)
- ushort `xshort` (ushort x)

Variables

- uint `freeblock`
- uint `freeinode` = 1
- int `fsfd`
- int `nbitmap` = `FSSIZE` / (`BFSIZE` * 8) + 1
- int `nblocks`
- int `ninodeblocks` = `NINODES` / `IPB` + 1
- int `nlog` = `LOGSIZE`
- int `nmeta`
- struct `superblock` `sb`
- char `zeroes` [`BFSIZE`]

5.125.1 Macro Definition Documentation

5.125.1.1 min

```
#define min(  
    a,  
    b ) ((a) < (b) ? (a) : (b))
```

Definition at line 253 of file `mkfs.c`.

5.125.1.2 NINODES

```
#define NINODES 200
```

Definition at line 18 of file `mkfs.c`.

5.125.1.3 stat

```
#define stat xv6_stat
```

Definition at line 8 of file [mkfs.c](#).

5.125.1.4 static_assert

```
#define static_assert(  
    a,  
    b ) do { switch (0) case 0: case (a): ; } while (0)
```

Definition at line 15 of file [mkfs.c](#).

5.125.2 Function Documentation

5.125.2.1 balloc()

```
void balloc (  
    int used )
```

Definition at line 238 of file [mkfs.c](#).

```
00239 {  
00240     uchar buf[BSIZE];  
00241     int i;  
00242  
00243     printf("balloc: first %d blocks have been allocated\n", used);  
00244     assert(used < BSIZE*8);  
00245     bzero(buf, BSIZE);  
00246     for(i = 0; i < used; i++){  
00247         buf[i/8] = buf[i/8] | (0x1 << (i%8));  
00248     }  
00249     printf("balloc: write bitmap block at sector %d\n", sb.bmapstart);  
00250     wsect(sb.bmapstart, buf);  
00251 }
```

Referenced by [main\(\)](#).

5.125.2.2 ialloc()

```
uint ialloc (  
    ushort type )
```

Definition at line 224 of file [mkfs.c](#).

```
00225 {  
00226     uint inum = freeinode++;  
00227     struct dinode din;  
00228  
00229     bzero(&din, sizeof(din));  
00230     din.type = xshort(type);  
00231     din.nlink = xshort(1);  
00232     din.size = xint(0);  
00233     winode(inum, &din);  
00234     return inum;  
00235 }
```

Referenced by [main\(\)](#).

5.125.2.3 iappend()

```
void iappend (
    uint inum,
    void * p,
    int n )
```

Definition at line 256 of file [mkfs.c](#).

```
00257 {
00258     char *p = (char*)xp;
00259     uint fbn, off, nl;
00260     struct dinode din;
00261     char buf[BSIZE];
00262     uint indirect[NINDIRECT];
00263     uint x;
00264
00265     rinode(inum, &din);
00266     off = xint(din.size);
00267     // printf("append inum %d at off %d sz %d\n", inum, off, n);
00268     while(n > 0){
00269         fbn = off / BSIZE;
00270         assert(fbn < MAXFILE);
00271         if(fbn < NDIRECT){
00272             if(xint(din.addrs[fbn]) == 0){
00273                 din.addrs[fbn] = xint(freeblock++);
00274             }
00275             x = xint(din.addrs[fbn]);
00276         } else {
00277             if(xint(din.addrs[NDIRECT]) == 0){
00278                 din.addrs[NDIRECT] = xint(freeblock++);
00279             }
00280             rsect(xint(din.addrs[NDIRECT]), (char*)indirect);
00281             if(indirect[fbn - NDIRECT] == 0){
00282                 indirect[fbn - NDIRECT] = xint(freeblock++);
00283                 wsect(xint(din.addrs[NDIRECT]), (char*)indirect);
00284             }
00285             x = xint(indirect[fbn-NDIRECT]);
00286         }
00287         nl = min(n, (fbn + 1) * BSIZE - off);
00288         rsect(x, buf);
00289         bcopy(p, buf + off - (fbn * BSIZE), nl);
00290         wsect(x, buf);
00291         n -= nl;
00292         off += nl;
00293         p += nl;
00294     }
00295     din.size = xint(off);
00296     winode(inum, &din);
00297 }
```

Referenced by [main\(\)](#).

5.125.2.4 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 68 of file [mkfs.c](#).

```
00069 {
00070     int i, cc, fd;
00071     uint rootino, inum, off;
00072     struct dirent de;
00073     char buf[BSIZE];
00074     struct dinode din;
00075
00076
00077     static_assert(sizeof(int) == 4, "Integers must be 4 bytes!");
00078
00079     if(argc < 2){
00080         fprintf(stderr, "Usage: mkfs fs.img files...\n");
00081         exit(1);
00082     }
```

```

00082     }
00083
00084     assert((BSIZE % sizeof(struct dinode)) == 0);
00085     assert((BSIZE % sizeof(struct dirent)) == 0);
00086
00087     fsfd = open(argv[1], O_RDWR|O_CREAT|O_TRUNC, 0666);
00088     if(fsfd < 0){
00089         perror(argv[1]);
00090         exit(1);
00091     }
00092
00093     // 1 fs block = 1 disk sector
00094     nmeta = 2 + nlog + ninodeblocks + nbitmap;
00095     nblocks = FSSIZE - nmeta;
00096
00097     sb.size = xint(FSSIZE);
00098     sb.nblocks = xint(nblocks);
00099     sb.ninodes = xint(NINODES);
00100     sb.nlog = xint(nlog);
00101     sb.logstart = xint(2);
00102     sb.inodestart = xint(2+nlog);
00103     sb.bmapstart = xint(2+nlog+ninodeblocks);
00104
00105     printf("nmeta %d (boot, super, log blocks %u inode blocks %u, bitmap blocks %u) blocks %d total\n",
00106           nmeta, nlog, ninodeblocks, nbitmap, nblocks, FSSIZE);
00107
00108     freeblock = nmeta;    // the first free block that we can allocate
00109
00110     for(i = 0; i < FSSIZE; i++){
00111         wsect(i, zeroes);
00112
00113         memset(buf, 0, sizeof(buf));
00114         memmove(buf, &sb, sizeof(sb));
00115         wsect(1, buf);
00116
00117         rootino = ialloc(T_DIR);
00118         assert(rootino == ROOTINO);
00119
00120         bzero(&de, sizeof(de));
00121         de.inum = xshort(rootino);
00122         strcpy(de.name, ".");
00123         iappend(rootino, &de, sizeof(de));
00124
00125         bzero(&de, sizeof(de));
00126         de.inum = xshort(rootino);
00127         strcpy(de.name, "..");
00128         iappend(rootino, &de, sizeof(de));
00129
00130         for(i = 2; i < argc; i++){
00131             assert(index(argv[i], '/') == 0);
00132
00133             if((fd = open(argv[i], 0)) < 0){
00134                 perror(argv[i]);
00135                 exit(1);
00136             }
00137
00138             // Skip leading _ in name when writing to file system.
00139             // The binaries are named _rm, _cat, etc. to keep the
00140             // build operating system from trying to execute them
00141             // in place of system binaries like rm and cat.
00142             if(argv[i][0] == '_')
00143                 ++argv[i];
00144
00145             inum = ialloc(T_FILE);
00146
00147             bzero(&de, sizeof(de));
00148             de.inum = xshort(inum);
00149             strncpy(de.name, argv[i], DIRSIZ);
00150             iappend(rootino, &de, sizeof(de));
00151
00152             while((cc = read(fd, buf, sizeof(buf))) > 0)
00153                 iappend(inum, buf, cc);
00154
00155             close(fd);
00156         }
00157
00158         // fix size of root inode dir
00159         rinode(rootino, &din);
00160         off = xint(din.size);
00161         off = ((off/BSIZE) + 1) * BSIZE;
00162         din.size = xint(off);
00163         winode(rootino, &din);
00164
00165         balloc(freeblock);
00166
00167         exit(0);

```

```
00168 }
```

5.125.2.5 rinode()

```
void rinode (
    uint inum,
    struct dinode * ip )
```

Definition at line 198 of file [mkfs.c](#).

```
00199 {
00200     char buf[BSIZE];
00201     uint bn;
00202     struct dinode *dip;
00203
00204     bn = IBLOCK(inum, sb);
00205     rsect(bn, buf);
00206     dip = ((struct dinode*)buf) + (inum % IPB);
00207     *ip = *dip;
00208 }
```

Referenced by [iappend\(\)](#), and [main\(\)](#).

5.125.2.6 rsect()

```
void rsect (
    uint sec,
    void * buf )
```

Definition at line 211 of file [mkfs.c](#).

```
00212 {
00213     if(lseek(fsfd, sec * BSIZE, 0) != sec * BSIZE){
00214         perror("lseek");
00215         exit(1);
00216     }
00217     if(read(fsfd, buf, BSIZE) != BSIZE){
00218         perror("read");
00219         exit(1);
00220     }
00221 }
```

Referenced by [iappend\(\)](#), [rinode\(\)](#), and [winode\(\)](#).

5.125.2.7 winode()

```
void winode (
    uint inum,
    struct dinode * ip )
```

Definition at line 184 of file [mkfs.c](#).

```
00185 {
00186     char buf[BSIZE];
00187     uint bn;
00188     struct dinode *dip;
00189
00190     bn = IBLOCK(inum, sb);
00191     rsect(bn, buf);
00192     dip = ((struct dinode*)buf) + (inum % IPB);
00193     *dip = *ip;
00194     wsect(bn, buf);
00195 }
```

Referenced by [ialloc\(\)](#), [iappend\(\)](#), and [main\(\)](#).

5.125.2.8 wsect()

```
void wsect (
    uint sec,
    void * buf )
```

Definition at line 171 of file [mkfs.c](#).

```
00172 {
00173     if(lseek(fsfd, sec * BSIZE, 0) != sec * BSIZE){
00174         perror("lseek");
00175         exit(1);
00176     }
00177     if(write(fsfd, buf, BSIZE) != BSIZE){
00178         perror("write");
00179         exit(1);
00180     }
00181 }
```

Referenced by [balloc\(\)](#), [iappend\(\)](#), [main\(\)](#), and [winode\(\)](#).

5.125.2.9 xint()

```
uint xint (
    uint x )
```

Definition at line 56 of file [mkfs.c](#).

```
00057 {
00058     uint y;
00059     uchar *a = (uchar*)&y;
00060     a[0] = x;
00061     a[1] = x » 8;
00062     a[2] = x » 16;
00063     a[3] = x » 24;
00064     return y;
00065 }
```

Referenced by [ialloc\(\)](#), [iappend\(\)](#), and [main\(\)](#).

5.125.2.10 xshort()

```
ushort xshort (
    ushort x )
```

Definition at line 46 of file [mkfs.c](#).

```
00047 {
00048     ushort y;
00049     uchar *a = (uchar*)&y;
00050     a[0] = x;
00051     a[1] = x » 8;
00052     return y;
00053 }
```

Referenced by [ialloc\(\)](#), and [main\(\)](#).

5.125.3 Variable Documentation

5.125.3.1 freeblock

```
uint freeblock
```

Definition at line 33 of file [mkfs.c](#).

Referenced by [iappend\(\)](#), and [main\(\)](#).

5.125.3.2 freeinode

```
uint freeinode = 1
```

Definition at line 32 of file [mkfs.c](#).

Referenced by [ialloc\(\)](#).

5.125.3.3 fsfd

```
int fsfd
```

Definition at line 29 of file [mkfs.c](#).

Referenced by [main\(\)](#), [rsect\(\)](#), and [wsect\(\)](#).

5.125.3.4 nbitmap

```
int nbitmap = FSSIZE/(BSIZE*8) + 1
```

Definition at line 23 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.125.3.5 nblocks

```
int nblocks
```

Definition at line 27 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.125.3.6 ninodeblocks

```
int ninodeblocks = NINODES / IPB + 1
```

Definition at line 24 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.125.3.7 nlog

```
int nlog = LOGSIZE
```

Definition at line 25 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.125.3.8 nmeta

```
int nmeta
```

Definition at line 26 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.125.3.9 sb

```
struct superblock sb
```

Definition at line 30 of file [mkfs.c](#).

Referenced by [balloc\(\)](#), [main\(\)](#), [rinode\(\)](#), and [winode\(\)](#).

5.125.3.10 zeroes

```
char zeroes[BSIZE]
```

Definition at line 31 of file [mkfs.c](#).

Referenced by [main\(\)](#).

5.126 mkfs.c

[Go to the documentation of this file.](#)

```

00001 #include <stdio.h>
00002 #include <unistd.h>
00003 #include <stdlib.h>
00004 #include <string.h>
00005 #include <fcntl.h>
00006 #include <assert.h>
00007
00008 #define stat xv6_stat  // avoid clash with host struct stat
00009 #include "types.h"
00010 #include "fs.h"
00011 #include "stat.h"
00012 #include "param.h"
00013
00014 #ifndef static_assert
00015 #define static_assert(a, b) do { switch (0) case 0: case (a): ; } while (0)
00016 #endif
00017
00018 #define NINODES 200
00019
00020 // Disk layout:
00021 // [ boot block | sb block | log | inode blocks | free bit map | data blocks ]
00022
00023 int nbitmap = FSSIZE/(BSIZE*8) + 1;
00024 int ninodeblocks = NINODES / IPB + 1;
00025 int nlog = LOGSIZE;
00026 int nmeta; // Number of meta blocks (boot, sb, nlog, inode, bitmap)
00027 int nblocks; // Number of data blocks
00028
00029 int fsfd;
00030 struct superblock sb;
00031 char zeroes[BSIZE];
00032 uint freeinode = 1;
00033 uint freeblock;
00034
00035 void balloc(int);
00036 void wsect(uint, void*);
00037 void winode(uint, struct dinode*);
00038 void rinode(uint inum, struct dinode *ip);
00039 void rsect(uint sec, void *buf);
00040 uint ialloc(ushort type);
00041 void iappend(uint inum, void *p, int n);
00042
00043 // convert to intel byte order
00044 ushort
00045 xshort(ushort x)
00046 {
00047     ushort y;
00048     uchar *a = (uchar*)&y;
00049     a[0] = x;
00050     a[1] = x >> 8;
00051     return y;
00052 }
00053
00054 uint
00055 xint(uint x)
00056 {
00057     uint y;
00058     uchar *a = (uchar*)&y;
00059     a[0] = x;
00060     a[1] = x >> 8;
00061     a[2] = x >> 16;
00062     a[3] = x >> 24;
00063     return y;
00064 }
00065
00066 int
00067 main(int argc, char *argv[])
00068 {
00069     int i, cc, fd;
00070     uint rootino, inum, off;
00071     struct dirent de;
00072     char buf[BSIZE];
00073     struct dinode din;
00074
00075     static_assert(sizeof(int) == 4, "Integers must be 4 bytes!");
00076
00077     if(argc < 2){
00078         fprintf(stderr, "Usage: mkfs fs.img files...\n");
00079         exit(1);
00080     }
00081 }

```

```

00083
00084     assert((BSIZE % sizeof(struct dinode)) == 0);
00085     assert((BSIZE % sizeof(struct dirent)) == 0);
00086
00087     fsfd = open(argv[1], O_RDWR|O_CREAT|O_TRUNC, 0666);
00088     if(fsfd < 0){
00089         perror(argv[1]);
00090         exit(1);
00091     }
00092
00093     // 1 fs block = 1 disk sector
00094     nmeta = 2 + nlog + ninodeblocks + nbitmap;
00095     nblocks = FSSIZE - nmeta;
00096
00097     sb.size = xint(FSSIZE);
00098     sb.nblocks = xint(nblocks);
00099     sb.ninodes = xint(NINODES);
00100     sb.nlog = xint(nlog);
00101     sb.logstart = xint(2);
00102     sb.inodestart = xint(2+nlog);
00103     sb.bmapstart = xint(2+nlog+ninodeblocks);
00104
00105     printf("nmeta %d (boot, super, log blocks %u inode blocks %u, bitmap blocks %u) blocks %d total
%d\n",
00106           nmeta, nlog, ninodeblocks, nbitmap, nblocks, FSSIZE);
00107
00108     freeblock = nmeta;    // the first free block that we can allocate
00109
00110     for(i = 0; i < FSSIZE; i++){
00111         wsect(i, zeroes);
00112     }
00113     memset(buf, 0, sizeof(buf));
00114     memmove(buf, &sb, sizeof(sb));
00115     wsect(1, buf);
00116
00117     rootino = ialloc(T_DIR);
00118     assert(rootino == ROOTINO);
00119
00120     bzero(&de, sizeof(de));
00121     de.inum = xshort(rootino);
00122     strcpy(de.name, ".");
00123     iappend(rootino, &de, sizeof(de));
00124
00125     bzero(&de, sizeof(de));
00126     de.inum = xshort(rootino);
00127     strcpy(de.name, "..");
00128     iappend(rootino, &de, sizeof(de));
00129
00130     for(i = 2; i < argc; i++){
00131         assert(index(argv[i], '/') == 0);
00132
00133         if((fd = open(argv[i], 0)) < 0){
00134             perror(argv[i]);
00135             exit(1);
00136         }
00137
00138         // Skip leading _ in name when writing to file system.
00139         // The binaries are named _rm, _cat, etc. to keep the
00140         // build operating system from trying to execute them
00141         // in place of system binaries like rm and cat.
00142         if(argv[i][0] == '_')
00143             ++argv[i];
00144
00145         inum = ialloc(T_FILE);
00146
00147         bzero(&de, sizeof(de));
00148         de.inum = xshort(inum);
00149         strncpy(de.name, argv[i], DIRSIZ);
00150         iappend(rootino, &de, sizeof(de));
00151
00152         while((cc = read(fd, buf, sizeof(buf))) > 0)
00153             iappend(inum, buf, cc);
00154
00155         close(fd);
00156     }
00157
00158     // fix size of root inode dir
00159     rinode(rootino, &din);
00160     off = xint(din.size);
00161     off = ((off/BSIZE) + 1) * BSIZE;
00162     din.size = xint(off);
00163     winode(rootino, &din);
00164
00165     balloc(freeblock);
00166
00167     exit(0);
00168 }

```

```

00169
00170 void
00171 wsect(uint sec, void *buf)
00172 {
00173     if(lseek(fsfd, sec * BSIZE, 0) != sec * BSIZE){
00174         perror("lseek");
00175         exit(1);
00176     }
00177     if(write(fsfd, buf, BSIZE) != BSIZE){
00178         perror("write");
00179         exit(1);
00180     }
00181 }
00182
00183 void
00184 winode(uint inum, struct dinode *ip)
00185 {
00186     char buf[BSIZE];
00187     uint bn;
00188     struct dinode *dip;
00189
00190     bn = IBLOCK(inum, sb);
00191     rsect(bn, buf);
00192     dip = ((struct dinode*)buf) + (inum % IPB);
00193     *dip = *ip;
00194     wsect(bn, buf);
00195 }
00196
00197 void
00198 rinode(uint inum, struct dinode *ip)
00199 {
00200     char buf[BSIZE];
00201     uint bn;
00202     struct dinode *dip;
00203
00204     bn = IBLOCK(inum, sb);
00205     rsect(bn, buf);
00206     dip = ((struct dinode*)buf) + (inum % IPB);
00207     *ip = *dip;
00208 }
00209
00210 void
00211 rsect(uint sec, void *buf)
00212 {
00213     if(lseek(fsfd, sec * BSIZE, 0) != sec * BSIZE){
00214         perror("lseek");
00215         exit(1);
00216     }
00217     if(read(fsfd, buf, BSIZE) != BSIZE){
00218         perror("read");
00219         exit(1);
00220     }
00221 }
00222
00223 uint
00224 ialloc(ushort type)
00225 {
00226     uint inum = freeinode++;
00227     struct dinode din;
00228
00229     bzero(&din, sizeof(din));
00230     din.type = xshort(type);
00231     din.nlink = xshort(1);
00232     din.size = xint(0);
00233     winode(inum, &din);
00234     return inum;
00235 }
00236
00237 void
00238 balloc(int used)
00239 {
00240     uchar buf[BSIZE];
00241     int i;
00242
00243     printf("balloc: first %d blocks have been allocated\n", used);
00244     assert(used < BSIZE*8);
00245     bzero(buf, BSIZE);
00246     for(i = 0; i < used; i++){
00247         buf[i/8] = buf[i/8] | (0x1 << (i%8));
00248     }
00249     printf("balloc: write bitmap block at sector %d\n", sb.bmapstart);
00250     wsect(sb.bmapstart, buf);
00251 }
00252
00253 #define min(a, b) ((a) < (b) ? (a) : (b))
00254
00255 void

```

```

00256 iappend(uint inum, void *xp, int n)
00257 {
00258     char *p = (char*)xp;
00259     uint fbn, off, nl;
00260     struct dinode din;
00261     char buf[BSIZE];
00262     uint indirect[NINDIRECT];
00263     uint x;
00264
00265     rinode(inum, &din);
00266     off = xint(din.size);
00267     // printf("append inum %d at off %d sz %d\n", inum, off, n);
00268     while(n > 0){
00269         fbn = off / BSIZE;
00270         assert(fbn < MAXFILE);
00271         if(fbn < NDIRECT){
00272             if(xint(din.addrs[fbn]) == 0){
00273                 din.addrs[fbn] = xint(freeblock++);
00274             }
00275             x = xint(din.addrs[fbn]);
00276         } else {
00277             if(xint(din.addrs[NDIRECT]) == 0){
00278                 din.addrs[NDIRECT] = xint(freeblock++);
00279             }
00280             rsect(xint(din.addrs[NDIRECT]), (char*)indirect);
00281             if(indirect[fbn - NDIRECT] == 0){
00282                 indirect[fbn - NDIRECT] = xint(freeblock++);
00283                 wsect(xint(din.addrs[NDIRECT]), (char*)indirect);
00284             }
00285             x = xint(indirect[fbn-NDIRECT]);
00286         }
00287         nl = min(n, (fbn + 1) * BSIZE - off);
00288         rsect(x, buf);
00289         bcopy(p, buf + off - (fbn * BSIZE), nl);
00290         wsect(x, buf);
00291         n -= nl;
00292         off += nl;
00293         p += nl;
00294     }
00295     din.size = xint(off);
00296     winode(inum, &din);
00297 }

```

5.127 mmu.h File Reference

Classes

- struct [gatedesc](#)
- struct [segdesc](#)
- struct [taskstate](#)

Macros

- #define [CRO_PE](#) 0x00000001
- #define [CRO_PG](#) 0x80000000
- #define [CRO_WP](#) 0x00010000
- #define [CR4_PSE](#) 0x00000010
- #define [DPL_USER](#) 0x3
- #define [FL_IF](#) 0x00000200
- #define [NPENTRIES](#) 1024
- #define [NPTENTRIES](#) 1024
- #define [NSEGS](#) 6
- #define [PDX\(va\)](#) (((uint)(va) >> [PDXSHIFT](#)) & 0x3FF)
- #define [PDXSHIFT](#) 22
- #define [PGADDR\(d, t, o\)](#) (((uint)((d) << [PDXSHIFT](#) | (t) << [PTXSHIFT](#) | (o)))
- #define [PGROUNDDOWN\(a\)](#) (((a) & ~(PGSIZE-1))
- #define [PGROUNDUP\(sz\)](#) (((sz)+PGSIZE-1) & ~(PGSIZE-1))

- `#define PGSIZE 4096`
- `#define PTE_ADDR(pte) (((uint)(pte) & ~0xFFF)`
- `#define PTE_FLAGS(pte) (((uint)(pte) & 0xFFF)`
- `#define PTE_P 0x001`
- `#define PTE_PS 0x080`
- `#define PTE_U 0x004`
- `#define PTE_W 0x002`
- `#define PTX(va) (((uint)(va) >> PTXSHIFT) & 0x3FF)`
- `#define PTXSHIFT 12`
- `#define SEG(type, base, lim, dpl)`
- `#define SEG16(type, base, lim, dpl)`
- `#define SEG_KCODE 1`
- `#define SEG_KDATA 2`
- `#define SEG_TSS 5`
- `#define SEG_UCODE 3`
- `#define SEG_UDATA 4`
- `#define SETGATE(gate, istrap, sel, off, d)`
- `#define STA_R 0x2`
- `#define STA_W 0x2`
- `#define STA_X 0x8`
- `#define STS_IG32 0xE`
- `#define STS_T32A 0x9`
- `#define STS_TG32 0xF`

Typedefs

- `typedef uint pte_t`

5.127.1 Macro Definition Documentation

5.127.1.1 CR0_PE

```
#define CR0_PE 0x00000001
```

Definition at line 8 of file [mmu.h](#).

5.127.1.2 CR0_PG

```
#define CR0_PG 0x80000000
```

Definition at line 10 of file [mmu.h](#).

5.127.1.3 CR0_WP

```
#define CR0_WP 0x00010000
```

Definition at line 9 of file [mmu.h](#).

5.127.1.4 CR4_PSE

```
#define CR4_PSE 0x00000010
```

Definition at line 12 of file [mmu.h](#).

5.127.1.5 DPL_USER

```
#define DPL_USER 0x3
```

Definition at line 53 of file [mmu.h](#).

5.127.1.6 FL_IF

```
#define FL_IF 0x00000200
```

Definition at line 5 of file [mmu.h](#).

5.127.1.7 NPENTRIES

```
#define NPENTRIES 1024
```

Definition at line 83 of file [mmu.h](#).

5.127.1.8 NPTENTRIES

```
#define NPTENTRIES 1024
```

Definition at line 84 of file [mmu.h](#).

5.127.1.9 NSEGS

```
#define NSEGS 6
```

Definition at line 22 of file [mmu.h](#).

5.127.1.10 PDX

```
#define PDX(  
    va ) ((uint)(va) >> PDXSHIFT) & 0x3FF)
```

Definition at line 74 of file [mmu.h](#).

5.127.1.11 PDXSHIFT

```
#define PDXSHIFT 22
```

Definition at line 88 of file [mmu.h](#).

5.127.1.12 PGADDR

```
#define PGADDR(  
    d,  
    t,  
    o ) ((uint)((d) << PDXSHIFT | (t) << PTXSHIFT | (o)))
```

Definition at line 80 of file [mmu.h](#).

5.127.1.13 PGROUNDDOWN

```
#define PGROUNDDOWN(  
    a ) (((a)) & ~(PGSIZE-1))
```

Definition at line 91 of file [mmu.h](#).

5.127.1.14 PGROUNDUP

```
#define PGROUNDUP(  
    sz ) (((sz)+PGSIZE-1) & ~(PGSIZE-1))
```

Definition at line 90 of file [mmu.h](#).

5.127.1.15 PGSIZE

```
#define PGSIZE 4096
```

Definition at line 85 of file [mmu.h](#).

5.127.1.16 PTE_ADDR

```
#define PTE_ADDR(  
    pte ) ((uint) (pte) & ~0xFFF)
```

Definition at line 100 of file [mmu.h](#).

5.127.1.17 PTE_FLAGS

```
#define PTE_FLAGS(  
    pte ) ((uint) (pte) & 0xFFF)
```

Definition at line 101 of file [mmu.h](#).

5.127.1.18 PTE_P

```
#define PTE_P 0x001
```

Definition at line 94 of file [mmu.h](#).

5.127.1.19 PTE_PS

```
#define PTE_PS 0x080
```

Definition at line 97 of file [mmu.h](#).

5.127.1.20 PTE_U

```
#define PTE_U 0x004
```

Definition at line 96 of file [mmu.h](#).

5.127.1.21 PTE_W

```
#define PTE_W 0x002
```

Definition at line 95 of file [mmu.h](#).

5.127.1.22 PTX

```
#define PTX(  
    va ) ((uint)(va) >> PTXSHIFT) & 0x3FF)
```

Definition at line 77 of file [mmu.h](#).

5.127.1.23 PTXSHIFT

```
#define PTXSHIFT 12
```

Definition at line 87 of file [mmu.h](#).

5.127.1.24 SEG

```
#define SEG(  
    type,  
    base,  
    lim,  
    dpl )
```

Value:

```
(struct segdesc) \
{ ((lim) >> 12) & 0xffff, (uint)(base) & 0xffff, \
  ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
  (uint)(lim) >> 28, 0, 0, 1, 1, (uint)(base) >> 24 }
```

Definition at line 43 of file [mmu.h](#).

5.127.1.25 SEG16

```
#define SEG16(  
    type,  
    base,  
    lim,  
    dpl )
```

Value:

```
(struct segdesc) \
{ (lim) & 0xffff, (uint)(base) & 0xffff, \
  ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
  (uint)(lim) >> 16, 0, 0, 1, 0, (uint)(base) >> 24 }
```

Definition at line 47 of file [mmu.h](#).

5.127.1.26 SEG_KCODE

```
#define SEG_KCODE 1
```

Definition at line 15 of file [mmu.h](#).

5.127.1.27 SEG_KDATA

```
#define SEG_KDATA 2
```

Definition at line 16 of file [mmu.h](#).

5.127.1.28 SEG_TSS

```
#define SEG_TSS 5
```

Definition at line 19 of file [mmu.h](#).

5.127.1.29 SEG_UCODE

```
#define SEG_UCODE 3
```

Definition at line 17 of file [mmu.h](#).

5.127.1.30 SEG_UDATA

```
#define SEG_UDATA 4
```

Definition at line 18 of file [mmu.h](#).

5.127.1.31 SETGATE

```
#define SETGATE(
    gate,
    istrap,
    sel,
    off,
    d )
```

Value:

```
{
    (gate).off_15_0 = (uint)(off) & 0xffff;
    (gate).cs = (sel);
    (gate).args = 0;
    (gate).rsv1 = 0;
    (gate).type = (istrap) ? STS_TG32 : STS_IG32;
    (gate).s = 0;
    (gate).dpl = (d);
    (gate).p = 1;
    (gate).off_31_16 = (uint)(off) >> 16;
}
```

Definition at line 168 of file [mmu.h](#).

5.127.1.32 STA_R

```
#define STA_R 0x2
```

Definition at line 58 of file [mmu.h](#).

5.127.1.33 STA_W

```
#define STA_W 0x2
```

Definition at line 57 of file [mmu.h](#).

5.127.1.34 STA_X

```
#define STA_X 0x8
```

Definition at line 56 of file [mmu.h](#).

5.127.1.35 STS_IG32

```
#define STS_IG32 0xE
```

Definition at line 62 of file [mmu.h](#).

5.127.1.36 STS_T32A

```
#define STS_T32A 0x9
```

Definition at line 61 of file [mmu.h](#).

5.127.1.37 STS_TG32

```
#define STS_TG32 0xF
```

Definition at line 63 of file [mmu.h](#).

5.127.2 Typedef Documentation**5.127.2.1 pte_t**

```
typedef uint pte_t
```

Definition at line 104 of file [mmu.h](#).

5.128 mmu.h

[Go to the documentation of this file.](#)

```
00001 // This file contains definitions for the
00002 // x86 memory management unit (MMU).
00003
00004 // Eflags register
00005 #define FL_IF          0x00000200    // Interrupt Enable
00006
00007 // Control Register flags
00008 #define CR0_PE         0x00000001    // Protection Enable
00009 #define CR0_WP         0x00010000    // Write Protect
00010 #define CR0_PG         0x80000000    // Paging
00011
00012 #define CR4_PSE        0x00000010    // Page size extension
00013
00014 // various segment selectors.
00015 #define SEG_KCODE 1    // kernel code
00016 #define SEG_KDATA 2    // kernel data+stack
00017 #define SEG_UCODE 3    // user code
00018 #define SEG_UDATA 4    // user data+stack
00019 #define SEG_TSS   5    // this process's task state
00020
00021 // cpu->gdt[NSEGS] holds the above segments.
00022 #define NSEGS      6
00023
00024 #ifndef __ASSEMBLER__
00025 // Segment Descriptor
00026 struct segdesc {
00027     uint lim_15_0 : 16; // Low bits of segment limit
00028     uint base_15_0 : 16; // Low bits of segment base address
00029     uint base_23_16 : 8; // Middle bits of segment base address
00030     uint type : 4;       // Segment type (see STS_ constants)
00031     uint s : 1;         // 0 = system, 1 = application
00032     uint dpl : 2;       // Descriptor Privilege Level
00033     uint p : 1;         // Present
00034     uint lim_19_16 : 4; // High bits of segment limit
```

```

00035  uint avl : 1;          // Unused (available for software use)
00036  uint rsv1 : 1;         // Reserved
00037  uint db : 1;           // 0 = 16-bit segment, 1 = 32-bit segment
00038  uint g : 1;            // Granularity: limit scaled by 4K when set
00039  uint base_31_24 : 8;   // High bits of segment base address
00040 };
00041
00042 // Normal segment
00043 #define SEG(type, base, lim, dpl) (struct segdesc) \
00044 { ((lim) >> 12) & 0xffff, (uint)(base) & 0xffff, \
00045   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
00046   (uint)(lim) >> 28, 0, 0, 1, 1, (uint)(base) >> 24 }
00047 #define SEG16(type, base, lim, dpl) (struct segdesc) \
00048 { (lim) & 0xffff, (uint)(base) & 0xffff, \
00049   ((uint)(base) >> 16) & 0xff, type, 1, dpl, 1, \
00050   (uint)(lim) >> 16, 0, 0, 1, 0, (uint)(base) >> 24 }
00051 #endif
00052
00053 #define DPL_USER 0x3      // User DPL
00054
00055 // Application segment type bits
00056 #define STA_X 0x8         // Executable segment
00057 #define STA_W 0x2         // Writeable (non-executable segments)
00058 #define STA_R 0x2         // Readable (executable segments)
00059
00060 // System segment type bits
00061 #define STS_T32A 0x9      // Available 32-bit TSS
00062 #define STS_IG32 0xE      // 32-bit Interrupt Gate
00063 #define STS_TG32 0xF      // 32-bit Trap Gate
00064
00065 // A virtual address 'la' has a three-part structure as follows:
00066 //
00067 // +-----10-----+-----10-----+-----12-----+
00068 // | Page Directory | Page Table   | Offset within Page |
00069 // |      Index      |      Index   |                   |
00070 // +-----+-----+-----+
00071 // \--- PDX(va) ---/ \--- PTX(va) ---/
00072
00073 // page directory index
00074 #define PDX(va)          (((uint)(va) >> PDXSHIFT) & 0x3FF)
00075
00076 // page table index
00077 #define PTX(va)          (((uint)(va) >> PTXSHIFT) & 0x3FF)
00078
00079 // construct virtual address from indexes and offset
00080 #define PGADDR(d, t, o) ((uint)((d) << PDXSHIFT | (t) << PTXSHIFT | (o)))
00081
00082 // Page directory and page table constants.
00083 #define NPENTRIES 1024    // # directory entries per page directory
00084 #define NPTEENTRIES 1024 // # PTEs per page table
00085 #define PGSIZE 4096      // bytes mapped by a page
00086
00087 #define PTXSHIFT 12       // offset of PTX in a linear address
00088 #define PDXSHIFT 22       // offset of PDX in a linear address
00089
00090 #define PGROUNDUP(sz) (((sz)+PGSIZE-1) & ~(PGSIZE-1))
00091 #define PGROUNDDOWN(a) (((a)) & ~(PGSIZE-1))
00092
00093 // Page table/directory entry flags.
00094 #define PTE_P 0x001      // Present
00095 #define PTE_W 0x002      // Writeable
00096 #define PTE_U 0x004      // User
00097 #define PTE_PS 0x080     // Page Size
00098
00099 // Address in page table or page directory entry
00100 #define PTE_ADDR(pte) ((uint)(pte) & ~0xFFF)
00101 #define PTE_FLAGS(pte) ((uint)(pte) & 0xFFF)
00102
00103 #ifndef __ASSEMBLER__
00104 typedef uint pte_t;
00105
00106 // Task state segment format
00107 struct taskstate {
00108   uint link;          // Old ts selector
00109   uint esp0;          // Stack pointers and segment selectors
00110   ushort ss0;         // after an increase in privilege level
00111   ushort padding1;
00112   uint *esp1;
00113   ushort ss1;
00114   ushort padding2;
00115   uint *esp2;
00116   ushort ss2;
00117   ushort padding3;
00118   void *cr3;          // Page directory base
00119   uint *eip;          // Saved state from last task switch
00120   uint eflags;
00121   uint eax;           // More saved state (registers)

```



```

00122  uint ecx;
00123  uint edx;
00124  uint ebx;
00125  uint *esp;
00126  uint *ebp;
00127  uint esi;
00128  uint edi;
00129  ushort es;          // Even more saved state (segment selectors)
00130  ushort padding4;
00131  ushort cs;
00132  ushort padding5;
00133  ushort ss;
00134  ushort padding6;
00135  ushort ds;
00136  ushort padding7;
00137  ushort fs;
00138  ushort padding8;
00139  ushort gs;
00140  ushort padding9;
00141  ushort ldt;
00142  ushort padding10;
00143  ushort t;           // Trap on task switch
00144  ushort iomb;        // I/O map base address
00145 };
00146
00147 // Gate descriptors for interrupts and traps
00148 struct gatedesc {
00149     uint off_15_0 : 16;  // low 16 bits of offset in segment
00150     uint cs : 16;         // code segment selector
00151     uint args : 5;       // # args, 0 for interrupt/trap gates
00152     uint rsv1 : 3;       // reserved(should be zero I guess)
00153     uint type : 4;       // type(STS_{IG32,TG32})
00154     uint s : 1;         // must be 0 (system)
00155     uint dpl : 2;       // descriptor(meaning new) privilege level
00156     uint p : 1;         // Present
00157     uint off_31_16 : 16; // high bits of offset in segment
00158 };
00159
00160 // Set up a normal interrupt/trap gate descriptor.
00161 // - istrap: 1 for a trap (= exception) gate, 0 for an interrupt gate.
00162 // - sel: Code segment selector for interrupt/trap handler
00163 // - off: Offset in code segment for interrupt/trap handler
00164 // - dpl: Descriptor Privilege Level -
00165 //       the privilege level required for software to invoke
00166 //       this interrupt/trap gate explicitly using an int instruction.
00167 #define SETGATE(gate, istrap, sel, off, d) \
00168 { \
00169     (gate).off_15_0 = (uint)(off) & 0xffff; \
00170     (gate).cs = (sel); \
00171     (gate).args = 0; \
00172     (gate).rsv1 = 0; \
00173     (gate).type = (istrap) ? STS_TG32 : STS_IG32; \
00174     (gate).s = 0; \
00175     (gate).dpl = (d); \
00176     (gate).p = 1; \
00177     (gate).off_31_16 = (uint)(off) >> 16; \
00178 } \
00179
00180
00181 #endif

```

5.129 mp.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mp.h"
#include "x86.h"
#include "mmu.h"
#include "proc.h"

```

Functions

- static struct `mpconf` * `mpconfig` (struct `mp` **`pmp`)

- void [mpinit](#) (void)
- static struct [mp](#) * [mpsearch](#) (void)
- static struct [mp](#) * [mpsearch1](#) (uint a, int len)
- static [uchar](#) [sum](#) ([uchar](#) *addr, int len)

Variables

- struct [cpu](#) [cpus](#) [[NCPU](#)]
- [uchar](#) [ioapicid](#)
- int [ncpu](#)

5.129.1 Function Documentation

5.129.1.1 mpconfig()

```
static struct mpconf * mpconfig (
    struct mp ** pmp ) [static]
```

Definition at line 73 of file [mp.c](#).

```
00074 {
00075     struct mpconf *conf;
00076     struct mp *mp;
00077
00078     if((mp = mpsearch()) == 0 || mp->physaddr == 0)
00079         return 0;
00080     conf = (struct mpconf*) P2V((uint) mp->physaddr);
00081     if(memcmp(conf, "PCMP", 4) != 0)
00082         return 0;
00083     if(conf->version != 1 && conf->version != 4)
00084         return 0;
00085     if(sum((uchar*)conf, conf->length) != 0)
00086         return 0;
00087     *pmp = mp;
00088     return conf;
00089 }
```

Referenced by [mpinit](#)).

5.129.1.2 mpinit()

```
void mpinit (
    void )
```

Definition at line 92 of file [mp.c](#).

```
00093 {
00094     uchar *p, *e;
00095     int ismp;
00096     struct mp *mp;
00097     struct mpconf *conf;
00098     struct mproc *proc;
00099     struct mpioapic *ioapic;
00100
00101     if((conf = mpconfig(&mp)) == 0)
00102         panic("Expect to run on an SMP");
00103     ismp = 1;
00104     lapic = (uint*)conf->lapicaddr;
00105     for(p=(uchar*) (conf+1), e=(uchar*) conf+conf->length; p<e; ){
```

```

00106     switch(*p){
00107     case MPPROC:
00108         proc = (struct mpproc*)p;
00109         if(ncpu < NCPU) {
00110             cpus[ncpu].apicid = proc->apicid; // apicid may differ from ncpu
00111             ncpu++;
00112         }
00113         p += sizeof(struct mpproc);
00114         continue;
00115     case MPIOAPIC:
00116         ioapic = (struct mpioapic*)p;
00117         ioapicid = ioapic->apicno;
00118         p += sizeof(struct mpioapic);
00119         continue;
00120     case MPBUS:
00121     case MPIOINTR:
00122     case MPLINTR:
00123         p += 8;
00124         continue;
00125     default:
00126         ismp = 0;
00127         break;
00128     }
00129 }
00130 if(!ismp)
00131     panic("Didn't find a suitable machine");
00132
00133 if(mp->imcrp) {
00134     // Bochs doesn't support IMCR, so this doesn't run on Bochs.
00135     // But it would on real hardware.
00136     outb(0x22, 0x70); // Select IMCR
00137     outb(0x23, inb(0x23) | 1); // Mask external interrupts.
00138 }
00139 }

```

5.129.1.3 mpsearch()

```

static struct mp * mpsearch (
    void ) [static]

```

Definition at line 49 of file [mp.c](#).

```

00050 {
00051     uchar *bda;
00052     uint p;
00053     struct mp *mp;
00054
00055     bda = (uchar *) P2V(0x400);
00056     if((p = ((bda[0x0F] << 8) | bda[0x0E]) << 4)){
00057         if((mp = mpsearch1(p, 1024))
00058             return mp;
00059     } else {
00060         p = ((bda[0x14] << 8) | bda[0x13]) * 1024;
00061         if((mp = mpsearch1(p-1024, 1024))
00062             return mp;
00063     }
00064     return mpsearch1(0xF0000, 0x10000);
00065 }

```

Referenced by [mpconfig\(\)](#).

5.129.1.4 mpsearch1()

```

static struct mp * mpsearch1 (
    uint a,
    int len ) [static]

```

Definition at line 31 of file [mp.c](#).

```

00032 {

```

```

00033     uchar *e, *p, *addr;
00034
00035     addr = P2V(a);
00036     e = addr+len;
00037     for(p = addr; p < e; p += sizeof(struct mp))
00038         if(memcmp(p, "_MP_", 4) == 0 && sum(p, sizeof(struct mp)) == 0)
00039             return (struct mp*)p;
00040     return 0;
00041 }

```

Referenced by [mpsearch\(\)](#).

5.129.1.5 sum()

```

static uchar sum (
    uchar * addr,
    int len ) [static]

```

Definition at line 19 of file [mp.c](#).

```

00020 {
00021     int i, sum;
00022
00023     sum = 0;
00024     for(i=0; i<len; i++)
00025         sum += addr[i];
00026     return sum;
00027 }

```

Referenced by [mpconfig\(\)](#), [mpsearch1\(\)](#), and [sum\(\)](#).

5.129.2 Variable Documentation

5.129.2.1 cpus

```
struct cpu cpus[NCPU]
```

Definition at line 14 of file [mp.c](#).

Referenced by [cpuid\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), [seginit\(\)](#), and [startothers\(\)](#).

5.129.2.2 ioapicid

```
uchar ioapicid
```

Definition at line 16 of file [mp.c](#).

Referenced by [ioapicinit\(\)](#), and [mpinit\(\)](#).

5.129.2.3 ncpu

```
int ncpu
```

Definition at line 15 of file [mp.c](#).

Referenced by [ideinit\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), and [startothers\(\)](#).

5.130 mp.c

[Go to the documentation of this file.](#)

```
00001 // Multiprocessor support
00002 // Search memory for MP description structures.
00003 // http://developer.intel.com/design/pentium/datashts/24201606.pdf
00004
00005 #include "types.h"
00006 #include "defs.h"
00007 #include "param.h"
00008 #include "memlayout.h"
00009 #include "mp.h"
00010 #include "x86.h"
00011 #include "mmu.h"
00012 #include "proc.h"
00013
00014 struct cpu cpus[NCPU];
00015 int ncpu;
00016 uchar ioapicid;
00017
00018 static uchar
00019 sum(uchar *addr, int len)
00020 {
00021     int i, sum;
00022
00023     sum = 0;
00024     for(i=0; i<len; i++)
00025         sum += addr[i];
00026     return sum;
00027 }
00028
00029 // Look for an MP structure in the len bytes at addr.
00030 static struct mp*
00031 mpsearch1(uint a, int len)
00032 {
00033     uchar *e, *p, *addr;
00034
00035     addr = P2V(a);
00036     e = addr+len;
00037     for(p = addr; p < e; p += sizeof(struct mp))
00038         if(memcmp(p, "_MP_", 4) == 0 && sum(p, sizeof(struct mp)) == 0)
00039             return (struct mp*)p;
00040     return 0;
00041 }
00042
00043 // Search for the MP Floating Pointer Structure, which according to the
00044 // spec is in one of the following three locations:
00045 // 1) in the first KB of the EBDA;
00046 // 2) in the last KB of system base memory;
00047 // 3) in the BIOS ROM between 0xE0000 and 0xFFFFF.
00048 static struct mp*
00049 mpsearch(void)
00050 {
00051     uchar *bda;
00052     uint p;
00053     struct mp *mp;
00054
00055     bda = (uchar *) P2V(0x400);
00056     if((p = ((bda[0x0F]«8) | bda[0x0E]) « 4)){
00057         if((mp = mpsearch1(p, 1024)))
00058             return mp;
00059     } else {
00060         p = ((bda[0x14]«8) | bda[0x13]) * 1024;
00061         if((mp = mpsearch1(p-1024, 1024)))
00062             return mp;
00063     }
00064     return mpsearch1(0xF0000, 0x10000);
00065 }
00066
```

```

00067 // Search for an MP configuration table. For now,
00068 // don't accept the default configurations (physaddr == 0).
00069 // Check for correct signature, calculate the checksum and,
00070 // if correct, check the version.
00071 // To do: check extended table checksum.
00072 static struct mpconf*
00073 mpconfig(struct mp **pmp)
00074 {
00075     struct mpconf *conf;
00076     struct mp *mp;
00077
00078     if((mp = mpsearch()) == 0 || mp->physaddr == 0)
00079         return 0;
00080     conf = (struct mpconf*) P2V((uint) mp->physaddr);
00081     if(memcmp(conf, "PCMP", 4) != 0)
00082         return 0;
00083     if(conf->version != 1 && conf->version != 4)
00084         return 0;
00085     if(sum((uchar*)conf, conf->length) != 0)
00086         return 0;
00087     *pmp = mp;
00088     return conf;
00089 }
00090
00091 void
00092 mpinit(void)
00093 {
00094     uchar *p, *e;
00095     int ismp;
00096     struct mp *mp;
00097     struct mpconf *conf;
00098     struct mpproc *proc;
00099     struct mpioapic *ioapic;
00100
00101     if((conf = mpconfig(&mp)) == 0)
00102         panic("Expect to run on an SMP");
00103     ismp = 1;
00104     lapic = (uint*)conf->lapicaddr;
00105     for(p=(uchar*)(conf+1), e=(uchar*)conf+conf->length; p<e; ){
00106         switch(*p){
00107             case MPPROC:
00108                 proc = (struct mpproc*)p;
00109                 if(ncpu < NCPU) {
00110                     cpus[ncpu].apicid = proc->apicid; // apicid may differ from ncpu
00111                     ncpu++;
00112                 }
00113                 p += sizeof(struct mpproc);
00114                 continue;
00115             case MPIOAPIC:
00116                 ioapic = (struct mpioapic*)p;
00117                 ioapicid = ioapic->apicno;
00118                 p += sizeof(struct mpioapic);
00119                 continue;
00120             case MPBUS:
00121             case MPIOINTR:
00122             case MPLINTR:
00123                 p += 8;
00124                 continue;
00125             default:
00126                 ismp = 0;
00127                 break;
00128         }
00129     }
00130     if(!ismp)
00131         panic("Didn't find a suitable machine");
00132
00133     if(mp->imcrp){
00134         // Bochs doesn't support IMCR, so this doesn't run on Bochs.
00135         // But it would on real hardware.
00136         outb(0x22, 0x70); // Select IMCR
00137         outb(0x23, inb(0x23) | 1); // Mask external interrupts.
00138     }
00139 }

```

5.131 mp.d File Reference

5.132 mp.d

[Go to the documentation of this file.](#)

```

00001 mp.o: mp.c /usr/include/stdc-predef.h types.h defs.h param.h memlayout.h \
00002 mp.h x86.h mmu.h proc.h

```

5.133 mp.h File Reference

Classes

- struct [mp](#)
- struct [mpconf](#)
- struct [mpioapic](#)
- struct [mpproc](#)

Macros

- #define [MPBOOT](#) 0x02
- #define [MPBUS](#) 0x01
- #define [MPIOAPIC](#) 0x02
- #define [MPIINTR](#) 0x03
- #define [MPLINTR](#) 0x04
- #define [MPPROC](#) 0x00

5.133.1 Macro Definition Documentation

5.133.1.1 MPBOOT

```
#define MPBOOT 0x02
```

Definition at line [34](#) of file [mp.h](#).

5.133.1.2 MPBUS

```
#define MPBUS 0x01
```

Definition at line [50](#) of file [mp.h](#).

5.133.1.3 MPIOAPIC

```
#define MPIOAPIC 0x02
```

Definition at line [51](#) of file [mp.h](#).

5.133.1.4 MPIINTR

```
#define MPIINTR 0x03
```

Definition at line 52 of file [mp.h](#).

5.133.1.5 MPLINTR

```
#define MPLINTR 0x04
```

Definition at line 53 of file [mp.h](#).

5.133.1.6 MPPROC

```
#define MPPROC 0x00
```

Definition at line 49 of file [mp.h](#).

5.134 mp.h

[Go to the documentation of this file.](#)

```
00001 // See MultiProcessor Specification Version 1.[14]
00002
00003 struct mp {                // floating pointer
00004     uchar signature[4];     // "_MP_"
00005     void *physaddr;         // phys addr of MP config table
00006     uchar length;          // 1
00007     uchar specrev;         // [14]
00008     uchar checksum;        // all bytes must add up to 0
00009     uchar type;            // MP system config type
00010     uchar imcrp;
00011     uchar reserved[3];
00012 };
00013
00014 struct mpconf {            // configuration table header
00015     uchar signature[4];     // "PCMP"
00016     ushort length;         // total table length
00017     uchar version;         // [14]
00018     uchar checksum;        // all bytes must add up to 0
00019     uchar product[20];     // product id
00020     uint *oemtable;        // OEM table pointer
00021     ushort oemlength;      // OEM table length
00022     ushort entry;          // entry count
00023     uint *lapicaddr;       // address of local APIC
00024     ushort xlength;        // extended table length
00025     uchar xchecksum;       // extended table checksum
00026     uchar reserved;
00027 };
00028
00029 struct mpproc {            // processor table entry
00030     uchar type;            // entry type (0)
00031     uchar apicid;          // local APIC id
00032     uchar version;         // local APIC verison
00033     uchar flags;           // CPU flags
00034     #define MPBOOT 0x02    // This proc is the bootstrap processor.
00035     uchar signature[4];    // CPU signature
00036     uint feature;          // feature flags from CPUID instruction
00037     uchar reserved[8];
00038 };
00039
00040 struct mpioapic {         // I/O APIC table entry
```



```

00041  uchar type;                // entry type (2)
00042  uchar apicno;              // I/O APIC id
00043  uchar version;            // I/O APIC version
00044  uchar flags;              // I/O APIC flags
00045  uint *addr;               // I/O APIC address
00046 };
00047
00048 // Table entry types
00049 #define MPPROC 0x00 // One per processor
00050 #define MPBUS 0x01 // One per bus
00051 #define MPIOAPIC 0x02 // One per I/O APIC
00052 #define MPIOINTR 0x03 // One per bus interrupt source
00053 #define MPLINTR 0x04 // One per system interrupt source
00054
00055 //PAGEBREAK!
00056 // Blank page.

```

5.135 nice.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

```

Functions

- int [main](#) (int argc, char *argv[])

5.135.1 Function Documentation

5.135.1.1 main()

```

int main (
    int argc,
    char * argv[] )

```

Definition at line 7 of file [nice.c](#).

```

00008 {
00009     int priority, pid;
00010     if(argc < 3){
00011         printf(2, "Usage: nice pid priority\n");
00012         exit();
00013     }
00014     pid = atoi(argv[1]);
00015     priority = atoi(argv[2]);
00016     if (priority < 0 || priority > 20){
00017         printf(2, "Invalid priority (0-20)!\n");
00018         exit();
00019     }
00020     chpr(pid, priority);
00021     exit();
00022 }

```

5.136 nice.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "fcntl.h"
00005
00006 int
00007 main(int argc, char *argv[])
00008 {
00009     int priority, pid;
00010     if(argc < 3){
00011         printf(2, "Usage: nice pid priority\n");
00012         exit();
00013     }
00014     pid = atoi(argv[1]);
00015     priority = atoi(argv[2]);
00016     if (priority < 0 || priority > 20){
00017         printf(2, "Invalid priority (0-20)!\n");
00018         exit();
00019     }
00020     chpr(pid, priority);
00021     exit();
00022 }
```

5.137 nice.d File Reference

5.138 nice.d

[Go to the documentation of this file.](#)

```
00001 nice.o: nice.c /usr/include/stdc-predef.h types.h stat.h user.h fcntl.h
```

5.139 param.h File Reference

Macros

- #define [FSSIZE](#) 1000
- #define [KSTACKSIZE](#) 4096
- #define [LOGSIZE](#) (MAXOPBLOCKS*3)
- #define [MAXARG](#) 32
- #define [MAXOPBLOCKS](#) 10
- #define [NBUF](#) (MAXOPBLOCKS*3)
- #define [NCPU](#) 8
- #define [NDEV](#) 10
- #define [NFILE](#) 100
- #define [NINODE](#) 50
- #define [NOFILE](#) 16
- #define [NPROC](#) 64
- #define [ROOTDEV](#) 1

5.139.1 Macro Definition Documentation

5.139.1.1 FSSIZE

```
#define FSSIZE 1000
```

Definition at line 13 of file [param.h](#).

5.139.1.2 KSTACKSIZE

```
#define KSTACKSIZE 4096
```

Definition at line 2 of file [param.h](#).

5.139.1.3 LOGSIZE

```
#define LOGSIZE (MAXOPBLOCKS*3)
```

Definition at line 11 of file [param.h](#).

5.139.1.4 MAXARG

```
#define MAXARG 32
```

Definition at line 9 of file [param.h](#).

5.139.1.5 MAXOPBLOCKS

```
#define MAXOPBLOCKS 10
```

Definition at line 10 of file [param.h](#).

5.139.1.6 NBUF

```
#define NBUF (MAXOPBLOCKS*3)
```

Definition at line 12 of file [param.h](#).

5.139.1.7 NCPU

```
#define NCPU 8
```

Definition at line 3 of file [param.h](#).

5.139.1.8 NDEV

```
#define NDEV 10
```

Definition at line 7 of file [param.h](#).

5.139.1.9 NFILE

```
#define NFILE 100
```

Definition at line 5 of file [param.h](#).

5.139.1.10 NINODE

```
#define NINODE 50
```

Definition at line 6 of file [param.h](#).

5.139.1.11 NOFILE

```
#define NOFILE 16
```

Definition at line 4 of file [param.h](#).

5.139.1.12 NPROC

```
#define NPROC 64
```

Definition at line 1 of file [param.h](#).

5.139.1.13 ROOTDEV

```
#define ROOTDEV 1
```

Definition at line 8 of file [param.h](#).

5.140 param.h

[Go to the documentation of this file.](#)

```
00001 #define NPROC      64 // maximum number of processes
00002 #define KSTACKSIZE 4096 // size of per-process kernel stack
00003 #define NCPU        8 // maximum number of CPUs
00004 #define NOFILE      16 // open files per process
00005 #define NFILE       100 // open files per system
00006 #define NINODE       50 // maximum number of active i-nodes
00007 #define NDEV        10 // maximum major device number
00008 #define ROOTDEV      1 // device number of file system root disk
00009 #define MAXARG       32 // max exec arguments
00010 #define MAXOPBLOCKS  10 // max # of blocks any FS op writes
00011 #define LOGSIZE      (MAXOPBLOCKS*3) // max data blocks in on-disk log
00012 #define NBUF         (MAXOPBLOCKS*3) // size of disk block cache
00013 #define FSSIZE       1000 // size of file system in blocks
00014
```

5.141 picirq.c File Reference

```
#include "types.h"
#include "x86.h"
#include "traps.h"
```

Macros

- `#define IO_PIC1 0x20`
- `#define IO_PIC2 0xA0`

Functions

- void [picinit](#) (void)

5.141.1 Macro Definition Documentation

5.141.1.1 IO_PIC1

```
#define IO_PIC1 0x20
```

Definition at line 6 of file [picirq.c](#).

5.141.1.2 IO_PIC2

```
#define IO_PIC2 0xA0
```

Definition at line 7 of file [picirq.c](#).

5.141.2 Function Documentation

5.141.2.1 picinit()

```
void picinit (  
    void )
```

Definition at line 11 of file [picirq.c](#).

```
00012 {  
00013     // mask all interrupts  
00014     outb(IO_PIC1+1, 0xFF);  
00015     outb(IO_PIC2+1, 0xFF);  
00016 }
```

5.142 picirq.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"  
00002 #include "x86.h"  
00003 #include "traps.h"  
00004  
00005 // I/O Addresses of the two programmable interrupt controllers  
00006 #define IO_PIC1          0x20    // Master (IRQs 0-7)  
00007 #define IO_PIC2          0xA0    // Slave (IRQs 8-15)  
00008  
00009 // Don't use the 8259A interrupt controllers.  Xv6 assumes SMP hardware.  
00010 void  
00011 picinit(void)  
00012 {  
00013     // mask all interrupts  
00014     outb(IO_PIC1+1, 0xFF);  
00015     outb(IO_PIC2+1, 0xFF);  
00016 }  
00017  
00018 //PAGEBREAK!  
00019 // Blank page.
```

5.143 picirq.d File Reference

5.144 picirq.d

[Go to the documentation of this file.](#)

```
00001 picirq.o: picirq.c /usr/include/stdc-predef.h types.h x86.h traps.h
```

5.145 pipe.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "mmu.h"
#include "proc.h"
#include "fs.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "file.h"
```

Classes

- struct [pipe](#)

Macros

- `#define` [PIPESIZE](#) 512

Functions

- int [pipealloc](#) (struct [file](#) **f0, struct [file](#) **f1)
- void [pipeclose](#) (struct [pipe](#) *p, int writable)
- int [piperead](#) (struct [pipe](#) *p, char *addr, int n)
- int [pipewrite](#) (struct [pipe](#) *p, char *addr, int n)

5.145.1 Macro Definition Documentation

5.145.1.1 PIPESIZE

```
#define PIPESIZE 512
```

Definition at line 11 of file [pipe.c](#).

5.145.2 Function Documentation

5.145.2.1 pipealloc()

```
int pipealloc (
    struct file ** f0,
    struct file ** f1 )
```

Definition at line 23 of file [pipe.c](#).

```
00024 {
00025     struct pipe *p;
00026
00027     p = 0;
00028     *f0 = *f1 = 0;
00029     if ((*f0 = filealloc()) == 0 || (*f1 = filealloc()) == 0)
00030         goto bad;
00031     if ((p = (struct pipe*)kalloc()) == 0)
00032         goto bad;
00033     p->readopen = 1;
00034     p->writeopen = 1;
00035     p->nwrite = 0;
00036     p->nread = 0;
00037     initlock(&p->lock, "pipe");
00038     (*f0)->type = FD_PIPE;
00039     (*f0)->readable = 1;
00040     (*f0)->writable = 0;
00041     (*f0)->pipe = p;
00042     (*f1)->type = FD_PIPE;
00043     (*f1)->readable = 0;
00044     (*f1)->writable = 1;
00045     (*f1)->pipe = p;
00046     return 0;
00047
00048 //PAGEBREAK: 20
00049 bad:
00050     if (p)
00051         kfree((char*)p);
00052     if (*f0)
00053         fileclose(*f0);
00054     if (*f1)
00055         fileclose(*f1);
00056     return -1;
00057 }
```

Referenced by [sys_pipe\(\)](#).

5.145.2.2 pipeclose()

```
void pipeclose (
    struct pipe * p,
    int writable )
```

Definition at line 60 of file [pipe.c](#).

```
00061 {
00062     acquire(&p->lock);
00063     if (writable) {
00064         p->writeopen = 0;
00065         wakeup(&p->nread);
00066     } else {
00067         p->readopen = 0;
00068         wakeup(&p->nwrite);
00069     }
00070     if (p->readopen == 0 && p->writeopen == 0) {
00071         release(&p->lock);
00072         kfree((char*)p);
00073     } else
00074         release(&p->lock);
00075 }
```

Referenced by [fileclose\(\)](#).

5.145.2.3 `piperead()`

```
int piperead (
    struct pipe * p,
    char * addr,
    int n )
```

Definition at line 101 of file `pipe.c`.

```
00102 {
00103     int i;
00104
00105     acquire(&p->lock);
00106     while(p->nread == p->nwrite && p->writeopen){ //DOC: pipe-empty
00107         if(myproc()->killed){
00108             release(&p->lock);
00109             return -1;
00110         }
00111         sleep(&p->nread, &p->lock); //DOC: piperead-sleep
00112     }
00113     for(i = 0; i < n; i++){ //DOC: piperead-copy
00114         if(p->nread == p->nwrite)
00115             break;
00116         addr[i] = p->data[p->nread++ % PIPESIZE];
00117     }
00118     wakeup(&p->nwrite); //DOC: piperead-wakeup
00119     release(&p->lock);
00120     return i;
00121 }
```

Referenced by `fileread()`.

5.145.2.4 `pipewrite()`

```
int pipewrite (
    struct pipe * p,
    char * addr,
    int n )
```

Definition at line 79 of file `pipe.c`.

```
00080 {
00081     int i;
00082
00083     acquire(&p->lock);
00084     for(i = 0; i < n; i++){
00085         while(p->nwrite == p->nread + PIPESIZE){ //DOC: pipewrite-full
00086             if(p->readopen == 0 || myproc()->killed){
00087                 release(&p->lock);
00088                 return -1;
00089             }
00090             wakeup(&p->nread);
00091             sleep(&p->nwrite, &p->lock); //DOC: pipewrite-sleep
00092         }
00093         p->data[p->nwrite++ % PIPESIZE] = addr[i];
00094     }
00095     wakeup(&p->nread); //DOC: pipewrite-wakeup1
00096     release(&p->lock);
00097     return n;
00098 }
```

Referenced by `filewrite()`.

5.146 pipe.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "mmu.h"
00005 #include "proc.h"
00006 #include "fs.h"
00007 #include "spinlock.h"
00008 #include "sleeplock.h"
00009 #include "file.h"
00010
00011 #define PIPESIZE 512
00012
00013 struct pipe {
00014     struct spinlock lock;
00015     char data[PIPESIZE];
00016     uint nread; // number of bytes read
00017     uint nwrite; // number of bytes written
00018     int readopen; // read fd is still open
00019     int writeopen; // write fd is still open
00020 };
00021
00022 int
00023 pipealloc(struct file **f0, struct file **f1)
00024 {
00025     struct pipe *p;
00026
00027     p = 0;
00028     *f0 = *f1 = 0;
00029     if((*f0 = filealloc()) == 0 || (*f1 = filealloc()) == 0)
00030         goto bad;
00031     if((p = (struct pipe*)kalloc()) == 0)
00032         goto bad;
00033     p->readopen = 1;
00034     p->writeopen = 1;
00035     p->nwrite = 0;
00036     p->nread = 0;
00037     initlock(&p->lock, "pipe");
00038     (*f0)->type = FD_PIPE;
00039     (*f0)->readable = 1;
00040     (*f0)->writable = 0;
00041     (*f0)->pipe = p;
00042     (*f1)->type = FD_PIPE;
00043     (*f1)->readable = 0;
00044     (*f1)->writable = 1;
00045     (*f1)->pipe = p;
00046     return 0;
00047
00048 //PAGEBREAK: 20
00049 bad:
00050     if(p)
00051         kfree((char*)p);
00052     if(*f0)
00053         fileclose(*f0);
00054     if(*f1)
00055         fileclose(*f1);
00056     return -1;
00057 }
00058
00059 void
00060 pipeclose(struct pipe *p, int writable)
00061 {
00062     acquire(&p->lock);
00063     if(writable){
00064         p->writeopen = 0;
00065         wakeup(&p->nread);
00066     } else {
00067         p->readopen = 0;
00068         wakeup(&p->nwrite);
00069     }
00070     if(p->readopen == 0 && p->writeopen == 0){
00071         release(&p->lock);
00072         kfree((char*)p);
00073     } else
00074         release(&p->lock);
00075 }
00076
00077 //PAGEBREAK: 40
00078 int
00079 pipewrite(struct pipe *p, char *addr, int n)
00080 {
00081     int i;
00082

```

```

00083     acquire(&p->lock);
00084     for(i = 0; i < n; i++){
00085         while(p->nwrite == p->nread + PIPESIZE){ //DOC: pipewrite-full
00086             if(p->readopen == 0 || myproc()->killed){
00087                 release(&p->lock);
00088                 return -1;
00089             }
00090             wakeup(&p->nread);
00091             sleep(&p->nwrite, &p->lock); //DOC: pipewrite-sleep
00092         }
00093         p->data[p->nwrite++ % PIPESIZE] = addr[i];
00094     }
00095     wakeup(&p->nread); //DOC: pipewrite-wakeup1
00096     release(&p->lock);
00097     return n;
00098 }
00099
00100 int
00101 piperead(struct pipe *p, char *addr, int n)
00102 {
00103     int i;
00104
00105     acquire(&p->lock);
00106     while(p->nread == p->nwrite && p->writeopen){ //DOC: pipe-empty
00107         if(myproc()->killed){
00108             release(&p->lock);
00109             return -1;
00110         }
00111         sleep(&p->nread, &p->lock); //DOC: piperead-sleep
00112     }
00113     for(i = 0; i < n; i++){ //DOC: piperead-copy
00114         if(p->nread == p->nwrite)
00115             break;
00116         addr[i] = p->data[p->nread++ % PIPESIZE];
00117     }
00118     wakeup(&p->nwrite); //DOC: piperead-wakeup
00119     release(&p->lock);
00120     return i;
00121 }

```

5.147 pipe.d File Reference

5.148 pipe.d

[Go to the documentation of this file.](#)

```

00001 pipe.o: pipe.c /usr/include/stdc-predef.h types.h defs.h param.h mmu.h \
00002 proc.h fs.h spinlock.h sleeplock.h file.h

```

5.149 printf.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"

```

Functions

- void [printf](#) (int fd, const char *fmt,...)
- static void [printint](#) (int fd, int xx, int base, int sgn)
- static void [putc](#) (int fd, char c)

5.149.1 Function Documentation

5.149.1.1 printf()

```
void printf (
    int fd,
    const char * fmt,
    ... )
```

Definition at line 40 of file [printf.c](#).

```
00041 {
00042     char *s;
00043     int c, i, state;
00044     uint *ap;
00045
00046     state = 0;
00047     ap = (uint*)(void*)&fmt + 1;
00048     for(i = 0; fmt[i]; i++){
00049         c = fmt[i] & 0xff;
00050         if(state == 0){
00051             if(c == '%'){
00052                 state = '%';
00053             } else {
00054                 putc(fd, c);
00055             }
00056         } else if(state == '%'){
00057             if(c == 'd'){
00058                 printint(fd, *ap, 10, 1);
00059                 ap++;
00060             } else if(c == 'x' || c == 'p'){
00061                 printint(fd, *ap, 16, 0);
00062                 ap++;
00063             } else if(c == 's'){
00064                 s = (char*)*ap;
00065                 ap++;
00066                 if(s == 0)
00067                     s = "(null)";
00068                 while(*s != 0){
00069                     putc(fd, *s);
00070                     s++;
00071                 }
00072             } else if(c == 'c'){
00073                 putc(fd, *ap);
00074                 ap++;
00075             } else if(c == '%'){
00076                 putc(fd, c);
00077             } else {
00078                 // Unknown % sequence. Print it to draw attention.
00079                 putc(fd, '%');
00080                 putc(fd, c);
00081             }
00082             state = 0;
00083         }
00084     }
00085 }
```

5.149.1.2 printint()

```
static void printint (
    int fd,
    int xx,
    int base,
    int sgn ) [static]
```

Definition at line 12 of file [printf.c](#).

```
00013 {
```

```

00014     static char digits[] = "0123456789ABCDEF";
00015     char buf[16];
00016     int i, neg;
00017     uint x;
00018
00019     neg = 0;
00020     if(sgn && xx < 0){
00021         neg = 1;
00022         x = -xx;
00023     } else {
00024         x = xx;
00025     }
00026
00027     i = 0;
00028     do{
00029         buf[i++] = digits[x % base];
00030     }while((x /= base) != 0);
00031     if(neg)
00032         buf[i++] = '-';
00033
00034     while(--i >= 0)
00035         putc(fd, buf[i]);
00036 }

```

Referenced by [printf\(\)](#).

5.149.1.3 putc()

```

static void putc (
    int fd,
    char c ) [static]

```

Definition at line 6 of file [printf.c](#).

```

00007 {
00008     write(fd, &c, 1);
00009 }

```

Referenced by [printf\(\)](#), and [printint\(\)](#).

5.150 printf.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 static void
00006 putc(int fd, char c)
00007 {
00008     write(fd, &c, 1);
00009 }
00010
00011 static void
00012 printint(int fd, int xx, int base, int sgn)
00013 {
00014     static char digits[] = "0123456789ABCDEF";
00015     char buf[16];
00016     int i, neg;
00017     uint x;
00018
00019     neg = 0;
00020     if(sgn && xx < 0){
00021         neg = 1;
00022         x = -xx;
00023     } else {
00024         x = xx;
00025     }
00026
00027     i = 0;
00028     do{

```

```

00029     buf[i++] = digits[x % base];
00030 }while((x /= base) != 0);
00031 if(neg)
00032     buf[i++] = '-';
00033
00034 while(--i >= 0)
00035     putc(fd, buf[i]);
00036 }
00037
00038 // Print to the given fd. Only understands %d, %x, %p, %s.
00039 void
00040 printf(int fd, const char *fmt, ...)
00041 {
00042     char *s;
00043     int c, i, state;
00044     uint *ap;
00045
00046     state = 0;
00047     ap = (uint*)(void*)&fmt + 1;
00048     for(i = 0; fmt[i]; i++){
00049         c = fmt[i] & 0xff;
00050         if(state == 0){
00051             if(c == '%'){
00052                 state = '%';
00053             } else {
00054                 putc(fd, c);
00055             }
00056         } else if(state == '%'){
00057             if(c == 'd'){
00058                 printint(fd, *ap, 10, 1);
00059                 ap++;
00060             } else if(c == 'x' || c == 'p'){
00061                 printint(fd, *ap, 16, 0);
00062                 ap++;
00063             } else if(c == 's'){
00064                 s = (char*)*ap;
00065                 ap++;
00066                 if(s == 0)
00067                     s = "(null)";
00068                 while(*s != 0){
00069                     putc(fd, *s);
00070                     s++;
00071                 }
00072             } else if(c == 'c'){
00073                 putc(fd, *ap);
00074                 ap++;
00075             } else if(c == '%'){
00076                 putc(fd, c);
00077             } else {
00078                 // Unknown % sequence. Print it to draw attention.
00079                 putc(fd, '%');
00080                 putc(fd, c);
00081             }
00082             state = 0;
00083         }
00084     }
00085 }

```

5.151 printf.d File Reference

5.152 printf.d

[Go to the documentation of this file.](#)

00001 printf.o: printf.c /usr/include/stdc-predef.h types.h stat.h user.h

5.153 proc.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"

```

```
#include "mmu.h"
#include "x86.h"
#include "proc.h"
#include "spinlock.h"
```

Functions

- static struct [proc](#) * [allocproc](#) (void)
- int [chpr](#) (int pid, int priority)
- int [cps](#) ()
- int [cpuid](#) ()
- void [exit](#) (void)
- int [fork](#) (void)
- void [forkret](#) (void)
- int [growproc](#) (int n)
- int [kill](#) (int pid)
- struct [cpu](#) * [mycpu](#) (void)
- struct [proc](#) * [myproc](#) (void)
- void [pinit](#) (void)
- void [procdump](#) (void)
- void [sched](#) (void)
- void [scheduler](#) (void)
- void [sleep](#) (void *chan, struct [spinlock](#) *lk)
- void [trapret](#) (void)
- void [userinit](#) (void)
- int [wait](#) (void)
- void [wakeup](#) (void *chan)
- static void [wakeup1](#) (void *chan)
- void [yield](#) (void)

Variables

- static struct [proc](#) * [initproc](#)
- int [nextpid](#) = 1
- struct {
 - struct [spinlock](#) [lock](#)
 - struct [proc](#) [proc](#) [[NPROC](#)]
- } [ptable](#)

5.153.1 Function Documentation

5.153.1.1 allocproc()

```
static struct proc * allocproc (
    void ) [static]
```

Definition at line 74 of file [proc.c](#).

```
00075 {
00076     struct proc *p;
00077     char *sp;
00078
00079     acquire(&ptable.lock);
00080
00081     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
00082         if(p->state == UNUSED)
00083             goto found;
00084
00085     release(&ptable.lock);
00086     return 0;
00087
00088 found:
00089     p->state = EMBRYO;
00090     p->pid = nextpid++;
00091
00092     release(&ptable.lock);
00093
00094     // Allocate kernel stack.
00095     if((p->kstack = kalloc()) == 0){
00096         p->state = UNUSED;
00097         return 0;
00098     }
00099     sp = p->kstack + KSTACKSIZE;
00100
00101     // Leave room for trap frame.
00102     sp -= sizeof *p->tf;
00103     p->tf = (struct trapframe*)sp;
00104
00105     // Set up new context to start executing at forkret,
00106     // which returns to trapret.
00107     sp -= 4;
00108     *(uint*)sp = (uint)trapret;
00109
00110     sp -= sizeof *p->context;
00111     p->context = (struct context*)sp;
00112     memset(p->context, 0, sizeof *p->context);
00113     p->context->eip = (uint)forkret;
00114
00115     return p;
00116 }
```

Referenced by [fork\(\)](#), and [userinit\(\)](#).

5.153.1.2 chpr()

```
int chpr (
    int pid,
    int priority )
```

Definition at line 559 of file [proc.c](#).

```
00560 {
00561     struct proc *p;
00562     acquire(&ptable.lock);
00563     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00564         if(p->pid == pid){
00565             p->priority = priority;
00566             break;
00567         }
00568     }
00569     release(&ptable.lock);
00570     return pid;
00571 }
```

Referenced by [main\(\)](#), and [sys_chpr\(\)](#).

5.153.1.3 cps()

```
int cps (
    void )
```

Definition at line 537 of file [proc.c](#).

```
00538 {
00539     struct proc *p;
00540     //Enables interrupts on this processor.
00541     sti();
00542
00543     //Loop over process table looking for process with pid.
00544     acquire(&ptable.lock);
00545     cprintf("name \t pid \t state \t priority \n");
00546     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00547         if(p->state == SLEEPING)
00548             cprintf("%s \t %d \t SLEEPING \t %d \n ", p->name, p->pid, p->priority);
00549         else if(p->state == RUNNING)
00550             cprintf("%s \t %d \t RUNNING \t %d \n ", p->name, p->pid, p->priority);
00551         else if(p->state == RUNNABLE)
00552             cprintf("%s \t %d \t RUNNABLE \t %d \n ", p->name, p->pid, p->priority);
00553     }
00554     release(&ptable.lock);
00555     return 23;
00556 }
```

Referenced by [main\(\)](#), and [sys_cps\(\)](#).

5.153.1.4 cpuid()

```
int cpuid (
    void )
```

Definition at line 31 of file [proc.c](#).

```
00031 {
00032     return mycpu() - cpus;
00033 }
```

Referenced by [seginit\(\)](#), and [trap\(\)](#).

5.153.1.5 exit()

```
void exit (
    void )
```

Definition at line 228 of file [proc.c](#).

```
00229 {
00230     struct proc *curproc = myproc();
00231     struct proc *p;
00232     int fd;
00233
00234     if(curproc == initproc)
00235         panic("init exiting");
00236
00237     // Close all open files.
00238     for(fd = 0; fd < NOFILE; fd++){
00239         if(curproc->ofile[fd]){
00240             fileclose(curproc->ofile[fd]);
00241             curproc->ofile[fd] = 0;
00242         }
00243     }
00244
00245     begin_op();
00246     iput(curproc->cwd);
```

```

00247     end_op();
00248     curproc->cwd = 0;
00249
00250     acquire(&ptable.lock);
00251
00252     // Parent might be sleeping in wait().
00253     wakeup1(curproc->parent);
00254
00255     // Pass abandoned children to init.
00256     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00257         if(p->parent == curproc){
00258             p->parent = initproc;
00259             if(p->state == ZOMBIE)
00260                 wakeup1(initproc);
00261         }
00262     }
00263
00264     // Jump into the scheduler, never to return.
00265     curproc->state = ZOMBIE;
00266     sched();
00267     panic("zombie exit");
00268 }

```

Referenced by [argptest\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [bsstest\(\)](#), [cat\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [dirfile\(\)](#), [dirtest\(\)](#), [exectest\(\)](#), [exitputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [panic\(\)](#), [pipe1\(\)](#), [rmdot\(\)](#), [rsect\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [sys_exit\(\)](#), [trap\(\)](#), [uio\(\)](#), [unlinkread\(\)](#), [validatestest\(\)](#), [wc\(\)](#), [writetest\(\)](#), [writetest1\(\)](#), and [wsect\(\)](#).

5.153.1.6 fork()

```

int fork (
    void )

```

Definition at line 181 of file [proc.c](#).

```

00182 {
00183     int i, pid;
00184     struct proc *np;
00185     struct proc *curproc = myproc();
00186
00187     // Allocate process.
00188     if((np = allocproc()) == 0){
00189         return -1;
00190     }
00191
00192     // Copy process state from proc.
00193     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
00194         kfree(np->kstack);
00195         np->kstack = 0;
00196         np->state = UNUSED;
00197         return -1;
00198     }
00199     np->sz = curproc->sz;
00200     np->parent = curproc;
00201     *np->tf = *curproc->tf;
00202
00203     // Clear %eax so that fork returns 0 in the child.
00204     np->tf->eax = 0;
00205
00206     for(i = 0; i < NOFILE; i++){
00207         if(curproc->ofile[i])
00208             np->ofile[i] = filedup(curproc->ofile[i]);
00209     }
00210     np->cwd = idup(curproc->cwd);
00211     safestrncpy(np->name, curproc->name, sizeof(curproc->name));
00212
00213     pid = np->pid;
00214     acquire(&ptable.lock);
00215     np->state = RUNNABLE;
00216     release(&ptable.lock);
00217
00218     return pid;
00219 }

```

Referenced by [bigargtest\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [exitiputtest\(\)](#), [exitwait\(\)](#), [fork1\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [sys_fork\(\)](#), [uio\(\)](#), and [validateetest\(\)](#).

5.153.1.7 forkret()

```
void forkret (
    void )
```

Definition at line 397 of file [proc.c](#).

```
00398 {
00399     static int first = 1;
00400     // Still holding ptable.lock from scheduler.
00401     release(&ptable.lock);
00402
00403     if (first) {
00404         // Some initialization functions must be run in the context
00405         // of a regular process (e.g., they call sleep), and thus cannot
00406         // be run from main().
00407         first = 0;
00408         iinit(ROOTDEV);
00409         initlog(ROOTDEV);
00410     }
00411
00412     // Return to "caller", actually trapret (see allocproc).
00413 }
```

Referenced by [allocproc\(\)](#).

5.153.1.8 growproc()

```
int growproc (
    int n )
```

Definition at line 159 of file [proc.c](#).

```
00160 {
00161     uint sz;
00162     struct proc *curproc = myproc();
00163
00164     sz = curproc->sz;
00165     if (n > 0) {
00166         if ((sz = allocuvmm(curproc->pgdir, sz, sz + n)) == 0)
00167             return -1;
00168     } else if (n < 0) {
00169         if ((sz = deallocuvmm(curproc->pgdir, sz, sz + n)) == 0)
00170             return -1;
00171     }
00172     curproc->sz = sz;
00173     switchvmm(curproc);
00174     return 0;
00175 }
```

Referenced by [sys_sbrk\(\)](#).

5.153.1.9 kill()

```
int kill (
    int pid )
```

Definition at line 480 of file [proc.c](#).

```
00481 {
00482     struct proc *p;
00483
00484     acquire(&ptable.lock);
00485     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00486         if(p->pid == pid){
00487             p->killed = 1;
00488             // Wake process from sleep if necessary.
00489             if(p->state == SLEEPING)
00490                 p->state = RUNNABLE;
00491             release(&ptable.lock);
00492             return 0;
00493         }
00494     }
00495     release(&ptable.lock);
00496     return -1;
00497 }
```

Referenced by [main\(\)](#), [mem\(\)](#), [preempt\(\)](#), [sbrktest\(\)](#), [sys_kill\(\)](#), and [validatetest\(\)](#).

5.153.1.10 mycpu()

```
struct cpu * mycpu (
    void )
```

Definition at line 38 of file [proc.c](#).

```
00039 {
00040     int apicid, i;
00041
00042     if(readeflags() & FL_IF)
00043         panic("mycpu called with interrupts enabled\n");
00044
00045     apicid = lapicid();
00046     // APIC IDs are not guaranteed to be contiguous. Maybe we should have
00047     // a reverse map, or reserve a register to store &cpus[i].
00048     for (i = 0; i < ncpu; ++i) {
00049         if (cpus[i].apicid == apicid)
00050             return &cpus[i];
00051     }
00052     panic("unknown apicid\n");
00053 }
```

Referenced by [acquire\(\)](#), [cupid\(\)](#), [holding\(\)](#), [myproc\(\)](#), [popcli\(\)](#), [pushcli\(\)](#), [sched\(\)](#), [scheduler\(\)](#), [startothers\(\)](#), and [switchvm\(\)](#).

5.153.1.11 myproc()

```
struct proc * myproc (
    void )
```

Definition at line 58 of file [proc.c](#).

```
00058     {
00059         struct cpu *c;
00060         struct proc *p;
00061         pushcli();
00062         c = mycpu();
00063         p = c->proc;
00064         popcli();
00065         return p;
00066     }
```

Referenced by [acquiresleep\(\)](#), [argfd\(\)](#), [argint\(\)](#), [argptr\(\)](#), [consoleread\(\)](#), [exec\(\)](#), [exit\(\)](#), [fdalloc\(\)](#), [fetchint\(\)](#), [fetchstr\(\)](#), [fork\(\)](#), [growproc\(\)](#), [holdingsleep\(\)](#), [namex\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [sched\(\)](#), [sleep\(\)](#), [sys_chdir\(\)](#), [sys_close\(\)](#), [sys_getpid\(\)](#), [sys_pipe\(\)](#), [sys_sbrk\(\)](#), [sys_sleep\(\)](#), [syscall\(\)](#), [trap\(\)](#), [wait\(\)](#), and [yield\(\)](#).

5.153.1.12 pinit()

```
void pinit (
    void )
```

Definition at line 24 of file [proc.c](#).

```
00025 {
00026     initlock(&ptable.lock, "ptable");
00027 }
```

5.153.1.13 procdump()

```
void procdump (
    void )
```

Definition at line 504 of file [proc.c](#).

```
00505 {
00506     static char *states[] = {
00507         [UNUSED]    "unused",
00508         [EMBRYO]    "embryo",
00509         [SLEEPING]  "sleep ",
00510         [RUNNABLE]  "runble",
00511         [RUNNING]   "run   ",
00512         [ZOMBIE]    "zombie"
00513     };
00514     int i;
00515     struct proc *p;
00516     char *state;
00517     uint pc[10];
00518
00519     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00520         if(p->state == UNUSED)
00521             continue;
00522         if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
00523             state = states[p->state];
00524         else
00525             state = "???";
00526         cprintf("%d %s %s", p->pid, state, p->name);
00527         if(p->state == SLEEPING){
00528             getcallerpcs((uint*)p->context->ebp+2, pc);
00529             for(i=0; i<10 && pc[i] != 0; i++)
00530                 cprintf(" %p", pc[i]);
00531         }
00532         cprintf("\n");
00533     }
00534 }
```

Referenced by [consoleintr\(\)](#).

5.153.1.14 sched()

```
void sched (
    void )
```

Definition at line 366 of file [proc.c](#).

```
00367 {
00368     int intena;
00369     struct proc *p = myproc();
00370
00371     if(!holding(&ptable.lock))
00372         panic("sched ptable.lock");
00373     if(mycpu()->ncli != 1)
00374         panic("sched locks");
00375     if(p->state == RUNNING)
00376         panic("sched running");
00377     if(readeflags() & FL_IF)
00378         panic("sched interruptible");
00379     intena = mycpu()->intena;
00380     swtch(&p->context, mycpu()->scheduler);
00381     mycpu()->intena = intena;
00382 }
```

Referenced by [exit\(\)](#), [sleep\(\)](#), and [yield\(\)](#).

5.153.1.15 scheduler()

```
void scheduler (
    void )
```

Definition at line 323 of file [proc.c](#).

```
00324 {
00325     struct proc *p;
00326     struct cpu *c = mycpu();
00327     c->proc = 0;
00328
00329     for(;;){
00330         // Enable interrupts on this processor.
00331         sti();
00332
00333         // Loop over process table looking for process to run.
00334         acquire(&ptable.lock);
00335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00336             if(p->state != RUNNABLE)
00337                 continue;
00338
00339             // Switch to chosen process. It is the process's job
00340             // to release ptable.lock and then reacquire it
00341             // before jumping back to us.
00342             c->proc = p;
00343             switchvm(p);
00344             p->state = RUNNING;
00345
00346             swtch(&(c->scheduler), p->context);
00347             switchkvm();
00348
00349             // Process is done running for now.
00350             // It should have changed its p->state before coming back.
00351             c->proc = 0;
00352         }
00353         release(&ptable.lock);
00354     }
00355 }
00356 }
```

Referenced by [sched\(\)](#).

5.153.1.16 sleep()

```
void sleep (
    void * chan,
    struct spinlock * lk )
```

Definition at line 418 of file [proc.c](#).

```
00419 {
00420     struct proc *p = myproc();
00421
00422     if(p == 0)
00423         panic("sleep");
00424
00425     if(lk == 0)
00426         panic("sleep without lk");
00427
00428     // Must acquire ptable.lock in order to
00429     // change p->state and then call sched.
00430     // Once we hold ptable.lock, we can be
00431     // guaranteed that we won't miss any wakeup
00432     // (wakeup runs with ptable.lock locked),
00433     // so it's okay to release lk.
00434     if(lk != &ptable.lock){ //DOC: sleeplock0
00435         acquire(&ptable.lock); //DOC: sleeplock1
00436         release(lk);
00437     }
00438     // Go to sleep.
00439     p->chan = chan;
00440     p->state = SLEEPING;
00441
00442     sched();
```

```

00443
00444 // Tidy up.
00445 p->chan = 0;
00446
00447 // Reacquire original lock.
00448 if(lk != &ptable.lock){ //DOC: sleeplock2
00449     release(&ptable.lock);
00450     acquire(lk);
00451 }
00452 }

```

Referenced by [acquiresleep\(\)](#), [begin_op\(\)](#), [consoleread\(\)](#), [iderw\(\)](#), [main\(\)](#), [openiputtest\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [sbrktest\(\)](#), [sys_sleep\(\)](#), [validatetest\(\)](#), and [wait\(\)](#).

5.153.1.17 trapret()

```

void trapret (
    void )

```

Referenced by [allocproc\(\)](#).

5.153.1.18 userinit()

```

void userinit (
    void )

```

Definition at line 121 of file [proc.c](#).

```

00122 {
00123     struct proc *p;
00124     extern char _binary_initcode_start[], _binary_initcode_size[];
00125
00126     p = allocproc();
00127
00128     initproc = p;
00129     if((p->pgdir = setupkvm()) == 0)
00130         panic("userinit: out of memory?");
00131     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
00132     p->sz = PGSIZE;
00133     memset(p->tf, 0, sizeof(*p->tf));
00134     p->tf->cs = (SEG_UCODE « 3) | DPL_USER;
00135     p->tf->ds = (SEG_UDATA « 3) | DPL_USER;
00136     p->tf->es = p->tf->ds;
00137     p->tf->ss = p->tf->ds;
00138     p->tf->eflags = FL_IF;
00139     p->tf->esp = PGSIZE;
00140     p->tf->eip = 0; // beginning of initcode.S
00141
00142     safestrcpy(p->name, "initcode", sizeof(p->name));
00143     p->cwd = namei("/");
00144
00145     // this assignment to p->state lets other cores
00146     // run this process. the acquire forces the above
00147     // writes to be visible, and the lock is also needed
00148     // because the assignment might not be atomic.
00149     acquire(&ptable.lock);
00150
00151     p->state = RUNNABLE;
00152
00153     release(&ptable.lock);
00154 }

```

5.153.1.19 wait()

```
int wait (
    void )
```

Definition at line 273 of file [proc.c](#).

```
00274 {
00275     struct proc *p;
00276     int havekids, pid;
00277     struct proc *curproc = myproc();
00278
00279     acquire(&ptable.lock);
00280     for(;;){
00281         // Scan through table looking for exited children.
00282         havekids = 0;
00283         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00284             if(p->parent != curproc)
00285                 continue;
00286             havekids = 1;
00287             if(p->state == ZOMBIE){
00288                 // Found one.
00289                 pid = p->pid;
00290                 kfree(p->kstack);
00291                 p->kstack = 0;
00292                 freevm(p->pgdir);
00293                 p->pid = 0;
00294                 p->parent = 0;
00295                 p->name[0] = 0;
00296                 p->killed = 0;
00297                 p->state = UNUSED;
00298                 release(&ptable.lock);
00299                 return pid;
00300             }
00301         }
00302
00303         // No point waiting if we don't have any children.
00304         if(!havekids || curproc->killed){
00305             release(&ptable.lock);
00306             return -1;
00307         }
00308
00309         // Wait for children to exit. (See wakeupl call in proc_exit.)
00310         sleep(curproc, &ptable.lock); //DOC: wait-sleep
00311     }
00312 }
```

Referenced by [bigargtest\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [exitiputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [sys_wait\(\)](#), [uio\(\)](#), and [validatestest\(\)](#).

5.153.1.20 wakeup()

```
void wakeup (
    void * chan )
```

Definition at line 469 of file [proc.c](#).

```
00470 {
00471     acquire(&ptable.lock);
00472     wakeupl(chan);
00473     release(&ptable.lock);
00474 }
```

Referenced by [consoleintr\(\)](#), [end_op\(\)](#), [ideintr\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), and [trap\(\)](#).

5.153.1.21 wakeup1()

```
static void wakeup1 (
    void * chan ) [static]
```

Definition at line 458 of file [proc.c](#).

```
00459 {
00460     struct proc *p;
00461
00462     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
00463         if(p->state == SLEEPING && p->chan == chan)
00464             p->state = RUNNABLE;
00465 }
```

Referenced by [exit\(\)](#), and [wakeup\(\)](#).

5.153.1.22 yield()

```
void yield (
    void )
```

Definition at line 386 of file [proc.c](#).

```
00387 {
00388     acquire(&ptable.lock); //DOC: yieldlock
00389     myproc()->state = RUNNABLE;
00390     sched();
00391     release(&ptable.lock);
00392 }
```

Referenced by [trap\(\)](#).

5.153.2 Variable Documentation

5.153.2.1 initproc

```
struct proc* initproc [static]
```

Definition at line 15 of file [proc.c](#).

Referenced by [exit\(\)](#), and [userinit\(\)](#).

5.153.2.2 lock

```
struct spinlock lock
```

Definition at line 11 of file [proc.c](#).

5.153.2.3 nextpid

```
int nextpid = 1
```

Definition at line 17 of file [proc.c](#).

Referenced by [allocproc\(\)](#).

5.153.2.4 proc

```
struct proc proc[NPROC]
```

Definition at line 12 of file [proc.c](#).

5.153.2.5

```
struct { ... } ptable
```

Referenced by [allocproc\(\)](#), [chpr\(\)](#), [cps\(\)](#), [exit\(\)](#), [fork\(\)](#), [forkret\(\)](#), [kill\(\)](#), [pinit\(\)](#), [procdump\(\)](#), [sched\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup\(\)](#), [wakeup1\(\)](#), and [yield\(\)](#).

5.154 proc.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "memlayout.h"
00005 #include "mmu.h"
00006 #include "x86.h"
00007 #include "proc.h"
00008 #include "spinlock.h"
00009
00010 struct {
00011     struct spinlock lock;
00012     struct proc proc[NPROC];
00013 } ptable;
00014
00015 static struct proc *initproc;
00016
00017 int nextpid = 1;
00018 extern void forkret(void);
00019 extern void trapret(void);
00020
00021 static void wakeup1(void *chan);
00022
00023 void
00024 pinit(void)
00025 {
00026     initlock(&ptable.lock, "ptable");
00027 }
00028
00029 // Must be called with interrupts disabled
00030 int
00031 cpuid() {
00032     return mycpu() - cpus;
00033 }
00034
00035 // Must be called with interrupts disabled to avoid the caller being
00036 // rescheduled between reading lapicid and running through the loop.
```

```

00037 struct cpu*
00038 mycpu(void)
00039 {
00040     int apicid, i;
00041
00042     if(readeflags() & FL_IF)
00043         panic("mycpu called with interrupts enabled\n");
00044
00045     apicid = lapicid();
00046     // APIC IDs are not guaranteed to be contiguous. Maybe we should have
00047     // a reverse map, or reserve a register to store &cpus[i].
00048     for (i = 0; i < ncpu; ++i) {
00049         if (cpus[i].apicid == apicid)
00050             return &cpus[i];
00051     }
00052     panic("unknown apicid\n");
00053 }
00054
00055 // Disable interrupts so that we are not rescheduled
00056 // while reading proc from the cpu structure
00057 struct proc*
00058 myproc(void) {
00059     struct cpu *c;
00060     struct proc *p;
00061     pushcli();
00062     c = mycpu();
00063     p = c->proc;
00064     popcli();
00065     return p;
00066 }
00067
00068 //PAGEBREAK: 32
00069 // Look in the process table for an UNUSED proc.
00070 // If found, change state to EMBRYO and initialize
00071 // state required to run in the kernel.
00072 // Otherwise return 0.
00073 static struct proc*
00074 allocproc(void)
00075 {
00076     struct proc *p;
00077     char *sp;
00078
00079     acquire(&ptable.lock);
00080
00081     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)
00082         if(p->state == UNUSED)
00083             goto found;
00084
00085     release(&ptable.lock);
00086     return 0;
00087
00088 found:
00089     p->state = EMBRYO;
00090     p->pid = nextpid++;
00091
00092     release(&ptable.lock);
00093
00094     // Allocate kernel stack.
00095     if((p->kstack = kalloc()) == 0){
00096         p->state = UNUSED;
00097         return 0;
00098     }
00099     sp = p->kstack + KSTACKSIZE;
00100
00101     // Leave room for trap frame.
00102     sp -= sizeof *p->tf;
00103     p->tf = (struct trapframe*)sp;
00104
00105     // Set up new context to start executing at forkret,
00106     // which returns to trapret.
00107     sp -= 4;
00108     *(uint*)sp = (uint)trapret;
00109
00110     sp -= sizeof *p->context;
00111     p->context = (struct context*)sp;
00112     memset(p->context, 0, sizeof *p->context);
00113     p->context->eip = (uint)forkret;
00114
00115     return p;
00116 }
00117
00118 //PAGEBREAK: 32
00119 // Set up first user process.
00120 void
00121 userinit(void)
00122 {
00123     struct proc *p;

```

```

00124  extern char _binary_initcode_start[], _binary_initcode_size[];
00125
00126  p = allocproc();
00127
00128  initproc = p;
00129  if((p->pgdir = setupkvm()) == 0)
00130      panic("userinit: out of memory?");
00131  inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
00132  p->sz = PGSIZE;
00133  memset(p->tf, 0, sizeof(*p->tf));
00134  p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
00135  p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
00136  p->tf->es = p->tf->ds;
00137  p->tf->ss = p->tf->ds;
00138  p->tf->eflags = FL_IF;
00139  p->tf->esp = PGSIZE;
00140  p->tf->eip = 0;  // beginning of initcode.S
00141
00142  safestrncpy(p->name, "initcode", sizeof(p->name));
00143  p->cwd = namei("/");
00144
00145  // this assignment to p->state lets other cores
00146  // run this process. the acquire forces the above
00147  // writes to be visible, and the lock is also needed
00148  // because the assignment might not be atomic.
00149  acquire(&ptable.lock);
00150
00151  p->state = RUNNABLE;
00152
00153  release(&ptable.lock);
00154 }
00155
00156 // Grow current process's memory by n bytes.
00157 // Return 0 on success, -1 on failure.
00158 int
00159 growproc(int n)
00160 {
00161     uint sz;
00162     struct proc *curproc = myproc();
00163
00164     sz = curproc->sz;
00165     if(n > 0){
00166         if((sz = allocuvm(curproc->pgdir, sz, sz + n)) == 0)
00167             return -1;
00168     } else if(n < 0){
00169         if((sz = deallocuvm(curproc->pgdir, sz, sz + n)) == 0)
00170             return -1;
00171     }
00172     curproc->sz = sz;
00173     switchuvm(curproc);
00174     return 0;
00175 }
00176
00177 // Create a new process copying p as the parent.
00178 // Sets up stack to return as if from system call.
00179 // Caller must set state of returned proc to RUNNABLE.
00180 int
00181 fork(void)
00182 {
00183     int i, pid;
00184     struct proc *np;
00185     struct proc *curproc = myproc();
00186
00187     // Allocate process.
00188     if((np = allocproc()) == 0){
00189         return -1;
00190     }
00191
00192     // Copy process state from proc.
00193     if((np->pgdir = copyuvm(curproc->pgdir, curproc->sz)) == 0){
00194         kfree(np->kstack);
00195         np->kstack = 0;
00196         np->state = UNUSED;
00197         return -1;
00198     }
00199     np->sz = curproc->sz;
00200     np->parent = curproc;
00201     *np->tf = *curproc->tf;
00202
00203     // Clear %eax so that fork returns 0 in the child.
00204     np->tf->eax = 0;
00205
00206     for(i = 0; i < NOFILE; i++)
00207         if(curproc->ofile[i])
00208             np->ofile[i] = filedup(curproc->ofile[i]);
00209     np->cwd = idup(curproc->cwd);
00210

```

```

00211     safestrcpy(np->name, curproc->name, sizeof(curproc->name));
00212
00213     pid = np->pid;
00214
00215     acquire(&ptable.lock);
00216
00217     np->state = RUNNABLE;
00218
00219     release(&ptable.lock);
00220
00221     return pid;
00222 }
00223
00224 // Exit the current process. Does not return.
00225 // An exited process remains in the zombie state
00226 // until its parent calls wait() to find out it exited.
00227 void
00228 exit(void)
00229 {
00230     struct proc *curproc = myproc();
00231     struct proc *p;
00232     int fd;
00233
00234     if(curproc == initproc)
00235         panic("init exiting");
00236
00237     // Close all open files.
00238     for(fd = 0; fd < NOFILE; fd++){
00239         if(curproc->ofile[fd]){
00240             fileclose(curproc->ofile[fd]);
00241             curproc->ofile[fd] = 0;
00242         }
00243     }
00244
00245     begin_op();
00246     iput(curproc->cwd);
00247     end_op();
00248     curproc->cwd = 0;
00249
00250     acquire(&ptable.lock);
00251
00252     // Parent might be sleeping in wait().
00253     wakeup1(curproc->parent);
00254
00255     // Pass abandoned children to init.
00256     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00257         if(p->parent == curproc){
00258             p->parent = initproc;
00259             if(p->state == ZOMBIE)
00260                 wakeup1(initproc);
00261         }
00262     }
00263
00264     // Jump into the scheduler, never to return.
00265     curproc->state = ZOMBIE;
00266     sched();
00267     panic("zombie exit");
00268 }
00269
00270 // Wait for a child process to exit and return its pid.
00271 // Return -1 if this process has no children.
00272 int
00273 wait(void)
00274 {
00275     struct proc *p;
00276     int havekids, pid;
00277     struct proc *curproc = myproc();
00278
00279     acquire(&ptable.lock);
00280     for(;;){
00281         // Scan through table looking for exited children.
00282         havekids = 0;
00283         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00284             if(p->parent != curproc)
00285                 continue;
00286             havekids = 1;
00287             if(p->state == ZOMBIE){
00288                 // Found one.
00289                 pid = p->pid;
00290                 kfree(p->kstack);
00291                 p->kstack = 0;
00292                 freevm(p->pgdir);
00293                 p->pid = 0;
00294                 p->parent = 0;
00295                 p->name[0] = 0;
00296                 p->killed = 0;
00297                 p->state = UNUSED;

```

```

00298         release(&ptable.lock);
00299         return pid;
00300     }
00301 }
00302
00303 // No point waiting if we don't have any children.
00304 if(!havekids || curproc->killed){
00305     release(&ptable.lock);
00306     return -1;
00307 }
00308
00309 // Wait for children to exit.  (See wakeup1 call in proc_exit.)
00310 sleep(curproc, &ptable.lock); //DOC: wait-sleep
00311 }
00312 }
00313
00314 //PAGEBREAK: 42
00315 // Per-CPU process scheduler.
00316 // Each CPU calls scheduler() after setting itself up.
00317 // Scheduler never returns.  It loops, doing:
00318 //  - choose a process to run
00319 //  - switch to start running that process
00320 //  - eventually that process transfers control
00321 //    via swtch back to the scheduler.
00322 void
00323 scheduler(void)
00324 {
00325     struct proc *p;
00326     struct cpu *c = mycpu();
00327     c->proc = 0;
00328
00329     for(;;){
00330         // Enable interrupts on this processor.
00331         sti();
00332
00333         // Loop over process table looking for process to run.
00334         acquire(&ptable.lock);
00335         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00336             if(p->state != RUNNABLE)
00337                 continue;
00338
00339             // Switch to chosen process.  It is the process's job
00340             // to release ptable.lock and then reacquire it
00341             // before jumping back to us.
00342             c->proc = p;
00343             switchvm(p);
00344             p->state = RUNNING;
00345
00346             swtch(&(c->scheduler), p->context);
00347             switchkvm();
00348
00349             // Process is done running for now.
00350             // It should have changed its p->state before coming back.
00351             c->proc = 0;
00352         }
00353         release(&ptable.lock);
00354     }
00355 }
00356 }
00357
00358 // Enter scheduler.  Must hold only ptable.lock
00359 // and have changed proc->state.  Saves and restores
00360 // intena because intena is a property of this
00361 // kernel thread, not this CPU.  It should
00362 // be proc->intena and proc->ncli, but that would
00363 // break in the few places where a lock is held but
00364 // there's no process.
00365 void
00366 sched(void)
00367 {
00368     int intena;
00369     struct proc *p = myproc();
00370
00371     if(!holding(&ptable.lock))
00372         panic("sched ptable.lock");
00373     if(mycpu()->ncli != 1)
00374         panic("sched locks");
00375     if(p->state == RUNNING)
00376         panic("sched running");
00377     if(readeflags() & FL_IF)
00378         panic("sched interruptible");
00379     intena = mycpu()->intena;
00380     swtch(&p->context, mycpu()->scheduler);
00381     mycpu()->intena = intena;
00382 }
00383
00384 // Give up the CPU for one scheduling round.

```

```

00385 void
00386 yield(void)
00387 {
00388     acquire(&ptable.lock); //DOC: yieldlock
00389     myproc()->state = RUNNABLE;
00390     sched();
00391     release(&ptable.lock);
00392 }
00393
00394 // A fork child's very first scheduling by scheduler()
00395 // will switch here. "Return" to user space.
00396 void
00397 forkret(void)
00398 {
00399     static int first = 1;
00400     // Still holding ptable.lock from scheduler.
00401     release(&ptable.lock);
00402
00403     if (first) {
00404         // Some initialization functions must be run in the context
00405         // of a regular process (e.g., they call sleep), and thus cannot
00406         // be run from main().
00407         first = 0;
00408         iinit(ROOTDEV);
00409         initlog(ROOTDEV);
00410     }
00411
00412     // Return to "caller", actually trapret (see allocproc).
00413 }
00414
00415 // Atomically release lock and sleep on chan.
00416 // Reacquires lock when awakened.
00417 void
00418 sleep(void *chan, struct spinlock *lk)
00419 {
00420     struct proc *p = myproc();
00421
00422     if (p == 0)
00423         panic("sleep");
00424
00425     if (lk == 0)
00426         panic("sleep without lk");
00427
00428     // Must acquire ptable.lock in order to
00429     // change p->state and then call sched.
00430     // Once we hold ptable.lock, we can be
00431     // guaranteed that we won't miss any wakeup
00432     // (wakeup runs with ptable.lock locked),
00433     // so it's okay to release lk.
00434     if (lk != &ptable.lock) { //DOC: sleeplock0
00435         acquire(&ptable.lock); //DOC: sleeplock1
00436         release(lk);
00437     }
00438     // Go to sleep.
00439     p->chan = chan;
00440     p->state = SLEEPING;
00441
00442     sched();
00443
00444     // Tidy up.
00445     p->chan = 0;
00446
00447     // Reacquire original lock.
00448     if (lk != &ptable.lock) { //DOC: sleeplock2
00449         release(&ptable.lock);
00450         acquire(lk);
00451     }
00452 }
00453
00454 //PAGEBREAK!
00455 // Wake up all processes sleeping on chan.
00456 // The ptable lock must be held.
00457 static void
00458 wakeup1(void *chan)
00459 {
00460     struct proc *p;
00461
00462     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)
00463         if (p->state == SLEEPING && p->chan == chan)
00464             p->state = RUNNABLE;
00465 }
00466
00467 // Wake up all processes sleeping on chan.
00468 void
00469 wakeup(void *chan)
00470 {
00471     acquire(&ptable.lock);

```

```

00472     wakeup1(chan);
00473     release(&ptable.lock);
00474 }
00475
00476 // Kill the process with the given pid.
00477 // Process won't exit until it returns
00478 // to user space (see trap in trap.c).
00479 int
00480 kill(int pid)
00481 {
00482     struct proc *p;
00483
00484     acquire(&ptable.lock);
00485     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00486         if(p->pid == pid){
00487             p->killed = 1;
00488             // Wake process from sleep if necessary.
00489             if(p->state == SLEEPING)
00490                 p->state = RUNNABLE;
00491             release(&ptable.lock);
00492             return 0;
00493         }
00494     }
00495     release(&ptable.lock);
00496     return -1;
00497 }
00498
00499 //PAGEBREAK: 36
00500 // Print a process listing to console.  For debugging.
00501 // Runs when user types ^P on console.
00502 // No lock to avoid wedging a stuck machine further.
00503 void
00504 procdump(void)
00505 {
00506     static char *states[] = {
00507         [UNUSED]    "unused",
00508         [EMBRYO]    "embryo",
00509         [SLEEPING]  "sleep ",
00510         [RUNNABLE]  "runble",
00511         [RUNNING]   "run   ",
00512         [ZOMBIE]    "zombie"
00513     };
00514     int i;
00515     struct proc *p;
00516     char *state;
00517     uint pc[10];
00518
00519     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00520         if(p->state == UNUSED)
00521             continue;
00522         if(p->state >= 0 && p->state < NELEM(states) && states[p->state])
00523             state = states[p->state];
00524         else
00525             state = "???";
00526         cprintf("%d %s %s", p->pid, state, p->name);
00527         if(p->state == SLEEPING){
00528             getcallerpcs((uint*)p->context->ebp+2, pc);
00529             for(i=0; i<10 && pc[i] != 0; i++)
00530                 cprintf(" %p", pc[i]);
00531         }
00532         cprintf("\n");
00533     }
00534 }
00535
00536 int
00537 cps()
00538 {
00539     struct proc *p;
00540     //Enables interrupts on this processor.
00541     sti();
00542
00543     //Loop over process table looking for process with pid.
00544     acquire(&ptable.lock);
00545     cprintf("name \t pid \t state \t priority \n");
00546     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00547         if(p->state == SLEEPING)
00548             cprintf("%s \t %d \t SLEEPING \t %d \n ", p->name, p->pid, p->priority);
00549         else if(p->state == RUNNING)
00550             cprintf("%s \t %d \t RUNNING \t %d \n ", p->name, p->pid, p->priority);
00551         else if(p->state == RUNNABLE)
00552             cprintf("%s \t %d \t RUNNABLE \t %d \n ", p->name, p->pid, p->priority);
00553     }
00554     release(&ptable.lock);
00555     return 23;
00556 }
00557
00558 int

```



```

00559 chpr(int pid, int priority)
00560 {
00561     struct proc *p;
00562     acquire(&ptable.lock);
00563     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00564         if(p->pid == pid){
00565             p->priority = priority;
00566             break;
00567         }
00568     }
00569     release(&ptable.lock);
00570     return pid;
00571 }

```

5.155 proc.d File Reference

5.156 proc.d

[Go to the documentation of this file.](#)

```

00001 proc.o: proc.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 memlayout.h mmu.h x86.h proc.h spinlock.h

```

5.157 proc.h File Reference

Classes

- struct [context](#)
- struct [cpu](#)
- struct [proc](#)

Enumerations

- enum [procstate](#) {
[UNUSED](#) , [EMBRYO](#) , [SLEEPING](#) , [RUNNABLE](#) ,
[RUNNING](#) , [ZOMBIE](#) }

Variables

- struct [cpu](#) [cpus](#) [[NCPU](#)]
- int [ncpu](#)

5.157.1 Enumeration Type Documentation

5.157.1.1 procstate

```
enum procstate
```

Enumerator

UNUSED	
EMBRYO	
SLEEPING	
RUNNABLE	
RUNNING	
ZOMBIE	

Definition at line 35 of file [proc.h](#).

```
00035 { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
```

5.157.2 Variable Documentation

5.157.2.1 cpus

```
struct cpu cpus[NCPU] [extern]
```

Definition at line 14 of file [mp.c](#).

Referenced by [cpuid\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), [seginit\(\)](#), and [startothers\(\)](#).

5.157.2.2 ncpu

```
int ncpu [extern]
```

Definition at line 15 of file [mp.c](#).

Referenced by [ideinit\(\)](#), [mpinit\(\)](#), [mycpu\(\)](#), and [startothers\(\)](#).

5.158 proc.h

[Go to the documentation of this file.](#)

```
00001 // Per-CPU state
00002 struct cpu {
00003     uchar apicid;           // Local APIC ID
00004     struct context *scheduler; // swtch() here to enter scheduler
00005     struct taskstate ts;     // Used by x86 to find stack for interrupt
00006     struct segdesc gdt[NSEGS]; // x86 global descriptor table
00007     volatile uint started;   // Has the CPU started?
00008     int ncli;                // Depth of pushcli nesting.
00009     int intena;              // Were interrupts enabled before pushcli?
00010     struct proc *proc;       // The process running on this cpu or null
00011 };
00012
00013 extern struct cpu cpus[NCPU];
00014 extern int ncpu;
00015
00016 //PAGEBREAK: 17
00017 // Saved registers for kernel context switches.
00018 // Don't need to save all the segment registers (%cs, etc),
```

```

00019 // because they are constant across kernel contexts.
00020 // Don't need to save %eax, %ecx, %edx, because the
00021 // x86 convention is that the caller has saved them.
00022 // Contexts are stored at the bottom of the stack they
00023 // describe; the stack pointer is the address of the context.
00024 // The layout of the context matches the layout of the stack in swtch.S
00025 // at the "Switch stacks" comment. Switch doesn't save eip explicitly,
00026 // but it is on the stack and allocproc() manipulates it.
00027 struct context {
00028     uint edi;
00029     uint esi;
00030     uint ebx;
00031     uint ebp;
00032     uint eip;
00033 };
00034
00035 enum procstate { UNUSED, EMBRYO, SLEEPING, RUNNABLE, RUNNING, ZOMBIE };
00036
00037 // Per-process state
00038 struct proc {
00039     uint sz; // Size of process memory (bytes)
00040     pde_t* pgdir; // Page table
00041     char *kstack; // Bottom of kernel stack for this process
00042     enum procstate state; // Process state
00043     int pid; // Process ID
00044     struct proc *parent; // Parent process
00045     struct trapframe *tf; // Trap frame for current syscall
00046     struct context *context; // swtch() here to run process
00047     void *chan; // If non-zero, sleeping on chan
00048     int killed; // If non-zero, have been killed
00049     struct file *ofile[NOFILE]; // Open files
00050     struct inode *cwd; // Current directory
00051     char name[16]; // Process name (debugging)
00052     int priority; // Process priority
00053 };
00054
00055 // Process memory is laid out contiguously, low addresses first:
00056 // text
00057 // original data and bss
00058 // fixed-size stack
00059 // expandable heap

```

5.159 ps.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"

```

Functions

- int [main](#) (void)

5.159.1 Function Documentation

5.159.1.1 main()

```

int main (
    void )

```

Definition at line 6 of file [ps.c](#).

```

00006     {
00007     cps();
00008     exit();
00009 }

```

5.160 ps.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "fcntl.h"
00005
00006 int main(void){
00007     cps();
00008     exit();
00009 }
```

5.161 ps.d File Reference

5.162 ps.d

[Go to the documentation of this file.](#)

```
00001 ps.o: ps.c /usr/include/stdc-predef.h types.h stat.h user.h fcntl.h
```

5.163 pstree.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "uproc.h"
```

Macros

- #define `MAXPROC` 64

Functions

- int `main` (int argc, char *argv[])

5.163.1 Macro Definition Documentation

5.163.1.1 MAXPROC

```
#define MAXPROC 64
```

Definition at line 5 of file [pstree.c](#).

5.163.2 Function Documentation

5.163.2.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 8 of file [pstree.c](#).

```
00008      {
00009      // declare a uproc struct with an allocated size of 64
00010      struct uproc *prs = malloc(MAXPROC*sizeof(struct uproc));
00011
00012      // call getprocs in kernel space and get the amount of processes running back
00013      int num = getprocs(MAXPROC, prs);
00014
00015      // getprocs will return -1 if an error is received. Handle errors here
00016      if(num == -1){
00017          printf(1, "Kernel returned an error when getting processes. \n");
00018          exit();
00019      }
00020
00021      // create a counter
00022      int i=0;
00023
00024      // loop through uproc structs that were created in the kernel space for each process
00025      for(; i< num; i++) {
00026          // format output to the user
00027          if(prs[i].pid == 1){
00028              // if the pid is 1 then this is init
00029              printf(1, "%s[%d] \n", prs[i].name, prs[i].pid-1);
00030          }else if(prs[i].ppid == 1){
00031              // if parent pid is 1, then it is a child to parent init
00032              printf(1, "  %s[%d] \n", prs[i].name, prs[i].pid-1);
00033          }else {
00034              // if parent pid is not 1, then it is a grandchild to init and child to another process
00035              printf(1, "    %s[%d] \n", prs[i].name, prs[i].pid-1);
00036          }
00037      }
00038
00039      // exit after pstree command is done
00040      exit();
00041  }
```

5.164 pstree.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "uproc.h"
00005 #define MAXPROC 64
00006
00007 int
00008 main(int argc, char *argv[]){
00009     // declare a uproc struct with an allocated size of 64
00010     struct uproc *prs = malloc(MAXPROC*sizeof(struct uproc));
00011
00012     // call getprocs in kernel space and get the amount of processes running back
00013     int num = getprocs(MAXPROC, prs);
00014
00015     // getprocs will return -1 if an error is received. Handle errors here
00016     if(num == -1){
00017         printf(1, "Kernel returned an error when getting processes. \n");
00018         exit();
00019     }
00020
00021     // create a counter
00022     int i=0;
00023
00024     // loop through uproc structs that were created in the kernel space for each process
```

```
00025     for( i< num; i++) {
00026         // format output to the user
00027         if(prs[i].pid == 1){
00028             // if the pid is 1 then this is init
00029             printf(1, "%s[%d] \n", prs[i].name, prs[i].pid-1);
00030         }else if(prs[i].ppid == 1){
00031             // if parent pid is 1, then it is a child to parent init
00032             printf(1, "    %s[%d] \n", prs[i].name, prs[i].pid-1);
00033         }else {
00034             // if parent pid is not 1, then it is a grandchild to init and child to another process
00035             printf(1, "        %s[%d] \n", prs[i].name, prs[i].pid-1);
00036         }
00037     }
00038
00039     // exit after pstree command is done
00040     exit();
00041 }
```

5.165 pstree.d File Reference

5.166 pstree.d

[Go to the documentation of this file.](#)

```
00001 pstree.o: pstree.c /usr/include/stdc-predef.h types.h stat.h user.h \
00002 uproc.h
```

5.167 README.md File Reference

5.168 rm.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- int [main](#) (int argc, char *argv[])

5.168.1 Function Documentation

5.168.1.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 6 of file [rm.c](#).

```
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "Usage: rm files...\n");
00012         exit();
00013     }
00014
00015     for(i = 1; i < argc; i++){
00016         if(unlink(argv[i]) < 0){
00017             printf(2, "rm: %s failed to delete\n", argv[i]);
00018             break;
00019         }
00020     }
00021
00022     exit();
00023 }
```

5.169 rm.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 int
00006 main(int argc, char *argv[])
00007 {
00008     int i;
00009
00010     if(argc < 2){
00011         printf(2, "Usage: rm files...\n");
00012         exit();
00013     }
00014
00015     for(i = 1; i < argc; i++){
00016         if(unlink(argv[i]) < 0){
00017             printf(2, "rm: %s failed to delete\n", argv[i]);
00018             break;
00019         }
00020     }
00021
00022     exit();
00023 }
```

5.170 rm.d File Reference**5.171 rm.d**

[Go to the documentation of this file.](#)

```
00001 rm.o: rm.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.172 sh.c File Reference

```
#include "types.h"
#include "user.h"
#include "fcntl.h"
```

Classes

- struct [backcmd](#)
- struct [cmd](#)
- struct [execcmd](#)
- struct [listcmd](#)
- struct [pipecmd](#)
- struct [redircmd](#)

Macros

- #define [BACK](#) 5
- #define [EXEC](#) 1
- #define [LIST](#) 4
- #define [MAXARGS](#) 10
- #define [PIPE](#) 3
- #define [REDIR](#) 2

Functions

- struct [cmd](#) * [backcmd](#) (struct [cmd](#) *subcmd)
- struct [cmd](#) * [execcmd](#) (void)
- int [fork1](#) (void)
- int [getcmd](#) (char *buf, int nbuf)
- int [gettoken](#) (char **ps, char *es, char **q, char **eq)
- struct [cmd](#) * [listcmd](#) (struct [cmd](#) *left, struct [cmd](#) *right)
- int [main](#) (void)
- struct [cmd](#) * [nulterminate](#) (struct [cmd](#) *)
- void [panic](#) (char *)
- struct [cmd](#) * [parseblock](#) (char **ps, char *es)
- struct [cmd](#) * [parsecmd](#) (char *)
- struct [cmd](#) * [parseexec](#) (char **, char *)
- struct [cmd](#) * [parseline](#) (char **, char *)
- struct [cmd](#) * [parsepipe](#) (char **, char *)
- struct [cmd](#) * [parseredirs](#) (struct [cmd](#) *cmd, char **ps, char *es)
- int [peek](#) (char **ps, char *es, char *toks)
- struct [cmd](#) * [pipecmd](#) (struct [cmd](#) *left, struct [cmd](#) *right)
- struct [cmd](#) * [redircmd](#) (struct [cmd](#) *subcmd, char *file, char *efile, int mode, int fd)
- void [runcmd](#) (struct [cmd](#) *cmd)

Variables

- char [symbols](#) [] = "<|>&();"
- char [whitespace](#) [] = "\t\r\n\v"

5.172.1 Macro Definition Documentation

5.172.1.1 BACK

```
#define BACK 5
```

Definition at line 12 of file [sh.c](#).

5.172.1.2 EXEC

```
#define EXEC 1
```

Definition at line 8 of file [sh.c](#).

5.172.1.3 LIST

```
#define LIST 4
```

Definition at line 11 of file [sh.c](#).

5.172.1.4 MAXARGS

```
#define MAXARGS 10
```

Definition at line 14 of file [sh.c](#).

5.172.1.5 PIPE

```
#define PIPE 3
```

Definition at line 10 of file [sh.c](#).

5.172.1.6 REDIR

```
#define REDIR 2
```

Definition at line 9 of file [sh.c](#).

5.172.2 Function Documentation

5.172.2.1 backcmd()

```
struct cmd * backcmd (  
    struct cmd * subcmd )
```

Definition at line 249 of file [sh.c](#).

```
00250 {  
00251     struct backcmd *cmd;  
00252  
00253     cmd = malloc(sizeof(*cmd));  
00254     memset(cmd, 0, sizeof(*cmd));  
00255     cmd->type = BACK;  
00256     cmd->cmd = subcmd;  
00257     return (struct cmd*)cmd;  
00258 }
```

5.172.2.2 execcmd()

```
struct cmd * execcmd (  
    void )
```

Definition at line 196 of file [sh.c](#).

```
00197 {  
00198     struct execcmd *cmd;  
00199  
00200     cmd = malloc(sizeof(*cmd));  
00201     memset(cmd, 0, sizeof(*cmd));  
00202     cmd->type = EXEC;  
00203     return (struct cmd*)cmd;  
00204 }
```

5.172.2.3 fork1()

```
int fork1 (  
    void )
```

Definition at line 182 of file [sh.c](#).

```
00183 {  
00184     int pid;  
00185  
00186     pid = fork();  
00187     if(pid == -1)  
00188         panic("fork");  
00189     return pid;  
00190 }
```

Referenced by [main\(\)](#), and [runcmd\(\)](#).

5.172.2.4 getcmd()

```
int getcmd (
    char * buf,
    int nbuf )
```

Definition at line 134 of file [sh.c](#).

```
00135 {
00136     printf(2, "$ ");
00137     memset(buf, 0, nbuf);
00138     gets(buf, nbuf);
00139     if(buf[0] == 0) // EOF
00140         return -1;
00141     return 0;
00142 }
```

Referenced by [main\(\)](#).

5.172.2.5 gettoken()

```
int gettoken (
    char ** ps,
    char * es,
    char ** q,
    char ** eq )
```

Definition at line 266 of file [sh.c](#).

```
00267 {
00268     char *s;
00269     int ret;
00270
00271     s = *ps;
00272     while(s < es && strchr(whitespace, *s))
00273         s++;
00274     if(q)
00275         *q = s;
00276     ret = *s;
00277     switch(*s){
00278     case 0:
00279         break;
00280     case '|':
00281     case '(':
00282     case ')':
00283     case ';':
00284     case '&':
00285     case '<':
00286         s++;
00287         break;
00288     case '>':
00289         s++;
00290         if(*s == '>'){
00291             ret = '+';
00292             s++;
00293         }
00294         break;
00295     default:
00296         ret = 'a';
00297         while(s < es && !strchr(whitespace, *s) && !strchr(symbols, *s))
00298             s++;
00299         break;
00300     }
00301     if(eq)
00302         *eq = s;
00303
00304     while(s < es && strchr(whitespace, *s))
00305         s++;
00306     *ps = s;
00307     return ret;
00308 }
```

Referenced by [parseblock\(\)](#), [parseexec\(\)](#), [parseline\(\)](#), [parsepipe\(\)](#), and [parseredirs\(\)](#).

5.172.2.6 listcmd()

```
struct cmd * listcmd (
    struct cmd * left,
    struct cmd * right )
```

Definition at line 236 of file [sh.c](#).

```
00237 {
00238     struct listcmd *cmd;
00239
00240     cmd = malloc(sizeof(*cmd));
00241     memset(cmd, 0, sizeof(*cmd));
00242     cmd->type = LIST;
00243     cmd->left = left;
00244     cmd->right = right;
00245     return (struct cmd*)cmd;
00246 }
```

5.172.2.7 main()

```
int main (
    void )
```

Definition at line 145 of file [sh.c](#).

```
00146 {
00147     static char buf[100];
00148     int fd;
00149
00150     // Ensure that three file descriptors are open.
00151     while((fd = open("console", O_RDWR)) >= 0){
00152         if(fd >= 3){
00153             close(fd);
00154             break;
00155         }
00156     }
00157
00158     // Read and run input commands.
00159     while(getcmd(buf, sizeof(buf)) >= 0){
00160         if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
00161             // Chdir must be called by the parent, not the child.
00162             buf[strlen(buf)-1] = 0; // chop \n
00163             if(chdir(buf+3) < 0)
00164                 printf(2, "cannot cd %s\n", buf+3);
00165             continue;
00166         }
00167         if(fork1() == 0)
00168             runcmd(parsecmd(buf));
00169         wait();
00170     }
00171     exit();
00172 }
```

5.172.2.8 nulterminate()

```
struct cmd * nulterminate (
    struct cmd * cmd )
```

Definition at line 450 of file [sh.c](#).

```
00451 {
00452     int i;
00453     struct backcmd *bcmd;
00454     struct execcmd *ecmd;
00455     struct listcmd *lcmd;
00456     struct pipecmd *pcmd;
00457     struct redircmd *rcmd;
```

```

00458
00459     if(cmd == 0)
00460         return 0;
00461
00462     switch(cmd->type) {
00463     case EXEC:
00464         ecmd = (struct execcmd*)cmd;
00465         for(i=0; ecmd->argv[i]; i++)
00466             *ecmd->eargv[i] = 0;
00467         break;
00468
00469     case REDIR:
00470         rcmd = (struct redircmd*)cmd;
00471         nulterminate(rcmd->cmd);
00472         *rcmd->efile = 0;
00473         break;
00474
00475     case PIPE:
00476         pcmd = (struct pipecmd*)cmd;
00477         nulterminate(pcmd->left);
00478         nulterminate(pcmd->right);
00479         break;
00480
00481     case LIST:
00482         lcmd = (struct listcmd*)cmd;
00483         nulterminate(lcmd->left);
00484         nulterminate(lcmd->right);
00485         break;
00486
00487     case BACK:
00488         bcmd = (struct backcmd*)cmd;
00489         nulterminate(bcmd->cmd);
00490         break;
00491     }
00492     return cmd;
00493 }

```

Referenced by [nulterminate\(\)](#), and [parsecmd\(\)](#).

5.172.2.9 panic()

```

void panic (
    char * s )

```

Definition at line 175 of file [sh.c](#).

```

00176 {
00177     printf(2, "%s\n", s);
00178     exit();
00179 }

```

Referenced by [fork1\(\)](#), [parseblock\(\)](#), [parsecmd\(\)](#), [parseexec\(\)](#), [parseredirs\(\)](#), and [runcmd\(\)](#).

5.172.2.10 parseblock()

```

struct cmd * parseblock (
    char ** ps,
    char * es )

```

Definition at line 400 of file [sh.c](#).

```

00401 {
00402     struct cmd *cmd;
00403
00404     if(!peek(ps, es, "("))
00405         panic("parseblock");
00406     gettoken(ps, es, 0, 0);
00407     cmd = parseline(ps, es);
00408     if(!peek(ps, es, " "))
00409         panic("syntax - missing ");
00410     gettoken(ps, es, 0, 0);
00411     cmd = parseredirs(cmd, ps, es);
00412     return cmd;
00413 }

```

Referenced by [parseexec\(\)](#).

5.172.2.11 parsecmd()

```
struct cmd * parsecmd (
    char * s )
```

Definition at line 328 of file [sh.c](#).

```
00329 {
00330     char *es;
00331     struct cmd *cmd;
00332
00333     es = s + strlen(s);
00334     cmd = parseline(&s, es);
00335     peek(&s, es, "");
00336     if(s != es){
00337         printf(2, "leftovers: %s\n", s);
00338         panic("syntax");
00339     }
00340     nulterminate(cmd);
00341     return cmd;
00342 }
```

Referenced by [main\(\)](#).

5.172.2.12 parseexec()

```
struct cmd * parseexec (
    char ** ps,
    char * es )
```

Definition at line 416 of file [sh.c](#).

```
00417 {
00418     char *q, *eq;
00419     int tok, argc;
00420     struct execcmd *cmd;
00421     struct cmd *ret;
00422
00423     if(peek(ps, es, "("))
00424         return parseblock(ps, es);
00425
00426     ret = execcmd();
00427     cmd = (struct execcmd*)ret;
00428
00429     argc = 0;
00430     ret = parseredirs(ret, ps, es);
00431     while(!peek(ps, es, "|)&;")){
00432         if((tok=gettoken(ps, es, &q, &eq)) == 0)
00433             break;
00434         if(tok != 'a')
00435             panic("syntax");
00436         cmd->argv[argc] = q;
00437         cmd->eargv[argc] = eq;
00438         argc++;
00439         if(argc >= MAXARGS)
00440             panic("too many args");
00441         ret = parseredirs(ret, ps, es);
00442     }
00443     cmd->argv[argc] = 0;
00444     cmd->eargv[argc] = 0;
00445     return ret;
00446 }
```

Referenced by [parsepipe\(\)](#).

5.172.2.13 parseline()

```
struct cmd * parseline (
    char ** ps,
    char * es )
```

Definition at line 345 of file [sh.c](#).

```
00346 {
00347     struct cmd *cmd;
00348
00349     cmd = parsepipe(ps, es);
00350     while(peek(ps, es, "&")){
00351         gettoken(ps, es, 0, 0);
00352         cmd = backcmd(cmd);
00353     }
00354     if(peek(ps, es, ";")){
00355         gettoken(ps, es, 0, 0);
00356         cmd = listcmd(cmd, parseline(ps, es));
00357     }
00358     return cmd;
00359 }
```

Referenced by [parseblock\(\)](#), [parsecmd\(\)](#), and [parseline\(\)](#).

5.172.2.14 parsepipe()

```
struct cmd * parsepipe (
    char ** ps,
    char * es )
```

Definition at line 362 of file [sh.c](#).

```
00363 {
00364     struct cmd *cmd;
00365
00366     cmd = parseexec(ps, es);
00367     if(peek(ps, es, "|")){
00368         gettoken(ps, es, 0, 0);
00369         cmd = pipecmd(cmd, parsepipe(ps, es));
00370     }
00371     return cmd;
00372 }
```

Referenced by [parseline\(\)](#), and [parsepipe\(\)](#).

5.172.2.15 parseredirs()

```
struct cmd * parseredirs (
    struct cmd * cmd,
    char ** ps,
    char * es )
```

Definition at line 375 of file [sh.c](#).

```
00376 {
00377     int tok;
00378     char *q, *eq;
00379
00380     while(peek(ps, es, "<>")){
00381         tok = gettoken(ps, es, 0, 0);
00382         if(gettoken(ps, es, &q, &eq) != 'a')
00383             panic("missing file for redirection");
00384         switch(tok) {
00385             case '<':
```

```

00386     cmd = redircmd(cmd, q, eq, O_RDONLY, 0);
00387     break;
00388     case '>':
00389     cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
00390     break;
00391     case '+': // »
00392     cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
00393     break;
00394     }
00395 }
00396 return cmd;
00397 }

```

Referenced by [parseblock\(\)](#), and [parseexec\(\)](#).

5.172.2.16 peek()

```

int peek (
    char ** ps,
    char * es,
    char * toks )

```

Definition at line 311 of file [sh.c](#).

```

00312 {
00313     char *s;
00314
00315     s = *ps;
00316     while(s < es && strchr(whitespace, *s))
00317         s++;
00318     *ps = s;
00319     return *s && strchr(toks, *s);
00320 }

```

Referenced by [parseblock\(\)](#), [parsecmd\(\)](#), [parseexec\(\)](#), [parseline\(\)](#), [parsepipe\(\)](#), and [parseredirs\(\)](#).

5.172.2.17 pipecmd()

```

struct cmd * pipecmd (
    struct cmd * left,
    struct cmd * right )

```

Definition at line 223 of file [sh.c](#).

```

00224 {
00225     struct pipecmd *cmd;
00226
00227     cmd = malloc(sizeof(*cmd));
00228     memset(cmd, 0, sizeof(*cmd));
00229     cmd->type = PIPE;
00230     cmd->left = left;
00231     cmd->right = right;
00232     return (struct cmd*)cmd;
00233 }

```


5.172.2.18 redircmd()

```

struct cmd * redircmd (
    struct cmd * subcmd,
    char * file,
    char * efile,
    int mode,
    int fd )

```

Definition at line 207 of file [sh.c](#).

```

00208 {
00209     struct redircmd *cmd;
00210
00211     cmd = malloc(sizeof(*cmd));
00212     memset(cmd, 0, sizeof(*cmd));
00213     cmd->type = REDIR;
00214     cmd->cmd = subcmd;
00215     cmd->file = file;
00216     cmd->efile = efile;
00217     cmd->mode = mode;
00218     cmd->fd = fd;
00219     return (struct cmd*)cmd;
00220 }

```

5.172.2.19 runcmd()

```

void runcmd (
    struct cmd * cmd )

```

Definition at line 58 of file [sh.c](#).

```

00059 {
00060     int p[2];
00061     struct backcmd *bcmd;
00062     struct execcmd *ecmd;
00063     struct listcmd *lcmd;
00064     struct pipecmd *pcmd;
00065     struct redircmd *rcmd;
00066
00067     if(cmd == 0)
00068         exit();
00069
00070     switch(cmd->type){
00071     default:
00072         panic("runcmd");
00073
00074     case EXEC:
00075         ecmd = (struct execcmd*)cmd;
00076         if(ecmd->argv[0] == 0)
00077             exit();
00078         exec(ecmd->argv[0], ecmd->argv);
00079         printf(2, "exec %s failed\n", ecmd->argv[0]);
00080         break;
00081
00082     case REDIR:
00083         rcmd = (struct redircmd*)cmd;
00084         close(rcmd->fd);
00085         if(open(rcmd->file, rcmd->mode) < 0){
00086             printf(2, "open %s failed\n", rcmd->file);
00087             exit();
00088         }
00089         runcmd(rcmd->cmd);
00090         break;
00091
00092     case LIST:
00093         lcmd = (struct listcmd*)cmd;
00094         if(fork1() == 0)
00095             runcmd(lcmd->left);
00096         wait();
00097         runcmd(lcmd->right);
00098         break;
00099
00100     case PIPE:
00101         pcmd = (struct pipecmd*)cmd;

```

```
00102     if(pipe(p) < 0)
00103         panic("pipe");
00104     if(fork1() == 0){
00105         close(1);
00106         dup(p[1]);
00107         close(p[0]);
00108         close(p[1]);
00109         runcmd(pcmd->left);
00110     }
00111     if(fork1() == 0){
00112         close(0);
00113         dup(p[0]);
00114         close(p[0]);
00115         close(p[1]);
00116         runcmd(pcmd->right);
00117     }
00118     close(p[0]);
00119     close(p[1]);
00120     wait();
00121     wait();
00122     break;
00123
00124     case BACK:
00125         bcmd = (struct backcmd*)cmd;
00126         if(fork1() == 0)
00127             runcmd(bcmd->cmd);
00128         break;
00129     }
00130     exit();
00131 }
```

Referenced by [main\(\)](#), and [runcmd\(\)](#).

5.172.3 Variable Documentation

5.172.3.1 symbols

```
char symbols[] = "<|>&; () "
```

Definition at line 263 of file [sh.c](#).

Referenced by [gettoken\(\)](#).

5.172.3.2 whitespace

```
char whitespace[] = " \t\r\n\v"
```

Definition at line 262 of file [sh.c](#).

Referenced by [gettoken\(\)](#), and [peek\(\)](#).

5.173 sh.c

[Go to the documentation of this file.](#)

```

00001 // Shell.
00002
00003 #include "types.h"
00004 #include "user.h"
00005 #include "fcntl.h"
00006
00007 // Parsed command representation
00008 #define EXEC 1
00009 #define REDIR 2
00010 #define PIPE 3
00011 #define LIST 4
00012 #define BACK 5
00013
00014 #define MAXARGS 10
00015
00016 struct cmd {
00017     int type;
00018 };
00019
00020 struct execcmd {
00021     int type;
00022     char *argv[MAXARGS];
00023     char *eargv[MAXARGS];
00024 };
00025
00026 struct redircmd {
00027     int type;
00028     struct cmd *cmd;
00029     char *file;
00030     char *efile;
00031     int mode;
00032     int fd;
00033 };
00034
00035 struct pipecmd {
00036     int type;
00037     struct cmd *left;
00038     struct cmd *right;
00039 };
00040
00041 struct listcmd {
00042     int type;
00043     struct cmd *left;
00044     struct cmd *right;
00045 };
00046
00047 struct backcmd {
00048     int type;
00049     struct cmd *cmd;
00050 };
00051
00052 int fork1(void); // Fork but panics on failure.
00053 void panic(char*);
00054 struct cmd *parsecmd(char*);
00055
00056 // Execute cmd. Never returns.
00057 void
00058 runcmd(struct cmd *cmd)
00059 {
00060     int p[2];
00061     struct backcmd *bcmd;
00062     struct execcmd *ecmd;
00063     struct listcmd *lcmd;
00064     struct pipecmd *pcmd;
00065     struct redircmd *rcmd;
00066
00067     if(cmd == 0)
00068         exit();
00069
00070     switch(cmd->type){
00071     default:
00072         panic("runcmd");
00073
00074     case EXEC:
00075         ecmd = (struct execcmd*)cmd;
00076         if(ecmd->argv[0] == 0)
00077             exit();
00078         exec(ecmd->argv[0], ecmd->argv);
00079         printf(2, "exec %s failed\n", ecmd->argv[0]);
00080         break;
00081
00082     case REDIR:

```

```

00083     rcmd = (struct redircmd*)cmd;
00084     close(rcmd->fd);
00085     if(open(rcmd->file, rcmd->mode) < 0){
00086         printf(2, "open %s failed\n", rcmd->file);
00087         exit();
00088     }
00089     runcmd(rcmd->cmd);
00090     break;
00091
00092 case LIST:
00093     lcm = (struct listcmd*)cmd;
00094     if(fork1() == 0)
00095         runcmd(lcmd->left);
00096     wait();
00097     runcmd(lcmd->right);
00098     break;
00099
00100 case PIPE:
00101     pcmd = (struct pipecmd*)cmd;
00102     if(pipe(p) < 0)
00103         panic("pipe");
00104     if(fork1() == 0){
00105         close(1);
00106         dup(p[1]);
00107         close(p[0]);
00108         close(p[1]);
00109         runcmd(pcmd->left);
00110     }
00111     if(fork1() == 0){
00112         close(0);
00113         dup(p[0]);
00114         close(p[0]);
00115         close(p[1]);
00116         runcmd(pcmd->right);
00117     }
00118     close(p[0]);
00119     close(p[1]);
00120     wait();
00121     wait();
00122     break;
00123
00124 case BACK:
00125     bcm = (struct backcmd*)cmd;
00126     if(fork1() == 0)
00127         runcmd(bcmd->cmd);
00128     break;
00129 }
00130 exit();
00131 }
00132
00133 int
00134 getcmd(char *buf, int nbuf)
00135 {
00136     printf(2, "$ ");
00137     memset(buf, 0, nbuf);
00138     gets(buf, nbuf);
00139     if(buf[0] == 0) // EOF
00140         return -1;
00141     return 0;
00142 }
00143
00144 int
00145 main(void)
00146 {
00147     static char buf[100];
00148     int fd;
00149
00150     // Ensure that three file descriptors are open.
00151     while((fd = open("console", O_RDWR)) >= 0){
00152         if(fd >= 3){
00153             close(fd);
00154             break;
00155         }
00156     }
00157
00158     // Read and run input commands.
00159     while(getcmd(buf, sizeof(buf)) >= 0){
00160         if(buf[0] == 'c' && buf[1] == 'd' && buf[2] == ' '){
00161             // Chdir must be called by the parent, not the child.
00162             buf[strlen(buf)-1] = 0; // chop \n
00163             if(chdir(buf+3) < 0)
00164                 printf(2, "cannot cd %s\n", buf+3);
00165             continue;
00166         }
00167         if(fork1() == 0)
00168             runcmd(parsecmd(buf));
00169         wait();

```

```
00170     }
00171     exit();
00172 }
00173
00174 void
00175 panic(char *s)
00176 {
00177     printf(2, "%s\n", s);
00178     exit();
00179 }
00180
00181 int
00182 fork1(void)
00183 {
00184     int pid;
00185
00186     pid = fork();
00187     if(pid == -1)
00188         panic("fork");
00189     return pid;
00190 }
00191
00192 //PAGEBREAK!
00193 // Constructors
00194
00195 struct cmd*
00196 execcmd(void)
00197 {
00198     struct execcmd *cmd;
00199
00200     cmd = malloc(sizeof(*cmd));
00201     memset(cmd, 0, sizeof(*cmd));
00202     cmd->type = EXEC;
00203     return (struct cmd*)cmd;
00204 }
00205
00206 struct cmd*
00207 redircmd(struct cmd *subcmd, char *file, char *efile, int mode, int fd)
00208 {
00209     struct redircmd *cmd;
00210
00211     cmd = malloc(sizeof(*cmd));
00212     memset(cmd, 0, sizeof(*cmd));
00213     cmd->type = REDIR;
00214     cmd->cmd = subcmd;
00215     cmd->file = file;
00216     cmd->efile = efile;
00217     cmd->mode = mode;
00218     cmd->fd = fd;
00219     return (struct cmd*)cmd;
00220 }
00221
00222 struct cmd*
00223 pipecmd(struct cmd *left, struct cmd *right)
00224 {
00225     struct pipecmd *cmd;
00226
00227     cmd = malloc(sizeof(*cmd));
00228     memset(cmd, 0, sizeof(*cmd));
00229     cmd->type = PIPE;
00230     cmd->left = left;
00231     cmd->right = right;
00232     return (struct cmd*)cmd;
00233 }
00234
00235 struct cmd*
00236 listcmd(struct cmd *left, struct cmd *right)
00237 {
00238     struct listcmd *cmd;
00239
00240     cmd = malloc(sizeof(*cmd));
00241     memset(cmd, 0, sizeof(*cmd));
00242     cmd->type = LIST;
00243     cmd->left = left;
00244     cmd->right = right;
00245     return (struct cmd*)cmd;
00246 }
00247
00248 struct cmd*
00249 backcmd(struct cmd *subcmd)
00250 {
00251     struct backcmd *cmd;
00252
00253     cmd = malloc(sizeof(*cmd));
00254     memset(cmd, 0, sizeof(*cmd));
00255     cmd->type = BACK;
00256     cmd->cmd = subcmd;
```

```

00257     return (struct cmd*)cmd;
00258 }
00259 //PAGEBREAK!
00260 // Parsing
00261
00262 char whitespace[] = " \t\r\n\v";
00263 char symbols[] = "<|>&; ()";
00264
00265 int
00266 gettoken(char **ps, char *es, char **q, char **eq)
00267 {
00268     char *s;
00269     int ret;
00270
00271     s = *ps;
00272     while(s < es && strchr(whitespace, *s))
00273         s++;
00274     if(*s)
00275         *q = s;
00276     ret = *s;
00277     switch(*s){
00278     case 0:
00279         break;
00280     case '|':
00281     case '(':
00282     case ')':
00283     case ';':
00284     case '&':
00285     case '<':
00286         s++;
00287         break;
00288     case '>':
00289         s++;
00290         if(*s == '>'){
00291             ret = '+';
00292             s++;
00293         }
00294         break;
00295     default:
00296         ret = 'a';
00297         while(s < es && !strchr(whitespace, *s) && !strchr(symbols, *s))
00298             s++;
00299         break;
00300     }
00301     if(eq)
00302         *eq = s;
00303
00304     while(s < es && strchr(whitespace, *s))
00305         s++;
00306     *ps = s;
00307     return ret;
00308 }
00309
00310 int
00311 peek(char **ps, char *es, char *toks)
00312 {
00313     char *s;
00314
00315     s = *ps;
00316     while(s < es && strchr(whitespace, *s))
00317         s++;
00318     *ps = s;
00319     return *s && strchr(toks, *s);
00320 }
00321
00322 struct cmd *parseline(char**, char*);
00323 struct cmd *parsepipe(char**, char*);
00324 struct cmd *parseexec(char**, char*);
00325 struct cmd *nulterminate(struct cmd*);
00326
00327 struct cmd*
00328 parsecmd(char *s)
00329 {
00330     char *es;
00331     struct cmd *cmd;
00332
00333     es = s + strlen(s);
00334     cmd = parseline(&s, es);
00335     peek(&s, es, "");
00336     if(s != es){
00337         printf(2, "leftovers: %s\n", s);
00338         panic("syntax");
00339     }
00340     nulterminate(cmd);
00341     return cmd;
00342 }
00343

```

```

00344 struct cmd*
00345 parseline(char **ps, char *es)
00346 {
00347     struct cmd *cmd;
00348
00349     cmd = parsepipe(ps, es);
00350     while(peek(ps, es, "&")){
00351         gettoken(ps, es, 0, 0);
00352         cmd = backcmd(cmd);
00353     }
00354     if(peek(ps, es, ";")){
00355         gettoken(ps, es, 0, 0);
00356         cmd = listcmd(cmd, parseline(ps, es));
00357     }
00358     return cmd;
00359 }
00360
00361 struct cmd*
00362 parsepipe(char **ps, char *es)
00363 {
00364     struct cmd *cmd;
00365
00366     cmd = parseexec(ps, es);
00367     if(peek(ps, es, "|")){
00368         gettoken(ps, es, 0, 0);
00369         cmd = pipecmd(cmd, parsepipe(ps, es));
00370     }
00371     return cmd;
00372 }
00373
00374 struct cmd*
00375 parseredirs(struct cmd *cmd, char **ps, char *es)
00376 {
00377     int tok;
00378     char *q, *eq;
00379
00380     while(peek(ps, es, "<>")){
00381         tok = gettoken(ps, es, 0, 0);
00382         if(gettoken(ps, es, &q, &eq) != 'a')
00383             panic("missing file for redirection");
00384         switch(tok){
00385             case '<':
00386                 cmd = redircmd(cmd, q, eq, O_RDONLY, 0);
00387                 break;
00388             case '>':
00389                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
00390                 break;
00391             case '+': // »
00392                 cmd = redircmd(cmd, q, eq, O_WRONLY|O_CREATE, 1);
00393                 break;
00394         }
00395     }
00396     return cmd;
00397 }
00398
00399 struct cmd*
00400 parseblock(char **ps, char *es)
00401 {
00402     struct cmd *cmd;
00403
00404     if(!peek(ps, es, "("))
00405         panic("parseblock");
00406     gettoken(ps, es, 0, 0);
00407     cmd = parseline(ps, es);
00408     if(!peek(ps, es, ")"))
00409         panic("syntax - missing )");
00410     gettoken(ps, es, 0, 0);
00411     cmd = parseredirs(cmd, ps, es);
00412     return cmd;
00413 }
00414
00415 struct cmd*
00416 parseexec(char **ps, char *es)
00417 {
00418     char *q, *eq;
00419     int tok, argc;
00420     struct execcmd *cmd;
00421     struct cmd *ret;
00422
00423     if(peek(ps, es, "("))
00424         return parseblock(ps, es);
00425
00426     ret = execcmd();
00427     cmd = (struct execcmd*)ret;
00428
00429     argc = 0;
00430     ret = parseredirs(ret, ps, es);

```

```

00431 while(!peek(ps, es, "|)&;")){
00432     if((tok=gettoken(ps, es, &q, &eq)) == 0)
00433         break;
00434     if(tok != 'a')
00435         panic("syntax");
00436     cmd->argv[argc] = q;
00437     cmd->eargv[argc] = eq;
00438     argc++;
00439     if(argc >= MAXARGS)
00440         panic("too many args");
00441     ret = parseredirs(ret, ps, es);
00442 }
00443 cmd->argv[argc] = 0;
00444 cmd->eargv[argc] = 0;
00445 return ret;
00446 }
00447
00448 // NUL-terminate all the counted strings.
00449 struct cmd*
00450 nulterminate(struct cmd *cmd)
00451 {
00452     int i;
00453     struct backcmd *bcmd;
00454     struct execcmd *ecmd;
00455     struct listcmd *lcmd;
00456     struct pipecmd *pcmd;
00457     struct redircmd *rcmd;
00458
00459     if(cmd == 0)
00460         return 0;
00461
00462     switch(cmd->type){
00463     case EXEC:
00464         ecmd = (struct execcmd*)cmd;
00465         for(i=0; ecmd->argv[i]; i++)
00466             *ecmd->eargv[i] = 0;
00467         break;
00468
00469     case REDIR:
00470         rcmd = (struct redircmd*)cmd;
00471         nulterminate(rcmd->cmd);
00472         *rcmd->efile = 0;
00473         break;
00474
00475     case PIPE:
00476         pcmd = (struct pipecmd*)cmd;
00477         nulterminate(pcmd->left);
00478         nulterminate(pcmd->right);
00479         break;
00480
00481     case LIST:
00482         lcmd = (struct listcmd*)cmd;
00483         nulterminate(lcmd->left);
00484         nulterminate(lcmd->right);
00485         break;
00486
00487     case BACK:
00488         bcmd = (struct backcmd*)cmd;
00489         nulterminate(bcmd->cmd);
00490         break;
00491     }
00492     return cmd;
00493 }

```

5.174 sh.d File Reference

5.175 sh.d

[Go to the documentation of this file.](#)

```
00001 sh.o: sh.c /usr/include/stdc-predef.h types.h user.h fcntl.h
```

5.176 shutdown.c File Reference

```

#include "types.h"
#include "stat.h"

```



```
#include "user.h"
```

Functions

- int [main](#) (int argc, char *[argv](#)[])

5.176.1 Function Documentation

5.176.1.1 main()

```
int main (  
    int argc,  
    char * argv[] )
```

Definition at line 10 of file [shutdown.c](#).

```
00011 {  
00012     halt();  
00013     exit();  
00014 }
```

5.177 shutdown.c

[Go to the documentation of this file.](#)

```
00001 // Shuts down the system by using the halt() system call  
00002 // to send a special signal to QEMU.  
00003 // Added by Bill Katsak  
00004 // Copied from: http://pdos.csail.mit.edu/6.828/2012/homework/xv6-syscall.html  
00005  
00006 #include "types.h"  
00007 #include "stat.h"  
00008 #include "user.h"  
00009  
00010 int main(int argc, char *argv[])  
00011 {  
00012     halt();  
00013     exit();  
00014 }
```

5.178 shutdown.d File Reference

5.179 shutdown.d

[Go to the documentation of this file.](#)

```
00001 shutdown.o: shutdown.c /usr/include/stdc-predef.h types.h stat.h user.h
```

5.180 sleeplock.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "x86.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "spinlock.h"
#include "sleeplock.h"
```

Functions

- void [acquiresleep](#) (struct [sleeplock](#) *lk)
- int [holdingsleep](#) (struct [sleeplock](#) *lk)
- void [initsleeplock](#) (struct [sleeplock](#) *lk, char *name)
- void [releasesleep](#) (struct [sleeplock](#) *lk)

5.180.1 Function Documentation

5.180.1.1 [acquiresleep\(\)](#)

```
void acquiresleep (
    struct sleeplock * lk )
```

Definition at line 23 of file [sleeplock.c](#).

```
00024 {
00025     acquire(&lk->lk);
00026     while (lk->locked) {
00027         sleep(lk, &lk->lk);
00028     }
00029     lk->locked = 1;
00030     lk->pid = myproc()->pid;
00031     release(&lk->lk);
00032 }
```

Referenced by [bget\(\)](#), [ilock\(\)](#), and [iput\(\)](#).

5.180.1.2 [holdingsleep\(\)](#)

```
int holdingsleep (
    struct sleeplock * lk )
```

Definition at line 45 of file [sleeplock.c](#).

```
00046 {
00047     int r;
00048
00049     acquire(&lk->lk);
00050     r = lk->locked && (lk->pid == myproc()->pid);
00051     release(&lk->lk);
00052     return r;
00053 }
```

Referenced by [brelse\(\)](#), [bwrite\(\)](#), [iderw\(\)](#), and [iunlock\(\)](#).

5.180.1.3 initsleeplock()

```
void initsleeplock (
    struct sleeplock * lk,
    char * name )
```

Definition at line 14 of file [sleeplock.c](#).

```
00015 {
00016     initlock(&lk->lk, "sleep lock");
00017     lk->name = name;
00018     lk->locked = 0;
00019     lk->pid = 0;
00020 }
```

Referenced by [binit\(\)](#), and [iinit\(\)](#).

5.180.1.4 releasesleep()

```
void releasesleep (
    struct sleeplock * lk )
```

Definition at line 35 of file [sleeplock.c](#).

```
00036 {
00037     acquire(&lk->lk);
00038     lk->locked = 0;
00039     lk->pid = 0;
00040     wakeup(lk);
00041     release(&lk->lk);
00042 }
```

Referenced by [brelse\(\)](#), [iput\(\)](#), and [iunlock\(\)](#).

5.181 sleeplock.c

[Go to the documentation of this file.](#)

```
00001 // Sleeping locks
00002
00003 #include "types.h"
00004 #include "defs.h"
00005 #include "param.h"
00006 #include "x86.h"
00007 #include "memlayout.h"
00008 #include "mmu.h"
00009 #include "proc.h"
00010 #include "spinlock.h"
00011 #include "sleeplock.h"
00012
00013 void
00014 initsleeplock(struct sleeplock *lk, char *name)
00015 {
00016     initlock(&lk->lk, "sleep lock");
00017     lk->name = name;
00018     lk->locked = 0;
00019     lk->pid = 0;
00020 }
00021
00022 void
00023 acquiresleep(struct sleeplock *lk)
00024 {
00025     acquire(&lk->lk);
00026     while (lk->locked) {
00027         sleep(lk, &lk->lk);
00028     }
00029     lk->locked = 1;
00030     lk->pid = myproc()->pid;
00031     release(&lk->lk);
00032 }
```

```

00033
00034 void
00035 releasesleep(struct sleeplock *lk)
00036 {
00037     acquire(&lk->lk);
00038     lk->locked = 0;
00039     lk->pid = 0;
00040     wakeup(lk);
00041     release(&lk->lk);
00042 }
00043
00044 int
00045 holdingsleep(struct sleeplock *lk)
00046 {
00047     int r;
00048
00049     acquire(&lk->lk);
00050     r = lk->locked && (lk->pid == myproc()->pid);
00051     release(&lk->lk);
00052     return r;
00053 }
00054
00055
00056

```

5.182 sleeplock.d File Reference

5.183 sleeplock.d

[Go to the documentation of this file.](#)

```

00001 sleeplock.o: sleeplock.c /usr/include/stdc-predef.h types.h defs.h \
00002 param.h x86.h memlayout.h mmu.h proc.h spinlock.h sleeplock.h

```

5.184 sleeplock.h File Reference

Classes

- struct [sleeplock](#)

5.185 sleeplock.h

[Go to the documentation of this file.](#)

```

00001 // Long-term locks for processes
00002 struct sleeplock {
00003     uint locked;           // Is the lock held?
00004     struct spinlock lk;    // spinlock protecting this sleep lock
00005
00006     // For debugging:
00007     char *name;            // Name of lock.
00008     int pid;               // Process holding lock
00009 };
00010

```

5.186 spinlock.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "x86.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "spinlock.h"

```

Functions

- void [acquire](#) (struct [spinlock](#) *lk)
- void [getcallepcs](#) (void *v, uint pcs[])
- int [holding](#) (struct [spinlock](#) *lock)
- void [initlock](#) (struct [spinlock](#) *lk, char *name)
- void [popcli](#) (void)
- void [pushcli](#) (void)
- void [release](#) (struct [spinlock](#) *lk)

5.186.1 Function Documentation

5.186.1.1 [acquire\(\)](#)

```
void acquire (
    struct spinlock * lk )
```

Definition at line 25 of file [spinlock.c](#).

```
00026 {
00027     pushcli(); // disable interrupts to avoid deadlock.
00028     if(holding(lk))
00029         panic("acquire");
00030
00031     // The xchg is atomic.
00032     while(xchg(&lk->locked, 1) != 0)
00033         ;
00034
00035     // Tell the C compiler and the processor to not move loads or stores
00036     // past this point, to ensure that the critical section's memory
00037     // references happen after the lock is acquired.
00038     __sync_synchronize();
00039
00040     // Record info about lock acquisition for debugging.
00041     lk->cpu = mycpu();
00042     getcallepcs(&lk, lk->pcs);
00043 }
```

Referenced by [acquiresleep\(\)](#), [allocproc\(\)](#), [begin_op\(\)](#), [bget\(\)](#), [brelse\(\)](#), [chpr\(\)](#), [consoleintr\(\)](#), [consoleread\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), [cps\(\)](#), [end_op\(\)](#), [exit\(\)](#), [filealloc\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fork\(\)](#), [holdingsleep\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idup\(\)](#), [iget\(\)](#), [iput\(\)](#), [kalloc\(\)](#), [kfree\(\)](#), [kill\(\)](#), [log_write\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup\(\)](#), and [yield\(\)](#).

5.186.1.2 [getcallepcs\(\)](#)

```
void getcallepcs (
    void * v,
    uint pcs[] )
```

Definition at line 72 of file [spinlock.c](#).

```
00073 {
00074     uint *ebp;
00075     int i;
00076
00077     ebp = (uint*)v - 2;
00078     for(i = 0; i < 10; i++){
00079         if(ebp == 0 || ebp < (uint*)KERNBASE || ebp == (uint*)0xffffffff)
00080             break;
00081         pcs[i] = ebp[1]; // saved %eip
00082         ebp = (uint*)ebp[0]; // saved %ebp
00083     }
00084     for(; i < 10; i++)
00085         pcs[i] = 0;
00086 }
```

Referenced by [acquire\(\)](#).

5.186.1.3 holding()

```
int holding (
    struct spinlock * lock )
```

Definition at line 90 of file [spinlock.c](#).

```
00091 {
00092     int r;
00093     pushcli();
00094     r = lock->locked && lock->cpu == mycpu();
00095     popcli();
00096     return r;
00097 }
```

Referenced by [acquire\(\)](#), [release\(\)](#), and [sched\(\)](#).

5.186.1.4 initlock()

```
void initlock (
    struct spinlock * lk,
    char * name )
```

Definition at line 13 of file [spinlock.c](#).

```
00014 {
00015     lk->name = name;
00016     lk->locked = 0;
00017     lk->cpu = 0;
00018 }
```

Referenced by [binit\(\)](#), [consoleinit\(\)](#), [fileinit\(\)](#), [ideinit\(\)](#), [iinit\(\)](#), [initlog\(\)](#), [initsleeplock\(\)](#), [kinit1\(\)](#), [pinit\(\)](#), [pipealloc\(\)](#), and [tvinit\(\)](#).

5.186.1.5 popcli()

```
void popcli (
    void )
```

Definition at line 117 of file [spinlock.c](#).

```
00118 {
00119     if(readeflags() & FL\_IF)
00120         panic("popcli - interruptible");
00121     if(--mycpu()->ncli < 0)
00122         panic("popcli");
00123     if(mycpu()->ncli == 0 && mycpu()->intena)
00124         sti();
00125 }
```

Referenced by [holding\(\)](#), [myproc\(\)](#), [release\(\)](#), and [switchvm\(\)](#).

5.186.1.6 pushcli()

```
void pushcli (
    void )
```

Definition at line 105 of file [spinlock.c](#).

```
00106 {
00107     int eflags;
00108
00109     eflags = readeflags();
00110     cli();
00111     if(mycpu()->ncli == 0)
00112         mycpu()->intena = eflags & FL_IF;
00113     mycpu()->ncli += 1;
00114 }
```

Referenced by [acquire\(\)](#), [holding\(\)](#), [myproc\(\)](#), and [switchvm\(\)](#).

5.186.1.7 release()

```
void release (
    struct spinlock * lk )
```

Definition at line 47 of file [spinlock.c](#).

```
00048 {
00049     if(!holding(lk))
00050         panic("release");
00051
00052     lk->pcs[0] = 0;
00053     lk->cpu = 0;
00054
00055     // Tell the C compiler and the processor to not move loads or stores
00056     // past this point, to ensure that all the stores in the critical
00057     // section are visible to other cores before the lock is released.
00058     // Both the C compiler and the hardware may re-order loads and
00059     // stores; __sync_synchronize() tells them both not to.
00060     __sync_synchronize();
00061
00062     // Release the lock, equivalent to lk->locked = 0.
00063     // This code can't use a C assignment, since it might
00064     // not be atomic. A real OS would use C atomics here.
00065     asm volatile("movl $0, %0" : "+m" (lk->locked) : );
00066
00067     popcli();
00068 }
```

Referenced by [acquiresleep\(\)](#), [allocproc\(\)](#), [begin_op\(\)](#), [bget\(\)](#), [brelse\(\)](#), [chpr\(\)](#), [consoleintr\(\)](#), [consoleread\(\)](#), [consolewrite\(\)](#), [cprintf\(\)](#), [cps\(\)](#), [end_op\(\)](#), [filealloc\(\)](#), [fileclose\(\)](#), [filedup\(\)](#), [fork\(\)](#), [forkret\(\)](#), [holdingsleep\(\)](#), [ideintr\(\)](#), [iderw\(\)](#), [idup\(\)](#), [iget\(\)](#), [iput\(\)](#), [kalloc\(\)](#), [kfree\(\)](#), [kill\(\)](#), [log_write\(\)](#), [pipeclose\(\)](#), [piperead\(\)](#), [pipewrite\(\)](#), [releasesleep\(\)](#), [scheduler\(\)](#), [sleep\(\)](#), [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), [userinit\(\)](#), [wait\(\)](#), [wakeup\(\)](#), and [yield\(\)](#).

5.187 spinlock.c

[Go to the documentation of this file.](#)

```
00001 // Mutual exclusion spin locks.
00002
00003 #include "types.h"
00004 #include "defs.h"
00005 #include "param.h"
00006 #include "x86.h"
00007 #include "memlayout.h"
00008 #include "mmu.h"
00009 #include "proc.h"
00010 #include "spinlock.h"
00011
00012 void
```

```

00013 initlock(struct spinlock *lk, char *name)
00014 {
00015     lk->name = name;
00016     lk->locked = 0;
00017     lk->cpu = 0;
00018 }
00019
00020 // Acquire the lock.
00021 // Loops (spins) until the lock is acquired.
00022 // Holding a lock for a long time may cause
00023 // other CPUs to waste time spinning to acquire it.
00024 void
00025 acquire(struct spinlock *lk)
00026 {
00027     pushcli(); // disable interrupts to avoid deadlock.
00028     if(holding(lk))
00029         panic("acquire");
00030
00031     // The xchg is atomic.
00032     while(xchg(&lk->locked, 1) != 0)
00033         ;
00034
00035     // Tell the C compiler and the processor to not move loads or stores
00036     // past this point, to ensure that the critical section's memory
00037     // references happen after the lock is acquired.
00038     __sync_synchronize();
00039
00040     // Record info about lock acquisition for debugging.
00041     lk->cpu = mycpu();
00042     getcallerpcs(&lk, lk->pcs);
00043 }
00044
00045 // Release the lock.
00046 void
00047 release(struct spinlock *lk)
00048 {
00049     if(!holding(lk))
00050         panic("release");
00051
00052     lk->pcs[0] = 0;
00053     lk->cpu = 0;
00054
00055     // Tell the C compiler and the processor to not move loads or stores
00056     // past this point, to ensure that all the stores in the critical
00057     // section are visible to other cores before the lock is released.
00058     // Both the C compiler and the hardware may re-order loads and
00059     // stores; __sync_synchronize() tells them both not to.
00060     __sync_synchronize();
00061
00062     // Release the lock, equivalent to lk->locked = 0.
00063     // This code can't use a C assignment, since it might
00064     // not be atomic. A real OS would use C atomics here.
00065     asm volatile("movl $0, %0" : "+m" (lk->locked) : );
00066
00067     popcli();
00068 }
00069
00070 // Record the current call stack in pcs[] by following the %ebp chain.
00071 void
00072 getcallerpcs(void *v, uint pcs[])
00073 {
00074     uint *ebp;
00075     int i;
00076
00077     ebp = (uint*)v - 2;
00078     for(i = 0; i < 10; i++){
00079         if(ebp == 0 || ebp < (uint*)KERNBASE || ebp == (uint*)0xffffffff)
00080             break;
00081         pcs[i] = ebp[1]; // saved %eip
00082         ebp = (uint*)ebp[0]; // saved %ebp
00083     }
00084     for(; i < 10; i++)
00085         pcs[i] = 0;
00086 }
00087
00088 // Check whether this cpu is holding the lock.
00089 int
00090 holding(struct spinlock *lock)
00091 {
00092     int r;
00093     pushcli();
00094     r = lock->locked && lock->cpu == mycpu();
00095     popcli();
00096     return r;
00097 }
00098
00099

```



```

00100 // Pushcli/popcli are like cli/sti except that they are matched:
00101 // it takes two popcli to undo two pushcli. Also, if interrupts
00102 // are off, then pushcli, popcli leaves them off.
00103
00104 void
00105 pushcli(void)
00106 {
00107     int eflags;
00108
00109     eflags = readeflags();
00110     cli();
00111     if(mycpu()->ncli == 0)
00112         mycpu()->intena = eflags & FL_IF;
00113     mycpu()->ncli += 1;
00114 }
00115
00116 void
00117 popcli(void)
00118 {
00119     if(readeflags() & FL_IF)
00120         panic("popcli - interruptible");
00121     if(--mycpu()->ncli < 0)
00122         panic("popcli");
00123     if(mycpu()->ncli == 0 && mycpu()->intena)
00124         sti();
00125 }
00126

```

5.188 spinlock.d File Reference

5.189 spinlock.d

[Go to the documentation of this file.](#)

```

00001 spinlock.o: spinlock.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 x86.h memlayout.h mmu.h proc.h spinlock.h

```

5.190 spinlock.h File Reference

Classes

- struct [spinlock](#)

5.191 spinlock.h

[Go to the documentation of this file.](#)

```

00001 // Mutual exclusion lock.
00002 struct spinlock {
00003     uint locked;           // Is the lock held?
00004
00005     // For debugging:
00006     char *name;            // Name of lock.
00007     struct cpu *cpu;       // The cpu holding the lock.
00008     uint pcs[10];          // The call stack (an array of program counters)
00009                             // that locked the lock.
00010 };
00011

```

5.192 stat.h File Reference

Classes

- struct [stat](#)

Macros

- `#define T_DEV 3`
- `#define T_DIR 1`
- `#define T_FILE 2`

5.192.1 Macro Definition Documentation

5.192.1.1 T_DEV

```
#define T_DEV 3
```

Definition at line 3 of file [stat.h](#).

5.192.1.2 T_DIR

```
#define T_DIR 1
```

Definition at line 1 of file [stat.h](#).

5.192.1.3 T_FILE

```
#define T_FILE 2
```

Definition at line 2 of file [stat.h](#).

5.193 stat.h

[Go to the documentation of this file.](#)

```
00001 #define T_DIR 1 // Directory
00002 #define T_FILE 2 // File
00003 #define T_DEV 3 // Device
00004
00005 struct stat {
00006     short type; // Type of file
00007     int dev; // File system's disk device
00008     uint ino; // Inode number
00009     short nlink; // Number of links to file
00010     uint size; // Size of file in bytes
00011 };
```

5.194 stressfs.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fs.h"
#include "fcntl.h"
```

Functions

- int [main](#) (int argc, char *argv[])

5.194.1 Function Documentation

5.194.1.1 main()

```
int main (
    int argc,
    char * argv[] )
```

Definition at line 17 of file [stressfs.c](#).

```
00018 {
00019     int fd, i;
00020     char path[] = "stressfs0";
00021     char data[512];
00022
00023     printf(1, "stressfs starting\n");
00024     memset(data, 'a', sizeof(data));
00025
00026     for(i = 0; i < 4; i++)
00027         if(fork() > 0)
00028             break;
00029
00030     printf(1, "write %d\n", i);
00031
00032     path[8] += i;
00033     fd = open(path, O_CREATE | O_RDWR);
00034     for(i = 0; i < 20; i++)
00035         // printf(fd, "%d\n", i);
00036         write(fd, data, sizeof(data));
00037     close(fd);
00038
00039     printf(1, "read\n");
00040
00041     fd = open(path, O_RDONLY);
00042     for (i = 0; i < 20; i++)
00043         read(fd, data, sizeof(data));
00044     close(fd);
00045
00046     wait();
00047
00048     exit();
00049 }
```

5.195 stressfs.c

[Go to the documentation of this file.](#)

```
00001 // Demonstrate that moving the "acquire" in iderw after the loop that
00002 // appends to the idequeue results in a race.
00003
00004 // For this to work, you should also add a spin within iderw's
00005 // idequeue traversal loop. Adding the following demonstrated a panic
00006 // after about 5 runs of stressfs in QEMU on a 2.1GHz CPU:
00007 //     for (i = 0; i < 40000; i++)
00008 //         asm volatile("");
00009
00010 #include "types.h"
00011 #include "stat.h"
00012 #include "user.h"
00013 #include "fs.h"
00014 #include "fcntl.h"
00015
00016 int
00017 main(int argc, char *argv[])
00018 {
00019     int fd, i;
00020     char path[] = "stressfs0";
00021     char data[512];
00022
00023     printf(1, "stressfs starting\n");
00024     memset(data, 'a', sizeof(data));
00025
00026     for(i = 0; i < 4; i++)
00027         if(fork() > 0)
00028             break;
00029
00030     printf(1, "write %d\n", i);
00031
00032     path[8] += i;
00033     fd = open(path, O_CREATE | O_RDWR);
00034     for(i = 0; i < 20; i++)
00035         // printf(fd, "%d\n", i);
00036         write(fd, data, sizeof(data));
00037     close(fd);
00038
00039     printf(1, "read\n");
00040
00041     fd = open(path, O_RDONLY);
00042     for (i = 0; i < 20; i++)
00043         read(fd, data, sizeof(data));
00044     close(fd);
00045
00046     wait();
00047
00048     exit();
00049 }
```

5.196 stressfs.d File Reference

5.197 stressfs.d

[Go to the documentation of this file.](#)

```
00001 stressfs.o: stressfs.c /usr/include/stdc-predef.h types.h stat.h user.h \
00002 fs.h fcntl.h
```

5.198 string.c File Reference

```
#include "types.h"
#include "x86.h"
```

Functions

- int [memcmp](#) (const void *v1, const void *v2, [uint](#) n)
- void * [memcpy](#) (void *dst, const void *src, [uint](#) n)
- void * [memmove](#) (void *dst, const void *src, [uint](#) n)
- void * [memset](#) (void *dst, int c, [uint](#) n)
- char * [safestrcpy](#) (char *s, const char *t, int n)
- int [strlen](#) (const char *s)
- int [strncmp](#) (const char *p, const char *q, [uint](#) n)
- char * [strncpy](#) (char *s, const char *t, int n)

5.198.1 Function Documentation

5.198.1.1 memcmp()

```
int memcmp (
    const void * v1,
    const void * v2,
    uint n )
```

Definition at line 16 of file [string.c](#).

```
00017 {
00018     const uchar *s1, *s2;
00019
00020     s1 = v1;
00021     s2 = v2;
00022     while (n-- > 0) {
00023         if (*s1 != *s2)
00024             return *s1 - *s2;
00025         s1++, s2++;
00026     }
00027
00028     return 0;
00029 }
```

Referenced by [cmostime\(\)](#), [mpconfig\(\)](#), and [mpsearch1\(\)](#).

5.198.1.2 memcpy()

```
void * memcpy (
    void * dst,
    const void * src,
    uint n )
```

Definition at line 53 of file [string.c](#).

```
00054 {
00055     return memmove(dst, src, n);
00056 }
```

5.198.1.3 memmove()

```
void * memmove (
    void * dst,
    const void * src,
    uint n )
```

Definition at line 32 of file [string.c](#).

```
00033 {
00034     const char *s;
00035     char *d;
00036
00037     s = src;
00038     d = dst;
00039     if(s < d && s + n > d){
00040         s += n;
00041         d += n;
00042         while(n-- > 0)
00043             *--d = *--s;
00044     } else
00045         while(n-- > 0)
00046             *d++ = *s++;
00047     return dst;
00048 }
00049 }
```

Referenced by [cgaputc\(\)](#), [copyout\(\)](#), [copyuvm\(\)](#), [fmtname\(\)](#), [grep\(\)](#), [iderw\(\)](#), [ilock\(\)](#), [inituvm\(\)](#), [install_trans\(\)](#), [iupdate\(\)](#), [ls\(\)](#), [main\(\)](#), [memcpy\(\)](#), [readi\(\)](#), [readsb\(\)](#), [skipelem\(\)](#), [startothers\(\)](#), [write_log\(\)](#), and [writei\(\)](#).

5.198.1.4 memset()

```
void * memset (
    void * dst,
    int c,
    uint n )
```

Definition at line 5 of file [string.c](#).

```
00006 {
00007     if ((int)dst%4 == 0 && n%4 == 0){
00008         c &= 0xFF;
00009         stosl(dst, (c<<24) | (c<<16) | (c<<8) | c, n/4);
00010     } else
00011         stosb(dst, c, n);
00012     return dst;
00013 }
```

5.198.1.5 safestrcpy()

```
char * safestrcpy (
    char * s,
    const char * t,
    int n )
```

Definition at line 83 of file [string.c](#).

```
00084 {
00085     char *os;
00086
00087     os = s;
00088     if(n <= 0)
00089         return os;
00090     while(--n > 0 && (*s++ = *t++) != 0)
00091         ;
00092     *s = 0;
00093     return os;
00094 }
```

Referenced by [exec\(\)](#), [fork\(\)](#), and [userinit\(\)](#).

5.198.1.6 strlen()

```
int strlen (
    const char * s )
```

Definition at line 97 of file [string.c](#).

```
00098 {
00099     int n;
00100
00101     for(n = 0; s[n]; n++)
00102         ;
00103     return n;
00104 }
```

5.198.1.7 strncmp()

```
int strncmp (
    const char * p,
    const char * q,
    uint n )
```

Definition at line 59 of file [string.c](#).

```
00060 {
00061     while(n > 0 && *p && *p == *q)
00062         n--, p++, q++;
00063     if(n == 0)
00064         return 0;
00065     return (uchar)*p - (uchar)*q;
00066 }
```

Referenced by [namecmp\(\)](#).

5.198.1.8 strncpy()

```
char * strncpy (
    char * s,
    const char * t,
    int n )
```

Definition at line 69 of file [string.c](#).

```
00070 {
00071     char *os;
00072
00073     os = s;
00074     while(n-- > 0 && (*s++ = *t++) != 0)
00075         ;
00076     while(n-- > 0)
00077         *s++ = 0;
00078     return os;
00079 }
```

Referenced by [dirlink\(\)](#), [main\(\)](#), and [sys_getprocs\(\)](#).

5.199 string.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "x86.h"
00003
00004 void*
00005 memset(void *dst, int c, uint n)
00006 {
00007     if ((int)dst%4 == 0 && n%4 == 0){
00008         c &= 0xFF;
00009         stosl(dst, (c<<24)|(c<<16)|(c<<8)|c, n/4);
00010     } else
00011         stosb(dst, c, n);
00012     return dst;
00013 }
00014
00015 int
00016 memcmp(const void *v1, const void *v2, uint n)
00017 {
00018     const uchar *s1, *s2;
00019
00020     s1 = v1;
00021     s2 = v2;
00022     while(n-- > 0){
00023         if(*s1 != *s2)
00024             return *s1 - *s2;
00025         s1++, s2++;
00026     }
00027
00028     return 0;
00029 }
00030
00031 void*
00032 memmove(void *dst, const void *src, uint n)
00033 {
00034     const char *s;
00035     char *d;
00036
00037     s = src;
00038     d = dst;
00039     if(s < d && s + n > d){
00040         s += n;
00041         d += n;
00042         while(n-- > 0)
00043             *--d = *--s;
00044     } else
00045         while(n-- > 0)
00046             *d++ = *s++;
00047
00048     return dst;
00049 }
00050
00051 // memcpy exists to placate GCC. Use memmove.
00052 void*
00053 memcpy(void *dst, const void *src, uint n)
00054 {
00055     return memmove(dst, src, n);
00056 }
00057
00058 int
00059 strncmp(const char *p, const char *q, uint n)
00060 {
00061     while(n > 0 && *p && *p == *q)
00062         n--, p++, q++;
00063     if(n == 0)
00064         return 0;
00065     return (uchar)*p - (uchar)*q;
00066 }
00067
00068 char*
00069 strncpy(char *s, const char *t, int n)
00070 {
00071     char *os;
00072
00073     os = s;
00074     while(n-- > 0 && (*s++ = *t++) != 0)
00075         ;
00076     while(n-- > 0)
00077         *s++ = 0;
00078     return os;
00079 }
00080
00081 // Like strncpy but guaranteed to NUL-terminate.
00082 char*
```



```

00083 safestrcpy(char *s, const char *t, int n)
00084 {
00085     char *os;
00086
00087     os = s;
00088     if(n <= 0)
00089         return os;
00090     while(--n > 0 && (*s++ = *t++) != 0)
00091         ;
00092     *s = 0;
00093     return os;
00094 }
00095
00096 int
00097 strlen(const char *s)
00098 {
00099     int n;
00100
00101     for(n = 0; s[n]; n++)
00102         ;
00103     return n;
00104 }
00105

```

5.200 string.d File Reference

5.201 string.d

[Go to the documentation of this file.](#)

```
00001 string.o: string.c /usr/include/stdc-predef.h types.h x86.h
```

5.202 syscall.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"
#include "syscall.h"

```

Functions

- int [argint](#) (int n, int *ip)
- int [argptr](#) (int n, char **pp, int size)
- int [argstr](#) (int n, char **pp)
- int [fetchint](#) (uint addr, int *ip)
- int [fetchstr](#) (uint addr, char **pp)
- int [sys_chdir](#) (void)
- int [sys_chpr](#) (void)
- int [sys_close](#) (void)
- int [sys_cps](#) (void)
- int [sys_dup](#) (void)
- int [sys_exec](#) (void)
- int [sys_exit](#) (void)

- int [sys_fork](#) (void)
- int [sys_fstat](#) (void)
- int [sys_getpid](#) (void)
- int [sys_getprocs](#) (void)
- int [sys_halt](#) (void)
- int [sys_kill](#) (void)
- int [sys_link](#) (void)
- int [sys_mkdir](#) (void)
- int [sys_mknod](#) (void)
- int [sys_open](#) (void)
- int [sys_pipe](#) (void)
- int [sys_read](#) (void)
- int [sys_sbrk](#) (void)
- int [sys_sleep](#) (void)
- int [sys_unlink](#) (void)
- int [sys_uptime](#) (void)
- int [sys_wait](#) (void)
- int [sys_write](#) (void)
- void [syscall](#) (void)

Variables

- static int(* [syscalls](#) [])(void)

5.202.1 Function Documentation

5.202.1.1 argint()

```
int argint (  
    int n,  
    int * ip )
```

Definition at line 50 of file [syscall.c](#).

```
00051 {  
00052     return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);  
00053 }
```

Referenced by [argfd\(\)](#), [argptr\(\)](#), [argstr\(\)](#), [sys_chpr\(\)](#), [sys_exec\(\)](#), [sys_getprocs\(\)](#), [sys_kill\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), [sys_read\(\)](#), [sys_sbrk\(\)](#), [sys_sleep\(\)](#), and [sys_write\(\)](#).

5.202.1.2 argptr()

```
int argptr (
    int n,
    char ** pp,
    int size )
```

Definition at line 59 of file [syscall.c](#).

```
00060 {
00061     int i;
00062     struct proc *curproc = myproc();
00063
00064     if(argint(n, &i) < 0)
00065         return -1;
00066     if(size < 0 || (uint)i >= curproc->sz || (uint)i+size > curproc->sz)
00067         return -1;
00068     *pp = (char*)i;
00069     return 0;
00070 }
```

Referenced by [sys_fstat\(\)](#), [sys_getprocs\(\)](#), [sys_pipe\(\)](#), [sys_read\(\)](#), and [sys_write\(\)](#).

5.202.1.3 argstr()

```
int argstr (
    int n,
    char ** pp )
```

Definition at line 77 of file [syscall.c](#).

```
00078 {
00079     int addr;
00080     if(argint(n, &addr) < 0)
00081         return -1;
00082     return fetchstr(addr, pp);
00083 }
```

Referenced by [sys_chdir\(\)](#), [sys_exec\(\)](#), [sys_link\(\)](#), [sys_mkdir\(\)](#), [sys_mknod\(\)](#), [sys_open\(\)](#), and [sys_unlink\(\)](#).

5.202.1.4 fetchint()

```
int fetchint (
    uint addr,
    int * ip )
```

Definition at line 18 of file [syscall.c](#).

```
00019 {
00020     struct proc *curproc = myproc();
00021
00022     if(addr >= curproc->sz || addr+4 > curproc->sz)
00023         return -1;
00024     *ip = *(int*)(addr);
00025     return 0;
00026 }
```

Referenced by [argint\(\)](#), and [sys_exec\(\)](#).

5.202.1.5 fetchstr()

```
int fetchstr (
    uint addr,
    char ** pp )
```

Definition at line 32 of file [syscall.c](#).

```
00033 {
00034     char *s, *ep;
00035     struct proc *curproc = myproc();
00036
00037     if(addr >= curproc->sz)
00038         return -1;
00039     *pp = (char*)addr;
00040     ep = (char*)curproc->sz;
00041     for(s = *pp; s < ep; s++){
00042         if(*s == 0)
00043             return s - *pp;
00044     }
00045     return -1;
00046 }
```

Referenced by [argstr\(\)](#), and [sys_exec\(\)](#).

5.202.1.6 sys_chdir()

```
int sys_chdir (
    void )
```

Definition at line 372 of file [sysfile.c](#).

```
00373 {
00374     char *path;
00375     struct inode *ip;
00376     struct proc *curproc = myproc();
00377
00378     begin_op();
00379     if(argstr(0, &path) < 0 || (ip = namei(path)) == 0){
00380         end_op();
00381         return -1;
00382     }
00383     ilock(ip);
00384     if(ip->type != T_DIR){
00385         iunlockput(ip);
00386         end_op();
00387         return -1;
00388     }
00389     iunlock(ip);
00390     iput(curproc->cwd);
00391     end_op();
00392     curproc->cwd = ip;
00393     return 0;
00394 }
```

5.202.1.7 sys_chpr()

```
int sys_chpr (
    void )
```

Definition at line 125 of file [sysproc.c](#).

```
00126 {
00127     int pid, pr;
00128     if(argint(0, &pid) < 0)
00129         return -1;
00130     if(argint(1, &pr) < 0)
00131         return -1;
00132
00133     return chpr(pid, pr);
00134 }
```

5.202.1.8 sys_close()

```
int sys_close (  
    void )
```

Definition at line 94 of file [sysfile.c](#).

```
00095 {  
00096     int fd;  
00097     struct file *f;  
00098  
00099     if(argfd(0, &fd, &f) < 0)  
00100         return -1;  
00101     myproc()->ofile[fd] = 0;  
00102     fileclose(f);  
00103     return 0;  
00104 }
```

5.202.1.9 sys_cps()

```
int sys_cps (  
    void )
```

Definition at line 119 of file [sysproc.c](#).

```
00120 {  
00121     return cps();  
00122 }
```

5.202.1.10 sys_dup()

```
int sys_dup (  
    void )
```

Definition at line 56 of file [sysfile.c](#).

```
00057 {  
00058     struct file *f;  
00059     int fd;  
00060  
00061     if(argfd(0, 0, &f) < 0)  
00062         return -1;  
00063     if((fd=fdalloc(f)) < 0)  
00064         return -1;  
00065     filedup(f);  
00066     return fd;  
00067 }
```

5.202.1.11 sys_exec()

```
int sys_exec (
    void )
```

Definition at line 397 of file [sysfile.c](#).

```
00398 {
00399     char *path, *argv[MAXARG];
00400     int i;
00401     uint uargv, uarg;
00402
00403     if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
00404         return -1;
00405     }
00406     memset(argv, 0, sizeof(argv));
00407     for(i=0;; i++){
00408         if(i >= NELEM(argv))
00409             return -1;
00410         if(fetchint(uargv+4*i, (int*)&uarg) < 0)
00411             return -1;
00412         if(uarg == 0){
00413             argv[i] = 0;
00414             break;
00415         }
00416         if(fetchstr(uarg, &argv[i]) < 0)
00417             return -1;
00418     }
00419     return exec(path, argv);
00420 }
```

5.202.1.12 sys_exit()

```
int sys_exit (
    void )
```

Definition at line 25 of file [sysproc.c](#).

```
00026 {
00027     exit();
00028     return 0; // not reached
00029 }
```

5.202.1.13 sys_fork()

```
int sys_fork (
    void )
```

Definition at line 19 of file [sysproc.c](#).

```
00020 {
00021     return fork();
00022 }
```

5.202.1.14 sys_fstat()

```
int sys_fstat (
    void )
```

Definition at line 107 of file [sysfile.c](#).

```
00108 {
00109     struct file *f;
00110     struct stat *st;
00111
00112     if(argfd(0, 0, &f) < 0 || argptr(1, (void*)&st, sizeof(*st)) < 0)
00113         return -1;
00114     return filestat(f, st);
00115 }
```

5.202.1.15 sys_getpid()

```
int sys_getpid (
    void )
```

Definition at line 48 of file [sysproc.c](#).

```
00049 {
00050     return myproc()->pid;
00051 }
```

5.202.1.16 sys_getprocs()

```
int sys_getprocs (
    void )
```

Definition at line 137 of file [sysproc.c](#).

```
00137 {
00138     // declare local variables for max uproc size, struct uproc, and counter
00139     int max;
00140     struct uproc *p;
00141     int i=0;
00142
00143     // if argint has trouble, return error -1
00144     if(argint(0,&max) < 0)
00145         return -1;
00146
00147     // if argptr for allocating struct size has trouble, return -1
00148     if(argptr(1, (char**)&p, max*sizeof(struct uproc)) < 0)
00149         return -1;
00150
00151     // create pointer to ptable processes
00152     struct proc *ptr = ptable.proc;
00153
00154     // loop through ptable
00155     for(; ptr < &ptable.proc[NPROC]; ptr++)
00156     {
00157         if(!(ptr->state == UNUSED))
00158         {
00159             // if the process in ptable is not UNUSED, assign pid, parent pid, and name to uproc
00160             p[i].pid = ptr->pid;
00161             p[i].ppid = ptr->parent->pid;
00162             strncpy(p[i].name, ptr->name, 16);
00163             // add 1 to the process counter
00164             i++;
00165         }
00166     }
00167
00168     // return the number of processes that are not UNUSED
00169     return i;
00170 }
```

5.202.1.17 sys_halt()

```
int sys_halt (
    void )
```

Definition at line 111 of file [sysproc.c](#).

```
00112 {
00113     outb(0xf4, 0x00);
00114     return 0;
00115 }
```

5.202.1.18 sys_kill()

```
int sys_kill (
    void )
```

Definition at line 38 of file [sysproc.c](#).

```
00039 {
00040     int pid;
00041
00042     if(argint(0, &pid) < 0)
00043         return -1;
00044     return kill(pid);
00045 }
```

5.202.1.19 sys_link()

```
int sys_link (
    void )
```

Definition at line 119 of file [sysfile.c](#).

```
00120 {
00121     char name[DIRSIZ], *new, *old;
00122     struct inode *dp, *ip;
00123
00124     if(argstr(0, &old) < 0 || argstr(1, &new) < 0)
00125         return -1;
00126
00127     begin_op();
00128     if((ip = namei(old)) == 0){
00129         end_op();
00130         return -1;
00131     }
00132
00133     ilock(ip);
00134     if(ip->type == T_DIR){
00135         iunlockput(ip);
00136         end_op();
00137         return -1;
00138     }
00139
00140     ip->nlink++;
00141     iupdate(ip);
00142     iunlock(ip);
00143
00144     if((dp = nameiparent(new, name)) == 0)
00145         goto bad;
00146     ilock(dp);
00147     if(dp->dev != ip->dev || dirlink(dp, name, ip->inum) < 0){
00148         iunlockput(dp);
00149         goto bad;
00150     }
00151     iunlockput(dp);
00152     iput(ip);
00153
00154     end_op();
00155
00156     return 0;
00157
00158 bad:
00159     ilock(ip);
00160     ip->nlink--;
00161     iupdate(ip);
00162     iunlockput(ip);
00163     end_op();
00164     return -1;
00165 }
```


5.202.1.20 sys_mkdir()

```
int sys_mkdir (
    void )
```

Definition at line 336 of file [sysfile.c](#).

```
00337 {
00338     char *path;
00339     struct inode *ip;
00340
00341     begin_op();
00342     if(argstr(0, &path) < 0 || (ip = create(path, T_DIR, 0, 0)) == 0){
00343         end_op();
00344         return -1;
00345     }
00346     iunlockput(ip);
00347     end_op();
00348     return 0;
00349 }
```

5.202.1.21 sys_mknod()

```
int sys_mknod (
    void )
```

Definition at line 352 of file [sysfile.c](#).

```
00353 {
00354     struct inode *ip;
00355     char *path;
00356     int major, minor;
00357
00358     begin_op();
00359     if((argstr(0, &path)) < 0 ||
00360        argint(1, &major) < 0 ||
00361        argint(2, &minor) < 0 ||
00362        (ip = create(path, T_DEV, major, minor)) == 0){
00363         end_op();
00364         return -1;
00365     }
00366     iunlockput(ip);
00367     end_op();
00368     return 0;
00369 }
```

5.202.1.22 sys_open()

```
int sys_open (
    void )
```

Definition at line 286 of file [sysfile.c](#).

```
00287 {
00288     char *path;
00289     int fd, omode;
00290     struct file *f;
00291     struct inode *ip;
00292
00293     if(argstr(0, &path) < 0 || argint(1, &omode) < 0)
00294         return -1;
00295
00296     begin_op();
00297
00298     if(omode & O_CREATE){
00299         ip = create(path, T_FILE, 0, 0);
00300         if(ip == 0){
00301             end_op();
00302             return -1;
```

```

00303     }
00304 } else {
00305     if((ip = namei(path)) == 0){
00306         end_op();
00307         return -1;
00308     }
00309     ilock(ip);
00310     if(ip->type == T_DIR && omode != O_RDONLY){
00311         iunlockput(ip);
00312         end_op();
00313         return -1;
00314     }
00315 }
00316
00317 if((f = filealloc()) == 0 || (fd = fdalloc(f)) < 0){
00318     if(f)
00319         fileclose(f);
00320     iunlockput(ip);
00321     end_op();
00322     return -1;
00323 }
00324 iunlock(ip);
00325 end_op();
00326
00327 f->type = FD_INODE;
00328 f->ip = ip;
00329 f->off = 0;
00330 f->readable = !(omode & O_WRONLY);
00331 f->writable = (omode & O_WRONLY) || (omode & O_RDWR);
00332 return fd;
00333 }

```

5.202.1.23 sys_pipe()

```

int sys_pipe (
    void )

```

Definition at line 423 of file [sysfile.c](#).

```

00424 {
00425     int *fd;
00426     struct file *rf, *wf;
00427     int fd0, fd1;
00428
00429     if(argptr(0, (void*)&fd, 2*sizeof(fd[0])) < 0)
00430         return -1;
00431     if(pipealloc(&rf, &wf) < 0)
00432         return -1;
00433     fd0 = -1;
00434     if((fd0 = fdalloc(rf)) < 0 || (fd1 = fdalloc(wf)) < 0){
00435         if(fd0 >= 0)
00436             myproc()->ofile[fd0] = 0;
00437         fileclose(rf);
00438         fileclose(wf);
00439         return -1;
00440     }
00441     fd[0] = fd0;
00442     fd[1] = fd1;
00443     return 0;
00444 }

```

5.202.1.24 sys_read()

```

int sys_read (
    void )

```

Definition at line 70 of file [sysfile.c](#).

```

00071 {
00072     struct file *f;
00073     int n;

```

```
00074     char *p;
00075
00076     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
00077         return -1;
00078     return fileread(f, p, n);
00079 }
```

5.202.1.25 sys_sbrk()

```
int sys_sbrk (
    void )
```

Definition at line 54 of file [sysproc.c](#).

```
00055 {
00056     int addr;
00057     int n;
00058
00059     if(argint(0, &n) < 0)
00060         return -1;
00061     addr = myproc()->sz;
00062     if(growproc(n) < 0)
00063         return -1;
00064     return addr;
00065 }
```

5.202.1.26 sys_sleep()

```
int sys_sleep (
    void )
```

Definition at line 68 of file [sysproc.c](#).

```
00069 {
00070     int n;
00071     uint ticks0;
00072
00073     if(argint(0, &n) < 0)
00074         return -1;
00075     acquire(&tickslock);
00076     ticks0 = ticks;
00077     while(ticks - ticks0 < n){
00078         if(myproc()->killed){
00079             release(&tickslock);
00080             return -1;
00081         }
00082         sleep(&ticks, &tickslock);
00083     }
00084     release(&tickslock);
00085     return 0;
00086 }
```

5.202.1.27 sys_unlink()

```
int sys_unlink (
    void )
```

Definition at line 185 of file [sysfile.c](#).

```
00186 {
00187     struct inode *ip, *dp;
00188     struct dirent de;
00189     char name[DIRSIZ], *path;
00190     uint off;
00191
00192     if(argstr(0, &path) < 0)
00193         return -1;
00194
00195     begin_op();
00196     if((dp = nameiparent(path, name)) == 0){
00197         end_op();
00198         return -1;
00199     }
00200
00201     ilock(dp);
00202
00203     // Cannot unlink "." or "..".
00204     if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
00205         goto bad;
00206
00207     if((ip = dirlookup(dp, name, &off)) == 0)
00208         goto bad;
00209     ilock(ip);
00210
00211     if(ip->nlink < 1)
00212         panic("unlink: nlink < 1");
00213     if(ip->type == T_DIR && !isdirempty(ip)){
00214         iunlockput(ip);
00215         goto bad;
00216     }
00217
00218     memset(&de, 0, sizeof(de));
00219     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00220         panic("unlink: writei");
00221     if(ip->type == T_DIR){
00222         dp->nlink--;
00223         iupdate(dp);
00224     }
00225     iunlockput(dp);
00226
00227     ip->nlink--;
00228     iupdate(ip);
00229     iunlockput(ip);
00230
00231     end_op();
00232
00233     return 0;
00234
00235 bad:
00236     iunlockput(dp);
00237     end_op();
00238     return -1;
00239 }
```

5.202.1.28 sys_uptime()

```
int sys_uptime (
    void )
```

Definition at line 91 of file [sysproc.c](#).

```
00092 {
00093     uint xticks;
00094
00095     acquire(&tickslock);
00096     xticks = ticks;
00097     release(&tickslock);
00098     return xticks;
00099 }
```

5.202.1.29 sys_wait()

```
int sys_wait (  
    void )
```

Definition at line 32 of file [sysproc.c](#).

```
00033 {  
00034     return wait();  
00035 }
```

5.202.1.30 sys_write()

```
int sys_write (  
    void )
```

Definition at line 82 of file [sysfile.c](#).

```
00083 {  
00084     struct file *f;  
00085     int n;  
00086     char *p;  
00087  
00088     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)  
00089         return -1;  
00090     return filewrite(f, p, n);  
00091 }
```

5.202.1.31 syscall()

```
void syscall (  
    void )
```

Definition at line 140 of file [syscall.c](#).

```
00141 {  
00142     int num;  
00143     struct proc *curproc = myproc();  
00144  
00145     num = curproc->tf->eax;  
00146     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {  
00147         curproc->tf->eax = syscalls[num]();  
00148     } else {  
00149         cprintf("%d %s: unknown sys call %d\n",  
00150             curproc->pid, curproc->name, num);  
00151         curproc->tf->eax = -1;  
00152     }  
00153 }
```

Referenced by [trap\(\)](#).

5.202.2 Variable Documentation

5.202.2.1 syscalls

```
int(* syscalls[])(void) (
    void ) [static]
```

Initial value:

```
= {
[SYS_fork]      sys_fork,
[SYS_exit]     sys_exit,
[SYS_wait]     sys_wait,
[SYS_pipe]     sys_pipe,
[SYS_read]     sys_read,
[SYS_kill]     sys_kill,
[SYS_exec]     sys_exec,
[SYS_fstat]    sys_fstat,
[SYS_chdir]    sys_chdir,
[SYS_dup]      sys_dup,
[SYS_getpid]   sys_getpid,
[SYS_sbrk]     sys_sbrk,
[SYS_sleep]    sys_sleep,
[SYS_uptime]   sys_uptime,
[SYS_open]     sys_open,
[SYS_write]    sys_write,
[SYS_mknod]    sys_mknod,
[SYS_unlink]   sys_unlink,
[SYS_link]     sys_link,
[SYS_mkdir]    sys_mkdir,
[SYS_close]    sys_close,
[SYS_halt]     sys_halt,
[SYS_cps]      sys_cps,
[SYS_chpr]     sys_chpr,
[SYS_getprocs] sys_getprocs,
}
```

Definition at line 111 of file [syscall.c](#).

Referenced by [syscall\(\)](#).

5.203 syscall.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "memlayout.h"
00005 #include "mmu.h"
00006 #include "proc.h"
00007 #include "x86.h"
00008 #include "syscall.h"
00009
00010 // User code makes a system call with INT T_SYSCALL.
00011 // System call number in %eax.
00012 // Arguments on the stack, from the user call to the C
00013 // library system call function. The saved user %esp points
00014 // to a saved program counter, and then the first argument.
00015
00016 // Fetch the int at addr from the current process.
00017 int
00018 fetchint(uint addr, int *ip)
00019 {
00020     struct proc *curproc = myproc();
00021
00022     if(addr >= curproc->sz || addr+4 > curproc->sz)
00023         return -1;
00024     *ip = *(int*)(addr);
00025     return 0;
00026 }
00027
00028 // Fetch the nul-terminated string at addr from the current process.
00029 // Doesn't actually copy the string - just sets *pp to point at it.
00030 // Returns length of string, not including nul.
00031 int
00032 fetchstr(uint addr, char **pp)
00033 {
00034     char *s, *ep;
00035     struct proc *curproc = myproc();
```

```

00036
00037     if(addr >= curproc->sz)
00038         return -1;
00039     *pp = (char*)addr;
00040     ep = (char*)curproc->sz;
00041     for(s = *pp; s < ep; s++){
00042         if(*s == 0)
00043             return s - *pp;
00044     }
00045     return -1;
00046 }
00047
00048 // Fetch the nth 32-bit system call argument.
00049 int
00050 argint(int n, int *ip)
00051 {
00052     return fetchint((myproc()->tf->esp) + 4 + 4*n, ip);
00053 }
00054
00055 // Fetch the nth word-sized system call argument as a pointer
00056 // to a block of memory of size bytes. Check that the pointer
00057 // lies within the process address space.
00058 int
00059 argptr(int n, char **pp, int size)
00060 {
00061     int i;
00062     struct proc *curproc = myproc();
00063
00064     if(argint(n, &i) < 0)
00065         return -1;
00066     if(size < 0 || (uint)i >= curproc->sz || (uint)i+size > curproc->sz)
00067         return -1;
00068     *pp = (char*)i;
00069     return 0;
00070 }
00071
00072 // Fetch the nth word-sized system call argument as a string pointer.
00073 // Check that the pointer is valid and the string is nul-terminated.
00074 // (There is no shared writable memory, so the string can't change
00075 // between this check and being used by the kernel.)
00076 int
00077 argstr(int n, char **pp)
00078 {
00079     int addr;
00080     if(argint(n, &addr) < 0)
00081         return -1;
00082     return fetchstr(addr, pp);
00083 }
00084
00085 extern int sys_chdir(void);
00086 extern int sys_close(void);
00087 extern int sys_dup(void);
00088 extern int sys_exec(void);
00089 extern int sys_exit(void);
00090 extern int sys_fork(void);
00091 extern int sys_fstat(void);
00092 extern int sys_getpid(void);
00093 extern int sys_kill(void);
00094 extern int sys_link(void);
00095 extern int sys_mkdir(void);
00096 extern int sys_mknod(void);
00097 extern int sys_open(void);
00098 extern int sys_pipe(void);
00099 extern int sys_read(void);
00100 extern int sys_sbrk(void);
00101 extern int sys_sleep(void);
00102 extern int sys_unlink(void);
00103 extern int sys_wait(void);
00104 extern int sys_write(void);
00105 extern int sys_uptime(void);
00106 extern int sys_halt(void);
00107 extern int sys_cps(void);
00108 extern int sys_chpr(void);
00109 extern int sys_getprocs(void);
00110
00111 static int (*syscalls[]) (void) = {
00112     [SYS_fork] sys_fork,
00113     [SYS_exit] sys_exit,
00114     [SYS_wait] sys_wait,
00115     [SYS_pipe] sys_pipe,
00116     [SYS_read] sys_read,
00117     [SYS_kill] sys_kill,
00118     [SYS_exec] sys_exec,
00119     [SYS_fstat] sys_fstat,
00120     [SYS_chdir] sys_chdir,
00121     [SYS_dup] sys_dup,
00122     [SYS_getpid] sys_getpid,

```

```

00123 [SYS_sbrk]      sys_sbrk,
00124 [SYS_sleep]    sys_sleep,
00125 [SYS_uptime]   sys_uptime,
00126 [SYS_open]     sys_open,
00127 [SYS_write]    sys_write,
00128 [SYS_mknod]    sys_mknod,
00129 [SYS_unlink]   sys_unlink,
00130 [SYS_link]      sys_link,
00131 [SYS_mkdir]    sys_mkdir,
00132 [SYS_close]    sys_close,
00133 [SYS_halt]      sys_halt,
00134 [SYS_cps]      sys_cps,
00135 [SYS_chpr]     sys_chpr,
00136 [SYS_getprocs] sys_getprocs,
00137 };
00138
00139 void
00140 syscall(void)
00141 {
00142     int num;
00143     struct proc *curproc = myproc();
00144
00145     num = curproc->tf->eax;
00146     if(num > 0 && num < NELEM(syscalls) && syscalls[num]) {
00147         curproc->tf->eax = syscalls[num]();
00148     } else {
00149         cprintf("%d %s: unknown sys call %d\n",
00150             curproc->pid, curproc->name, num);
00151         curproc->tf->eax = -1;
00152     }
00153 }

```

5.204 syscall.d File Reference

5.205 syscall.d

[Go to the documentation of this file.](#)

```

00001 syscall.o: syscall.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 memlayout.h mmu.h proc.h x86.h syscall.h

```

5.206 syscall.h File Reference

Macros

- #define [SYS_chdir](#) 9
- #define [SYS_chpr](#) 24
- #define [SYS_close](#) 21
- #define [SYS_cps](#) 23
- #define [SYS_dup](#) 10
- #define [SYS_exec](#) 7
- #define [SYS_exit](#) 2
- #define [SYS_fork](#) 1
- #define [SYS_fstat](#) 8
- #define [SYS_getpid](#) 11
- #define [SYS_getprocs](#) 25
- #define [SYS_halt](#) 22
- #define [SYS_kill](#) 6
- #define [SYS_link](#) 19
- #define [SYS_mkdir](#) 20
- #define [SYS_mknod](#) 17
- #define [SYS_open](#) 15

- `#define SYS_pipe` 4
- `#define SYS_read` 5
- `#define SYS_sbrk` 12
- `#define SYS_sleep` 13
- `#define SYS_unlink` 18
- `#define SYS_uptime` 14
- `#define SYS_wait` 3
- `#define SYS_write` 16

5.206.1 Macro Definition Documentation

5.206.1.1 SYS_chdir

```
#define SYS_chdir 9
```

Definition at line 10 of file [syscall.h](#).

5.206.1.2 SYS_chpr

```
#define SYS_chpr 24
```

Definition at line 25 of file [syscall.h](#).

5.206.1.3 SYS_close

```
#define SYS_close 21
```

Definition at line 22 of file [syscall.h](#).

5.206.1.4 SYS_cps

```
#define SYS_cps 23
```

Definition at line 24 of file [syscall.h](#).

5.206.1.5 SYS_dup

```
#define SYS_dup 10
```

Definition at line 11 of file [syscall.h](#).

5.206.1.6 SYS_exec

```
#define SYS_exec 7
```

Definition at line 8 of file [syscall.h](#).

5.206.1.7 SYS_exit

```
#define SYS_exit 2
```

Definition at line 3 of file [syscall.h](#).

5.206.1.8 SYS_fork

```
#define SYS_fork 1
```

Definition at line 2 of file [syscall.h](#).

5.206.1.9 SYS_fstat

```
#define SYS_fstat 8
```

Definition at line 9 of file [syscall.h](#).

5.206.1.10 SYS_getpid

```
#define SYS_getpid 11
```

Definition at line 12 of file [syscall.h](#).

5.206.1.11 SYS_getprocs

```
#define SYS_getprocs 25
```

Definition at line 26 of file [syscall.h](#).

5.206.1.12 SYS_halt

```
#define SYS_halt 22
```

Definition at line 23 of file [syscall.h](#).

5.206.1.13 SYS_kill

```
#define SYS_kill 6
```

Definition at line 7 of file [syscall.h](#).

5.206.1.14 SYS_link

```
#define SYS_link 19
```

Definition at line 20 of file [syscall.h](#).

5.206.1.15 SYS_mkdir

```
#define SYS_mkdir 20
```

Definition at line 21 of file [syscall.h](#).

5.206.1.16 SYS_mknod

```
#define SYS_mknod 17
```

Definition at line 18 of file [syscall.h](#).

5.206.1.17 SYS_open

```
#define SYS_open 15
```

Definition at line 16 of file [syscall.h](#).

5.206.1.18 SYS_pipe

```
#define SYS_pipe 4
```

Definition at line 5 of file [syscall.h](#).

5.206.1.19 SYS_read

```
#define SYS_read 5
```

Definition at line 6 of file [syscall.h](#).

5.206.1.20 SYS_sbrk

```
#define SYS_sbrk 12
```

Definition at line 13 of file [syscall.h](#).

5.206.1.21 SYS_sleep

```
#define SYS_sleep 13
```

Definition at line 14 of file [syscall.h](#).

5.206.1.22 SYS_unlink

```
#define SYS_unlink 18
```

Definition at line 19 of file [syscall.h](#).

5.206.1.23 SYS_uptime

```
#define SYS_uptime 14
```

Definition at line 15 of file [syscall.h](#).

5.206.1.24 SYS_wait

```
#define SYS_wait 3
```

Definition at line 4 of file [syscall.h](#).

5.206.1.25 SYS_write

```
#define SYS_write 16
```

Definition at line 17 of file [syscall.h](#).

5.207 syscall.h

[Go to the documentation of this file.](#)

```
00001 // System call numbers
00002 #define SYS_fork      1
00003 #define SYS_exit      2
00004 #define SYS_wait      3
00005 #define SYS_pipe      4
00006 #define SYS_read      5
00007 #define SYS_kill      6
00008 #define SYS_exec      7
00009 #define SYS_fstat     8
00010 #define SYS_chdir     9
00011 #define SYS_dup      10
00012 #define SYS_getpid   11
00013 #define SYS_sbrk     12
00014 #define SYS_sleep    13
00015 #define SYS_uptime   14
00016 #define SYS_open     15
00017 #define SYS_write    16
00018 #define SYS_mknod    17
00019 #define SYS_unlink   18
00020 #define SYS_link     19
00021 #define SYS_mkdir    20
00022 #define SYS_close    21
00023 #define SYS_halt     22
00024 #define SYS_cps      23
00025 #define SYS_chpr     24
00026 #define SYS_getprocs 25
00027
```

5.208 sysfile.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "stat.h"
#include "mmu.h"
#include "proc.h"
#include "fs.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "file.h"
#include "fcntl.h"
```

Functions

- static int [argfd](#) (int n, int *pfd, struct [file](#) **pf)
- static struct [inode](#) * [create](#) (char *path, short type, short major, short minor)
- static int [fdalloc](#) (struct [file](#) *f)
- static int [isdirempty](#) (struct [inode](#) *dp)
- int [sys_chdir](#) (void)
- int [sys_close](#) (void)
- int [sys_dup](#) (void)
- int [sys_exec](#) (void)
- int [sys_fstat](#) (void)
- int [sys_link](#) (void)
- int [sys_mkdir](#) (void)
- int [sys_mknod](#) (void)
- int [sys_open](#) (void)
- int [sys_pipe](#) (void)
- int [sys_read](#) (void)
- int [sys_unlink](#) (void)
- int [sys_write](#) (void)

5.208.1 Function Documentation

5.208.1.1 argfd()

```
static int argfd (
    int n,
    int * pfd,
    struct file ** pf )  [static]
```

Definition at line 22 of file [sysfile.c](#).

```
00023 {
00024     int fd;
00025     struct file *f;
00026
00027     if(argint(n, &fd) < 0)
00028         return -1;
00029     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
```

```

00030     return -1;
00031     if(pfd)
00032         *pfd = fd;
00033     if(pf)
00034         *pf = f;
00035     return 0;
00036 }

```

Referenced by [sys_close\(\)](#), [sys_dup\(\)](#), [sys_fstat\(\)](#), [sys_read\(\)](#), and [sys_write\(\)](#).

5.208.1.2 create()

```

static struct inode * create (
    char * path,
    short type,
    short major,
    short minor ) [static]

```

Definition at line 242 of file [sysfile.c](#).

```

00243 {
00244     struct inode *ip, *dp;
00245     char name[DIRSIZ];
00246
00247     if((dp = nameiparent(path, name)) == 0)
00248         return 0;
00249     ilock(dp);
00250
00251     if((ip = dirlookup(dp, name, 0)) != 0){
00252         iunlockput(dp);
00253         ilock(ip);
00254         if(type == T_FILE && ip->type == T_FILE)
00255             return ip;
00256         iunlockput(ip);
00257         return 0;
00258     }
00259
00260     if((ip = ialloc(dp->dev, type)) == 0)
00261         panic("create: ialloc");
00262
00263     ilock(ip);
00264     ip->major = major;
00265     ip->minor = minor;
00266     ip->nlink = 1;
00267     iupdate(ip);
00268
00269     if(type == T_DIR){ // Create . and .. entries.
00270         dp->nlink++; // for "."
00271         iupdate(dp);
00272         // No ip->nlink++ for ".": avoid cyclic ref count.
00273         if(dirlink(ip, ".", ip->inum) < 0 || dirlink(ip, "..", dp->inum) < 0)
00274             panic("create dots");
00275     }
00276
00277     if(dirlink(dp, name, ip->inum) < 0)
00278         panic("create: dirlink");
00279
00280     iunlockput(dp);
00281
00282     return ip;
00283 }

```

Referenced by [sys_mkdir\(\)](#), [sys_mknod\(\)](#), and [sys_open\(\)](#).

5.208.1.3 fdalloc()

```
static int fdalloc (
    struct file * f ) [static]
```

Definition at line 41 of file [sysfile.c](#).

```
00042 {
00043     int fd;
00044     struct proc *curproc = myproc();
00045
00046     for(fd = 0; fd < NOFILE; fd++){
00047         if(curproc->ofile[fd] == 0){
00048             curproc->ofile[fd] = f;
00049             return fd;
00050         }
00051     }
00052     return -1;
00053 }
```

Referenced by [sys_dup\(\)](#), [sys_open\(\)](#), and [sys_pipe\(\)](#).

5.208.1.4 isdirempty()

```
static int isdirempty (
    struct inode * dp ) [static]
```

Definition at line 169 of file [sysfile.c](#).

```
00170 {
00171     int off;
00172     struct dirent de;
00173
00174     for(off=2*sizeof(de); off<dp->size; off+=sizeof(de)){
00175         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00176             panic("isdirempty: readi");
00177         if(de.inum != 0)
00178             return 0;
00179     }
00180     return 1;
00181 }
```

Referenced by [sys_unlink\(\)](#).

5.208.1.5 sys_chdir()

```
int sys_chdir (
    void )
```

Definition at line 372 of file [sysfile.c](#).

```
00373 {
00374     char *path;
00375     struct inode *ip;
00376     struct proc *curproc = myproc();
00377
00378     begin_op();
00379     if(argstr(0, &path) < 0 || (ip = namei(path)) == 0){
00380         end_op();
00381         return -1;
00382     }
00383     ilock(ip);
00384     if(ip->type != T_DIR){
00385         iunlockput(ip);
00386         end_op();
00387         return -1;
00388     }
00389     iunlock(ip);
00390     iput(curproc->cwd);
00391     end_op();
00392     curproc->cwd = ip;
00393     return 0;
00394 }
```


5.208.1.6 sys_close()

```
int sys_close (
    void )
```

Definition at line 94 of file [sysfile.c](#).

```
00095 {
00096     int fd;
00097     struct file *f;
00098
00099     if(argfd(0, &fd, &f) < 0)
00100         return -1;
00101     myproc()->ofile[fd] = 0;
00102     fileclose(f);
00103     return 0;
00104 }
```

5.208.1.7 sys_dup()

```
int sys_dup (
    void )
```

Definition at line 56 of file [sysfile.c](#).

```
00057 {
00058     struct file *f;
00059     int fd;
00060
00061     if(argfd(0, 0, &f) < 0)
00062         return -1;
00063     if((fd=fdalloc(f)) < 0)
00064         return -1;
00065     filedup(f);
00066     return fd;
00067 }
```

5.208.1.8 sys_exec()

```
int sys_exec (
    void )
```

Definition at line 397 of file [sysfile.c](#).

```
00398 {
00399     char *path, *argv[MAXARG];
00400     int i;
00401     uint uargv, uarg;
00402
00403     if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
00404         return -1;
00405     }
00406     memset(argv, 0, sizeof(argv));
00407     for(i=0; i++){
00408         if(i >= NELEM(argv))
00409             return -1;
00410         if(fetchint(uargv+4*i, (int*)&uarg) < 0)
00411             return -1;
00412         if(uarg == 0){
00413             argv[i] = 0;
00414             break;
00415         }
00416         if(fetchstr(uarg, &argv[i]) < 0)
00417             return -1;
00418     }
00419     return exec(path, argv);
00420 }
```

5.208.1.9 sys_fstat()

```
int sys_fstat (  
    void )
```

Definition at line 107 of file [sysfile.c](#).

```
00108 {  
00109     struct file *f;  
00110     struct stat *st;  
00111  
00112     if(argfd(0, 0, &f) < 0 || argptr(1, (void*)&st, sizeof(*st)) < 0)  
00113         return -1;  
00114     return filestat(f, st);  
00115 }
```

5.208.1.10 sys_link()

```
int sys_link (  
    void )
```

Definition at line 119 of file [sysfile.c](#).

```
00120 {  
00121     char name[DIRSIZ], *new, *old;  
00122     struct inode *dp, *ip;  
00123  
00124     if(argstr(0, &old) < 0 || argstr(1, &new) < 0)  
00125         return -1;  
00126  
00127     begin_op();  
00128     if((ip = namei(old)) == 0){  
00129         end_op();  
00130         return -1;  
00131     }  
00132  
00133     ilock(ip);  
00134     if(ip->type == T_DIR){  
00135         iunlockput(ip);  
00136         end_op();  
00137         return -1;  
00138     }  
00139  
00140     ip->nlink++;  
00141     iupdate(ip);  
00142     iunlock(ip);  
00143  
00144     if((dp = nameiparent(new, name)) == 0)  
00145         goto bad;  
00146     ilock(dp);  
00147     if(dp->dev != ip->dev || dirlink(dp, name, ip->inum) < 0){  
00148         iunlockput(dp);  
00149         goto bad;  
00150     }  
00151     iunlockput(dp);  
00152     iput(ip);  
00153  
00154     end_op();  
00155  
00156     return 0;  
00157  
00158 bad:  
00159     ilock(ip);  
00160     ip->nlink--;  
00161     iupdate(ip);  
00162     iunlockput(ip);  
00163     end_op();  
00164     return -1;  
00165 }
```

5.208.1.11 sys_mkdir()

```
int sys_mkdir (
    void )
```

Definition at line 336 of file [sysfile.c](#).

```
00337 {
00338     char *path;
00339     struct inode *ip;
00340
00341     begin_op();
00342     if(argstr(0, &path) < 0 || (ip = create(path, T_DIR, 0, 0)) == 0){
00343         end_op();
00344         return -1;
00345     }
00346     iunlockput(ip);
00347     end_op();
00348     return 0;
00349 }
```

5.208.1.12 sys_mknod()

```
int sys_mknod (
    void )
```

Definition at line 352 of file [sysfile.c](#).

```
00353 {
00354     struct inode *ip;
00355     char *path;
00356     int major, minor;
00357
00358     begin_op();
00359     if((argstr(0, &path)) < 0 ||
00360        argint(1, &major) < 0 ||
00361        argint(2, &minor) < 0 ||
00362        (ip = create(path, T_DEV, major, minor)) == 0){
00363         end_op();
00364         return -1;
00365     }
00366     iunlockput(ip);
00367     end_op();
00368     return 0;
00369 }
```

5.208.1.13 sys_open()

```
int sys_open (
    void )
```

Definition at line 286 of file [sysfile.c](#).

```
00287 {
00288     char *path;
00289     int fd, omode;
00290     struct file *f;
00291     struct inode *ip;
00292
00293     if(argstr(0, &path) < 0 || argint(1, &omode) < 0)
00294         return -1;
00295
00296     begin_op();
00297
00298     if(omode & O_CREATE){
00299         ip = create(path, T_FILE, 0, 0);
00300         if(ip == 0){
00301             end_op();
00302             return -1;
```

```

00303     }
00304 } else {
00305     if((ip = namei(path)) == 0){
00306         end_op();
00307         return -1;
00308     }
00309     ilock(ip);
00310     if(ip->type == T_DIR && omode != O_RDONLY){
00311         iunlockput(ip);
00312         end_op();
00313         return -1;
00314     }
00315 }
00316
00317 if((f = filealloc()) == 0 || (fd = fdalloc(f)) < 0){
00318     if(f)
00319         fileclose(f);
00320     iunlockput(ip);
00321     end_op();
00322     return -1;
00323 }
00324 iunlock(ip);
00325 end_op();
00326
00327 f->type = FD_INODE;
00328 f->ip = ip;
00329 f->off = 0;
00330 f->readable = !(omode & O_WRONLY);
00331 f->writable = (omode & O_WRONLY) || (omode & O_RDWR);
00332 return fd;
00333 }

```

5.208.1.14 sys_pipe()

```

int sys_pipe (
    void )

```

Definition at line 423 of file [sysfile.c](#).

```

00424 {
00425     int *fd;
00426     struct file *rf, *wf;
00427     int fd0, fd1;
00428
00429     if(argptr(0, (void*)&fd, 2*sizeof(fd[0])) < 0)
00430         return -1;
00431     if(pipealloc(&rf, &wf) < 0)
00432         return -1;
00433     fd0 = -1;
00434     if((fd0 = fdalloc(rf)) < 0 || (fd1 = fdalloc(wf)) < 0){
00435         if(fd0 >= 0)
00436             myproc()->ofile[fd0] = 0;
00437         fileclose(rf);
00438         fileclose(wf);
00439         return -1;
00440     }
00441     fd[0] = fd0;
00442     fd[1] = fd1;
00443     return 0;
00444 }

```

5.208.1.15 sys_read()

```

int sys_read (
    void )

```

Definition at line 70 of file [sysfile.c](#).

```

00071 {
00072     struct file *f;
00073     int n;

```

```

00074  char *p;
00075
00076  if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
00077      return -1;
00078  return fileread(f, p, n);
00079  }

```

5.208.1.16 sys_unlink()

```

int sys_unlink (
    void )

```

Definition at line 185 of file [sysfile.c](#).

```

00186 {
00187     struct inode *ip, *dp;
00188     struct dirent de;
00189     char name[DIRSIZ], *path;
00190     uint off;
00191
00192     if(argstr(0, &path) < 0)
00193         return -1;
00194
00195     begin_op();
00196     if((dp = nameiparent(path, name)) == 0){
00197         end_op();
00198         return -1;
00199     }
00200
00201     ilock(dp);
00202
00203     // Cannot unlink "." or "..".
00204     if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
00205         goto bad;
00206
00207     if((ip = dirlookup(dp, name, &off)) == 0)
00208         goto bad;
00209     ilock(ip);
00210
00211     if(ip->nlink < 1)
00212         panic("unlink: nlink < 1");
00213     if(ip->type == T_DIR && !isdirempty(ip)){
00214         iunlockput(ip);
00215         goto bad;
00216     }
00217
00218     memset(&de, 0, sizeof(de));
00219     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00220         panic("unlink: writei");
00221     if(ip->type == T_DIR){
00222         dp->nlink--;
00223         iupdate(dp);
00224     }
00225     iunlockput(dp);
00226
00227     ip->nlink--;
00228     iupdate(ip);
00229     iunlockput(ip);
00230
00231     end_op();
00232
00233     return 0;
00234
00235 bad:
00236     iunlockput(dp);
00237     end_op();
00238     return -1;
00239 }

```

5.208.1.17 sys_write()

```
int sys_write (
    void )
```

Definition at line 82 of file [sysfile.c](#).

```
00083 {
00084     struct file *f;
00085     int n;
00086     char *p;
00087
00088     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
00089         return -1;
00090     return filewrite(f, p, n);
00091 }
```

5.209 sysfile.c

[Go to the documentation of this file.](#)

```
00001 //
00002 // File-system system calls.
00003 // Mostly argument checking, since we don't trust
00004 // user code, and calls into file.c and fs.c.
00005 //
00006
00007 #include "types.h"
00008 #include "defs.h"
00009 #include "param.h"
00010 #include "stat.h"
00011 #include "mmu.h"
00012 #include "proc.h"
00013 #include "fs.h"
00014 #include "spinlock.h"
00015 #include "sleeplock.h"
00016 #include "file.h"
00017 #include "fcntl.h"
00018
00019 // Fetch the nth word-sized system call argument as a file descriptor
00020 // and return both the descriptor and the corresponding struct file.
00021 static int
00022 argfd(int n, int *pfd, struct file **pf)
00023 {
00024     int fd;
00025     struct file *f;
00026
00027     if(argint(n, &fd) < 0)
00028         return -1;
00029     if(fd < 0 || fd >= NOFILE || (f=myproc()->ofile[fd]) == 0)
00030         return -1;
00031     if(pfd)
00032         *pfd = fd;
00033     if(pf)
00034         *pf = f;
00035     return 0;
00036 }
00037
00038 // Allocate a file descriptor for the given file.
00039 // Takes over file reference from caller on success.
00040 static int
00041 fdalloc(struct file *f)
00042 {
00043     int fd;
00044     struct proc *curproc = myproc();
00045
00046     for(fd = 0; fd < NOFILE; fd++){
00047         if(curproc->ofile[fd] == 0){
00048             curproc->ofile[fd] = f;
00049             return fd;
00050         }
00051     }
00052     return -1;
00053 }
00054
00055 int
00056 sys_dup(void)
00057 {
00058     struct file *f;
00059     int fd;
```

```

00060
00061     if(argfd(0, 0, &f) < 0)
00062         return -1;
00063     if((fd=fdalloc(f)) < 0)
00064         return -1;
00065     filedup(f);
00066     return fd;
00067 }
00068
00069 int
00070 sys_read(void)
00071 {
00072     struct file *f;
00073     int n;
00074     char *p;
00075
00076     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
00077         return -1;
00078     return fileread(f, p, n);
00079 }
00080
00081 int
00082 sys_write(void)
00083 {
00084     struct file *f;
00085     int n;
00086     char *p;
00087
00088     if(argfd(0, 0, &f) < 0 || argint(2, &n) < 0 || argptr(1, &p, n) < 0)
00089         return -1;
00090     return filewrite(f, p, n);
00091 }
00092
00093 int
00094 sys_close(void)
00095 {
00096     int fd;
00097     struct file *f;
00098
00099     if(argfd(0, 0, &fd, &f) < 0)
00100         return -1;
00101     myproc()->ofile[fd] = 0;
00102     fileclose(f);
00103     return 0;
00104 }
00105
00106 int
00107 sys_fstat(void)
00108 {
00109     struct file *f;
00110     struct stat *st;
00111
00112     if(argfd(0, 0, &f) < 0 || argptr(1, (void*)&st, sizeof(*st)) < 0)
00113         return -1;
00114     return filestat(f, st);
00115 }
00116
00117 // Create the path new as a link to the same inode as old.
00118 int
00119 sys_link(void)
00120 {
00121     char name[DIRSIZ], *new, *old;
00122     struct inode *dp, *ip;
00123
00124     if(argstr(0, &old) < 0 || argstr(1, &new) < 0)
00125         return -1;
00126
00127     begin_op();
00128     if((ip = namei(old)) == 0){
00129         end_op();
00130         return -1;
00131     }
00132
00133     ilock(ip);
00134     if(ip->type == T_DIR){
00135         iunlockput(ip);
00136         end_op();
00137         return -1;
00138     }
00139
00140     ip->nlink++;
00141     iupdate(ip);
00142     iunlock(ip);
00143
00144     if((dp = nameiparent(new, name)) == 0)
00145         goto bad;
00146     ilock(dp);

```

```

00147     if(dp->dev != ip->dev || dirlink(dp, name, ip->inum) < 0){
00148         iunlockput(dp);
00149         goto bad;
00150     }
00151     iunlockput(dp);
00152     iput(ip);
00153
00154     end_op();
00155
00156     return 0;
00157
00158 bad:
00159     ilock(ip);
00160     ip->nlink--;
00161     iupdate(ip);
00162     iunlockput(ip);
00163     end_op();
00164     return -1;
00165 }
00166
00167 // Is the directory dp empty except for "." and ".." ?
00168 static int
00169 isdirempty(struct inode *dp)
00170 {
00171     int off;
00172     struct dirent de;
00173
00174     for(off=2*sizeof(de); off<dp->size; off+=sizeof(de)){
00175         if(readi(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00176             panic("isdirempty: readi");
00177         if(de.inum != 0)
00178             return 0;
00179     }
00180     return 1;
00181 }
00182
00183 //PAGEBREAK!
00184 int
00185 sys_unlink(void)
00186 {
00187     struct inode *ip, *dp;
00188     struct dirent de;
00189     char name[DIRSIZ], *path;
00190     uint off;
00191
00192     if(argstr(0, &path) < 0)
00193         return -1;
00194
00195     begin_op();
00196     if((dp = nameiparent(path, name)) == 0){
00197         end_op();
00198         return -1;
00199     }
00200
00201     ilock(dp);
00202
00203     // Cannot unlink "." or "..".
00204     if(namecmp(name, ".") == 0 || namecmp(name, "..") == 0)
00205         goto bad;
00206
00207     if((ip = dirlookup(dp, name, &off)) == 0)
00208         goto bad;
00209     ilock(ip);
00210
00211     if(ip->nlink < 1)
00212         panic("unlink: nlink < 1");
00213     if(ip->type == T_DIR && !isdirempty(ip)){
00214         iunlockput(ip);
00215         goto bad;
00216     }
00217
00218     memset(&de, 0, sizeof(de));
00219     if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
00220         panic("unlink: writei");
00221     if(ip->type == T_DIR){
00222         dp->nlink--;
00223         iupdate(dp);
00224     }
00225     iunlockput(dp);
00226
00227     ip->nlink--;
00228     iupdate(ip);
00229     iunlockput(ip);
00230
00231     end_op();
00232
00233     return 0;

```



```

00234
00235 bad:
00236     iunlockput(dp);
00237     end_op();
00238     return -1;
00239 }
00240
00241 static struct inode*
00242 create(char *path, short type, short major, short minor)
00243 {
00244     struct inode *ip, *dp;
00245     char name[DIRSIZ];
00246
00247     if((dp = nameiparent(path, name)) == 0)
00248         return 0;
00249     ilock(dp);
00250
00251     if((ip = dirlookup(dp, name, 0)) != 0){
00252         iunlockput(dp);
00253         ilock(ip);
00254         if(type == T_FILE && ip->type == T_FILE)
00255             return ip;
00256         iunlockput(ip);
00257         return 0;
00258     }
00259
00260     if((ip = ialloc(dp->dev, type)) == 0)
00261         panic("create: ialloc");
00262
00263     ilock(ip);
00264     ip->major = major;
00265     ip->minor = minor;
00266     ip->nlink = 1;
00267     iupdate(ip);
00268
00269     if(type == T_DIR){ // Create . and .. entries.
00270         dp->nlink++; // for ".."
00271         iupdate(dp);
00272         // No ip->nlink++ for ".": avoid cyclic ref count.
00273         if(dirlink(ip, ".", ip->inum) < 0 || dirlink(ip, "..", dp->inum) < 0)
00274             panic("create dots");
00275     }
00276
00277     if(dirlink(dp, name, ip->inum) < 0)
00278         panic("create: dirlink");
00279
00280     iunlockput(dp);
00281
00282     return ip;
00283 }
00284
00285 int
00286 sys_open(void)
00287 {
00288     char *path;
00289     int fd, omode;
00290     struct file *f;
00291     struct inode *ip;
00292
00293     if(argstr(0, &path) < 0 || argint(1, &omode) < 0)
00294         return -1;
00295
00296     begin_op();
00297
00298     if(omode & O_CREATE){
00299         ip = create(path, T_FILE, 0, 0);
00300         if(ip == 0){
00301             end_op();
00302             return -1;
00303         }
00304     } else {
00305         if((ip = namei(path)) == 0){
00306             end_op();
00307             return -1;
00308         }
00309         ilock(ip);
00310         if(ip->type == T_DIR && omode != O_RDONLY){
00311             iunlockput(ip);
00312             end_op();
00313             return -1;
00314         }
00315     }
00316
00317     if((f = filealloc()) == 0 || (fd = fdalloc(f)) < 0){
00318         if(f)
00319             fileclose(f);
00320         iunlockput(ip);

```

```

00321     end_op();
00322     return -1;
00323 }
00324 iunlock(ip);
00325 end_op();
00326
00327 f->type = FD_INODE;
00328 f->ip = ip;
00329 f->off = 0;
00330 f->readable = !(omode & O_WRONLY);
00331 f->writable = (omode & O_WRONLY) || (omode & O_RDWR);
00332 return fd;
00333 }
00334
00335 int
00336 sys_mkdir(void)
00337 {
00338     char *path;
00339     struct inode *ip;
00340
00341     begin_op();
00342     if(argstr(0, &path) < 0 || (ip = create(path, T_DIR, 0, 0)) == 0){
00343         end_op();
00344         return -1;
00345     }
00346     iunlockput(ip);
00347     end_op();
00348     return 0;
00349 }
00350
00351 int
00352 sys_mknod(void)
00353 {
00354     struct inode *ip;
00355     char *path;
00356     int major, minor;
00357
00358     begin_op();
00359     if((argstr(0, &path)) < 0 ||
00360        argint(1, &major) < 0 ||
00361        argint(2, &minor) < 0 ||
00362        (ip = create(path, T_DEV, major, minor)) == 0){
00363         end_op();
00364         return -1;
00365     }
00366     iunlockput(ip);
00367     end_op();
00368     return 0;
00369 }
00370
00371 int
00372 sys_chdir(void)
00373 {
00374     char *path;
00375     struct inode *ip;
00376     struct proc *curproc = myproc();
00377
00378     begin_op();
00379     if(argstr(0, &path) < 0 || (ip = namei(path)) == 0){
00380         end_op();
00381         return -1;
00382     }
00383     ilock(ip);
00384     if(ip->type != T_DIR){
00385         iunlockput(ip);
00386         end_op();
00387         return -1;
00388     }
00389     iunlock(ip);
00390     iput(curproc->cwd);
00391     end_op();
00392     curproc->cwd = ip;
00393     return 0;
00394 }
00395
00396 int
00397 sys_exec(void)
00398 {
00399     char *path, *argv[MAXARG];
00400     int i;
00401     uint uargv, uarg;
00402
00403     if(argstr(0, &path) < 0 || argint(1, (int*)&uargv) < 0){
00404         return -1;
00405     }
00406     memset(argv, 0, sizeof(argv));
00407     for(i=0;; i++){

```

```

00408     if(i >= NELEM(argv))
00409         return -1;
00410     if(fetchint(uargv+4*i, (int*)&uarg) < 0)
00411         return -1;
00412     if(uarg == 0){
00413         argv[i] = 0;
00414         break;
00415     }
00416     if(fetchstr(uarg, &argv[i]) < 0)
00417         return -1;
00418 }
00419 return exec(path, argv);
00420 }
00421
00422 int
00423 sys_pipe(void)
00424 {
00425     int *fd;
00426     struct file *rf, *wf;
00427     int fd0, fd1;
00428
00429     if(argptr(0, (void*)&fd, 2*sizeof(fd[0])) < 0)
00430         return -1;
00431     if(pipealloc(&rf, &wf) < 0)
00432         return -1;
00433     fd0 = -1;
00434     if((fd0 = fdalloc(rf)) < 0 || (fd1 = fdalloc(wf)) < 0){
00435         if(fd0 >= 0)
00436             myproc()->ofile[fd0] = 0;
00437         fileclose(rf);
00438         fileclose(wf);
00439         return -1;
00440     }
00441     fd[0] = fd0;
00442     fd[1] = fd1;
00443     return 0;
00444 }

```

5.210 sysfile.d File Reference

5.211 sysfile.d

[Go to the documentation of this file.](#)

```

00001 sysfile.o: sysfile.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 stat.h mmu.h proc.h fs.h spinlock.h sleeplock.h file.h fcntl.h

```

5.212 sysproc.c File Reference

```

#include "types.h"
#include "x86.h"
#include "defs.h"
#include "date.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "uproc.h"
#include "spinlock.h"

```

Functions

- int [sys_chpr](#) (void)
- int [sys_cps](#) (void)
- int [sys_exit](#) (void)
- int [sys_fork](#) (void)
- int [sys_getpid](#) (void)
- int [sys_getprocs](#) (void)
- int [sys_halt](#) (void)
- int [sys_kill](#) (void)
- int [sys_sbrk](#) (void)
- int [sys_sleep](#) (void)
- int [sys_uptime](#) (void)
- int [sys_wait](#) (void)

Variables

- struct {
 struct [spinlock](#) lock
 struct [proc](#) proc [NPROC]
} [ptable](#)

5.212.1 Function Documentation

5.212.1.1 sys_chpr()

```
int sys_chpr (  
    void )
```

Definition at line [125](#) of file [sysproc.c](#).

```
00126 {  
00127     int pid, pr;  
00128     if(argint(0, &pid) < 0)  
00129         return -1;  
00130     if(argint(1, &pr) < 0)  
00131         return -1;  
00132  
00133     return chpr(pid, pr);  
00134 }
```

5.212.1.2 sys_cps()

```
int sys_cps (  
    void )
```

Definition at line [119](#) of file [sysproc.c](#).

```
00120 {  
00121     return cps();  
00122 }
```

5.212.1.3 sys_exit()

```
int sys_exit (
    void )
```

Definition at line 25 of file [sysproc.c](#).

```
00026 {
00027     exit();
00028     return 0; // not reached
00029 }
```

5.212.1.4 sys_fork()

```
int sys_fork (
    void )
```

Definition at line 19 of file [sysproc.c](#).

```
00020 {
00021     return fork();
00022 }
```

5.212.1.5 sys_getpid()

```
int sys_getpid (
    void )
```

Definition at line 48 of file [sysproc.c](#).

```
00049 {
00050     return myproc()->pid;
00051 }
```

5.212.1.6 sys_getprocs()

```
int sys_getprocs (
    void )
```

Definition at line 137 of file [sysproc.c](#).

```
00137     {
00138         // declare local variables for max uproc size, struct uproc, and counter
00139         int max;
00140         struct uproc *p;
00141         int i=0;
00142
00143         // if argint has trouble, return error -1
00144         if(argint(0,&max) < 0)
00145             return -1;
00146
00147         // if argptr for allocating struct size has trouble, return -1
00148         if(argptr(1,(char**)&p, max*sizeof(struct uproc)) < 0)
00149             return -1;
00150
00151         // create pointer to ptable processes
00152         struct proc *ptr = ptable.proc;
00153
00154         // loop through ptable
00155         for(; ptr < &ptable.proc[NPROC]; ptr++)
00156         {
```

```
00157     if(!(ptr->state == UNUSED))
00158     {
00159         // if the process in ptable is not UNUSED, assign pid, parent pid, and name to uproc
00160         p[i].pid = ptr->pid;
00161         p[i].ppid = ptr->parent->pid;
00162         strncpy(p[i].name, ptr->name, 16);
00163         // add 1 to the process counter
00164         i++;
00165     }
00166 }
00167
00168 // return the number of processes that are not UNUSED
00169 return i;
00170 }
```

5.212.1.7 sys_halt()

```
int sys_halt (
    void )
```

Definition at line 111 of file [sysproc.c](#).

```
00112 {
00113     outb(0xf4, 0x00);
00114     return 0;
00115 }
```

5.212.1.8 sys_kill()

```
int sys_kill (
    void )
```

Definition at line 38 of file [sysproc.c](#).

```
00039 {
00040     int pid;
00041
00042     if(argint(0, &pid) < 0)
00043         return -1;
00044     return kill(pid);
00045 }
```

5.212.1.9 sys_sbrk()

```
int sys_sbrk (
    void )
```

Definition at line 54 of file [sysproc.c](#).

```
00055 {
00056     int addr;
00057     int n;
00058
00059     if(argint(0, &n) < 0)
00060         return -1;
00061     addr = myproc()->sz;
00062     if(growproc(n) < 0)
00063         return -1;
00064     return addr;
00065 }
```

5.212.1.10 sys_sleep()

```
int sys_sleep (
    void )
```

Definition at line 68 of file [sysproc.c](#).

```
00069 {
00070     int n;
00071     uint ticks0;
00072
00073     if(argint(0, &n) < 0)
00074         return -1;
00075     acquire(&tickslock);
00076     ticks0 = ticks;
00077     while(ticks - ticks0 < n){
00078         if(myproc()->killed){
00079             release(&tickslock);
00080             return -1;
00081         }
00082         sleep(&ticks, &tickslock);
00083     }
00084     release(&tickslock);
00085     return 0;
00086 }
```

5.212.1.11 sys_uptime()

```
int sys_uptime (
    void )
```

Definition at line 91 of file [sysproc.c](#).

```
00092 {
00093     uint xticks;
00094
00095     acquire(&tickslock);
00096     xticks = ticks;
00097     release(&tickslock);
00098     return xticks;
00099 }
```

5.212.1.12 sys_wait()

```
int sys_wait (
    void )
```

Definition at line 32 of file [sysproc.c](#).

```
00033 {
00034     return wait();
00035 }
```

5.212.2 Variable Documentation

5.212.2.1 lock

```
struct spinlock lock
```

Definition at line 13 of file [sysproc.c](#).

5.212.2.2 proc

```
struct proc proc[NPROC]
```

Definition at line 14 of file [sysproc.c](#).

5.212.2.3

```
struct { ... } ptable
```

Referenced by [sys_getprocs\(\)](#).

5.213 sysproc.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "x86.h"
00003 #include "defs.h"
00004 #include "date.h"
00005 #include "param.h"
00006 #include "memlayout.h"
00007 #include "mmu.h"
00008 #include "proc.h"
00009 #include "uproc.h"
00010 #include "spinlock.h"
00011
00012 struct {
00013     struct spinlock lock;
00014     struct proc proc[NPROC];
00015 } ptable;
00016
00017
00018 int
00019 sys_fork(void)
00020 {
00021     return fork();
00022 }
00023
00024 int
00025 sys_exit(void)
00026 {
00027     exit();
00028     return 0; // not reached
00029 }
00030
00031 int
00032 sys_wait(void)
00033 {
00034     return wait();
00035 }
00036
00037 int
00038 sys_kill(void)
00039 {
00040     int pid;
```



```

00041
00042     if(argint(0, &pid) < 0)
00043         return -1;
00044     return kill(pid);
00045 }
00046
00047 int
00048 sys_getpid(void)
00049 {
00050     return myproc()->pid;
00051 }
00052
00053 int
00054 sys_sbrk(void)
00055 {
00056     int addr;
00057     int n;
00058
00059     if(argint(0, &n) < 0)
00060         return -1;
00061     addr = myproc()->sz;
00062     if(growproc(n) < 0)
00063         return -1;
00064     return addr;
00065 }
00066
00067 int
00068 sys_sleep(void)
00069 {
00070     int n;
00071     uint ticks0;
00072
00073     if(argint(0, &n) < 0)
00074         return -1;
00075     acquire(&tickslock);
00076     ticks0 = ticks;
00077     while(ticks - ticks0 < n){
00078         if(myproc()->killed){
00079             release(&tickslock);
00080             return -1;
00081         }
00082         sleep(&ticks, &tickslock);
00083     }
00084     release(&tickslock);
00085     return 0;
00086 }
00087
00088 // return how many clock tick interrupts have occurred
00089 // since start.
00090 int
00091 sys_uptime(void)
00092 {
00093     uint xticks;
00094
00095     acquire(&tickslock);
00096     xticks = ticks;
00097     release(&tickslock);
00098     return xticks;
00099 }
00100
00101 /*int
00102 sys_halt(void)
00103 {
00104     char *p = "Shutdown";
00105     for( ; *p; p++)
00106         outw(0xB004, 0x2000);
00107     return 0;
00108 }*/
00109
00110 int
00111 sys_halt(void)
00112 {
00113     outb(0xf4, 0x00);
00114     return 0;
00115 }
00116
00117
00118 int
00119 sys_cps(void)
00120 {
00121     return cps();
00122 }
00123
00124 int
00125 sys_chpr(void)
00126 {
00127     int pid, pr;

```

```

00128     if(argint(0, &pid) < 0)
00129         return -1;
00130     if(argint(1, &pr) < 0)
00131         return -1;
00132
00133     return chpr(pid, pr);
00134 }
00135
00136 int
00137 sys_getprocs(void){
00138     // declare local variables for max uproc size, struct uproc, and counter
00139     int max;
00140     struct uproc *p;
00141     int i=0;
00142
00143     // if argint has trouble, return error -1
00144     if(argint(0, &max) < 0)
00145         return -1;
00146
00147     // if argptr for allocating struct size has trouble, return -1
00148     if(argptr(1, (char**)&p, max*sizeof(struct uproc)) < 0)
00149         return -1;
00150
00151     // create pointer to ptable processes
00152     struct proc *ptr = ptable.proc;
00153
00154     // loop through ptable
00155     for(; ptr < &ptable.proc[NPROC]; ptr++)
00156     {
00157         if(!(ptr->state == UNUSED))
00158             {
00159                 // if the process in ptable is not UNUSED, assign pid, parent pid, and name to uproc
00160                 p[i].pid = ptr->pid;
00161                 p[i].ppid = ptr->parent->pid;
00162                 strncpy(p[i].name, ptr->name, 16);
00163                 // add 1 to the process counter
00164                 i++;
00165             }
00166     }
00167
00168     // return the number of processes that are not UNUSED
00169     return i;
00170 }
00171
00172
00173

```

5.214 sysproc.d File Reference

5.215 sysproc.d

[Go to the documentation of this file.](#)

```

00001 sysproc.o: sysproc.c /usr/include/stdc-predef.h types.h x86.h defs.h \
00002     date.h param.h memlayout.h mmu.h proc.h uproc.h spinlock.h

```

5.216 trap.c File Reference

```

#include "types.h"
#include "defs.h"
#include "param.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"
#include "traps.h"
#include "spinlock.h"

```

Functions

- void `idtinit` (void)
- void `trap` (struct `trapframe` *tf)
- void `tvinit` (void)

Variables

- struct `gatedesc` `idt` [256]
- `uint` `ticks`
- struct `spinlock` `tickslock`
- `uint` `vectors` []

5.216.1 Function Documentation

5.216.1.1 idtinit()

```
void idtinit (
    void )
```

Definition at line 30 of file `trap.c`.

```
00031 {
00032     lidt(idt, sizeof(idt));
00033 }
```

5.216.1.2 trap()

```
void trap (
    struct trapframe * tf )
```

Definition at line 37 of file `trap.c`.

```
00038 {
00039     if(tf->trapno == T_SYSCALL) {
00040         if(myproc()->killed)
00041             exit();
00042         myproc()->tf = tf;
00043         syscall();
00044         if(myproc()->killed)
00045             exit();
00046         return;
00047     }
00048     switch(tf->trapno) {
00049     case T_IRQ0 + IRQ_TIMER:
00050         if(cpuid() == 0) {
00051             acquire(&tickslock);
00052             ticks++;
00053             wakeup(&ticks);
00054             release(&tickslock);
00055         }
00056         lapiceoi();
00057         break;
00058     case T_IRQ0 + IRQ_IDE:
00059         ideintr();
00060         lapiceoi();
00061         break;
00062     case T_IRQ0 + IRQ_IDE+1:
00063         // ...
```

```

00064     // Bochs generates spurious IDE1 interrupts.
00065     break;
00066 case T_IRQ0 + IRQ_KBD:
00067     kbdtintr();
00068     lapiceoi();
00069     break;
00070 case T_IRQ0 + IRQ_COM1:
00071     uartintr();
00072     lapiceoi();
00073     break;
00074 case T_IRQ0 + 7:
00075 case T_IRQ0 + IRQ_SPURIOUS:
00076     cprintf("cpu%d: spurious interrupt at %x:%x\n",
00077             cpuid(), tf->cs, tf->eip);
00078     lapiceoi();
00079     break;
00080 //PAGEBREAK: 13
00081 default:
00082     if(myproc() == 0 || (tf->cs&3) == 0){
00083         // In kernel, it must be our mistake.
00084         cprintf("unexpected trap %d from cpu %d eip %x (cr2=0x%x)\n",
00085                 tf->trapno, cpuid(), tf->eip, rcr2());
00086         panic("trap");
00087     }
00088     // In user space, assume process misbehaved.
00089     cprintf("pid %d %s: trap %d err %d on cpu %d "
00090             "eip 0x%x addr 0x%x--kill proc\n",
00091             myproc()->pid, myproc()->name, tf->trapno,
00092             tf->err, cpuid(), tf->eip, rcr2());
00093     myproc()->killed = 1;
00094 }
00095 }
00096 // Force process exit if it has been killed and is in user space.
00097 // (If it is still executing in the kernel, let it keep running
00098 // until it gets to the regular system call return.)
00099 if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
00100     exit();
00101 // Force process to give up CPU on clock tick.
00102 // If interrupts were on while locks held, would need to check nlock.
00103 if(myproc() && myproc()->state == RUNNING &&
00104     tf->trapno == T_IRQ0+IRQ_TIMER)
00105     yield();
00106 // Check if the process has been killed since we yielded
00107 if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
00108     exit();
00109 }

```

5.216.1.3 tvinit()

```

void tvinit (
    void )

```

Definition at line 18 of file [trap.c](#).

```

00019 {
00020     int i;
00021     for(i = 0; i < 256; i++)
00022         SETGATE(idt[i], 0, SEG_KCODE<3, vectors[i], 0);
00023     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<3, vectors[T_SYSCALL], DPL_USER);
00024     initlock(&tickslock, "time");
00025 }

```

5.216.2 Variable Documentation

5.216.2.1 idt

```
struct gatedesc idt[256]
```

Definition at line 12 of file [trap.c](#).

Referenced by [idtinit\(\)](#), and [tvinit\(\)](#).

5.216.2.2 ticks

```
uint ticks
```

Definition at line 15 of file [trap.c](#).

Referenced by [sys_sleep\(\)](#), [sys_uptime\(\)](#), and [trap\(\)](#).

5.216.2.3 tickslock

```
struct spinlock tickslock
```

Definition at line 14 of file [trap.c](#).

Referenced by [sys_sleep\(\)](#), [sys_uptime\(\)](#), [trap\(\)](#), and [tvinit\(\)](#).

5.216.2.4 vectors

```
uint vectors[] [extern]
```

Referenced by [tvinit\(\)](#).

5.217 trap.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "defs.h"
00003 #include "param.h"
00004 #include "memlayout.h"
00005 #include "mmu.h"
00006 #include "proc.h"
00007 #include "x86.h"
00008 #include "traps.h"
00009 #include "spinlock.h"
00010
00011 // Interrupt descriptor table (shared by all CPUs).
00012 struct gatedesc idt[256];
00013 extern uint vectors[]; // in vectors.S: array of 256 entry pointers
00014 struct spinlock tickslock;
00015 uint ticks;
00016
00017 void
00018 tvinit(void)
00019 {
00020     int i;
00021
00022     for(i = 0; i < 256; i++)
00023         SETGATE(idt[i], 0, SEG_KCODE<3, vectors[i], 0);
00024     SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<3, vectors[T_SYSCALL], DPL_USER);
00025
00026     initlock(&tickslock, "time");
00027 }
00028
00029 void
00030 idtinit(void)
00031 {
00032     lidt(idt, sizeof(idt));
00033 }
00034
00035 //PAGEBREAK: 41
00036 void
00037 trap(struct trapframe *tf)
00038 {
00039     if(tf->trapno == T_SYSCALL){
00040         if(myproc()->killed)
00041             exit();
00042         myproc()->tf = tf;
00043         syscall();
00044         if(myproc()->killed)
00045             exit();
00046         return;
00047     }
00048
00049     switch(tf->trapno){
00050     case T_IRQ0 + IRQ_TIMER:
00051         if(cpuid() == 0){
00052             acquire(&tickslock);
00053             ticks++;
00054             wakeup(&ticks);
00055             release(&tickslock);
00056         }
00057         lapiceoi();
00058         break;
00059     case T_IRQ0 + IRQ_IDE:
00060         ideintr();
00061         lapiceoi();
00062         break;
00063     case T_IRQ0 + IRQ_IDE+1:
00064         // Bochs generates spurious IDE1 interrupts.
00065         break;
00066     case T_IRQ0 + IRQ_KBD:
00067         kbdintr();
00068         lapiceoi();
00069         break;
00070     case T_IRQ0 + IRQ_COM1:
00071         uartintr();
00072         lapiceoi();
00073         break;
00074     case T_IRQ0 + 7:
00075     case T_IRQ0 + IRQ_SPURIOUS:
00076         printf("cpu%d: spurious interrupt at %x:%x\n",
00077             cpuid(), tf->cs, tf->eip);
00078         lapiceoi();
00079         break;
00080
00081     //PAGEBREAK: 13
00082     default:

```

```

00083     if(myproc() == 0 || (tf->cs&3) == 0){
00084         // In kernel, it must be our mistake.
00085         cprintf("unexpected trap %d from cpu %d eip %x (cr2=0x%x)\n",
00086             tf->trapno, cpuid(), tf->eip, rcr2());
00087         panic("trap");
00088     }
00089     // In user space, assume process misbehaved.
00090     cprintf("pid %d %s: trap %d err %d on cpu %d "
00091         "eip 0x%x addr 0x%x--kill proc\n",
00092         myproc()->pid, myproc()->name, tf->trapno,
00093         tf->err, cpuid(), tf->eip, rcr2());
00094     myproc()->killed = 1;
00095 }
00096
00097 // Force process exit if it has been killed and is in user space.
00098 // (If it is still executing in the kernel, let it keep running
00099 // until it gets to the regular system call return.)
00100 if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
00101     exit();
00102
00103 // Force process to give up CPU on clock tick.
00104 // If interrupts were on while locks held, would need to check nlock.
00105 if(myproc() && myproc()->state == RUNNING &&
00106     tf->trapno == T_IRQ0+IRQ_TIMER)
00107     yield();
00108
00109 // Check if the process has been killed since we yielded
00110 if(myproc() && myproc()->killed && (tf->cs&3) == DPL_USER)
00111     exit();
00112 }

```

5.218 trap.d File Reference

5.219 trap.d

[Go to the documentation of this file.](#)

```

00001 trap.o: trap.c /usr/include/stdc-predef.h types.h defs.h param.h \
00002 memlayout.h mmu.h proc.h x86.h traps.h spinlock.h

```

5.220 traps.h File Reference

Macros

- `#define IRQ_COM1 4`
- `#define IRQ_ERROR 19`
- `#define IRQ_IDE 14`
- `#define IRQ_KBD 1`
- `#define IRQ_SPURIOUS 31`
- `#define IRQ_TIMER 0`
- `#define T_ALIGN 17`
- `#define T_BOUND 5`
- `#define T_BRKPT 3`
- `#define T_DBLFLT 8`
- `#define T_DEBUG 1`
- `#define T_DEFAULT 500`
- `#define T_DEVICE 7`
- `#define T_DIVIDE 0`
- `#define T_FPERR 16`
- `#define T_GPFLT 13`
- `#define T_ILLOP 6`
- `#define T_IRQ0 32`

- `#define T_MCHK` 18
- `#define T_NMI` 2
- `#define T_OFLOW` 4
- `#define T_PGFLT` 14
- `#define T_SEGNP` 11
- `#define T_SIMDERR` 19
- `#define T_STACK` 12
- `#define T_SYSCALL` 64
- `#define T_TSS` 10

5.220.1 Macro Definition Documentation

5.220.1.1 IRQ_COM1

```
#define IRQ_COM1 4
```

Definition at line 34 of file [traps.h](#).

5.220.1.2 IRQ_ERROR

```
#define IRQ_ERROR 19
```

Definition at line 36 of file [traps.h](#).

5.220.1.3 IRQ_IDE

```
#define IRQ_IDE 14
```

Definition at line 35 of file [traps.h](#).

5.220.1.4 IRQ_KBD

```
#define IRQ_KBD 1
```

Definition at line 33 of file [traps.h](#).

5.220.1.5 IRQ_SPURIOUS

```
#define IRQ_SPURIOUS 31
```

Definition at line 37 of file [traps.h](#).

5.220.1.6 IRQ_TIMER

```
#define IRQ_TIMER 0
```

Definition at line 32 of file [traps.h](#).

5.220.1.7 T_ALIGN

```
#define T_ALIGN 17
```

Definition at line 21 of file [traps.h](#).

5.220.1.8 T_BOUND

```
#define T_BOUND 5
```

Definition at line 9 of file [traps.h](#).

5.220.1.9 T_BRKPT

```
#define T_BRKPT 3
```

Definition at line 7 of file [traps.h](#).

5.220.1.10 T_DBLFLT

```
#define T_DBLFLT 8
```

Definition at line 12 of file [traps.h](#).

5.220.1.11 T_DEBUG

```
#define T_DEBUG 1
```

Definition at line 5 of file [traps.h](#).

5.220.1.12 T_DEFAULT

```
#define T_DEFAULT 500
```

Definition at line 28 of file [traps.h](#).

5.220.1.13 T_DEVICE

```
#define T_DEVICE 7
```

Definition at line 11 of file [traps.h](#).

5.220.1.14 T_DIVIDE

```
#define T_DIVIDE 0
```

Definition at line 4 of file [traps.h](#).

5.220.1.15 T_FPERR

```
#define T_FPERR 16
```

Definition at line 20 of file [traps.h](#).

5.220.1.16 T_GPFLT

```
#define T_GPFLT 13
```

Definition at line 17 of file [traps.h](#).

5.220.1.17 T_ILLOP

```
#define T_ILLOP 6
```

Definition at line 10 of file [traps.h](#).

5.220.1.18 T_IRQ0

```
#define T_IRQ0 32
```

Definition at line 30 of file [traps.h](#).

5.220.1.19 T_MCHK

```
#define T_MCHK 18
```

Definition at line 22 of file [traps.h](#).

5.220.1.20 T_NMI

```
#define T_NMI 2
```

Definition at line 6 of file [traps.h](#).

5.220.1.21 T_OFLOW

```
#define T_OFLOW 4
```

Definition at line 8 of file [traps.h](#).

5.220.1.22 T_PGFLT

```
#define T_PGFLT 14
```

Definition at line 18 of file [traps.h](#).

5.220.1.23 T_SEGNP

```
#define T_SEGNP 11
```

Definition at line 15 of file [traps.h](#).

5.220.1.24 T_SIMDERR

```
#define T_SIMDERR 19
```

Definition at line 23 of file [traps.h](#).

5.220.1.25 T_STACK

```
#define T_STACK 12
```

Definition at line 16 of file [traps.h](#).

5.220.1.26 T_SYSCALL

```
#define T_SYSCALL 64
```

Definition at line 27 of file [traps.h](#).

5.220.1.27 T_TSS

```
#define T_TSS 10
```

Definition at line 14 of file [traps.h](#).

5.221 traps.h

[Go to the documentation of this file.](#)

```
00001 // x86 trap and interrupt constants.
00002
00003 // Processor-defined:
00004 #define T_DIVIDE      0      // divide error
00005 #define T_DEBUG      1      // debug exception
00006 #define T_NMI        2      // non-maskable interrupt
00007 #define T_BRKPT      3      // breakpoint
00008 #define T_OFLOW      4      // overflow
00009 #define T_BOUND      5      // bounds check
00010 #define T_ILLOP      6      // illegal opcode
00011 #define T_DEVICE      7      // device not available
00012 #define T_DBLFLT      8      // double fault
00013 // #define T_COPROC    9      // reserved (not used since 486)
00014 #define T_TSS        10     // invalid task switch segment
00015 #define T_SEGNP      11     // segment not present
00016 #define T_STACK      12     // stack exception
00017 #define T_GPFLT      13     // general protection fault
00018 #define T_PGFLT      14     // page fault
00019 // #define T_RES       15     // reserved
00020 #define T_FPERR      16     // floating point error
00021 #define T_ALIGN      17     // alignment check
00022 #define T_MCHK       18     // machine check
00023 #define T_SIMDERR     19     // SIMD floating point error
00024
00025 // These are arbitrarily chosen, but with care not to overlap
00026 // processor defined exceptions or interrupt vectors.
00027 #define T_SYSCALL      64     // system call
00028 #define T_DEFAULT      500    // catchall
00029
00030 #define T_IRQ0         32     // IRQ 0 corresponds to int T_IRQ
00031
00032 #define IRQ_TIMER      0
00033 #define IRQ_KBD        1
00034 #define IRQ_COM1       4
00035 #define IRQ_IDE        14
00036 #define IRQ_ERROR      19
00037 #define IRQ_SPURIOUS   31
00038
```

5.222 types.h File Reference

Typedefs

- typedef [uint pde_t](#)
- typedef unsigned char [uchar](#)
- typedef unsigned int [uint](#)
- typedef unsigned short [ushort](#)

5.222.1 Typedef Documentation

5.222.1.1 pde_t

```
typedef uint pde\_t
```

Definition at line 4 of file [types.h](#).

5.222.1.2 uchar

```
typedef unsigned char uchar
```

Definition at line 3 of file [types.h](#).

5.222.1.3 uint

```
typedef unsigned int uint
```

Definition at line 1 of file [types.h](#).

5.222.1.4 ushort

```
typedef unsigned short ushort
```

Definition at line 2 of file [types.h](#).

5.223 types.h

[Go to the documentation of this file.](#)

```
00001 typedef unsigned int    uint;
00002 typedef unsigned short   ushort;
00003 typedef unsigned char     uchar;
00004 typedef uint pde_t;
```

5.224 uart.c File Reference

```
#include "types.h"
#include "defs.h"
#include "param.h"
#include "traps.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "file.h"
#include "mmu.h"
#include "proc.h"
#include "x86.h"
```

Macros

- `#define COM1 0x3f8`

Functions

- static int [uartgetc](#) (void)
- void [uartinit](#) (void)
- void [uartintr](#) (void)
- void [uartputc](#) (int c)

Variables

- static int [uart](#)

5.224.1 Macro Definition Documentation

5.224.1.1 COM1

```
#define COM1 0x3f8
```

Definition at line 15 of file [uart.c](#).

5.224.2 Function Documentation

5.224.2.1 [uartgetc\(\)](#)

```
static int uartgetc (  
    void ) [static]
```

Definition at line 64 of file [uart.c](#).

```
00065 {  
00066     if(!uart)  
00067         return -1;  
00068     if(!(inb(COM1+5) & 0x01))  
00069         return -1;  
00070     return inb(COM1+0);  
00071 }
```

Referenced by [uartintr\(\)](#).

5.224.2.2 uartinit()

```
void uartinit (
    void )
```

Definition at line 20 of file [uart.c](#).

```
00021 {
00022     char *p;
00023
00024     // Turn off the FIFO
00025     outb(COM1+2, 0);
00026
00027     // 9600 baud, 8 data bits, 1 stop bit, parity off.
00028     outb(COM1+3, 0x80); // Unlock divisor
00029     outb(COM1+0, 115200/9600);
00030     outb(COM1+1, 0);
00031     outb(COM1+3, 0x03); // Lock divisor, 8 data bits.
00032     outb(COM1+4, 0);
00033     outb(COM1+1, 0x01); // Enable receive interrupts.
00034
00035     // If status is 0xFF, no serial port.
00036     if(inb(COM1+5) == 0xFF)
00037         return;
00038     uart = 1;
00039
00040     // Acknowledge pre-existing interrupt conditions;
00041     // enable interrupts.
00042     inb(COM1+2);
00043     inb(COM1+0);
00044     ioapicenable(IRQ_COM1, 0);
00045
00046     // Announce that we're here.
00047     for(p="xv6...\n"; *p; p++)
00048         uartputc(*p);
00049 }
```

5.224.2.3 uartintr()

```
void uartintr (
    void )
```

Definition at line 74 of file [uart.c](#).

```
00075 {
00076     consoleintr(uartgetc);
00077 }
```

Referenced by [trap\(\)](#).

5.224.2.4 uartputc()

```
void uartputc (
    int c )
```

Definition at line 52 of file [uart.c](#).

```
00053 {
00054     int i;
00055
00056     if(!uart)
00057         return;
00058     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)
00059         microdelay(10);
00060     outb(COM1+0, c);
00061 }
```

Referenced by [consputc\(\)](#), and [uartinit\(\)](#).

5.224.3 Variable Documentation

5.224.3.1 uart

```
int uart [static]
```

Definition at line 17 of file [uart.c](#).

Referenced by [uartgetc\(\)](#), [uartinit\(\)](#), and [uartputc\(\)](#).

5.225 uart.c

[Go to the documentation of this file.](#)

```
00001 // Intel 8250 serial port (UART).
00002
00003 #include "types.h"
00004 #include "defs.h"
00005 #include "param.h"
00006 #include "traps.h"
00007 #include "spinlock.h"
00008 #include "sleeplock.h"
00009 #include "fs.h"
00010 #include "file.h"
00011 #include "mmu.h"
00012 #include "proc.h"
00013 #include "x86.h"
00014
00015 #define COM1      0x3f8
00016
00017 static int uart;    // is there a uart?
00018
00019 void
00020 uartinit(void)
00021 {
00022     char *p;
00023
00024     // Turn off the FIFO
00025     outb(COM1+2, 0);
00026
00027     // 9600 baud, 8 data bits, 1 stop bit, parity off.
00028     outb(COM1+3, 0x80);    // Unlock divisor
00029     outb(COM1+0, 115200/9600);
00030     outb(COM1+1, 0);
00031     outb(COM1+3, 0x03);    // Lock divisor, 8 data bits.
00032     outb(COM1+4, 0);
00033     outb(COM1+1, 0x01);    // Enable receive interrupts.
00034
00035     // If status is 0xFF, no serial port.
00036     if(inb(COM1+5) == 0xFF)
00037         return;
00038     uart = 1;
00039
00040     // Acknowledge pre-existing interrupt conditions;
00041     // enable interrupts.
00042     inb(COM1+2);
00043     inb(COM1+0);
00044     ioapicenable(IRQ_COM1, 0);
00045
00046     // Announce that we're here.
00047     for(p="xv6...\n"; *p; p++)
00048         uartputc(*p);
00049 }
00050
00051 void
00052 uartputc(int c)
00053 {
00054     int i;
00055
00056     if(!uart)
00057         return;
00058     for(i = 0; i < 128 && !(inb(COM1+5) & 0x20); i++)
```

```

00059     microdelay(10);
00060     outb(COM1+0, c);
00061 }
00062
00063 static int
00064 uartgetc(void)
00065 {
00066     if(!uart)
00067         return -1;
00068     if(!(inb(COM1+5) & 0x01))
00069         return -1;
00070     return inb(COM1+0);
00071 }
00072
00073 void
00074 uartintr(void)
00075 {
00076     consoleintr(uartgetc);
00077 }

```

5.226 uart.d File Reference

5.227 uart.d

[Go to the documentation of this file.](#)

```

00001 uart.o: uart.c /usr/include/stdc-predef.h types.h defs.h param.h traps.h \
00002 spinlock.h sleeplock.h fs.h file.h mmu.h proc.h x86.h

```

5.228 ulib.c File Reference

```

#include "types.h"
#include "stat.h"
#include "fcntl.h"
#include "user.h"
#include "x86.h"

```

Functions

- int [atoi](#) (const char *s)
- char * [gets](#) (char *buf, int max)
- void * [memmove](#) (void *vdst, const void *vsrc, int n)
- void * [memset](#) (void *dst, int c, [uint](#) n)
- int [stat](#) (const char *n, struct [stat](#) *st)
- char * [strchr](#) (const char *s, char c)
- int [strcmp](#) (const char *p, const char *q)
- char * [strcpy](#) (char *s, const char *t)
- [uint](#) [strlen](#) (const char *s)

5.228.1 Function Documentation

5.228.1.1 atoi()

```
int atoi (
    const char * s )
```

Definition at line 85 of file [ulib.c](#).

```
00086 {
00087     int n;
00088     n = 0;
00089     while ('0' <= *s && *s <= '9')
00090         n = n*10 + *s++ - '0';
00091     return n;
00092 }
00093 }
```

Referenced by [main\(\)](#).

5.228.1.2 gets()

```
char * gets (
    char * buf,
    int max )
```

Definition at line 53 of file [ulib.c](#).

```
00054 {
00055     int i, cc;
00056     char c;
00057     for(i=0; i+1 < max; ){
00058         cc = read(0, &c, 1);
00059         if(cc < 1)
00060             break;
00061         buf[i++] = c;
00062         if(c == '\n' || c == '\r')
00063             break;
00064     }
00065     buf[i] = '\0';
00066     return buf;
00067 }
00068 }
```

Referenced by [getcmd\(\)](#).

5.228.1.3 memmove()

```
void * memmove (
    void * vdst,
    const void * vsrc,
    int n )
```

Definition at line 96 of file [ulib.c](#).

```
00097 {
00098     char *dst;
00099     const char *src;
00100     dst = vdst;
00101     src = vsrc;
00102     while(n-- > 0)
00103         *dst++ = *src++;
00104     return vdst;
00105 }
00106 }
```

5.228.1.4 `memset()`

```
void * memset (
    void * dst,
    int c,
    uint n )
```

Definition at line 37 of file `ulib.c`.

```
00038 {
00039     stosb(dst, c, n);
00040     return dst;
00041 }
```

Referenced by `allocproc()`, `allocuvn()`, `backcmd()`, `bigfile()`, `bzero()`, `cgaputc()`, `concreate()`, `execcmd()`, `fmtname()`, `fourfiles()`, `getcmd()`, `ialloc()`, `inituvm()`, `kfree()`, `listcmd()`, `main()`, `pipecmd()`, `redircmd()`, `setupkvm()`, `sharedfd()`, `sys_exec()`, `sys_unlink()`, `userinit()`, and `walkpgdir()`.

5.228.1.5 `stat()`

```
int stat (
    const char * n,
    struct stat * st )
```

Definition at line 71 of file `ulib.c`.

```
00072 {
00073     int fd;
00074     int r;
00075
00076     fd = open(n, O_RDONLY);
00077     if(fd < 0)
00078         return -1;
00079     r = fstat(fd, st);
00080     close(fd);
00081     return r;
00082 }
```

5.228.1.6 `strchr()`

```
char * strchr (
    const char * s,
    char c )
```

Definition at line 44 of file `ulib.c`.

```
00045 {
00046     for(; *s; s++)
00047         if(*s == c)
00048             return (char*)s;
00049     return 0;
00050 }
```

Referenced by `gettoken()`, `grep()`, `peek()`, and `wc()`.

5.228.1.7 strcmp()

```
int strcmp (
    const char * p,
    const char * q )
```

Definition at line 19 of file [ulib.c](#).

```
00020 {
00021     while(*p && *p == *q)
00022         p++, q++;
00023     return (uchar)*p - (uchar)*q;
00024 }
```

5.228.1.8 strcpy()

```
char * strcpy (
    char * s,
    const char * t )
```

Definition at line 8 of file [ulib.c](#).

```
00009 {
00010     char *os;
00011
00012     os = s;
00013     while((*s++ = *t++) != 0)
00014         ;
00015     return os;
00016 }
```

Referenced by [ls\(\)](#), and [main\(\)](#).

5.228.1.9 strlen()

```
uint strlen (
    const char * s )
```

Definition at line 27 of file [ulib.c](#).

```
00028 {
00029     int n;
00030
00031     for(n = 0; s[n]; n++)
00032         ;
00033     return n;
00034 }
```

Referenced by [exec\(\)](#), [fmtname\(\)](#), [ls\(\)](#), [main\(\)](#), [parsecmd\(\)](#), and [printf\(\)](#).

5.229 ulib.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "stat.h"
00003 #include "fcntl.h"
00004 #include "user.h"
00005 #include "x86.h"
00006
00007 char*
00008 strcpy(char *s, const char *t)
00009 {
00010     char *os;
00011
00012     os = s;
00013     while ((*s++ = *t++) != 0)
00014         ;
00015     return os;
00016 }
00017
00018 int
00019 strcmp(const char *p, const char *q)
00020 {
00021     while (*p && *p == *q)
00022         p++, q++;
00023     return (uchar)*p - (uchar)*q;
00024 }
00025
00026 uint
00027 strlen(const char *s)
00028 {
00029     int n;
00030
00031     for(n = 0; s[n]; n++)
00032         ;
00033     return n;
00034 }
00035
00036 void*
00037 memset(void *dst, int c, uint n)
00038 {
00039     stosb(dst, c, n);
00040     return dst;
00041 }
00042
00043 char*
00044 strchr(const char *s, char c)
00045 {
00046     for(; *s; s++)
00047         if(*s == c)
00048             return (char*)s;
00049     return 0;
00050 }
00051
00052 char*
00053 gets(char *buf, int max)
00054 {
00055     int i, cc;
00056     char c;
00057
00058     for(i=0; i+1 < max; ){
00059         cc = read(0, &c, 1);
00060         if(cc < 1)
00061             break;
00062         buf[i++] = c;
00063         if(c == '\n' || c == '\r')
00064             break;
00065     }
00066     buf[i] = '\0';
00067     return buf;
00068 }
00069
00070 int
00071 stat(const char *n, struct stat *st)
00072 {
00073     int fd;
00074     int r;
00075
00076     fd = open(n, O_RDONLY);
00077     if(fd < 0)
00078         return -1;
00079     r = fstat(fd, st);
00080     close(fd);
00081     return r;
00082 }

```

```

00083
00084 int
00085 atoi(const char *s)
00086 {
00087     int n;
00088
00089     n = 0;
00090     while('0' <= *s && *s <= '9')
00091         n = n*10 + *s++ - '0';
00092     return n;
00093 }
00094
00095 void*
00096 memmove(void *vdst, const void *vsrc, int n)
00097 {
00098     char *dst;
00099     const char *src;
00100
00101     dst = vdst;
00102     src = vsrc;
00103     while(n-- > 0)
00104         *dst++ = *src++;
00105     return vdst;
00106 }

```

5.230 ulib.d File Reference

5.231 ulib.d

[Go to the documentation of this file.](#)

```

00001 ulib.o: ulib.c /usr/include/stdc-predef.h types.h stat.h fcntl.h user.h \
00002     x86.h

```

5.232 umalloc.c File Reference

```

#include "types.h"
#include "stat.h"
#include "user.h"
#include "param.h"

```

Classes

- union [header](#)

Typedefs

- typedef long [Align](#)
- typedef union [header](#) [Header](#)

Functions

- void [free](#) (void *ap)
- void * [malloc](#) (uint nbytes)
- static [Header](#) * [morecore](#) (uint nu)

Variables

- static [Header](#) [base](#)
- static [Header](#) * [freep](#)

5.232.1 Typedef Documentation

5.232.1.1 Align

```
typedef long Align
```

Definition at line 9 of file [umalloc.c](#).

5.232.1.2 Header

```
typedef union header Header
```

Definition at line 19 of file [umalloc.c](#).

5.232.2 Function Documentation

5.232.2.1 free()

```
void free (
    void * ap )
```

Definition at line 25 of file [umalloc.c](#).

```
00026 {
00027     Header *bp, *p;
00028
00029     bp = (Header*)ap - 1;
00030     for(p = freep; !(bp > p && bp < p->s.ptr); p = p->s.ptr)
00031         if(p >= p->s.ptr && (bp > p || bp < p->s.ptr))
00032             break;
00033     if(bp + bp->s.size == p->s.ptr){
00034         bp->s.size += p->s.ptr->s.size;
00035         bp->s.ptr = p->s.ptr->s.ptr;
00036     } else
00037         bp->s.ptr = p->s.ptr;
00038     if(p + p->s.size == bp){
00039         p->s.size += bp->s.size;
00040         p->s.ptr = bp->s.ptr;
00041     } else
00042         p->s.ptr = bp;
00043     freep = p;
00044 }
```

Referenced by [mem\(\)](#), and [morecore\(\)](#).

5.232.2.2 malloc()

```
void * malloc (
    uint nbytes )
```

Definition at line 64 of file [umalloc.c](#).

```
00065 {
00066     Header *p, *prevp;
00067     uint nunits;
00068
00069     nunits = (nbytes + sizeof(Header) - 1)/sizeof(Header) + 1;
00070     if((prevp = freep) == 0){
00071         base.s.ptr = freep = prevp = &base;
00072         base.s.size = 0;
00073     }
00074     for(p = prevp->s.ptr; ; prevp = p, p = p->s.ptr){
00075         if(p->s.size >= nunits){
00076             if(p->s.size == nunits)
00077                 prevp->s.ptr = p->s.ptr;
00078             else {
00079                 p->s.size -= nunits;
00080                 p += p->s.size;
00081                 p->s.size = nunits;
00082             }
00083             freep = prevp;
00084             return (void*)(p + 1);
00085         }
00086         if(p == freep)
00087             if((p = morecore(nunits)) == 0)
00088                 return 0;
00089     }
00090 }
```

Referenced by [backcmd\(\)](#), [execcmd\(\)](#), [listcmd\(\)](#), [main\(\)](#), [mem\(\)](#), [pipecmd\(\)](#), and [redircmd\(\)](#).

5.232.2.3 morecore()

```
static Header * morecore (
    uint nu ) [static]
```

Definition at line 47 of file [umalloc.c](#).

```
00048 {
00049     char *p;
00050     Header *hp;
00051
00052     if(nu < 4096)
00053         nu = 4096;
00054     p = sbrk(nu * sizeof(Header));
00055     if(p == (char*)-1)
00056         return 0;
00057     hp = (Header*)p;
00058     hp->s.size = nu;
00059     free((void*)(hp + 1));
00060     return freep;
00061 }
```

Referenced by [malloc\(\)](#).

5.232.3 Variable Documentation

5.232.3.1 base

Header base [static]

Definition at line 21 of file [umalloc.c](#).

Referenced by [malloc\(\)](#), and [printint\(\)](#).

5.232.3.2 freep

Header* freep [static]

Definition at line 22 of file [umalloc.c](#).

Referenced by [free\(\)](#), [malloc\(\)](#), and [morecore\(\)](#).

5.233 umalloc.c

[Go to the documentation of this file.](#)

```

00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004 #include "param.h"
00005
00006 // Memory allocator by Kernighan and Ritchie,
00007 // The C programming Language, 2nd ed. Section 8.7.
00008
00009 typedef long Align;
00010
00011 union header {
00012     struct {
00013         union header *ptr;
00014         uint size;
00015     } s;
00016     Align x;
00017 };
00018
00019 typedef union header Header;
00020
00021 static Header base;
00022 static Header *freep;
00023
00024 void
00025 free(void *ap)
00026 {
00027     Header *bp, *p;
00028
00029     bp = (Header*)ap - 1;
00030     for(p = freep; !(bp > p && bp < p->s.ptr); p = p->s.ptr)
00031         if(p >= p->s.ptr && (bp > p || bp < p->s.ptr))
00032             break;
00033     if(bp + bp->s.size == p->s.ptr){
00034         bp->s.size += p->s.ptr->s.size;
00035         bp->s.ptr = p->s.ptr->s.ptr;
00036     } else
00037         bp->s.ptr = p->s.ptr;
00038     if(p + p->s.size == bp){
00039         p->s.size += bp->s.size;
00040         p->s.ptr = bp->s.ptr;
00041     } else
00042         p->s.ptr = bp;
00043     freep = p;
00044 }
00045
00046 static Header*
00047 morecore(uint nu)
00048 {
00049     char *p;

```

```

00050  Header *hp;
00051
00052  if(nu < 4096)
00053      nu = 4096;
00054  p = sbrk(nu * sizeof(Header));
00055  if(p == (char*)-1)
00056      return 0;
00057  hp = (Header*)p;
00058  hp->s.size = nu;
00059  free((void*)(hp + 1));
00060  return freep;
00061 }
00062
00063 void*
00064 malloc(uint nbytes)
00065 {
00066     Header *p, *prevp;
00067     uint nunits;
00068
00069     nunits = (nbytes + sizeof(Header) - 1)/sizeof(Header) + 1;
00070     if((prevp = freep) == 0){
00071         base.s.ptr = freep = prevp = &base;
00072         base.s.size = 0;
00073     }
00074     for(p = prevp->s.ptr; ; prevp = p, p = p->s.ptr){
00075         if(p->s.size >= nunits){
00076             if(p->s.size == nunits)
00077                 prevp->s.ptr = p->s.ptr;
00078             else {
00079                 p->s.size -= nunits;
00080                 p += p->s.size;
00081                 p->s.size = nunits;
00082             }
00083             freep = prevp;
00084             return (void*)(p + 1);
00085         }
00086         if(p == freep)
00087             if((p = morecore(nunits)) == 0)
00088                 return 0;
00089     }
00090 }

```

5.234 umalloc.d File Reference

5.235 umalloc.d

[Go to the documentation of this file.](#)

```

00001 umalloc.o: umalloc.c /usr/include/stdc-predef.h types.h stat.h user.h \
00002 param.h

```

5.236 uproc.h File Reference

Classes

- struct [uproc](#)

5.237 uproc.h

[Go to the documentation of this file.](#)

```

00001 struct uproc {
00002     int pid;
00003     int ppid;
00004     char name[16];
00005 };

```

5.238 user.h File Reference

Functions

- int [atoi](#) (const char *)
- int [chdir](#) (const char *)
- int [chpr](#) (int pid, int priority)
- int [close](#) (int)
- int [cps](#) (void)
- int [dup](#) (int)
- int [exec](#) (char *, char **)
- int [exit](#) (void) [__attribute__\(\(noreturn\)\)](#)
- int [fork](#) (void)
- void [free](#) (void *)
- int [fstat](#) (int fd, struct [stat](#) *)
- int [getpid](#) (void)
- int [getprocs](#) (int max, struct [uproc](#) *)
- char * [gets](#) (char *, int max)
- int [halt](#) (void)
- int [kill](#) (int)
- int [link](#) (const char *, const char *)
- void * [malloc](#) (uint)
- void * [memmove](#) (void *, const void *, int)
- void * [memset](#) (void *, int, uint)
- int [mkdir](#) (const char *)
- int [mknod](#) (const char *, short, short)
- int [open](#) (const char *, int)
- int [pipe](#) (int *)
- void [printf](#) (int, const char *,...)
- int [read](#) (int, void *, int)
- char * [sbrk](#) (int)
- int [sleep](#) (int)
- int [stat](#) (const char *, struct [stat](#) *)
- char * [strchr](#) (const char *, char c)
- int [strcmp](#) (const char *, const char *)
- char * [strcpy](#) (char *, const char *)
- uint [strlen](#) (const char *)
- int [unlink](#) (const char *)
- int [uptime](#) (void)
- int [wait](#) (void)
- int [write](#) (int, const void *, int)

5.238.1 Function Documentation

5.238.1.1 atoi()

```
int atoi (
    const char * s )
```

Definition at line 85 of file [ulib.c](#).

```
00086 {
00087     int n;
00088     n = 0;
00089     while ('0' <= *s && *s <= '9')
00090         n = n*10 + *s++ - '0';
00091     return n;
00092 }
00093 }
```

Referenced by [main\(\)](#).

5.238.1.2 chdir()

```
int chdir (
    const char * )
```

Referenced by [dirfile\(\)](#), [dirttest\(\)](#), [exitiputtest\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [main\(\)](#), [rmdot\(\)](#), and [subdir\(\)](#).

5.238.1.3 chpr()

```
int chpr (
    int pid,
    int priority )
```

Definition at line 559 of file [proc.c](#).

```
00560 {
00561     struct proc *p;
00562     acquire(&ptable.lock);
00563     for (p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
00564         if (p->pid == pid) {
00565             p->priority = priority;
00566             break;
00567         }
00568     }
00569     release(&ptable.lock);
00570     return pid;
00571 }
```

5.238.1.4 close()

```
int close (
    int )
```

Referenced by [argptest\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirfile\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [fsfull\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [ls\(\)](#), [main\(\)](#), [opentest\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [runcmd\(\)](#), [sharedfd\(\)](#), [stat\(\)](#), [subdir\(\)](#), [unlinkread\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.238.1.5 cps()

```
int cps (
    void )
```

Definition at line 537 of file [proc.c](#).

```
00538 {
00539     struct proc *p;
00540     //Enables interrupts on this processor.
00541     sti();
00542
00543     //Loop over process table looking for process with pid.
00544     acquire(&ptable.lock);
00545     cprintf("name \t pid \t state \t priority \n");
00546     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00547         if(p->state == SLEEPING)
00548             cprintf("%s \t %d \t SLEEPING \t %d \n ", p->name, p->pid, p->priority);
00549         else if(p->state == RUNNING)
00550             cprintf("%s \t %d \t RUNNING \t %d \n ", p->name, p->pid, p->priority);
00551         else if(p->state == RUNNABLE)
00552             cprintf("%s \t %d \t RUNNABLE \t %d \n ", p->name, p->pid, p->priority);
00553     }
00554     release(&ptable.lock);
00555     return 23;
00556 }
```

5.238.1.6 dup()

```
int dup (
    int )
```

Referenced by [main\(\)](#), and [runcmd\(\)](#).

5.238.1.7 exec()

```
int exec (
    char * path,
    char ** argv )
```

Definition at line 11 of file [exec.c](#).

```
00012 {
00013     char *s, *last;
00014     int i, off;
00015     uint argc, sz, sp, ustack[3+MAXARG+1];
00016     struct elfhdr elf;
00017     struct inode *ip;
00018     struct proghdr ph;
00019     pde_t *pgdir, *oldpgdir;
00020     struct proc *curproc = myproc();
00021
00022     begin_op();
00023
00024     if((ip = namei(path)) == 0){
00025         end_op();
00026         cprintf("exec: fail\n");
00027         return -1;
00028     }
00029     ilock(ip);
00030     pgdir = 0;
00031
00032     // Check ELF header
00033     if(readi(ip, (char*)&elf, 0, sizeof(elf)) != sizeof(elf))
00034         goto bad;
00035     if(elf.magic != ELF_MAGIC)
00036         goto bad;
```

```

00037
00038     if((pgdir = setupkvm()) == 0)
00039         goto bad;
00040
00041     // Load program into memory.
00042     sz = 0;
00043     for(i=0, off=elf.phoff; i<elf.phnum; i++, off+=sizeof(ph)){
00044         if(readi(ip, (char*)&ph, off, sizeof(ph)) != sizeof(ph))
00045             goto bad;
00046         if(ph.type != ELF_PROG_LOAD)
00047             continue;
00048         if(ph.memsz < ph.filesz)
00049             goto bad;
00050         if(ph.vaddr + ph.memsz < ph.vaddr)
00051             goto bad;
00052         if((sz = allocuvvm(pgdir, sz, ph.vaddr + ph.memsz)) == 0)
00053             goto bad;
00054         if(ph.vaddr % PGSIZE != 0)
00055             goto bad;
00056         if(loaduvvm(pgdir, (char*)ph.vaddr, ip, ph.off, ph.filesz) < 0)
00057             goto bad;
00058     }
00059     iunlockput(ip);
00060     end_op();
00061     ip = 0;
00062
00063     // Allocate two pages at the next page boundary.
00064     // Make the first inaccessible. Use the second as the user stack.
00065     sz = PGROUNDUP(sz);
00066     if((sz = allocuvvm(pgdir, sz, sz + 2*PGSIZE)) == 0)
00067         goto bad;
00068     clearpteu(pgdir, (char*)(sz - 2*PGSIZE));
00069     sp = sz;
00070
00071     // Push argument strings, prepare rest of stack in ustack.
00072     for(argc = 0; argv[argc]; argc++) {
00073         if(argc >= MAXARG)
00074             goto bad;
00075         sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
00076         if(copyout(pgdir, sp, argv[argc], strlen(argv[argc]) + 1) < 0)
00077             goto bad;
00078         ustack[3+argc] = sp;
00079     }
00080     ustack[3+argc] = 0;
00081
00082     ustack[0] = 0xffffffff; // fake return PC
00083     ustack[1] = argc;
00084     ustack[2] = sp - (argc+1)*4; // argv pointer
00085
00086     sp -= (3+argc+1) * 4;
00087     if(copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
00088         goto bad;
00089
00090     // Save program name for debugging.
00091     for(last=s=path; *s; s++)
00092         if(*s == '/')
00093             last = s+1;
00094     safestrcpy(curproc->name, last, sizeof(curproc->name));
00095
00096     // Commit to the user image.
00097     oldpgdir = curproc->pgdir;
00098     curproc->pgdir = pgdir;
00099     curproc->sz = sz;
00100     curproc->tf->eip = elf.entry; // main
00101     curproc->tf->esp = sp;
00102     switchvm(curproc);
00103     freevm(oldpgdir);
00104     return 0;
00105
00106 bad:
00107     if(pgdir)
00108         freevm(pgdir);
00109     if(ip){
00110         iunlockput(ip);
00111         end_op();
00112     }
00113     return -1;
00114 }

```

5.238.1.8 exit()

```
int exit (
```

```
void )
```

Definition at line 228 of file `proc.c`.

```
00229 {
00230     struct proc *curproc = myproc();
00231     struct proc *p;
00232     int fd;
00233
00234     if(curproc == initproc)
00235         panic("init exiting");
00236
00237     // Close all open files.
00238     for(fd = 0; fd < NOFILE; fd++){
00239         if(curproc->ofile[fd]){
00240             fileclose(curproc->ofile[fd]);
00241             curproc->ofile[fd] = 0;
00242         }
00243     }
00244
00245     begin_op();
00246     iput(curproc->cwd);
00247     end_op();
00248     curproc->cwd = 0;
00249
00250     acquire(&ptable.lock);
00251
00252     // Parent might be sleeping in wait().
00253     wakeup1(curproc->parent);
00254
00255     // Pass abandoned children to init.
00256     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00257         if(p->parent == curproc){
00258             p->parent = initproc;
00259             if(p->state == ZOMBIE)
00260                 wakeup1(initproc);
00261         }
00262     }
00263
00264     // Jump into the scheduler, never to return.
00265     curproc->state = ZOMBIE;
00266     sched();
00267     panic("zombie exit");
00268 }
```

5.238.1.9 fork()

```
int fork (
    void )
```

Definition at line 181 of file `proc.c`.

```
00182 {
00183     int i, pid;
00184     struct proc *np;
00185     struct proc *curproc = myproc();
00186
00187     // Allocate process.
00188     if((np = allocproc()) == 0){
00189         return -1;
00190     }
00191
00192     // Copy process state from proc.
00193     if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
00194         kfree(np->kstack);
00195         np->kstack = 0;
00196         np->state = UNUSED;
00197         return -1;
00198     }
00199     np->sz = curproc->sz;
00200     np->parent = curproc;
00201     *np->tf = *curproc->tf;
00202
00203     // Clear %eax so that fork returns 0 in the child.
00204     np->tf->eax = 0;
00205
00206     for(i = 0; i < NOFILE; i++)
00207         if(curproc->ofile[i])
00208             np->ofile[i] = filedup(curproc->ofile[i]);
```



```

00209  np->cwd = idup(curproc->cwd);
00210
00211  safestrcpy(np->name, curproc->name, sizeof(curproc->name));
00212
00213  pid = np->pid;
00214
00215  acquire(&ptable.lock);
00216
00217  np->state = RUNNABLE;
00218
00219  release(&ptable.lock);
00220
00221  return pid;
00222 }

```

5.238.1.10 free()

```

void free (
    void * ap )

```

Definition at line 25 of file [umalloc.c](#).

```

00026 {
00027     Header *bp, *p;
00028
00029     bp = (Header*)ap - 1;
00030     for(p = freep; !(bp > p && bp < p->s.ptr); p = p->s.ptr)
00031         if(p >= p->s.ptr && (bp > p || bp < p->s.ptr))
00032             break;
00033     if(bp + bp->s.size == p->s.ptr){
00034         bp->s.size += p->s.ptr->s.size;
00035         bp->s.ptr = p->s.ptr->s.ptr;
00036     } else
00037         bp->s.ptr = p->s.ptr;
00038     if(p + p->s.size == bp){
00039         p->s.size += bp->s.size;
00040         p->s.ptr = bp->s.ptr;
00041     } else
00042         p->s.ptr = bp;
00043     freep = p;
00044 }

```

Referenced by [mem\(\)](#), and [morecore\(\)](#).

5.238.1.11 fstat()

```

int fstat (
    int fd,
    struct stat * )

```

Referenced by [ls\(\)](#), and [stat\(\)](#).

5.238.1.12 getpid()

```

int getpid (
    void )

```

Referenced by [main\(\)](#), [mem\(\)](#), and [sbrktest\(\)](#).

5.238.1.13 getprocs()

```
int getprocs (
    int max,
    struct uproc * )
```

5.238.1.14 gets()

```
char * gets (
    char * buf,
    int max )
```

Definition at line 53 of file [ulib.c](#).

```
00054 {
00055     int i, cc;
00056     char c;
00057
00058     for(i=0; i+1 < max; ){
00059         cc = read(0, &c, 1);
00060         if(cc < 1)
00061             break;
00062         buf[i++] = c;
00063         if(c == '\n' || c == '\r')
00064             break;
00065     }
00066     buf[i] = '\0';
00067     return buf;
00068 }
```

Referenced by [getcmod\(\)](#).

5.238.1.15 halt()

```
int halt (
    void )
```

Referenced by [main\(\)](#).

5.238.1.16 kill()

```
int kill (
    int pid )
```

Definition at line 480 of file [proc.c](#).

```
00481 {
00482     struct proc *p;
00483
00484     acquire(&ptable.lock);
00485     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00486         if(p->pid == pid){
00487             p->killed = 1;
00488             // Wake process from sleep if necessary.
00489             if(p->state == SLEEPING)
00490                 p->state = RUNNABLE;
00491             release(&ptable.lock);
00492             return 0;
00493         }
00494     }
00495     release(&ptable.lock);
00496     return -1;
00497 }
```

5.238.1.17 link()

```
int link (
    const char * ,
    const char * )
```

Referenced by [bigdir\(\)](#), [concreate\(\)](#), [dirfile\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [subdir\(\)](#), and [validatetest\(\)](#).

5.238.1.18 malloc()

```
void * malloc (
    uint nbytes )
```

Definition at line 64 of file [umalloc.c](#).

```
00065 {
00066     Header *p, *prevp;
00067     uint nunits;
00068
00069     nunits = (nbytes + sizeof(Header) - 1)/sizeof(Header) + 1;
00070     if((prevp = freep) == 0){
00071         base.s.ptr = freep = prevp = &base;
00072         base.s.size = 0;
00073     }
00074     for(p = prevp->s.ptr; ; prevp = p, p = p->s.ptr){
00075         if(p->s.size >= nunits){
00076             if(p->s.size == nunits)
00077                 prevp->s.ptr = p->s.ptr;
00078             else {
00079                 p->s.size -= nunits;
00080                 p += p->s.size;
00081                 p->s.size = nunits;
00082             }
00083             freep = prevp;
00084             return (void*)(p + 1);
00085         }
00086         if(p == freep)
00087             if((p = morecore(nunits)) == 0)
00088                 return 0;
00089     }
00090 }
```

Referenced by [backcmd\(\)](#), [execcmd\(\)](#), [listcmd\(\)](#), [main\(\)](#), [mem\(\)](#), [pipecmd\(\)](#), and [redircmd\(\)](#).

5.238.1.19 memmove()

```
void * memmove (
    void * vdst,
    const void * vsrc,
    int n )
```

Definition at line 96 of file [ulib.c](#).

```
00097 {
00098     char *dst;
00099     const char *src;
00100
00101     dst = vdst;
00102     src = vsrc;
00103     while(n-- > 0)
00104         *dst++ = *src++;
00105     return vdst;
00106 }
```

5.238.1.20 `memset()`

```
void * memset (
    void * dst,
    int c,
    uint n )
```

Definition at line 5 of file [string.c](#).

```
00006 {
00007     if ((int)dst%4 == 0 && n%4 == 0){
00008         c &= 0xFF;
00009         stosl(dst, (c<<24)|(c<<16)|(c<<8)|c, n/4);
00010     } else
00011         stosb(dst, c, n);
00012     return dst;
00013 }
```

5.238.1.21 `mkdir()`

```
int mkdir (
    const char * )
```

Referenced by [dirfile\(\)](#), [dirtest\(\)](#), [exitiputtest\(\)](#), [fourteen\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [main\(\)](#), [openiputtest\(\)](#), [rmdot\(\)](#), and [subdir\(\)](#).

5.238.1.22 `mknod()`

```
int mknod (
    const char * ,
    short ,
    short )
```

Referenced by [main\(\)](#).

5.238.1.23 `open()`

```
int open (
    const char * ,
    int )
```

Referenced by [argptest\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirfile\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [fsfull\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [ls\(\)](#), [main\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [runcmd\(\)](#), [sharedfd\(\)](#), [stat\(\)](#), [subdir\(\)](#), [unlinkread\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.238.1.24 pipe()

```
int pipe (  
    int * )
```

5.238.1.25 printf()

```
void printf (  
    int fd,  
    const char * s,  
    ... )
```

Definition at line 11 of file [forktest.c](#).

```
00012 {  
00013     write(fd, s, strlen(s));  
00014 }
```

Referenced by [argptest\(\)](#), [balloc\(\)](#), [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [bsstest\(\)](#), [cat\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirfile\(\)](#), [dirtest\(\)](#), [exectest\(\)](#), [exitputtest\(\)](#), [exitwait\(\)](#), [forktest\(\)](#), [fourfiles\(\)](#), [fourteen\(\)](#), [fsfull\(\)](#), [getcmd\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [ls\(\)](#), [main\(\)](#), [mem\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [panic\(\)](#), [parsecmd\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [rmdot\(\)](#), [runcmd\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [uio\(\)](#), [unlinkread\(\)](#), [validatetest\(\)](#), [wc\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.238.1.26 read()

```
int read (  
    int ,  
    void * ,  
    int )
```

Referenced by [argptest\(\)](#), [bigfile\(\)](#), [cat\(\)](#), [concreate\(\)](#), [fourfiles\(\)](#), [gets\(\)](#), [grep\(\)](#), [linktest\(\)](#), [ls\(\)](#), [main\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [rsect\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [unlinkread\(\)](#), [wc\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.238.1.27 sbrk()

```
char * sbrk (  
    int )
```

Referenced by [argptest\(\)](#), [morecore\(\)](#), and [sbrktest\(\)](#).

5.238.1.28 sleep()

```
int sleep (  
    int )
```

5.238.1.29 stat()

```
int stat (
    const char * n,
    struct stat * st )
```

Definition at line 71 of file [ulib.c](#).

```
00072 {
00073     int fd;
00074     int r;
00075
00076     fd = open(n, O_RDONLY);
00077     if(fd < 0)
00078         return -1;
00079     r = fstat(fd, st);
00080     close(fd);
00081     return r;
00082 }
```

5.238.1.30 strchr()

```
char * strchr (
    const char * s,
    char c )
```

Definition at line 44 of file [ulib.c](#).

```
00045 {
00046     for(; *s; s++)
00047         if(*s == c)
00048             return (char*)s;
00049     return 0;
00050 }
```

Referenced by [gettoken\(\)](#), [grep\(\)](#), [peek\(\)](#), and [wc\(\)](#).

5.238.1.31 strcmp()

```
int strcmp (
    const char * p,
    const char * q )
```

Definition at line 19 of file [ulib.c](#).

```
00020 {
00021     while(*p && *p == *q)
00022         p++, q++;
00023     return (uchar)*p - (uchar)*q;
00024 }
```

5.238.1.32 strcpy()

```
char * strcpy (
    char * s,
    const char * t )
```

Definition at line 8 of file [ulib.c](#).

```
00009 {
00010     char *os;
00011
00012     os = s;
00013     while ((*s++ = *t++) != 0)
00014         ;
00015     return os;
00016 }
```

Referenced by [ls\(\)](#), and [main\(\)](#).

5.238.1.33 strlen()

```
uint strlen (
    const char * s )
```

Definition at line 97 of file [string.c](#).

```
00098 {
00099     int n;
00100
00101     for (n = 0; s[n]; n++)
00102         ;
00103     return n;
00104 }
```

5.238.1.34 unlink()

```
int unlink (
    const char * )
```

Referenced by [bigargtest\(\)](#), [bigdir\(\)](#), [bigfile\(\)](#), [bigwrite\(\)](#), [concreate\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirfile\(\)](#), [dirttest\(\)](#), [exitiputtest\(\)](#), [fourfiles\(\)](#), [fsfull\(\)](#), [iputtest\(\)](#), [iref\(\)](#), [linktest\(\)](#), [linkunlink\(\)](#), [main\(\)](#), [openiputtest\(\)](#), [rmdot\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [unlinkread\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.238.1.35 uptime()

```
int uptime (
    void )
```

5.238.1.36 wait()

```
int wait (
    void )
```

Definition at line 273 of file [proc.c](#).

```
00274 {
00275     struct proc *p;
00276     int havekids, pid;
00277     struct proc *curproc = myproc();
00278
00279     acquire(&ptable.lock);
00280     for(;;){
00281         // Scan through table looking for exited children.
00282         havekids = 0;
00283         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
00284             if(p->parent != curproc)
00285                 continue;
00286             havekids = 1;
00287             if(p->state == ZOMBIE){
00288                 // Found one.
00289                 pid = p->pid;
00290                 kfree(p->kstack);
00291                 p->kstack = 0;
00292                 freevm(p->pgdir);
00293                 p->pid = 0;
00294                 p->parent = 0;
00295                 p->name[0] = 0;
00296                 p->killed = 0;
00297                 p->state = UNUSED;
00298                 release(&ptable.lock);
00299                 return pid;
00300             }
00301         }
00302
00303         // No point waiting if we don't have any children.
00304         if(!havekids || curproc->killed){
00305             release(&ptable.lock);
00306             return -1;
00307         }
00308
00309         // Wait for children to exit.  (See wakeup1 call in proc_exit.)
00310         sleep(curproc, &ptable.lock); //DOC: wait-sleep
00311     }
00312 }
```

5.238.1.37 write()

```
int write (
    int ,
    const void * ,
    int )
```

Referenced by [bigfile\(\)](#), [bigwrite\(\)](#), [cat\(\)](#), [dirfile\(\)](#), [fourfiles\(\)](#), [fsfull\(\)](#), [grep\(\)](#), [linktest\(\)](#), [main\(\)](#), [pipe1\(\)](#), [preempt\(\)](#), [printf\(\)](#), [putc\(\)](#), [sbrktest\(\)](#), [sharedfd\(\)](#), [subdir\(\)](#), [unlinkread\(\)](#), [writetest\(\)](#), [writetest1\(\)](#), and [wsect\(\)](#).

5.239 user.h

[Go to the documentation of this file.](#)

```
00001 struct stat;
00002 struct rtcdate;
00003 struct uproc;
00004
00005 // system calls
00006 int fork(void);
00007 int exit(void) __attribute__((noreturn));
00008 int wait(void);
```



```

00009 int pipe(int*);
00010 int write(int, const void*, int);
00011 int read(int, void*, int);
00012 int close(int);
00013 int kill(int);
00014 int exec(char*, char**);
00015 int open(const char*, int);
00016 int mknod(const char*, short, short);
00017 int unlink(const char*);
00018 int fstat(int fd, struct stat*);
00019 int link(const char*, const char*);
00020 int mkdir(const char*);
00021 int chdir(const char*);
00022 int dup(int);
00023 int getpid(void);
00024 char* sbrk(int);
00025 int sleep(int);
00026 int uptime(void);
00027 int halt(void);
00028 int cps(void);
00029 int chpr(int pid, int priority);
00030 int getprocs(int max, struct uproc*);
00031
00032 // ulib.c
00033 int stat(const char*, struct stat*);
00034 char* strcpy(char*, const char*);
00035 void *memmove(void*, const void*, int);
00036 char* strchr(const char*, char c);
00037 int strcmp(const char*, const char*);
00038 void printf(int, const char*, ...);
00039 char* gets(char*, int max);
00040 uint strlen(const char*);
00041 void* memset(void*, int, uint);
00042 void* malloc(uint);
00043 void free(void*);
00044 int atoi(const char*);

```

5.240 usertests.c File Reference

```

#include "param.h"
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fs.h"
#include "fcntl.h"
#include "syscall.h"
#include "traps.h"
#include "memlayout.h"

```

Macros

- #define [BIG](#) (100*1024*1024)
- #define [RTC_ADDR](#) 0x70
- #define [RTC_DATA](#) 0x71

Functions

- void [argptest](#) ()
- void [bigargtest](#) (void)
- void [bigdir](#) (void)
- void [bigfile](#) (void)
- void [bigwrite](#) (void)
- void [bsstest](#) (void)

- void [concreate](#) (void)
- void [createdelete](#) (void)
- void [createtest](#) (void)
- void [dirfile](#) (void)
- void [dirtest](#) (void)
- void [exectest](#) (void)
- void [exitputtest](#) (void)
- void [exitwait](#) (void)
- void [forktest](#) (void)
- void [fourfiles](#) (void)
- void [fourteen](#) (void)
- void [fsfull](#) ()
- void [iputtest](#) (void)
- void [iref](#) (void)
- void [linktest](#) (void)
- void [linkunlink](#) ()
- int [main](#) (int argc, char *[argv](#)[])
- void [mem](#) (void)
- void [openiputtest](#) (void)
- void [opentest](#) (void)
- void [pipe1](#) (void)
- void [preempt](#) (void)
- unsigned int [rand](#) ()
- void [rmdot](#) (void)
- void [sbrktest](#) (void)
- void [sharedfd](#) (void)
- void [subdir](#) (void)
- void [uio](#) ()
- void [unlinkread](#) (void)
- void [validateint](#) (int *p)
- void [validatetest](#) (void)
- void [writetest](#) (void)
- void [writetest1](#) (void)

Variables

- char [buf](#) [8192]
- char * [echoargv](#) [] = { "echo", "ALL", "TESTS", "PASSED", 0 }
- char [name](#) [3]
- unsigned long [randstate](#) = 1
- int [stdout](#) = 1
- char [uninit](#) [10000]

5.240.1 Macro Definition Documentation

5.240.1.1 BIG

```
#define BIG (100*1024*1024)
```

5.240.1.2 RTC_ADDR

```
#define RTC_ADDR 0x70
```

5.240.1.3 RTC_DATA

```
#define RTC_DATA 0x71
```

5.240.2 Function Documentation

5.240.2.1 argptest()

```
void argptest ( )
```

Definition at line 1727 of file [usertests.c](#).

```
01728 {
01729     int fd;
01730     fd = open("init", O_RDONLY);
01731     if (fd < 0) {
01732         printf(2, "open failed\n");
01733         exit();
01734     }
01735     read(fd, sbrk(0) - 1, -1);
01736     close(fd);
01737     printf(1, "arg test passed\n");
01738 }
```

Referenced by [main\(\)](#).

5.240.2.2 bigargtest()

```
void bigargtest (
    void )
```

Definition at line 1613 of file [usertests.c](#).

```
01614 {
01615     int pid, fd;
01616
01617     unlink("bigarg-ok");
01618     pid = fork();
01619     if(pid == 0){
01620         static char *args[MAXARG];
01621         int i;
01622         for(i = 0; i < MAXARG-1; i++)
01623             args[i] = "bigargs test: failed\n";
01624
01625         args[MAXARG-1] = 0;
01626         printf(stdout, "bigarg test\n");
01627         exec("echo", args);
01628         printf(stdout, "bigarg test ok\n");
01629         fd = open("bigarg-ok", O_CREATE);
01630         close(fd);
01631         exit();
01632     } else if(pid < 0){
```

```

01632     printf(stdout, "bigargtest: fork failed\n");
01633     exit();
01634 }
01635 wait();
01636 fd = open("bigarg-ok", 0);
01637 if(fd < 0){
01638     printf(stdout, "bigarg test failed!\n");
01639     exit();
01640 }
01641 close(fd);
01642 unlink("bigarg-ok");
01643 }

```

Referenced by [main\(\)](#).

5.240.2.3 bigdir()

```

void bigdir (
    void )

```

Definition at line 893 of file [usertests.c](#).

```

00894 {
00895     int i, fd;
00896     char name[10];
00897
00898     printf(1, "bigdir test\n");
00899     unlink("bd");
00900
00901     fd = open("bd", O_CREATE);
00902     if(fd < 0){
00903         printf(1, "bigdir create failed\n");
00904         exit();
00905     }
00906     close(fd);
00907
00908     for(i = 0; i < 500; i++){
00909         name[0] = 'x';
00910         name[1] = '0' + (i / 64);
00911         name[2] = '0' + (i % 64);
00912         name[3] = '\\0';
00913         if(link("bd", name) != 0){
00914             printf(1, "bigdir link failed\n");
00915             exit();
00916         }
00917     }
00918
00919     unlink("bd");
00920     for(i = 0; i < 500; i++){
00921         name[0] = 'x';
00922         name[1] = '0' + (i / 64);
00923         name[2] = '0' + (i % 64);
00924         name[3] = '\\0';
00925         if(unlink(name) != 0){
00926             printf(1, "bigdir unlink failed");
00927             exit();
00928         }
00929     }
00930
00931     printf(1, "bigdir ok\n");
00932 }

```

Referenced by [main\(\)](#).

5.240.2.4 bigfile()

```
void bigfile (
    void )
```

Definition at line 1148 of file [usertests.c](#).

```
01149 {
01150     int fd, i, total, cc;
01151
01152     printf(1, "bigfile test\n");
01153
01154     unlink("bigfile");
01155     fd = open("bigfile", O_CREATE | O_RDWR);
01156     if(fd < 0){
01157         printf(1, "cannot create bigfile");
01158         exit();
01159     }
01160     for(i = 0; i < 20; i++){
01161         memset(buf, i, 600);
01162         if(write(fd, buf, 600) != 600){
01163             printf(1, "write bigfile failed\n");
01164             exit();
01165         }
01166     }
01167     close(fd);
01168
01169     fd = open("bigfile", 0);
01170     if(fd < 0){
01171         printf(1, "cannot open bigfile\n");
01172         exit();
01173     }
01174     total = 0;
01175     for(i = 0; ; i++){
01176         cc = read(fd, buf, 300);
01177         if(cc < 0){
01178             printf(1, "read bigfile failed\n");
01179             exit();
01180         }
01181         if(cc == 0)
01182             break;
01183         if(cc != 300){
01184             printf(1, "short read bigfile\n");
01185             exit();
01186         }
01187         if(buf[0] != i/2 || buf[299] != i/2){
01188             printf(1, "read bigfile wrong data\n");
01189             exit();
01190         }
01191         total += cc;
01192     }
01193     close(fd);
01194     if(total != 20*600){
01195         printf(1, "read bigfile wrong total\n");
01196         exit();
01197     }
01198     unlink("bigfile");
01199
01200     printf(1, "bigfile test ok\n");
01201 }
```

Referenced by [main\(\)](#).

5.240.2.5 bigwrite()

```
void bigwrite (
    void )
```

Definition at line 1119 of file [usertests.c](#).

```
01120 {
01121     int fd, sz;
01122
01123     printf(1, "bigwrite test\n");
01124
01125     unlink("bigwrite");
```

```

01126     for(sz = 499; sz < 12*512; sz += 471){
01127         fd = open("bigwrite", O_CREATE | O_RDWR);
01128         if(fd < 0){
01129             printf(1, "cannot create bigwrite\n");
01130             exit();
01131         }
01132         int i;
01133         for(i = 0; i < 2; i++){
01134             int cc = write(fd, buf, sz);
01135             if(cc != sz){
01136                 printf(1, "write(%d) ret %d\n", sz, cc);
01137                 exit();
01138             }
01139         }
01140         close(fd);
01141         unlink("bigwrite");
01142     }
01143     printf(1, "bigwrite ok\n");
01144 }
01145 }

```

Referenced by [main\(\)](#).

5.240.2.6 bsstest()

```

void bsstest (
    void )

```

Definition at line 1595 of file [usertests.c](#).

```

01596 {
01597     int i;
01598
01599     printf(stdout, "bss test\n");
01600     for(i = 0; i < sizeof(uninit); i++){
01601         if(uninit[i] != '\0'){
01602             printf(stdout, "bss test failed\n");
01603             exit();
01604         }
01605     }
01606     printf(stdout, "bss test ok\n");
01607 }

```

Referenced by [main\(\)](#).

5.240.2.7 concreate()

```

void concreate (
    void )

```

Definition at line 765 of file [usertests.c](#).

```

00766 {
00767     char file[3];
00768     int i, pid, n, fd;
00769     char fa[40];
00770     struct {
00771         ushort inum;
00772         char name[14];
00773     } de;
00774
00775     printf(1, "concreate test\n");
00776     file[0] = 'C';
00777     file[2] = '\0';
00778     for(i = 0; i < 40; i++){
00779         file[1] = '0' + i;
00780         unlink(file);
00781         pid = fork();
00782         if(pid && (i % 3) == 1){

```

```

00783     link("C0", file);
00784 } else if(pid == 0 && (i % 5) == 1){
00785     link("C0", file);
00786 } else {
00787     fd = open(file, O_CREATE | O_RDWR);
00788     if(fd < 0){
00789         printf(1, "concreate create %s failed\n", file);
00790         exit();
00791     }
00792     close(fd);
00793 }
00794 if(pid == 0)
00795     exit();
00796 else
00797     wait();
00798 }
00799
00800 memset(fa, 0, sizeof(fa));
00801 fd = open(".", 0);
00802 n = 0;
00803 while(read(fd, &de, sizeof(de)) > 0){
00804     if(de.inum == 0)
00805         continue;
00806     if(de.name[0] == 'C' && de.name[2] == '\0'){
00807         i = de.name[1] - '0';
00808         if(i < 0 || i >= sizeof(fa)){
00809             printf(1, "concreate weird file %s\n", de.name);
00810             exit();
00811         }
00812         if(fa[i]){
00813             printf(1, "concreate duplicate file %s\n", de.name);
00814             exit();
00815         }
00816         fa[i] = 1;
00817         n++;
00818     }
00819 }
00820 close(fd);
00821
00822 if(n != 40){
00823     printf(1, "concreate not enough files in directory listing\n");
00824     exit();
00825 }
00826
00827 for(i = 0; i < 40; i++){
00828     file[1] = '0' + i;
00829     pid = fork();
00830     if(pid < 0){
00831         printf(1, "fork failed\n");
00832         exit();
00833     }
00834     if(((i % 3) == 0 && pid == 0) ||
00835        ((i % 3) == 1 && pid != 0)){
00836         close(open(file, 0));
00837         close(open(file, 0));
00838         close(open(file, 0));
00839         close(open(file, 0));
00840     } else {
00841         unlink(file);
00842         unlink(file);
00843         unlink(file);
00844         unlink(file);
00845     }
00846     if(pid == 0)
00847         exit();
00848     else
00849         wait();
00850 }
00851
00852 printf(1, "concreate ok\n");
00853 }

```

Referenced by [main\(\)](#).

5.240.2.8 createdelete()

```

void createdelete (
    void )

```

Definition at line 584 of file `usertests.c`.

```

00585 {
00586     enum { N = 20 };
00587     int pid, i, fd, pi;
00588     char name[32];
00589
00590     printf(1, "createdelete test\n");
00591
00592     for(pi = 0; pi < 4; pi++){
00593         pid = fork();
00594         if(pid < 0){
00595             printf(1, "fork failed\n");
00596             exit();
00597         }
00598
00599         if(pid == 0){
00600             name[0] = 'p' + pi;
00601             name[2] = '\0';
00602             for(i = 0; i < N; i++){
00603                 name[1] = '0' + i;
00604                 fd = open(name, O_CREATE | O_RDWR);
00605                 if(fd < 0){
00606                     printf(1, "create failed\n");
00607                     exit();
00608                 }
00609                 close(fd);
00610                 if(i > 0 && (i % 2) == 0){
00611                     name[1] = '0' + (i / 2);
00612                     if(unlink(name) < 0){
00613                         printf(1, "unlink failed\n");
00614                         exit();
00615                     }
00616                 }
00617             }
00618             exit();
00619         }
00620     }
00621
00622     for(pi = 0; pi < 4; pi++){
00623         wait();
00624     }
00625
00626     name[0] = name[1] = name[2] = 0;
00627     for(i = 0; i < N; i++){
00628         for(pi = 0; pi < 4; pi++){
00629             name[0] = 'p' + pi;
00630             name[1] = '0' + i;
00631             fd = open(name, 0);
00632             if((i == 0 || i >= N/2) && fd < 0){
00633                 printf(1, "oops createdelete %s didn't exist\n", name);
00634                 exit();
00635             } else if((i >= 1 && i < N/2) && fd >= 0){
00636                 printf(1, "oops createdelete %s did exist\n", name);
00637                 exit();
00638             }
00639             if(fd >= 0)
00640                 close(fd);
00641         }
00642     }
00643
00644     for(i = 0; i < N; i++){
00645         for(pi = 0; pi < 4; pi++){
00646             name[0] = 'p' + i;
00647             name[1] = '0' + i;
00648             unlink(name);
00649         }
00650     }
00651
00652     printf(1, "createdelete ok\n");
00653 }

```

Referenced by `main()`.

5.240.2.9 createtest()

```

void createtest (
    void )

```


Definition at line 245 of file [usertests.c](#).

```

00246 {
00247     int i, fd;
00248
00249     printf(stdout, "many creates, followed by unlink test\n");
00250
00251     name[0] = 'a';
00252     name[2] = '\0';
00253     for(i = 0; i < 52; i++){
00254         name[1] = '0' + i;
00255         fd = open(name, O_CREATE|O_RDWR);
00256         close(fd);
00257     }
00258     name[0] = 'a';
00259     name[2] = '\0';
00260     for(i = 0; i < 52; i++){
00261         name[1] = '0' + i;
00262         unlink(name);
00263     }
00264     printf(stdout, "many creates, followed by unlink; ok\n");
00265 }
```

Referenced by [main\(\)](#).

5.240.2.10 dirfile()

```

void dirfile (
    void )
```

Definition at line 1284 of file [usertests.c](#).

```

01285 {
01286     int fd;
01287
01288     printf(1, "dir vs file\n");
01289
01290     fd = open("dirfile", O_CREATE);
01291     if(fd < 0){
01292         printf(1, "create dirfile failed\n");
01293         exit();
01294     }
01295     close(fd);
01296     if(chdir("dirfile") == 0){
01297         printf(1, "chdir dirfile succeeded!\n");
01298         exit();
01299     }
01300     fd = open("dirfile/xx", 0);
01301     if(fd >= 0){
01302         printf(1, "create dirfile/xx succeeded!\n");
01303         exit();
01304     }
01305     fd = open("dirfile/xx", O_CREATE);
01306     if(fd >= 0){
01307         printf(1, "create dirfile/xx succeeded!\n");
01308         exit();
01309     }
01310     if(mkdir("dirfile/xx") == 0){
01311         printf(1, "mkdir dirfile/xx succeeded!\n");
01312         exit();
01313     }
01314     if(unlink("dirfile/xx") == 0){
01315         printf(1, "unlink dirfile/xx succeeded!\n");
01316         exit();
01317     }
01318     if(link("README", "dirfile/xx") == 0){
01319         printf(1, "link to dirfile/xx succeeded!\n");
01320         exit();
01321     }
01322     if(unlink("dirfile") != 0){
01323         printf(1, "unlink dirfile failed!\n");
01324         exit();
01325     }
01326
01327     fd = open(".", O_RDWR);
01328     if(fd >= 0){
01329         printf(1, "open . for writing succeeded!\n");
01330         exit();
01331     }
```

```
01332 fd = open(".", 0);
01333 if(write(fd, "x", 1) > 0){
01334     printf(1, "write . succeeded!\n");
01335     exit();
01336 }
01337 close(fd);
01338
01339 printf(1, "dir vs file OK\n");
01340 }
```

Referenced by [main\(\)](#).

5.240.2.11 dirtest()

```
void dirtest (
    void )
```

Definition at line 267 of file [usertests.c](#).

```
00268 {
00269     printf(stdout, "mkdir test\n");
00270
00271     if(mkdir("dir0") < 0){
00272         printf(stdout, "mkdir failed\n");
00273         exit();
00274     }
00275
00276     if(chdir("dir0") < 0){
00277         printf(stdout, "chdir dir0 failed\n");
00278         exit();
00279     }
00280
00281     if(chdir("../") < 0){
00282         printf(stdout, "chdir .. failed\n");
00283         exit();
00284     }
00285
00286     if(unlink("dir0") < 0){
00287         printf(stdout, "unlink dir0 failed\n");
00288         exit();
00289     }
00290     printf(stdout, "mkdir test ok\n");
00291 }
```

5.240.2.12 exectest()

```
void exectest (
    void )
```

Definition at line 294 of file [usertests.c](#).

```
00295 {
00296     printf(stdout, "exec test\n");
00297     if(exec("echo", echoargv) < 0){
00298         printf(stdout, "exec echo failed\n");
00299         exit();
00300     }
00301 }
```

Referenced by [main\(\)](#).

5.240.2.13 exitiputtest()

```
void exitiputtest (  
    void )
```

Definition at line 43 of file [usertests.c](#).

```
00044 {  
00045     int pid;  
00046  
00047     printf(stdout, "exitiput test\n");  
00048  
00049     pid = fork();  
00050     if(pid < 0){  
00051         printf(stdout, "fork failed\n");  
00052         exit();  
00053     }  
00054     if(pid == 0){  
00055         if(mkdir("iputdir") < 0){  
00056             printf(stdout, "mkdir failed\n");  
00057             exit();  
00058         }  
00059         if(chdir("iputdir") < 0){  
00060             printf(stdout, "child chdir failed\n");  
00061             exit();  
00062         }  
00063         if(unlink("../iputdir") < 0){  
00064             printf(stdout, "unlink ../iputdir failed\n");  
00065             exit();  
00066         }  
00067         exit();  
00068     }  
00069     wait();  
00070     printf(stdout, "exitiput test ok\n");  
00071 }
```

Referenced by [main\(\)](#).

5.240.2.14 exitwait()

```
void exitwait (  
    void )
```

Definition at line 405 of file [usertests.c](#).

```
00406 {  
00407     int i, pid;  
00408  
00409     for(i = 0; i < 100; i++){  
00410         pid = fork();  
00411         if(pid < 0){  
00412             printf(1, "fork failed\n");  
00413             return;  
00414         }  
00415         if(pid){  
00416             if(wait() != pid){  
00417                 printf(1, "wait wrong pid\n");  
00418                 return;  
00419             }  
00420         } else {  
00421             exit();  
00422         }  
00423     }  
00424     printf(1, "exitwait ok\n");  
00425 }
```

Referenced by [main\(\)](#).

5.240.2.15 forktest()

```
void forktest (
    void )
```

Definition at line 1380 of file [usertests.c](#).

```
01381 {
01382     int n, pid;
01383
01384     printf(1, "fork test\n");
01385
01386     for(n=0; n<1000; n++){
01387         pid = fork();
01388         if(pid < 0)
01389             break;
01390         if(pid == 0)
01391             exit();
01392     }
01393
01394     if(n == 1000){
01395         printf(1, "fork claimed to work 1000 times!\n");
01396         exit();
01397     }
01398
01399     for(; n > 0; n--){
01400         if(wait() < 0){
01401             printf(1, "wait stopped early\n");
01402             exit();
01403         }
01404     }
01405
01406     if(wait() != -1){
01407         printf(1, "wait got too many\n");
01408         exit();
01409     }
01410
01411     printf(1, "fork test OK\n");
01412 }
```

Referenced by [main\(\)](#).

5.240.2.16 fourfiles()

```
void fourfiles (
    void )
```

Definition at line 518 of file [usertests.c](#).

```
00519 {
00520     int fd, pid, i, j, n, total, pi;
00521     char *names[] = { "f0", "f1", "f2", "f3" };
00522     char *fname;
00523
00524     printf(1, "fourfiles test\n");
00525
00526     for(pi = 0; pi < 4; pi++){
00527         fname = names[pi];
00528         unlink(fname);
00529
00530         pid = fork();
00531         if(pid < 0){
00532             printf(1, "fork failed\n");
00533             exit();
00534         }
00535
00536         if(pid == 0){
00537             fd = open(fname, O_CREATE | O_RDWR);
00538             if(fd < 0){
00539                 printf(1, "create failed\n");
00540                 exit();
00541             }
00542
00543             memset(buf, '0'+pi, 512);
00544             for(i = 0; i < 12; i++){
00545                 if((n = write(fd, buf, 500)) != 500){
```

```

00546         printf(1, "write failed %d\n", n);
00547         exit();
00548     }
00549 }
00550     exit();
00551 }
00552 }
00553
00554 for(pi = 0; pi < 4; pi++){
00555     wait();
00556 }
00557
00558 for(i = 0; i < 2; i++){
00559     fname = names[i];
00560     fd = open(fname, 0);
00561     total = 0;
00562     while((n = read(fd, buf, sizeof(buf))) > 0){
00563         for(j = 0; j < n; j++){
00564             if(buf[j] != '0'+i){
00565                 printf(1, "wrong char\n");
00566                 exit();
00567             }
00568         }
00569         total += n;
00570     }
00571     close(fd);
00572     if(total != 12*500){
00573         printf(1, "wrong length %d\n", total);
00574         exit();
00575     }
00576     unlink(fname);
00577 }
00578
00579 printf(1, "fourfiles ok\n");
00580 }

```

Referenced by [main\(\)](#).

5.240.2.17 fourteen()

```

void fourteen (
    void )

```

Definition at line 1204 of file [usertests.c](#).

```

01205 {
01206     int fd;
01207
01208     // DIRSIZ is 14.
01209     printf(1, "fourteen test\n");
01210
01211     if(mkdir("12345678901234") != 0){
01212         printf(1, "mkdir 12345678901234 failed\n");
01213         exit();
01214     }
01215     if(mkdir("12345678901234/123456789012345") != 0){
01216         printf(1, "mkdir 12345678901234/123456789012345 failed\n");
01217         exit();
01218     }
01219     fd = open("123456789012345/123456789012345/123456789012345", O_CREATE);
01220     if(fd < 0){
01221         printf(1, "create 123456789012345/123456789012345/123456789012345 failed\n");
01222         exit();
01223     }
01224     close(fd);
01225     fd = open("12345678901234/12345678901234/12345678901234", 0);
01226     if(fd < 0){
01227         printf(1, "open 12345678901234/12345678901234/12345678901234 failed\n");
01228         exit();
01229     }
01230     close(fd);
01231
01232     if(mkdir("12345678901234/12345678901234") == 0){
01233         printf(1, "mkdir 12345678901234/12345678901234 succeeded!\n");
01234         exit();
01235     }
01236     if(mkdir("123456789012345/12345678901234") == 0){
01237         printf(1, "mkdir 12345678901234/123456789012345 succeeded!\n");

```

```

01238     exit();
01239 }
01240
01241 printf(1, "fourteen ok\n");
01242 }

```

Referenced by [main\(\)](#).

5.240.2.18 fsfull()

```
void fsfull ( )
```

Definition at line 1648 of file [usertests.c](#).

```

01649 {
01650     int nfiles;
01651     int fsblocks = 0;
01652
01653     printf(1, "fsfull test\n");
01654
01655     for(nfiles = 0; ; nfiles++){
01656         char name[64];
01657         name[0] = 'f';
01658         name[1] = '0' + nfiles / 1000;
01659         name[2] = '0' + (nfiles % 1000) / 100;
01660         name[3] = '0' + (nfiles % 100) / 10;
01661         name[4] = '0' + (nfiles % 10);
01662         name[5] = '\0';
01663         printf(1, "writing %s\n", name);
01664         int fd = open(name, O_CREATE|O_RDWR);
01665         if(fd < 0){
01666             printf(1, "open %s failed\n", name);
01667             break;
01668         }
01669         int total = 0;
01670         while(1){
01671             int cc = write(fd, buf, 512);
01672             if(cc < 512)
01673                 break;
01674             total += cc;
01675             fsblocks++;
01676         }
01677         printf(1, "wrote %d bytes\n", total);
01678         close(fd);
01679         if(total == 0)
01680             break;
01681     }
01682
01683     while(nfiles >= 0){
01684         char name[64];
01685         name[0] = 'f';
01686         name[1] = '0' + nfiles / 1000;
01687         name[2] = '0' + (nfiles % 1000) / 100;
01688         name[3] = '0' + (nfiles % 100) / 10;
01689         name[4] = '0' + (nfiles % 10);
01690         name[5] = '\0';
01691         unlink(name);
01692         nfiles--;
01693     }
01694
01695     printf(1, "fsfull test finished\n");
01696 }

```

5.240.2.19 iputtest()

```
void iputtest (
    void )
```

Definition at line 18 of file [usertests.c](#).

```
00019 {
```

```

00020     printf(stdout, "iput test\n");
00021
00022     if(mkdir("iputdir") < 0){
00023         printf(stdout, "mkdir failed\n");
00024         exit();
00025     }
00026     if(chdir("iputdir") < 0){
00027         printf(stdout, "chdir iputdir failed\n");
00028         exit();
00029     }
00030     if(unlink("../iputdir") < 0){
00031         printf(stdout, "unlink ../iputdir failed\n");
00032         exit();
00033     }
00034     if(chdir("/") < 0){
00035         printf(stdout, "chdir / failed\n");
00036         exit();
00037     }
00038     printf(stdout, "iput test ok\n");
00039 }

```

Referenced by [main\(\)](#).

5.240.2.20 iref()

```

void iref (
    void )

```

Definition at line 1344 of file [usertests.c](#).

```

01345 {
01346     int i, fd;
01347
01348     printf(1, "empty file name\n");
01349
01350     // the 50 is NINODE
01351     for(i = 0; i < 50 + 1; i++){
01352         if(mkdir("irefd") != 0){
01353             printf(1, "mkdir irefd failed\n");
01354             exit();
01355         }
01356         if(chdir("irefd") != 0){
01357             printf(1, "chdir irefd failed\n");
01358             exit();
01359         }
01360
01361         mkdir("");
01362         link("README", "");
01363         fd = open("", O_CREATE);
01364         if(fd >= 0)
01365             close(fd);
01366         fd = open("xx", O_CREATE);
01367         if(fd >= 0)
01368             close(fd);
01369         unlink("xx");
01370     }
01371
01372     chdir("/");
01373     printf(1, "empty file name OK\n");
01374 }

```

Referenced by [main\(\)](#).

5.240.2.21 linktest()

```
void linktest (
    void )
```

Definition at line 702 of file [usertests.c](#).

```
00703 {
00704     int fd;
00705
00706     printf(1, "linktest\n");
00707
00708     unlink("lf1");
00709     unlink("lf2");
00710
00711     fd = open("lf1", O_CREATE|O_RDWR);
00712     if(fd < 0){
00713         printf(1, "create lf1 failed\n");
00714         exit();
00715     }
00716     if(write(fd, "hello", 5) != 5){
00717         printf(1, "write lf1 failed\n");
00718         exit();
00719     }
00720     close(fd);
00721
00722     if(link("lf1", "lf2") < 0){
00723         printf(1, "link lf1 lf2 failed\n");
00724         exit();
00725     }
00726     unlink("lf1");
00727
00728     if(open("lf1", 0) >= 0){
00729         printf(1, "unlinked lf1 but it is still there!\n");
00730         exit();
00731     }
00732
00733     fd = open("lf2", 0);
00734     if(fd < 0){
00735         printf(1, "open lf2 failed\n");
00736         exit();
00737     }
00738     if(read(fd, buf, sizeof(buf)) != 5){
00739         printf(1, "read lf2 failed\n");
00740         exit();
00741     }
00742     close(fd);
00743
00744     if(link("lf2", "lf2") >= 0){
00745         printf(1, "link lf2 lf2 succeeded! oops\n");
00746         exit();
00747     }
00748
00749     unlink("lf2");
00750     if(link("lf2", "lf1") >= 0){
00751         printf(1, "link non-existant succeeded! oops\n");
00752         exit();
00753     }
00754
00755     if(link(".", "lf1") >= 0){
00756         printf(1, "link . lf1 succeeded! oops\n");
00757         exit();
00758     }
00759
00760     printf(1, "linktest ok\n");
00761 }
```

Referenced by [main\(\)](#).

5.240.2.22 linkunlink()

```
void linkunlink ( )
```

Definition at line 858 of file [usertests.c](#).

```
00859 {
```



```

00860  int pid, i;
00861
00862  printf(1, "linkunlink test\n");
00863
00864  unlink("x");
00865  pid = fork();
00866  if(pid < 0){
00867      printf(1, "fork failed\n");
00868      exit();
00869  }
00870
00871  unsigned int x = (pid ? 1 : 97);
00872  for(i = 0; i < 100; i++){
00873      x = x * 1103515245 + 12345;
00874      if((x % 3) == 0){
00875          close(open("x", O_RDWR | O_CREATE));
00876      } else if((x % 3) == 1){
00877          link("cat", "x");
00878      } else {
00879          unlink("x");
00880      }
00881  }
00882
00883  if(pid)
00884      wait();
00885  else
00886      exit();
00887
00888  printf(1, "linkunlink ok\n");
00889 }

```

Referenced by [main\(\)](#).

5.240.2.23 main()

```

int main (
    int argc,
    char * argv[] )

```

Definition at line 1749 of file [usertests.c](#).

```

01750 {
01751     printf(1, "usertests starting\n");
01752
01753     if(open("usertests.ran", 0) >= 0){
01754         printf(1, "already ran user tests -- rebuild fs.img\n");
01755         exit();
01756     }
01757     close(open("usertests.ran", O_CREATE));
01758
01759     argptest();
01760     createdelete();
01761     linkunlink();
01762     concreate();
01763     fourfiles();
01764     sharedfd();
01765
01766     bigargtest();
01767     bigwrite();
01768     bigargtest();
01769     bsstest();
01770     sbrktest();
01771     validatetest();
01772
01773     opentest();
01774     writetest();
01775     writetest1();
01776     createtest();
01777
01778     openiputtest();
01779     exitiputtest();
01780     iputtest();
01781
01782     mem();
01783     pipel();
01784     preempt();
01785     exitwait();
01786 }

```

```

01787     rmdot();
01788     fourteen();
01789     bigfile();
01790     subdir();
01791     linktest();
01792     unlinkread();
01793     dirfile();
01794     iref();
01795     forktest();
01796     bigdir(); // slow
01797
01798     uio();
01799
01800     exectest();
01801
01802     exit();
01803 }

```

5.240.2.24 mem()

```

void mem (
    void )

```

Definition at line 428 of file [usertests.c](#).

```

00429 {
00430     void *m1, *m2;
00431     int pid, ppid;
00432
00433     printf(1, "mem test\n");
00434     ppid = getpid();
00435     if((pid = fork()) == 0){
00436         m1 = 0;
00437         while((m2 = malloc(10001)) != 0){
00438             *(char**)m2 = m1;
00439             m1 = m2;
00440         }
00441         while(m1){
00442             m2 = *(char**)m1;
00443             free(m1);
00444             m1 = m2;
00445         }
00446         m1 = malloc(1024*20);
00447         if(m1 == 0){
00448             printf(1, "couldn't allocate mem?!\n");
00449             kill(ppid);
00450             exit();
00451         }
00452         free(m1);
00453         printf(1, "mem ok\n");
00454         exit();
00455     } else {
00456         wait();
00457     }
00458 }

```

Referenced by [allocuvm\(\)](#), [copyuvm\(\)](#), [inituvm\(\)](#), and [main\(\)](#).

5.240.2.25 openiputtest()

```

void openiputtest (
    void )

```

Definition at line 85 of file [usertests.c](#).

```

00086 {
00087     int pid;
00088
00089     printf(stdout, "openiput test\n");
00090     if(mkdir("oidir") < 0){

```

```

00091     printf(stdout, "mkdir oidir failed\n");
00092     exit();
00093 }
00094 pid = fork();
00095 if(pid < 0){
00096     printf(stdout, "fork failed\n");
00097     exit();
00098 }
00099 if(pid == 0){
00100     int fd = open("oidir", O_RDWR);
00101     if(fd >= 0){
00102         printf(stdout, "open directory for write succeeded\n");
00103         exit();
00104     }
00105     exit();
00106 }
00107 sleep(1);
00108 if(unlink("oidir") != 0){
00109     printf(stdout, "unlink failed\n");
00110     exit();
00111 }
00112 wait();
00113 printf(stdout, "openiput test ok\n");
00114 }

```

Referenced by [main\(\)](#).

5.240.2.26 opentest()

```

void opentest (
    void )

```

Definition at line 119 of file [usertests.c](#).

```

00120 {
00121     int fd;
00122
00123     printf(stdout, "open test\n");
00124     fd = open("echo", 0);
00125     if(fd < 0){
00126         printf(stdout, "open echo failed!\n");
00127         exit();
00128     }
00129     close(fd);
00130     fd = open("doesnotexist", 0);
00131     if(fd >= 0){
00132         printf(stdout, "open doesnotexist succeeded!\n");
00133         exit();
00134     }
00135     printf(stdout, "open test ok\n");
00136 }

```

Referenced by [main\(\)](#).

5.240.2.27 pipe1()

```

void pipe1 (
    void )

```

Definition at line 306 of file [usertests.c](#).

```

00307 {
00308     int fds[2], pid;
00309     int seq, i, n, cc, total;
00310
00311     if(pipe(fds) != 0){
00312         printf(1, "pipe() failed\n");
00313         exit();
00314     }

```

```

00315 pid = fork();
00316 seq = 0;
00317 if(pid == 0){
00318     close(fds[0]);
00319     for(n = 0; n < 5; n++){
00320         for(i = 0; i < 1033; i++)
00321             buf[i] = seq++;
00322         if(write(fds[1], buf, 1033) != 1033){
00323             printf(1, "pipel oops 1\n");
00324             exit();
00325         }
00326     }
00327     exit();
00328 } else if(pid > 0){
00329     close(fds[1]);
00330     total = 0;
00331     cc = 1;
00332     while((n = read(fds[0], buf, cc)) > 0){
00333         for(i = 0; i < n; i++){
00334             if((buf[i] & 0xff) != (seq++ & 0xff)){
00335                 printf(1, "pipel oops 2\n");
00336                 return;
00337             }
00338         }
00339         total += n;
00340         cc = cc * 2;
00341         if(cc > sizeof(buf))
00342             cc = sizeof(buf);
00343     }
00344     if(total != 5 * 1033){
00345         printf(1, "pipel oops 3 total %d\n", total);
00346         exit();
00347     }
00348     close(fds[0]);
00349     wait();
00350 } else {
00351     printf(1, "fork() failed\n");
00352     exit();
00353 }
00354 printf(1, "pipel ok\n");
00355 }

```

Referenced by [main\(\)](#).

5.240.2.28 preempt()

```

void preempt (
    void )

```

Definition at line 359 of file [usertests.c](#).

```

00360 {
00361     int pid1, pid2, pid3;
00362     int pfds[2];
00363
00364     printf(1, "preempt: ");
00365     pid1 = fork();
00366     if(pid1 == 0)
00367         for(;;)
00368             ;
00369
00370     pid2 = fork();
00371     if(pid2 == 0)
00372         for(;;)
00373             ;
00374
00375     pipe(pfds);
00376     pid3 = fork();
00377     if(pid3 == 0){
00378         close(pfds[0]);
00379         if(write(pfds[1], "x", 1) != 1)
00380             printf(1, "preempt write error");
00381         close(pfds[1]);
00382         for(;;)
00383             ;
00384     }
00385
00386     close(pfds[1]);

```

```

00387     if(read(pfds[0], buf, sizeof(buf)) != 1){
00388         printf(1, "preempt read error");
00389         return;
00390     }
00391     close(pfds[0]);
00392     printf(1, "kill... ");
00393     kill(pid1);
00394     kill(pid2);
00395     kill(pid3);
00396     printf(1, "wait... ");
00397     wait();
00398     wait();
00399     wait();
00400     printf(1, "preempt ok\n");
00401 }

```

Referenced by [main\(\)](#).

5.240.2.29 rand()

```
unsigned int rand ( )
```

Definition at line 1742 of file [usertests.c](#).

```

01743 {
01744     randstate = randstate * 1664525 + 1013904223;
01745     return randstate;
01746 }

```

5.240.2.30 rmdot()

```
void rmdot (
    void )
```

Definition at line 1245 of file [usertests.c](#).

```

01246 {
01247     printf(1, "rmdot test\n");
01248     if(mkdir("dots") != 0){
01249         printf(1, "mkdir dots failed\n");
01250         exit();
01251     }
01252     if(chdir("dots") != 0){
01253         printf(1, "chdir dots failed\n");
01254         exit();
01255     }
01256     if(unlink(".") == 0){
01257         printf(1, "rm . worked!\n");
01258         exit();
01259     }
01260     if(unlink("..") == 0){
01261         printf(1, "rm .. worked!\n");
01262         exit();
01263     }
01264     if(chdir("/") != 0){
01265         printf(1, "chdir / failed\n");
01266         exit();
01267     }
01268     if(unlink("dots/.") == 0){
01269         printf(1, "unlink dots/. worked!\n");
01270         exit();
01271     }
01272     if(unlink("dots/..") == 0){
01273         printf(1, "unlink dots/.. worked!\n");
01274         exit();
01275     }
01276     if(unlink("dots") != 0){
01277         printf(1, "unlink dots failed!\n");
01278         exit();
01279     }
01280     printf(1, "rmdot ok\n");
01281 }

```

Referenced by [main\(\)](#).

5.240.2.31 sbrktest()

```
void sbrktest (
    void )
```

Definition at line 1415 of file `usertests.c`.

```
01416 {
01417     int fds[2], pid, pids[10], ppid;
01418     char *a, *b, *c, *lastaddr, *oldbrk, *p, scratch;
01419     uint amt;
01420
01421     printf(stdout, "sbrk test\n");
01422     oldbrk = sbrk(0);
01423
01424     // can one sbrk() less than a page?
01425     a = sbrk(0);
01426     int i;
01427     for(i = 0; i < 5000; i++){
01428         b = sbrk(1);
01429         if(b != a){
01430             printf(stdout, "sbrk test failed %d %x %x\n", i, a, b);
01431             exit();
01432         }
01433         *b = 1;
01434         a = b + 1;
01435     }
01436     pid = fork();
01437     if(pid < 0){
01438         printf(stdout, "sbrk test fork failed\n");
01439         exit();
01440     }
01441     c = sbrk(1);
01442     c = sbrk(1);
01443     if(c != a + 1){
01444         printf(stdout, "sbrk test failed post-fork\n");
01445         exit();
01446     }
01447     if(pid == 0)
01448         exit();
01449     wait();
01450
01451     // can one grow address space to something big?
01452     #define BIG (100*1024*1024)
01453     a = sbrk(0);
01454     amt = (BIG) - (uint)a;
01455     p = sbrk(amt);
01456     if (p != a) {
01457         printf(stdout, "sbrk test failed to grow big address space; enough phys mem?\n");
01458         exit();
01459     }
01460     lastaddr = (char*) (BIG-1);
01461     *lastaddr = 99;
01462
01463     // can one de-allocate?
01464     a = sbrk(0);
01465     c = sbrk(-4096);
01466     if(c == (char*)0xffffffff){
01467         printf(stdout, "sbrk could not deallocate\n");
01468         exit();
01469     }
01470     c = sbrk(0);
01471     if(c != a - 4096){
01472         printf(stdout, "sbrk deallocation produced wrong address, a %x c %x\n", a, c);
01473         exit();
01474     }
01475
01476     // can one re-allocate that page?
01477     a = sbrk(0);
01478     c = sbrk(4096);
01479     if(c != a || sbrk(0) != a + 4096){
01480         printf(stdout, "sbrk re-allocation failed, a %x c %x\n", a, c);
01481         exit();
01482     }
01483     if(*lastaddr == 99){
01484         // should be zero
01485         printf(stdout, "sbrk de-allocation didn't really deallocate\n");
01486         exit();
01487     }
01488
01489     a = sbrk(0);
01490     c = sbrk(-(sbrk(0) - oldbrk));
01491     if(c != a){
01492         printf(stdout, "sbrk downsize failed, a %x c %x\n", a, c);
01493         exit();
01494     }
01495 }
```

```

01494     }
01495
01496     // can we read the kernel's memory?
01497     for(a = (char*)(KERNBASE); a < (char*) (KERNBASE+2000000); a += 50000){
01498         ppid = getpid();
01499         pid = fork();
01500         if(pid < 0){
01501             printf(stdout, "fork failed\n");
01502             exit();
01503         }
01504         if(pid == 0){
01505             printf(stdout, "oops could read %x = %x\n", a, *a);
01506             kill(ppid);
01507             exit();
01508         }
01509         wait();
01510     }
01511
01512     // if we run the system out of memory, does it clean up the last
01513     // failed allocation?
01514     if(pipe(fds) != 0){
01515         printf(1, "pipe() failed\n");
01516         exit();
01517     }
01518     for(i = 0; i < sizeof(pids)/sizeof(pids[0]); i++){
01519         if((pids[i] = fork()) == 0){
01520             // allocate a lot of memory
01521             sbrk(BIG - (uint)sbrk(0));
01522             write(fds[1], "x", 1);
01523             // sit around until killed
01524             for(;;) sleep(1000);
01525         }
01526         if(pids[i] != -1)
01527             read(fds[0], &scratch, 1);
01528     }
01529     // if those failed allocations freed up the pages they did allocate,
01530     // we'll be able to allocate here
01531     c = sbrk(4096);
01532     for(i = 0; i < sizeof(pids)/sizeof(pids[0]); i++){
01533         if(pids[i] == -1)
01534             continue;
01535         kill(pids[i]);
01536         wait();
01537     }
01538     if(c == (char*)0xffffffff){
01539         printf(stdout, "failed sbrk leaked memory\n");
01540         exit();
01541     }
01542
01543     if(sbrk(0) > oldbrk)
01544         sbrk(-(sbrk(0) - oldbrk));
01545
01546     printf(stdout, "sbrk test OK\n");
01547 }

```

Referenced by [main\(\)](#).

5.240.2.32 sharedfd()

```

void sharedfd (
    void )

```

Definition at line 465 of file [usertests.c](#).

```

00466 {
00467     int fd, pid, i, n, nc, np;
00468     char buf[10];
00469
00470     printf(1, "sharedfd test\n");
00471
00472     unlink("sharedfd");
00473     fd = open("sharedfd", O_CREATE|O_RDWR);
00474     if(fd < 0){
00475         printf(1, "fstests: cannot open sharedfd for writing");
00476         return;
00477     }
00478     pid = fork();
00479     memset(buf, pid==0?'c':'p', sizeof(buf));

```

```

00480     for(i = 0; i < 1000; i++){
00481         if(write(fd, buf, sizeof(buf)) != sizeof(buf)){
00482             printf(1, "fstests: write sharedfd failed\n");
00483             break;
00484         }
00485     }
00486     if(pid == 0)
00487         exit();
00488     else
00489         wait();
00490     close(fd);
00491     fd = open("sharedfd", 0);
00492     if(fd < 0){
00493         printf(1, "fstests: cannot open sharedfd for reading\n");
00494         return;
00495     }
00496     nc = np = 0;
00497     while((n = read(fd, buf, sizeof(buf))) > 0){
00498         for(i = 0; i < sizeof(buf); i++){
00499             if(buf[i] == 'c')
00500                 nc++;
00501             if(buf[i] == 'p')
00502                 np++;
00503         }
00504     }
00505     close(fd);
00506     unlink("sharedfd");
00507     if(nc == 10000 && np == 10000){
00508         printf(1, "sharedfd ok\n");
00509     } else {
00510         printf(1, "sharedfd oops %d %d\n", nc, np);
00511         exit();
00512     }
00513 }

```

Referenced by [main\(\)](#).

5.240.2.33 subdird()

```

void subdird (
    void )

```

Definition at line 935 of file [usertests.c](#).

```

00936 {
00937     int fd, cc;
00938
00939     printf(1, "subdird test\n");
00940
00941     unlink("ff");
00942     if(mkdir("dd") != 0){
00943         printf(1, "subdird mkdir dd failed\n");
00944         exit();
00945     }
00946
00947     fd = open("dd/ff", O_CREATE | O_RDWR);
00948     if(fd < 0){
00949         printf(1, "create dd/ff failed\n");
00950         exit();
00951     }
00952     write(fd, "ff", 2);
00953     close(fd);
00954
00955     if(unlink("dd") >= 0){
00956         printf(1, "unlink dd (non-empty dir) succeeded!\n");
00957         exit();
00958     }
00959
00960     if(mkdir("/dd/dd") != 0){
00961         printf(1, "subdird mkdir dd/dd failed\n");
00962         exit();
00963     }
00964
00965     fd = open("dd/dd/ff", O_CREATE | O_RDWR);
00966     if(fd < 0){
00967         printf(1, "create dd/dd/ff failed\n");
00968         exit();
00969     }

```



```
00970     write(fd, "FF", 2);
00971     close(fd);
00972
00973     fd = open("dd/dd/../../ff", 0);
00974     if(fd < 0){
00975         printf(1, "open dd/dd/../../ff failed\n");
00976         exit();
00977     }
00978     cc = read(fd, buf, sizeof(buf));
00979     if(cc != 2 || buf[0] != 'f'){
00980         printf(1, "dd/dd/../../ff wrong content\n");
00981         exit();
00982     }
00983     close(fd);
00984
00985     if(link("dd/dd/ff", "dd/dd/ffff") != 0){
00986         printf(1, "link dd/dd/ff dd/dd/ffff failed\n");
00987         exit();
00988     }
00989
00990     if(unlink("dd/dd/ff") != 0){
00991         printf(1, "unlink dd/dd/ff failed\n");
00992         exit();
00993     }
00994     if(open("dd/dd/ff", O_RDONLY) >= 0){
00995         printf(1, "open (unlinked) dd/dd/ff succeeded\n");
00996         exit();
00997     }
00998
00999     if(chdir("dd") != 0){
01000         printf(1, "chdir dd failed\n");
01001         exit();
01002     }
01003     if(chdir("dd/../../dd") != 0){
01004         printf(1, "chdir dd/../../dd failed\n");
01005         exit();
01006     }
01007     if(chdir("dd/../../../dd") != 0){
01008         printf(1, "chdir dd/../../../dd failed\n");
01009         exit();
01010     }
01011     if(chdir("../..") != 0){
01012         printf(1, "chdir ../.. failed\n");
01013         exit();
01014     }
01015
01016     fd = open("dd/dd/ffff", 0);
01017     if(fd < 0){
01018         printf(1, "open dd/dd/ffff failed\n");
01019         exit();
01020     }
01021     if(read(fd, buf, sizeof(buf)) != 2){
01022         printf(1, "read dd/dd/ffff wrong len\n");
01023         exit();
01024     }
01025     close(fd);
01026
01027     if(open("dd/dd/ff", O_RDONLY) >= 0){
01028         printf(1, "open (unlinked) dd/dd/ff succeeded!\n");
01029         exit();
01030     }
01031
01032     if(open("dd/ff/ff", O_CREATE|O_RDWR) >= 0){
01033         printf(1, "create dd/ff/ff succeeded!\n");
01034         exit();
01035     }
01036     if(open("dd/xx/ff", O_CREATE|O_RDWR) >= 0){
01037         printf(1, "create dd/xx/ff succeeded!\n");
01038         exit();
01039     }
01040     if(open("dd", O_CREATE) >= 0){
01041         printf(1, "create dd succeeded!\n");
01042         exit();
01043     }
01044     if(open("dd", O_RDWR) >= 0){
01045         printf(1, "open dd rdwr succeeded!\n");
01046         exit();
01047     }
01048     if(open("dd", O_WRONLY) >= 0){
01049         printf(1, "open dd wronly succeeded!\n");
01050         exit();
01051     }
01052     if(link("dd/ff/ff", "dd/dd/xx") == 0){
01053         printf(1, "link dd/ff/ff dd/dd/xx succeeded!\n");
01054         exit();
01055     }
01056     if(link("dd/xx/ff", "dd/dd/xx") == 0){
```

```

01057     printf(1, "link dd/xx/ff dd/dd/xx succeeded!\n");
01058     exit();
01059 }
01060 if(link("dd/ff", "dd/dd/ffff") == 0){
01061     printf(1, "link dd/ff dd/dd/ffff succeeded!\n");
01062     exit();
01063 }
01064 if(mkdir("dd/ff/ff") == 0){
01065     printf(1, "mkdir dd/ff/ff succeeded!\n");
01066     exit();
01067 }
01068 if(mkdir("dd/xx/ff") == 0){
01069     printf(1, "mkdir dd/xx/ff succeeded!\n");
01070     exit();
01071 }
01072 if(mkdir("dd/dd/ffff") == 0){
01073     printf(1, "mkdir dd/dd/ffff succeeded!\n");
01074     exit();
01075 }
01076 if(unlink("dd/xx/ff") == 0){
01077     printf(1, "unlink dd/xx/ff succeeded!\n");
01078     exit();
01079 }
01080 if(unlink("dd/ff/ff") == 0){
01081     printf(1, "unlink dd/ff/ff succeeded!\n");
01082     exit();
01083 }
01084 if(chdir("dd/ff") == 0){
01085     printf(1, "chdir dd/ff succeeded!\n");
01086     exit();
01087 }
01088 if(chdir("dd/xx") == 0){
01089     printf(1, "chdir dd/xx succeeded!\n");
01090     exit();
01091 }
01092 }
01093 if(unlink("dd/dd/ffff") != 0){
01094     printf(1, "unlink dd/dd/ff failed\n");
01095     exit();
01096 }
01097 if(unlink("dd/ff") != 0){
01098     printf(1, "unlink dd/ff failed\n");
01099     exit();
01100 }
01101 if(unlink("dd") == 0){
01102     printf(1, "unlink non-empty dd succeeded!\n");
01103     exit();
01104 }
01105 if(unlink("dd/dd") < 0){
01106     printf(1, "unlink dd/dd failed\n");
01107     exit();
01108 }
01109 if(unlink("dd") < 0){
01110     printf(1, "unlink dd failed\n");
01111     exit();
01112 }
01113 }
01114 printf(1, "subdir ok\n");
01115 }

```

Referenced by [main\(\)](#).

5.240.2.34 uio()

```
void uio ( )
```

Definition at line 1699 of file [usertests.c](#).

```

01700 {
01701     #define RTC_ADDR 0x70
01702     #define RTC_DATA 0x71
01703
01704     ushort port = 0;
01705     uchar val = 0;
01706     int pid;
01707
01708     printf(1, "uio test\n");
01709     pid = fork();
01710     if(pid == 0){

```

```

01711     port = RTC_ADDR;
01712     val = 0x09; /* year */
01713     /* http://wiki.osdev.org/Inline_Assembly/Examples */
01714     asm volatile("outb %0,%1::"a"(val), "d" (port));
01715     port = RTC_DATA;
01716     asm volatile("inb %1,%0" : "=a" (val) : "d" (port));
01717     printf(1, "uio: uio succeeded; test FAILED\n");
01718     exit();
01719 } else if(pid < 0){
01720     printf(1, "fork failed\n");
01721     exit();
01722 }
01723 wait();
01724 printf(1, "uio test done\n");
01725 }

```

Referenced by [main\(\)](#).

5.240.2.35 unlinkread()

```

void unlinkread (
    void )

```

Definition at line 657 of file [usertests.c](#).

```

00658 {
00659     int fd, fd1;
00660
00661     printf(1, "unlinkread test\n");
00662     fd = open("unlinkread", O_CREATE | O_RDWR);
00663     if(fd < 0){
00664         printf(1, "create unlinkread failed\n");
00665         exit();
00666     }
00667     write(fd, "hello", 5);
00668     close(fd);
00669
00670     fd = open("unlinkread", O_RDWR);
00671     if(fd < 0){
00672         printf(1, "open unlinkread failed\n");
00673         exit();
00674     }
00675     if(unlink("unlinkread") != 0){
00676         printf(1, "unlink unlinkread failed\n");
00677         exit();
00678     }
00679
00680     fd1 = open("unlinkread", O_CREATE | O_RDWR);
00681     write(fd1, "yyy", 3);
00682     close(fd1);
00683
00684     if(read(fd, buf, sizeof(buf)) != 5){
00685         printf(1, "unlinkread read failed");
00686         exit();
00687     }
00688     if(buf[0] != 'h'){
00689         printf(1, "unlinkread wrong data\n");
00690         exit();
00691     }
00692     if(write(fd, buf, 10) != 10){
00693         printf(1, "unlinkread write failed\n");
00694         exit();
00695     }
00696     close(fd);
00697     unlink("unlinkread");
00698     printf(1, "unlinkread ok\n");
00699 }

```

Referenced by [main\(\)](#).

5.240.2.36 validateint()

```
void validateint (
    int * p )
```

Definition at line 1550 of file [usertests.c](#).

```
01551 {
01552     int res;
01553     asm("mov %%esp, %%ebx\n\t"
01554         "mov %3, %%esp\n\t"
01555         "int %2\n\t"
01556         "mov %%ebx, %%esp" :
01557         "=a" (res) :
01558         "a" (SYS_sleep), "n" (T_SYSCALL), "c" (p) :
01559         "ebx");
01560 }
```

Referenced by [validatetest\(\)](#).

5.240.2.37 validatetest()

```
void validatetest (
    void )
```

Definition at line 1563 of file [usertests.c](#).

```
01564 {
01565     int hi, pid;
01566     uint p;
01567
01568     printf(stdout, "validate test\n");
01569     hi = 1100*1024;
01570
01571     for(p = 0; p <= (uint)hi; p += 4096){
01572         if((pid = fork()) == 0){
01573             // try to crash the kernel by passing in a badly placed integer
01574             validateint((int*)p);
01575             exit();
01576         }
01577         sleep(0);
01578         sleep(0);
01579         kill(pid);
01580         wait();
01581
01582         // try to crash the kernel by passing in a bad string pointer
01583         if(link("nosuchfile", (char*)p) != -1){
01584             printf(stdout, "link should not succeed\n");
01585             exit();
01586         }
01587     }
01588
01589     printf(stdout, "validate ok\n");
01590 }
```

Referenced by [main\(\)](#).

5.240.2.38 writetest()

```
void writetest (
    void )
```

Definition at line 139 of file [usertests.c](#).

```
00140 {
00141     int fd;
00142     int i;
00143
00144     printf(stdout, "small file test\n");
00145     fd = open("small", O_CREATE|O_RDWR);
00146     if(fd >= 0){
00147         printf(stdout, "creat small succeeded; ok\n");
00148     } else {
00149         printf(stdout, "error: creat small failed!\n");
00150         exit();
00151     }
00152     for(i = 0; i < 100; i++){
00153         if(write(fd, "aaaaaaaaa", 10) != 10){
00154             printf(stdout, "error: write aa %d new file failed\n", i);
00155             exit();
00156         }
00157         if(write(fd, "bbbbbbbbbbb", 10) != 10){
00158             printf(stdout, "error: write bb %d new file failed\n", i);
00159             exit();
00160         }
00161     }
00162     printf(stdout, "writes ok\n");
00163     close(fd);
00164     fd = open("small", O_RDONLY);
00165     if(fd >= 0){
00166         printf(stdout, "open small succeeded ok\n");
00167     } else {
00168         printf(stdout, "error: open small failed!\n");
00169         exit();
00170     }
00171     i = read(fd, buf, 2000);
00172     if(i == 2000){
00173         printf(stdout, "read succeeded ok\n");
00174     } else {
00175         printf(stdout, "read failed\n");
00176         exit();
00177     }
00178     close(fd);
00179
00180     if(unlink("small") < 0){
00181         printf(stdout, "unlink small failed\n");
00182         exit();
00183     }
00184     printf(stdout, "small file test ok\n");
00185 }
```

Referenced by [main\(\)](#).

5.240.2.39 writetest1()

```
void writetest1 (
    void )
```

Definition at line 188 of file [usertests.c](#).

```
00189 {
00190     int i, fd, n;
00191
00192     printf(stdout, "big files test\n");
00193
00194     fd = open("big", O_CREATE|O_RDWR);
00195     if(fd < 0){
00196         printf(stdout, "error: creat big failed!\n");
00197         exit();
00198     }
00199
00200     for(i = 0; i < MAXFILE; i++){
00201         ((int*)buf)[0] = i;
```

```

00202     if(write(fd, buf, 512) != 512){
00203         printf(stdout, "error: write big file failed\n", i);
00204         exit();
00205     }
00206 }
00207
00208 close(fd);
00209
00210 fd = open("big", O_RDONLY);
00211 if(fd < 0){
00212     printf(stdout, "error: open big failed!\n");
00213     exit();
00214 }
00215
00216 n = 0;
00217 for(;;){
00218     i = read(fd, buf, 512);
00219     if(i == 0){
00220         if(n == MAXFILE - 1){
00221             printf(stdout, "read only %d blocks from big", n);
00222             exit();
00223         }
00224         break;
00225     } else if(i != 512){
00226         printf(stdout, "read failed %d\n", i);
00227         exit();
00228     }
00229     if(((int*)buf)[0] != n){
00230         printf(stdout, "read content of block %d is %d\n",
00231             n, ((int*)buf)[0]);
00232         exit();
00233     }
00234     n++;
00235 }
00236 close(fd);
00237 if(unlink("big") < 0){
00238     printf(stdout, "unlink big failed\n");
00239     exit();
00240 }
00241 printf(stdout, "big files ok\n");
00242 }

```

Referenced by [main\(\)](#).

5.240.3 Variable Documentation

5.240.3.1 buf

```
char buf[8192]
```

Definition at line 11 of file [usertests.c](#).

Referenced by [pipe1\(\)](#), and [sharedfd\(\)](#).

5.240.3.2 echoargv

```
char* echoargv[] = { "echo", "ALL", "TESTS", "PASSED", 0 }
```

Definition at line 13 of file [usertests.c](#).

Referenced by [exectest\(\)](#).

5.240.3.3 name

```
char name[3]
```

Definition at line 12 of file [usertests.c](#).

Referenced by [bigdir\(\)](#), [concreate\(\)](#), [create\(\)](#), [createdelete\(\)](#), [createtest\(\)](#), [dirlink\(\)](#), [dirlookup\(\)](#), [fsfull\(\)](#), [initlock\(\)](#), [initsleeplock\(\)](#), [main\(\)](#), [namei\(\)](#), [nameiparent\(\)](#), [namex\(\)](#), [skipelem\(\)](#), [sys_getprocs\(\)](#), [sys_link\(\)](#), [sys_unlink\(\)](#), [trap\(\)](#), and [wc\(\)](#).

5.240.3.4 randstate

```
unsigned long randstate = 1
```

Definition at line 1740 of file [usertests.c](#).

Referenced by [rand\(\)](#).

5.240.3.5 stdout

```
int stdout = 1
```

Definition at line 14 of file [usertests.c](#).

Referenced by [bigargtest\(\)](#), [bsstest\(\)](#), [createtest\(\)](#), [dirtest\(\)](#), [exectest\(\)](#), [exitiputtest\(\)](#), [iputtest\(\)](#), [openiputtest\(\)](#), [opentest\(\)](#), [sbrktest\(\)](#), [validatetest\(\)](#), [writetest\(\)](#), and [writetest1\(\)](#).

5.240.3.6 uninit

```
char uninit[10000]
```

Definition at line 1593 of file [usertests.c](#).

Referenced by [bsstest\(\)](#).

5.241 usertests.c

[Go to the documentation of this file.](#)

```

00001 #include "param.h"
00002 #include "types.h"
00003 #include "stat.h"
00004 #include "user.h"
00005 #include "fs.h"
00006 #include "fcntl.h"
00007 #include "syscall.h"
00008 #include "traps.h"
00009 #include "memlayout.h"
00010
00011 char buf[8192];
00012 char name[3];
00013 char *echoargv[] = { "echo", "ALL", "TESTS", "PASSED", 0 };
00014 int stdout = 1;
00015
00016 // does chdir() call iput(p->cwd) in a transaction?
00017 void
00018 iputtest(void)
00019 {
00020     printf(stdout, "iput test\n");
00021
00022     if(mkdir("iputdir") < 0){
00023         printf(stdout, "mkdir failed\n");
00024         exit();
00025     }
00026     if(chdir("iputdir") < 0){
00027         printf(stdout, "chdir iputdir failed\n");
00028         exit();
00029     }
00030     if(unlink("../iputdir") < 0){
00031         printf(stdout, "unlink ../iputdir failed\n");
00032         exit();
00033     }
00034     if(chdir("/") < 0){
00035         printf(stdout, "chdir / failed\n");
00036         exit();
00037     }
00038     printf(stdout, "iput test ok\n");
00039 }
00040
00041 // does exit() call iput(p->cwd) in a transaction?
00042 void
00043 exitiputtest(void)
00044 {
00045     int pid;
00046
00047     printf(stdout, "exitiput test\n");
00048
00049     pid = fork();
00050     if(pid < 0){
00051         printf(stdout, "fork failed\n");
00052         exit();
00053     }
00054     if(pid == 0){
00055         if(mkdir("iputdir") < 0){
00056             printf(stdout, "mkdir failed\n");
00057             exit();
00058         }
00059         if(chdir("iputdir") < 0){
00060             printf(stdout, "child chdir failed\n");
00061             exit();
00062         }
00063         if(unlink("../iputdir") < 0){
00064             printf(stdout, "unlink ../iputdir failed\n");
00065             exit();
00066         }
00067         exit();
00068     }
00069     wait();
00070     printf(stdout, "exitiput test ok\n");
00071 }
00072
00073 // does the error path in open() for attempt to write a
00074 // directory call iput() in a transaction?
00075 // needs a hacked kernel that pauses just after the namei()
00076 // call in sys_open():
00077 //     if((ip = namei(path)) == 0)
00078 //         return -1;
00079 //     {
00080 //         int i;
00081 //         for(i = 0; i < 10000; i++)
00082 //             yield();

```



```

00083 //    }
00084 void
00085 openiputtest(void)
00086 {
00087     int pid;
00088
00089     printf(stdout, "openiput test\n");
00090     if(mkdir("oidir") < 0){
00091         printf(stdout, "mkdir oidir failed\n");
00092         exit();
00093     }
00094     pid = fork();
00095     if(pid < 0){
00096         printf(stdout, "fork failed\n");
00097         exit();
00098     }
00099     if(pid == 0){
00100         int fd = open("oidir", O_RDWR);
00101         if(fd >= 0){
00102             printf(stdout, "open directory for write succeeded\n");
00103             exit();
00104         }
00105         exit();
00106     }
00107     sleep(1);
00108     if(unlink("oidir") != 0){
00109         printf(stdout, "unlink failed\n");
00110         exit();
00111     }
00112     wait();
00113     printf(stdout, "openiput test ok\n");
00114 }
00115
00116 // simple file system tests
00117 void
00118 opentest(void)
00119 {
00120     int fd;
00121
00122     printf(stdout, "open test\n");
00123     fd = open("echo", 0);
00124     if(fd < 0){
00125         printf(stdout, "open echo failed!\n");
00126         exit();
00127     }
00128     close(fd);
00129     fd = open("doesnotexist", 0);
00130     if(fd >= 0){
00131         printf(stdout, "open doesnotexist succeeded!\n");
00132         exit();
00133     }
00134     printf(stdout, "open test ok\n");
00135 }
00136
00137 void
00138 writetest(void)
00139 {
00140     int fd;
00141     int i;
00142
00143     printf(stdout, "small file test\n");
00144     fd = open("small", O_CREATE|O_RDWR);
00145     if(fd >= 0){
00146         printf(stdout, "creat small succeeded; ok\n");
00147     } else {
00148         printf(stdout, "error: creat small failed!\n");
00149         exit();
00150     }
00151     for(i = 0; i < 100; i++){
00152         if(write(fd, "aaaaaaaaaa", 10) != 10){
00153             printf(stdout, "error: write aa %d new file failed\n", i);
00154             exit();
00155         }
00156         if(write(fd, "bbbbbbbbbb", 10) != 10){
00157             printf(stdout, "error: write bb %d new file failed\n", i);
00158             exit();
00159         }
00160     }
00161     printf(stdout, "writes ok\n");
00162     close(fd);
00163     fd = open("small", O_RDONLY);
00164     if(fd >= 0){
00165         printf(stdout, "open small succeeded ok\n");
00166     } else {
00167         printf(stdout, "error: open small failed!\n");
00168         exit();
00169     }

```

```

00170     }
00171     i = read(fd, buf, 2000);
00172     if(i == 2000){
00173         printf(stdout, "read succeeded ok\n");
00174     } else {
00175         printf(stdout, "read failed\n");
00176         exit();
00177     }
00178     close(fd);
00179
00180     if(unlink("small") < 0){
00181         printf(stdout, "unlink small failed\n");
00182         exit();
00183     }
00184     printf(stdout, "small file test ok\n");
00185 }
00186
00187 void
00188 writetest1(void)
00189 {
00190     int i, fd, n;
00191
00192     printf(stdout, "big files test\n");
00193
00194     fd = open("big", O_CREATE|O_RDWR);
00195     if(fd < 0){
00196         printf(stdout, "error: creat big failed!\n");
00197         exit();
00198     }
00199
00200     for(i = 0; i < MAXFILE; i++){
00201         ((int*)buf)[0] = i;
00202         if(write(fd, buf, 512) != 512){
00203             printf(stdout, "error: write big file failed\n", i);
00204             exit();
00205         }
00206     }
00207
00208     close(fd);
00209
00210     fd = open("big", O_RDONLY);
00211     if(fd < 0){
00212         printf(stdout, "error: open big failed!\n");
00213         exit();
00214     }
00215
00216     n = 0;
00217     for(;;){
00218         i = read(fd, buf, 512);
00219         if(i == 0){
00220             if(n == MAXFILE - 1){
00221                 printf(stdout, "read only %d blocks from big", n);
00222                 exit();
00223             }
00224             break;
00225         } else if(i != 512){
00226             printf(stdout, "read failed %d\n", i);
00227             exit();
00228         }
00229         if(((int*)buf)[0] != n){
00230             printf(stdout, "read content of block %d is %d\n",
00231                 n, ((int*)buf)[0]);
00232             exit();
00233         }
00234         n++;
00235     }
00236     close(fd);
00237     if(unlink("big") < 0){
00238         printf(stdout, "unlink big failed\n");
00239         exit();
00240     }
00241     printf(stdout, "big files ok\n");
00242 }
00243
00244 void
00245 createtest(void)
00246 {
00247     int i, fd;
00248
00249     printf(stdout, "many creates, followed by unlink test\n");
00250
00251     name[0] = 'a';
00252     name[2] = '\0';
00253     for(i = 0; i < 52; i++){
00254         name[1] = '0' + i;
00255         fd = open(name, O_CREATE|O_RDWR);
00256         close(fd);

```

```

00257     }
00258     name[0] = 'a';
00259     name[2] = '\\0';
00260     for(i = 0; i < 52; i++){
00261         name[1] = '0' + i;
00262         unlink(name);
00263     }
00264     printf(stdout, "many creates, followed by unlink; ok\n");
00265 }
00266
00267 void dirtest(void)
00268 {
00269     printf(stdout, "mkdir test\n");
00270
00271     if(mkdir("dir0") < 0){
00272         printf(stdout, "mkdir failed\n");
00273         exit();
00274     }
00275
00276     if(chdir("dir0") < 0){
00277         printf(stdout, "chdir dir0 failed\n");
00278         exit();
00279     }
00280
00281     if(chdir("..") < 0){
00282         printf(stdout, "chdir .. failed\n");
00283         exit();
00284     }
00285
00286     if(unlink("dir0") < 0){
00287         printf(stdout, "unlink dir0 failed\n");
00288         exit();
00289     }
00290     printf(stdout, "mkdir test ok\n");
00291 }
00292
00293 void
00294 exectest(void)
00295 {
00296     printf(stdout, "exec test\n");
00297     if(exec("echo", echoargv) < 0){
00298         printf(stdout, "exec echo failed\n");
00299         exit();
00300     }
00301 }
00302
00303 // simple fork and pipe read/write
00304
00305 void
00306 pipel(void)
00307 {
00308     int fds[2], pid;
00309     int seq, i, n, cc, total;
00310
00311     if(pipe(fds) != 0){
00312         printf(1, "pipe() failed\n");
00313         exit();
00314     }
00315     pid = fork();
00316     seq = 0;
00317     if(pid == 0){
00318         close(fds[0]);
00319         for(n = 0; n < 5; n++){
00320             for(i = 0; i < 1033; i++)
00321                 buf[i] = seq++;
00322             if(write(fds[1], buf, 1033) != 1033){
00323                 printf(1, "pipel oops 1\n");
00324                 exit();
00325             }
00326         }
00327         exit();
00328     } else if(pid > 0){
00329         close(fds[1]);
00330         total = 0;
00331         cc = 1;
00332         while((n = read(fds[0], buf, cc)) > 0){
00333             for(i = 0; i < n; i++){
00334                 if((buf[i] & 0xff) != (seq++ & 0xff)){
00335                     printf(1, "pipel oops 2\n");
00336                     return;
00337                 }
00338             }
00339             total += n;
00340             cc = cc * 2;
00341             if(cc > sizeof(buf))
00342                 cc = sizeof(buf);
00343         }

```

```

00344     if(total != 5 * 1033){
00345         printf(1, "pipe1 oops 3 total %d\n", total);
00346         exit();
00347     }
00348     close(fds[0]);
00349     wait();
00350 } else {
00351     printf(1, "fork() failed\n");
00352     exit();
00353 }
00354 printf(1, "pipe1 ok\n");
00355 }
00356
00357 // meant to be run w/ at most two CPUs
00358 void
00359 preempt(void)
00360 {
00361     int pid1, pid2, pid3;
00362     int pfds[2];
00363
00364     printf(1, "preempt: ");
00365     pid1 = fork();
00366     if(pid1 == 0)
00367         for(;;)
00368             ;
00369
00370     pid2 = fork();
00371     if(pid2 == 0)
00372         for(;;)
00373             ;
00374
00375     pipe(pfds);
00376     pid3 = fork();
00377     if(pid3 == 0){
00378         close(pfds[0]);
00379         if(write(pfds[1], "x", 1) != 1)
00380             printf(1, "preempt write error");
00381         close(pfds[1]);
00382         for(;;)
00383             ;
00384     }
00385
00386     close(pfds[1]);
00387     if(read(pfds[0], buf, sizeof(buf)) != 1){
00388         printf(1, "preempt read error");
00389         return;
00390     }
00391     close(pfds[0]);
00392     printf(1, "kill... ");
00393     kill(pid1);
00394     kill(pid2);
00395     kill(pid3);
00396     printf(1, "wait... ");
00397     wait();
00398     wait();
00399     wait();
00400     printf(1, "preempt ok\n");
00401 }
00402
00403 // try to find any races between exit and wait
00404 void
00405 exitwait(void)
00406 {
00407     int i, pid;
00408
00409     for(i = 0; i < 100; i++){
00410         pid = fork();
00411         if(pid < 0){
00412             printf(1, "fork failed\n");
00413             return;
00414         }
00415         if(pid){
00416             if(wait() != pid){
00417                 printf(1, "wait wrong pid\n");
00418                 return;
00419             }
00420         } else {
00421             exit();
00422         }
00423     }
00424     printf(1, "exitwait ok\n");
00425 }
00426
00427 void
00428 mem(void)
00429 {
00430     void *m1, *m2;

```

```

00431 int pid, ppid;
00432
00433 printf(1, "mem test\n");
00434 ppid = getpid();
00435 if((pid = fork()) == 0){
00436     m1 = 0;
00437     while((m2 = malloc(10001)) != 0){
00438         *(char**)m2 = m1;
00439         m1 = m2;
00440     }
00441     while(m1){
00442         m2 = *(char**)m1;
00443         free(m1);
00444         m1 = m2;
00445     }
00446     m1 = malloc(1024*20);
00447     if(m1 == 0){
00448         printf(1, "couldn't allocate mem?!!\n");
00449         kill(ppid);
00450         exit();
00451     }
00452     free(m1);
00453     printf(1, "mem ok\n");
00454     exit();
00455 } else {
00456     wait();
00457 }
00458 }
00459
00460 // More file system tests
00461
00462 // two processes write to the same file descriptor
00463 // is the offset shared? does inode locking work?
00464 void
00465 sharedfd(void)
00466 {
00467     int fd, pid, i, n, nc, np;
00468     char buf[10];
00469
00470     printf(1, "sharedfd test\n");
00471
00472     unlink("sharedfd");
00473     fd = open("sharedfd", O_CREATE|O_RDWR);
00474     if(fd < 0){
00475         printf(1, "fstests: cannot open sharedfd for writing");
00476         return;
00477     }
00478     pid = fork();
00479     memset(buf, pid==0?'c':'p', sizeof(buf));
00480     for(i = 0; i < 1000; i++){
00481         if(write(fd, buf, sizeof(buf)) != sizeof(buf)){
00482             printf(1, "fstests: write sharedfd failed\n");
00483             break;
00484         }
00485     }
00486     if(pid == 0)
00487         exit();
00488     else
00489         wait();
00490     close(fd);
00491     fd = open("sharedfd", 0);
00492     if(fd < 0){
00493         printf(1, "fstests: cannot open sharedfd for reading\n");
00494         return;
00495     }
00496     nc = np = 0;
00497     while((n = read(fd, buf, sizeof(buf))) > 0){
00498         for(i = 0; i < sizeof(buf); i++){
00499             if(buf[i] == 'c')
00500                 nc++;
00501             if(buf[i] == 'p')
00502                 np++;
00503         }
00504     }
00505     close(fd);
00506     unlink("sharedfd");
00507     if(nc == 10000 && np == 10000){
00508         printf(1, "sharedfd ok\n");
00509     } else {
00510         printf(1, "sharedfd oops %d %d\n", nc, np);
00511         exit();
00512     }
00513 }
00514
00515 // four processes write different files at the same
00516 // time, to test block allocation.
00517 void

```

```

00518 fourfiles(void)
00519 {
00520     int fd, pid, i, j, n, total, pi;
00521     char *names[] = { "f0", "f1", "f2", "f3" };
00522     char *fname;
00523
00524     printf(1, "fourfiles test\n");
00525
00526     for(pi = 0; pi < 4; pi++){
00527         fname = names[pi];
00528         unlink(fname);
00529
00530         pid = fork();
00531         if(pid < 0){
00532             printf(1, "fork failed\n");
00533             exit();
00534         }
00535
00536         if(pid == 0){
00537             fd = open(fname, O_CREATE | O_RDWR);
00538             if(fd < 0){
00539                 printf(1, "create failed\n");
00540                 exit();
00541             }
00542
00543             memset(buf, '0'+pi, 512);
00544             for(i = 0; i < 12; i++){
00545                 if((n = write(fd, buf, 500)) != 500){
00546                     printf(1, "write failed %d\n", n);
00547                     exit();
00548                 }
00549             }
00550             exit();
00551         }
00552     }
00553
00554     for(pi = 0; pi < 4; pi++){
00555         wait();
00556     }
00557
00558     for(i = 0; i < 2; i++){
00559         fname = names[i];
00560         fd = open(fname, 0);
00561         total = 0;
00562         while((n = read(fd, buf, sizeof(buf))) > 0){
00563             for(j = 0; j < n; j++){
00564                 if(buf[j] != '0'+i){
00565                     printf(1, "wrong char\n");
00566                     exit();
00567                 }
00568             }
00569             total += n;
00570         }
00571         close(fd);
00572         if(total != 12*500){
00573             printf(1, "wrong length %d\n", total);
00574             exit();
00575         }
00576         unlink(fname);
00577     }
00578
00579     printf(1, "fourfiles ok\n");
00580 }
00581
00582 // four processes create and delete different files in same directory
00583 void
00584 createdelete(void)
00585 {
00586     enum { N = 20 };
00587     int pid, i, fd, pi;
00588     char name[32];
00589
00590     printf(1, "createdelete test\n");
00591
00592     for(pi = 0; pi < 4; pi++){
00593         pid = fork();
00594         if(pid < 0){
00595             printf(1, "fork failed\n");
00596             exit();
00597         }
00598
00599         if(pid == 0){
00600             name[0] = 'p' + pi;
00601             name[2] = '\0';
00602             for(i = 0; i < N; i++){
00603                 name[1] = '0' + i;
00604                 fd = open(name, O_CREATE | O_RDWR);

```

```

00605         if(fd < 0){
00606             printf(1, "create failed\n");
00607             exit();
00608         }
00609         close(fd);
00610         if(i > 0 && (i % 2) == 0){
00611             name[1] = '0' + (i / 2);
00612             if(unlink(name) < 0){
00613                 printf(1, "unlink failed\n");
00614                 exit();
00615             }
00616         }
00617     }
00618     exit();
00619 }
00620 }
00621
00622 for(pi = 0; pi < 4; pi++){
00623     wait();
00624 }
00625
00626 name[0] = name[1] = name[2] = 0;
00627 for(i = 0; i < N; i++){
00628     for(pi = 0; pi < 4; pi++){
00629         name[0] = 'p' + pi;
00630         name[1] = '0' + i;
00631         fd = open(name, 0);
00632         if((i == 0 || i >= N/2) && fd < 0){
00633             printf(1, "oops createdelete %s didn't exist\n", name);
00634             exit();
00635         } else if((i >= 1 && i < N/2) && fd >= 0){
00636             printf(1, "oops createdelete %s did exist\n", name);
00637             exit();
00638         }
00639         if(fd >= 0)
00640             close(fd);
00641     }
00642 }
00643
00644 for(i = 0; i < N; i++){
00645     for(pi = 0; pi < 4; pi++){
00646         name[0] = 'p' + i;
00647         name[1] = '0' + i;
00648         unlink(name);
00649     }
00650 }
00651
00652 printf(1, "createdelete ok\n");
00653 }
00654
00655 // can I unlink a file and still read it?
00656 void
00657 unlinkread(void)
00658 {
00659     int fd, fd1;
00660
00661     printf(1, "unlinkread test\n");
00662     fd = open("unlinkread", O_CREATE | O_RDWR);
00663     if(fd < 0){
00664         printf(1, "create unlinkread failed\n");
00665         exit();
00666     }
00667     write(fd, "hello", 5);
00668     close(fd);
00669
00670     fd = open("unlinkread", O_RDWR);
00671     if(fd < 0){
00672         printf(1, "open unlinkread failed\n");
00673         exit();
00674     }
00675     if(unlink("unlinkread") != 0){
00676         printf(1, "unlink unlinkread failed\n");
00677         exit();
00678     }
00679
00680     fd1 = open("unlinkread", O_CREATE | O_RDWR);
00681     write(fd1, "yyy", 3);
00682     close(fd1);
00683
00684     if(read(fd, buf, sizeof(buf)) != 5){
00685         printf(1, "unlinkread read failed");
00686         exit();
00687     }
00688     if(buf[0] != 'h'){
00689         printf(1, "unlinkread wrong data\n");
00690         exit();
00691     }

```

```

00692     if(write(fd, buf, 10) != 10){
00693         printf(1, "unlinkread write failed\n");
00694         exit();
00695     }
00696     close(fd);
00697     unlink("unlinkread");
00698     printf(1, "unlinkread ok\n");
00699 }
00700
00701 void
00702 linktest(void)
00703 {
00704     int fd;
00705
00706     printf(1, "linktest\n");
00707
00708     unlink("lf1");
00709     unlink("lf2");
00710
00711     fd = open("lf1", O_CREATE|O_RDWR);
00712     if(fd < 0){
00713         printf(1, "create lf1 failed\n");
00714         exit();
00715     }
00716     if(write(fd, "hello", 5) != 5){
00717         printf(1, "write lf1 failed\n");
00718         exit();
00719     }
00720     close(fd);
00721
00722     if(link("lf1", "lf2") < 0){
00723         printf(1, "link lf1 lf2 failed\n");
00724         exit();
00725     }
00726     unlink("lf1");
00727
00728     if(open("lf1", 0) >= 0){
00729         printf(1, "unlinked lf1 but it is still there!\n");
00730         exit();
00731     }
00732
00733     fd = open("lf2", 0);
00734     if(fd < 0){
00735         printf(1, "open lf2 failed\n");
00736         exit();
00737     }
00738     if(read(fd, buf, sizeof(buf)) != 5){
00739         printf(1, "read lf2 failed\n");
00740         exit();
00741     }
00742     close(fd);
00743
00744     if(link("lf2", "lf2") >= 0){
00745         printf(1, "link lf2 lf2 succeeded! oops\n");
00746         exit();
00747     }
00748
00749     unlink("lf2");
00750     if(link("lf2", "lf1") >= 0){
00751         printf(1, "link non-existant succeeded! oops\n");
00752         exit();
00753     }
00754
00755     if(link(".", "lf1") >= 0){
00756         printf(1, "link . lf1 succeeded! oops\n");
00757         exit();
00758     }
00759
00760     printf(1, "linktest ok\n");
00761 }
00762
00763 // test concurrent create/link/unlink of the same file
00764 void
00765 concreate(void)
00766 {
00767     char file[3];
00768     int i, pid, n, fd;
00769     char fa[40];
00770     struct {
00771         ushort inum;
00772         char name[14];
00773     } de;
00774
00775     printf(1, "concreate test\n");
00776     file[0] = 'C';
00777     file[2] = '\0';
00778     for(i = 0; i < 40; i++){

```



```

00779     file[1] = '0' + i;
00780     unlink(file);
00781     pid = fork();
00782     if(pid && (i % 3) == 1){
00783         link("C0", file);
00784     } else if(pid == 0 && (i % 5) == 1){
00785         link("C0", file);
00786     } else {
00787         fd = open(file, O_CREATE | O_RDWR);
00788         if(fd < 0){
00789             printf(1, "concreate create %s failed\n", file);
00790             exit();
00791         }
00792         close(fd);
00793     }
00794     if(pid == 0)
00795         exit();
00796     else
00797         wait();
00798 }
00799
00800 memset(fa, 0, sizeof(fa));
00801 fd = open(".", 0);
00802 n = 0;
00803 while(read(fd, &de, sizeof(de)) > 0){
00804     if(de.inum == 0)
00805         continue;
00806     if(de.name[0] == 'C' && de.name[2] == '\0'){
00807         i = de.name[1] - '0';
00808         if(i < 0 || i >= sizeof(fa)){
00809             printf(1, "concreate weird file %s\n", de.name);
00810             exit();
00811         }
00812         if(fa[i]){
00813             printf(1, "concreate duplicate file %s\n", de.name);
00814             exit();
00815         }
00816         fa[i] = 1;
00817         n++;
00818     }
00819 }
00820 close(fd);
00821
00822 if(n != 40){
00823     printf(1, "concreate not enough files in directory listing\n");
00824     exit();
00825 }
00826
00827 for(i = 0; i < 40; i++){
00828     file[1] = '0' + i;
00829     pid = fork();
00830     if(pid < 0){
00831         printf(1, "fork failed\n");
00832         exit();
00833     }
00834     if(((i % 3) == 0 && pid == 0) ||
00835        ((i % 3) == 1 && pid != 0)){
00836         close(open(file, 0));
00837         close(open(file, 0));
00838         close(open(file, 0));
00839         close(open(file, 0));
00840     } else {
00841         unlink(file);
00842         unlink(file);
00843         unlink(file);
00844         unlink(file);
00845     }
00846     if(pid == 0)
00847         exit();
00848     else
00849         wait();
00850 }
00851
00852 printf(1, "concreate ok\n");
00853 }
00854
00855 // another concurrent link/unlink/create test,
00856 // to look for deadlocks.
00857 void
00858 linkunlink()
00859 {
00860     int pid, i;
00861
00862     printf(1, "linkunlink test\n");
00863
00864     unlink("x");
00865     pid = fork();

```

```

00866     if(pid < 0){
00867         printf(1, "fork failed\n");
00868         exit();
00869     }
00870
00871     unsigned int x = (pid ? 1 : 97);
00872     for(i = 0; i < 100; i++){
00873         x = x * 1103515245 + 12345;
00874         if((x % 3) == 0){
00875             close(open("x", O_RDWR | O_CREATE));
00876         } else if((x % 3) == 1){
00877             link("cat", "x");
00878         } else {
00879             unlink("x");
00880         }
00881     }
00882
00883     if(pid)
00884         wait();
00885     else
00886         exit();
00887
00888     printf(1, "linkunlink ok\n");
00889 }
00890
00891 // directory that uses indirect blocks
00892 void
00893 bigdir(void)
00894 {
00895     int i, fd;
00896     char name[10];
00897
00898     printf(1, "bigdir test\n");
00899     unlink("bd");
00900
00901     fd = open("bd", O_CREATE);
00902     if(fd < 0){
00903         printf(1, "bigdir create failed\n");
00904         exit();
00905     }
00906     close(fd);
00907
00908     for(i = 0; i < 500; i++){
00909         name[0] = 'x';
00910         name[1] = '0' + (i / 64);
00911         name[2] = '0' + (i % 64);
00912         name[3] = '\0';
00913         if(link("bd", name) != 0){
00914             printf(1, "bigdir link failed\n");
00915             exit();
00916         }
00917     }
00918
00919     unlink("bd");
00920     for(i = 0; i < 500; i++){
00921         name[0] = 'x';
00922         name[1] = '0' + (i / 64);
00923         name[2] = '0' + (i % 64);
00924         name[3] = '\0';
00925         if(unlink(name) != 0){
00926             printf(1, "bigdir unlink failed");
00927             exit();
00928         }
00929     }
00930
00931     printf(1, "bigdir ok\n");
00932 }
00933
00934 void
00935 subdir(void)
00936 {
00937     int fd, cc;
00938
00939     printf(1, "subdir test\n");
00940
00941     unlink("ff");
00942     if(mkdir("dd") != 0){
00943         printf(1, "subdir mkdir dd failed\n");
00944         exit();
00945     }
00946
00947     fd = open("dd/ff", O_CREATE | O_RDWR);
00948     if(fd < 0){
00949         printf(1, "create dd/ff failed\n");
00950         exit();
00951     }
00952     write(fd, "ff", 2);

```

```
00953     close(fd);
00954
00955     if(unlink("dd") >= 0){
00956         printf(1, "unlink dd (non-empty dir) succeeded!\n");
00957         exit();
00958     }
00959
00960     if(mkdir("/dd/dd") != 0){
00961         printf(1, "subdir mkdir dd/dd failed\n");
00962         exit();
00963     }
00964
00965     fd = open("dd/dd/ff", O_CREATE | O_RDWR);
00966     if(fd < 0){
00967         printf(1, "create dd/dd/ff failed\n");
00968         exit();
00969     }
00970     write(fd, "FF", 2);
00971     close(fd);
00972
00973     fd = open("dd/dd/../ff", 0);
00974     if(fd < 0){
00975         printf(1, "open dd/dd/../ff failed\n");
00976         exit();
00977     }
00978     cc = read(fd, buf, sizeof(buf));
00979     if(cc != 2 || buf[0] != 'f'){
00980         printf(1, "dd/dd/../ff wrong content\n");
00981         exit();
00982     }
00983     close(fd);
00984
00985     if(link("dd/dd/ff", "dd/dd/ffff") != 0){
00986         printf(1, "link dd/dd/ff dd/dd/ffff failed\n");
00987         exit();
00988     }
00989
00990     if(unlink("dd/dd/ff") != 0){
00991         printf(1, "unlink dd/dd/ff failed\n");
00992         exit();
00993     }
00994     if(open("dd/dd/ff", O_RDONLY) >= 0){
00995         printf(1, "open (unlinked) dd/dd/ff succeeded\n");
00996         exit();
00997     }
00998
00999     if(chdir("dd") != 0){
01000         printf(1, "chdir dd failed\n");
01001         exit();
01002     }
01003     if(chdir("dd/../../dd") != 0){
01004         printf(1, "chdir dd/../../dd failed\n");
01005         exit();
01006     }
01007     if(chdir("dd/../../..") != 0){
01008         printf(1, "chdir dd/../../.. failed\n");
01009         exit();
01010     }
01011     if(chdir("../..") != 0){
01012         printf(1, "chdir ../.. failed\n");
01013         exit();
01014     }
01015
01016     fd = open("dd/dd/ffff", 0);
01017     if(fd < 0){
01018         printf(1, "open dd/dd/ffff failed\n");
01019         exit();
01020     }
01021     if(read(fd, buf, sizeof(buf)) != 2){
01022         printf(1, "read dd/dd/ffff wrong len\n");
01023         exit();
01024     }
01025     close(fd);
01026
01027     if(open("dd/dd/ff", O_RDONLY) >= 0){
01028         printf(1, "open (unlinked) dd/dd/ff succeeded!\n");
01029         exit();
01030     }
01031
01032     if(open("dd/ff/ff", O_CREATE|O_RDWR) >= 0){
01033         printf(1, "create dd/ff/ff succeeded!\n");
01034         exit();
01035     }
01036     if(open("dd/xx/ff", O_CREATE|O_RDWR) >= 0){
01037         printf(1, "create dd/xx/ff succeeded!\n");
01038         exit();
01039     }
```

```

01040     if(open("dd", O_CREATE) >= 0){
01041         printf(1, "create dd succeeded!\n");
01042         exit();
01043     }
01044     if(open("dd", O_RDWR) >= 0){
01045         printf(1, "open dd rdwr succeeded!\n");
01046         exit();
01047     }
01048     if(open("dd", O_WRONLY) >= 0){
01049         printf(1, "open dd wronly succeeded!\n");
01050         exit();
01051     }
01052     if(link("dd/ff/ff", "dd/dd/xx") == 0){
01053         printf(1, "link dd/ff/ff dd/dd/xx succeeded!\n");
01054         exit();
01055     }
01056     if(link("dd/xx/ff", "dd/dd/xx") == 0){
01057         printf(1, "link dd/xx/ff dd/dd/xx succeeded!\n");
01058         exit();
01059     }
01060     if(link("dd/ff", "dd/dd/ffff") == 0){
01061         printf(1, "link dd/ff dd/dd/ffff succeeded!\n");
01062         exit();
01063     }
01064     if(mkdir("dd/ff/ff") == 0){
01065         printf(1, "mkdir dd/ff/ff succeeded!\n");
01066         exit();
01067     }
01068     if(mkdir("dd/xx/ff") == 0){
01069         printf(1, "mkdir dd/xx/ff succeeded!\n");
01070         exit();
01071     }
01072     if(mkdir("dd/dd/ffff") == 0){
01073         printf(1, "mkdir dd/dd/ffff succeeded!\n");
01074         exit();
01075     }
01076     if(unlink("dd/xx/ff") == 0){
01077         printf(1, "unlink dd/xx/ff succeeded!\n");
01078         exit();
01079     }
01080     if(unlink("dd/ff/ff") == 0){
01081         printf(1, "unlink dd/ff/ff succeeded!\n");
01082         exit();
01083     }
01084     if(chdir("dd/ff") == 0){
01085         printf(1, "chdir dd/ff succeeded!\n");
01086         exit();
01087     }
01088     if(chdir("dd/xx") == 0){
01089         printf(1, "chdir dd/xx succeeded!\n");
01090         exit();
01091     }
01092 }
01093 if(unlink("dd/dd/ffff") != 0){
01094     printf(1, "unlink dd/dd/ff failed\n");
01095     exit();
01096 }
01097 if(unlink("dd/ff") != 0){
01098     printf(1, "unlink dd/ff failed\n");
01099     exit();
01100 }
01101 if(unlink("dd") == 0){
01102     printf(1, "unlink non-empty dd succeeded!\n");
01103     exit();
01104 }
01105 if(unlink("dd/dd") < 0){
01106     printf(1, "unlink dd/dd failed\n");
01107 }
01108 }
01109 if(unlink("dd") < 0){
01110     printf(1, "unlink dd failed\n");
01111     exit();
01112 }
01113 }
01114 printf(1, "subdir ok\n");
01115 }
01116
01117 // test writes that are larger than the log.
01118 void
01119 bigwrite(void)
01120 {
01121     int fd, sz;
01122
01123     printf(1, "bigwrite test\n");
01124
01125     unlink("bigwrite");
01126     for(sz = 499; sz < 12*512; sz += 471){

```

```
01127     fd = open("bigwrite", O_CREATE | O_RDWR);
01128     if(fd < 0){
01129         printf(1, "cannot create bigwrite\n");
01130         exit();
01131     }
01132     int i;
01133     for(i = 0; i < 2; i++){
01134         int cc = write(fd, buf, sz);
01135         if(cc != sz){
01136             printf(1, "write(%d) ret %d\n", sz, cc);
01137             exit();
01138         }
01139     }
01140     close(fd);
01141     unlink("bigwrite");
01142 }
01143
01144 printf(1, "bigwrite ok\n");
01145 }
01146
01147 void
01148 bigfile(void)
01149 {
01150     int fd, i, total, cc;
01151
01152     printf(1, "bigfile test\n");
01153
01154     unlink("bigfile");
01155     fd = open("bigfile", O_CREATE | O_RDWR);
01156     if(fd < 0){
01157         printf(1, "cannot create bigfile");
01158         exit();
01159     }
01160     for(i = 0; i < 20; i++){
01161         memset(buf, i, 600);
01162         if(write(fd, buf, 600) != 600){
01163             printf(1, "write bigfile failed\n");
01164             exit();
01165         }
01166     }
01167     close(fd);
01168
01169     fd = open("bigfile", 0);
01170     if(fd < 0){
01171         printf(1, "cannot open bigfile\n");
01172         exit();
01173     }
01174     total = 0;
01175     for(i = 0; ; i++){
01176         cc = read(fd, buf, 300);
01177         if(cc < 0){
01178             printf(1, "read bigfile failed\n");
01179             exit();
01180         }
01181         if(cc == 0)
01182             break;
01183         if(cc != 300){
01184             printf(1, "short read bigfile\n");
01185             exit();
01186         }
01187         if(buf[0] != i/2 || buf[299] != i/2){
01188             printf(1, "read bigfile wrong data\n");
01189             exit();
01190         }
01191         total += cc;
01192     }
01193     close(fd);
01194     if(total != 20*600){
01195         printf(1, "read bigfile wrong total\n");
01196         exit();
01197     }
01198     unlink("bigfile");
01199
01200     printf(1, "bigfile test ok\n");
01201 }
01202
01203 void
01204 fourteen(void)
01205 {
01206     int fd;
01207
01208     // DIRSIZ is 14.
01209     printf(1, "fourteen test\n");
01210
01211     if(mkdir("12345678901234") != 0){
01212         printf(1, "mkdir 12345678901234 failed\n");
01213         exit();
01214     }
```

```

01214     }
01215     if(mkdir("12345678901234/123456789012345") != 0){
01216         printf(1, "mkdir 12345678901234/123456789012345 failed\n");
01217         exit();
01218     }
01219     fd = open("123456789012345/123456789012345/123456789012345", O_CREATE);
01220     if(fd < 0){
01221         printf(1, "create 123456789012345/123456789012345/123456789012345 failed\n");
01222         exit();
01223     }
01224     close(fd);
01225     fd = open("12345678901234/12345678901234/12345678901234", 0);
01226     if(fd < 0){
01227         printf(1, "open 12345678901234/12345678901234/12345678901234 failed\n");
01228         exit();
01229     }
01230     close(fd);
01231
01232     if(mkdir("12345678901234/12345678901234") == 0){
01233         printf(1, "mkdir 12345678901234/12345678901234 succeeded!\n");
01234         exit();
01235     }
01236     if(mkdir("123456789012345/12345678901234") == 0){
01237         printf(1, "mkdir 12345678901234/123456789012345 succeeded!\n");
01238         exit();
01239     }
01240
01241     printf(1, "fourteen ok\n");
01242 }
01243
01244 void
01245 rmdot(void)
01246 {
01247     printf(1, "rmdot test\n");
01248     if(mkdir("dots") != 0){
01249         printf(1, "mkdir dots failed\n");
01250         exit();
01251     }
01252     if(chdir("dots") != 0){
01253         printf(1, "chdir dots failed\n");
01254         exit();
01255     }
01256     if(unlink(".") == 0){
01257         printf(1, "rm . worked!\n");
01258         exit();
01259     }
01260     if(unlink("..") == 0){
01261         printf(1, "rm .. worked!\n");
01262         exit();
01263     }
01264     if(chdir("/") != 0){
01265         printf(1, "chdir / failed\n");
01266         exit();
01267     }
01268     if(unlink("dots/.") == 0){
01269         printf(1, "unlink dots/. worked!\n");
01270         exit();
01271     }
01272     if(unlink("dots/..") == 0){
01273         printf(1, "unlink dots/.. worked!\n");
01274         exit();
01275     }
01276     if(unlink("dots") != 0){
01277         printf(1, "unlink dots failed!\n");
01278         exit();
01279     }
01280     printf(1, "rmdot ok\n");
01281 }
01282
01283 void
01284 dirfile(void)
01285 {
01286     int fd;
01287
01288     printf(1, "dir vs file\n");
01289
01290     fd = open("dirfile", O_CREATE);
01291     if(fd < 0){
01292         printf(1, "create dirfile failed\n");
01293         exit();
01294     }
01295     close(fd);
01296     if(chdir("dirfile") == 0){
01297         printf(1, "chdir dirfile succeeded!\n");
01298         exit();
01299     }
01300     fd = open("dirfile/xx", 0);

```

```

01301     if(fd >= 0){
01302         printf(1, "create dirfile/xx succeeded!\n");
01303         exit();
01304     }
01305     fd = open("dirfile/xx", O_CREATE);
01306     if(fd >= 0){
01307         printf(1, "create dirfile/xx succeeded!\n");
01308         exit();
01309     }
01310     if(mkdir("dirfile/xx") == 0){
01311         printf(1, "mkdir dirfile/xx succeeded!\n");
01312         exit();
01313     }
01314     if(unlink("dirfile/xx") == 0){
01315         printf(1, "unlink dirfile/xx succeeded!\n");
01316         exit();
01317     }
01318     if(link("README", "dirfile/xx") == 0){
01319         printf(1, "link to dirfile/xx succeeded!\n");
01320         exit();
01321     }
01322     if(unlink("dirfile") != 0){
01323         printf(1, "unlink dirfile failed!\n");
01324         exit();
01325     }
01326
01327     fd = open(".", O_RDWR);
01328     if(fd >= 0){
01329         printf(1, "open . for writing succeeded!\n");
01330         exit();
01331     }
01332     fd = open(".", 0);
01333     if(write(fd, "x", 1) > 0){
01334         printf(1, "write . succeeded!\n");
01335         exit();
01336     }
01337     close(fd);
01338
01339     printf(1, "dir vs file OK\n");
01340 }
01341
01342 // test that iput() is called at the end of _namei()
01343 void
01344 iref(void)
01345 {
01346     int i, fd;
01347
01348     printf(1, "empty file name\n");
01349
01350     // the 50 is NINODE
01351     for(i = 0; i < 50 + 1; i++){
01352         if(mkdir("irefd") != 0){
01353             printf(1, "mkdir irefd failed\n");
01354             exit();
01355         }
01356         if(chdir("irefd") != 0){
01357             printf(1, "chdir irefd failed\n");
01358             exit();
01359         }
01360
01361         mkdir("");
01362         link("README", "");
01363         fd = open("", O_CREATE);
01364         if(fd >= 0)
01365             close(fd);
01366         fd = open("xx", O_CREATE);
01367         if(fd >= 0)
01368             close(fd);
01369         unlink("xx");
01370     }
01371
01372     chdir("/");
01373     printf(1, "empty file name OK\n");
01374 }
01375
01376 // test that fork fails gracefully
01377 // the forktest binary also does this, but it runs out of proc entries first.
01378 // inside the bigger usertests binary, we run out of memory first.
01379 void
01380 forktest(void)
01381 {
01382     int n, pid;
01383
01384     printf(1, "fork test\n");
01385
01386     for(n=0; n<1000; n++){
01387         pid = fork();

```

```

01388     if(pid < 0)
01389         break;
01390     if(pid == 0)
01391         exit();
01392 }
01393
01394 if(n == 1000){
01395     printf(1, "fork claimed to work 1000 times!\n");
01396     exit();
01397 }
01398
01399 for(; n > 0; n--){
01400     if(wait() < 0){
01401         printf(1, "wait stopped early\n");
01402         exit();
01403     }
01404 }
01405
01406 if(wait() != -1){
01407     printf(1, "wait got too many\n");
01408     exit();
01409 }
01410
01411 printf(1, "fork test OK\n");
01412 }
01413
01414 void
01415 sbrktest(void)
01416 {
01417     int fds[2], pid, pids[10], ppid;
01418     char *a, *b, *c, *lastaddr, *oldbrk, *p, scratch;
01419     uint amt;
01420
01421     printf(stdout, "sbrk test\n");
01422     oldbrk = sbrk(0);
01423
01424     // can one sbrk() less than a page?
01425     a = sbrk(0);
01426     int i;
01427     for(i = 0; i < 5000; i++){
01428         b = sbrk(1);
01429         if(b != a){
01430             printf(stdout, "sbrk test failed %d %x %x\n", i, a, b);
01431             exit();
01432         }
01433         *b = 1;
01434         a = b + 1;
01435     }
01436     pid = fork();
01437     if(pid < 0){
01438         printf(stdout, "sbrk test fork failed\n");
01439         exit();
01440     }
01441     c = sbrk(1);
01442     c = sbrk(1);
01443     if(c != a + 1){
01444         printf(stdout, "sbrk test failed post-fork\n");
01445         exit();
01446     }
01447     if(pid == 0)
01448         exit();
01449     wait();
01450
01451     // can one grow address space to something big?
01452     #define BIG (100*1024*1024)
01453     a = sbrk(0);
01454     amt = (BIG) - (uint)a;
01455     p = sbrk(amt);
01456     if (p != a) {
01457         printf(stdout, "sbrk test failed to grow big address space; enough phys mem?\n");
01458         exit();
01459     }
01460     lastaddr = (char*) (BIG-1);
01461     *lastaddr = 99;
01462
01463     // can one de-allocate?
01464     a = sbrk(0);
01465     c = sbrk(-4096);
01466     if(c == (char*)0xffffffff){
01467         printf(stdout, "sbrk could not deallocate\n");
01468         exit();
01469     }
01470     c = sbrk(0);
01471     if(c != a - 4096){
01472         printf(stdout, "sbrk deallocation produced wrong address, a %x c %x\n", a, c);
01473         exit();
01474     }

```



```

01475
01476 // can one re-allocate that page?
01477 a = sbrk(0);
01478 c = sbrk(4096);
01479 if(c != a || sbrk(0) != a + 4096){
01480     printf(stdout, "sbrk re-allocation failed, a %x c %x\n", a, c);
01481     exit();
01482 }
01483 if(*lastaddr == 99){
01484     // should be zero
01485     printf(stdout, "sbrk de-allocation didn't really deallocate\n");
01486     exit();
01487 }
01488
01489 a = sbrk(0);
01490 c = sbrk(-(sbrk(0) - oldbrk));
01491 if(c != a){
01492     printf(stdout, "sbrk downsize failed, a %x c %x\n", a, c);
01493     exit();
01494 }
01495
01496 // can we read the kernel's memory?
01497 for(a = (char*)(KERNBASE); a < (char*)(KERNBASE+2000000); a += 50000){
01498     ppid = getpid();
01499     pid = fork();
01500     if(pid < 0){
01501         printf(stdout, "fork failed\n");
01502         exit();
01503     }
01504     if(pid == 0){
01505         printf(stdout, "oops could read %x = %x\n", a, *a);
01506         kill(ppid);
01507         exit();
01508     }
01509     wait();
01510 }
01511
01512 // if we run the system out of memory, does it clean up the last
01513 // failed allocation?
01514 if(pipe(fds) != 0){
01515     printf(1, "pipe() failed\n");
01516     exit();
01517 }
01518 for(i = 0; i < sizeof(pids)/sizeof(pids[0]); i++){
01519     if((pids[i] = fork()) == 0){
01520         // allocate a lot of memory
01521         sbrk(BIG - (uint)sbrk(0));
01522         write(fds[1], "x", 1);
01523         // sit around until killed
01524         for(;;) sleep(1000);
01525     }
01526     if(pids[i] != -1)
01527         read(fds[0], &scratch, 1);
01528 }
01529 // if those failed allocations freed up the pages they did allocate,
01530 // we'll be able to allocate here
01531 c = sbrk(4096);
01532 for(i = 0; i < sizeof(pids)/sizeof(pids[0]); i++){
01533     if(pids[i] == -1)
01534         continue;
01535     kill(pids[i]);
01536     wait();
01537 }
01538 if(c == (char*)0xffffffff){
01539     printf(stdout, "failed sbrk leaked memory\n");
01540     exit();
01541 }
01542
01543 if(sbrk(0) > oldbrk)
01544     sbrk(-(sbrk(0) - oldbrk));
01545
01546 printf(stdout, "sbrk test OK\n");
01547 }
01548
01549 void
01550 validateint(int *p)
01551 {
01552     int res;
01553     asm("mov %esp, %%ebx\n\t"
01554         "mov %3, %%esp\n\t"
01555         "int %2\n\t"
01556         "mov %%ebx, %%esp" :
01557         "=a" (res) :
01558         "a" (SYS_sleep), "n" (T_SYSCALL), "c" (p) :
01559         "ebx");
01560 }
01561

```

```

01562 void
01563 validate_test(void)
01564 {
01565     int hi, pid;
01566     uint p;
01567
01568     printf(stdout, "validate test\n");
01569     hi = 1100*1024;
01570
01571     for(p = 0; p <= (uint)hi; p += 4096){
01572         if((pid = fork()) == 0){
01573             // try to crash the kernel by passing in a badly placed integer
01574             validateint((int*)p);
01575             exit();
01576         }
01577         sleep(0);
01578         sleep(0);
01579         kill(pid);
01580         wait();
01581
01582         // try to crash the kernel by passing in a bad string pointer
01583         if(link("nosuchfile", (char*)p) != -1){
01584             printf(stdout, "link should not succeed\n");
01585             exit();
01586         }
01587     }
01588
01589     printf(stdout, "validate ok\n");
01590 }
01591
01592 // does uninitialized data start out zero?
01593 char uninit[10000];
01594 void
01595 bsstest(void)
01596 {
01597     int i;
01598
01599     printf(stdout, "bss test\n");
01600     for(i = 0; i < sizeof(uninit); i++){
01601         if(uninit[i] != '\0'){
01602             printf(stdout, "bss test failed\n");
01603             exit();
01604         }
01605     }
01606     printf(stdout, "bss test ok\n");
01607 }
01608
01609 // does exec return an error if the arguments
01610 // are larger than a page? or does it write
01611 // below the stack and wreck the instructions/data?
01612 void
01613 bigargtest(void)
01614 {
01615     int pid, fd;
01616
01617     unlink("bigarg-ok");
01618     pid = fork();
01619     if(pid == 0){
01620         static char *args[MAXARG];
01621         int i;
01622         for(i = 0; i < MAXARG-1; i++){
01623             args[i] = "bigargs test: failed\n";
01624
01625             ";
01626             args[MAXARG-1] = 0;
01627             printf(stdout, "bigarg test\n");
01628             exec("echo", args);
01629             printf(stdout, "bigarg test ok\n");
01630             fd = open("bigarg-ok", O_CREATE);
01631             close(fd);
01632             exit();
01633         } else if(pid < 0){
01634             printf(stdout, "bigargtest: fork failed\n");
01635             exit();
01636         }
01637         wait();
01638         fd = open("bigarg-ok", 0);
01639         if(fd < 0){
01640             printf(stdout, "bigarg test failed!\n");
01641             exit();
01642         }
01643         close(fd);
01644         unlink("bigarg-ok");
01645     }
01646 }
01647
01648 // what happens when the file system runs out of blocks?
01649 // answer: ballocc panics, so this test is not useful.

```

```

01647 void
01648 fsfull()
01649 {
01650     int nfiles;
01651     int fsblocks = 0;
01652
01653     printf(1, "fsfull test\n");
01654
01655     for(nfiles = 0; ; nfiles++){
01656         char name[64];
01657         name[0] = 'f';
01658         name[1] = '0' + nfiles / 1000;
01659         name[2] = '0' + (nfiles % 1000) / 100;
01660         name[3] = '0' + (nfiles % 100) / 10;
01661         name[4] = '0' + (nfiles % 10);
01662         name[5] = '\0';
01663         printf(1, "writing %s\n", name);
01664         int fd = open(name, O_CREATE|O_RDWR);
01665         if(fd < 0){
01666             printf(1, "open %s failed\n", name);
01667             break;
01668         }
01669         int total = 0;
01670         while(1){
01671             int cc = write(fd, buf, 512);
01672             if(cc < 512)
01673                 break;
01674             total += cc;
01675             fsblocks++;
01676         }
01677         printf(1, "wrote %d bytes\n", total);
01678         close(fd);
01679         if(total == 0)
01680             break;
01681     }
01682
01683     while(nfiles >= 0){
01684         char name[64];
01685         name[0] = 'f';
01686         name[1] = '0' + nfiles / 1000;
01687         name[2] = '0' + (nfiles % 1000) / 100;
01688         name[3] = '0' + (nfiles % 100) / 10;
01689         name[4] = '0' + (nfiles % 10);
01690         name[5] = '\0';
01691         unlink(name);
01692         nfiles--;
01693     }
01694
01695     printf(1, "fsfull test finished\n");
01696 }
01697
01698 void
01699 uio()
01700 {
01701     #define RTC_ADDR 0x70
01702     #define RTC_DATA 0x71
01703
01704     ushort port = 0;
01705     uchar val = 0;
01706     int pid;
01707
01708     printf(1, "uio test\n");
01709     pid = fork();
01710     if(pid == 0){
01711         port = RTC_ADDR;
01712         val = 0x09; /* year */
01713         /* http://wiki.osdev.org/Inline_Assembly/Examples */
01714         asm volatile("outb %0,%1::" : "a" (val), "d" (port));
01715         port = RTC_DATA;
01716         asm volatile("inb %1,%0 : " : "a" (val) : "d" (port));
01717         printf(1, "uio: uio succeeded; test FAILED\n");
01718         exit();
01719     } else if(pid < 0){
01720         printf(1, "fork failed\n");
01721         exit();
01722     }
01723     wait();
01724     printf(1, "uio test done\n");
01725 }
01726
01727 void argptest()
01728 {
01729     int fd;
01730     fd = open("init", O_RDONLY);
01731     if (fd < 0) {
01732         printf(2, "open failed\n");
01733         exit();
01734     }

```

```

01734     }
01735     read(fd, sbrk(0) - 1, -1);
01736     close(fd);
01737     printf(1, "arg test passed\n");
01738 }
01739
01740 unsigned long randstate = 1;
01741 unsigned int
01742 rand()
01743 {
01744     randstate = randstate * 1664525 + 1013904223;
01745     return randstate;
01746 }
01747
01748 int
01749 main(int argc, char *argv[])
01750 {
01751     printf(1, "usertests starting\n");
01752
01753     if(open("usertests.ran", 0) >= 0){
01754         printf(1, "already ran user tests -- rebuild fs.img\n");
01755         exit();
01756     }
01757     close(open("usertests.ran", O_CREATE));
01758
01759     argptest();
01760     createdelete();
01761     linkunlink();
01762     concreate();
01763     fourfiles();
01764     sharedfd();
01765
01766     bigargtest();
01767     bigwrite();
01768     bigargtest();
01769     bsstest();
01770     sbrktest();
01771     validatetest();
01772
01773     opentest();
01774     writetest();
01775     writetest1();
01776     createtest();
01777
01778     openiputtest();
01779     exitiputtest();
01780     iputtest();
01781
01782     mem();
01783     pipel();
01784     preempt();
01785     exitwait();
01786
01787     rmdot();
01788     fourteen();
01789     bigfile();
01790     subdir();
01791     linktest();
01792     unlinkread();
01793     dirfile();
01794     iref();
01795     forktest();
01796     bigdir(); // slow
01797
01798     uio();
01799
01800     exectest();
01801
01802     exit();
01803 }

```

5.242 usertests.d File Reference

5.243 usertests.d

[Go to the documentation of this file.](#)

```

00001 usertests.o: usertests.c /usr/include/stdc-predef.h param.h types.h \
00002 stat.h user.h fs.h fcntl.h syscall.h traps.h memlayout.h

```

5.244 vm.c File Reference

```
#include "param.h"
#include "types.h"
#include "defs.h"
#include "x86.h"
#include "memlayout.h"
#include "mmu.h"
#include "proc.h"
#include "elf.h"
```

Classes

- struct [kmap](#)

Functions

- int [allocuvm](#) ([pde_t](#) *pgdir, [uint](#) oldsz, [uint](#) newsz)
- void [clearpteu](#) ([pde_t](#) *pgdir, char *uva)
- int [copyout](#) ([pde_t](#) *pgdir, [uint](#) va, void *p, [uint](#) len)
- [pde_t](#) * [copyuvm](#) ([pde_t](#) *pgdir, [uint](#) sz)
- int [deallocuvm](#) ([pde_t](#) *pgdir, [uint](#) oldsz, [uint](#) newsz)
- void [freevm](#) ([pde_t](#) *pgdir)
- void [inituvm](#) ([pde_t](#) *pgdir, char *init, [uint](#) sz)
- void [kvmalloc](#) (void)
- int [loaduvm](#) ([pde_t](#) *pgdir, char *addr, struct [inode](#) *ip, [uint](#) offset, [uint](#) sz)
- static int [mappages](#) ([pde_t](#) *pgdir, void *va, [uint](#) size, [uint](#) pa, int perm)
- void [seginit](#) (void)
- [pde_t](#) * [setupkvm](#) (void)
- void [switchkvm](#) (void)
- void [switchuvm](#) (struct [proc](#) *p)
- char * [uva2ka](#) ([pde_t](#) *pgdir, char *uva)
- static [pte_t](#) * [walkpgdir](#) ([pde_t](#) *pgdir, const void *va, int alloc)

Variables

- char [data](#) []
- static struct [kmap](#) [kmap](#) []
- [pde_t](#) * [kpgdir](#)

5.244.1 Function Documentation

5.244.1.1 allocvm()

```
int allocvm (
    pde_t * pgdir,
    uint oldsz,
    uint newsz )
```

Definition at line 222 of file [vm.c](#).

```
00223 {
00224     char *mem;
00225     uint a;
00226
00227     if(newsz >= KERNBASE)
00228         return 0;
00229     if(newsz < oldsz)
00230         return oldsz;
00231
00232     a = PGROUNDUP(oldsz);
00233     for(; a < newsz; a += PGSIZE){
00234         mem = kalloc();
00235         if(mem == 0){
00236             cprintf("allocvm out of memory\n");
00237             deallocvm(pgdir, newsz, oldsz);
00238             return 0;
00239         }
00240         memset(mem, 0, PGSIZE);
00241         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
00242             cprintf("allocvm out of memory (2)\n");
00243             deallocvm(pgdir, newsz, oldsz);
00244             kfree(mem);
00245             return 0;
00246         }
00247     }
00248     return newsz;
00249 }
```

Referenced by [exec\(\)](#), and [growproc\(\)](#).

5.244.1.2 clearpteu()

```
void clearpteu (
    pde_t * pgdir,
    char * uva )
```

Definition at line 303 of file [vm.c](#).

```
00304 {
00305     pte_t *pte;
00306
00307     pte = walkpgdir(pgdir, uva, 0);
00308     if(pte == 0)
00309         panic("clearpteu");
00310     *pte &= ~PTE_U;
00311 }
```

Referenced by [exec\(\)](#).

5.244.1.3 copyout()

```
int copyout (
    pde_t * pgdir,
    uint va,
    void * p,
    uint len )
```

Definition at line 366 of file [vm.c](#).

```
00367 {
00368     char *buf, *pa0;
00369     uint n, va0;
00370
00371     buf = (char*)p;
00372     while(len > 0){
00373         va0 = (uint)PGROUNDDOWN(va);
00374         pa0 = uva2ka(pgdir, (char*)va0);
00375         if(pa0 == 0)
00376             return -1;
00377         n = PGSIZE - (va - va0);
00378         if(n > len)
00379             n = len;
00380         memmove(pa0 + (va - va0), buf, n);
00381         len -= n;
00382         buf += n;
00383         va = va0 + PGSIZE;
00384     }
00385     return 0;
00386 }
```

Referenced by [exec\(\)](#).

5.244.1.4 copyvm()

```
pde_t * copyvm (
    pde_t * pgdir,
    uint sz )
```

Definition at line 316 of file [vm.c](#).

```
00317 {
00318     pde_t *d;
00319     pte_t *pte;
00320     uint pa, i, flags;
00321     char *mem;
00322
00323     if((d = setupkvm()) == 0)
00324         return 0;
00325     for(i = 0; i < sz; i += PGSIZE){
00326         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
00327             panic("copyvm: pte should exist");
00328         if(!(*pte & PTE_P))
00329             panic("copyvm: page not present");
00330         pa = PTE_ADDR(*pte);
00331         flags = PTE_FLAGS(*pte);
00332         if((mem = kalloc()) == 0)
00333             goto bad;
00334         memmove(mem, (char*)P2V(pa), PGSIZE);
00335         if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
00336             kfree(mem);
00337             goto bad;
00338         }
00339     }
00340     return d;
00341
00342 bad:
00343     freevm(d);
00344     return 0;
00345 }
```

Referenced by [fork\(\)](#).

5.244.1.5 deallocvm()

```
int deallocvm (
    pde_t * pgdir,
    uint oldsz,
    uint newsz )
```

Definition at line 256 of file [vm.c](#).

```
00257 {
00258     pte_t *pte;
00259     uint a, pa;
00260
00261     if(newsz >= oldsz)
00262         return oldsz;
00263
00264     a = PGROUNDUP(newsz);
00265     for(; a < oldsz; a += PGSIZE){
00266         pte = walkpgdir(pgdir, (char*)a, 0);
00267         if(!pte)
00268             a = PGADDR(PDX(a) + 1, 0, 0) - PGSIZE;
00269         else if((*pte & PTE_P) != 0){
00270             pa = PTE_ADDR(*pte);
00271             if(pa == 0)
00272                 panic("kfree");
00273             char *v = P2V(pa);
00274             kfree(v);
00275             *pte = 0;
00276         }
00277     }
00278     return newsz;
00279 }
```

Referenced by [allocvm\(\)](#), [freevm\(\)](#), and [growproc\(\)](#).

5.244.1.6 freevm()

```
void freevm (
    pde_t * pgdir )
```

Definition at line 284 of file [vm.c](#).

```
00285 {
00286     uint i;
00287
00288     if(pgdir == 0)
00289         panic("freevm: no pgdir");
00290     deallocvm(pgdir, KERNBASE, 0);
00291     for(i = 0; i < NPENTRIES; i++){
00292         if(pgdir[i] & PTE_P){
00293             char *v = P2V(PTE_ADDR(pgdir[i]));
00294             kfree(v);
00295         }
00296     }
00297     kfree((char*)pgdir);
00298 }
```

Referenced by [copyvm\(\)](#), [exec\(\)](#), [setupkvm\(\)](#), and [wait\(\)](#).

5.244.1.7 inituvm()

```
void inituvm (
    pde_t * pgdir,
    char * init,
    uint sz )
```

Definition at line 183 of file [vm.c](#).

```
00184 {
00185     char *mem;
00186
00187     if(sz >= PGSIZE)
00188         panic("inituvm: more than a page");
00189     mem = kalloc();
00190     memset(mem, 0, PGSIZE);
00191     mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_W|PTE_U);
00192     memmove(mem, init, sz);
00193 }
```

Referenced by [userinit\(\)](#).

5.244.1.8 kvmalloc()

```
void kvmalloc (
    void )
```

Definition at line 141 of file [vm.c](#).

```
00142 {
00143     kpgdir = setupkvm();
00144     switchkvm();
00145 }
```

5.244.1.9 loaduvm()

```
int loaduvm (
    pde_t * pgdir,
    char * addr,
    struct inode * ip,
    uint offset,
    uint sz )
```

Definition at line 198 of file [vm.c](#).

```
00199 {
00200     uint i, pa, n;
00201     pte_t *pte;
00202
00203     if((uint) addr % PGSIZE != 0)
00204         panic("loaduvm: addr must be page aligned");
00205     for(i = 0; i < sz; i += PGSIZE){
00206         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
00207             panic("loaduvm: address should exist");
00208         pa = PTE_ADDR(*pte);
00209         if(sz - i < PGSIZE)
00210             n = sz - i;
00211         else
00212             n = PGSIZE;
00213         if(readi(ip, P2V(pa), offset+i, n) != n)
00214             return -1;
00215     }
00216     return 0;
00217 }
```

Referenced by [exec\(\)](#).

5.244.1.10 mappages()

```
static int mappages (
    pde_t * pgdir,
    void * va,
    uint size,
    uint pa,
    int perm ) [static]
```

Definition at line 61 of file [vm.c](#).

```
00062 {
00063     char *a, *last;
00064     pte_t *pte;
00065
00066     a = (char*)PGROUNDDOWN((uint)va);
00067     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
00068     for(;;){
00069         if((pte = walkpgdir(pgdir, a, 1)) == 0)
00070             return -1;
00071         if(*pte & PTE_P)
00072             panic("remap");
00073         *pte = pa | perm | PTE_P;
00074         if(a == last)
00075             break;
00076         a += PGSIZE;
00077         pa += PGSIZE;
00078     }
00079     return 0;
00080 }
```

Referenced by [allocuvn\(\)](#), [copyuvn\(\)](#), [inituvn\(\)](#), and [setupkvm\(\)](#).

5.244.1.11 seginit()

```
void seginit (
    void )
```

Definition at line 16 of file [vm.c](#).

```
00017 {
00018     struct cpu *c;
00019
00020     // Map "logical" addresses to virtual addresses using identity map.
00021     // Cannot share a CODE descriptor for both kernel and user
00022     // because it would have to have DPL_USR, but the CPU forbids
00023     // an interrupt from CPL=0 to DPL=3.
00024     c = &cpus[cuid()];
00025     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
00026     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
00027     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
00028     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
00029     lgdt(c->gdt, sizeof(c->gdt));
00030 }
```

Referenced by [mpenter\(\)](#).

5.244.1.12 setupkvm()

```
pde_t * setupkvm (
    void )
```

Definition at line 119 of file `vm.c`.

```
00120 {
00121     pde_t *pgdir;
00122     struct kmap *k;
00123
00124     if((pgdir = (pde_t*)kalloc()) == 0)
00125         return 0;
00126     memset(pgdir, 0, PGSIZE);
00127     if (P2V(PHYSTOP) > (void*)DEVSPACE)
00128         panic("PHYSTOP too high");
00129     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
00130         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
00131             (uint)k->phys_start, k->perm) < 0) {
00132             freevm(pgdir);
00133             return 0;
00134         }
00135     return pgdir;
00136 }
```

Referenced by `copyuvm()`, `exec()`, `kvmalloc()`, and `userinit()`.

5.244.1.13 switchkvm()

```
void switchkvm (
    void )
```

Definition at line 150 of file `vm.c`.

```
00151 {
00152     lcr3(V2P(kpgdir));    // switch to the kernel page table
00153 }
```

Referenced by `kvmalloc()`, `mpenter()`, and `scheduler()`.

5.244.1.14 switchuvm()

```
void switchuvm (
    struct proc * p )
```

Definition at line 157 of file `vm.c`.

```
00158 {
00159     if(p == 0)
00160         panic("switchuvm: no process");
00161     if(p->kstack == 0)
00162         panic("switchuvm: no kstack");
00163     if(p->pgdir == 0)
00164         panic("switchuvm: no pgdir");
00165
00166     pushcli();
00167     mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A, &mycpu()->ts,
00168         sizeof(mycpu()->ts)-1, 0);
00169     mycpu()->gdt[SEG_TSS].s = 0;
00170     mycpu()->ts.ss0 = SEG_KDATA << 3;
00171     mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;
00172     // setting IOPL=0 in eflags *and* iomb beyond the tss segment limit
00173     // forbids I/O instructions (e.g., inb and outb) from user space
00174     mycpu()->ts.iomb = (ushort) 0xFFFF;
00175     ltr(SEG_TSS << 3);
00176     lcr3(V2P(p->pgdir));    // switch to process's address space
00177     popcli();
00178 }
```

Referenced by `exec()`, `growproc()`, and `scheduler()`.

5.244.1.15 uva2ka()

```
char * uva2ka (
    pde_t * pgdir,
    char * uva )
```

Definition at line 350 of file [vm.c](#).

```
00351 {
00352     pte_t *pte;
00353
00354     pte = walkpgdir(pgdir, uva, 0);
00355     if((*pte & PTE_P) == 0)
00356         return 0;
00357     if((*pte & PTE_U) == 0)
00358         return 0;
00359     return (char*)P2V(PTE_ADDR(*pte));
00360 }
```

Referenced by [copyout\(\)](#).

5.244.1.16 walkpgdir()

```
static pte_t * walkpgdir (
    pde_t * pgdir,
    const void * va,
    int alloc ) [static]
```

Definition at line 36 of file [vm.c](#).

```
00037 {
00038     pde_t *pde;
00039     pte_t *pgtab;
00040
00041     pde = &pgdir[PDX(va)];
00042     if(*pde & PTE_P){
00043         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
00044     } else {
00045         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
00046             return 0;
00047         // Make sure all those PTE_P bits are zero.
00048         memset(pgtab, 0, PGSIZE);
00049         // The permissions here are overly generous, but they can
00050         // be further restricted by the permissions in the page table
00051         // entries, if necessary.
00052         *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
00053     }
00054     return &pgtab[PTX(va)];
00055 }
```

Referenced by [clearpte\(\)](#), [copyuvn\(\)](#), [deallocuvn\(\)](#), [loaduvn\(\)](#), [mappages\(\)](#), and [uva2ka\(\)](#).

5.244.2 Variable Documentation

5.244.2.1 data

```
char data[] [extern]
```

Referenced by [inb\(\)](#), [ioapicwrite\(\)](#), [kbdgetc\(\)](#), [main\(\)](#), [outb\(\)](#), [outw\(\)](#), [stosb\(\)](#), and [stosl\(\)](#).

5.244.2.2 kmap

```
struct kmap kmap[] [static]
```

Initial value:

```
= {
{ (void*)KERNBASE, 0,          EXTMEM,    PTE_W},
{ (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0},
{ (void*)data,     V2P(data),   PHYSTOP,  PTE_W},
{ (void*)DEVSPACE, DEVSPACE,    0,        PTE_W},
}
```

5.244.2.3 kpgdir

```
pde_t* kpgdir
```

Definition at line 11 of file `vm.c`.

Referenced by `kvmalloc()`, and `switchkvm()`.

5.245 vm.c

[Go to the documentation of this file.](#)

```
00001 #include "param.h"
00002 #include "types.h"
00003 #include "defs.h"
00004 #include "x86.h"
00005 #include "memlayout.h"
00006 #include "mmu.h"
00007 #include "proc.h"
00008 #include "elf.h"
00009
00010 extern char data[]; // defined by kernel.ld
00011 pde_t *kpgdir; // for use in scheduler()
00012
00013 // Set up CPU's kernel segment descriptors.
00014 // Run once on entry on each CPU.
00015 void
00016 seginit(void)
00017 {
00018     struct cpu *c;
00019
00020     // Map "logical" addresses to virtual addresses using identity map.
00021     // Cannot share a CODE descriptor for both kernel and user
00022     // because it would have to have DPL_USR, but the CPU forbids
00023     // an interrupt from CPL=0 to DPL=3.
00024     c = &cpus[cpuuid()];
00025     c->gdt[SEG_KCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, 0);
00026     c->gdt[SEG_KDATA] = SEG(STA_W, 0, 0xffffffff, 0);
00027     c->gdt[SEG_UCODE] = SEG(STA_X|STA_R, 0, 0xffffffff, DPL_USER);
00028     c->gdt[SEG_UDATA] = SEG(STA_W, 0, 0xffffffff, DPL_USER);
00029     lgdt(c->gdt, sizeof(c->gdt));
00030 }
00031
00032 // Return the address of the PTE in page table pgdir
00033 // that corresponds to virtual address va. If alloc!=0,
00034 // create any required page table pages.
00035 static pte_t *
00036 walkpgdir(pde_t *pgdir, const void *va, int alloc)
00037 {
00038     pde_t *pde;
00039     pte_t *pgtab;
00040
00041     pde = &pgdir[PDX(va)];
00042     if(*pde & PTE_P){
00043         pgtab = (pte_t*)P2V(PTE_ADDR(*pde));
00044     } else {
00045         if(!alloc || (pgtab = (pte_t*)kalloc()) == 0)
00046             return 0;
```

```

00047 // Make sure all those PTE_P bits are zero.
00048 memset(pgtab, 0, PGSIZE);
00049 // The permissions here are overly generous, but they can
00050 // be further restricted by the permissions in the page table
00051 // entries, if necessary.
00052 *pde = V2P(pgtab) | PTE_P | PTE_W | PTE_U;
00053 }
00054 return &pgtab[PTX(va)];
00055 }
00056
00057 // Create PTEs for virtual addresses starting at va that refer to
00058 // physical addresses starting at pa. va and size might not
00059 // be page-aligned.
00060 static int
00061 mappages(pde_t *pgdir, void *va, uint size, uint pa, int perm)
00062 {
00063     char *a, *last;
00064     pte_t *pte;
00065
00066     a = (char*)PGROUNDDOWN((uint)va);
00067     last = (char*)PGROUNDDOWN(((uint)va) + size - 1);
00068     for(;;){
00069         if((pte = walkpgdir(pgdir, a, 1)) == 0)
00070             return -1;
00071         if(*pte & PTE_P)
00072             panic("remap");
00073         *pte = pa | perm | PTE_P;
00074         if(a == last)
00075             break;
00076         a += PGSIZE;
00077         pa += PGSIZE;
00078     }
00079     return 0;
00080 }
00081
00082 // There is one page table per process, plus one that's used when
00083 // a CPU is not running any process (kpgdir). The kernel uses the
00084 // current process's page table during system calls and interrupts;
00085 // page protection bits prevent user code from using the kernel's
00086 // mappings.
00087 //
00088 // setupkvm() and exec() set up every page table like this:
00089 //
00090 // 0..KERNBASE: user memory (text+data+stack+heap), mapped to
00091 //                phys memory allocated by the kernel
00092 // KERNBASE..KERNBASE+EXTMEM: mapped to 0..EXTMEM (for I/O space)
00093 // KERNBASE+EXTMEM..data: mapped to EXTMEM..V2P(data)
00094 //                for the kernel's instructions and r/o data
00095 // data..KERNBASE+PHYSTOP: mapped to V2P(data)..PHYSTOP,
00096 //                rw data + free physical memory
00097 // 0xfe000000..0: mapped direct (devices such as ioapic)
00098 //
00099 // The kernel allocates physical memory for its heap and for user memory
00100 // between V2P(end) and the end of physical memory (PHYSTOP)
00101 // (directly addressable from end..P2V(PHYSTOP)).
00102 //
00103 // This table defines the kernel's mappings, which are present in
00104 // every process's page table.
00105 static struct kmap {
00106     void *virt;
00107     uint phys_start;
00108     uint phys_end;
00109     int perm;
00110 } kmap[] = {
00111     { (void*)KERNBASE, 0,          EXTMEM,    PTE_W}, // I/O space
00112     { (void*)KERNLINK, V2P(KERNLINK), V2P(data), 0}, // kern text+rodata
00113     { (void*)data,     V2P(data),    PHYSTOP,   PTE_W}, // kern data+memory
00114     { (void*)DEVSPACE, DEVSPACE,     0,        PTE_W}, // more devices
00115 };
00116
00117 // Set up kernel part of a page table.
00118 pde_t*
00119 setupkvm(void)
00120 {
00121     pde_t *pgdir;
00122     struct kmap *k;
00123
00124     if((pgdir = (pde_t*)kalloc()) == 0)
00125         return 0;
00126     memset(pgdir, 0, PGSIZE);
00127     if (P2V(PHYSTOP) > (void*)DEVSPACE)
00128         panic("PHYSTOP too high");
00129     for(k = kmap; k < &kmap[NELEM(kmap)]; k++)
00130         if(mappages(pgdir, k->virt, k->phys_end - k->phys_start,
00131                     (uint)k->phys_start, k->perm) < 0) {
00132             freevm(pgdir);
00133             return 0;
00134         }
00135 }

```

```

00134     }
00135     return pgdir;
00136 }
00137
00138 // Allocate one page table for the machine for the kernel address
00139 // space for scheduler processes.
00140 void
00141 kvmalloc(void)
00142 {
00143     kpgdir = setupkvm();
00144     switchkvm();
00145 }
00146
00147 // Switch h/w page table register to the kernel-only page table,
00148 // for when no process is running.
00149 void
00150 switchkvm(void)
00151 {
00152     lcr3(V2P(kpgdir)); // switch to the kernel page table
00153 }
00154
00155 // Switch TSS and h/w page table to correspond to process p.
00156 void
00157 switchvm(struct proc *p)
00158 {
00159     if(p == 0)
00160         panic("switchvm: no process");
00161     if(p->kstack == 0)
00162         panic("switchvm: no kstack");
00163     if(p->pgdir == 0)
00164         panic("switchvm: no pgdir");
00165
00166     pushcli();
00167     mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A, &mycpu()->ts,
00168                                   sizeof(mycpu()->ts)-1, 0);
00169     mycpu()->gdt[SEG_TSS].s = 0;
00170     mycpu()->ts.ss0 = SEG_KDATA << 3;
00171     mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;
00172     // setting IOPL=0 in eflags *and* iomb beyond the tss segment limit
00173     // forbids I/O instructions (e.g., inb and outb) from user space
00174     mycpu()->ts.iomb = (ushort) 0xFFFF;
00175     ltr(SEG_TSS << 3);
00176     lcr3(V2P(p->pgdir)); // switch to process's address space
00177     popcli();
00178 }
00179
00180 // Load the initcode into address 0 of pgdir.
00181 // sz must be less than a page.
00182 void
00183 initvm(pde_t *pgdir, char *init, uint sz)
00184 {
00185     char *mem;
00186
00187     if(sz >= PGSIZE)
00188         panic("initvm: more than a page");
00189     mem = kalloc();
00190     memset(mem, 0, PGSIZE);
00191     mappages(pgdir, 0, PGSIZE, V2P(mem), PTE_W|PTE_U);
00192     memmove(mem, init, sz);
00193 }
00194
00195 // Load a program segment into pgdir. addr must be page-aligned
00196 // and the pages from addr to addr+sz must already be mapped.
00197 int
00198 loadvm(pde_t *pgdir, char *addr, struct inode *ip, uint offset, uint sz)
00199 {
00200     uint i, pa, n;
00201     pte_t *pte;
00202
00203     if((uint) addr % PGSIZE != 0)
00204         panic("loadvm: addr must be page aligned");
00205     for(i = 0; i < sz; i += PGSIZE){
00206         if((pte = walkpgdir(pgdir, addr+i, 0)) == 0)
00207             panic("loadvm: address should exist");
00208         pa = PTE_ADDR(*pte);
00209         if(sz - i < PGSIZE)
00210             n = sz - i;
00211         else
00212             n = PGSIZE;
00213         if(readi(ip, P2V(pa), offset+i, n) != n)
00214             return -1;
00215     }
00216     return 0;
00217 }
00218
00219 // Allocate page tables and physical memory to grow process from oldsz to
00220 // newsz, which need not be page aligned. Returns new size or 0 on error.

```

```

00221 int
00222 allocvm(pde_t *pgdir, uint oldsz, uint newsz)
00223 {
00224     char *mem;
00225     uint a;
00226
00227     if(newsz >= KERNBASE)
00228         return 0;
00229     if(newsz < oldsz)
00230         return oldsz;
00231
00232     a = PGROUNDUP(oldsz);
00233     for(; a < newsz; a += PGSIZE){
00234         mem = kalloc();
00235         if(mem == 0){
00236             cprintf("allocvm out of memory\n");
00237             deallocvm(pgdir, newsz, oldsz);
00238             return 0;
00239         }
00240         memset(mem, 0, PGSIZE);
00241         if(mappages(pgdir, (char*)a, PGSIZE, V2P(mem), PTE_W|PTE_U) < 0){
00242             cprintf("allocvm out of memory (2)\n");
00243             deallocvm(pgdir, newsz, oldsz);
00244             kfree(mem);
00245             return 0;
00246         }
00247     }
00248     return newsz;
00249 }
00250
00251 // Deallocate user pages to bring the process size from oldsz to
00252 // newsz. oldsz and newsz need not be page-aligned, nor does newsz
00253 // need to be less than oldsz. oldsz can be larger than the actual
00254 // process size. Returns the new process size.
00255 int
00256 deallocvm(pde_t *pgdir, uint oldsz, uint newsz)
00257 {
00258     pte_t *pte;
00259     uint a, pa;
00260
00261     if(newsz >= oldsz)
00262         return oldsz;
00263
00264     a = PGROUNDUP(newsz);
00265     for(; a < oldsz; a += PGSIZE){
00266         pte = walkpgdir(pgdir, (char*)a, 0);
00267         if(!pte)
00268             a = PGADDR(PDX(a) + 1, 0, 0) - PGSIZE;
00269         else if ((*pte & PTE_P) != 0){
00270             pa = PTE_ADDR(*pte);
00271             if(pa == 0)
00272                 panic("kfree");
00273             char *v = P2V(pa);
00274             kfree(v);
00275             *pte = 0;
00276         }
00277     }
00278     return newsz;
00279 }
00280
00281 // Free a page table and all the physical memory pages
00282 // in the user part.
00283 void
00284 freevm(pde_t *pgdir)
00285 {
00286     uint i;
00287
00288     if(pgdir == 0)
00289         panic("freevm: no pgdir");
00290     deallocvm(pgdir, KERNBASE, 0);
00291     for(i = 0; i < NPENTRIES; i++){
00292         if(pgdir[i] & PTE_P){
00293             char *v = P2V(PTE_ADDR(pgdir[i]));
00294             kfree(v);
00295         }
00296     }
00297     kfree((char*)pgdir);
00298 }
00299
00300 // Clear PTE_U on a page. Used to create an inaccessible
00301 // page beneath the user stack.
00302 void
00303 clearpteu(pde_t *pgdir, char *uva)
00304 {
00305     pte_t *pte;
00306
00307     pte = walkpgdir(pgdir, uva, 0);

```



```

00308     if(pte == 0)
00309         panic("clearpteu");
00310     *pte &= ~PTE_U;
00311 }
00312
00313 // Given a parent process's page table, create a copy
00314 // of it for a child.
00315 pde_t*
00316 copyuvm(pde_t *pgdir, uint sz)
00317 {
00318     pde_t *d;
00319     pte_t *pte;
00320     uint pa, i, flags;
00321     char *mem;
00322
00323     if((d = setupkvm()) == 0)
00324         return 0;
00325     for(i = 0; i < sz; i += PGSIZE){
00326         if((pte = walkpgdir(pgdir, (void *) i, 0)) == 0)
00327             panic("copyuvm: pte should exist");
00328         if(!(*pte & PTE_P))
00329             panic("copyuvm: page not present");
00330         pa = PTE_ADDR(*pte);
00331         flags = PTE_FLAGS(*pte);
00332         if((mem = kalloc()) == 0)
00333             goto bad;
00334         memmove(mem, (char*)P2V(pa), PGSIZE);
00335         if(mappages(d, (void*)i, PGSIZE, V2P(mem), flags) < 0) {
00336             kfree(mem);
00337             goto bad;
00338         }
00339     }
00340     return d;
00341
00342 bad:
00343     freevm(d);
00344     return 0;
00345 }
00346
00347 //PAGEBREAK!
00348 // Map user virtual address to kernel address.
00349 char*
00350 uva2ka(pde_t *pgdir, char *uva)
00351 {
00352     pte_t *pte;
00353
00354     pte = walkpgdir(pgdir, uva, 0);
00355     if((*pte & PTE_P) == 0)
00356         return 0;
00357     if((*pte & PTE_U) == 0)
00358         return 0;
00359     return (char*)P2V(PTE_ADDR(*pte));
00360 }
00361
00362 // Copy len bytes from p to user address va in page table pgdir.
00363 // Most useful when pgdir is not the current page table.
00364 // uva2ka ensures this only works for PTE_U pages.
00365 int
00366 copyout(pde_t *pgdir, uint va, void *p, uint len)
00367 {
00368     char *buf, *pa0;
00369     uint n, va0;
00370
00371     buf = (char*)p;
00372     while(len > 0){
00373         va0 = (uint)PGROUNDDOWN(va);
00374         pa0 = uva2ka(pgdir, (char*)va0);
00375         if(pa0 == 0)
00376             return -1;
00377         n = PGSIZE - (va - va0);
00378         if(n > len)
00379             n = len;
00380         memmove(pa0 + (va - va0), buf, n);
00381         len -= n;
00382         buf += n;
00383         va = va0 + PGSIZE;
00384     }
00385     return 0;
00386 }
00387
00388 //PAGEBREAK!
00389 // Blank page.
00390 //PAGEBREAK!
00391 // Blank page.
00392 //PAGEBREAK!
00393 // Blank page.
00394

```

5.246 vm.d File Reference

5.247 vm.d

[Go to the documentation of this file.](#)

```
00001 vm.o: vm.c /usr/include/stdc-predef.h param.h types.h defs.h x86.h \  
00002 memlayout.h mmu.h proc.h elf.h
```

5.248 wc.c File Reference

```
#include "types.h"  
#include "stat.h"  
#include "user.h"
```

Functions

- int [main](#) (int argc, char *[argv](#)[])
- void [wc](#) (int fd, char *[name](#))

Variables

- char [buf](#) [512]

5.248.1 Function Documentation

5.248.1.1 main()

```
int main (  
    int argc,  
    char * argv[] )
```

Definition at line 36 of file [wc.c](#).

```
00037 {  
00038     int fd, i;  
00039  
00040     if(argc <= 1){  
00041         wc(0, "");  
00042         exit();  
00043     }  
00044  
00045     for(i = 1; i < argc; i++){  
00046         if((fd = open(argv[i], 0)) < 0){  
00047             printf(1, "wc: cannot open %s\n", argv[i]);  
00048             exit();  
00049         }  
00050         wc(fd, argv[i]);  
00051         close(fd);  
00052     }  
00053     exit();  
00054 }
```

5.248.1.2 wc()

```
void wc (
    int fd,
    char * name )
```

Definition at line 8 of file [wc.c](#).

```
00009 {
00010     int i, n;
00011     int l, w, c, inword;
00012
00013     l = w = c = 0;
00014     inword = 0;
00015     while((n = read(fd, buf, sizeof(buf))) > 0){
00016         for(i=0; i<n; i++){
00017             c++;
00018             if(buf[i] == '\n')
00019                 l++;
00020             if(strchr(" \r\t\n\v", buf[i]))
00021                 inword = 0;
00022             else if(!inword){
00023                 w++;
00024                 inword = 1;
00025             }
00026         }
00027     }
00028     if(n < 0){
00029         printf(l, "wc: read error\n");
00030         exit();
00031     }
00032     printf(l, "%d %d %d %s\n", l, w, c, name);
00033 }
```

Referenced by [main\(\)](#).

5.248.2 Variable Documentation

5.248.2.1 buf

```
char buf[512]
```

Definition at line 5 of file [wc.c](#).

5.249 wc.c

[Go to the documentation of this file.](#)

```
00001 #include "types.h"
00002 #include "stat.h"
00003 #include "user.h"
00004
00005 char buf[512];
00006
00007 void
00008 wc(int fd, char *name)
00009 {
00010     int i, n;
00011     int l, w, c, inword;
00012
00013     l = w = c = 0;
00014     inword = 0;
00015     while((n = read(fd, buf, sizeof(buf))) > 0){
00016         for(i=0; i<n; i++){
00017             c++;
00018             if(buf[i] == '\n')
```

```

00019         l++;
00020         if(strchr(" \r\t\n\v", buf[i]))
00021             inword = 0;
00022         else if(!inword){
00023             w++;
00024             inword = 1;
00025         }
00026     }
00027 }
00028 if(n < 0){
00029     printf(1, "wc: read error\n");
00030     exit();
00031 }
00032 printf(1, "%d %d %d %s\n", l, w, c, name);
00033 }
00034
00035 int
00036 main(int argc, char *argv[])
00037 {
00038     int fd, i;
00039
00040     if(argc <= 1){
00041         wc(0, "");
00042         exit();
00043     }
00044
00045     for(i = 1; i < argc; i++){
00046         if((fd = open(argv[i], 0)) < 0){
00047             printf(1, "wc: cannot open %s\n", argv[i]);
00048             exit();
00049         }
00050         wc(fd, argv[i]);
00051         close(fd);
00052     }
00053     exit();
00054 }

```

5.250 wc.d File Reference

5.251 wc.d

[Go to the documentation of this file.](#)

00001 wc.o: wc.c /usr/include/stdc-predef.h types.h stat.h user.h

5.252 x86.h File Reference

Classes

- struct [trapframe](#)

Functions

- static void [cli](#) (void)
- static [uchar](#) [inb](#) ([ushort](#) port)
- static void [insl](#) (int port, void *addr, int cnt)
- static void [lcr3](#) ([uint](#) val)
- static void [lgdt](#) (struct [segdesc](#) *p, int size)
- static void [lidt](#) (struct [gatedesc](#) *p, int size)
- static void [loadgs](#) ([ushort](#) v)
- static void [ltr](#) ([ushort](#) sel)
- static void [outb](#) ([ushort](#) port, [uchar](#) data)

- static void [outsl](#) (int port, const void *addr, int cnt)
- static void [outw](#) ([ushort](#) port, [ushort](#) data)
- static [uint](#) [rcr2](#) (void)
- static [uint](#) [readeflags](#) (void)
- static void [sti](#) (void)
- static void [stosb](#) (void *addr, int data, int cnt)
- static void [stosl](#) (void *addr, int data, int cnt)
- static [uint](#) [xchg](#) (volatile [uint](#) *addr, [uint](#) newval)

5.252.1 Function Documentation

5.252.1.1 cli()

```
static void cli (
    void ) [inline], [static]
```

Definition at line 109 of file [x86.h](#).

```
00110 {
00111     asm volatile("cli");
00112 }
```

Referenced by [consputc\(\)](#), [panic\(\)](#), and [pushcli\(\)](#).

5.252.1.2 inb()

```
static uchar inb (
    ushort port ) [inline], [static]
```

Definition at line 4 of file [x86.h](#).

```
00005 {
00006     uchar data;
00007
00008     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
00009     return data;
00010 }
```

Referenced by [cgaputc\(\)](#), [cmos_read\(\)](#), [ideinit\(\)](#), [idewait\(\)](#), [kbdgetc\(\)](#), [mpinit\(\)](#), [uartgetc\(\)](#), [uartinit\(\)](#), [uartputc\(\)](#), and [waitdisk\(\)](#).

5.252.1.3 insl()

```
static void insl (
    int port,
    void * addr,
    int cnt ) [inline], [static]
```

Definition at line 13 of file [x86.h](#).

```
00014 {
00015     asm volatile("cld; rep insl" :
00016         "=D" (addr), "=c" (cnt) :
00017         "d" (port), "0" (addr), "1" (cnt) :
00018         "memory", "cc");
00019 }
```

Referenced by [ideintr\(\)](#), and [readsect\(\)](#).

5.252.1.4 lcr3()

```
static void lcr3 (
    uint val ) [inline], [static]
```

Definition at line 142 of file [x86.h](#).

```
00143 {
00144     asm volatile("movl %0,%%cr3" : : "r" (val));
00145 }
```

Referenced by [switchkvm\(\)](#), and [switchuvm\(\)](#).

5.252.1.5 lgdt()

```
static void lgdt (
    struct segdesc * p,
    int size ) [inline], [static]
```

Definition at line 63 of file [x86.h](#).

```
00064 {
00065     volatile ushort pd[3];
00066
00067     pd[0] = size-1;
00068     pd[1] = (uint)p;
00069     pd[2] = (uint)p » 16;
00070
00071     asm volatile("lgdt (%0)" : : "r" (pd));
00072 }
```

Referenced by [seginit\(\)](#).

5.252.1.6 lidt()

```
static void lidt (
    struct gatedesc * p,
    int size ) [inline], [static]
```

Definition at line 77 of file [x86.h](#).

```
00078 {
00079     volatile ushort pd[3];
00080
00081     pd[0] = size-1;
00082     pd[1] = (uint)p;
00083     pd[2] = (uint)p » 16;
00084
00085     asm volatile("lidt (%0)" : : "r" (pd));
00086 }
```

Referenced by [idtinit\(\)](#).

5.252.1.7 loadgs()

```
static void loadgs (  
    ushort v ) [inline], [static]
```

Definition at line 103 of file [x86.h](#).

```
00104 {  
00105     asm volatile("movw %0, %%gs" : : "r" (v));  
00106 }
```

5.252.1.8 ltr()

```
static void ltr (  
    ushort sel ) [inline], [static]
```

Definition at line 89 of file [x86.h](#).

```
00090 {  
00091     asm volatile("ltr %0" : : "r" (sel));  
00092 }
```

Referenced by [switchuvmm\(\)](#).

5.252.1.9 outb()

```
static void outb (  
    ushort port,  
    uchar data ) [inline], [static]
```

Definition at line 22 of file [x86.h](#).

```
00023 {  
00024     asm volatile("out %0,%1" : : "a" (data), "d" (port));  
00025 }
```

Referenced by [cgaputc\(\)](#), [cmos_read\(\)](#), [ideinit\(\)](#), [idestart\(\)](#), [lapiestartap\(\)](#), [mpinit\(\)](#), [picinit\(\)](#), [readsect\(\)](#), [sys_halt\(\)](#), [uartinit\(\)](#), and [uartputc\(\)](#).

5.252.1.10 outsl()

```
static void outsl (  
    int port,  
    const void * addr,  
    int cnt ) [inline], [static]
```

Definition at line 34 of file [x86.h](#).

```
00035 {  
00036     asm volatile("cld; rep outsl" :  
00037         "=S" (addr), "=c" (cnt) :  
00038         "d" (port), "0" (addr), "1" (cnt) :  
00039         "cc");  
00040 }
```

Referenced by [idestart\(\)](#).

5.252.1.11 outw()

```
static void outw (
    ushort port,
    ushort data ) [inline], [static]
```

Definition at line 28 of file [x86.h](#).

```
00029 {
00030     asm volatile("out %0,%1" : : "a" (data), "d" (port));
00031 }
```

5.252.1.12 rcr2()

```
static uint rcr2 (
    void ) [inline], [static]
```

Definition at line 134 of file [x86.h](#).

```
00135 {
00136     uint val;
00137     asm volatile("movl %%cr2,%0" : "=r" (val));
00138     return val;
00139 }
```

Referenced by [trap\(\)](#).

5.252.1.13 readeflags()

```
static uint readeflags (
    void ) [inline], [static]
```

Definition at line 95 of file [x86.h](#).

```
00096 {
00097     uint eflags;
00098     asm volatile("pushfl; popl %0" : "=r" (eflags));
00099     return eflags;
00100 }
```

Referenced by [mycpu\(\)](#), [popcli\(\)](#), [pushcli\(\)](#), and [sched\(\)](#).

5.252.1.14 sti()

```
static void sti (
    void ) [inline], [static]
```

Definition at line 115 of file [x86.h](#).

```
00116 {
00117     asm volatile("sti");
00118 }
```

Referenced by [cps\(\)](#), [popcli\(\)](#), and [scheduler\(\)](#).

5.252.1.15 stosb()

```
static void stosb (  
    void * addr,  
    int data,  
    int cnt ) [inline], [static]
```

Definition at line 43 of file [x86.h](#).

```
00044 {  
00045     asm volatile("cld; rep stosb" :  
00046         "=D" (addr), "=c" (cnt) :  
00047         "0" (addr), "1" (cnt), "a" (data) :  
00048         "memory", "cc");  
00049 }
```

Referenced by [bootmain\(\)](#), and [memset\(\)](#).

5.252.1.16 stosl()

```
static void stosl (  
    void * addr,  
    int data,  
    int cnt ) [inline], [static]
```

Definition at line 52 of file [x86.h](#).

```
00053 {  
00054     asm volatile("cld; rep stosl" :  
00055         "=D" (addr), "=c" (cnt) :  
00056         "0" (addr), "1" (cnt), "a" (data) :  
00057         "memory", "cc");  
00058 }
```

Referenced by [memset\(\)](#).

5.252.1.17 xchg()

```
static uint xchg (  
    volatile uint * addr,  
    uint newval ) [inline], [static]
```

Definition at line 121 of file [x86.h](#).

```
00122 {  
00123     uint result;  
00124  
00125     // The + in "+m" denotes a read-modify-write operand.  
00126     asm volatile("lock; xchgl %0, %1" :  
00127         "+m" (*addr), "=a" (result) :  
00128         "1" (newval) :  
00129         "cc");  
00130     return result;  
00131 }
```

Referenced by [acquire\(\)](#).

5.253 x86.h

[Go to the documentation of this file.](#)

```

00001 // Routines to let C code use special x86 instructions.
00002
00003 static inline uchar
00004 inb(ushort port)
00005 {
00006     uchar data;
00007
00008     asm volatile("in %1,%0" : "=a" (data) : "d" (port));
00009     return data;
00010 }
00011
00012 static inline void
00013 insl(int port, void *addr, int cnt)
00014 {
00015     asm volatile("cld; rep insl" :
00016                 "=D" (addr), "=c" (cnt) :
00017                 "d" (port), "0" (addr), "1" (cnt) :
00018                 "memory", "cc");
00019 }
00020
00021 static inline void
00022 outb(ushort port, uchar data)
00023 {
00024     asm volatile("out %0,%1" : : "a" (data), "d" (port));
00025 }
00026
00027 static inline void
00028 outw(ushort port, ushort data)
00029 {
00030     asm volatile("out %0,%1" : : "a" (data), "d" (port));
00031 }
00032
00033 static inline void
00034 outsl(int port, const void *addr, int cnt)
00035 {
00036     asm volatile("cld; rep outsl" :
00037                 "=S" (addr), "=c" (cnt) :
00038                 "d" (port), "0" (addr), "1" (cnt) :
00039                 "cc");
00040 }
00041
00042 static inline void
00043 stosb(void *addr, int data, int cnt)
00044 {
00045     asm volatile("cld; rep stosb" :
00046                 "=D" (addr), "=c" (cnt) :
00047                 "0" (addr), "1" (cnt), "a" (data) :
00048                 "memory", "cc");
00049 }
00050
00051 static inline void
00052 stosl(void *addr, int data, int cnt)
00053 {
00054     asm volatile("cld; rep stosl" :
00055                 "=D" (addr), "=c" (cnt) :
00056                 "0" (addr), "1" (cnt), "a" (data) :
00057                 "memory", "cc");
00058 }
00059
00060 struct segdesc;
00061
00062 static inline void
00063 lgdt(struct segdesc *p, int size)
00064 {
00065     volatile ushort pd[3];
00066
00067     pd[0] = size-1;
00068     pd[1] = (uint)p;
00069     pd[2] = (uint)p >> 16;
00070
00071     asm volatile("lgdt (%0)" : : "r" (pd));
00072 }
00073
00074 struct gatedesc;
00075
00076 static inline void
00077 lidt(struct gatedesc *p, int size)
00078 {
00079     volatile ushort pd[3];
00080
00081     pd[0] = size-1;
00082     pd[1] = (uint)p;

```

```

00083     pd[2] = (uint)p >> 16;
00084
00085     asm volatile("lidt (%0)" : : "r" (pd));
00086 }
00087
00088 static inline void
00089 ltr(ushort sel)
00090 {
00091     asm volatile("ltr %0" : : "r" (sel));
00092 }
00093
00094 static inline uint
00095 readeflags(void)
00096 {
00097     uint eflags;
00098     asm volatile("pushfl; popl %0" : "=r" (eflags));
00099     return eflags;
00100 }
00101
00102 static inline void
00103 loadgs(ushort v)
00104 {
00105     asm volatile("movw %0, %%gs" : : "r" (v));
00106 }
00107
00108 static inline void
00109 cli(void)
00110 {
00111     asm volatile("cli");
00112 }
00113
00114 static inline void
00115 sti(void)
00116 {
00117     asm volatile("sti");
00118 }
00119
00120 static inline uint
00121 xchg(volatile uint *addr, uint newval)
00122 {
00123     uint result;
00124
00125     // The + in "+m" denotes a read-modify-write operand.
00126     asm volatile("lock; xchgl %0, %1" :
00127                 "+m" (*addr), "=a" (result) :
00128                 "1" (newval) :
00129                 "cc");
00130     return result;
00131 }
00132
00133 static inline uint
00134 rcr2(void)
00135 {
00136     uint val;
00137     asm volatile("movl %%cr2,%0" : "=r" (val));
00138     return val;
00139 }
00140
00141 static inline void
00142 lcr3(uint val)
00143 {
00144     asm volatile("movl %0,%%cr3" : : "r" (val));
00145 }
00146
00147 //PAGEBREAK: 36
00148 // Layout of the trap frame built on the stack by the
00149 // hardware and by trapasm.S, and passed to trap().
00150 struct trapframe {
00151     // registers as pushed by pusha
00152     uint edi;
00153     uint esi;
00154     uint ebp;
00155     uint oesp;      // useless & ignored
00156     uint ebx;
00157     uint edx;
00158     uint ecx;
00159     uint eax;
00160
00161     // rest of trap frame
00162     ushort gs;
00163     ushort padding1;
00164     ushort fs;
00165     ushort padding2;
00166     ushort es;
00167     ushort padding3;
00168     ushort ds;
00169     ushort padding4;

```

```
00170     uint trapno;
00171
00172     // below here defined by x86 hardware
00173     uint err;
00174     uint eip;
00175     ushort cs;
00176     ushort padding5;
00177     uint eflags;
00178
00179     // below here only when crossing rings, such as from user to kernel
00180     uint esp;
00181     ushort ss;
00182     ushort padding6;
00183 };
```

5.254 xv6.dox File Reference

5.255 zombie.c File Reference

```
#include "types.h"
#include "stat.h"
#include "user.h"
```

Functions

- int [main](#) (void)

5.255.1 Function Documentation

5.255.1.1 main()

```
int main (
    void )
```

Definition at line 9 of file [zombie.c](#).

```
00010 {
00011     if(fork() > 0)
00012         sleep(5); // Let child exit before parent.
00013     exit();
00014 }
```

5.256 zombie.c

[Go to the documentation of this file.](#)

```
00001 // Create a zombie process that
00002 // must be reparented at exit.
00003
00004 #include "types.h"
00005 #include "stat.h"
00006 #include "user.h"
00007
00008 int
00009 main(void)
00010 {
00011     if(fork() > 0)
00012         sleep(5); // Let child exit before parent.
00013     exit();
00014 }
```

5.257 zombie.d File Reference

5.258 zombie.d

[Go to the documentation of this file.](#)

```
00001 zombie.o: zombie.c /usr/include/stdc-predef.h types.h stat.h user.h
```


Index

- `__attribute__`
 - `main.c`, [271](#)
 - `_binary_fs_img_size`
 - `memide.c`, [276](#)
 - `_binary_fs_img_start`
 - `memide.c`, [276](#)
- acquire
 - `defs.h`, [110](#)
 - `spinlock.c`, [371](#)
- acquiresleep
 - `defs.h`, [111](#)
 - `sleeplock.c`, [368](#)
- addr
 - `mpioapic`, [43](#)
- addrs
 - `dinode`, [17](#)
 - `inode`, [30](#)
- Align
 - `umalloc.c`, [446](#)
- align
 - `proghdr`, [52](#)
- allocproc
 - `proc.c`, [325](#)
- allocvm
 - `defs.h`, [111](#)
 - `vm.c`, [515](#)
- ALT
 - `kbd.h`, [235](#)
- apicid
 - `cpu`, [14](#)
 - `mpproc`, [45](#)
- apicno
 - `mpioapic`, [43](#)
- argfd
 - `sysfile.c`, [404](#)
- argint
 - `defs.h`, [112](#)
 - `syscall.c`, [384](#)
- argptest
 - `usertests.c`, [465](#)
- argptr
 - `defs.h`, [112](#)
 - `syscall.c`, [384](#)
- args
 - `gatedesc`, [27](#)
- argstr
 - `defs.h`, [112](#)
 - `syscall.c`, [385](#)
- argv
 - `execcmd`, [23](#)
 - `init.c`, [222](#)
- asm.h, [81](#)
 - `SEG_ASM`, [81](#)
 - `SEG_NULLASM`, [81](#)
 - `STA_R`, [81](#)
 - `STA_W`, [82](#)
 - `STA_X`, [82](#)
- ASSERT
 - `lapic.c`, [245](#)
- atoi
 - `ulib.c`, [440](#)
 - `user.h`, [450](#)
- avl
 - `segdesc`, [59](#)
- B_DIRTY
 - `buf.h`, [92](#)
- B_VALID
 - `buf.h`, [92](#)
- BACK
 - `sh.c`, [350](#)
- backcmd, [9](#)
 - `cmd`, [9](#)
 - `sh.c`, [352](#)
 - `type`, [9](#)
- BACKSPACE
 - `console.c`, [96](#)
- ballocc
 - `fs.c`, [186](#)
 - `mkfs.c`, [283](#)
- base
 - `umalloc.c`, [447](#)
- base_15_0
 - `segdesc`, [59](#)
- base_23_16
 - `segdesc`, [59](#)
- base_31_24
 - `segdesc`, [59](#)
- BBLOCK
 - `fs.h`, [207](#)
- bcache
 - `bio.c`, [85](#)
- BCAST
 - `lapic.c`, [245](#)
- begin_op
 - `defs.h`, [113](#)
 - `log.c`, [260](#)
- bfree
 - `fs.c`, [186](#)

- bget
 - bio.c, 83
- BIG
 - usertests.c, 464
- bigargtest
 - usertests.c, 465
- bigdir
 - usertests.c, 466
- bigfile
 - usertests.c, 466
- bigwrite
 - usertests.c, 467
- binit
 - bio.c, 83
 - defs.h, 113
- bio.c, 82
 - bcache, 85
 - bget, 83
 - binit, 83
 - bread, 84
 - breelse, 84
 - buf, 85
 - bwrite, 85
 - head, 85
 - lock, 85
- bio.d, 88
- block
 - logheader, 38
- blockno
 - buf, 10
- bmap
 - fs.c, 187
- bmapstart
 - superblock, 65
- bootasm.d, 88
- bootmain
 - bootmain.c, 89
- bootmain.c, 88
 - bootmain, 89
 - readsect, 89
 - readseg, 89
 - SECTSIZE, 88
 - waitdisk, 90
- bootmain.d, 91
- BPB
 - fs.h, 207
- bread
 - bio.c, 84
 - defs.h, 113
- breelse
 - bio.c, 84
 - defs.h, 114
- BSIZE
 - fs.h, 207
- bsstest
 - usertests.c, 468
- buf, 10
 - bio.c, 85
- blockno, 10
- cat.c, 94
- console.c, 101
- data, 10
- dev, 10
- flags, 11
- grep.c, 212
- lock, 11
- next, 11
- prev, 11
- qnext, 11
- refcnt, 12
- usertests.c, 492
- wc.c, 529
- buf.h, 92
 - B_DIRTY, 92
 - B_VALID, 92
- BUSY
 - lapic.c, 245
- bwrite
 - bio.c, 85
 - defs.h, 114
- bzero
 - fs.c, 187
- C
 - console.c, 96
 - kbd.h, 235
- CAPSLOCK
 - kbd.h, 235
- cat
 - cat.c, 93
- cat.c, 93
 - buf, 94
 - cat, 93
 - main, 93
- cat.d, 95
- cgaputc
 - console.c, 97
- chan
 - proc, 49
- chdir
 - user.h, 451
- checksum
 - mp, 39
 - mpconf, 41
- chpr
 - defs.h, 114
 - proc.c, 326
 - user.h, 451
- clearpteu
 - defs.h, 115
 - vm.c, 516
- cli
 - x86.h, 531
- close
 - user.h, 451
- cmd, 12
 - backcmd, 9

- redircmd, [54](#)
- type, [12](#)
- CMOS_PORT
 - lapic.c, [246](#)
- cmos_read
 - lapic.c, [252](#)
- CMOS_RETURN
 - lapic.c, [246](#)
- CMOS_STATA
 - lapic.c, [246](#)
- CMOS_STATB
 - lapic.c, [246](#)
- CMOS_UIP
 - lapic.c, [246](#)
- cmostime
 - defs.h, [115](#)
 - lapic.c, [252](#)
- COM1
 - uart.c, [437](#)
- commit
 - log.c, [261](#)
- committing
 - log, [36](#)
- concreate
 - usertests.c, [468](#)
- cons
 - console.c, [101](#)
- CONSOLE
 - file.h, [179](#)
- console.c, [95](#)
 - BACKSPACE, [96](#)
 - buf, [101](#)
 - C, [96](#)
 - cgaputc, [97](#)
 - cons, [101](#)
 - consoleinit, [97](#)
 - consoleintr, [97](#)
 - consoleread, [98](#)
 - consolewrite, [99](#)
 - consputc, [99](#)
 - cprintf, [99](#)
 - crt, [102](#)
 - CRTPORT, [96](#)
 - e, [102](#)
 - input, [102](#)
 - INPUT_BUF, [96](#)
 - lock, [102](#)
 - locking, [102](#)
 - panic, [100](#)
 - panicked, [102](#)
 - printint, [101](#)
 - r, [103](#)
 - w, [103](#)
- console.d, [107](#)
- consoleinit
 - console.c, [97](#)
 - defs.h, [116](#)
- consoleintr
 - console.c, [97](#)
 - defs.h, [116](#)
- consoleread
 - console.c, [98](#)
- consolewrite
 - console.c, [99](#)
- consputc
 - console.c, [99](#)
- context, [13](#)
 - ebp, [13](#)
 - ebx, [13](#)
 - edi, [13](#)
 - eip, [13](#)
 - esi, [14](#)
 - proc, [49](#)
- CONV
 - lapic.c, [246](#)
- copyout
 - defs.h, [117](#)
 - vm.c, [516](#)
- copyuvms
 - defs.h, [117](#)
 - vm.c, [517](#)
- cprintf
 - console.c, [99](#)
 - defs.h, [118](#)
- cps
 - defs.h, [118](#)
 - proc.c, [326](#)
 - user.h, [451](#)
- cpu, [14](#)
 - apicid, [14](#)
 - gdt, [15](#)
 - intena, [15](#)
 - ncli, [15](#)
 - proc, [15](#)
 - scheduler, [15](#)
 - spinlock, [63](#)
 - started, [16](#)
 - ts, [16](#)
- cpuid
 - defs.h, [119](#)
 - proc.c, [327](#)
- cpus
 - mp.c, [306](#)
 - proc.h, [344](#)
- CR0_PE
 - mmu.h, [294](#)
- CR0_PG
 - mmu.h, [294](#)
- CR0_WP
 - mmu.h, [294](#)
- cr3
 - taskstate, [68](#)
- CR4_PSE
 - mmu.h, [295](#)
- create
 - sysfile.c, [405](#)

- createdelete
 - usertests.c, [469](#)
- createtest
 - usertests.c, [470](#)
- crt
 - console.c, [102](#)
- CRTPORT
 - console.c, [96](#)
- cs
 - gatedesc, [27](#)
 - taskstate, [68](#)
 - trapframe, [75](#)
- CTL
 - kbd.h, [236](#)
- ctlmap
 - kbd.h, [239](#)
- cwd
 - proc, [49](#)
- data
 - buf, [10](#)
 - ioapic, [33](#)
 - pipe, [46](#)
 - vm.c, [522](#)
- date.h, [107](#)
- DAY
 - lapic.c, [247](#)
- day
 - rtcdate, [56](#)
- db
 - segdesc, [59](#)
- deallocvm
 - defs.h, [119](#)
 - vm.c, [517](#)
- DEASSERT
 - lapic.c, [247](#)
- defs.h, [107](#)
 - acquire, [110](#)
 - acquiresleep, [111](#)
 - allocvm, [111](#)
 - argint, [112](#)
 - argptr, [112](#)
 - argstr, [112](#)
 - begin_op, [113](#)
 - binit, [113](#)
 - bread, [113](#)
 - brese, [114](#)
 - bwrite, [114](#)
 - chpr, [114](#)
 - clearpteu, [115](#)
 - cmostime, [115](#)
 - consoleinit, [116](#)
 - consoleintr, [116](#)
 - copyout, [117](#)
 - copyvm, [117](#)
 - cprintf, [118](#)
 - cps, [118](#)
 - cpuid, [119](#)
 - deallocvm, [119](#)
 - dirlink, [120](#)
 - dirlookup, [120](#)
 - end_op, [121](#)
 - exec, [121](#)
 - exit, [123](#)
 - fetchint, [123](#)
 - fetchstr, [124](#)
 - filealloc, [124](#)
 - fileclose, [124](#)
 - filedup, [125](#)
 - fileinit, [125](#)
 - fileread, [125](#)
 - filestat, [126](#)
 - filewrite, [126](#)
 - fork, [127](#)
 - freevm, [128](#)
 - getcallerpcs, [128](#)
 - getprocs, [128](#)
 - growproc, [128](#)
 - holding, [129](#)
 - holdingsleep, [129](#)
 - ialloc, [129](#)
 - ideinit, [130](#)
 - ideintr, [130](#)
 - iderw, [131](#)
 - idtinit, [131](#)
 - idup, [132](#)
 - iinit, [132](#)
 - ilock, [132](#)
 - initlock, [133](#)
 - initlog, [133](#)
 - initsleeplock, [133](#)
 - initvm, [134](#)
 - ioapicenable, [134](#)
 - ioapicid, [160](#)
 - ioapicinit, [134](#)
 - iput, [135](#)
 - ismp, [161](#)
 - iunlock, [135](#)
 - iunlockput, [135](#)
 - iupdate, [136](#)
 - kalloc, [136](#)
 - kbdintr, [136](#)
 - kfree, [137](#)
 - kill, [137](#)
 - kinit1, [137](#)
 - kinit2, [138](#)
 - kvmalloc, [138](#)
 - lapic, [161](#)
 - lapiceoi, [138](#)
 - lapicid, [138](#)
 - lapicinit, [139](#)
 - lapicstartap, [139](#)
 - loadvm, [140](#)
 - log_write, [140](#)
 - memcmp, [141](#)
 - memmove, [141](#)
 - memset, [142](#)

- microdelay, 142
- mpinit, 142
- mycpu, 143
- myproc, 143
- namecmp, 144
- namei, 144
- nameiparent, 144
- NELEM, 110
- panic, 144
- picenable, 145
- picinit, 145
- pinit, 145
- pipealloc, 145
- pipeclose, 146
- piperead, 146
- pipewrite, 147
- popcli, 147
- procdump, 148
- pushcli, 148
- readi, 148
- readsb, 149
- release, 149
- releasesleep, 150
- safestrcpy, 150
- sched, 150
- scheduler, 151
- seginit, 151
- setproc, 152
- setupkvm, 152
- sleep, 152
- stati, 153
- strlen, 153
- strncmp, 154
- strncpy, 154
- switchkvm, 154
- switchvm, 155
- swtch, 155
- syscall, 155
- ticks, 161
- tickslock, 161
- timerinit, 156
- tvinit, 156
- uartinit, 156
- uartintr, 157
- uartputc, 157
- userinit, 157
- uva2ka, 158
- wait, 158
- wakeup, 159
- writei, 159
- yield, 160
- DELIVS
 - lapic.c, 247
- dev
 - buf, 10
 - inode, 30
 - log, 36
 - stat, 64
- DEVSPACE
 - memlayout.h, 278
- devsw, 16
 - file.c, 176
 - file.h, 180
 - read, 17
 - write, 17
- dinode, 17
 - addrs, 17
 - major, 18
 - minor, 18
 - nlink, 18
 - size, 18
 - type, 18
- dirent, 19
 - inum, 19
 - name, 19
- dirfile
 - usertests.c, 471
- dirlink
 - defs.h, 120
 - fs.c, 188
- dirlookup
 - defs.h, 120
 - fs.c, 188
- DIRSIZ
 - fs.h, 207
- dirtest
 - usertests.c, 472
- disksize
 - memide.c, 277
- dpl
 - gatedesc, 27
 - segdesc, 59
- DPL_USER
 - mmu.h, 295
- ds
 - taskstate, 68
 - trapframe, 75
- dup
 - user.h, 452
- e
 - console.c, 102
- E0ESC
 - kbd.h, 236
- eargv
 - execcmd, 23
- eax
 - taskstate, 68
 - trapframe, 76
- ebp
 - context, 13
 - taskstate, 69
 - trapframe, 76
- ebx
 - context, 13
 - taskstate, 69
 - trapframe, 76

- echo.c, [164](#)
 - main, [164](#)
- echo.d, [165](#)
- echoargv
 - usertests.c, [492](#)
- ecx
 - taskstate, [69](#)
 - trapframe, [76](#)
- edi
 - context, [13](#)
 - taskstate, [69](#)
 - trapframe, [76](#)
- edx
 - taskstate, [69](#)
 - trapframe, [76](#)
- efile
 - redircmd, [55](#)
- eflags
 - taskstate, [69](#)
 - trapframe, [77](#)
- ehsize
 - elfhdr, [20](#)
- eip
 - context, [13](#)
 - taskstate, [70](#)
 - trapframe, [77](#)
- elf
 - elfhdr, [20](#)
- elf.h, [165](#)
 - ELF_MAGIC, [165](#)
 - ELF_PROG_FLAG_EXEC, [165](#)
 - ELF_PROG_FLAG_READ, [166](#)
 - ELF_PROG_FLAG_WRITE, [166](#)
 - ELF_PROG_LOAD, [166](#)
- ELF_MAGIC
 - elf.h, [165](#)
- ELF_PROG_FLAG_EXEC
 - elf.h, [165](#)
- ELF_PROG_FLAG_READ
 - elf.h, [166](#)
- ELF_PROG_FLAG_WRITE
 - elf.h, [166](#)
- ELF_PROG_LOAD
 - elf.h, [166](#)
- elfhdr, [20](#)
 - ehsize, [20](#)
 - elf, [20](#)
 - entry, [20](#)
 - flags, [21](#)
 - machine, [21](#)
 - magic, [21](#)
 - phentsize, [21](#)
 - phnum, [21](#)
 - phoff, [21](#)
 - shentsize, [22](#)
 - shnum, [22](#)
 - shoff, [22](#)
 - shstrndx, [22](#)
- type, [22](#)
- version, [22](#)
- EMBRYO
 - proc.h, [344](#)
- ENABLE
 - lapic.c, [247](#)
- end
 - kalloc.c, [230](#)
- end_op
 - defs.h, [121](#)
 - log.c, [261](#)
- entry
 - elfhdr, [20](#)
 - mpconf, [41](#)
- entryother.d, [167](#)
- entrypgdir
 - main.c, [273](#)
- EOI
 - lapic.c, [247](#)
- err
 - trapframe, [77](#)
- ERROR
 - lapic.c, [247](#)
- es
 - taskstate, [70](#)
 - trapframe, [77](#)
- esi
 - context, [14](#)
 - taskstate, [70](#)
 - trapframe, [77](#)
- esp
 - taskstate, [70](#)
 - trapframe, [78](#)
- esp0
 - taskstate, [70](#)
- esp1
 - taskstate, [70](#)
- esp2
 - taskstate, [71](#)
- ESR
 - lapic.c, [248](#)
- EXEC
 - sh.c, [351](#)
- exec
 - defs.h, [121](#)
 - exec.c, [167](#)
 - user.h, [452](#)
- exec.c, [167](#)
 - exec, [167](#)
- exec.d, [170](#)
- execcmd, [23](#)
 - argv, [23](#)
 - eargv, [23](#)
 - sh.c, [352](#)
 - type, [23](#)
- exectest
 - usertests.c, [472](#)
- exit

- defs.h, 123
- proc.c, 327
- user.h, 453
- exit.c, 171
 - main, 171
- exit.d, 171
- exitiputtest
 - usertests.c, 472
- exitwait
 - usertests.c, 473
- EXTMEM
 - memlayout.h, 278
- fcntl.h, 172
 - O_CREATE, 172
 - O_RDONLY, 172
 - O_RDWR, 172
 - O_WRONLY, 172
- fd
 - redircmd, 55
- FD_INODE
 - file, 24
- FD_NONE
 - file, 24
- FD_PIPE
 - file, 24
- fdalloc
 - sysfile.c, 405
- feature
 - mpproc, 45
- fetchint
 - defs.h, 123
 - syscall.c, 385
- fetchstr
 - defs.h, 124
 - syscall.c, 385
- file, 24
 - FD_INODE, 24
 - FD_NONE, 24
 - FD_PIPE, 24
 - file.c, 176
 - ip, 25
 - off, 25
 - pipe, 25
 - readable, 25
 - redircmd, 55
 - ref, 25
 - type, 26
 - writable, 26
- file.c, 173
 - devsw, 176
 - file, 176
 - filealloc, 173
 - fileclose, 174
 - filedup, 174
 - fileinit, 174
 - fileread, 175
 - filestat, 175
 - filewrite, 175
 - ftable, 176
 - lock, 177
- file.d, 179
- file.h, 179
 - CONSOLE, 179
 - devsw, 180
- filealloc
 - defs.h, 124
 - file.c, 173
- fileclose
 - defs.h, 124
 - file.c, 174
- filedup
 - defs.h, 125
 - file.c, 174
- fileinit
 - defs.h, 125
 - file.c, 174
- fileread
 - defs.h, 125
 - file.c, 175
- filestat
 - defs.h, 126
 - file.c, 175
- filesz
 - proghdr, 53
- filewrite
 - defs.h, 126
 - file.c, 175
- fill_rtcdade
 - lapic.c, 253
- FIXED
 - lapic.c, 248
- FL_IF
 - mmu.h, 295
- flags
 - buf, 11
 - elfhdr, 21
 - mpioapic, 44
 - mpproc, 45
 - proghdr, 53
- fmtname
 - ls.c, 268
- foo.c, 180
 - main, 181
- foo.d, 182
- fork
 - defs.h, 127
 - proc.c, 328
 - user.h, 454
- fork1
 - sh.c, 352
- forkret
 - proc.c, 329
- forktest
 - forktest.c, 183
 - usertests.c, 473
- forktest.c, 182

- forktest, 183
- main, 183
- N, 182
- printf, 183
- forktest.d, 185
- fourfiles
 - usertests.c, 474
- fourteen
 - usertests.c, 475
- free
 - umalloc.c, 446
 - user.h, 455
- freeblock
 - mkfs.c, 287
- freeinode
 - mkfs.c, 288
- freelist
 - kalloc.c, 230
- freep
 - umalloc.c, 448
- freerange
 - kalloc.c, 228
- freevm
 - defs.h, 128
 - vm.c, 518
- fs
 - taskstate, 71
 - trapframe, 78
- fs.c, 185
 - balloc, 186
 - bfree, 186
 - bmap, 187
 - bzero, 187
 - dirlink, 188
 - dirlookup, 188
 - ialloc, 189
 - icache, 197
 - idup, 189
 - iget, 190
 - iinit, 190
 - ilock, 191
 - inode, 198
 - iput, 191
 - itrunc, 192
 - iunlock, 192
 - iunlockput, 193
 - iupdate, 193
 - lock, 198
 - min, 186
 - namecmp, 193
 - namei, 194
 - nameiparent, 194
 - namex, 194
 - readi, 195
 - readsb, 196
 - sb, 198
 - skipelem, 196
 - stati, 196
 - writei, 197
- fs.d, 206
- fs.h, 206
 - BBLOCK, 207
 - BPB, 207
 - BSIZE, 207
 - DIRSIZ, 207
 - IBLOCK, 207
 - IPB, 207
 - MAXFILE, 208
 - NDIRECT, 208
 - NINDIRECT, 208
 - ROOTINO, 208
- fsfd
 - mkfs.c, 288
- fsfull
 - usertests.c, 476
- FSSIZE
 - param.h, 312
- fstat
 - user.h, 455
- ftable
 - file.c, 176
- g
 - segdesc, 60
- gatedesc, 26
 - args, 27
 - cs, 27
 - dpl, 27
 - off_15_0, 27
 - off_31_16, 27
 - p, 27
 - rsv1, 28
 - s, 28
 - type, 28
- gdt
 - cpu, 15
- getcallerpcs
 - defs.h, 128
 - spinlock.c, 371
- getcmd
 - sh.c, 352
- getpid
 - user.h, 455
- getprocs
 - defs.h, 128
 - user.h, 455
- gets
 - ulib.c, 441
 - user.h, 456
- gettoken
 - sh.c, 353
- grep
 - grep.c, 210
- grep.c, 209
 - buf, 212
 - grep, 210
 - main, 210

- match, 211
- matchhere, 211
- matchstar, 211
- grep.d, 213
- growproc
 - defs.h, 128
 - proc.c, 329
- gs
 - taskstate, 71
 - trapframe, 78
- halt
 - user.h, 456
- havedisk1
 - ide.c, 218
- head
 - bio.c, 85
- Header
 - umalloc.c, 446
- header, 28
 - ptr, 29
 - s, 29
 - size, 29
 - x, 29
- holding
 - defs.h, 129
 - spinlock.c, 371
- holdingsleep
 - defs.h, 129
 - sleeplock.c, 368
- hour
 - rtcd, 56
- HOURS
 - lapic.c, 248
- ialloc
 - defs.h, 129
 - fs.c, 189
 - mkfs.c, 283
- iappend
 - mkfs.c, 283
- IBLOCK
 - fs.h, 207
- icache
 - fs.c, 197
- ICRHI
 - lapic.c, 248
- ICRLO
 - lapic.c, 248
- ID
 - lapic.c, 248
- ide.c, 214
 - havedisk1, 218
 - IDE_BSY, 214
 - IDE_CMD_RDMUL, 215
 - IDE_CMD_READ, 215
 - IDE_CMD_WRITE, 215
 - IDE_CMD_WRMUL, 215
 - IDE_DF, 215
 - IDE_DRDY, 215
 - IDE_ERR, 216
 - ideinit, 216
 - ideintr, 216
 - idelock, 218
 - idequeue, 219
 - iderw, 217
 - idestart, 217
 - idewait, 218
 - SECTOR_SIZE, 216
- ide.d, 221
- IDE_BSY
 - ide.c, 214
- IDE_CMD_RDMUL
 - ide.c, 215
- IDE_CMD_READ
 - ide.c, 215
- IDE_CMD_WRITE
 - ide.c, 215
- IDE_CMD_WRMUL
 - ide.c, 215
- IDE_DF
 - ide.c, 215
- IDE_DRDY
 - ide.c, 215
- IDE_ERR
 - ide.c, 216
- ideinit
 - defs.h, 130
 - ide.c, 216
 - memide.c, 275
- ideintr
 - defs.h, 130
 - ide.c, 216
 - memide.c, 275
- idelock
 - ide.c, 218
- idequeue
 - ide.c, 219
- iderw
 - defs.h, 131
 - ide.c, 217
 - memide.c, 276
- idestart
 - ide.c, 217
- idewait
 - ide.c, 218
- idt
 - trap.c, 426
- idtinit
 - defs.h, 131
 - trap.c, 425
- idup
 - defs.h, 132
 - fs.c, 189
- iget
 - fs.c, 190
- iinit

- defs.h, [132](#)
 - fs.c, [190](#)
- ilock
 - defs.h, [132](#)
 - fs.c, [191](#)
- imcrp
 - mp, [39](#)
- inb
 - x86.h, [531](#)
- INIT
 - lapic.c, [249](#)
- init.c, [221](#)
 - argv, [222](#)
 - main, [221](#)
- init.d, [223](#)
- initcode.d, [223](#)
- initlock
 - defs.h, [133](#)
 - spinlock.c, [372](#)
- initlog
 - defs.h, [133](#)
 - log.c, [262](#)
- initproc
 - proc.c, [335](#)
- initsleeplock
 - defs.h, [133](#)
 - sleeplock.c, [368](#)
- initvm
 - defs.h, [134](#)
 - vm.c, [518](#)
- ino
 - stat, [64](#)
- inode, [30](#)
 - addrs, [30](#)
 - dev, [30](#)
 - fs.c, [198](#)
 - inum, [30](#)
 - lock, [31](#)
 - major, [31](#)
 - minor, [31](#)
 - nlink, [31](#)
 - ref, [31](#)
 - size, [32](#)
 - type, [32](#)
 - valid, [32](#)
- inodestart
 - superblock, [66](#)
- input
 - console.c, [102](#)
- INPUT_BUF
 - console.c, [96](#)
- insl
 - x86.h, [531](#)
- install_trans
 - log.c, [262](#)
- INT_ACTIVELOW
 - ioapic.c, [224](#)
- INT_DISABLED
 - ioapic.c, [224](#)
- INT_LEVEL
 - ioapic.c, [224](#)
- INT_LOGICAL
 - ioapic.c, [224](#)
- intena
 - cpu, [15](#)
- inum
 - dirent, [19](#)
 - inode, [30](#)
- IO_PIC1
 - picirq.c, [315](#)
- IO_PIC2
 - picirq.c, [315](#)
- IOAPIC
 - ioapic.c, [224](#)
- ioapic, [32](#)
 - data, [33](#)
 - ioapic.c, [226](#)
 - pad, [33](#)
 - reg, [33](#)
- ioapic.c, [223](#)
 - INT_ACTIVELOW, [224](#)
 - INT_DISABLED, [224](#)
 - INT_LEVEL, [224](#)
 - INT_LOGICAL, [224](#)
 - IOAPIC, [224](#)
 - ioapic, [226](#)
 - ioapicenable, [225](#)
 - ioapicinit, [225](#)
 - ioapicread, [226](#)
 - ioapicwrite, [226](#)
 - REG_ID, [225](#)
 - REG_TABLE, [225](#)
 - REG_VER, [225](#)
- ioapic.d, [228](#)
- ioapicenable
 - defs.h, [134](#)
 - ioapic.c, [225](#)
- ioapicid
 - defs.h, [160](#)
 - mp.c, [306](#)
- ioapicinit
 - defs.h, [134](#)
 - ioapic.c, [225](#)
- ioapicread
 - ioapic.c, [226](#)
- ioapicwrite
 - ioapic.c, [226](#)
- iomb
 - taskstate, [71](#)
- ip
 - file, [25](#)
- IPB
 - fs.h, [207](#)
- iput
 - defs.h, [135](#)
 - fs.c, [191](#)

- iputtest
 - usertests.c, 476
- iref
 - usertests.c, 477
- IRQ_COM1
 - traps.h, 430
- IRQ_ERROR
 - traps.h, 430
- IRQ_IDE
 - traps.h, 430
- IRQ_KBD
 - traps.h, 430
- IRQ_SPURIOUS
 - traps.h, 430
- IRQ_TIMER
 - traps.h, 431
- isdirempty
 - sysfile.c, 406
- ismp
 - defs.h, 161
- itrunc
 - fs.c, 192
- iunlock
 - defs.h, 135
 - fs.c, 192
- iunlockput
 - defs.h, 135
 - fs.c, 193
- iupdate
 - defs.h, 136
 - fs.c, 193
- kalloc
 - defs.h, 136
 - kalloc.c, 229
- kalloc.c, 228
 - end, 230
 - freelist, 230
 - freerange, 228
 - kalloc, 229
 - kfree, 229
 - kinit1, 229
 - kinit2, 230
 - kmem, 230
 - lock, 230
 - use_lock, 231
- kalloc.d, 232
- kbd.c, 232
 - kbdgetc, 233
 - kbdintr, 233
- kbd.d, 234
- kbd.h, 234
 - ALT, 235
 - C, 235
 - CAPSLOCK, 235
 - CTL, 236
 - ctlmap, 239
 - E0ESC, 236
 - KBDATAP, 236
 - KBS_DIB, 236
 - KBSTATP, 236
 - KEY_DEL, 236
 - KEY_DN, 237
 - KEY_END, 237
 - KEY_HOME, 237
 - KEY_INS, 237
 - KEY_LF, 237
 - KEY_PGDN, 237
 - KEY_PGUP, 238
 - KEY_RT, 238
 - KEY_UP, 238
 - NO, 238
 - normalmap, 239
 - NUMLOCK, 238
 - SCROLLLOCK, 238
 - SHIFT, 239
 - shiftcode, 240
 - shiftmap, 240
 - togglecode, 241
- KBDATAP
 - kbd.h, 236
- kbdgetc
 - kbd.c, 233
- kbdintr
 - defs.h, 136
 - kbd.c, 233
- KBS_DIB
 - kbd.h, 236
- KBSTATP
 - kbd.h, 236
- KERNBASE
 - memlayout.h, 278
- KERNLINK
 - memlayout.h, 279
- KEY_DEL
 - kbd.h, 236
- KEY_DN
 - kbd.h, 237
- KEY_END
 - kbd.h, 237
- KEY_HOME
 - kbd.h, 237
- KEY_INS
 - kbd.h, 237
- KEY_LF
 - kbd.h, 237
- KEY_PGDN
 - kbd.h, 237
- KEY_PGUP
 - kbd.h, 238
- KEY_RT
 - kbd.h, 238
- KEY_UP
 - kbd.h, 238
- kfree
 - defs.h, 137
 - kalloc.c, 229

- kill
 - defs.h, 137
 - proc.c, 329
 - user.h, 456
- kill.c, 243
 - main, 243
- kill.d, 244
- killed
 - proc, 50
- kinit1
 - defs.h, 137
 - kalloc.c, 229
- kinit2
 - defs.h, 138
 - kalloc.c, 230
- kmap, 33
 - perm, 34
 - phys_end, 34
 - phys_start, 34
 - virt, 34
 - vm.c, 522
- kmem
 - kalloc.c, 230
- kpgdir
 - vm.c, 523
- kstack
 - proc, 50
- KSTACKSIZE
 - param.h, 313
- kvmalloc
 - defs.h, 138
 - vm.c, 519
- lapic
 - defs.h, 161
 - lapic.c, 256
- lapic.c, 244
 - ASSERT, 245
 - BCAST, 245
 - BUSY, 245
 - CMOS_PORT, 246
 - cmos_read, 252
 - CMOS_RETURN, 246
 - CMOS_STATA, 246
 - CMOS_STATB, 246
 - CMOS_UIP, 246
 - cmostime, 252
 - CONV, 246
 - DAY, 247
 - DEASSERT, 247
 - DELIVS, 247
 - ENABLE, 247
 - EOI, 247
 - ERROR, 247
 - ESR, 248
 - fill_rtcdade, 253
 - FIXED, 248
 - HOURS, 248
 - ICRHI, 248
 - ICRLO, 248
 - ID, 248
 - INIT, 249
 - lapic, 256
 - lapiceoi, 253
 - lapicid, 253
 - lapicinit, 254
 - lapicstartap, 254
 - lapicw, 255
 - LEVEL, 249
 - LINT0, 249
 - LINT1, 249
 - MASKED, 249
 - microdelay, 255
 - MINS, 249
 - MONTH, 250
 - PCINT, 250
 - PERIODIC, 250
 - SECS, 250
 - STARTUP, 250
 - SVR, 250
 - TCCR, 251
 - TDCR, 251
 - TICR, 251
 - TIMER, 251
 - TPR, 251
 - VER, 251
 - X1, 252
 - YEAR, 252
- lapic.d, 259
- lapicaddr
 - mpconf, 41
- lapiceoi
 - defs.h, 138
 - lapic.c, 253
- lapicid
 - defs.h, 138
 - lapic.c, 253
- lapicinit
 - defs.h, 139
 - lapic.c, 254
- lapicstartap
 - defs.h, 139
 - lapic.c, 254
- lapicw
 - lapic.c, 255
- lcr3
 - x86.h, 531
- ldt
 - taskstate, 71
- left
 - listcmd, 35
 - pipecmd, 48
- length
 - mp, 39
 - mpconf, 41
- LEVEL
 - lapic.c, 249

- lgdt
 - x86.h, 532
- lh
 - log, 36
- lidt
 - x86.h, 532
- lim_15_0
 - segdesc, 60
- lim_19_16
 - segdesc, 60
- link
 - taskstate, 71
 - user.h, 456
- linktest
 - usertests.c, 477
- linkunlink
 - usertests.c, 478
- LINT0
 - lapic.c, 249
- LINT1
 - lapic.c, 249
- LIST
 - sh.c, 351
- listcmd, 35
 - left, 35
 - right, 35
 - sh.c, 353
 - type, 35
- lk
 - sleeplock, 61
- ln.c, 259
 - main, 259
- ln.d, 260
- loadgs
 - x86.h, 532
- loadvm
 - defs.h, 140
 - vm.c, 519
- lock
 - bio.c, 85
 - buf, 11
 - console.c, 102
 - file.c, 177
 - fs.c, 198
 - inode, 31
 - kalloc.c, 230
 - log, 36
 - pipe, 46
 - proc.c, 335
 - sysproc.c, 421
- locked
 - sleeplock, 61
 - spinlock, 63
- locking
 - console.c, 102
- log, 36
 - committing, 36
 - dev, 36
 - lh, 36
 - lock, 36
 - log.c, 264
 - outstanding, 37
 - size, 37
 - start, 37
- log.c, 260
 - begin_op, 260
 - commit, 261
 - end_op, 261
 - initlog, 262
 - install_trans, 262
 - log, 264
 - log_write, 262
 - read_head, 263
 - recover_from_log, 263
 - write_head, 263
 - write_log, 264
- log.d, 267
- log_write
 - defs.h, 140
 - log.c, 262
- logheader, 37
 - block, 38
 - n, 38
- LOGSIZE
 - param.h, 313
- logstart
 - superblock, 66
- ls
 - ls.c, 268
- ls.c, 268
 - fmtname, 268
 - ls, 268
 - main, 269
- ls.d, 271
- ltr
 - x86.h, 533
- machine
 - elfhdr, 21
- magic
 - elfhdr, 21
- main
 - cat.c, 93
 - echo.c, 164
 - exit.c, 171
 - foo.c, 181
 - forktest.c, 183
 - grep.c, 210
 - init.c, 221
 - kill.c, 243
 - ln.c, 259
 - ls.c, 269
 - mkdir.c, 280
 - mkfs.c, 284
 - nice.c, 311
 - ps.c, 345
 - pstree.c, 347

- rm.c, 348
- sh.c, 354
- shutdown.c, 367
- stressfs.c, 377
- usertests.c, 479
- wc.c, 528
- zombie.c, 538
- main.c, 271
 - __attribute__, 271
 - entrypgdir, 273
 - mpenter, 271
 - mpmain, 272
 - startothers, 272
- main.d, 275
- major
 - dinode, 18
 - inode, 31
- malloc
 - umalloc.c, 446
 - user.h, 457
- mappages
 - vm.c, 519
- MASKED
 - lapic.c, 249
- match
 - grep.c, 211
- matchhere
 - grep.c, 211
- matchstar
 - grep.c, 211
- MAXARG
 - param.h, 313
- MAXARGS
 - sh.c, 351
- MAXFILE
 - fs.h, 208
- MAXOPBLOCKS
 - param.h, 313
- MAXPROC
 - pstree.c, 346
- mem
 - usertests.c, 480
- memcmp
 - defs.h, 141
 - string.c, 379
- memcpy
 - string.c, 379
- memdisk
 - memide.c, 277
- memide.c, 275
 - _binary_fs_img_size, 276
 - _binary_fs_img_start, 276
 - disksize, 277
 - ideinit, 275
 - ideintr, 275
 - iderw, 276
 - memdisk, 277
- memlayout.h, 278
- DEVSPACE, 278
- EXTMEM, 278
- KERNBASE, 278
- KERNLINK, 279
- P2V, 279
- P2V_WO, 279
- PHYSTOP, 279
- V2P, 279
- V2P_WO, 279
- memmove
 - defs.h, 141
 - string.c, 379
 - ulib.c, 441
 - user.h, 457
- memset
 - defs.h, 142
 - string.c, 380
 - ulib.c, 441
 - user.h, 457
- memset
 - proghdr, 53
- microdelay
 - defs.h, 142
 - lapic.c, 255
- min
 - fs.c, 186
 - mkfs.c, 282
- minor
 - dinode, 18
 - inode, 31
- MINS
 - lapic.c, 249
- minute
 - rtcdate, 57
- mkdir
 - user.h, 458
- mkdir.c, 280
 - main, 280
- mkdir.d, 281
- mkfs.c, 281
 - ballocc, 283
 - freeblock, 287
 - freeinode, 288
 - fsfd, 288
 - iallocc, 283
 - iappend, 283
 - main, 284
 - min, 282
 - nbitmap, 288
 - nblocks, 288
 - ninodeblocks, 288
 - NINODES, 282
 - nlog, 289
 - nmeta, 289
 - rinode, 286
 - rsect, 286
 - sb, 289
 - stat, 282

- static_assert, [283](#)
- winode, [286](#)
- wsect, [286](#)
- xint, [287](#)
- xshort, [287](#)
- zeroes, [289](#)
- mknod
 - user.h, [458](#)
- mmu.h, [293](#)
 - CR0_PE, [294](#)
 - CR0_PG, [294](#)
 - CR0_WP, [294](#)
 - CR4_PSE, [295](#)
 - DPL_USER, [295](#)
 - FL_IF, [295](#)
 - NPENTRIES, [295](#)
 - NPTENTRIES, [295](#)
 - NSEGS, [295](#)
 - PDX, [296](#)
 - PDXSHIFT, [296](#)
 - PGADDR, [296](#)
 - PGROUNDNDOWN, [296](#)
 - PGROUNDUP, [296](#)
 - PGSIZE, [296](#)
 - PTE_ADDR, [297](#)
 - PTE_FLAGS, [297](#)
 - PTE_P, [297](#)
 - PTE_PS, [297](#)
 - pte_t, [301](#)
 - PTE_U, [297](#)
 - PTE_W, [297](#)
 - PTX, [298](#)
 - PTXSHIFT, [298](#)
 - SEG, [298](#)
 - SEG16, [298](#)
 - SEG_KCODE, [298](#)
 - SEG_KDATA, [299](#)
 - SEG_TSS, [299](#)
 - SEG_UCODE, [299](#)
 - SEG_UDATA, [299](#)
 - SETGATE, [299](#)
 - STA_R, [300](#)
 - STA_W, [300](#)
 - STA_X, [300](#)
 - STS_IG32, [300](#)
 - STS_T32A, [300](#)
 - STS_TG32, [301](#)
- mode
 - redircmd, [55](#)
- MONTH
 - lapic.c, [250](#)
- month
 - rtcdate, [57](#)
- morecore
 - umalloc.c, [447](#)
- mp, [38](#)
 - checksum, [39](#)
 - imcrp, [39](#)
 - length, [39](#)
 - physaddr, [39](#)
 - reserved, [39](#)
 - signature, [40](#)
 - specrev, [40](#)
 - type, [40](#)
- mp.c, [303](#)
 - cpus, [306](#)
 - ioapicid, [306](#)
 - mpconfig, [304](#)
 - mpinit, [304](#)
 - mpsearch, [305](#)
 - mpsearch1, [305](#)
 - ncpu, [306](#)
 - sum, [306](#)
- mp.d, [308](#)
- mp.h, [309](#)
 - MPBOOT, [309](#)
 - MPBUS, [309](#)
 - MPIOAPIC, [309](#)
 - MPIOINTR, [309](#)
 - MPLINTR, [310](#)
 - MPPROC, [310](#)
- MPBOOT
 - mp.h, [309](#)
- MPBUS
 - mp.h, [309](#)
- mpconf, [40](#)
 - checksum, [41](#)
 - entry, [41](#)
 - lapicaddr, [41](#)
 - length, [41](#)
 - oemlength, [41](#)
 - oemtable, [42](#)
 - product, [42](#)
 - reserved, [42](#)
 - signature, [42](#)
 - version, [42](#)
 - xchecksum, [42](#)
 - xlength, [43](#)
- mpconfig
 - mp.c, [304](#)
- mpenter
 - main.c, [271](#)
- mpinit
 - defs.h, [142](#)
 - mp.c, [304](#)
- MPIOAPIC
 - mp.h, [309](#)
- mpioapic, [43](#)
 - addr, [43](#)
 - apicno, [43](#)
 - flags, [44](#)
 - type, [44](#)
 - version, [44](#)
- MPIOINTR
 - mp.h, [309](#)
- MPLINTR

- mp.h, 310
- mpmain
 - main.c, 272
- MPPROC
 - mp.h, 310
- mpproc, 44
 - apicid, 45
 - feature, 45
 - flags, 45
 - reserved, 45
 - signature, 45
 - type, 45
 - version, 46
- mpsearch
 - mp.c, 305
- mpsearch1
 - mp.c, 305
- mycpu
 - defs.h, 143
 - proc.c, 330
- myproc
 - defs.h, 143
 - proc.c, 330
- N
 - forktest.c, 182
- n
 - logheader, 38
- name
 - dirent, 19
 - proc, 50
 - sleeplock, 62
 - spinlock, 63
 - uproc, 80
 - usertests.c, 492
- namecmp
 - defs.h, 144
 - fs.c, 193
- namei
 - defs.h, 144
 - fs.c, 194
- nameiparent
 - defs.h, 144
 - fs.c, 194
- namex
 - fs.c, 194
- nbitmap
 - mkfs.c, 288
- nblocks
 - mkfs.c, 288
 - superblock, 66
- NBUF
 - param.h, 313
- ncli
 - cpu, 15
- NCPU
 - param.h, 313
- ncpu
 - mp.c, 306
- proc.h, 344
- NDEV
 - param.h, 314
- NDIRECT
 - fs.h, 208
- NELEM
 - defs.h, 110
- next
 - buf, 11
 - run, 58
- nextpid
 - proc.c, 335
- NFILE
 - param.h, 314
- nice.c, 311
 - main, 311
- nice.d, 312
- NINDIRECT
 - fs.h, 208
- NINODE
 - param.h, 314
- ninodeblocks
 - mkfs.c, 288
- NINODES
 - mkfs.c, 282
- ninodes
 - superblock, 66
- nlink
 - dinode, 18
 - inode, 31
 - stat, 64
- nlog
 - mkfs.c, 289
 - superblock, 66
- nmeta
 - mkfs.c, 289
- NO
 - kbd.h, 238
- NOFILE
 - param.h, 314
- normalmap
 - kbd.h, 239
- NPENTRIES
 - mmu.h, 295
- NPROC
 - param.h, 314
- NPTENTRIES
 - mmu.h, 295
- nread
 - pipe, 47
- NSEGS
 - mmu.h, 295
- nulterminate
 - sh.c, 354
- NUMLOCK
 - kbd.h, 238
- nwrite
 - pipe, 47

- O_CREATE
 - fcntl.h, [172](#)
- O_RDONLY
 - fcntl.h, [172](#)
- O_RDWR
 - fcntl.h, [172](#)
- O_WRONLY
 - fcntl.h, [172](#)
- oemlength
 - mpconf, [41](#)
- oemtable
 - mpconf, [42](#)
- oesp
 - trapframe, [78](#)
- off
 - file, [25](#)
 - proghdr, [53](#)
- off_15_0
 - gatedesc, [27](#)
- off_31_16
 - gatedesc, [27](#)
- ofile
 - proc, [50](#)
- open
 - user.h, [458](#)
- openiputtest
 - usertests.c, [480](#)
- opentest
 - usertests.c, [481](#)
- outb
 - x86.h, [533](#)
- outsl
 - x86.h, [533](#)
- outstanding
 - log, [37](#)
- outw
 - x86.h, [533](#)
- p
 - gatedesc, [27](#)
 - segdesc, [60](#)
- P2V
 - memlayout.h, [279](#)
- P2V_WO
 - memlayout.h, [279](#)
- pad
 - ioapic, [33](#)
- padding1
 - taskstate, [72](#)
 - trapframe, [78](#)
- padding10
 - taskstate, [72](#)
- padding2
 - taskstate, [72](#)
 - trapframe, [78](#)
- padding3
 - taskstate, [72](#)
 - trapframe, [79](#)
- padding4
 - taskstate, [72](#)
 - trapframe, [79](#)
- padding5
 - taskstate, [72](#)
 - trapframe, [79](#)
- padding6
 - taskstate, [73](#)
 - trapframe, [79](#)
- padding7
 - taskstate, [73](#)
- padding8
 - taskstate, [73](#)
- padding9
 - taskstate, [73](#)
- paddr
 - proghdr, [53](#)
- panic
 - console.c, [100](#)
 - defs.h, [144](#)
 - sh.c, [355](#)
- panicked
 - console.c, [102](#)
- param.h, [312](#)
 - FSSIZE, [312](#)
 - KSTACKSIZE, [313](#)
 - LOGSIZE, [313](#)
 - MAXARG, [313](#)
 - MAXOPBLOCKS, [313](#)
 - NBUF, [313](#)
 - NCPU, [313](#)
 - NDEV, [314](#)
 - NFILE, [314](#)
 - NINODE, [314](#)
 - NOFILE, [314](#)
 - NPROC, [314](#)
 - ROOTDEV, [314](#)
- parent
 - proc, [50](#)
- parseblock
 - sh.c, [355](#)
- parsecmd
 - sh.c, [355](#)
- parseexec
 - sh.c, [356](#)
- parseline
 - sh.c, [356](#)
- parsepipe
 - sh.c, [357](#)
- parseredirs
 - sh.c, [357](#)
- PCINT
 - lapic.c, [250](#)
- pcs
 - spinlock, [63](#)
- pde_t
 - types.h, [435](#)
- PDX
 - mmu.h, [296](#)

- PDXSHIFT
 - mmu.h, 296
- peek
 - sh.c, 358
- PERIODIC
 - lapic.c, 250
- perm
 - kmap, 34
- PGADDR
 - mmu.h, 296
- pgdir
 - proc, 51
- PGROUNDDOWN
 - mmu.h, 296
- PGROUNDUP
 - mmu.h, 296
- PGSIZE
 - mmu.h, 296
- phentsize
 - elfhdr, 21
- phnum
 - elfhdr, 21
- phoff
 - elfhdr, 21
- phys_end
 - kmap, 34
- phys_start
 - kmap, 34
- physaddr
 - mp, 39
- PHYSTOP
 - memlayout.h, 279
- picenable
 - defs.h, 145
- picinit
 - defs.h, 145
 - picirq.c, 316
- picirq.c, 315
 - IO_PIC1, 315
 - IO_PIC2, 315
 - picinit, 316
- picirq.d, 316
- pid
 - proc, 51
 - sleeplock, 62
 - uproc, 80
- pinit
 - defs.h, 145
 - proc.c, 330
- PIPE
 - sh.c, 351
- pipe, 46
 - data, 46
 - file, 25
 - lock, 46
 - nread, 47
 - nwrite, 47
 - readopen, 47
 - user.h, 458
 - writeopen, 47
- pipe.c, 317
 - pipealloc, 317
 - pipeclose, 318
 - piperead, 318
 - PIPESIZE, 317
 - pipewrite, 319
- pipe.d, 321
- pipe1
 - usertests.c, 481
- pipealloc
 - defs.h, 145
 - pipe.c, 317
- pipeclose
 - defs.h, 146
 - pipe.c, 318
- pipecmd, 48
 - left, 48
 - right, 48
 - sh.c, 358
 - type, 48
- piperead
 - defs.h, 146
 - pipe.c, 318
- PIPESIZE
 - pipe.c, 317
- pipewrite
 - defs.h, 147
 - pipe.c, 319
- popcli
 - defs.h, 147
 - spinlock.c, 372
- ppid
 - uproc, 80
- preempt
 - usertests.c, 482
- prev
 - buf, 11
- printf
 - forktest.c, 183
 - printf.c, 322
 - user.h, 459
- printf.c, 321
 - printf, 322
 - printint, 322
 - putc, 323
- printf.d, 324
- printint
 - console.c, 101
 - printf.c, 322
- priority
 - proc, 51
- proc, 49
 - chan, 49
 - context, 49
 - cpu, 15
 - cwd, 49

- killed, [50](#)
- kstack, [50](#)
- name, [50](#)
- ofile, [50](#)
- parent, [50](#)
- pgdir, [51](#)
- pid, [51](#)
- priority, [51](#)
- proc.c, [336](#)
- state, [51](#)
- sysproc.c, [422](#)
- sz, [51](#)
- tf, [52](#)
- proc.c, [324](#)
 - allocproc, [325](#)
 - chpr, [326](#)
 - cps, [326](#)
 - cpuid, [327](#)
 - exit, [327](#)
 - fork, [328](#)
 - forkret, [329](#)
 - growproc, [329](#)
 - initproc, [335](#)
 - kill, [329](#)
 - lock, [335](#)
 - mycpu, [330](#)
 - myproc, [330](#)
 - nextpid, [335](#)
 - pinit, [330](#)
 - proc, [336](#)
 - procdump, [331](#)
 - ptable, [336](#)
 - sched, [331](#)
 - scheduler, [331](#)
 - sleep, [332](#)
 - trapret, [333](#)
 - userinit, [333](#)
 - wait, [333](#)
 - wakeup, [334](#)
 - wakeup1, [334](#)
 - yield, [335](#)
- proc.d, [343](#)
- proc.h, [343](#)
 - cpus, [344](#)
 - EMBRYO, [344](#)
 - ncpu, [344](#)
 - procstate, [343](#)
 - RUNNABLE, [344](#)
 - RUNNING, [344](#)
 - SLEEPING, [344](#)
 - UNUSED, [344](#)
 - ZOMBIE, [344](#)
- procdump
 - defs.h, [148](#)
 - proc.c, [331](#)
- procstate
 - proc.h, [343](#)
- product
 - mpconf, [42](#)
- proghdr, [52](#)
 - align, [52](#)
 - filesz, [53](#)
 - flags, [53](#)
 - memsz, [53](#)
 - off, [53](#)
 - paddr, [53](#)
 - type, [54](#)
 - vaddr, [54](#)
- ps.c, [345](#)
 - main, [345](#)
- ps.d, [346](#)
- pstree.c, [346](#)
 - main, [347](#)
 - MAXPROC, [346](#)
- pstree.d, [348](#)
- ptable
 - proc.c, [336](#)
 - sysproc.c, [422](#)
- PTE_ADDR
 - mmu.h, [297](#)
- PTE_FLAGS
 - mmu.h, [297](#)
- PTE_P
 - mmu.h, [297](#)
- PTE_PS
 - mmu.h, [297](#)
- pte_t
 - mmu.h, [301](#)
- PTE_U
 - mmu.h, [297](#)
- PTE_W
 - mmu.h, [297](#)
- ptr
 - header, [29](#)
- PTX
 - mmu.h, [298](#)
- PTXSHIFT
 - mmu.h, [298](#)
- pushcli
 - defs.h, [148](#)
 - spinlock.c, [372](#)
- putc
 - printf.c, [323](#)
- qnext
 - buf, [11](#)
- r
 - console.c, [103](#)
- rand
 - usertests.c, [483](#)
- randstate
 - usertests.c, [493](#)
- rcr2
 - x86.h, [534](#)
- read
 - devsw, [17](#)

- user.h, 459
- read_head
 - log.c, 263
- readable
 - file, 25
- readeflags
 - x86.h, 534
- readi
 - defs.h, 148
 - fs.c, 195
- README.md, 348
- readopen
 - pipe, 47
- readsb
 - defs.h, 149
 - fs.c, 196
- readsect
 - bootmain.c, 89
- readseg
 - bootmain.c, 89
- recover_from_log
 - log.c, 263
- REDIR
 - sh.c, 351
- redircmd, 54
 - cmd, 54
 - efile, 55
 - fd, 55
 - file, 55
 - mode, 55
 - sh.c, 358
 - type, 55
- ref
 - file, 25
 - inode, 31
- refcnt
 - buf, 12
- reg
 - ioapic, 33
- REG_ID
 - ioapic.c, 225
- REG_TABLE
 - ioapic.c, 225
- REG_VER
 - ioapic.c, 225
- release
 - defs.h, 149
 - spinlock.c, 373
- releasesleep
 - defs.h, 150
 - sleeplock.c, 369
- reserved
 - mp, 39
 - mpconf, 42
 - mpproc, 45
- right
 - listcmd, 35
 - pipecmd, 48
- rinode
 - mkfs.c, 286
- rm.c, 348
 - main, 348
- rm.d, 349
- rmdot
 - usertests.c, 483
- ROOTDEV
 - param.h, 314
- ROOTINO
 - fs.h, 208
- rsect
 - mkfs.c, 286
- rsv1
 - gatedesc, 28
 - segdesc, 60
- RTC_ADDR
 - usertests.c, 464
- RTC_DATA
 - usertests.c, 465
- rtcdate, 56
 - day, 56
 - hour, 56
 - minute, 57
 - month, 57
 - second, 57
 - year, 57
- run, 58
 - next, 58
- runcmd
 - sh.c, 359
- RUNNABLE
 - proc.h, 344
- RUNNING
 - proc.h, 344
- s
 - gatedesc, 28
 - header, 29
 - segdesc, 60
- safestrcpy
 - defs.h, 150
 - string.c, 380
- sb
 - fs.c, 198
 - mkfs.c, 289
- sbrk
 - user.h, 459
- sbrktest
 - usertests.c, 483
- sched
 - defs.h, 150
 - proc.c, 331
- scheduler
 - cpu, 15
 - defs.h, 151
 - proc.c, 331
- SCROLLLOCK
 - kbd.h, 238

- second
 - rtcdatc, 57
- SECS
 - lapic.c, 250
- SECTOR_SIZE
 - ide.c, 216
- SECTSIZE
 - bootmain.c, 88
- SEG
 - mmu.h, 298
- SEG16
 - mmu.h, 298
- SEG_ASM
 - asm.h, 81
- SEG_KCODE
 - mmu.h, 298
- SEG_KDATA
 - mmu.h, 299
- SEG_NULLASM
 - asm.h, 81
- SEG_TSS
 - mmu.h, 299
- SEG_UCODE
 - mmu.h, 299
- SEG_UDATA
 - mmu.h, 299
- segdesc, 58
 - avl, 59
 - base_15_0, 59
 - base_23_16, 59
 - base_31_24, 59
 - db, 59
 - dpl, 59
 - g, 60
 - lim_15_0, 60
 - lim_19_16, 60
 - p, 60
 - rsv1, 60
 - s, 60
 - type, 61
- seginit
 - defs.h, 151
 - vm.c, 520
- SETGATE
 - mmu.h, 299
- setproc
 - defs.h, 152
- setupkvm
 - defs.h, 152
 - vm.c, 520
- sh.c, 349
 - BACK, 350
 - backcmd, 352
 - EXEC, 351
 - execcmd, 352
 - fork1, 352
 - getcnd, 352
 - gettoken, 353
 - LIST, 351
 - listcmd, 353
 - main, 354
 - MAXARGS, 351
 - nulterminate, 354
 - panic, 355
 - parseblock, 355
 - parsecmd, 355
 - parseexec, 356
 - parseline, 356
 - parsepipe, 357
 - parseredirs, 357
 - peek, 358
 - PIPE, 351
 - pipecmd, 358
 - REDIR, 351
 - redircmd, 358
 - runcmd, 359
 - symbols, 360
 - whitespace, 360
- sh.d, 366
- sharedfd
 - usertests.c, 485
- shentsize
 - elfhdr, 22
- SHIFT
 - kbd.h, 239
- shiftcode
 - kbd.h, 240
- shiftmap
 - kbd.h, 240
- shnum
 - elfhdr, 22
- shoff
 - elfhdr, 22
- shstrndx
 - elfhdr, 22
- shutdown.c, 366
 - main, 367
- shutdown.d, 367
- signature
 - mp, 40
 - mpconf, 42
 - mpproc, 45
- size
 - dinode, 18
 - header, 29
 - inode, 32
 - log, 37
 - stat, 64
 - superblock, 67
- skipelem
 - fs.c, 196
- sleep
 - defs.h, 152
 - proc.c, 332
 - user.h, 459
- SLEEPING

- proc.h, 344
- sleeplock, 61
 - lk, 61
 - locked, 61
 - name, 62
 - pid, 62
- sleeplock.c, 368
 - acquiresleep, 368
 - holdingsleep, 368
 - initsleeplock, 368
 - releasesleep, 369
- sleeplock.d, 370
- sleeplock.h, 370
- specrev
 - mp, 40
- spinlock, 62
 - cpu, 63
 - locked, 63
 - name, 63
 - pcs, 63
- spinlock.c, 370
 - acquire, 371
 - getcallerpcs, 371
 - holding, 371
 - initlock, 372
 - popcli, 372
 - pushcli, 372
 - release, 373
- spinlock.d, 375
- spinlock.h, 375
- ss
 - taskstate, 73
 - trapframe, 79
- ss0
 - taskstate, 73
- ss1
 - taskstate, 74
- ss2
 - taskstate, 74
- STA_R
 - asm.h, 81
 - mmu.h, 300
- STA_W
 - asm.h, 82
 - mmu.h, 300
- STA_X
 - asm.h, 82
 - mmu.h, 300
- start
 - log, 37
- started
 - cpu, 16
- startothers
 - main.c, 272
- STARTUP
 - lapic.c, 250
- stat, 63
 - dev, 64
 - ino, 64
 - mkfs.c, 282
 - nlink, 64
 - size, 64
 - type, 65
 - ulib.c, 442
 - user.h, 459
- stat.h, 375
 - T_DEV, 376
 - T_DIR, 376
 - T_FILE, 376
- state
 - proc, 51
- stati
 - defs.h, 153
 - fs.c, 196
- static_assert
 - mkfs.c, 283
- stdout
 - usertests.c, 493
- sti
 - x86.h, 534
- stosb
 - x86.h, 534
- stosl
 - x86.h, 535
- strchr
 - ulib.c, 442
 - user.h, 460
- strcmp
 - ulib.c, 442
 - user.h, 460
- strcpy
 - ulib.c, 443
 - user.h, 460
- stressfs.c, 377
 - main, 377
- stressfs.d, 378
- string.c, 378
 - memcmp, 379
 - memcpy, 379
 - memmove, 379
 - memset, 380
 - safestrncpy, 380
 - strlen, 380
 - strncmp, 381
 - strncpy, 381
- string.d, 383
- strlen
 - defs.h, 153
 - string.c, 380
 - ulib.c, 443
 - user.h, 461
- strncmp
 - defs.h, 154
 - string.c, 381
- strncpy
 - defs.h, 154

- string.c, 381
- STS_IG32
 - mmu.h, 300
- STS_T32A
 - mmu.h, 300
- STS_TG32
 - mmu.h, 301
- subdir
 - usertests.c, 486
- sum
 - mp.c, 306
- superblock, 65
 - bmapstart, 65
 - inodestart, 66
 - logstart, 66
 - nblocks, 66
 - ninodes, 66
 - nlog, 66
 - size, 67
- SVR
 - lapic.c, 250
- switchkvm
 - defs.h, 154
 - vm.c, 521
- switchvm
 - defs.h, 155
 - vm.c, 521
- switch
 - defs.h, 155
- symbols
 - sh.c, 360
- SYS_chdir
 - syscall.h, 399
- sys_chdir
 - syscall.c, 386
 - sysfile.c, 406
- SYS_chpr
 - syscall.h, 399
- sys_chpr
 - syscall.c, 386
 - sysproc.c, 418
- SYS_close
 - syscall.h, 399
- sys_close
 - syscall.c, 386
 - sysfile.c, 406
- SYS_cps
 - syscall.h, 399
- sys_cps
 - syscall.c, 387
 - sysproc.c, 418
- SYS_dup
 - syscall.h, 399
- sys_dup
 - syscall.c, 387
 - sysfile.c, 407
- SYS_exec
 - syscall.h, 400
- sys_exec
 - syscall.c, 387
 - sysfile.c, 407
- SYS_exit
 - syscall.h, 400
- sys_exit
 - syscall.c, 388
 - sysproc.c, 418
- SYS_fork
 - syscall.h, 400
- sys_fork
 - syscall.c, 388
 - sysproc.c, 419
- SYS_fstat
 - syscall.h, 400
- sys_fstat
 - syscall.c, 388
 - sysfile.c, 407
- SYS_getpid
 - syscall.h, 400
- sys_getpid
 - syscall.c, 388
 - sysproc.c, 419
- SYS_getprocs
 - syscall.h, 400
- sys_getprocs
 - syscall.c, 389
 - sysproc.c, 419
- SYS_halt
 - syscall.h, 401
- sys_halt
 - syscall.c, 389
 - sysproc.c, 420
- SYS_kill
 - syscall.h, 401
- sys_kill
 - syscall.c, 389
 - sysproc.c, 420
- SYS_link
 - syscall.h, 401
- sys_link
 - syscall.c, 390
 - sysfile.c, 408
- SYS_mkdir
 - syscall.h, 401
- sys_mkdir
 - syscall.c, 390
 - sysfile.c, 408
- SYS_mknod
 - syscall.h, 401
- sys_mknod
 - syscall.c, 391
 - sysfile.c, 409
- SYS_open
 - syscall.h, 401
- sys_open
 - syscall.c, 391
 - sysfile.c, 409

- SYS_pipe
 - syscall.h, 402
- sys_pipe
 - syscall.c, 392
 - sysfile.c, 410
- SYS_read
 - syscall.h, 402
- sys_read
 - syscall.c, 392
 - sysfile.c, 410
- SYS_sbrk
 - syscall.h, 402
- sys_sbrk
 - syscall.c, 393
 - sysproc.c, 420
- SYS_sleep
 - syscall.h, 402
- sys_sleep
 - syscall.c, 393
 - sysproc.c, 420
- SYS_unlink
 - syscall.h, 402
- sys_unlink
 - syscall.c, 393
 - sysfile.c, 411
- SYS_uptime
 - syscall.h, 402
- sys_uptime
 - syscall.c, 394
 - sysproc.c, 421
- SYS_wait
 - syscall.h, 403
- sys_wait
 - syscall.c, 394
 - sysproc.c, 421
- SYS_write
 - syscall.h, 403
- sys_write
 - syscall.c, 395
 - sysfile.c, 411
- syscall
 - defs.h, 155
 - syscall.c, 395
- syscall.c, 383
 - argint, 384
 - argptr, 384
 - argstr, 385
 - fetchint, 385
 - fetchstr, 385
 - sys_chdir, 386
 - sys_chpr, 386
 - sys_close, 386
 - sys_cps, 387
 - sys_dup, 387
 - sys_exec, 387
 - sys_exit, 388
 - sys_fork, 388
 - sys_fstat, 388
 - sys_getpid, 388
 - sys_getprocs, 389
 - sys_halt, 389
 - sys_kill, 389
 - sys_link, 390
 - sys_mkdir, 390
 - sys_mknod, 391
 - sys_open, 391
 - sys_pipe, 392
 - sys_read, 392
 - sys_sbrk, 393
 - sys_sleep, 393
 - sys_unlink, 393
 - sys_uptime, 394
 - sys_wait, 394
 - sys_write, 395
 - syscall, 395
 - syscalls, 395
- syscall.d, 398
- syscall.h, 398
 - SYS_chdir, 399
 - SYS_chpr, 399
 - SYS_close, 399
 - SYS_cps, 399
 - SYS_dup, 399
 - SYS_exec, 400
 - SYS_exit, 400
 - SYS_fork, 400
 - SYS_fstat, 400
 - SYS_getpid, 400
 - SYS_getprocs, 400
 - SYS_halt, 401
 - SYS_kill, 401
 - SYS_link, 401
 - SYS_mkdir, 401
 - SYS_mknod, 401
 - SYS_open, 401
 - SYS_pipe, 402
 - SYS_read, 402
 - SYS_sbrk, 402
 - SYS_sleep, 402
 - SYS_unlink, 402
 - SYS_uptime, 402
 - SYS_wait, 403
 - SYS_write, 403
- syscalls
 - syscall.c, 395
- sysfile.c, 404
 - argfd, 404
 - create, 405
 - fdalloc, 405
 - isdirempty, 406
 - sys_chdir, 406
 - sys_close, 406
 - sys_dup, 407
 - sys_exec, 407
 - sys_fstat, 407
 - sys_link, 408

- sys_mkdir, [408](#)
- sys_mknod, [409](#)
- sys_open, [409](#)
- sys_pipe, [410](#)
- sys_read, [410](#)
- sys_unlink, [411](#)
- sys_write, [411](#)
- sysfile.d, [417](#)
- sysproc.c, [417](#)
 - lock, [421](#)
 - proc, [422](#)
 - ptable, [422](#)
 - sys_chpr, [418](#)
 - sys_cps, [418](#)
 - sys_exit, [418](#)
 - sys_fork, [419](#)
 - sys_getpid, [419](#)
 - sys_getprocs, [419](#)
 - sys_halt, [420](#)
 - sys_kill, [420](#)
 - sys_sbrk, [420](#)
 - sys_sleep, [420](#)
 - sys_uptime, [421](#)
 - sys_wait, [421](#)
- sysproc.d, [424](#)
- sz
 - proc, [51](#)
- t
 - taskstate, [74](#)
- T_ALIGN
 - traps.h, [431](#)
- T_BOUND
 - traps.h, [431](#)
- T_BRKPT
 - traps.h, [431](#)
- T_DBLFLT
 - traps.h, [431](#)
- T_DEBUG
 - traps.h, [431](#)
- T_DEFAULT
 - traps.h, [432](#)
- T_DEV
 - stat.h, [376](#)
- T_DEVICE
 - traps.h, [432](#)
- T_DIR
 - stat.h, [376](#)
- T_DIVIDE
 - traps.h, [432](#)
- T_FILE
 - stat.h, [376](#)
- T_FPERR
 - traps.h, [432](#)
- T_GPFLT
 - traps.h, [432](#)
- T_ILLOP
 - traps.h, [432](#)
- T_IRQ0
 - traps.h, [433](#)
- T_MCHK
 - traps.h, [433](#)
- T_NMI
 - traps.h, [433](#)
- T_OFLOW
 - traps.h, [433](#)
- T_PGFLT
 - traps.h, [433](#)
- T_SEGNP
 - traps.h, [433](#)
- T_SIMDERR
 - traps.h, [434](#)
- T_STACK
 - traps.h, [434](#)
- T_SYSCALL
 - traps.h, [434](#)
- T_TSS
 - traps.h, [434](#)
- taskstate, [67](#)
 - cr3, [68](#)
 - cs, [68](#)
 - ds, [68](#)
 - eax, [68](#)
 - ebp, [69](#)
 - ebx, [69](#)
 - ecx, [69](#)
 - edi, [69](#)
 - edx, [69](#)
 - eflags, [69](#)
 - eip, [70](#)
 - es, [70](#)
 - esi, [70](#)
 - esp, [70](#)
 - esp0, [70](#)
 - esp1, [70](#)
 - esp2, [71](#)
 - fs, [71](#)
 - gs, [71](#)
 - iomb, [71](#)
 - ldt, [71](#)
 - link, [71](#)
 - padding1, [72](#)
 - padding10, [72](#)
 - padding2, [72](#)
 - padding3, [72](#)
 - padding4, [72](#)
 - padding5, [72](#)
 - padding6, [73](#)
 - padding7, [73](#)
 - padding8, [73](#)
 - padding9, [73](#)
 - ss, [73](#)
 - ss0, [73](#)
 - ss1, [74](#)
 - ss2, [74](#)
 - t, [74](#)
- TCCR

- lapic.c, [251](#)
- TDCR
 - lapic.c, [251](#)
- tf
 - proc, [52](#)
- ticks
 - defs.h, [161](#)
 - trap.c, [427](#)
- tickslock
 - defs.h, [161](#)
 - trap.c, [427](#)
- TICR
 - lapic.c, [251](#)
- TIMER
 - lapic.c, [251](#)
- timerinit
 - defs.h, [156](#)
- togglecode
 - kbd.h, [241](#)
- TPR
 - lapic.c, [251](#)
- trap
 - trap.c, [425](#)
- trap.c, [424](#)
 - idt, [426](#)
 - idtinit, [425](#)
 - ticks, [427](#)
 - tickslock, [427](#)
 - trap, [425](#)
 - tvinit, [426](#)
 - vectors, [427](#)
- trap.d, [429](#)
- trapframe, [74](#)
 - cs, [75](#)
 - ds, [75](#)
 - eax, [76](#)
 - ebp, [76](#)
 - ebx, [76](#)
 - ecx, [76](#)
 - edi, [76](#)
 - edx, [76](#)
 - eflags, [77](#)
 - eip, [77](#)
 - err, [77](#)
 - es, [77](#)
 - esi, [77](#)
 - esp, [78](#)
 - fs, [78](#)
 - gs, [78](#)
 - oesp, [78](#)
 - padding1, [78](#)
 - padding2, [78](#)
 - padding3, [79](#)
 - padding4, [79](#)
 - padding5, [79](#)
 - padding6, [79](#)
 - ss, [79](#)
 - trapno, [79](#)
- trapno
 - trapframe, [79](#)
- trapret
 - proc.c, [333](#)
- traps.h, [429](#)
 - IRQ_COM1, [430](#)
 - IRQ_ERROR, [430](#)
 - IRQ_IDE, [430](#)
 - IRQ_KBD, [430](#)
 - IRQ_SPURIOUS, [430](#)
 - IRQ_TIMER, [431](#)
 - T_ALIGN, [431](#)
 - T_BOUND, [431](#)
 - T_BRKPT, [431](#)
 - T_DBLFLT, [431](#)
 - T_DEBUG, [431](#)
 - T_DEFAULT, [432](#)
 - T_DEVICE, [432](#)
 - T_DIVIDE, [432](#)
 - T_FPERR, [432](#)
 - T_GPFLT, [432](#)
 - T_ILLOP, [432](#)
 - T_IRQ0, [433](#)
 - T_MCHK, [433](#)
 - T_NMI, [433](#)
 - T_OFLOW, [433](#)
 - T_PGFLT, [433](#)
 - T_SEGNP, [433](#)
 - T_SIMDERR, [434](#)
 - T_STACK, [434](#)
 - T_SYSCALL, [434](#)
 - T_TSS, [434](#)
- ts
 - cpu, [16](#)
- tvinit
 - defs.h, [156](#)
 - trap.c, [426](#)
- type
 - backcmd, [9](#)
 - cmd, [12](#)
 - dinode, [18](#)
 - elfhdr, [22](#)
 - execcmd, [23](#)
 - file, [26](#)
 - gatedesc, [28](#)
 - inode, [32](#)
 - listcmd, [35](#)
 - mp, [40](#)
 - mpioapic, [44](#)
 - mpproc, [45](#)
 - pipecmd, [48](#)
 - proghdr, [54](#)
 - redircmd, [55](#)
 - segdesc, [61](#)
 - stat, [65](#)
- types.h, [435](#)
 - pde_t, [435](#)
 - uchar, [435](#)

- uint, 436
- ushort, 436
- uart
 - uart.c, 439
- uart.c, 436
 - COM1, 437
 - uart, 439
 - uartgetc, 437
 - uartinit, 437
 - uartintr, 438
 - uartputc, 438
- uart.d, 440
- uartgetc
 - uart.c, 437
- uartinit
 - defs.h, 156
 - uart.c, 437
- uartintr
 - defs.h, 157
 - uart.c, 438
- uartputc
 - defs.h, 157
 - uart.c, 438
- uchar
 - types.h, 435
- uint
 - types.h, 436
- uio
 - usertests.c, 488
- ulib.c, 440
 - atoi, 440
 - gets, 441
 - memmove, 441
 - memset, 441
 - stat, 442
 - strchr, 442
 - strcmp, 442
 - strcpy, 443
 - strlen, 443
- ulib.d, 445
- umalloc.c, 445
 - Align, 446
 - base, 447
 - free, 446
 - freep, 448
 - Header, 446
 - malloc, 446
 - morecore, 447
- umalloc.d, 449
- uninit
 - usertests.c, 493
- unlink
 - user.h, 461
- unlinkread
 - usertests.c, 489
- UNUSED
 - proc.h, 344
- uproc, 80
 - name, 80
 - pid, 80
 - ppid, 80
- uproc.h, 449
- uptime
 - user.h, 461
- use_lock
 - kalloc.c, 231
- user.h, 450
 - atoi, 450
 - chdir, 451
 - chpr, 451
 - close, 451
 - cps, 451
 - dup, 452
 - exec, 452
 - exit, 453
 - fork, 454
 - free, 455
 - fstat, 455
 - getpid, 455
 - getprocs, 455
 - gets, 456
 - halt, 456
 - kill, 456
 - link, 456
 - malloc, 457
 - memmove, 457
 - memset, 457
 - mkdir, 458
 - mknod, 458
 - open, 458
 - pipe, 458
 - printf, 459
 - read, 459
 - sbrk, 459
 - sleep, 459
 - stat, 459
 - strchr, 460
 - strcmp, 460
 - strcpy, 460
 - strlen, 461
 - unlink, 461
 - uptime, 461
 - wait, 461
 - write, 462
- userinit
 - defs.h, 157
 - proc.c, 333
- usertests.c, 463
 - argptest, 465
 - BIG, 464
 - bigargtest, 465
 - bigdir, 466
 - bigfile, 466
 - bigwrite, 467
 - bsstest, 468
 - buf, 492

- concreate, 468
- createdelete, 469
- createtest, 470
- dirfile, 471
- dirtest, 472
- echoargv, 492
- exectest, 472
- exitiputtest, 472
- exitwait, 473
- forktest, 473
- fourfiles, 474
- fourteen, 475
- fsfull, 476
- iputtest, 476
- iref, 477
- linktest, 477
- linkunlink, 478
- main, 479
- mem, 480
- name, 492
- openiputtest, 480
- opentest, 481
- pipe1, 481
- preempt, 482
- rand, 483
- randstate, 493
- rmdot, 483
- RTC_ADDR, 464
- RTC_DATA, 465
- sbrktest, 483
- sharedfd, 485
- stdout, 493
- subdir, 486
- uio, 488
- uninit, 493
- unlinkread, 489
- validateint, 489
- validatetest, 490
- writetest, 490
- writetest1, 491
- usertests.d, 514
- ushort
 - types.h, 436
- uva2ka
 - defs.h, 158
 - vm.c, 521
- V2P
 - memlayout.h, 279
- V2P_WO
 - memlayout.h, 279
- vaddr
 - proghdr, 54
- valid
 - inode, 32
- validateint
 - usertests.c, 489
- validatetest
 - usertests.c, 490
- vectors
 - trap.c, 427
- VER
 - lapic.c, 251
- version
 - elfhdr, 22
 - mpconf, 42
 - mpioapic, 44
 - mpproc, 46
- virt
 - kmap, 34
- vm.c, 515
 - allocvm, 515
 - clearpteu, 516
 - copyout, 516
 - copyvm, 517
 - data, 522
 - deallocvm, 517
 - freevm, 518
 - initvm, 518
 - kmap, 522
 - kpgdir, 523
 - kvmalloc, 519
 - loadvm, 519
 - mappages, 519
 - seginit, 520
 - setupkvm, 520
 - switchkvm, 521
 - switchvm, 521
 - uva2ka, 521
 - walkpgdir, 522
- vm.d, 528
- w
 - console.c, 103
- wait
 - defs.h, 158
 - proc.c, 333
 - user.h, 461
- waitdisk
 - bootmain.c, 90
- wakeup
 - defs.h, 159
 - proc.c, 334
- wakeup1
 - proc.c, 334
- walkpgdir
 - vm.c, 522
- wc
 - wc.c, 528
- wc.c, 528
 - buf, 529
 - main, 528
 - wc, 528
- wc.d, 530
- whitespace
 - sh.c, 360
- winode
 - mkfs.c, 286

writable
 file, 26
write
 devsw, 17
 user.h, 462
write_head
 log.c, 263
write_log
 log.c, 264
writei
 defs.h, 159
 fs.c, 197
writeopen
 pipe, 47
writetest
 usertests.c, 490
writetest1
 usertests.c, 491
wsect
 mkfs.c, 286

x
 header, 29
X1
 lapic.c, 252
x86.h, 530
 cli, 531
 inb, 531
 insl, 531
 lcr3, 531
 lgdt, 532
 lidt, 532
 loadgs, 532
 ltr, 533
 outb, 533
 outsl, 533
 outw, 533
 rcr2, 534
 readeflags, 534
 sti, 534
 stosb, 534
 stosl, 535
 xchg, 535
xchecksum
 mpconf, 42
xchg
 x86.h, 535
xint
 mkfs.c, 287
xlength
 mpconf, 43
xshort
 mkfs.c, 287
xv6.dox, 538

YEAR
 lapic.c, 252
year
 rtcd, 57

yield
 defs.h, 160
 proc.c, 335

zeroes
 mkfs.c, 289
ZOMBIE
 proc.h, 344
zombie.c, 538
 main, 538
zombie.d, 539